

$GHC_{(STG, Cmm, asm)}$ illustrated

for hardware person

exploring some mental models and implementations

Takenobu T.

"Any sufficiently advanced technology is
indistinguishable from **magic**."

Arthur C. Clarke

NOTE

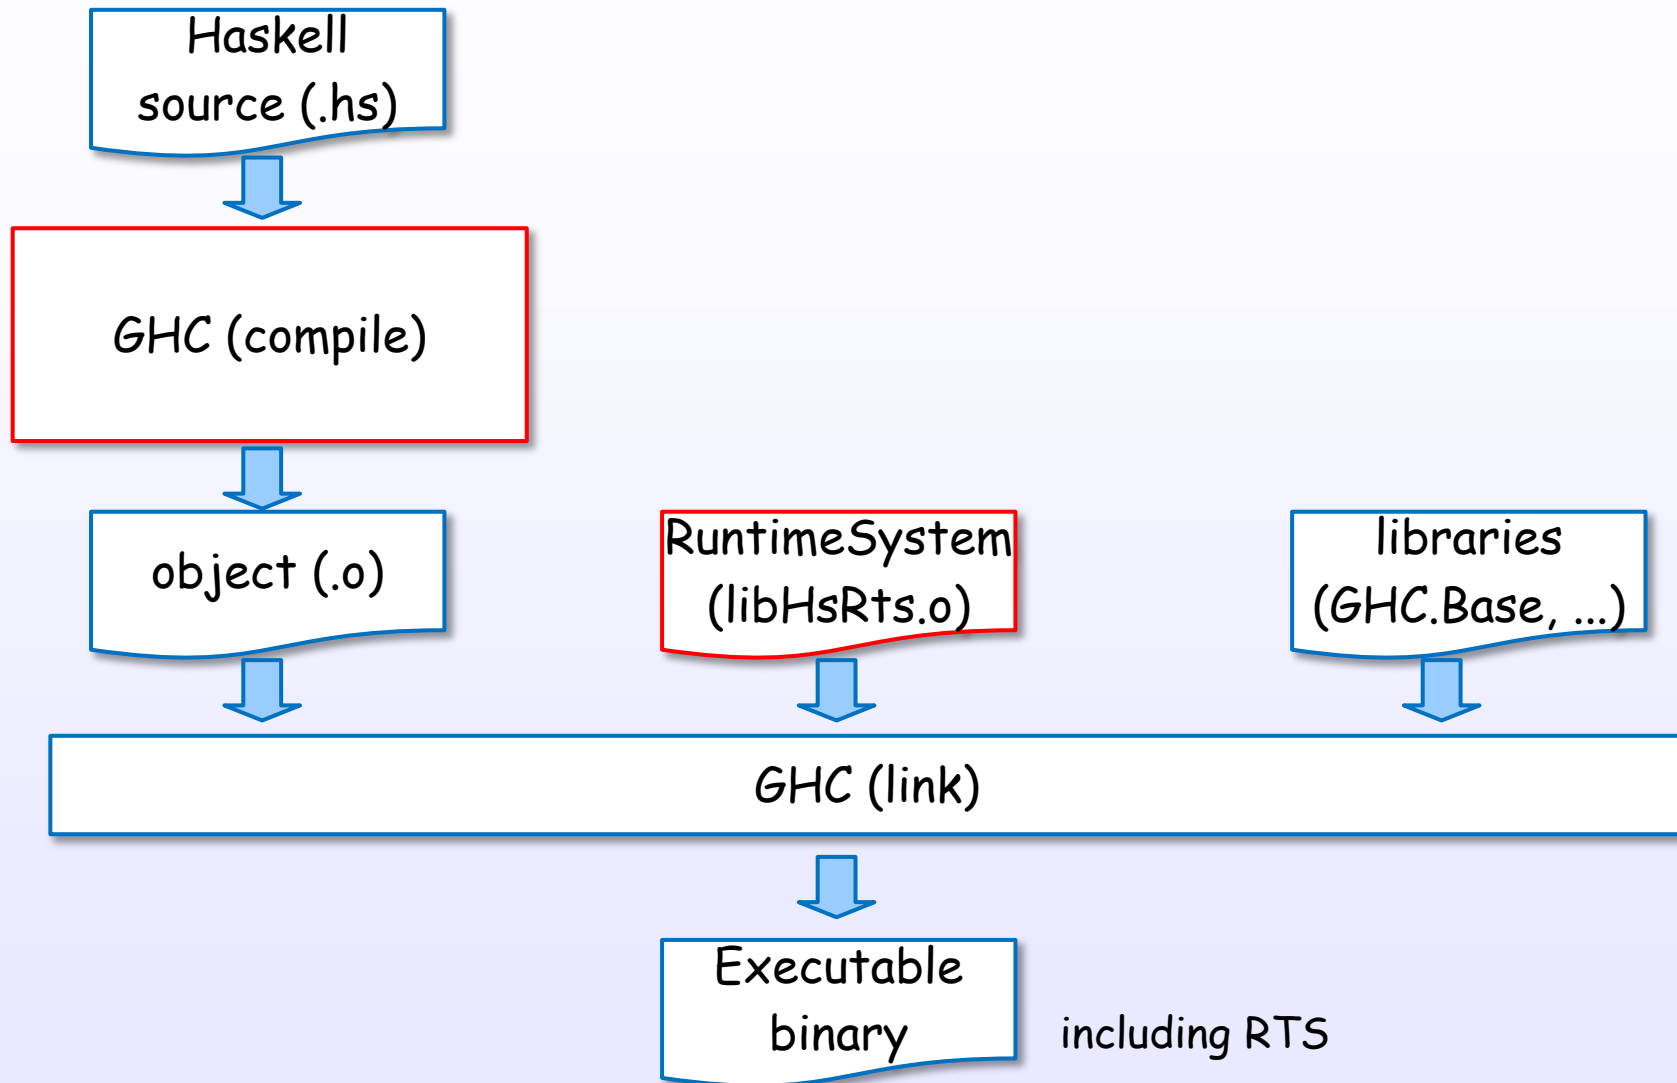
- This is not a official document by ghc development team.
- Please don't forget "semantics". It's very important.
- This is written for ghc 7.8.

Contents

- Executable binary
- Compile steps
- Runtime System
- Development languages
- Machine layer/models
- STG-machine
- Heap objects in STG-machine
- STG-machine evaluation
- Pointer tagging
- Thunk and update
- Allocate and free heap objects
- STG - C land interface
- Thread
- Thread context switch
- Creating main and sub threads
- Thread migration
- Heap and Threads
- Threads and GC
- Bound thread
- Spark
- Mvar
- Software transactional memory
- IO and FFI
- IO manager
- Bootstrap
- References

Executable binary

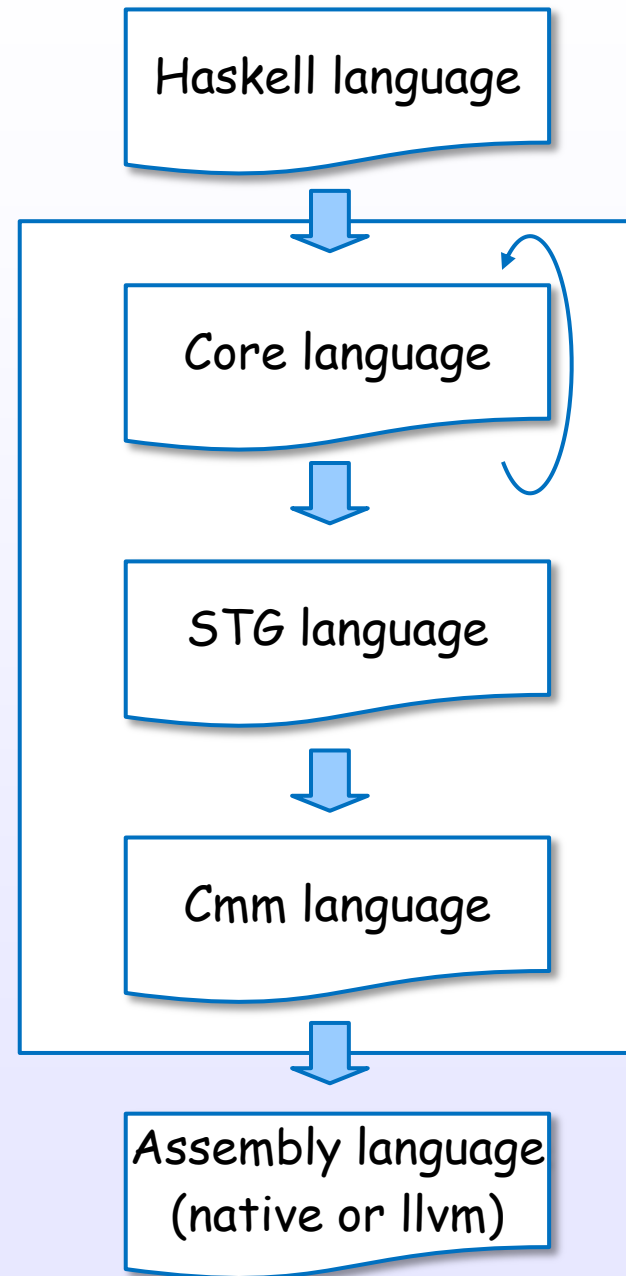
GHC = Compiler + Runtime System (RTS)



Compile steps

GHC transitions between five representations

GHC
compile
steps



*each intermediate code can
be dumped by :*

`% ghc -ddump-parsed`
`% ghc -ddump-rn`

`% ghc -ddump-ds`
`% ghc -ddump-simpl`
`% ghc -ddump-prep`

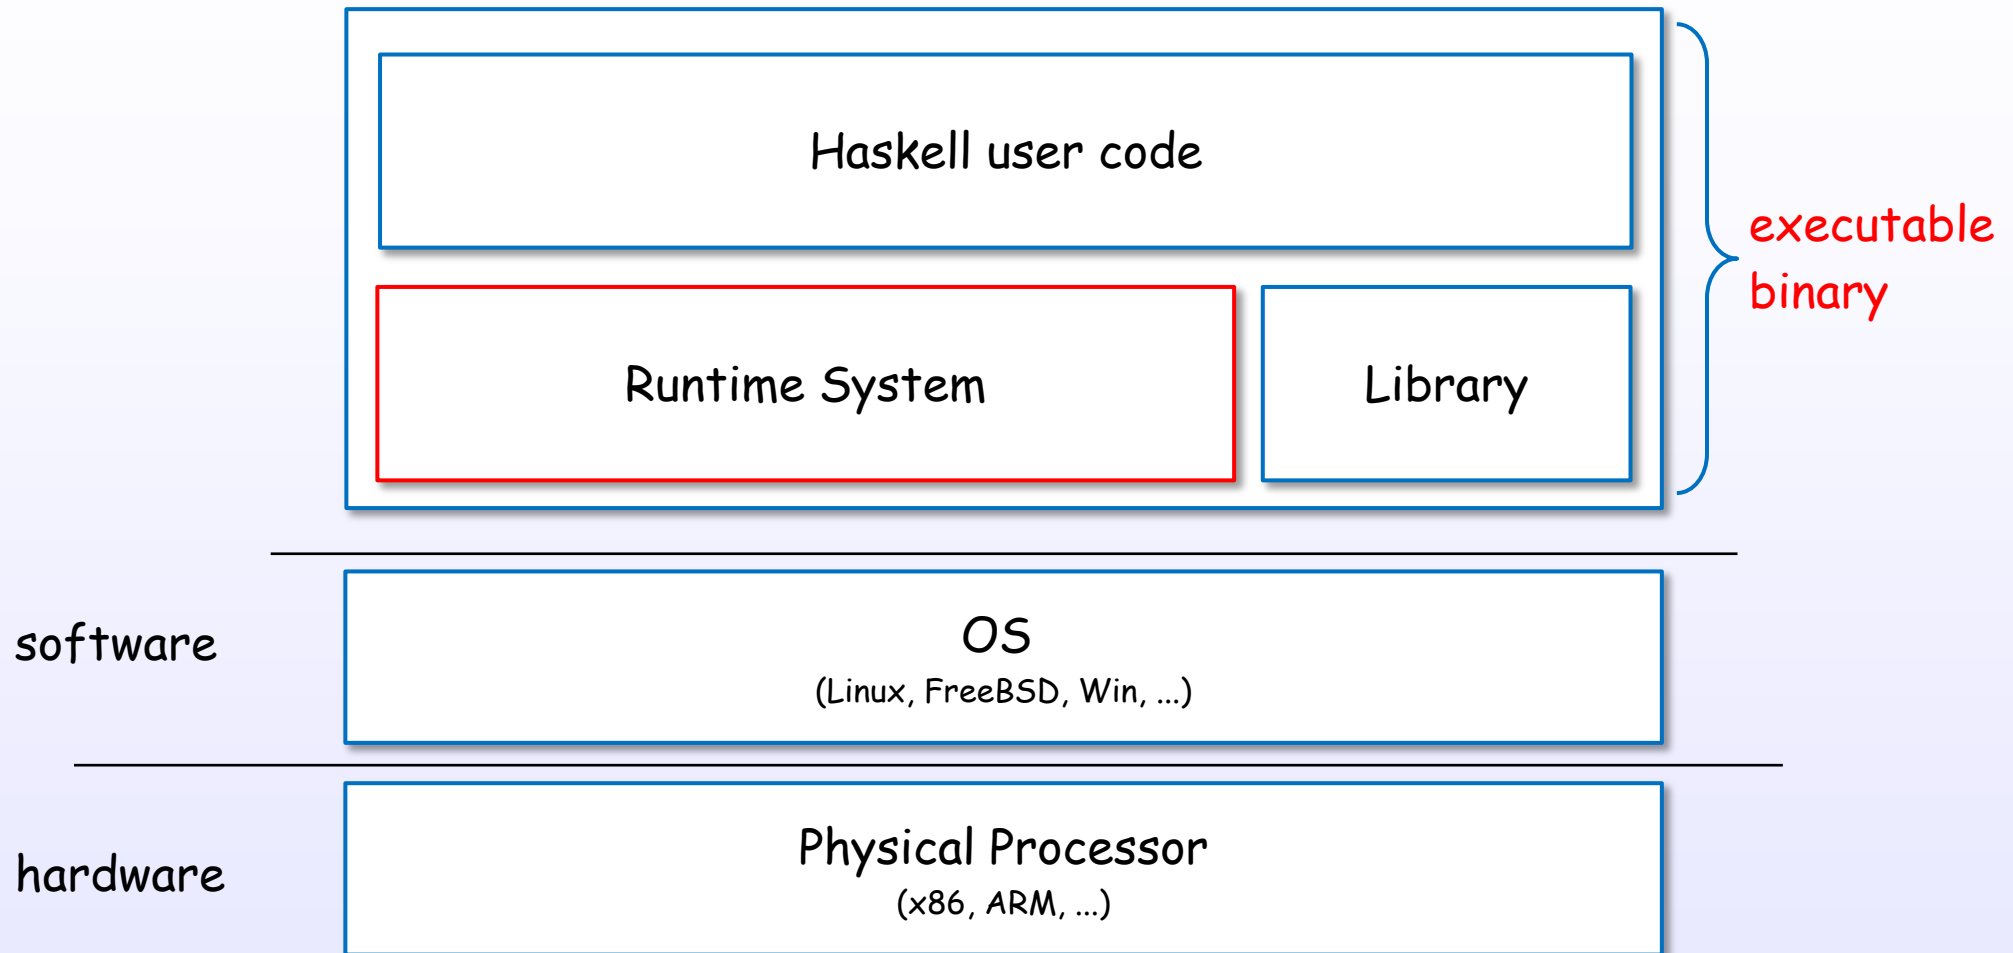
`% ghc -ddump-stg`

`% ghc -ddump-cmm`
`% ghc -ddump-opt-cmm`

`% ghc -ddump-llvm`
`% ghc -ddump-asm`

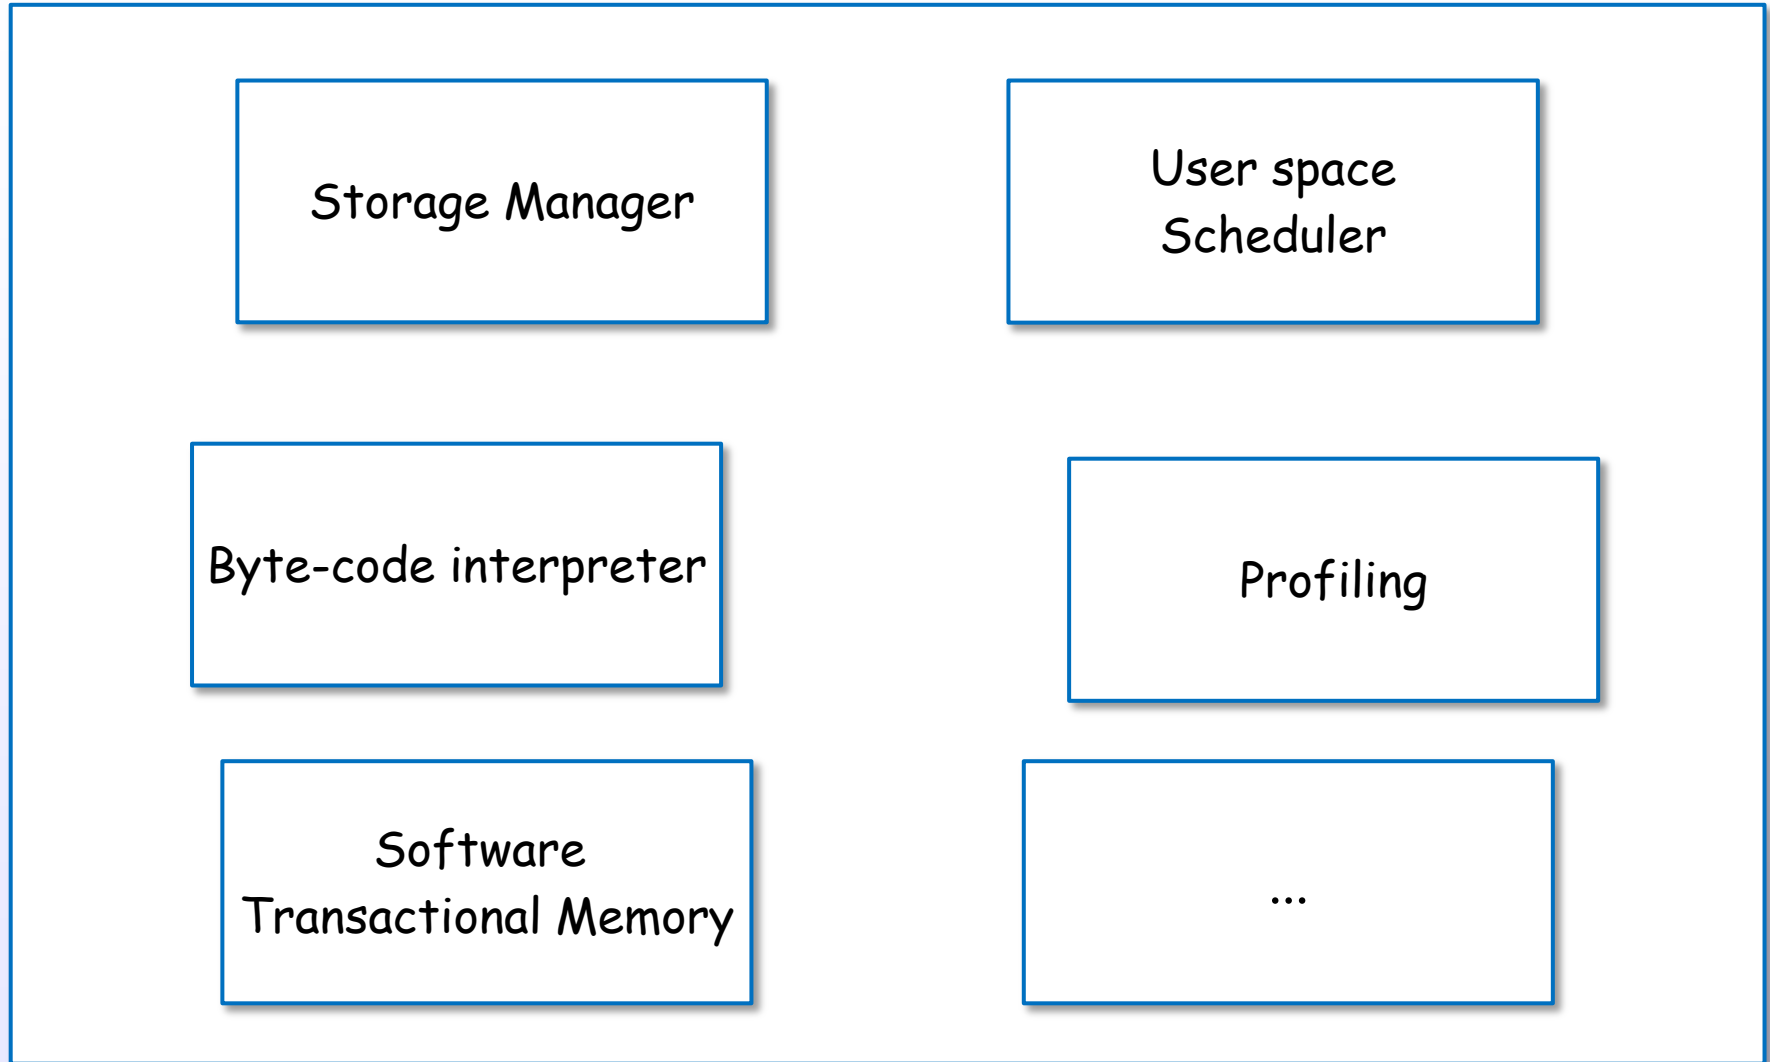
Runtime System

Generated binary includes the RTS



Runtime System includes ...

Runtime System



Development languages

GHC are developed by some languages

compiler

(\$(TOP)/**compiler**/*)

Haskell

+

Alex (lex)

Happy (yacc)

Cmm (C--)

Assembly

runtime system

(\$(TOP)/**rts**/*)

C

+

Cmm

Assembly

library

(\$(TOP)/**libraries**/*)

Haskell

+

C

Machine layer/models

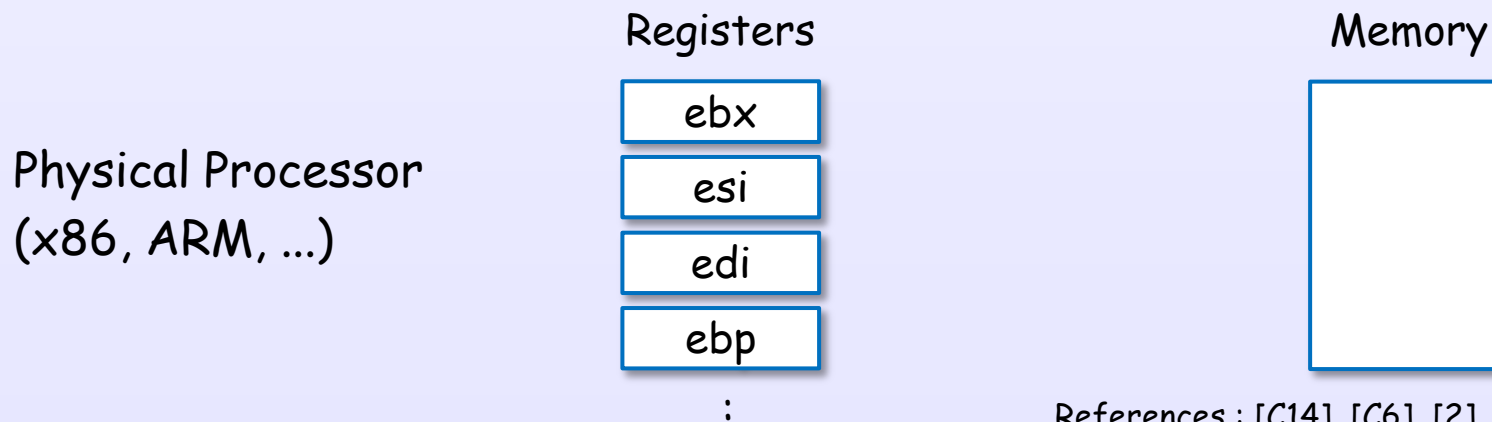
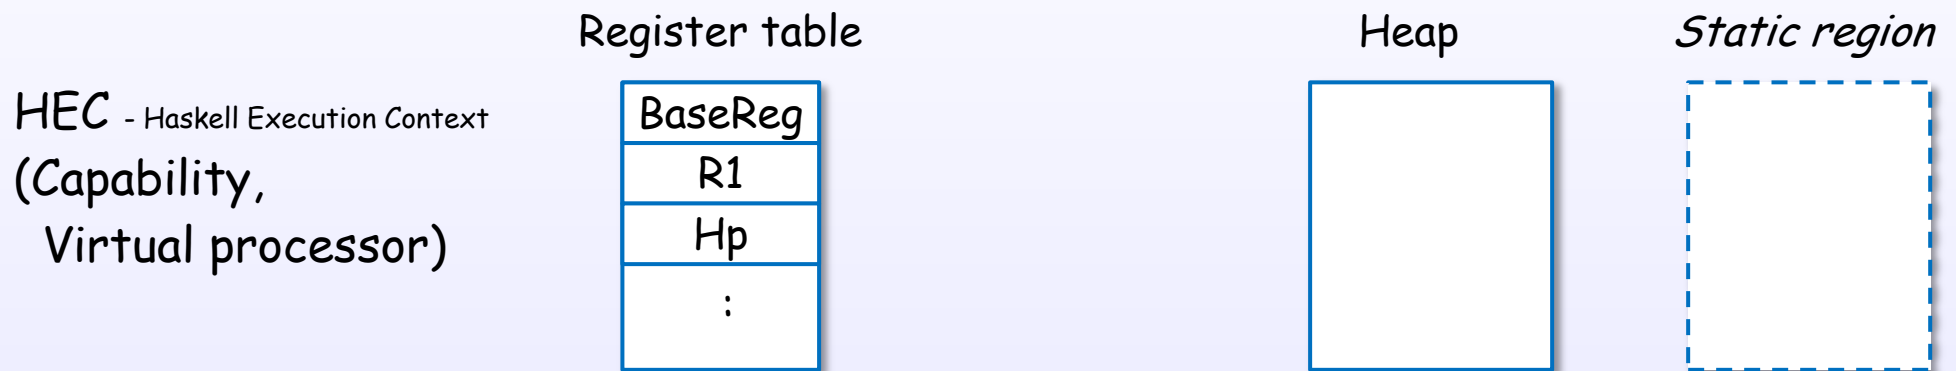
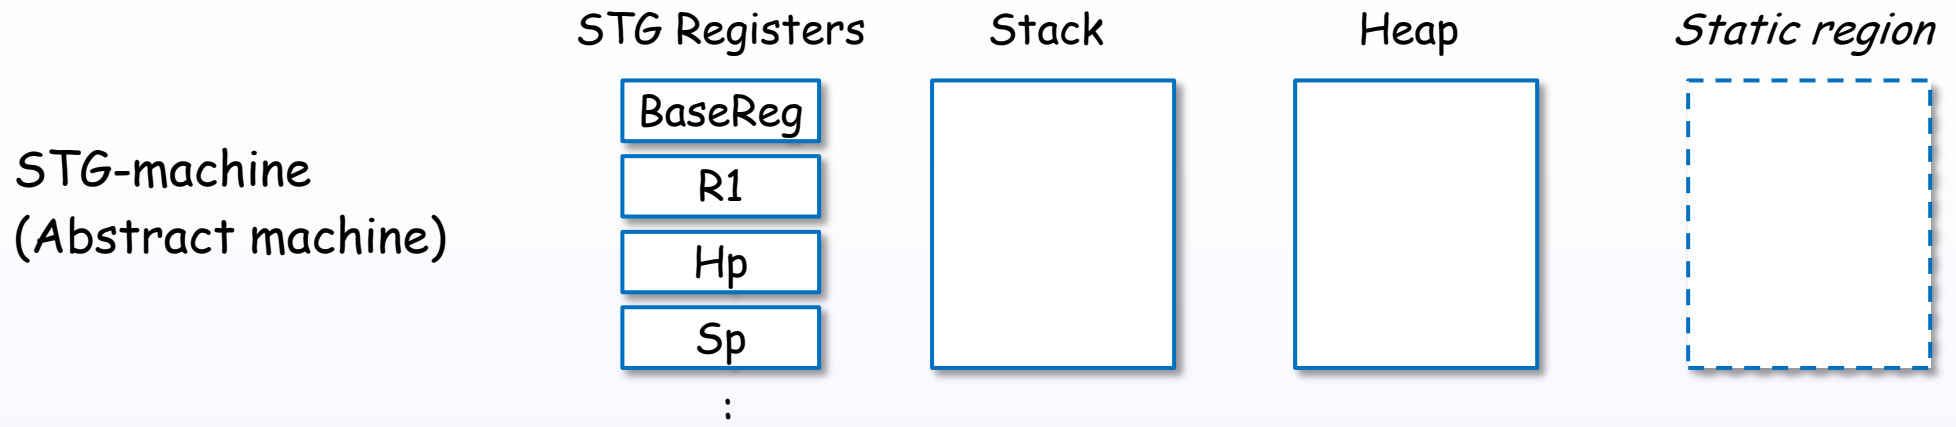
Machine layer

STG-machine
(Abstract machine)

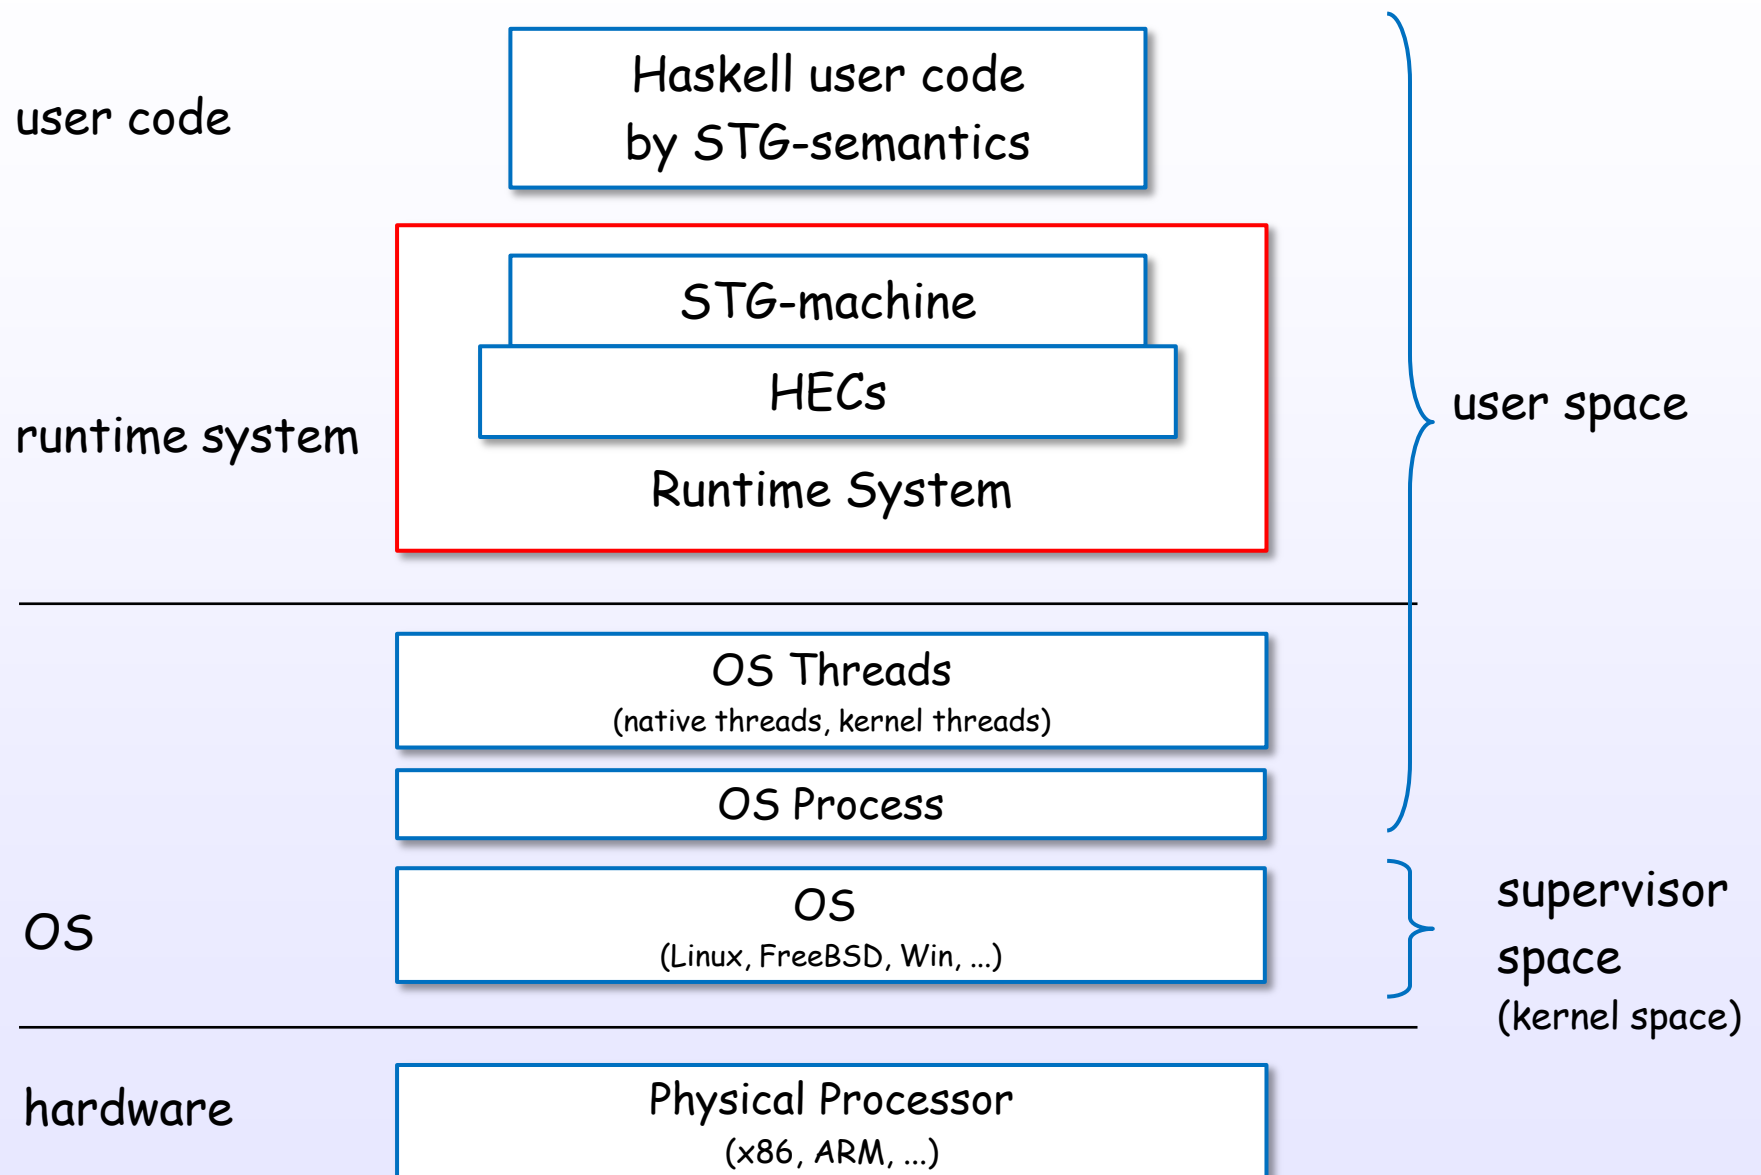
HEC - Haskell Execution Context
(Capability, Virtual processor)

Physical Processor
(x86, ARM, ...)

Machine layer



Runtime system and HEC

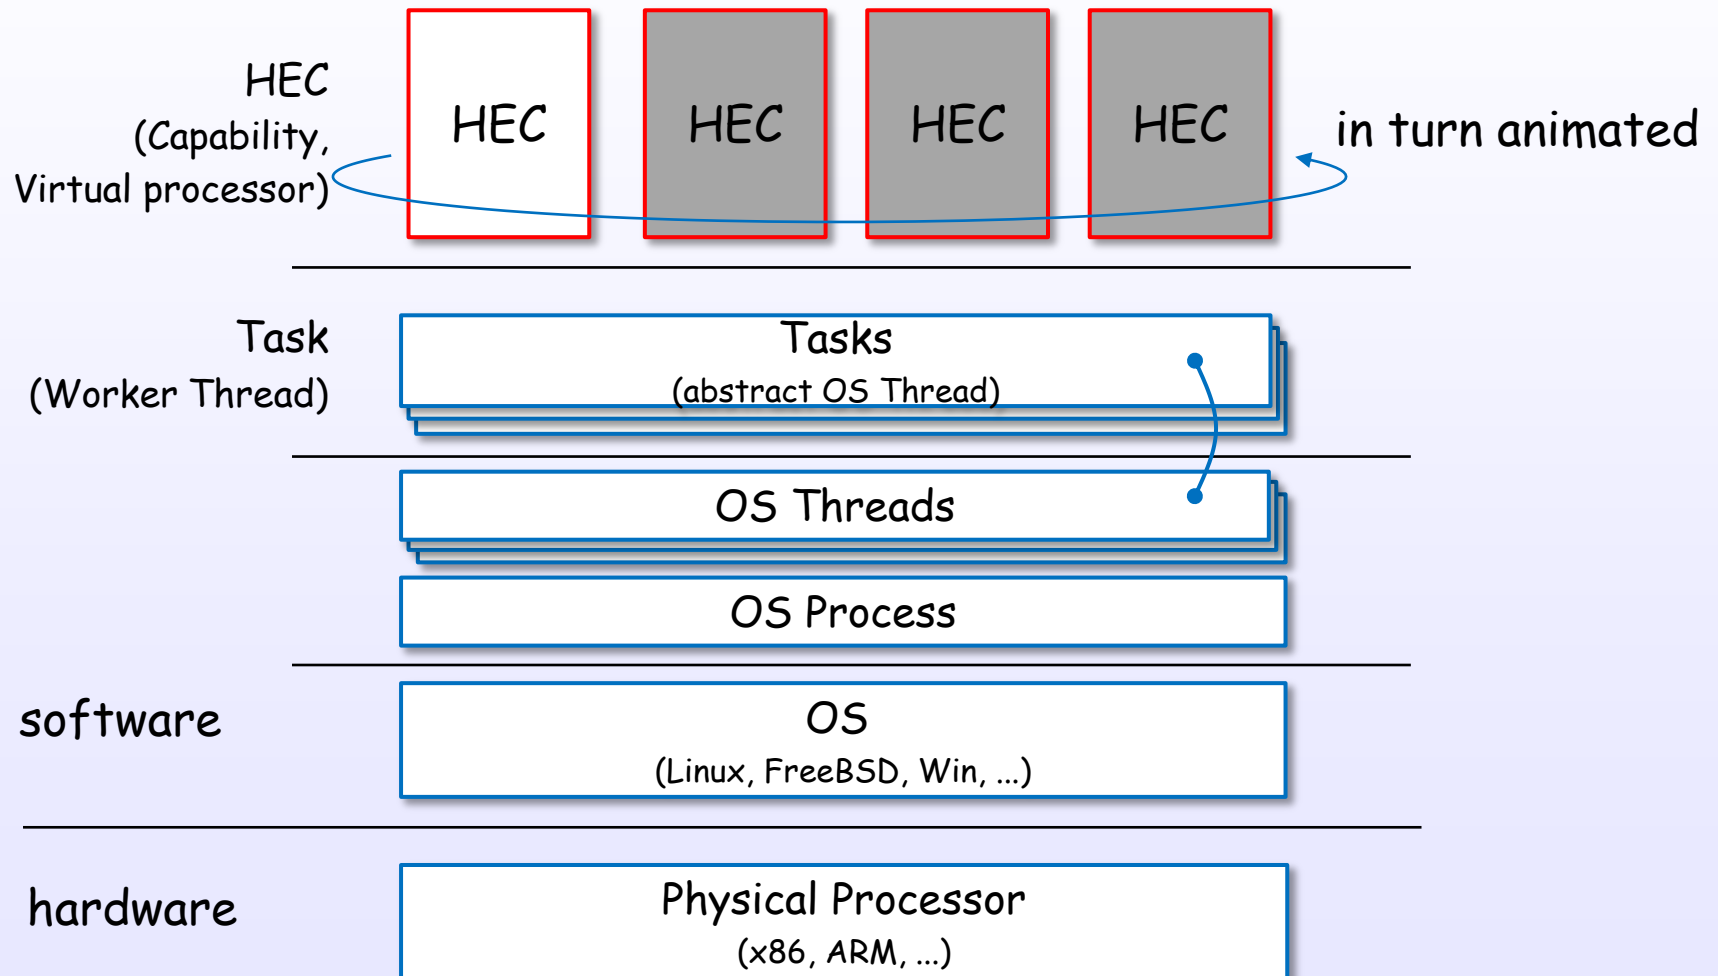


many HECs

Multi HECs can be generated by compile and runtime options :

```
$ ghc -rtsopts -threaded
```

```
$ ./xxx +RTS -N4
```



HEC (Capability) data structure

[rts/Capability.h]

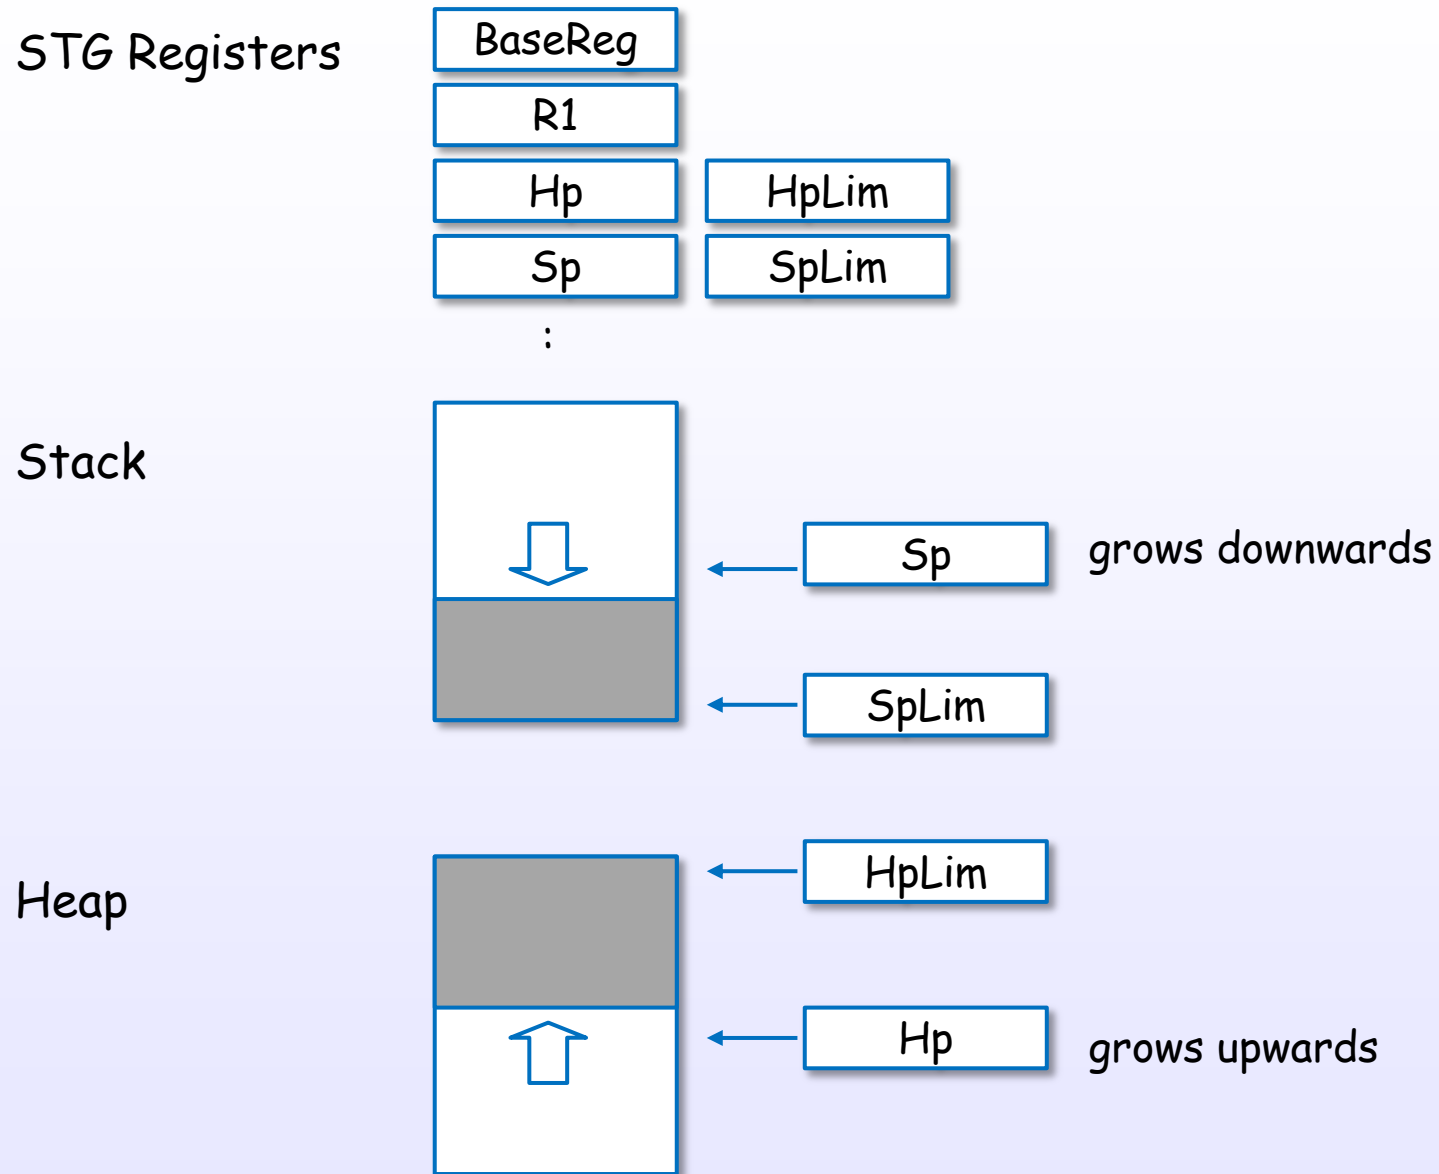
```
struct Capability_ {  
    StgFunTable f;  
    StgRegTable r;  
    nat no;  
    Task *running_task;  
    rtsBool in_haskell;  
    nat idle;  
    rtsBool disabled;  
    StgTSO *run_queue_hd;  
    StgTSO *run_queue_tl;  
    InCall *suspended_ccalls;  
    bdescr **mut_lists;  
    bdescr **saved_mut_lists;  
    bdescr *pinned_object_block;  
    bdescr *pinned_object_blocks;  
    int context_switch;  
    int interrupt;  
  
#if defined(THREADED_RTS)  
    Task *spare_workers;  
    nat n_spare_workers;  
    Mutex lock;  
    Task *returning_tasks_hd;  
    Task *returning_tasks_tl;  
    Message *inbox;  
    SparkPool *sparks;  
    SparkCounters spark_stats;  
#endif  
  
    W_ total_allocated;  
    StgTVarWatchQueue *free_tvar_watch_queues;  
    StgInvariantCheckQueue *free_invariant_check_queues;  
    StgTRecChunk *free_trec_chunks;  
    StgTRecHeader *free_trec_headers;  
    nat transaction_tokens;  
}
```

Each HEC (Capability) has a register table and a run queue and ...

Each HEC (Capability) is initialized at initCapabilities [rts/rts/Capability.c]

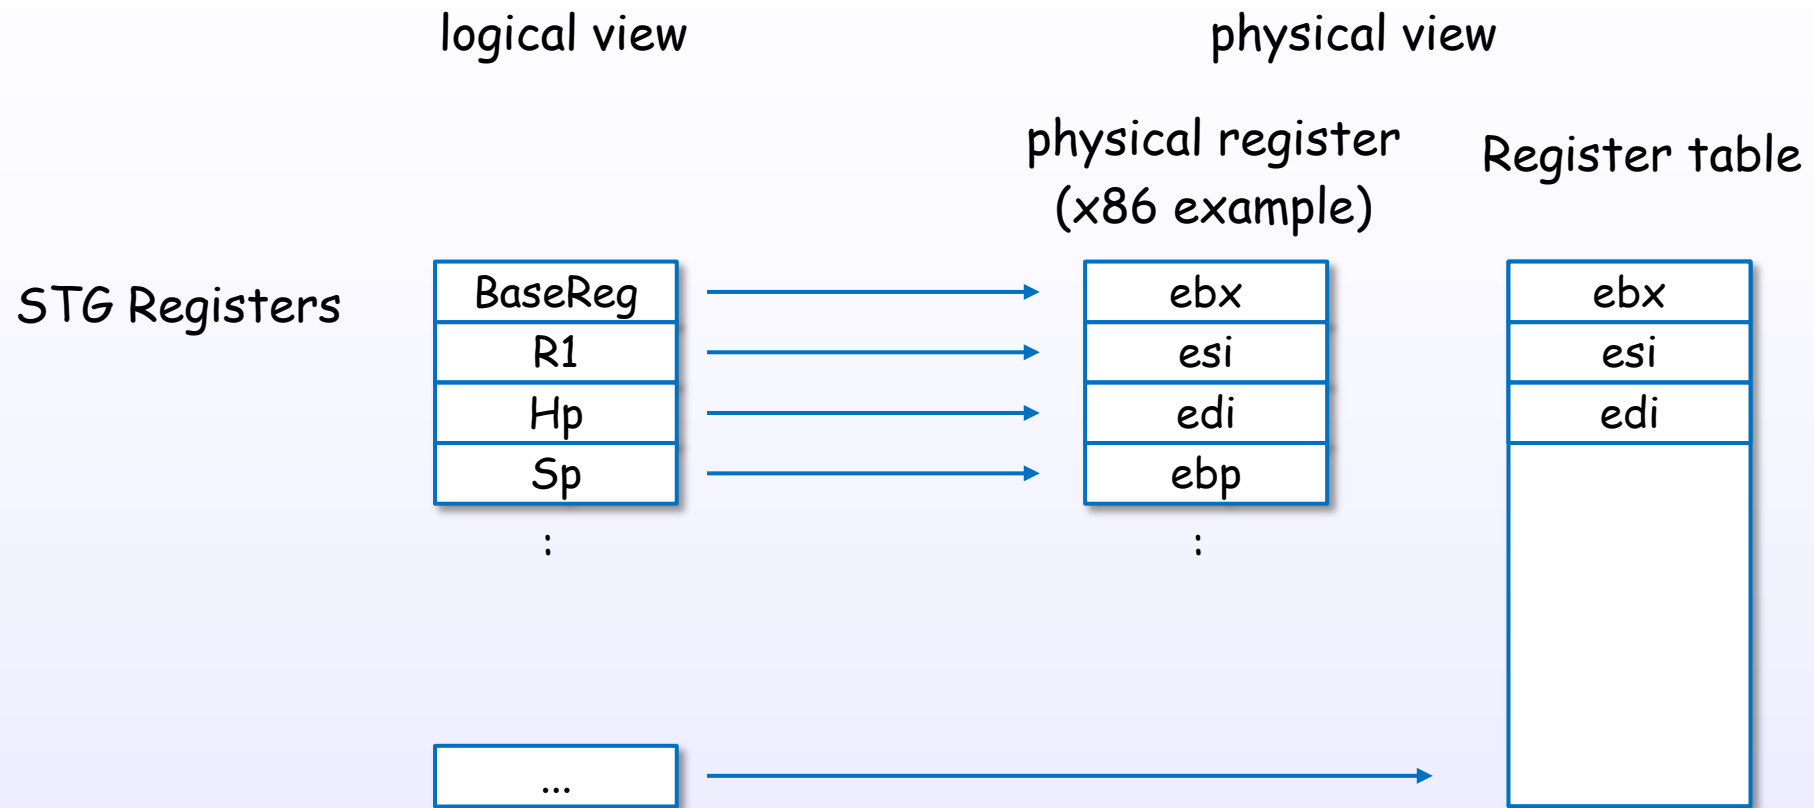
STG-machine

STG-machine

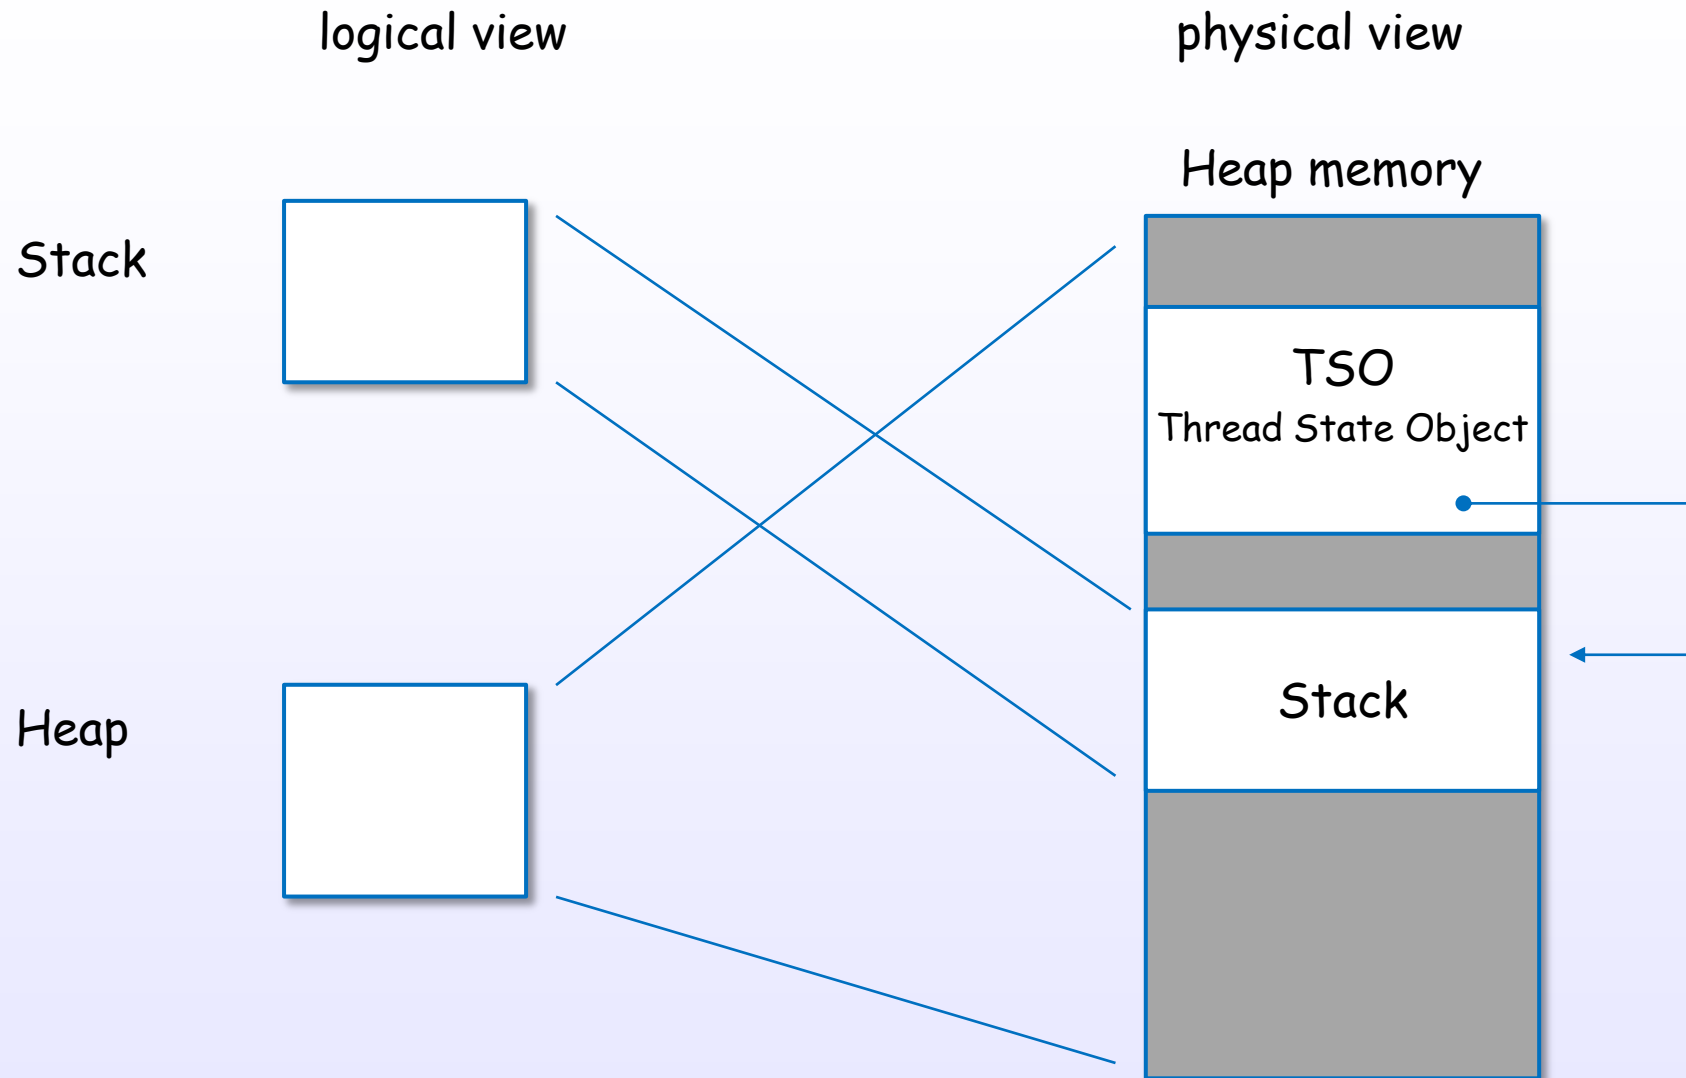


Each real Haskell code is executed in STG semantics.

STG-machine is mapped to physical processor



STG-machine is mapped to physical processor



A stack and a TSO object are in the heap.
The stack is stored separately from the TSO for size extension and GC.

TSO data structure

[includes/rts/storage/TSO.h]

```
typedef struct StgTSO_  
  StgHeader          header;  
  struct StgTSO_*     _link;  
  struct StgTSO_*     global_link;  
  struct StgStack_*   *stackobj;  
  StgWord16           what_next;  
  StgWord16           why_blocked;  
  StgWord32           flags;  
  StgTSOBlockInfo     block_info;  
  StgThreadID         id;  
  StgWord32           saved_errno;  
  StgWord32           dirty;  
  struct InCall_*      bound;  
  struct Capability_*  cap;  
  struct StgTRecHeader_* trec;  
  struct MessageThrowTo_* blocked_exceptions;  
  struct StgBlockingQueue_* bq;  
  StgWord32 tot_stack_size;  
} *StgTSOPtr;
```

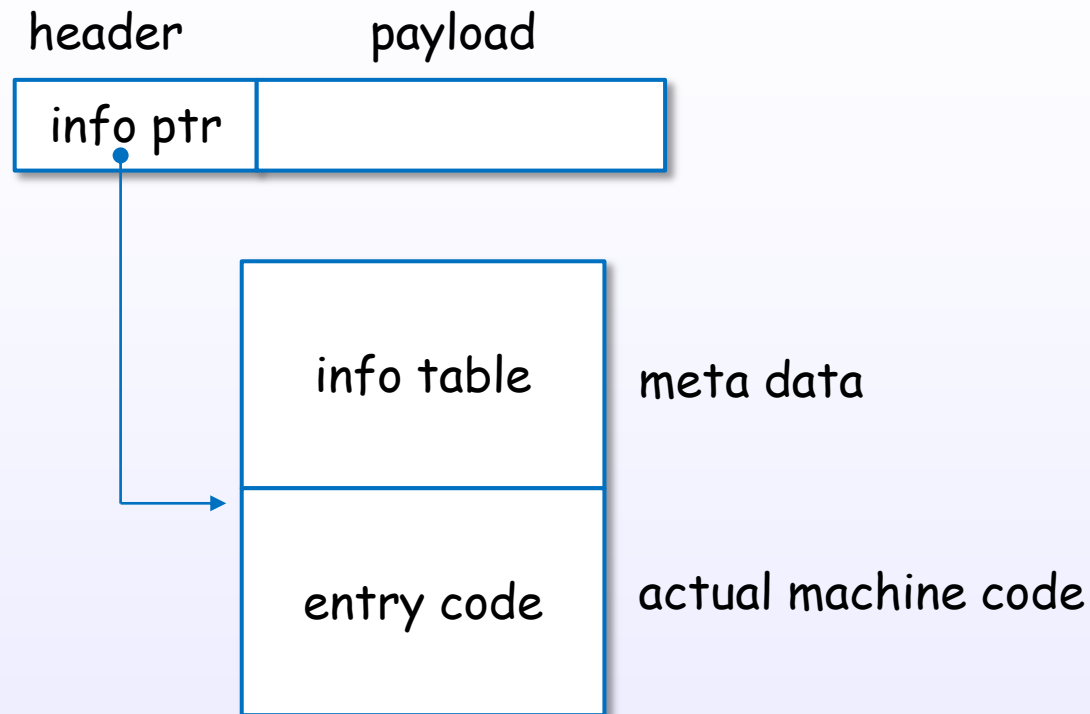
← link to stack object

A TSO object is **only ~17words + stack**. Lightweight!

Heap objects in STG-machine

Heap object (closure)

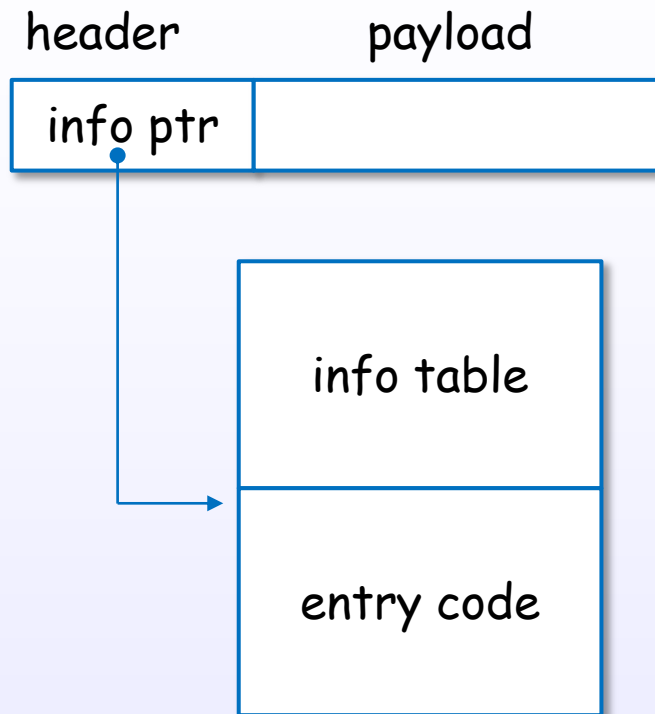
logical view




Closure (header + payload) + Info Table + Entry Code

Heap object (closure)

logical view

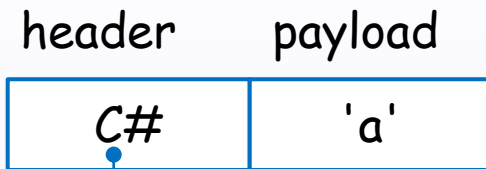


physical view



Closure examples : Char, Int

'a' :: Char



info
ptr

GHC.Types.C#_static_info

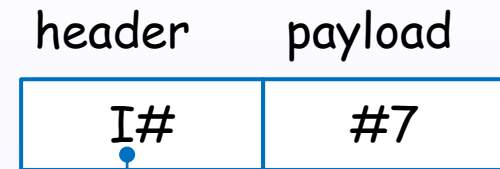
layout : 0_1
type : CONSTR
bitmap :

info table

inc %esi
jmp *0x0(%ebp)

entry code

7 :: Int



info
ptr

GHC.Types.I#_static_info

layout : 0_1
type : CONSTR
bitmap :

info table

inc %esi
jmp *0x0(%ebp)

entry code

Closure example code

[Example.hs]

```
module Example where
value1 :: Int
value1 = 7
```

STG

Cmm

[ghc -O -ddump-stg Example.hs]

```
Example.value1 :: GHC.Types.Int
[GblId, Caf=NoCafRefs, Str=DmdType m, Unf=OtherCon []] =
  NO_CCS GHC.Types.I#! [8];
```

[ghc -O -ddump-opt-cmm Example.hs]

```
section "data" { __stginit_main:Example:
}
```

```
section "data" {
  Example.value1 closure:
    const GHC.Types.I#_static_info;
    const 7;
}
```

asm

[ghc -O -ddump-asm Example.hs]

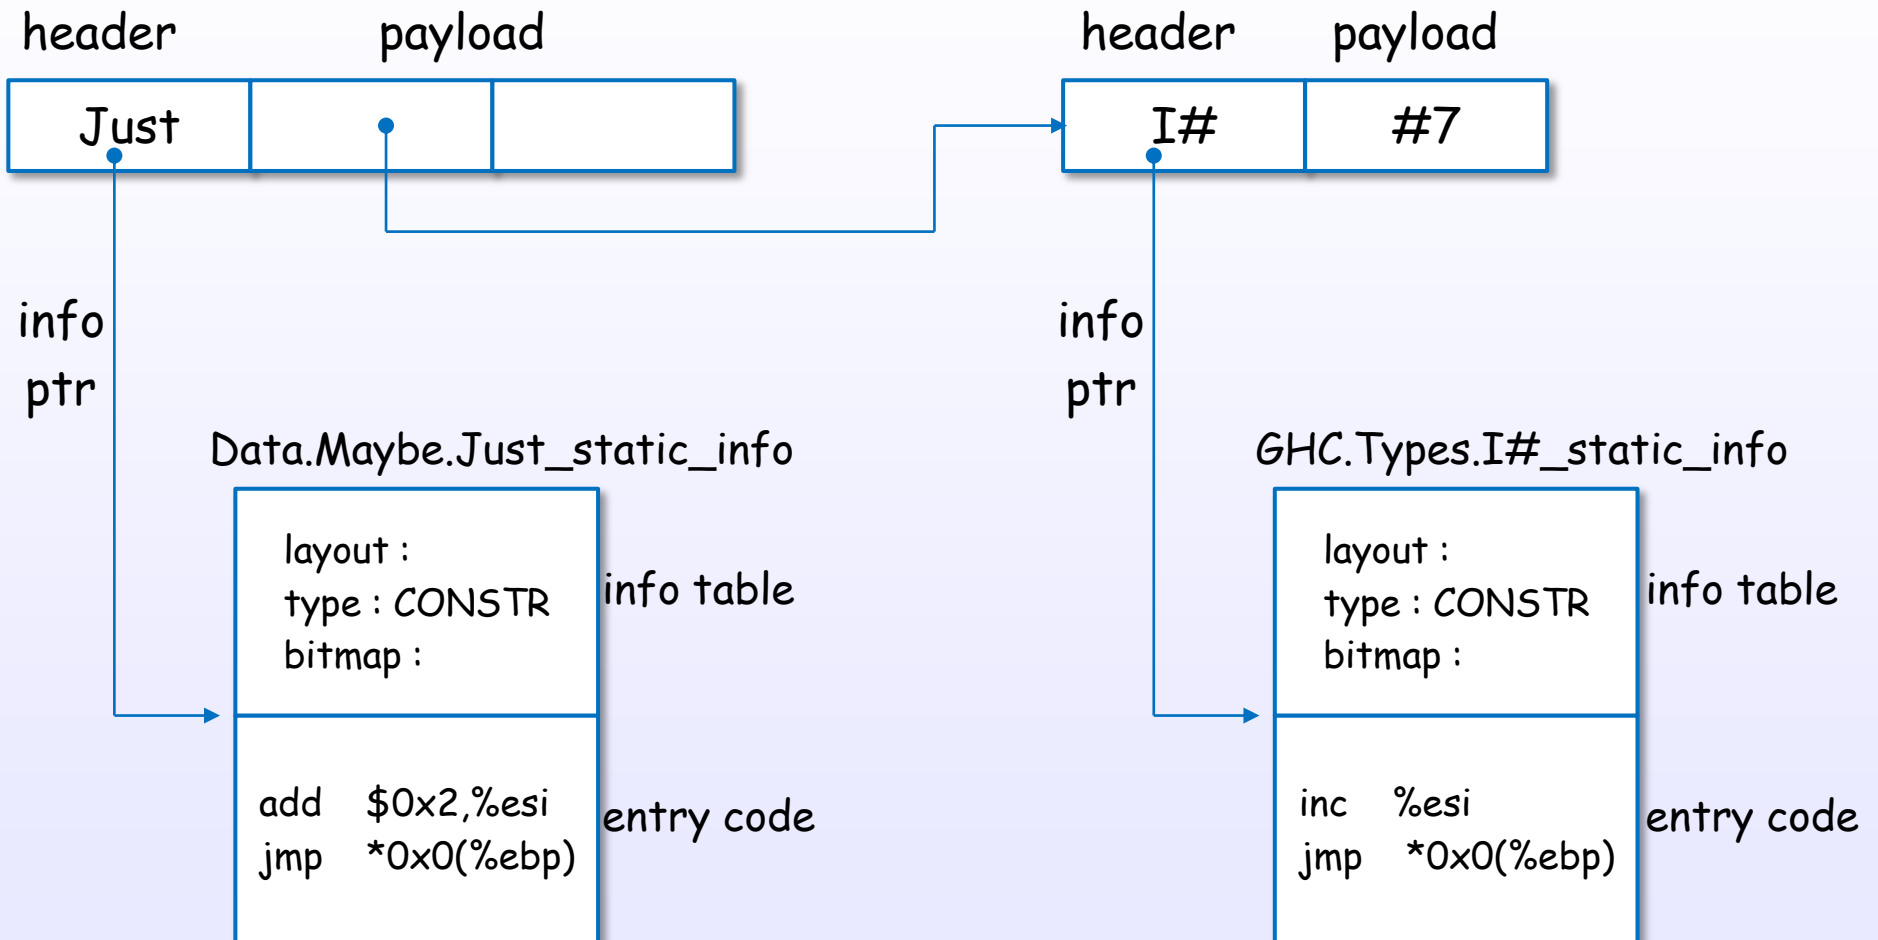
```
.data
    .align 4
    .align 1
    .globl __stginit_main:Example
    __stginit_main:Example:
    .data
        .align 4
        .align 1
        .globl Example.value1_closure
    Example.value1_closure:
        .long GHC.Types.I#_static_info
        long 7
```

```
.section .data
    .align 4
    .align 1
    SMd_srt:
```

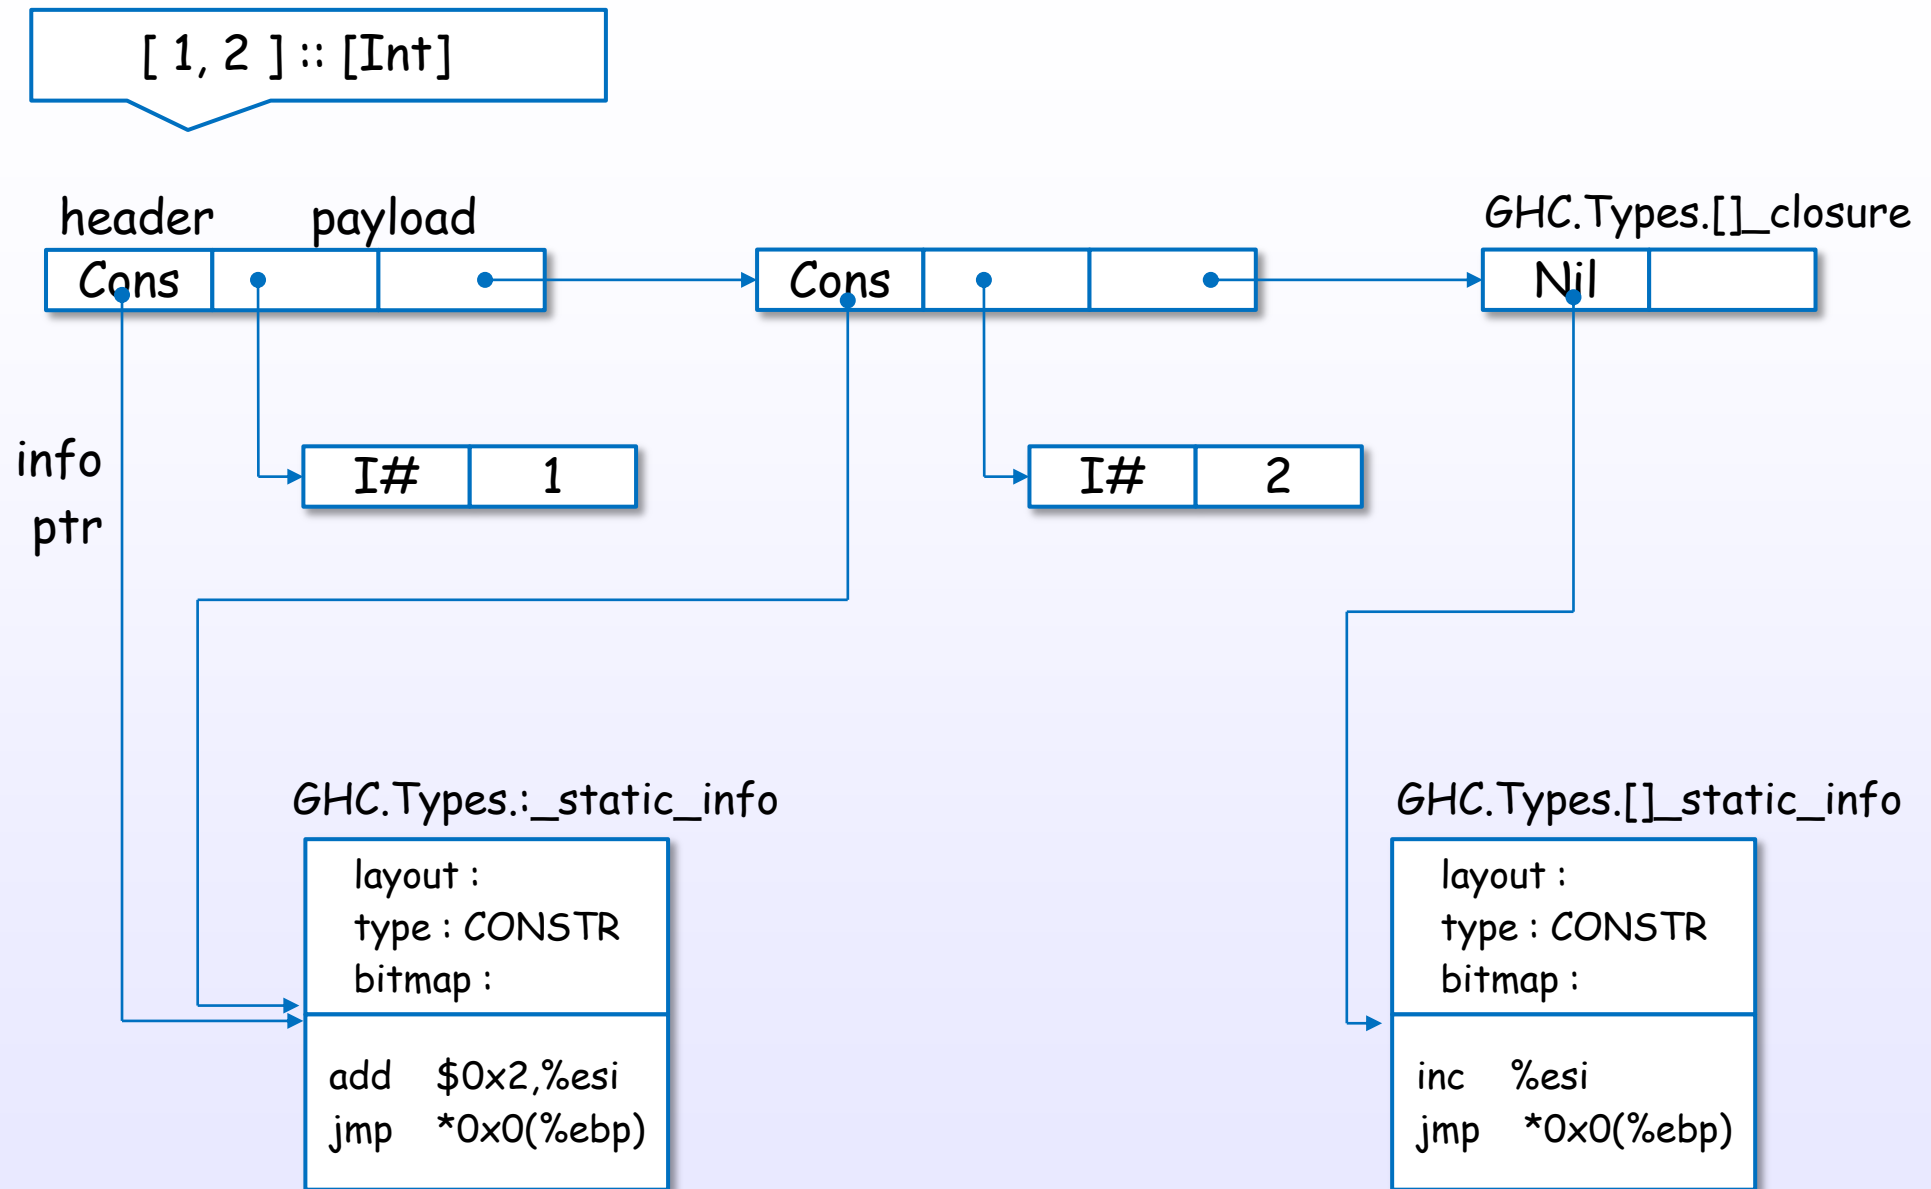
header		payload	
I#		#7	

Closure examples : Maybe

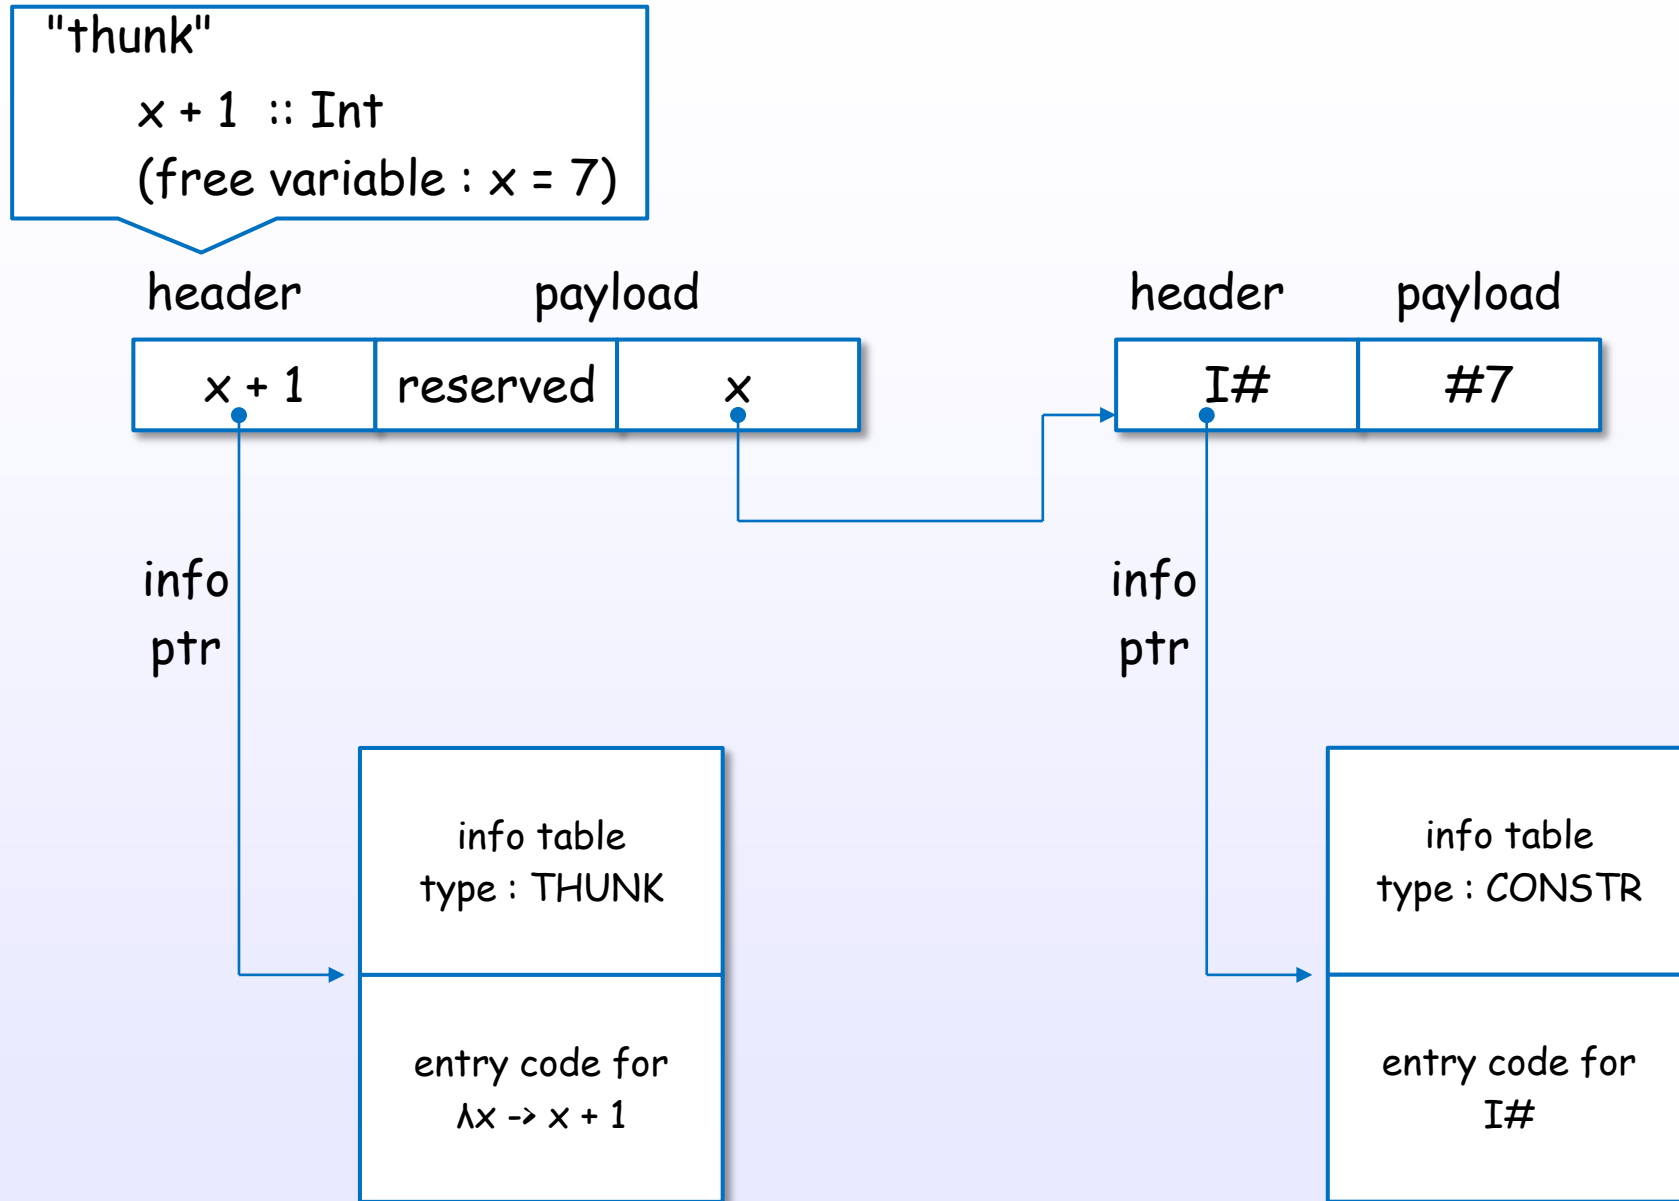
Just 7 :: Maybe Int



Closure examples : List

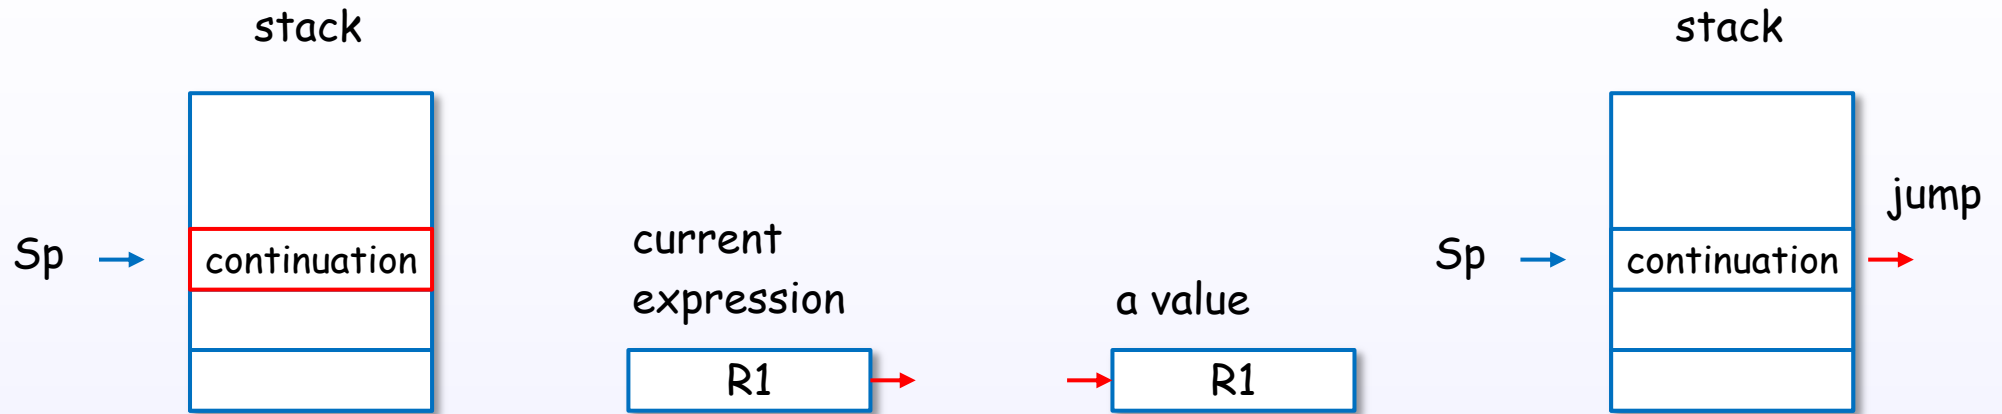


Closure examples : Thunk



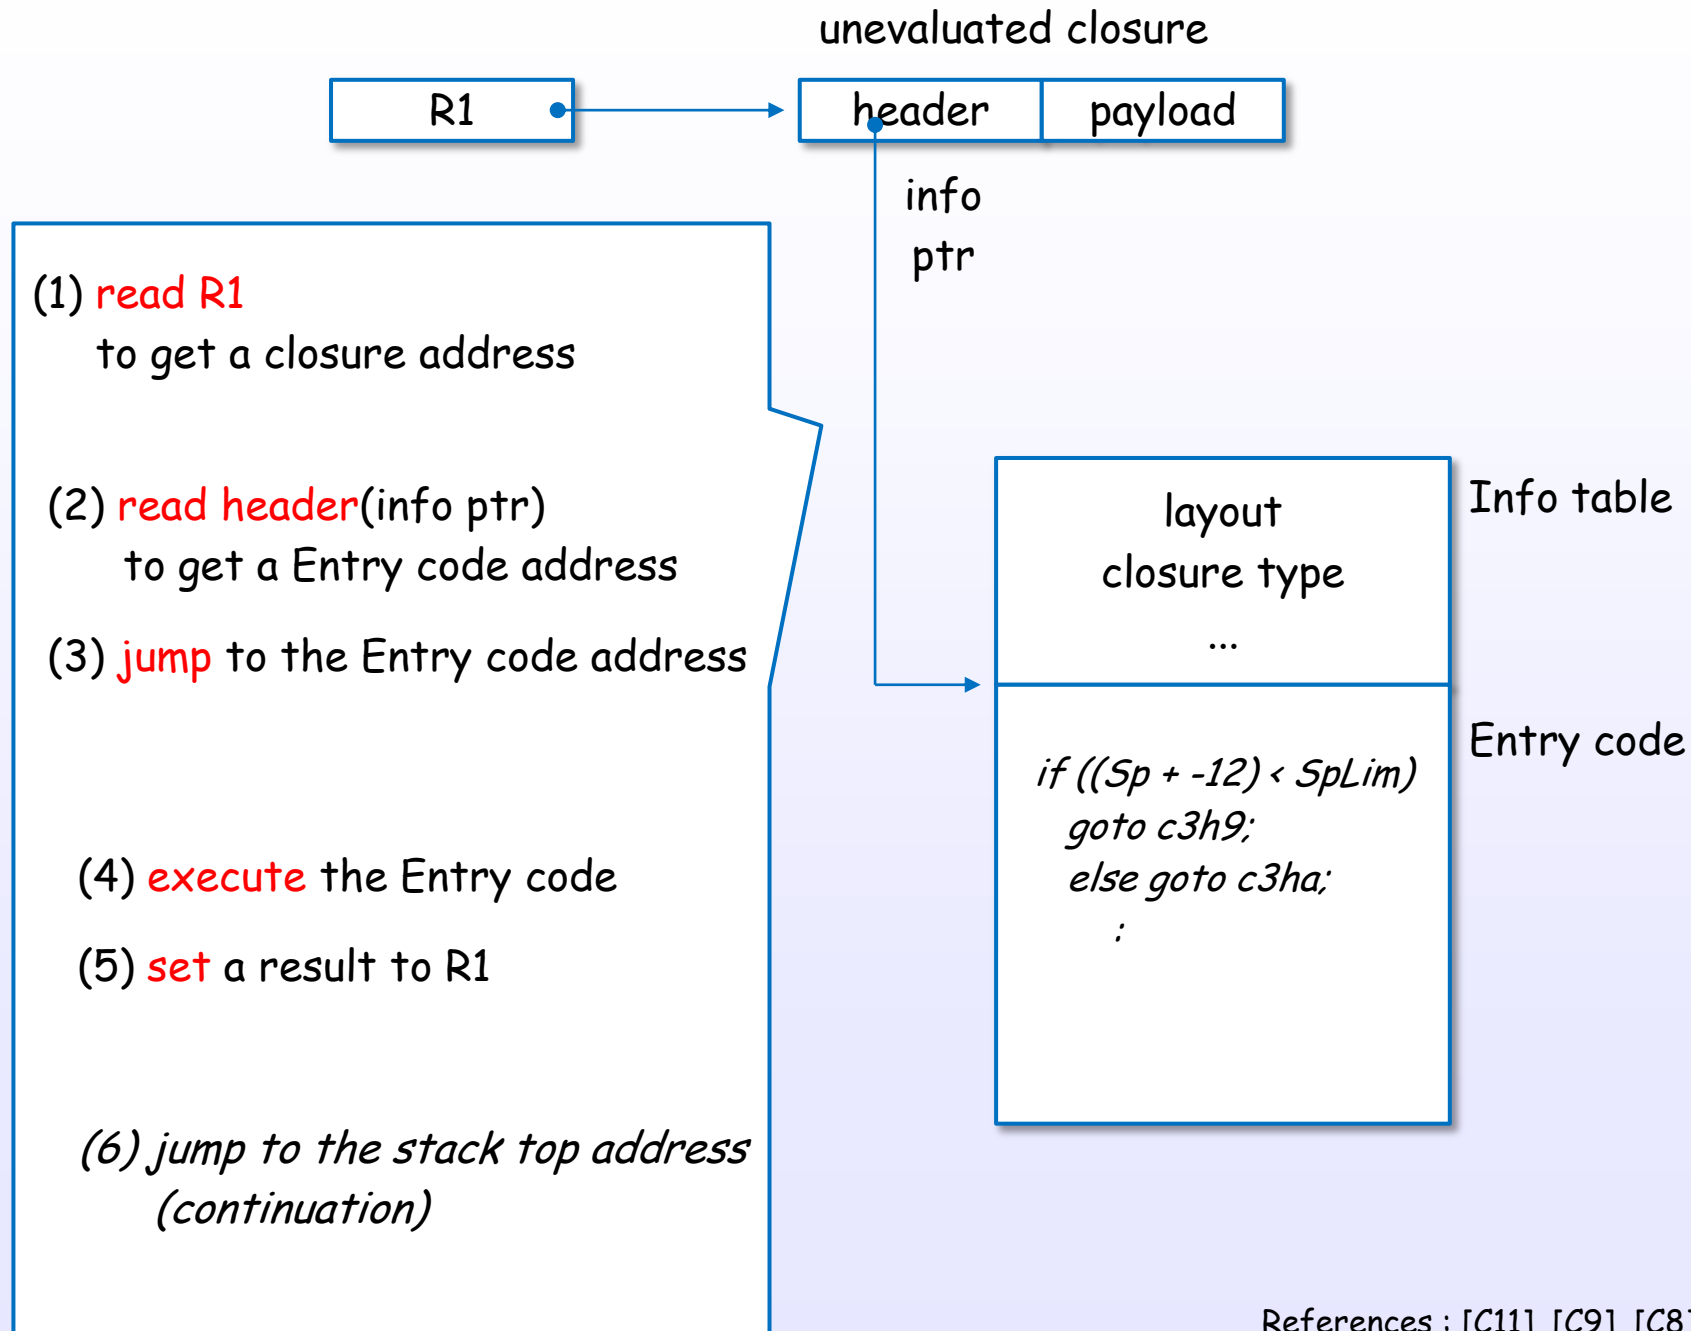
STG-machine evaluation

STG evaluation flow



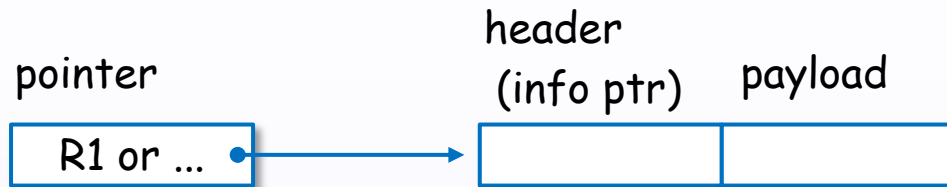
- (1) push a continuation code (next code) to the stack top
- (2) enter to R1 closure
- (3) set a result to R1
- (4) jump (return) to the stack top code
- (5) repeat from (1)

Enter to a closure



Pointer tagging

Pointer tagging



pointer



... an unevaluated closure



... an evaluated closure;
1st constructor value or evaluated.
(for instance: "Nothing")



... an evaluated closure; 2nd constructor value.
(for instance: "Just xx")



... an evaluated closure; 3rd constructor value.

* 32bit machine case

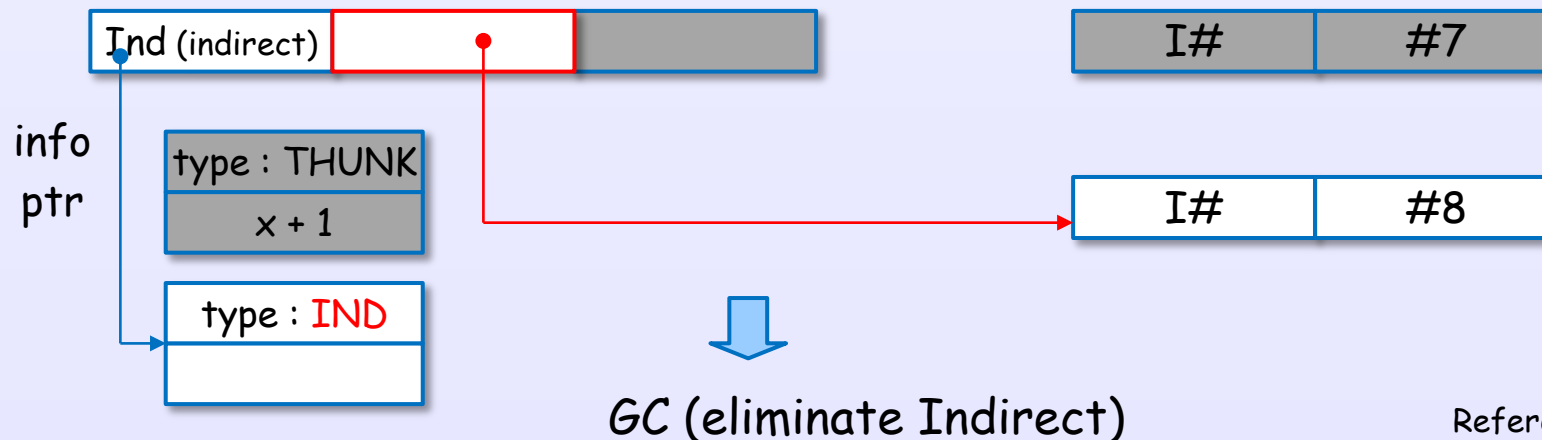

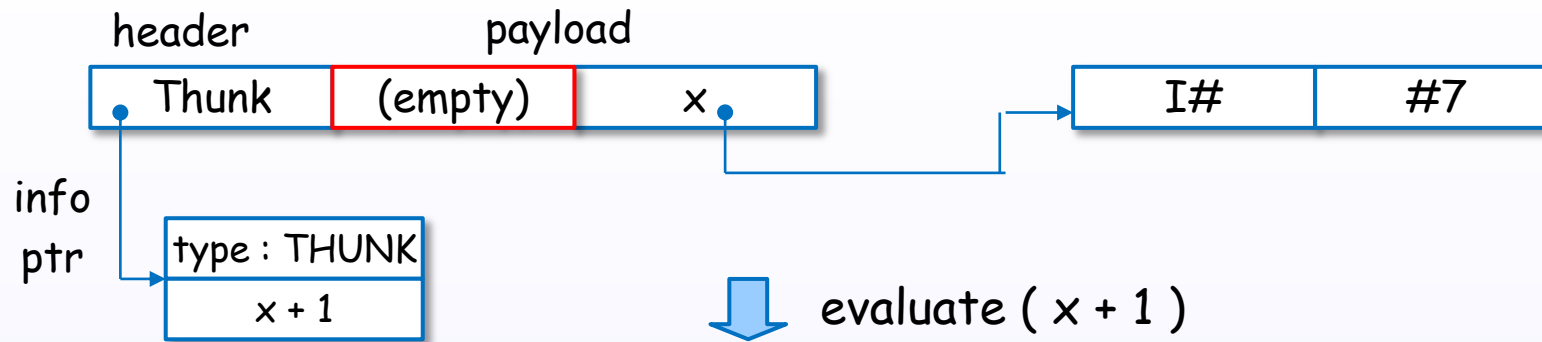
fast judgment!

checking only pointer's lower bits without evaluating the closure.

Think and update

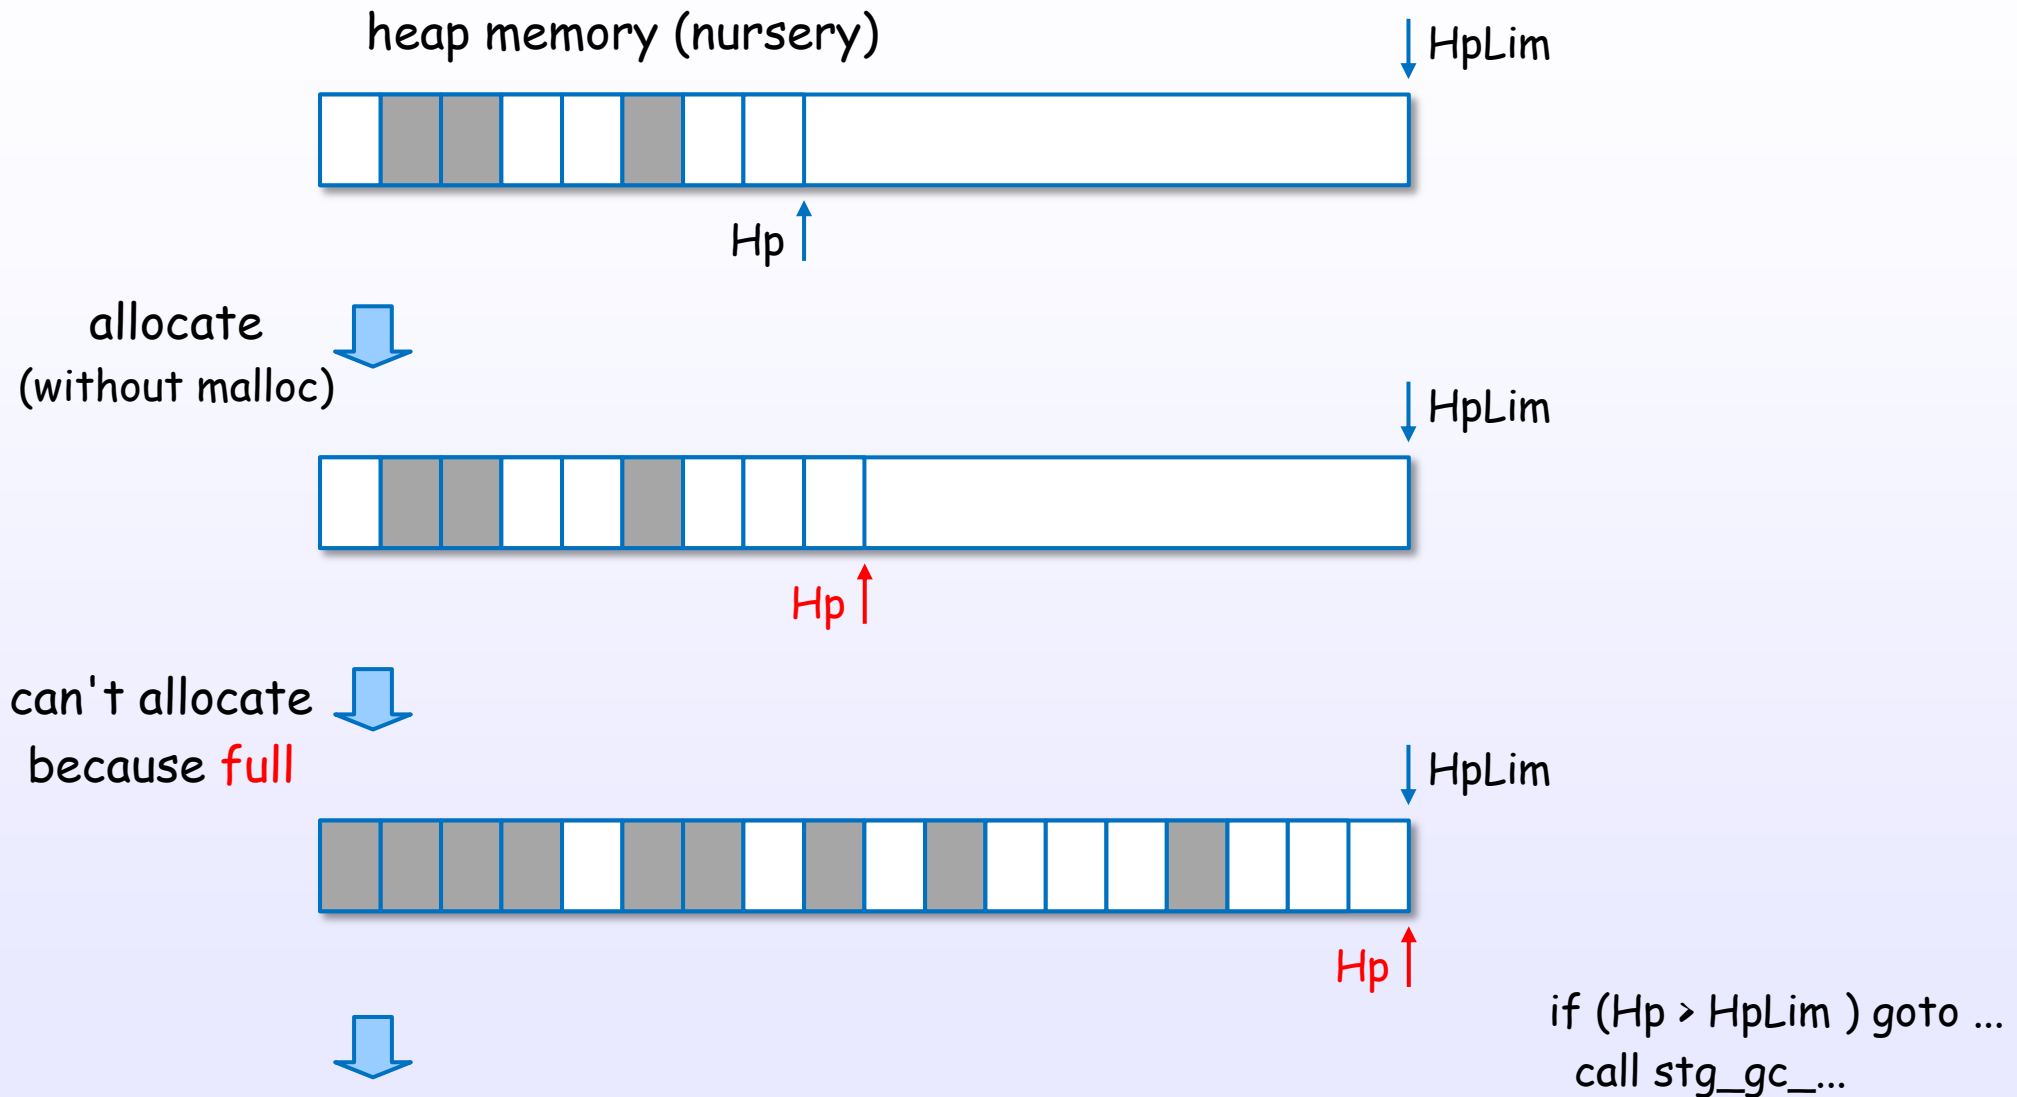
Thunk and update

"thunk" $x + 1 :: \text{Int}$ (free variable : $x = 7$)




Allocate and free heap objects

Allocate heap objects

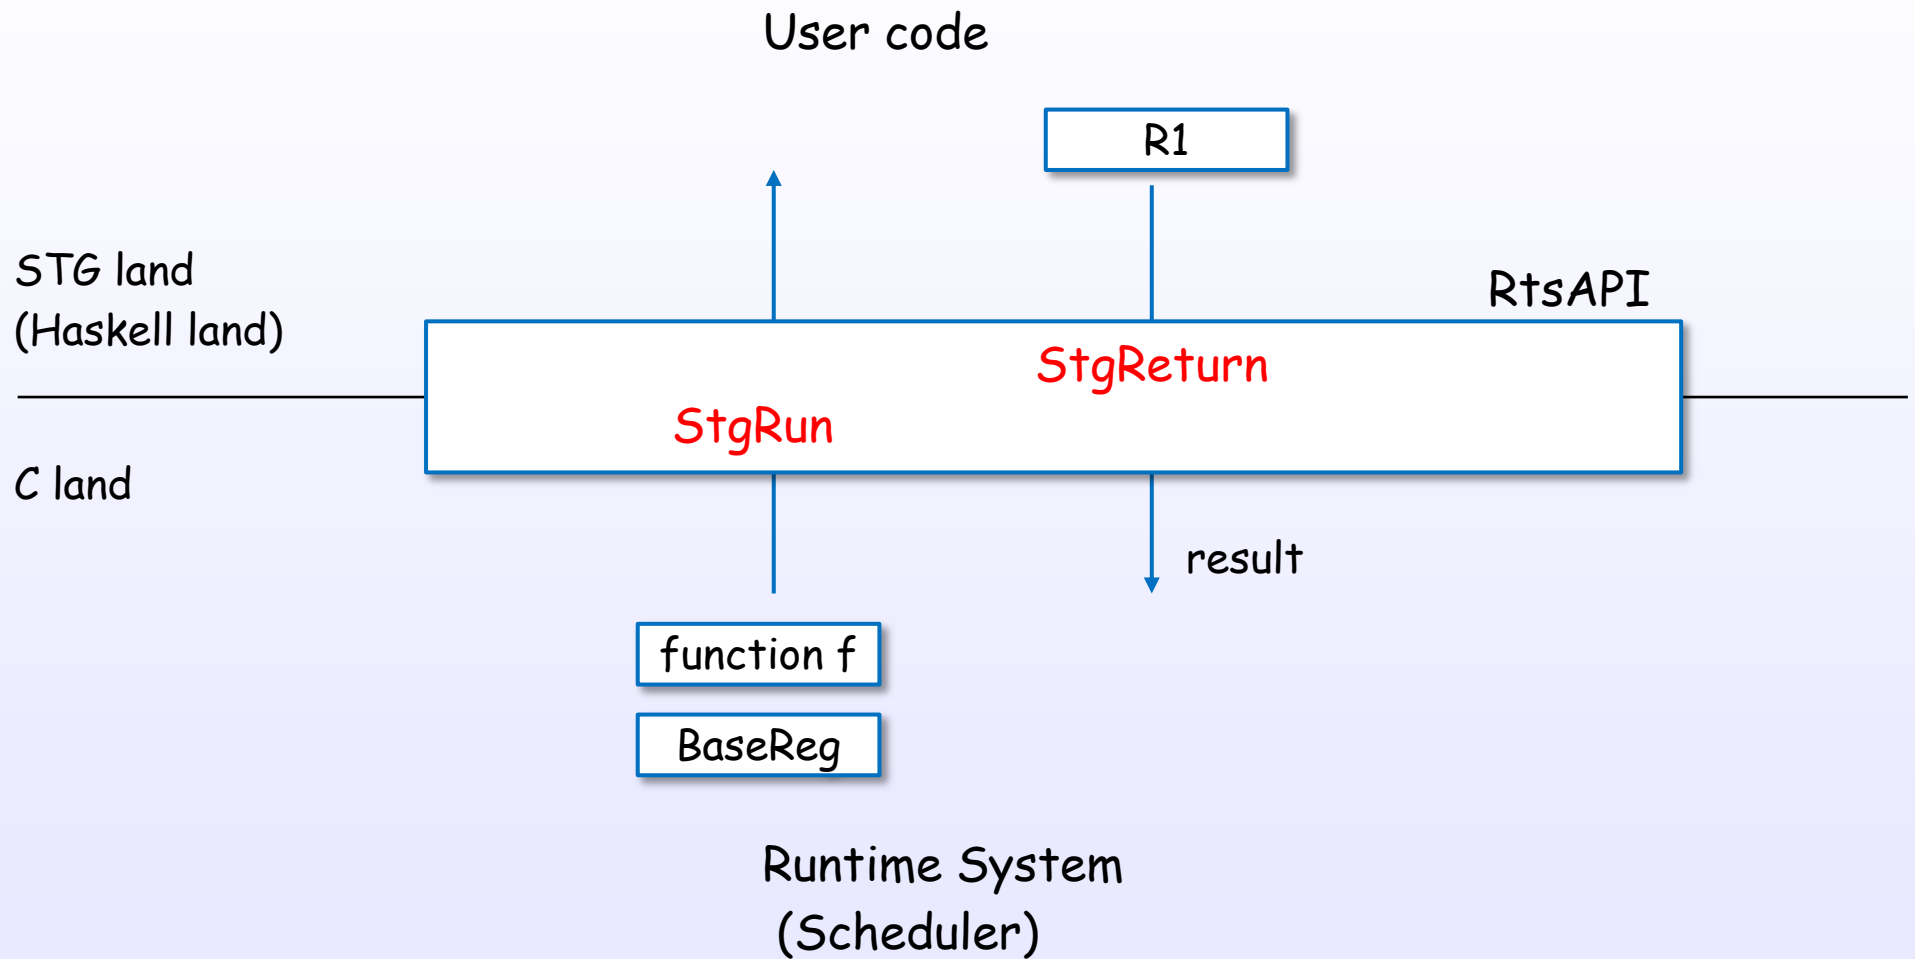


free and collect heap objects



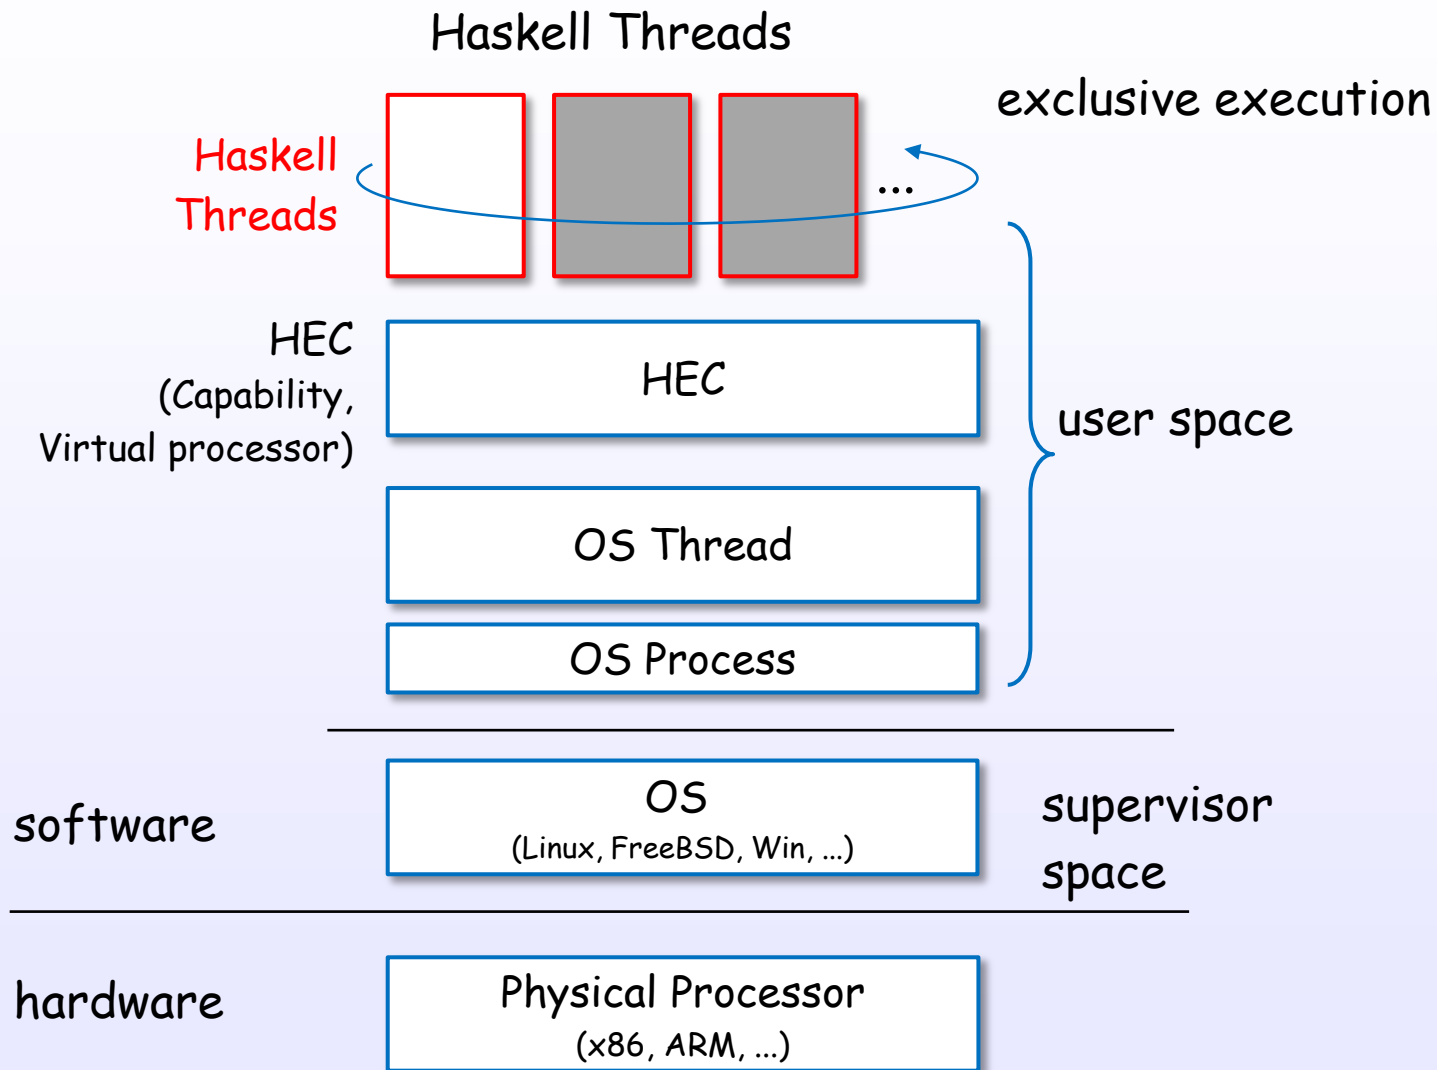
STG - C land interface

STG (Haskell) land - C land interface

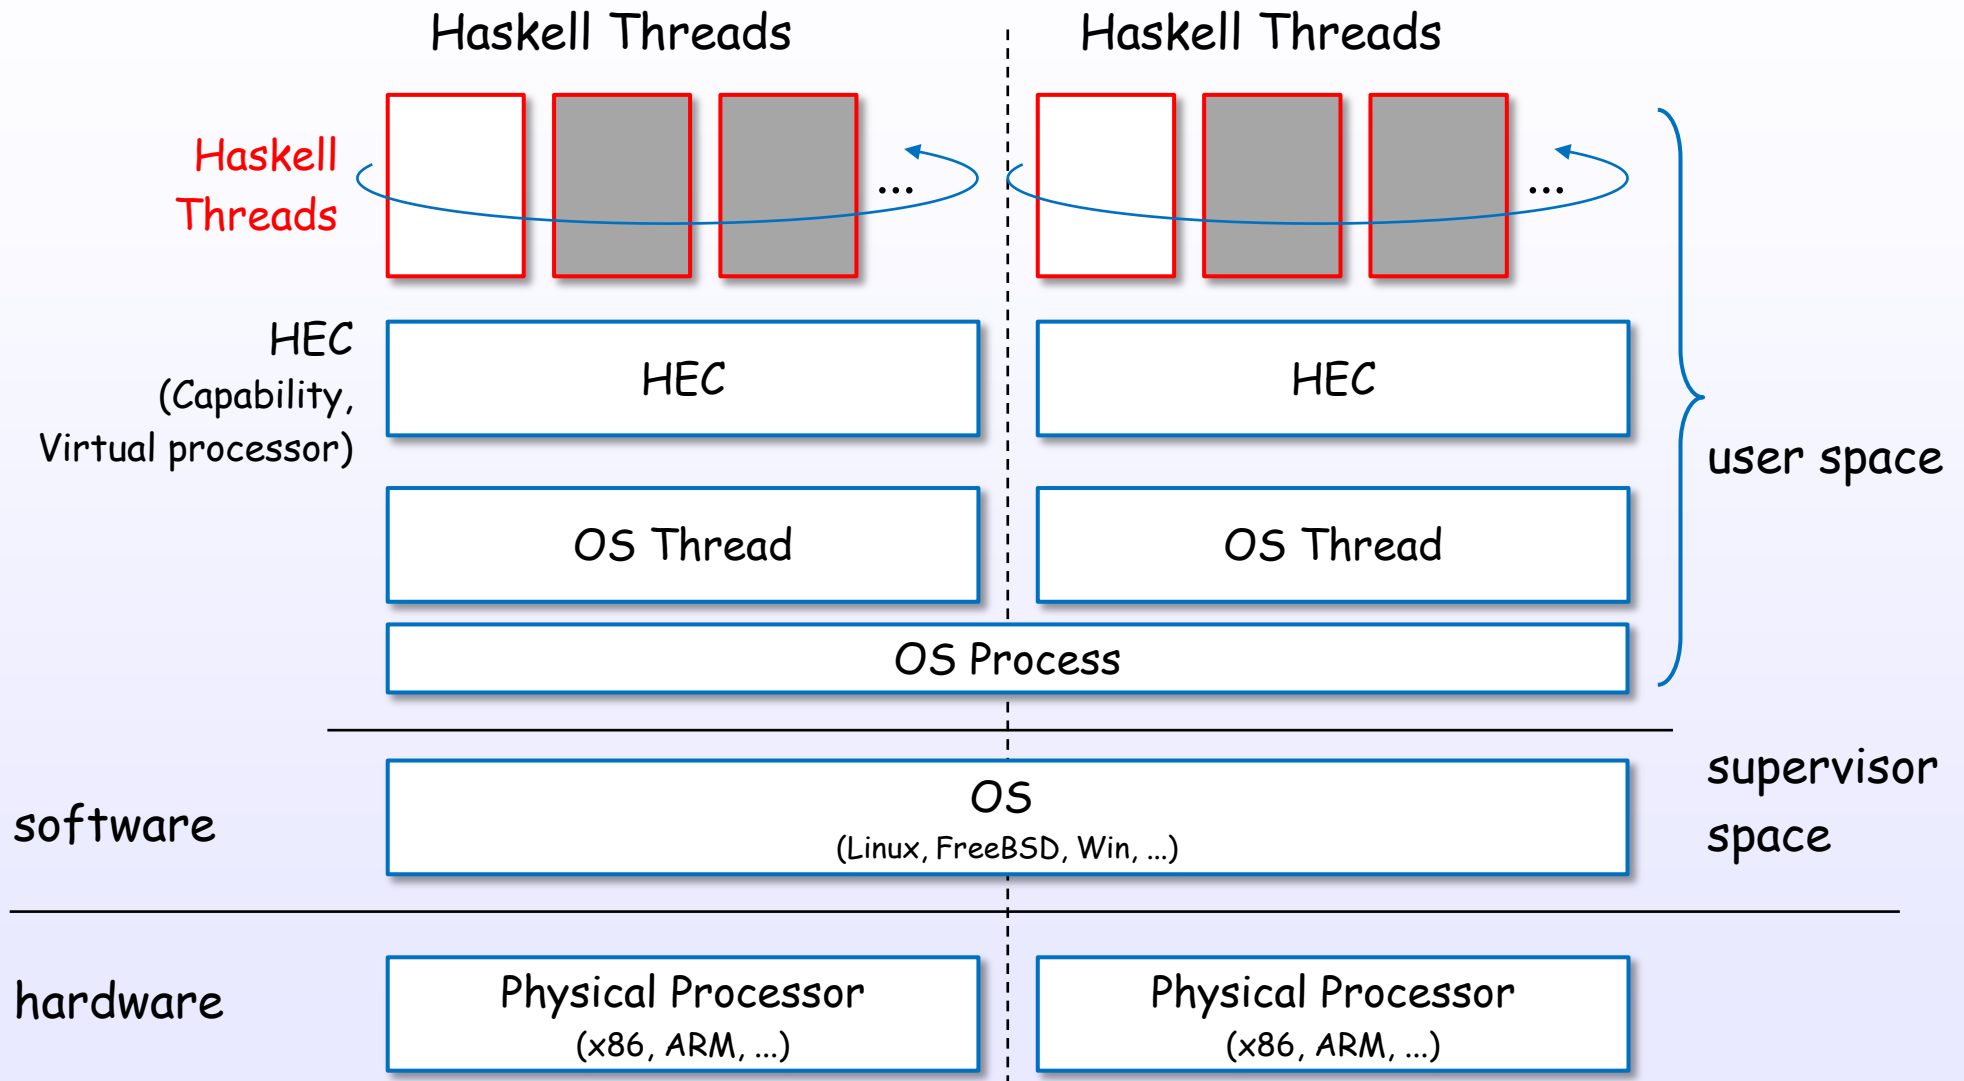


Thread

Thread layer (single core)



Thread layer (multi core)

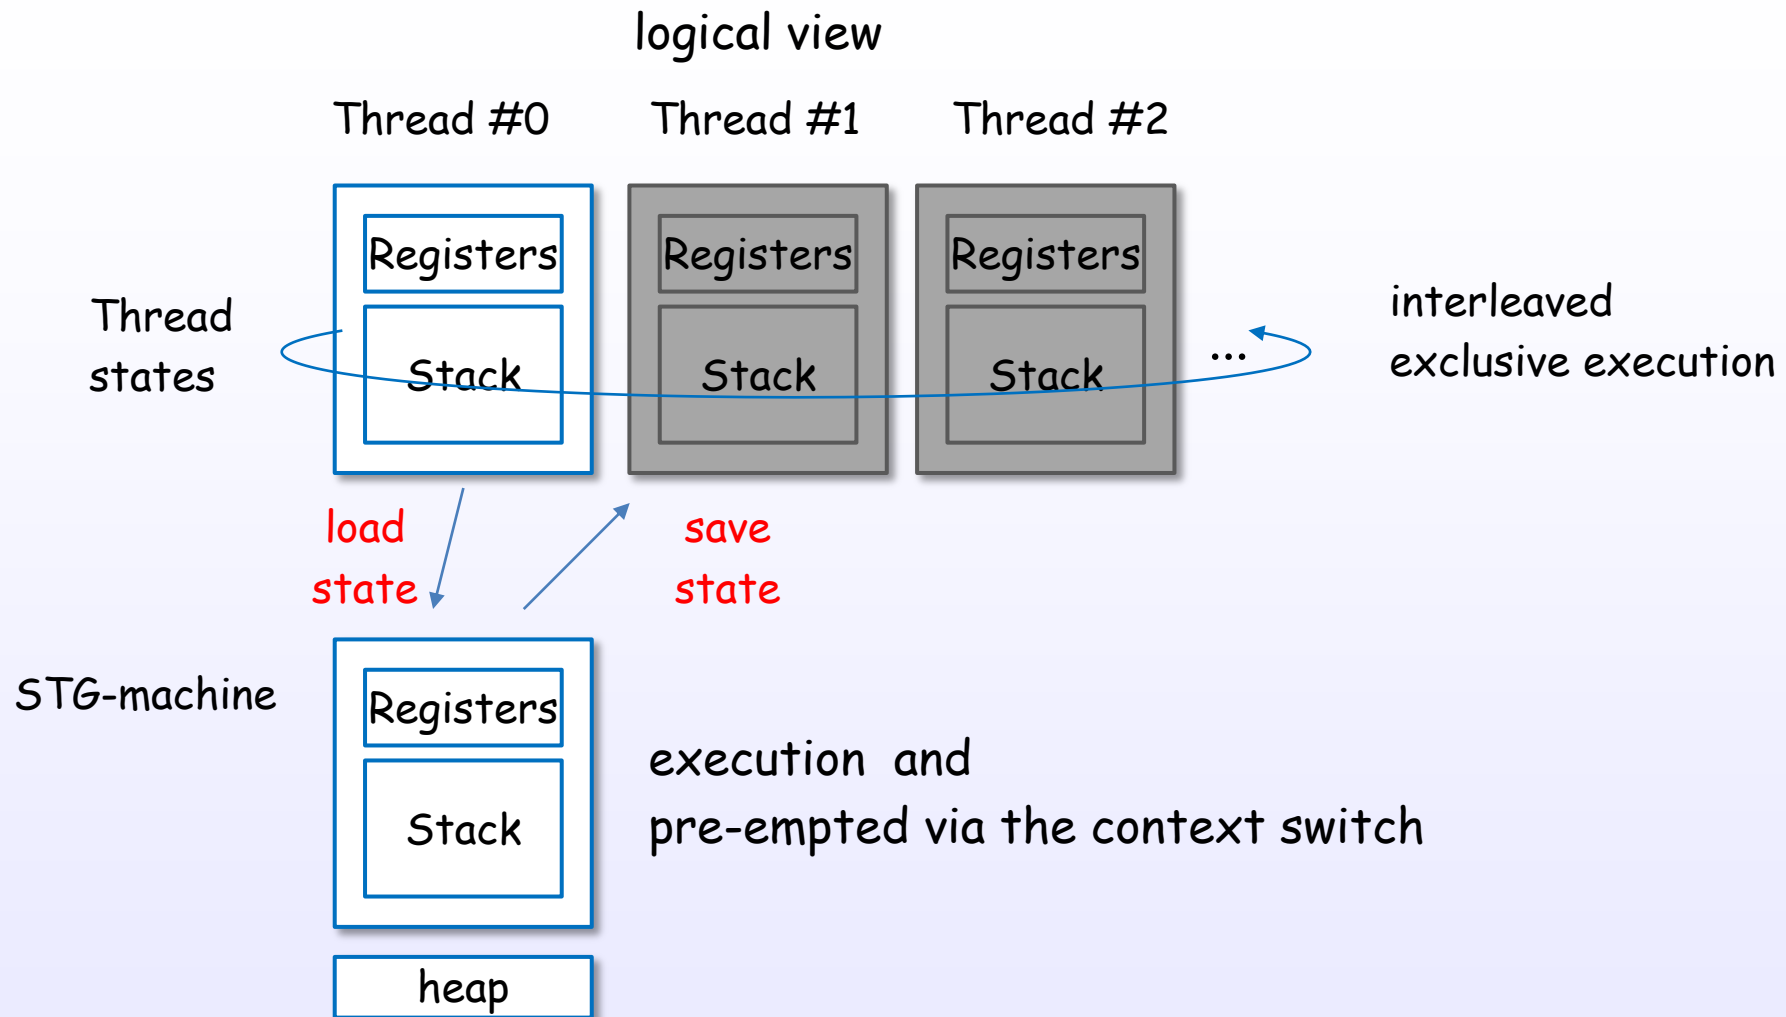


*Threaded option case (ghc -threaded)

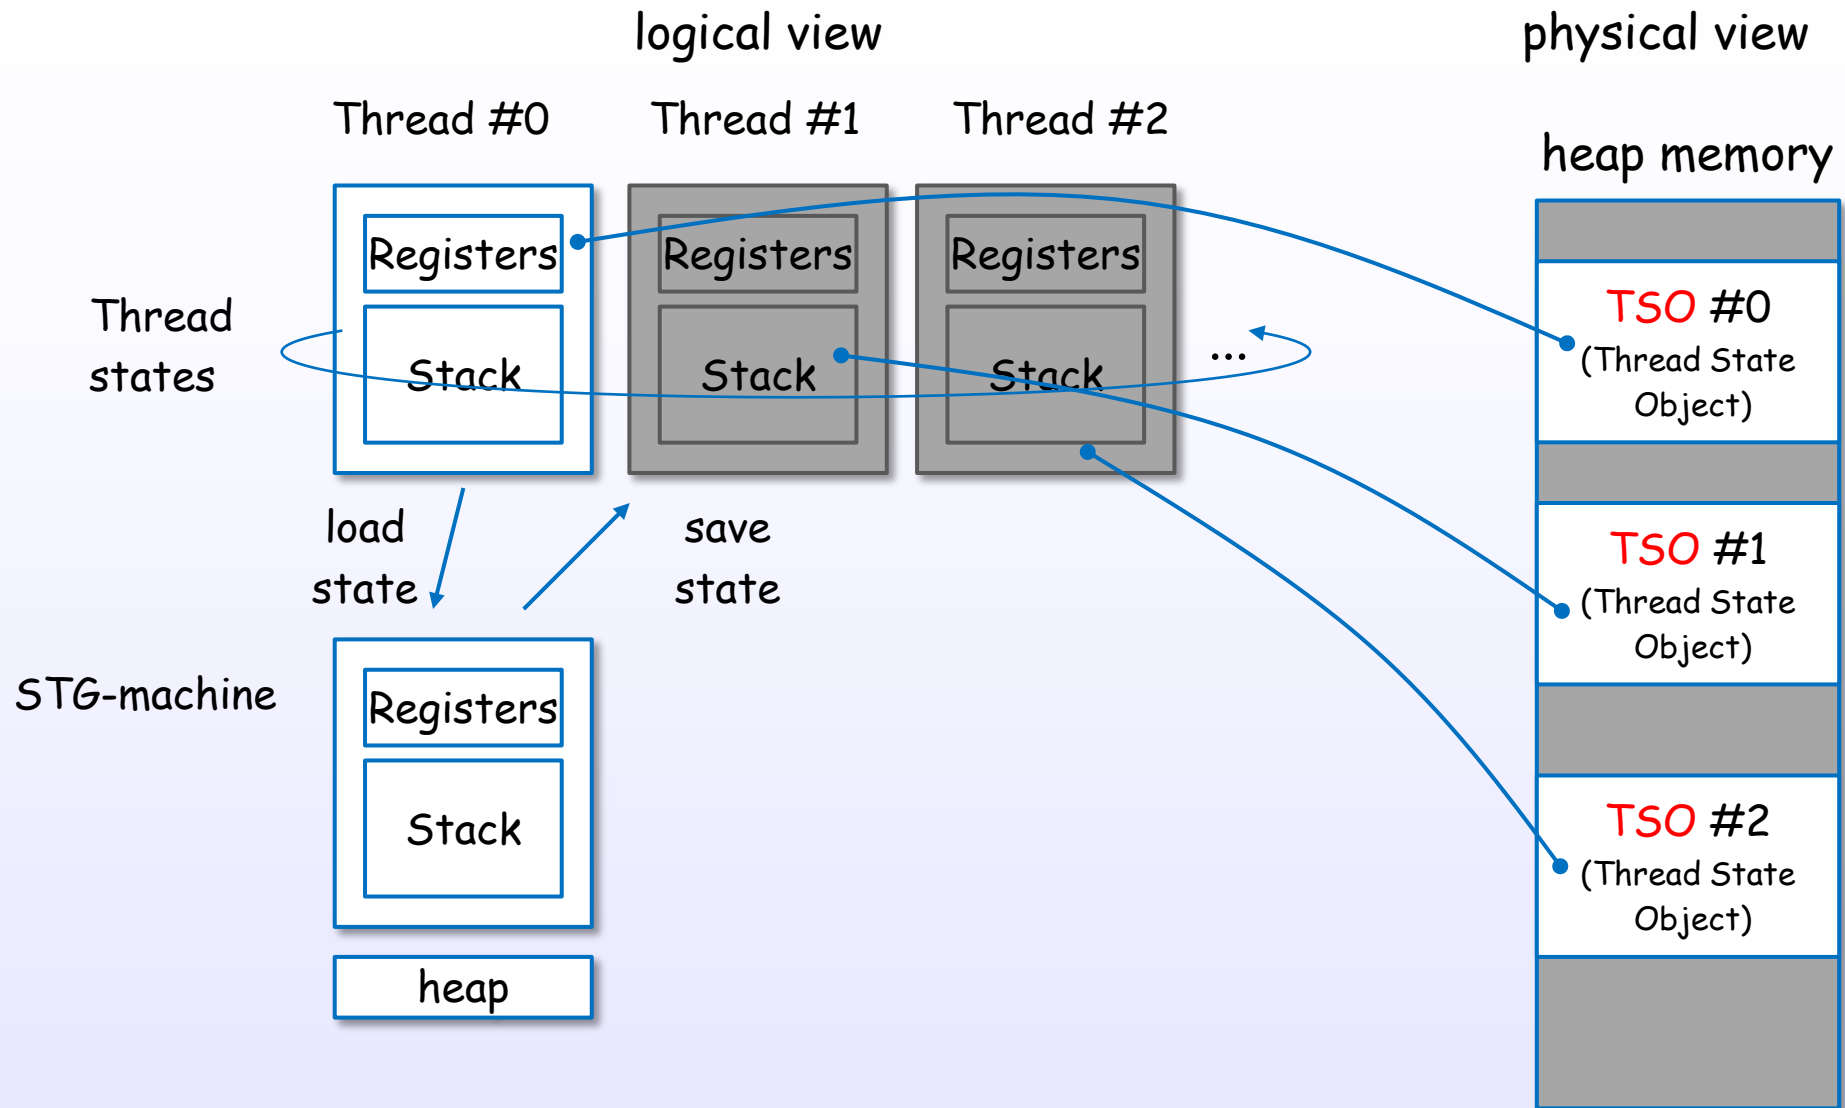
References : [5], [8], [9], [14], [C17], [C11], [19], [S17], [S16], [S23], [S22], [S14]

Thread context switch

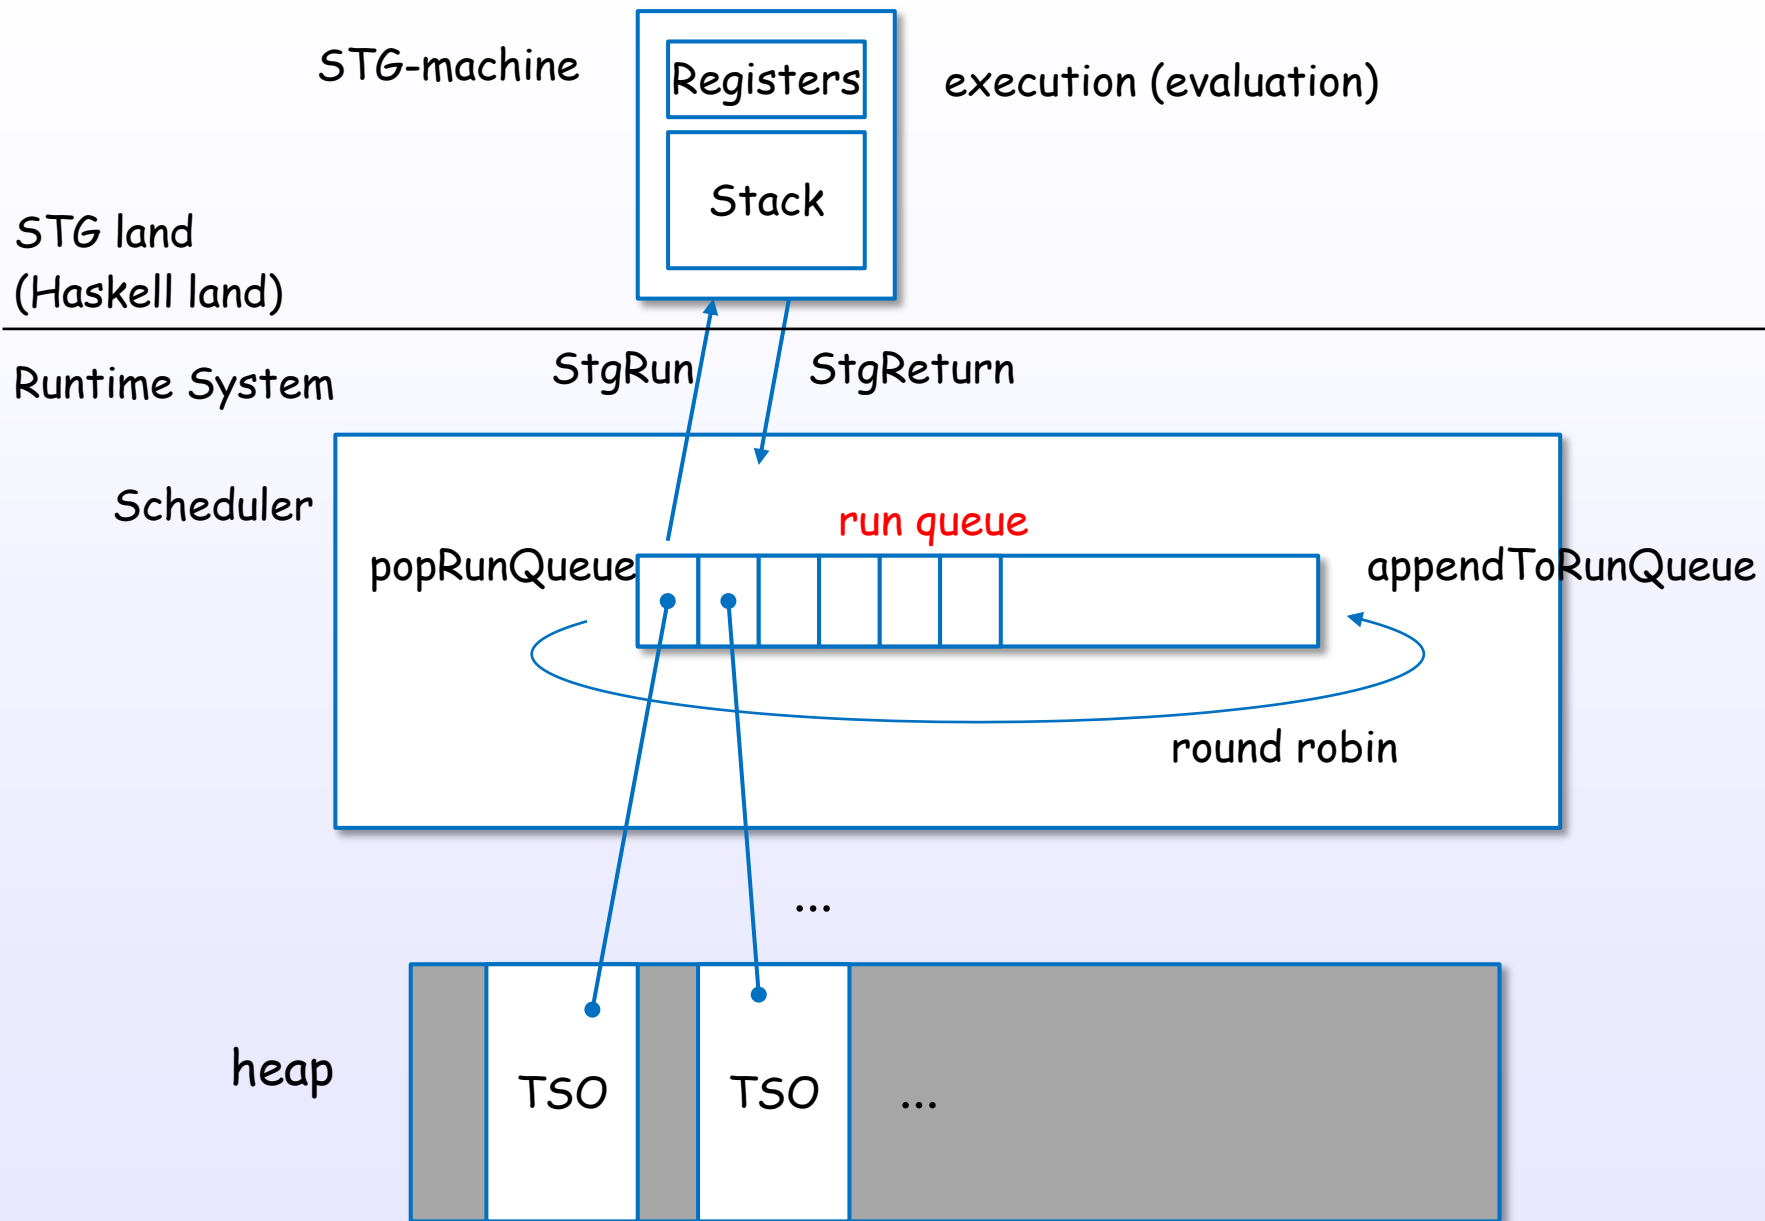
Threads and context switch



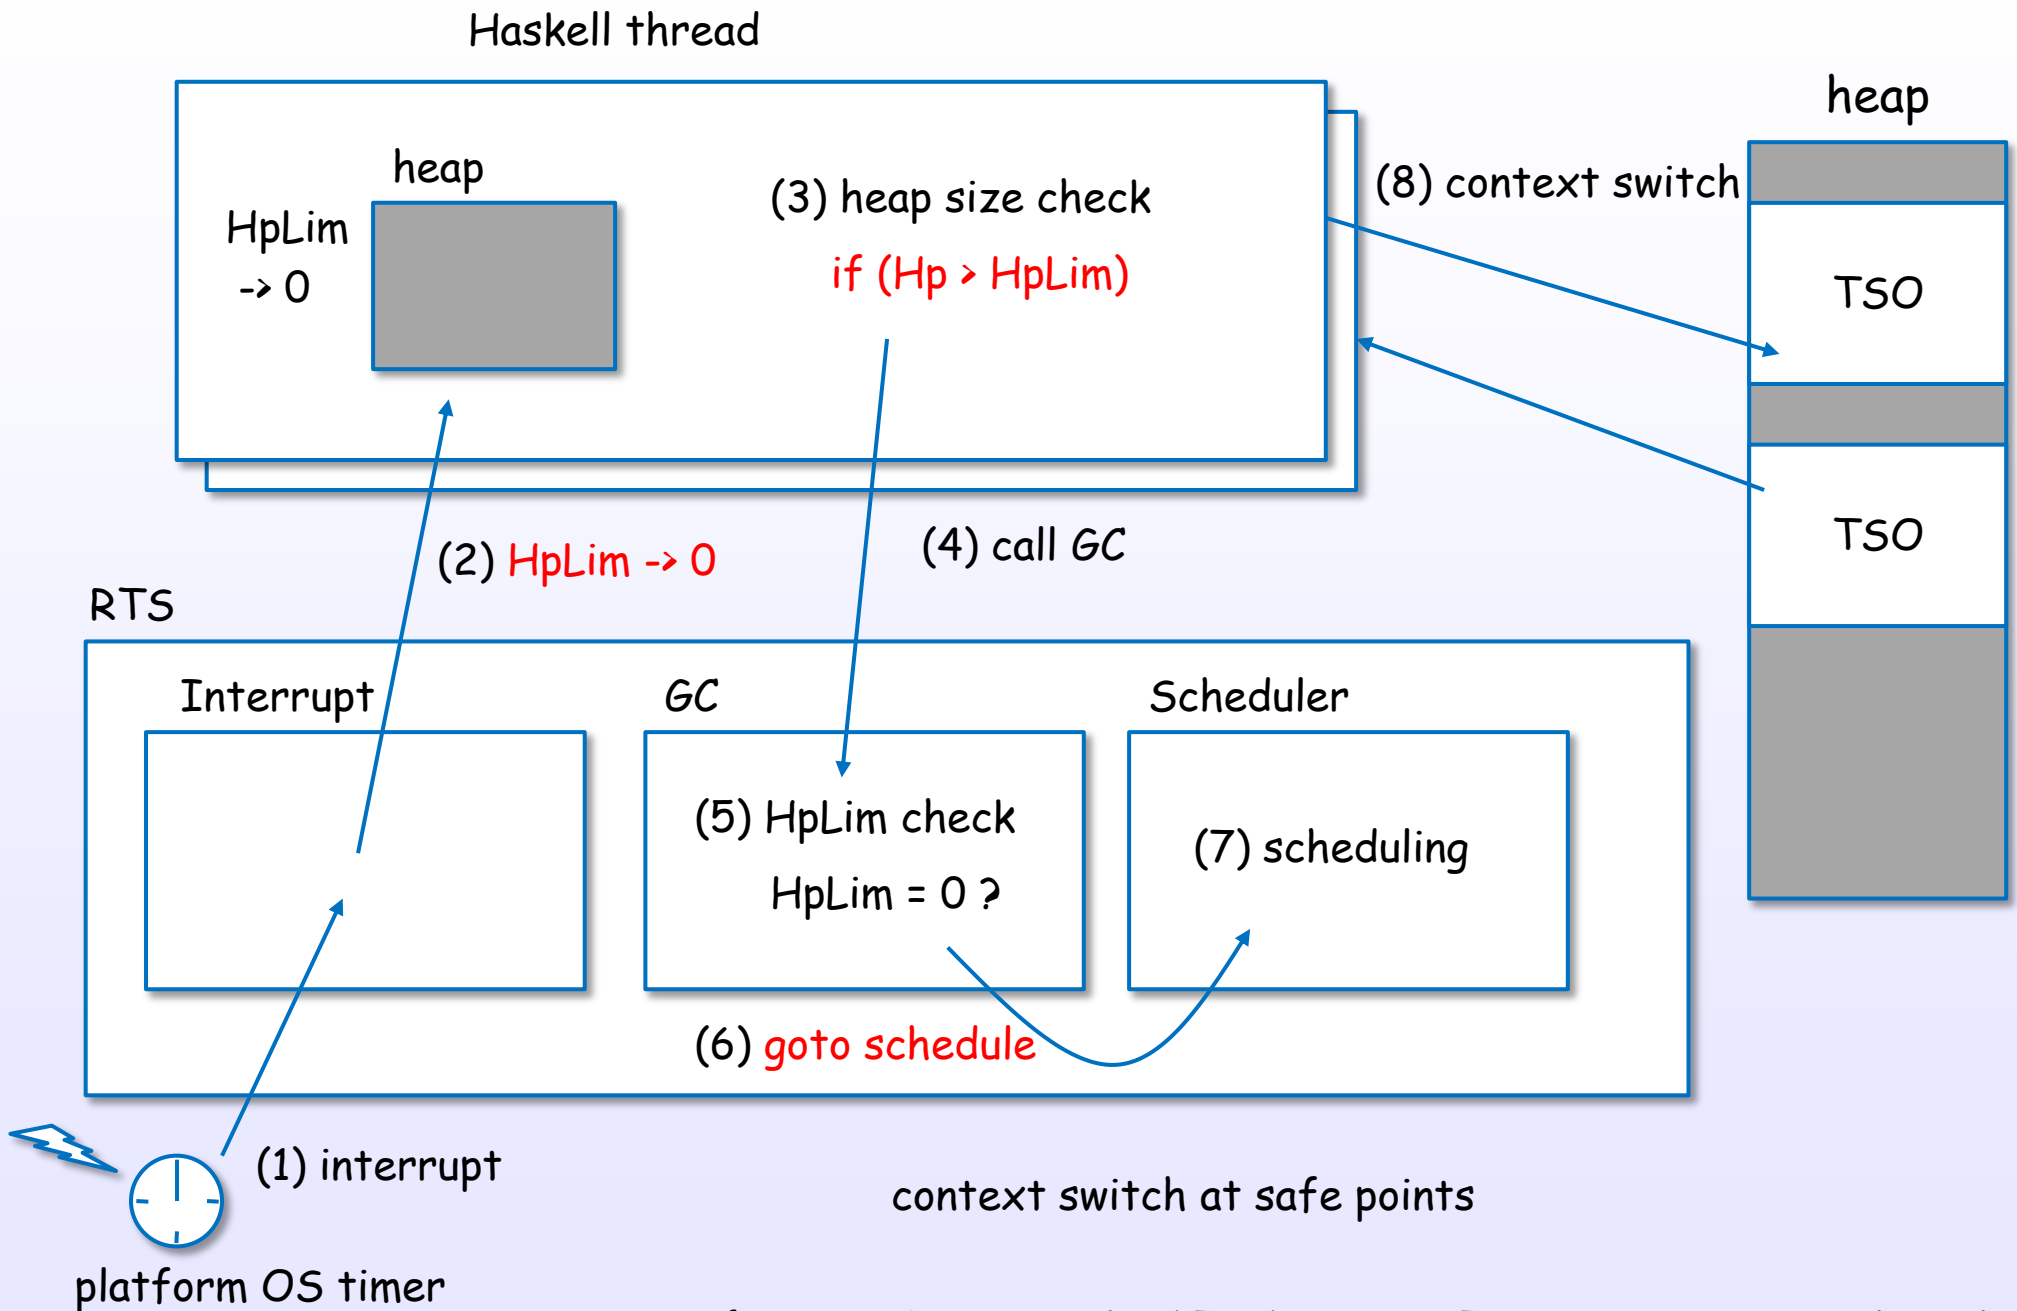
Threads and TSOs



Scheduling by run queue



Context switch flow



Context switch flow (code)

stg_gc_noregs

```
if (HpLim == 0) {  
    jump stg_returnToSched [R1];  
}
```

stg_returnToSched

```
W_ r1;  
r1 = R1; // foreign calls may clobber R1  
SAVE_THREAD_STATE();  
foreign "C" threadPaused(MyCapability()  
    "ptr", CurrentTSO);  
R1 = r1;  
jump StgReturn [R1];
```

STG land
(Haskell land)

C land

```
cap->r.rHpLim = NULL;
```

schedule

stopCapability
contextSwitchCapability
contextSwitchAllCapabilities
handle_tick

CreateTimerQueue

initTicker

initTimer

startTimer

hs_init_ghc

real_main

next
handle_tick ..

OS

*Windows case

References : [5], [8], [9], [14], [C17], [C11], [19], [S17], [S16], [S23], [S22], [S14], [S24]

Creating main and sub threads

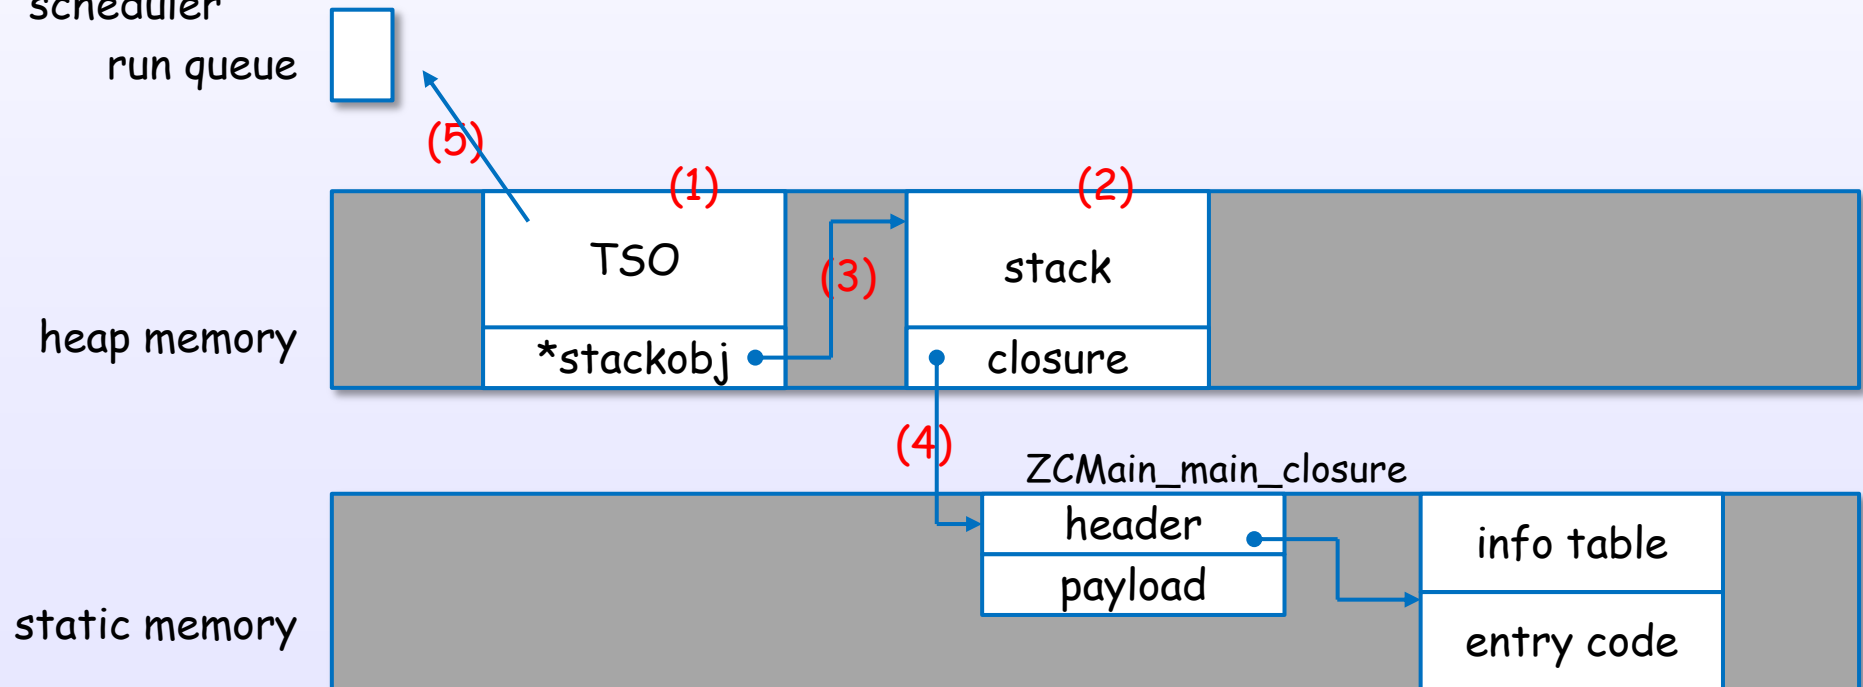
Create a main thread

C land

Runtime system bootstrap code [rts/RtsAPI.c]

```
rts_evalLazyIO
  createIOThread
    createThread ... (1), (2), (3)
    pushClosure ... (4)
  scheduleWaitThread
    appendToRunQueue ... (5)
```

scheduler
run queue



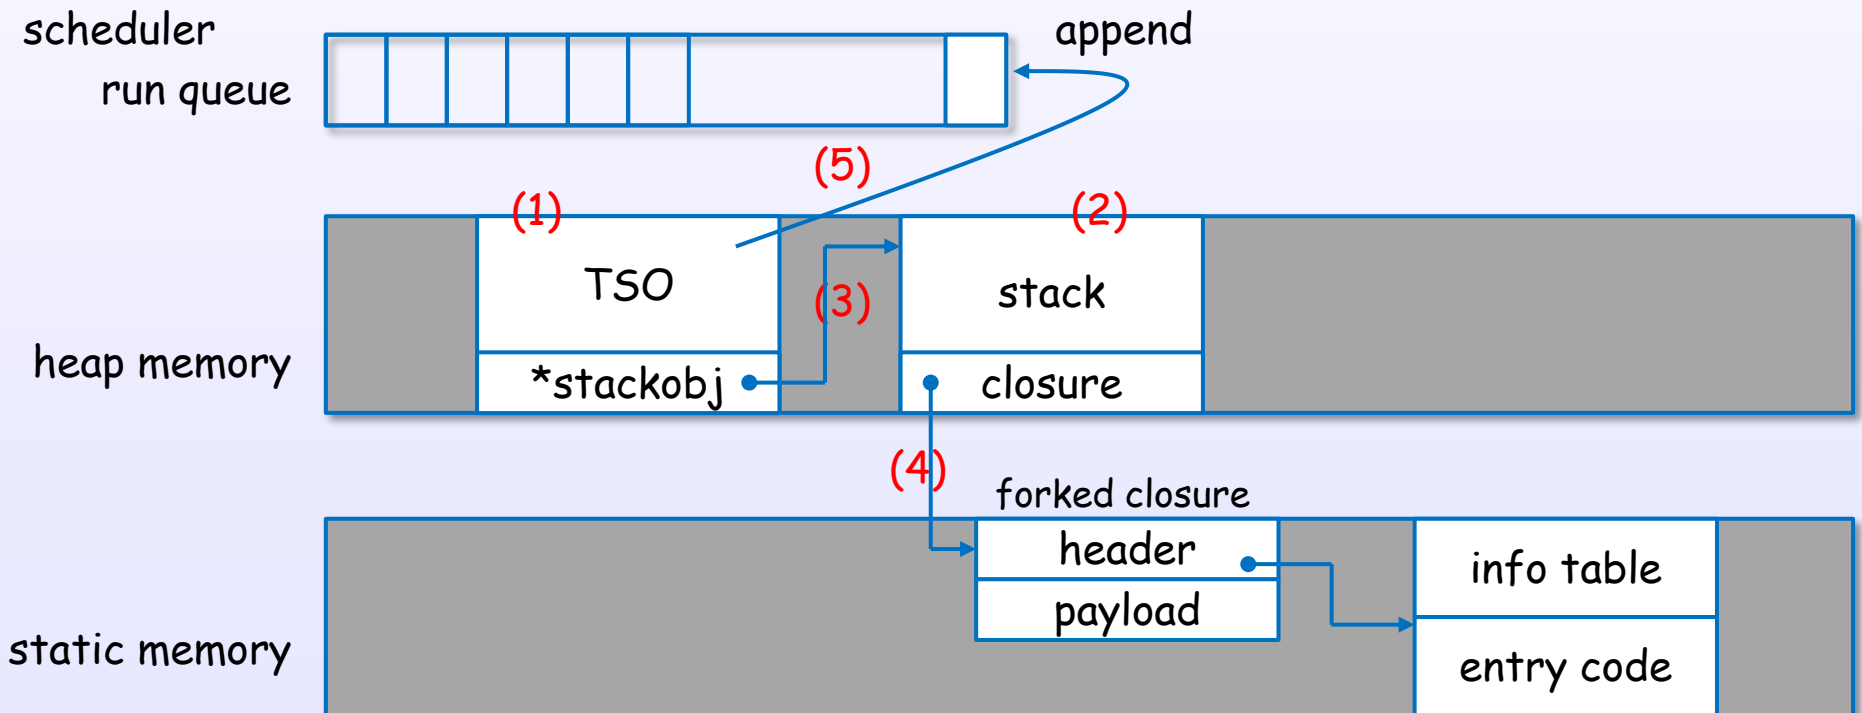
Create a sub thread by forkIO

Haskell Threads

```
forkIO
  stg_forkzh
    ccall createIOThread ... (1), (2), (3), (4)
    ccall scheduleThread ... (5)
```

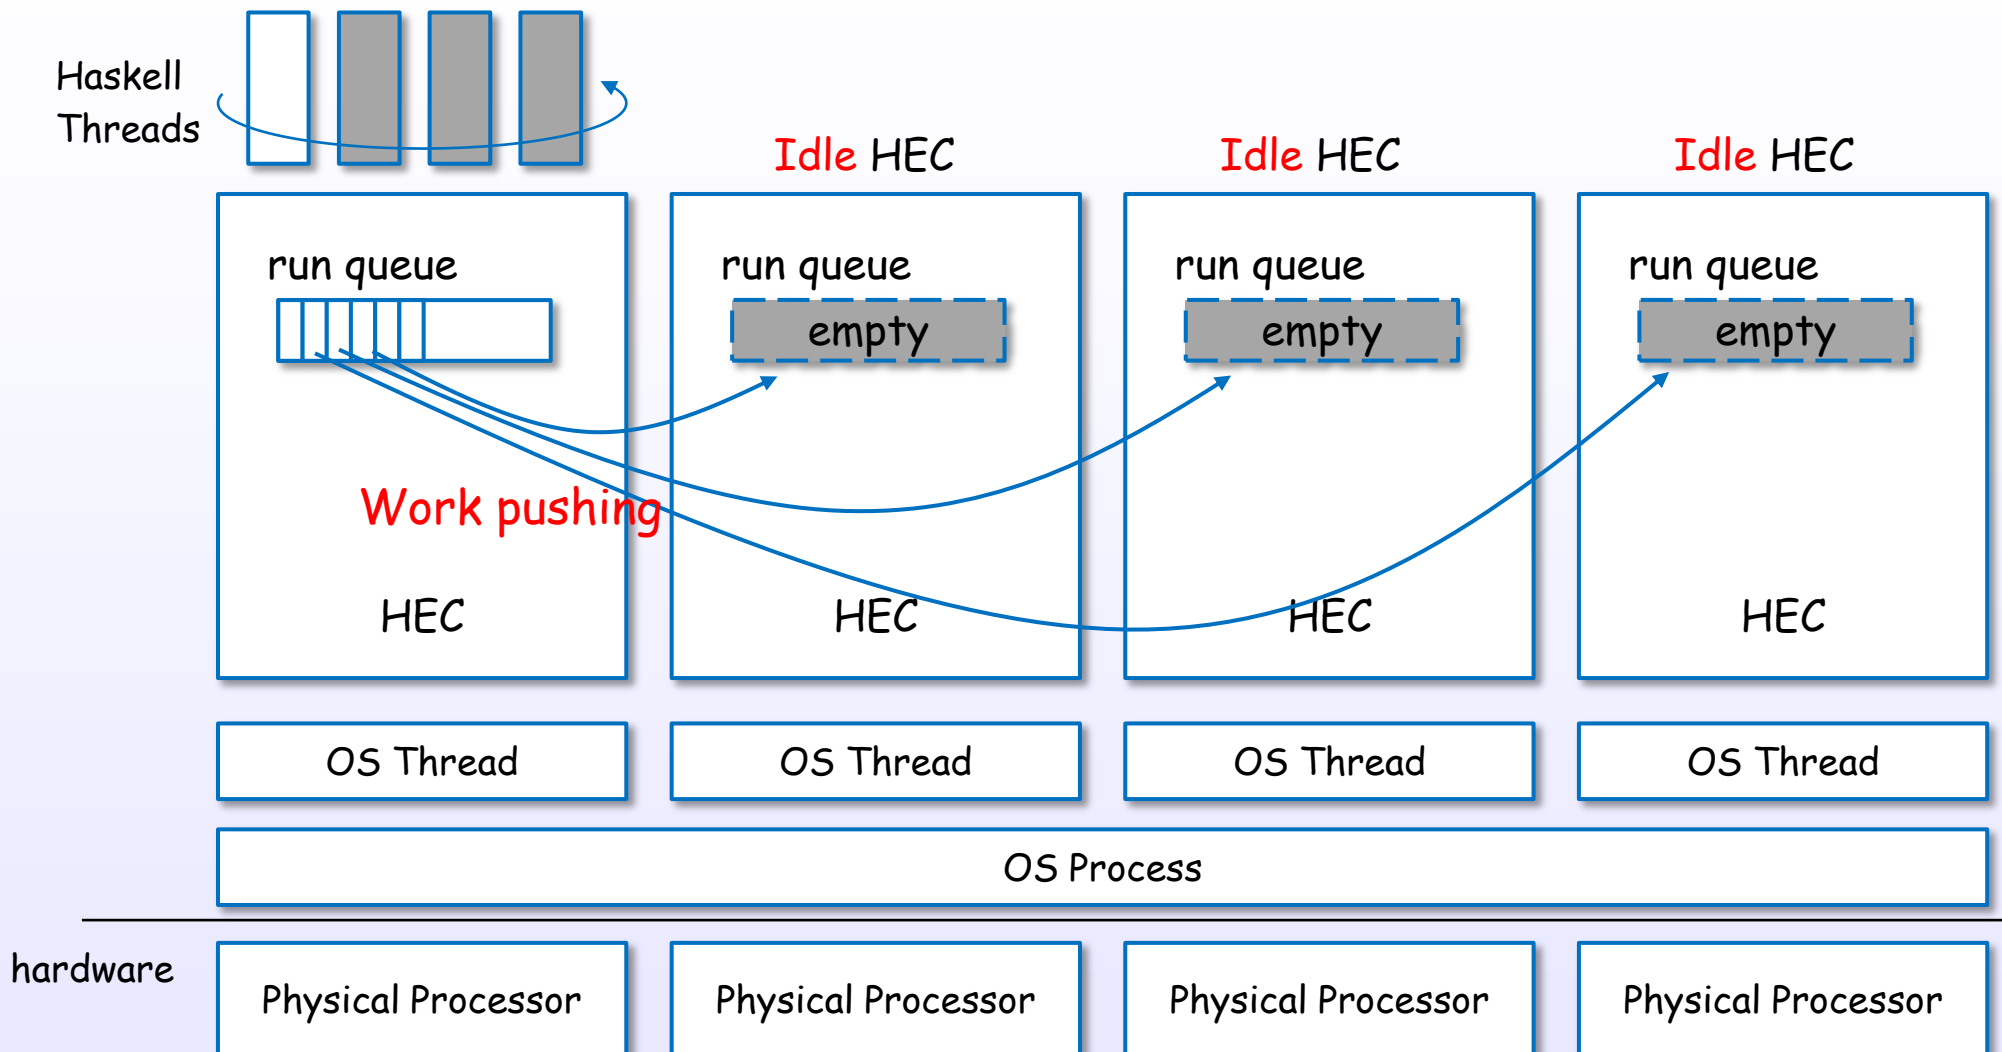
STG land
(Haskell land)

C land



Thread migration

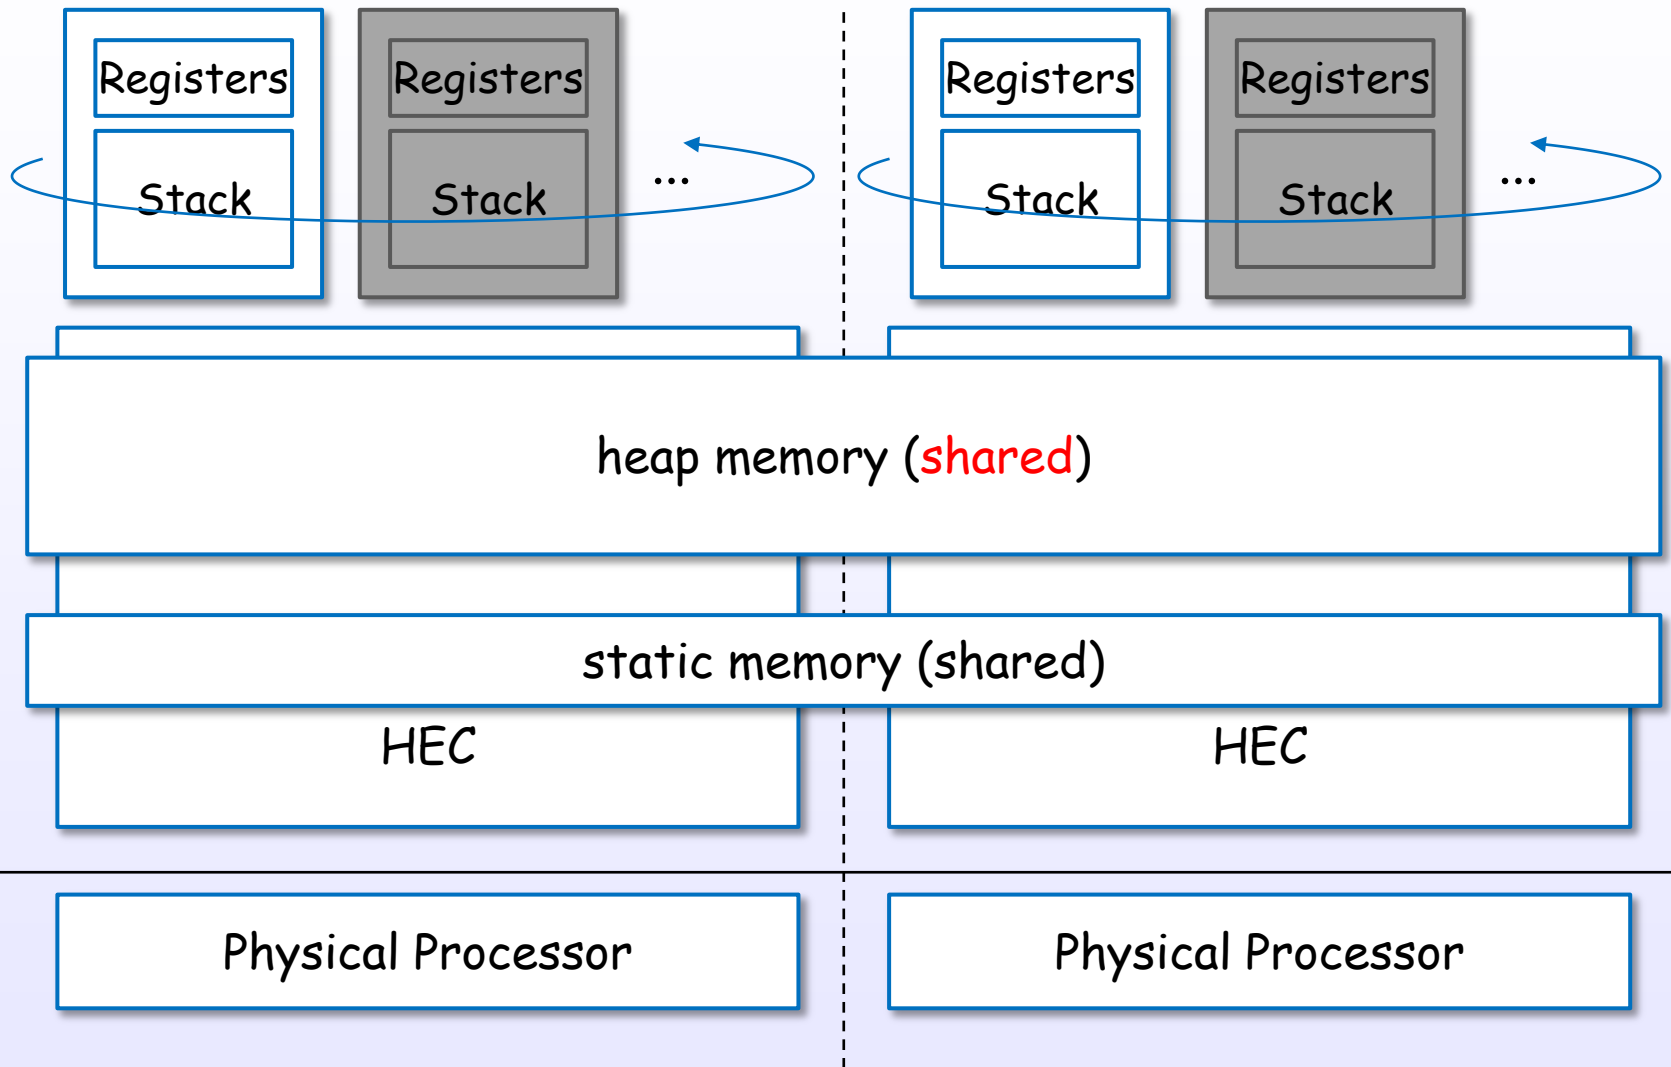
Threads are migrated to idle HECs



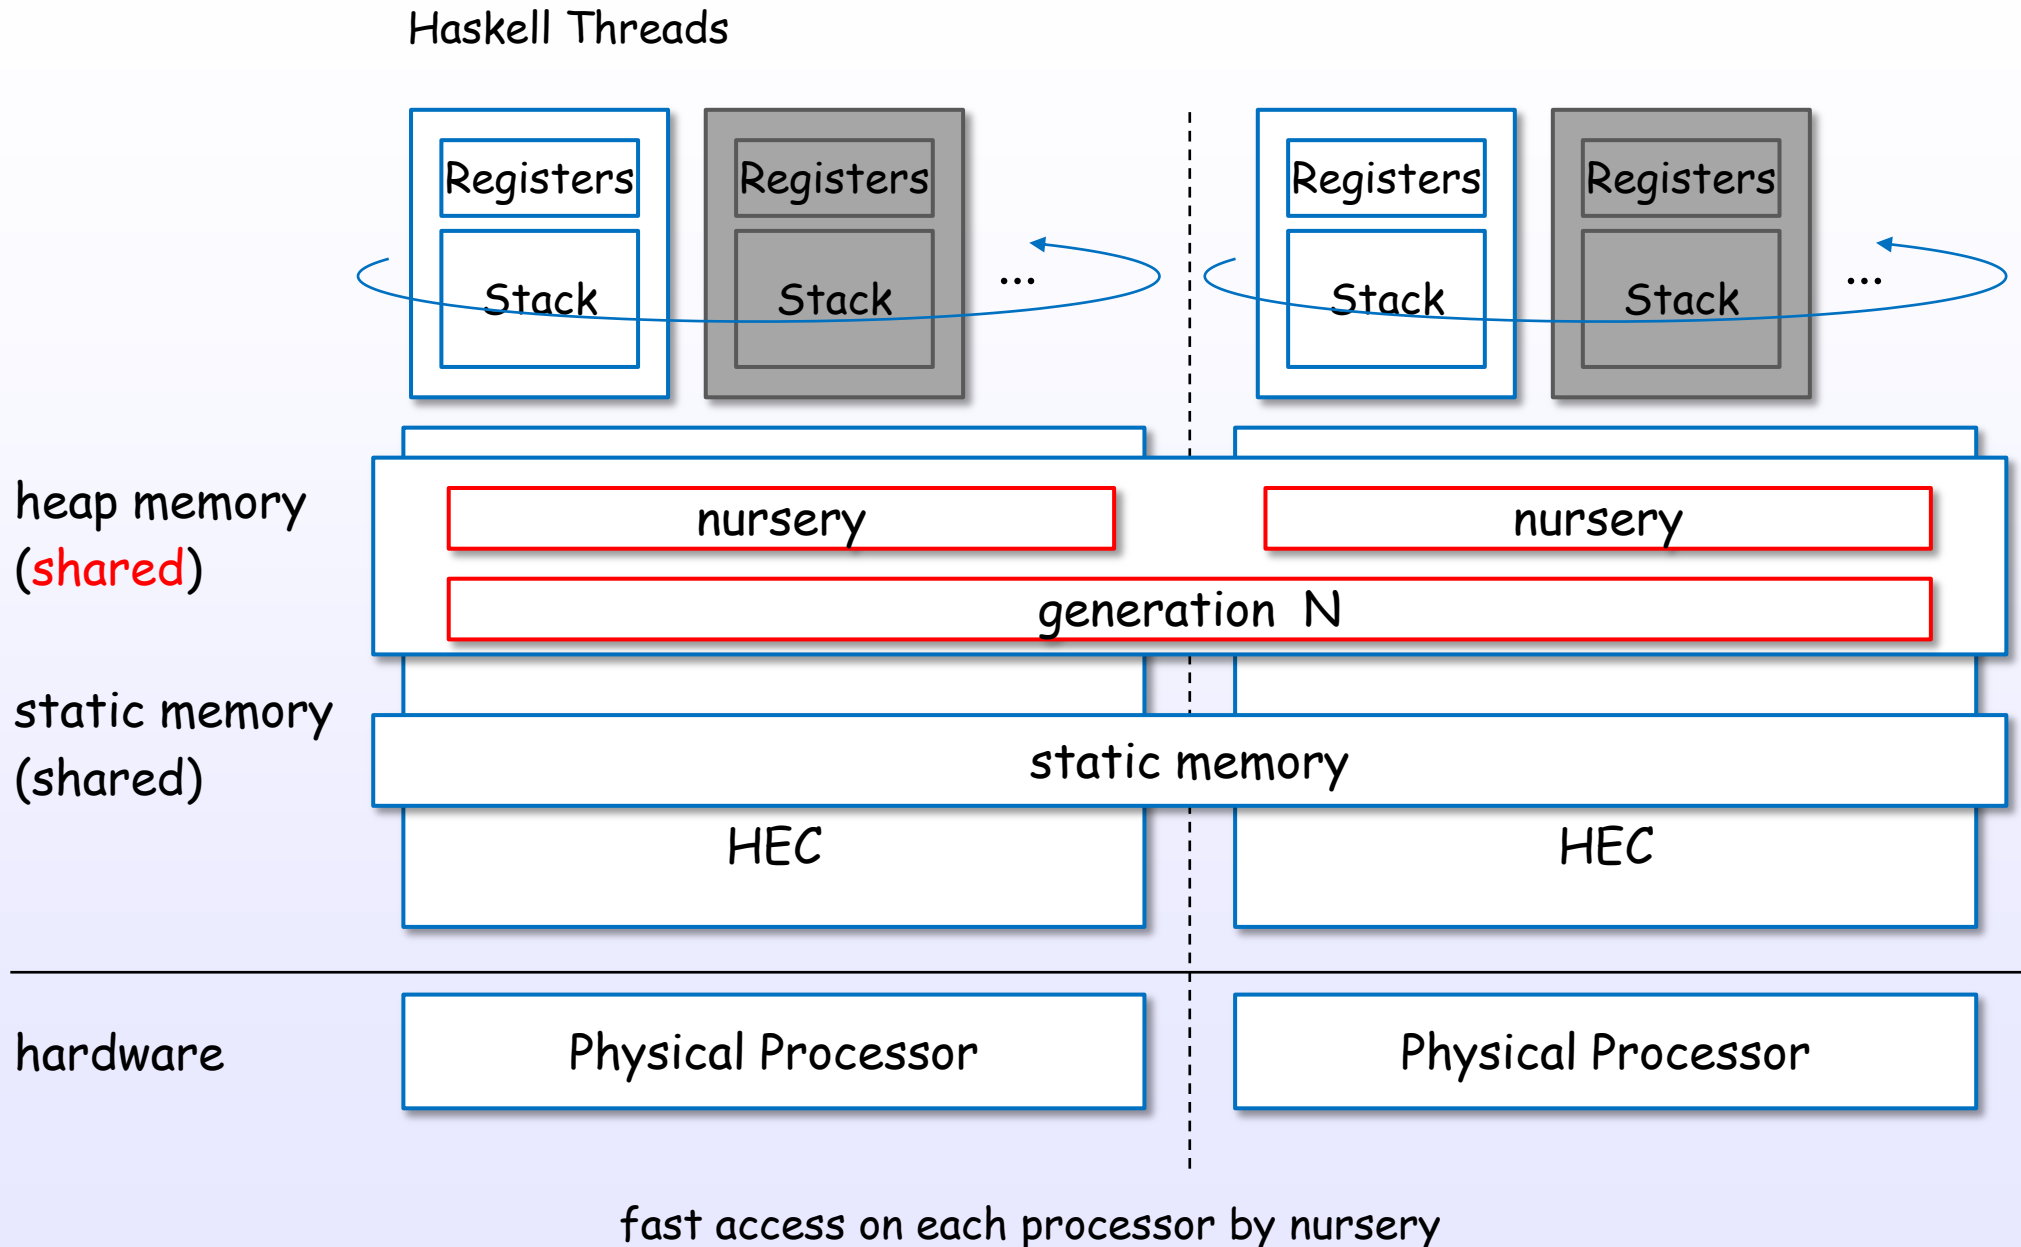
Heap and Threads

Threads and a shared heap

Haskell Threads

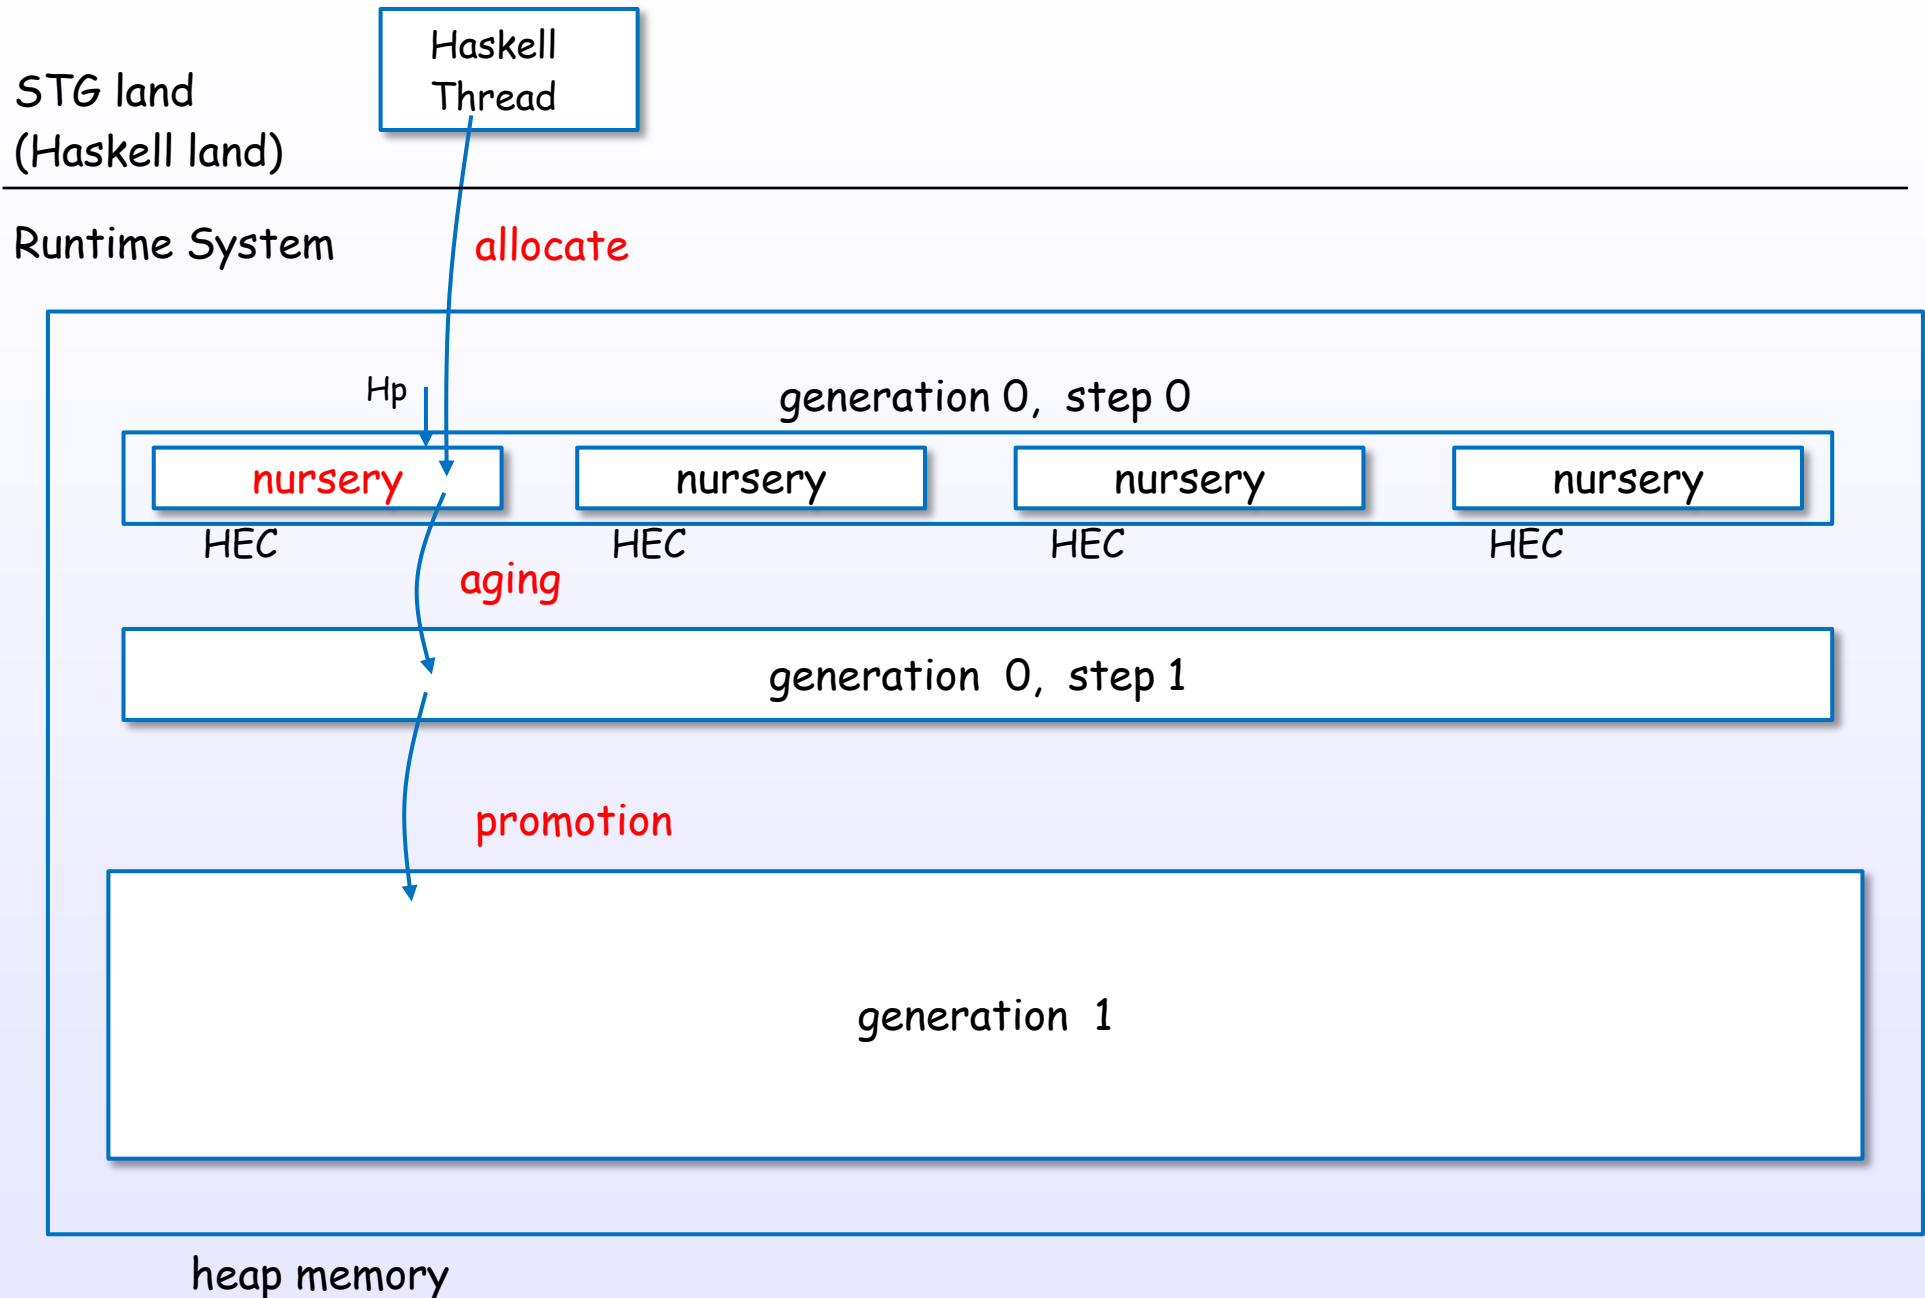


Local allocation area (nursery)



Threads and GC

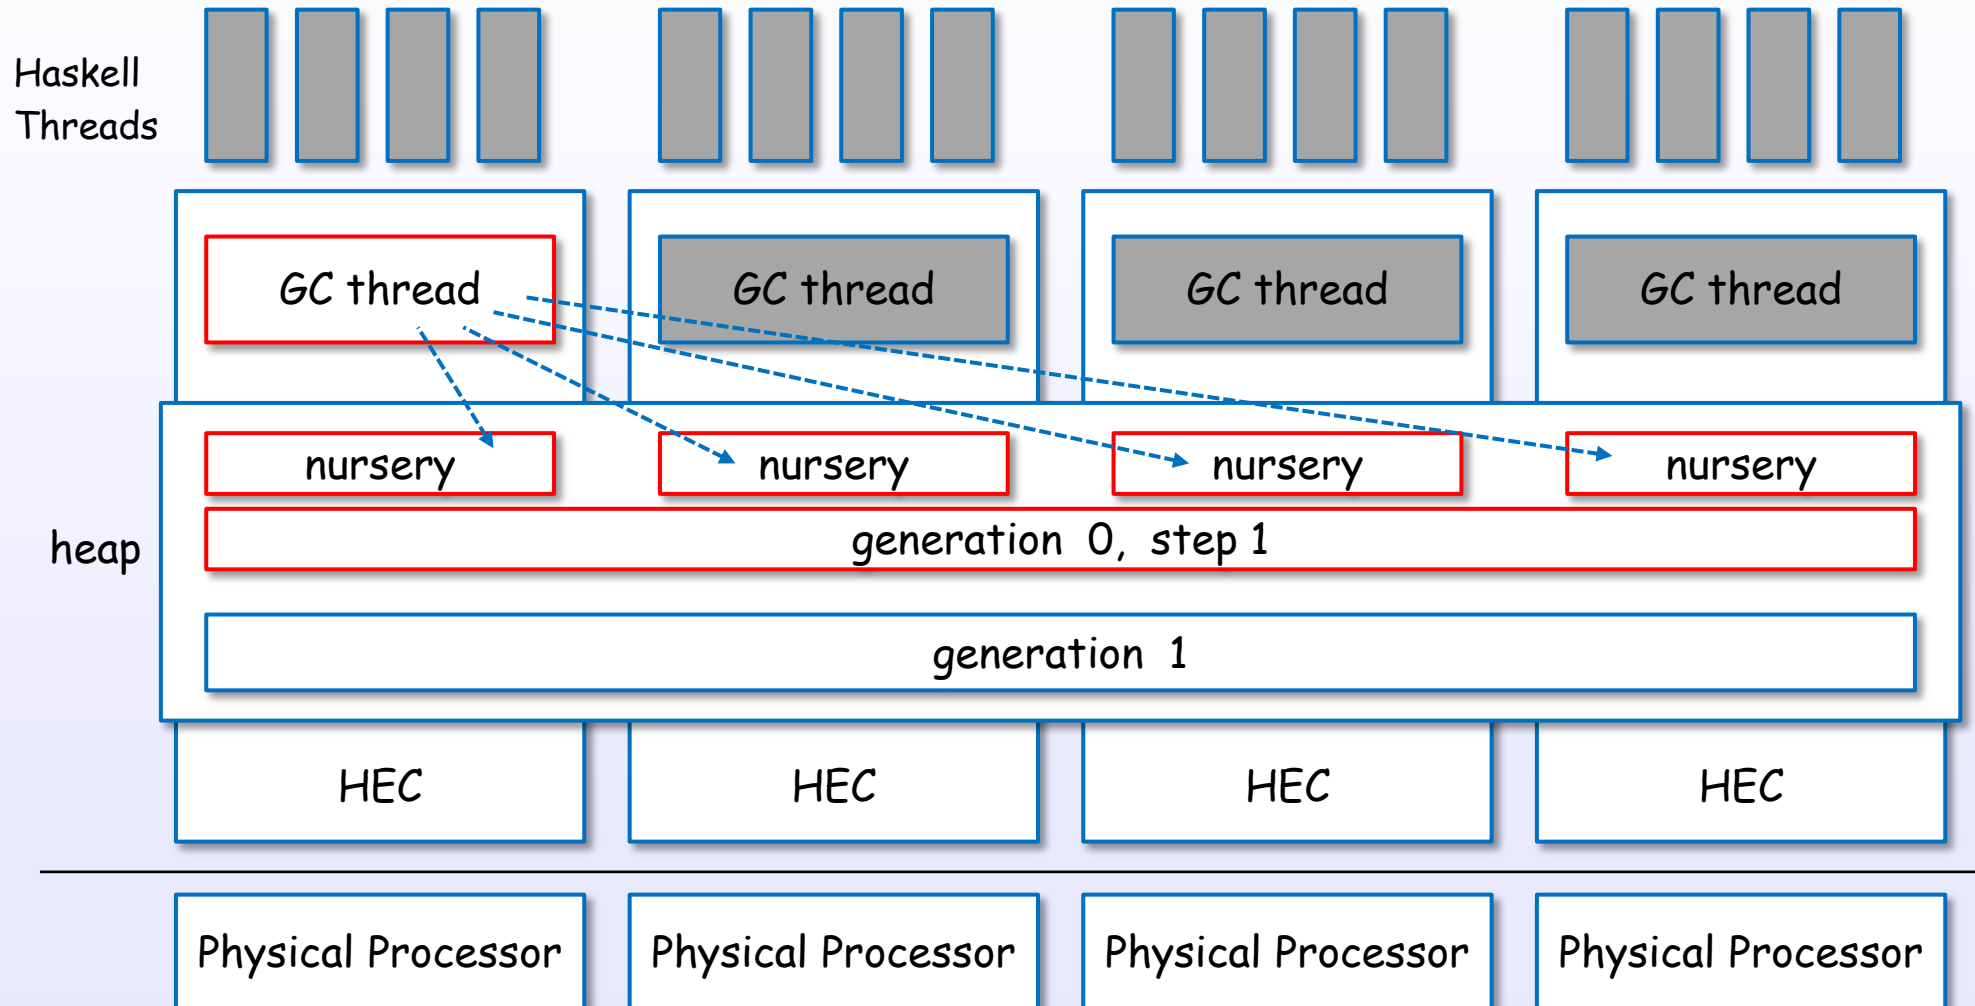
GC, nursery, generation, aging, promotion



Threads and minor GC

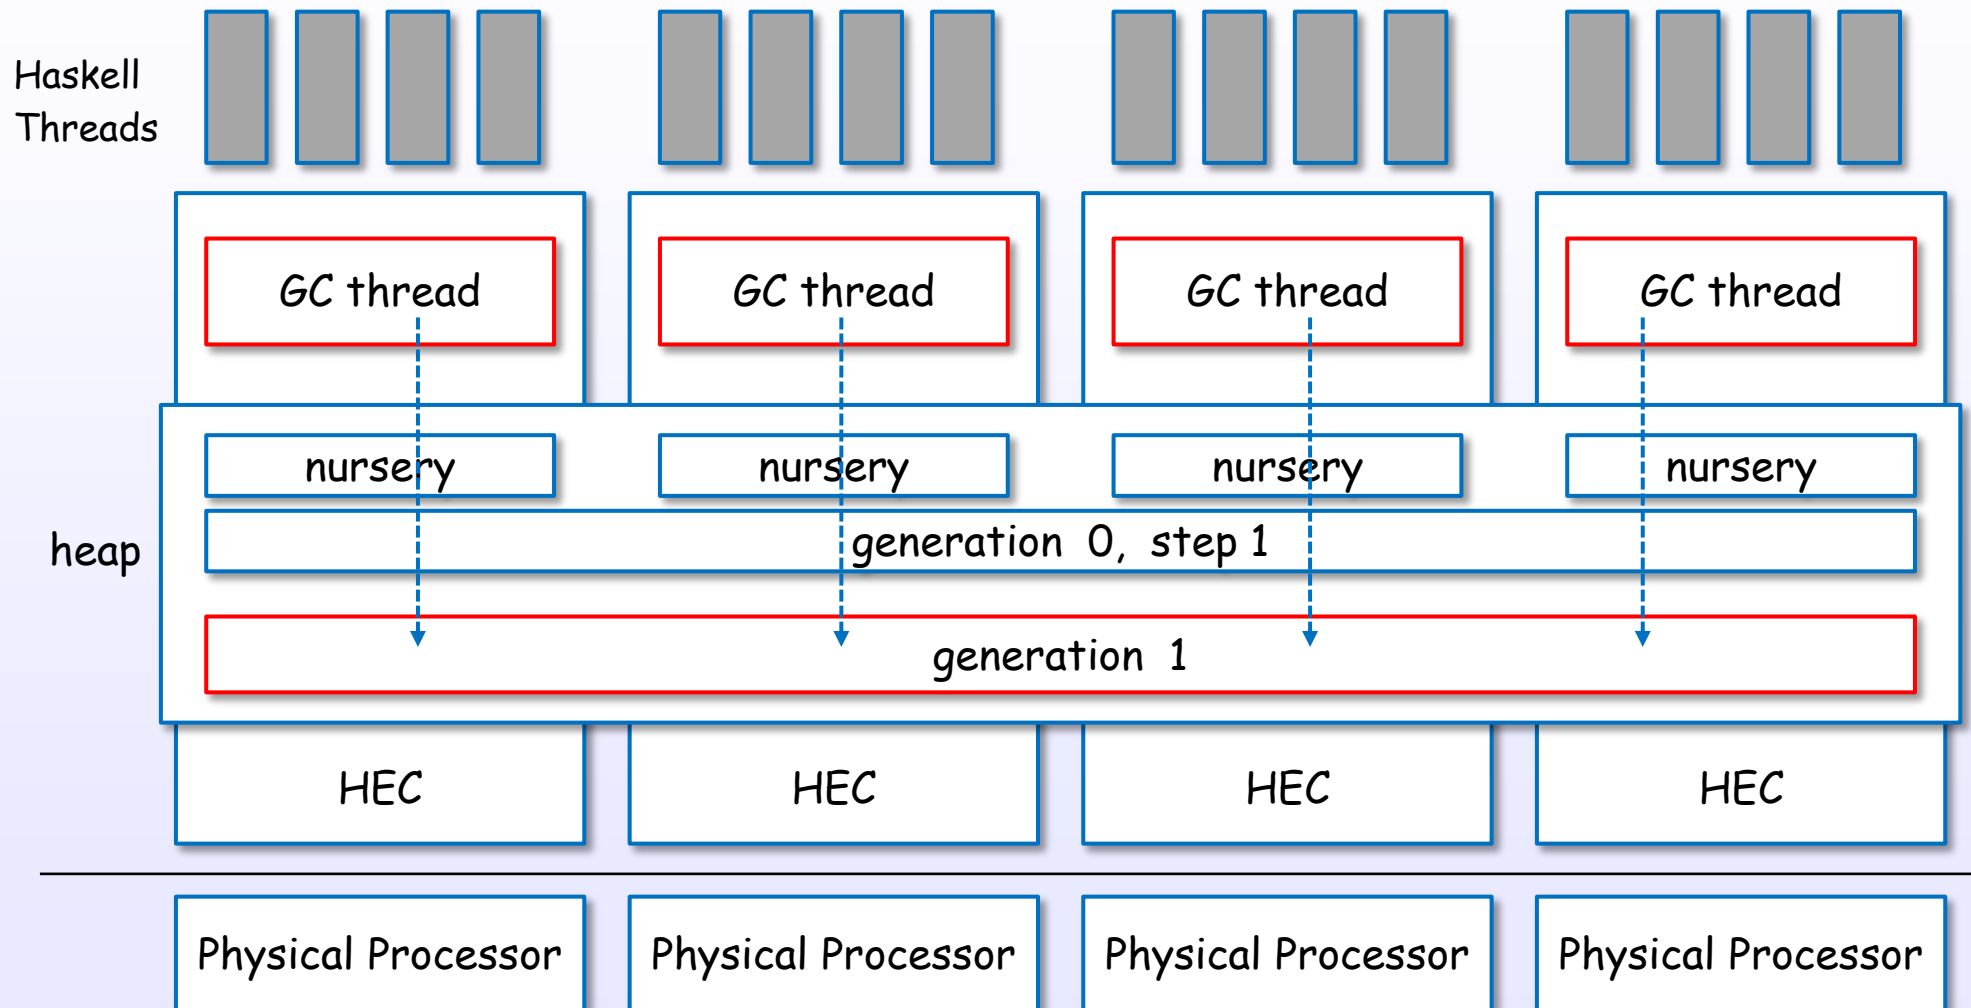
sequential GC for young generation (minor GC)

"stop-the-world" GC



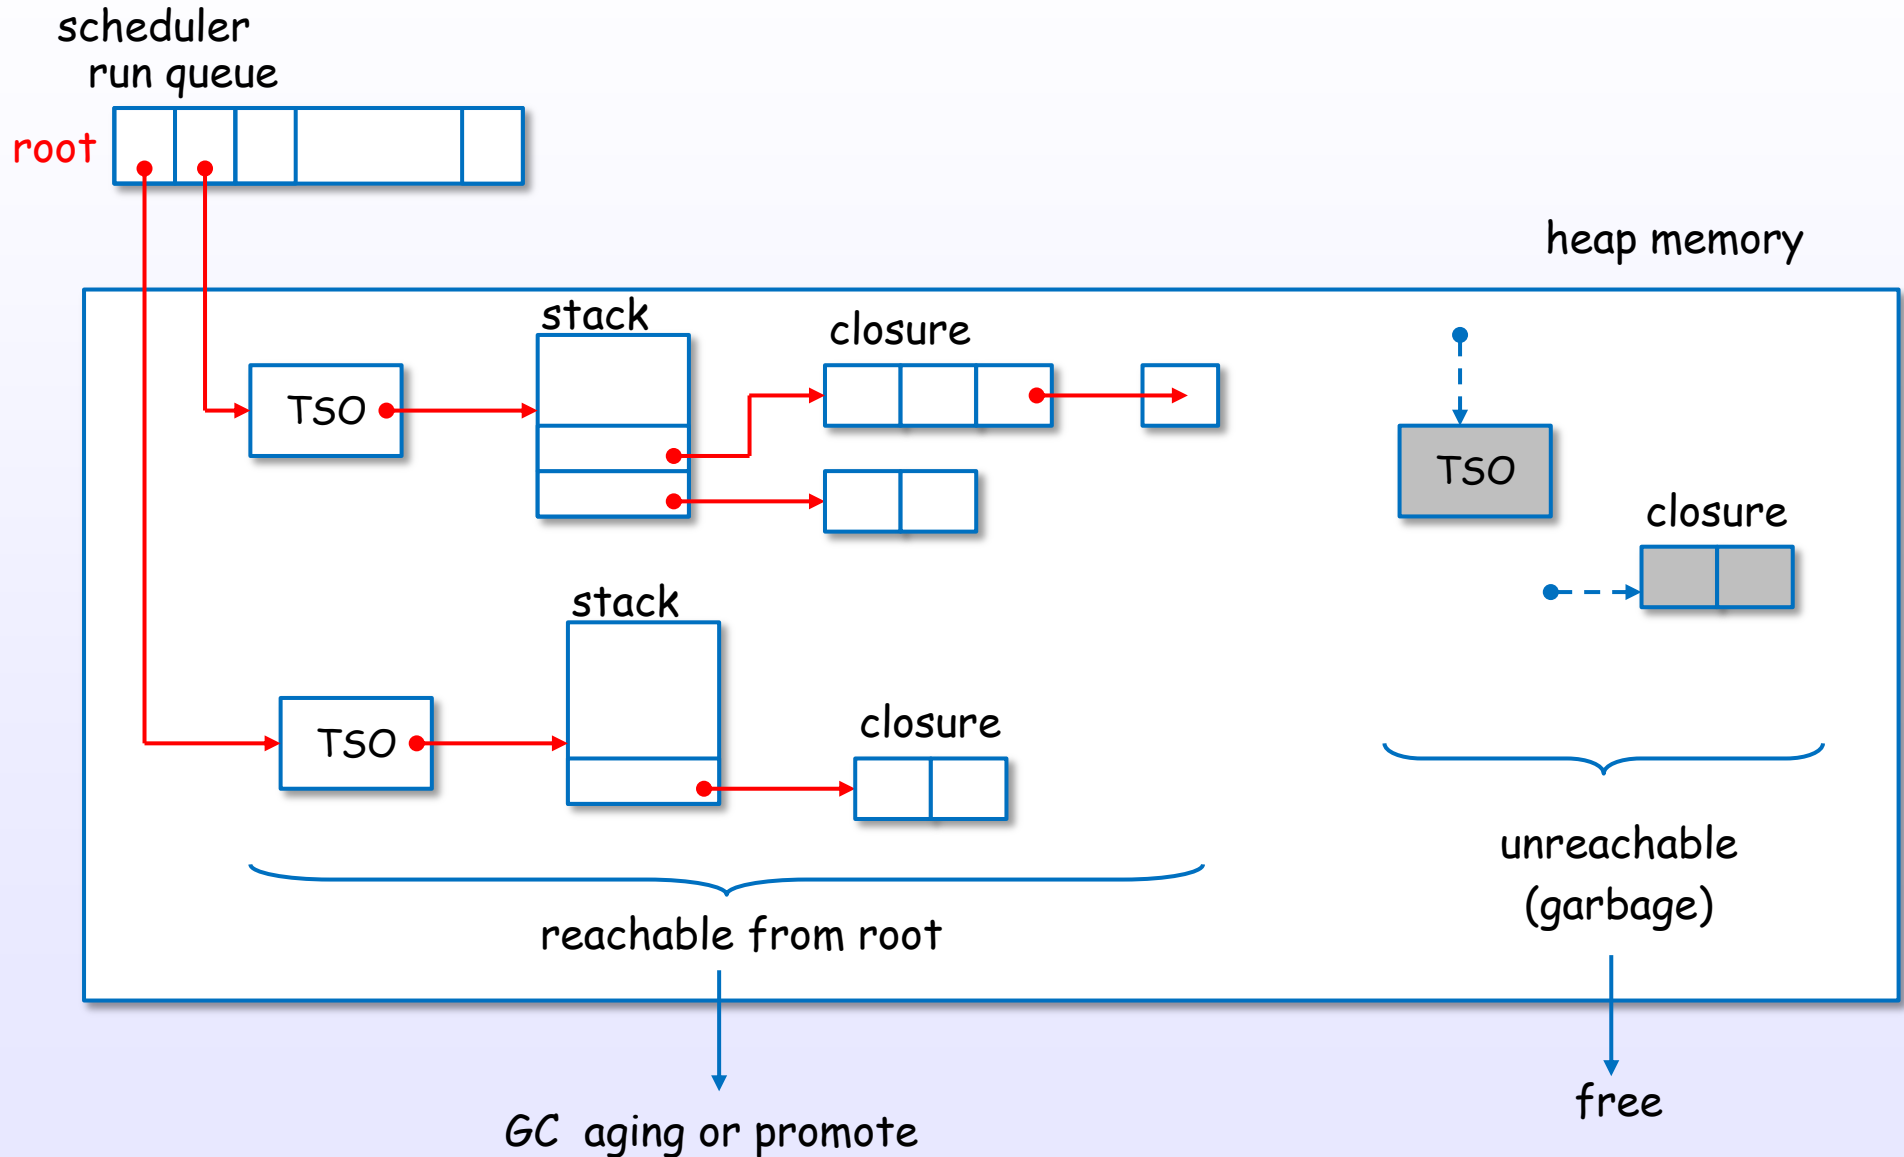
Threads and major GC

parallel GC for oldest generation (major GC)
"stop-the-world" GC



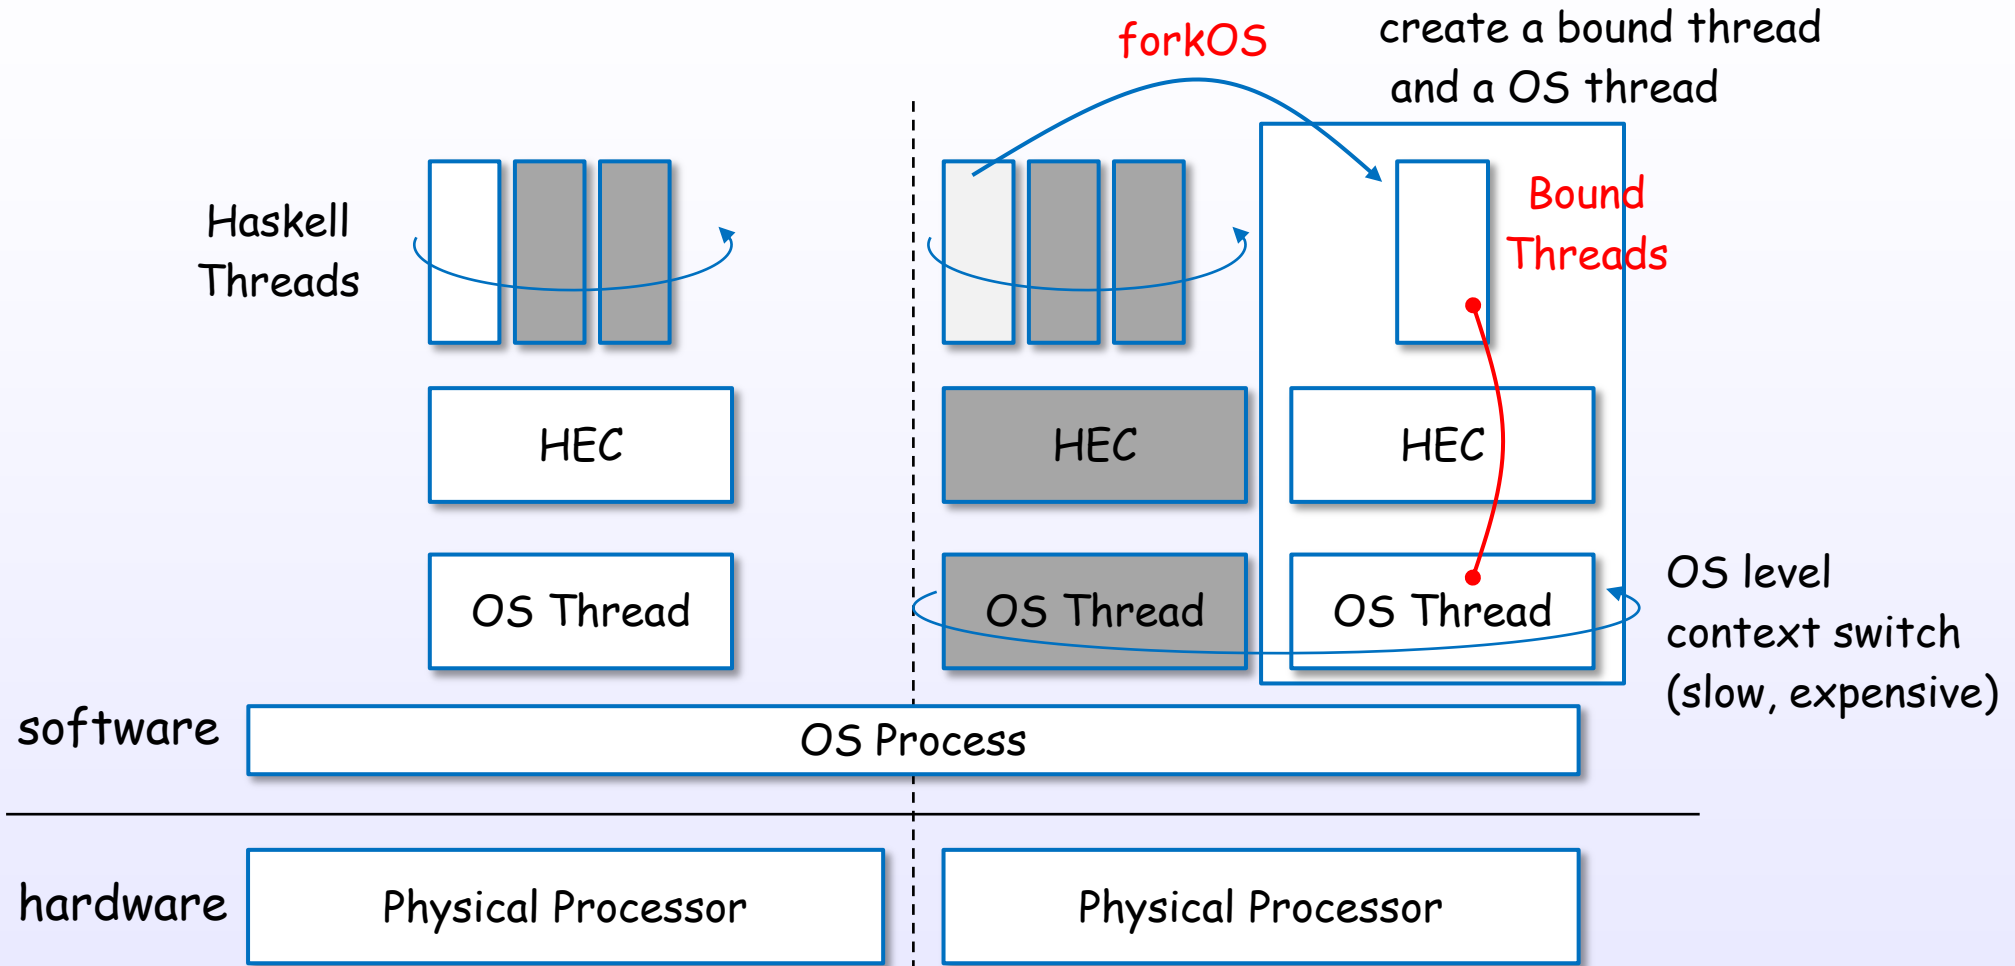
GC discover live objects from the root

Runtime System




Bound thread


Create a bound thread by forkOS




forkIO, forkOn, forkOS



create a haskell thread




create a haskell thread
on specified HEC



create a haskell thread
and a OS thread

Spark

Spark layer

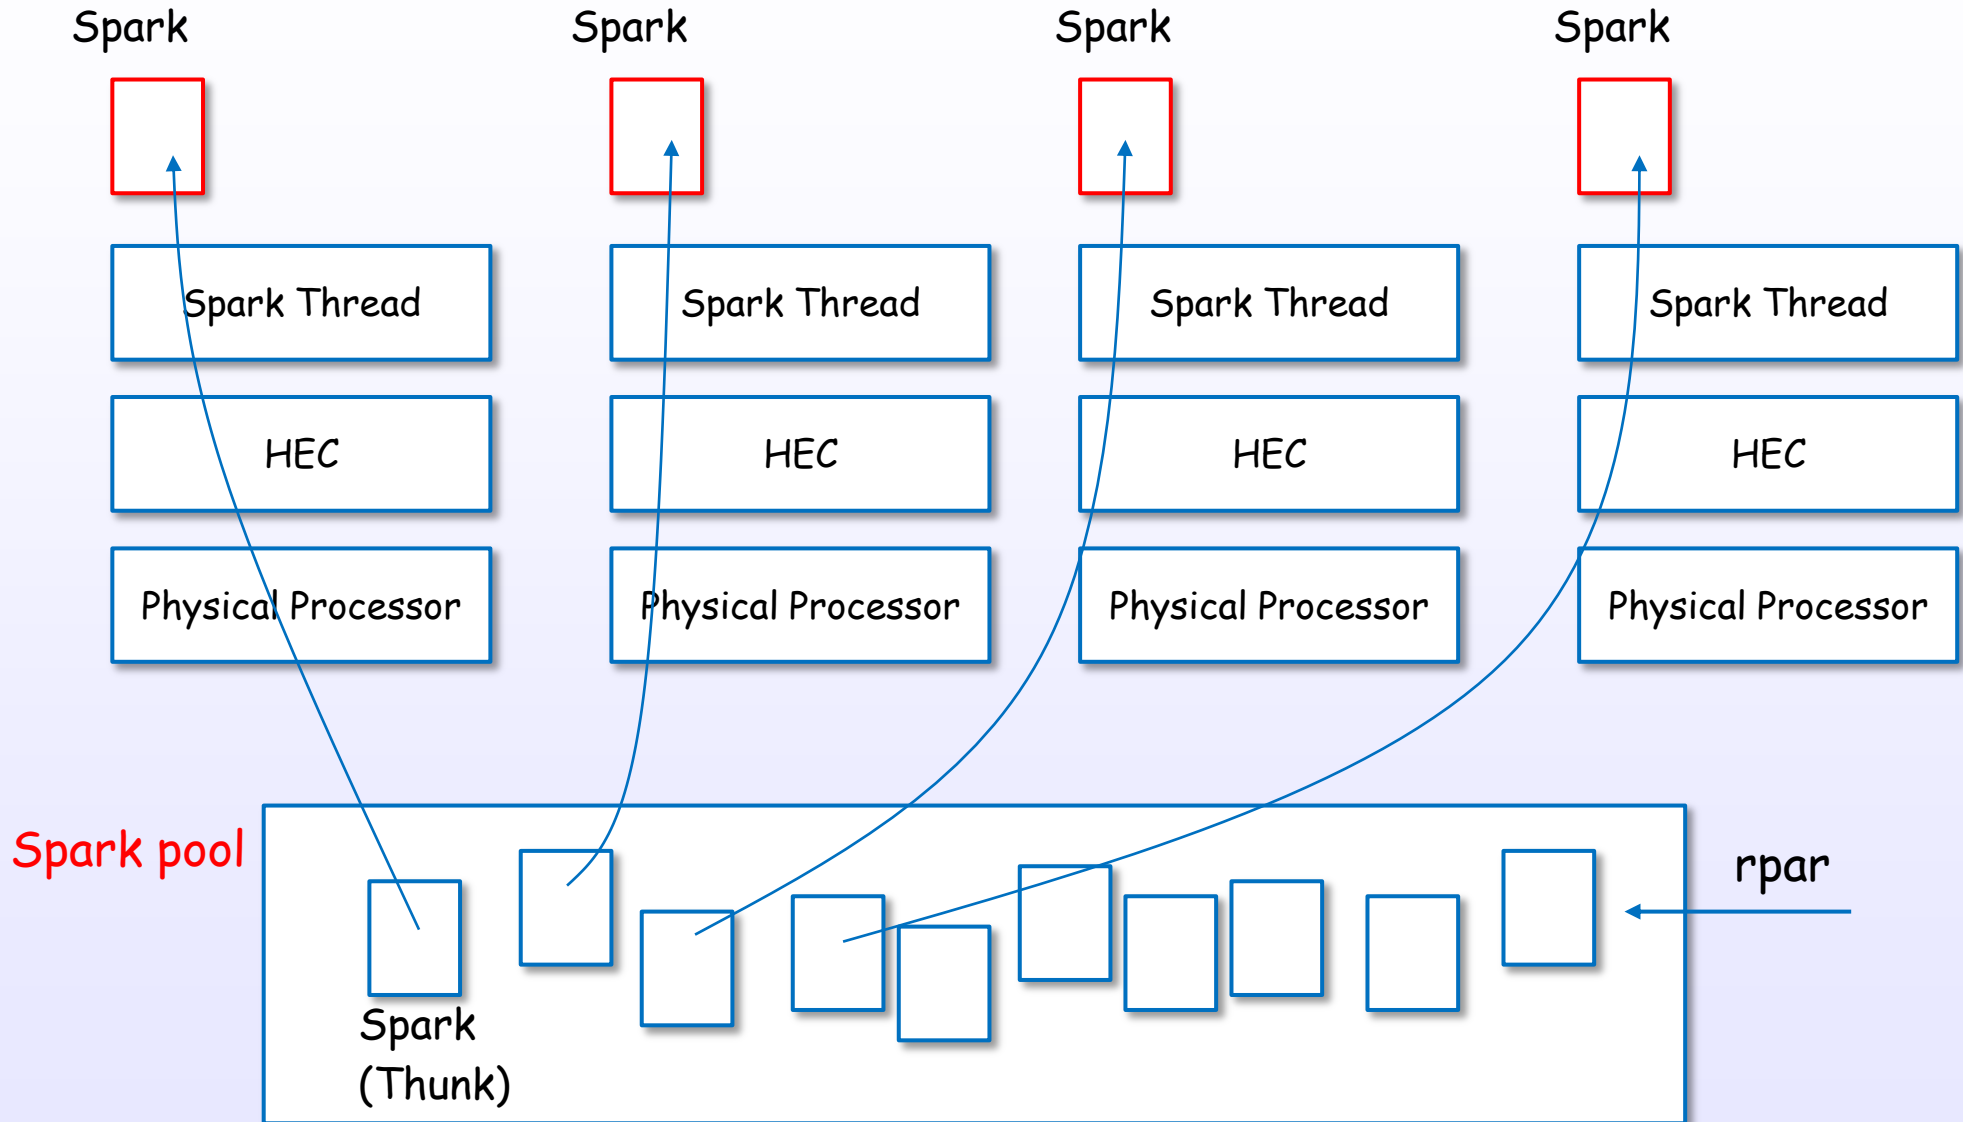


Spark Threads are generated on idle HECs.

References : [C17], [19], [S17], [S26], [S27], [S33], [S12]

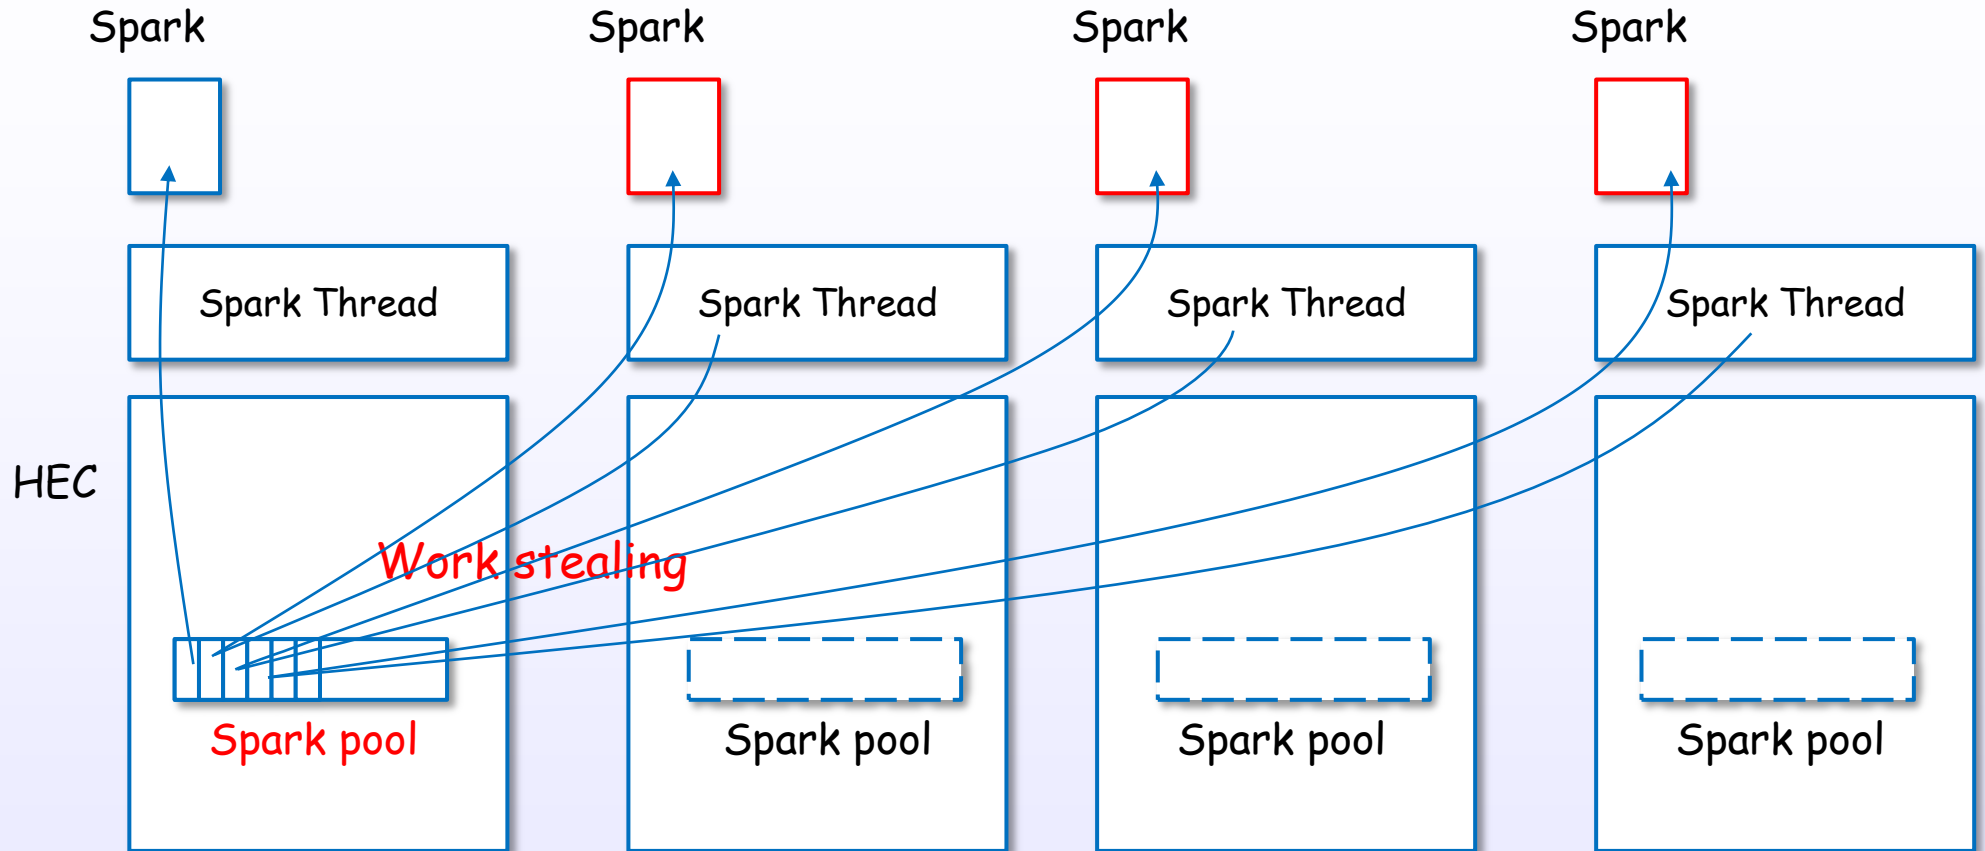
Sparks and Spark pool

logical view

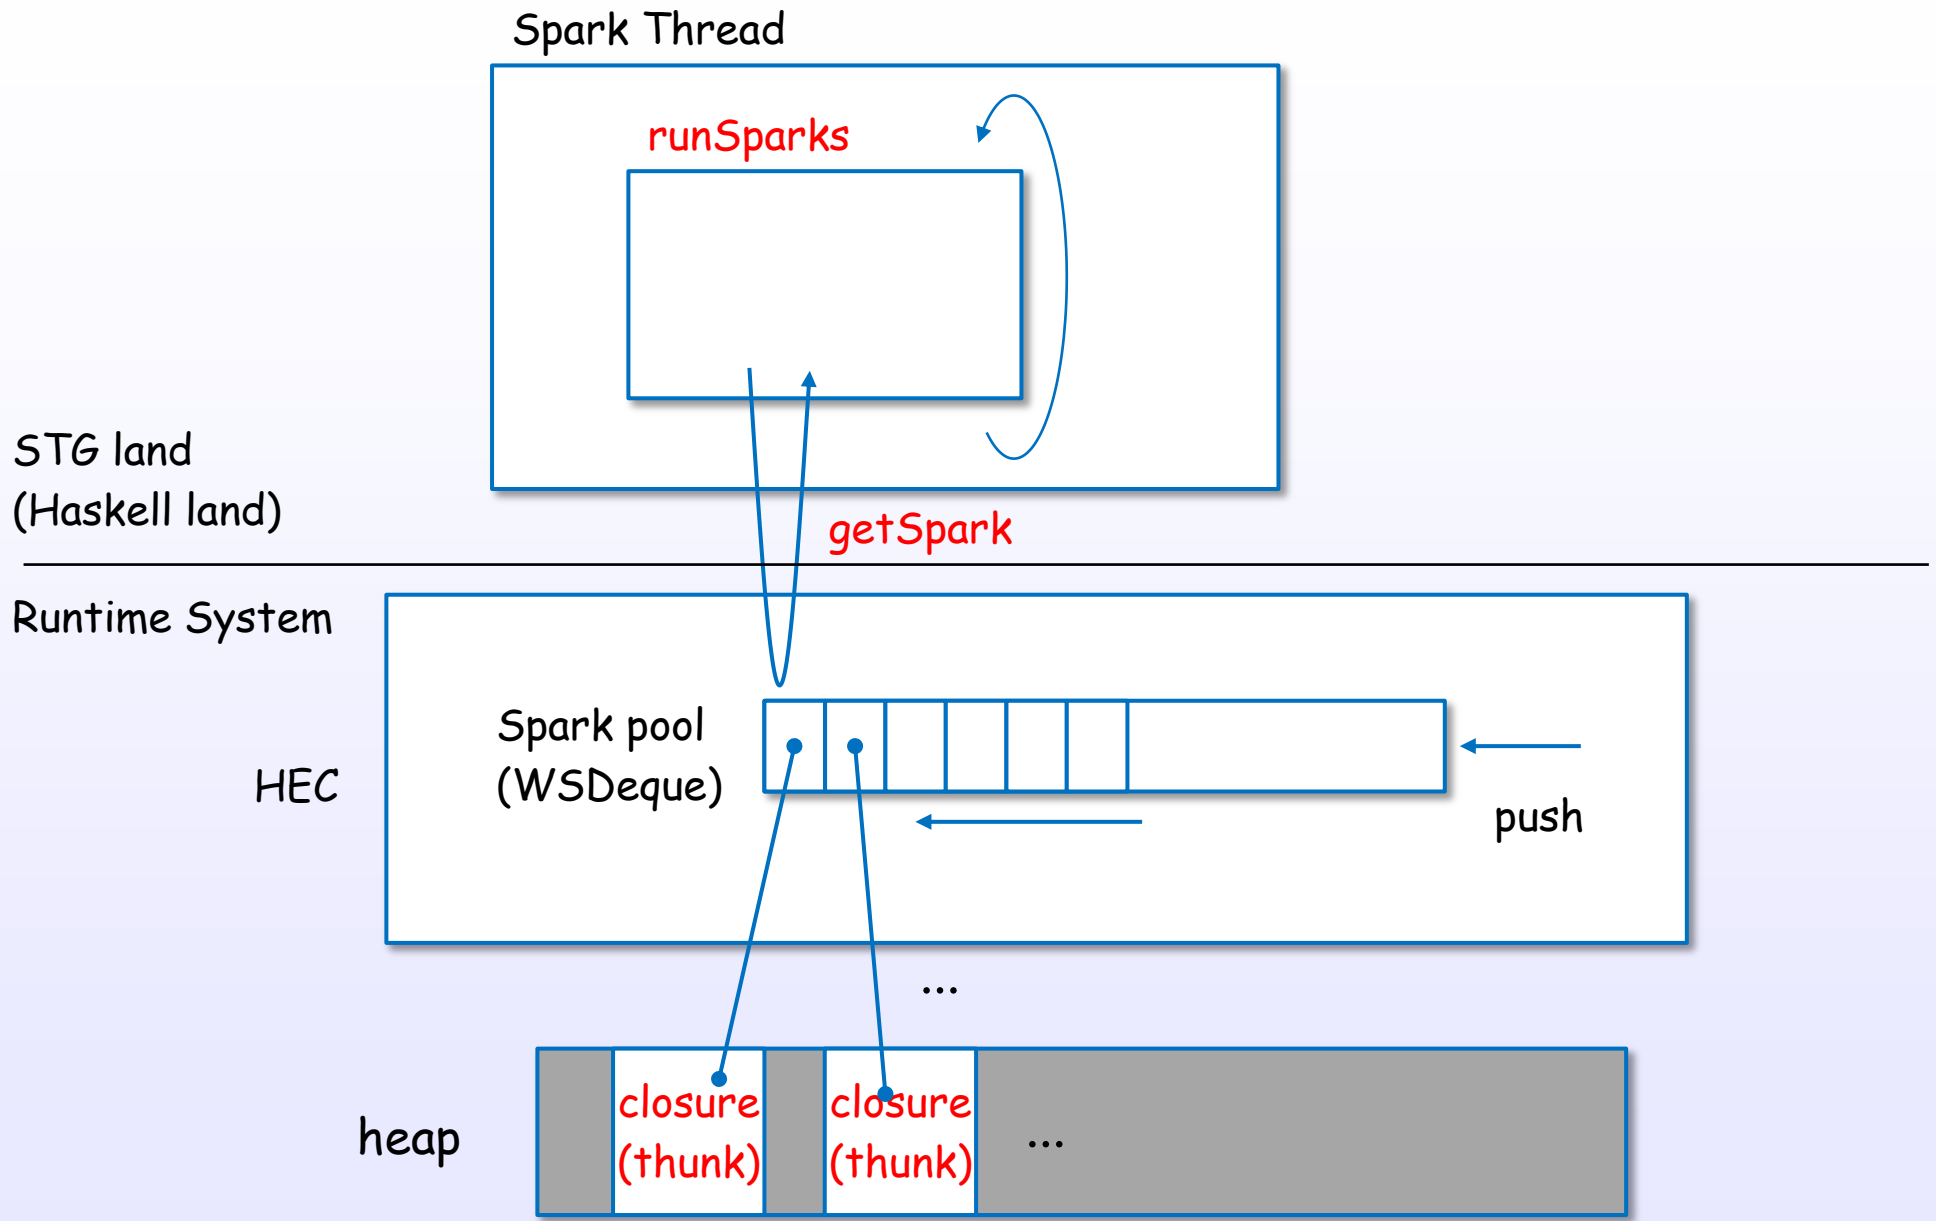


Spark pool and work stealing

physical view



Sparks and closures



(not TSO objects, but closures. therefore very lightweight)

MVar


MVar

Haskell Thread #0




putMVar

Haskell Thread #1



takeMVar



MVar

MVar and blocking


Haskell Thread



putMVar



BLOCKED
if full



MVar


Haskell Thread



takeMVar



BLOCKED
if empty

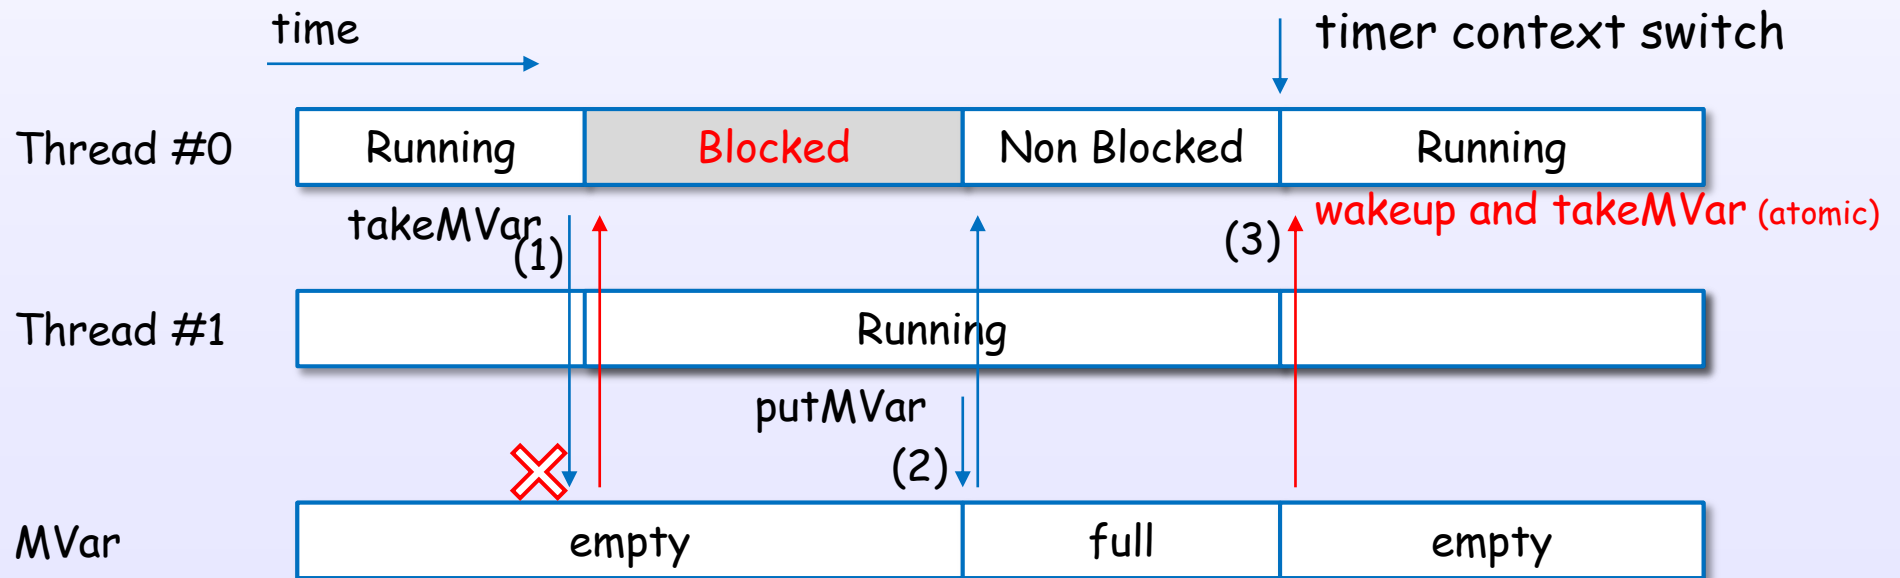



MVar

MVar example

Haskell Thread #0


Haskell Thread #1




MVar object view

User view


MVar



logical MVar object



physical MVar object



newEmptyMVar

Haskell Threads


```
newEmptyMVar  
newMVar#
```

(1) **call** the Runtime primitive

Runtime System

```
stg_newMVarzh  
  ALLOC_PRIM_  
  SET_HDR  
  StgMVar_head  
  StgMVar_tail  
  StgMVar_value
```

(2) **create a MVar** object



(3) **link** each fields

References : [16], [18], [19], [S31], [S12]

takeMVar (empty case)


Haskell Threads

```
takeMVar  
takeMVar#
```

Runtime System

```
stg_takeMVarzh  
  create StgMVarTSOQueue ... (1)  
  append      ... (2)  
  StgReturn   ... (3)
```

(3) return to the scheduler



MVar object

(1) create

StgMVarTSOQueue

(2) append

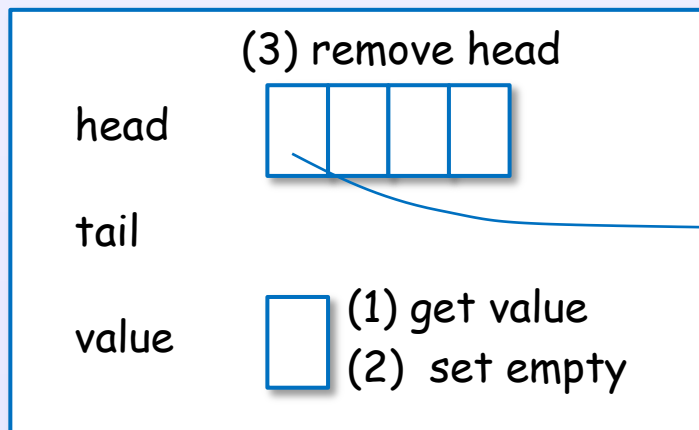
takeMVar (full case)

Haskell Threads

```
takeMVar  
takeMVar#
```

Runtime System

```
stg_takeMVarzh  
(1) get value  
(2) set empty  
(3) remove head  
(4) tryWakeupThread
```



MVar object

scheduler

run queue

fairness round robin



append

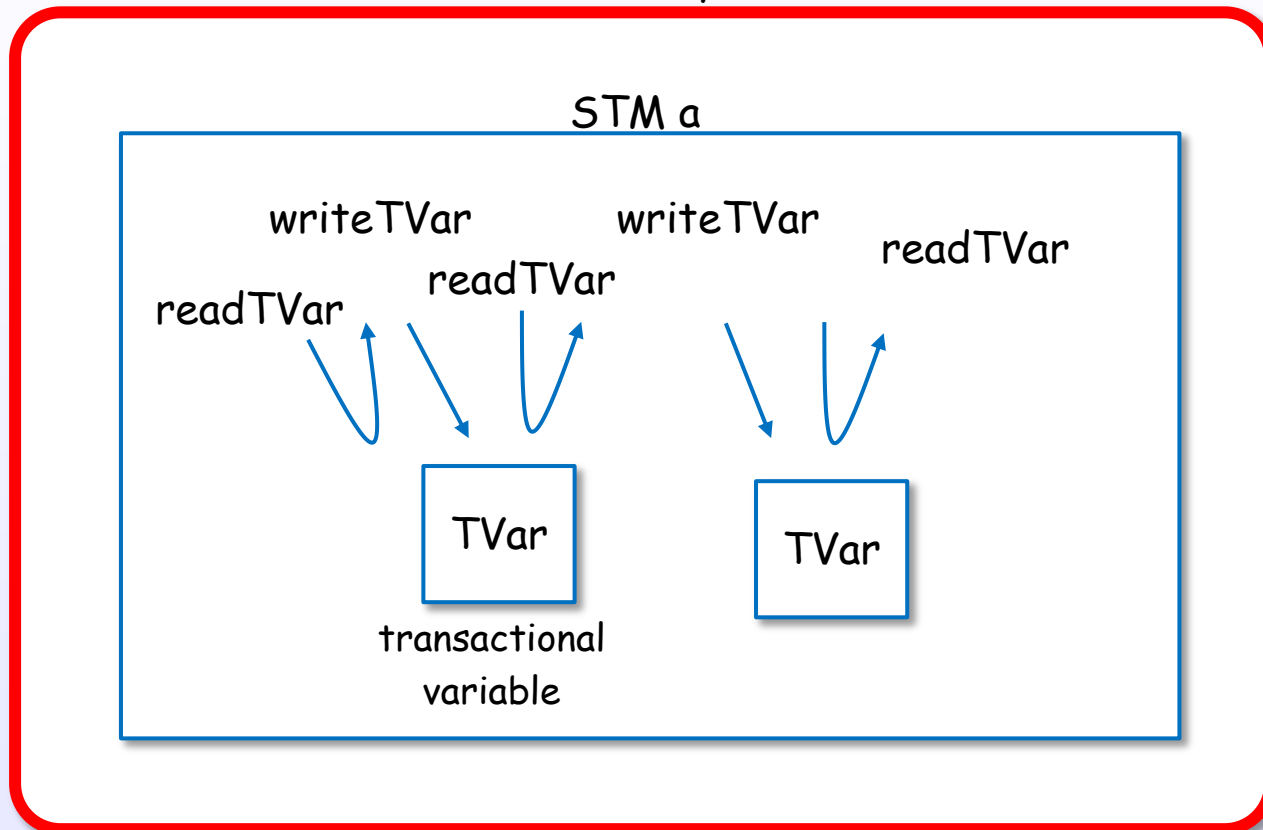
(4) wakeup the thread
blocked by putMVar

Software transactional memory

Create a atomic block by atomically

atomically :: STM a -> IO a

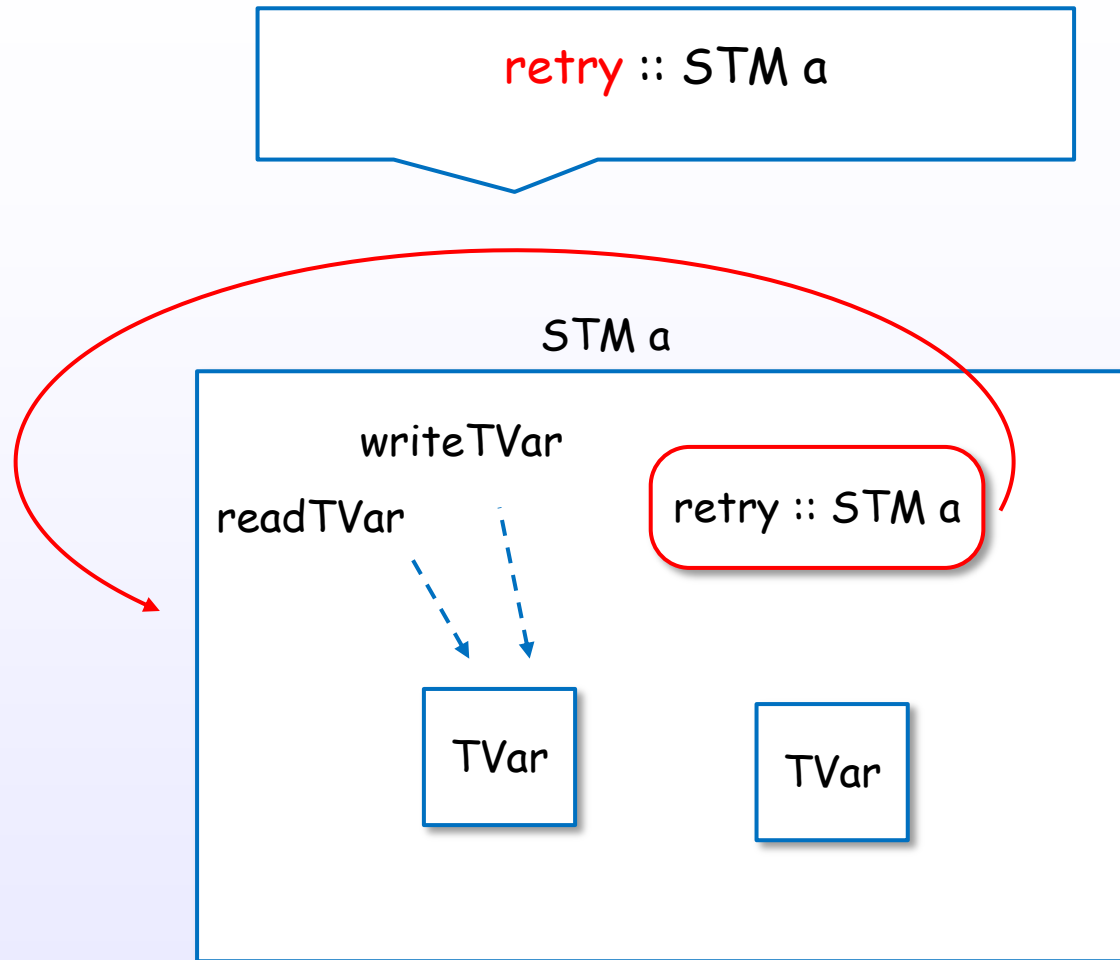
atomically



Create and evaluate a "**atomic block**"

Atomic block = All or Nothing

Rollback and blocking control by retry

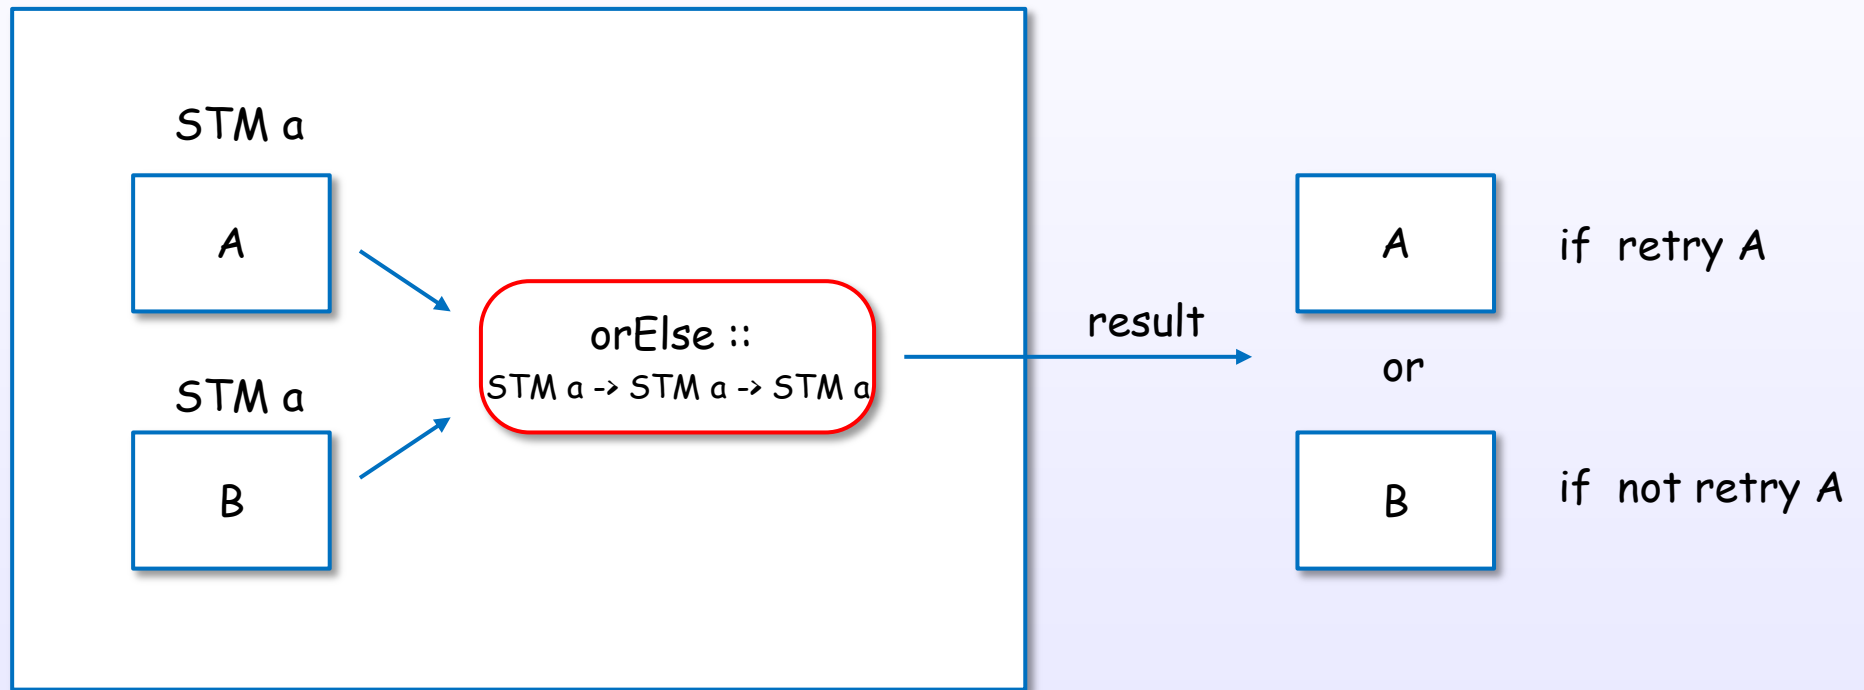


Discard, blocking and try again

Compose OR case by orElse

orElse :: STM a -> STM a -> STM a

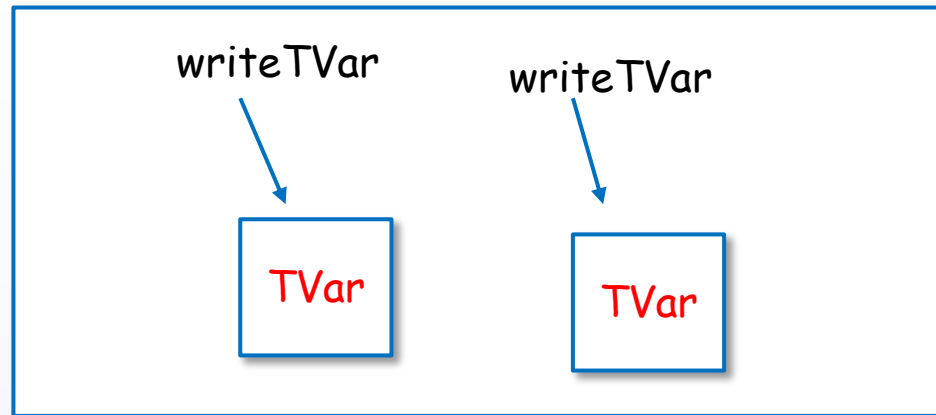
STM a



A **or** B or Nothing

STM, TVar example (normal case)

STM a




atomic block

time →

Thread #0



TVar #A




TVar #B

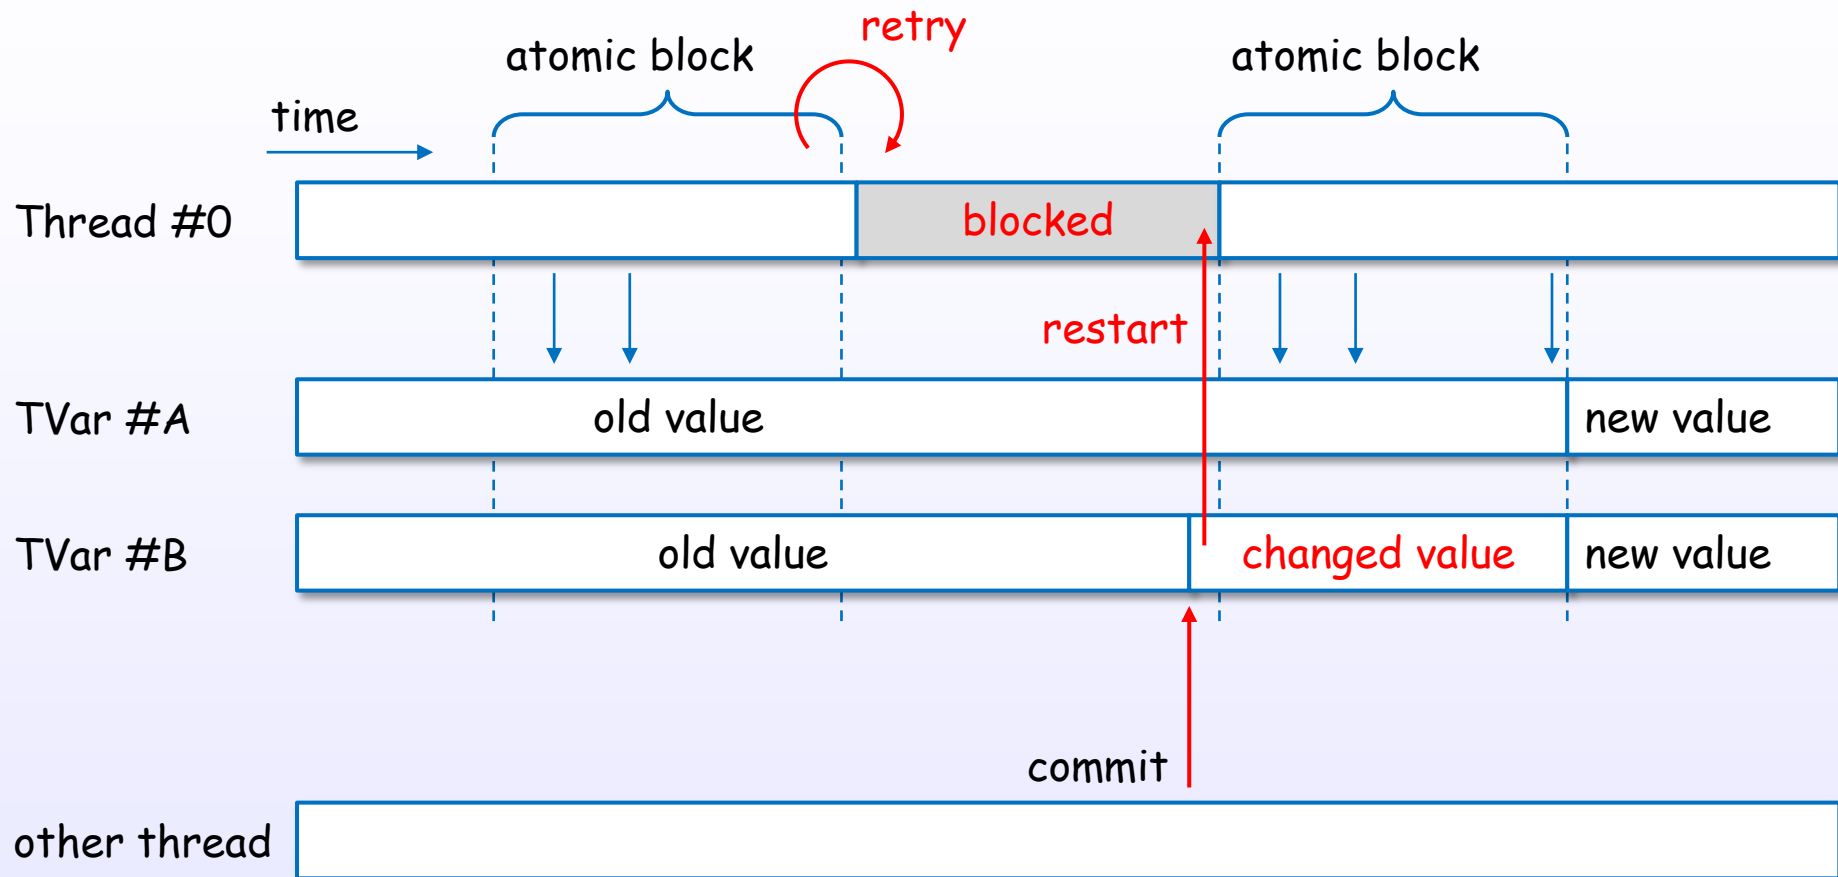


atomic update


STM, TVar example (conflict case)



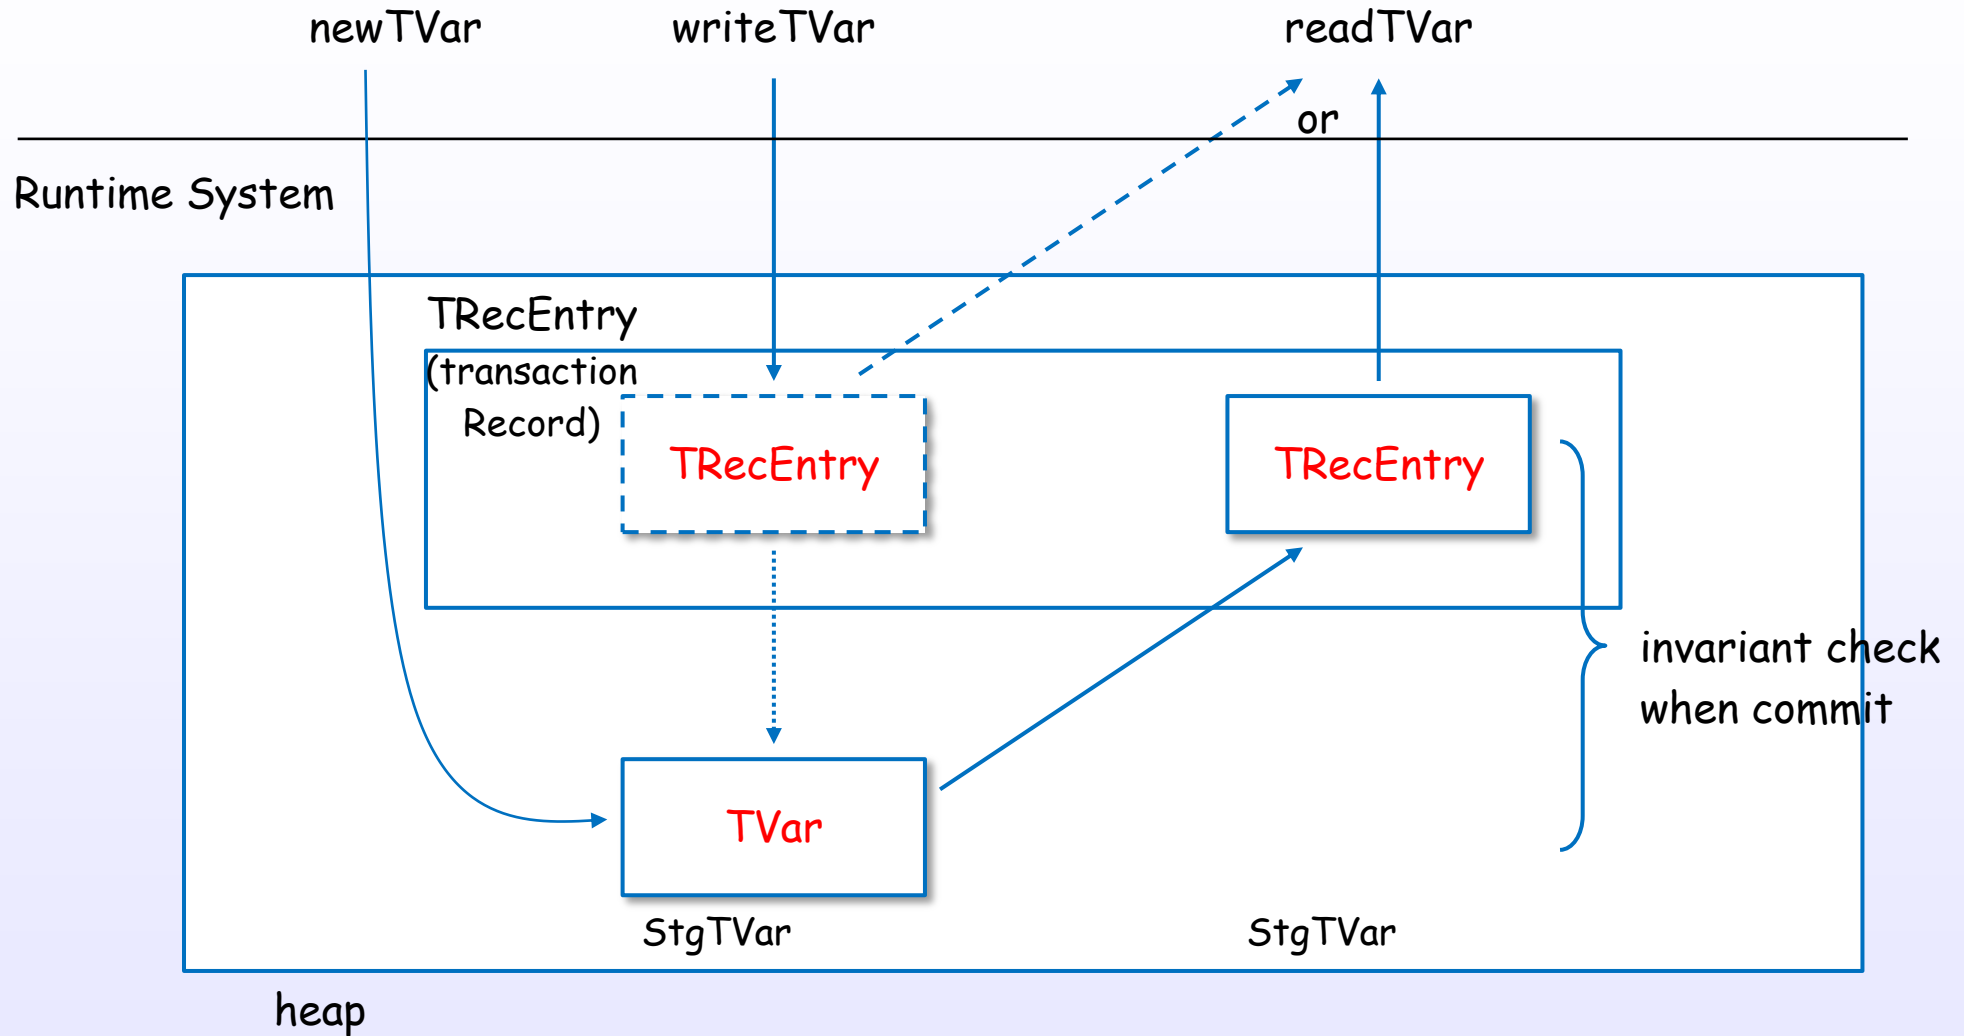
retry example



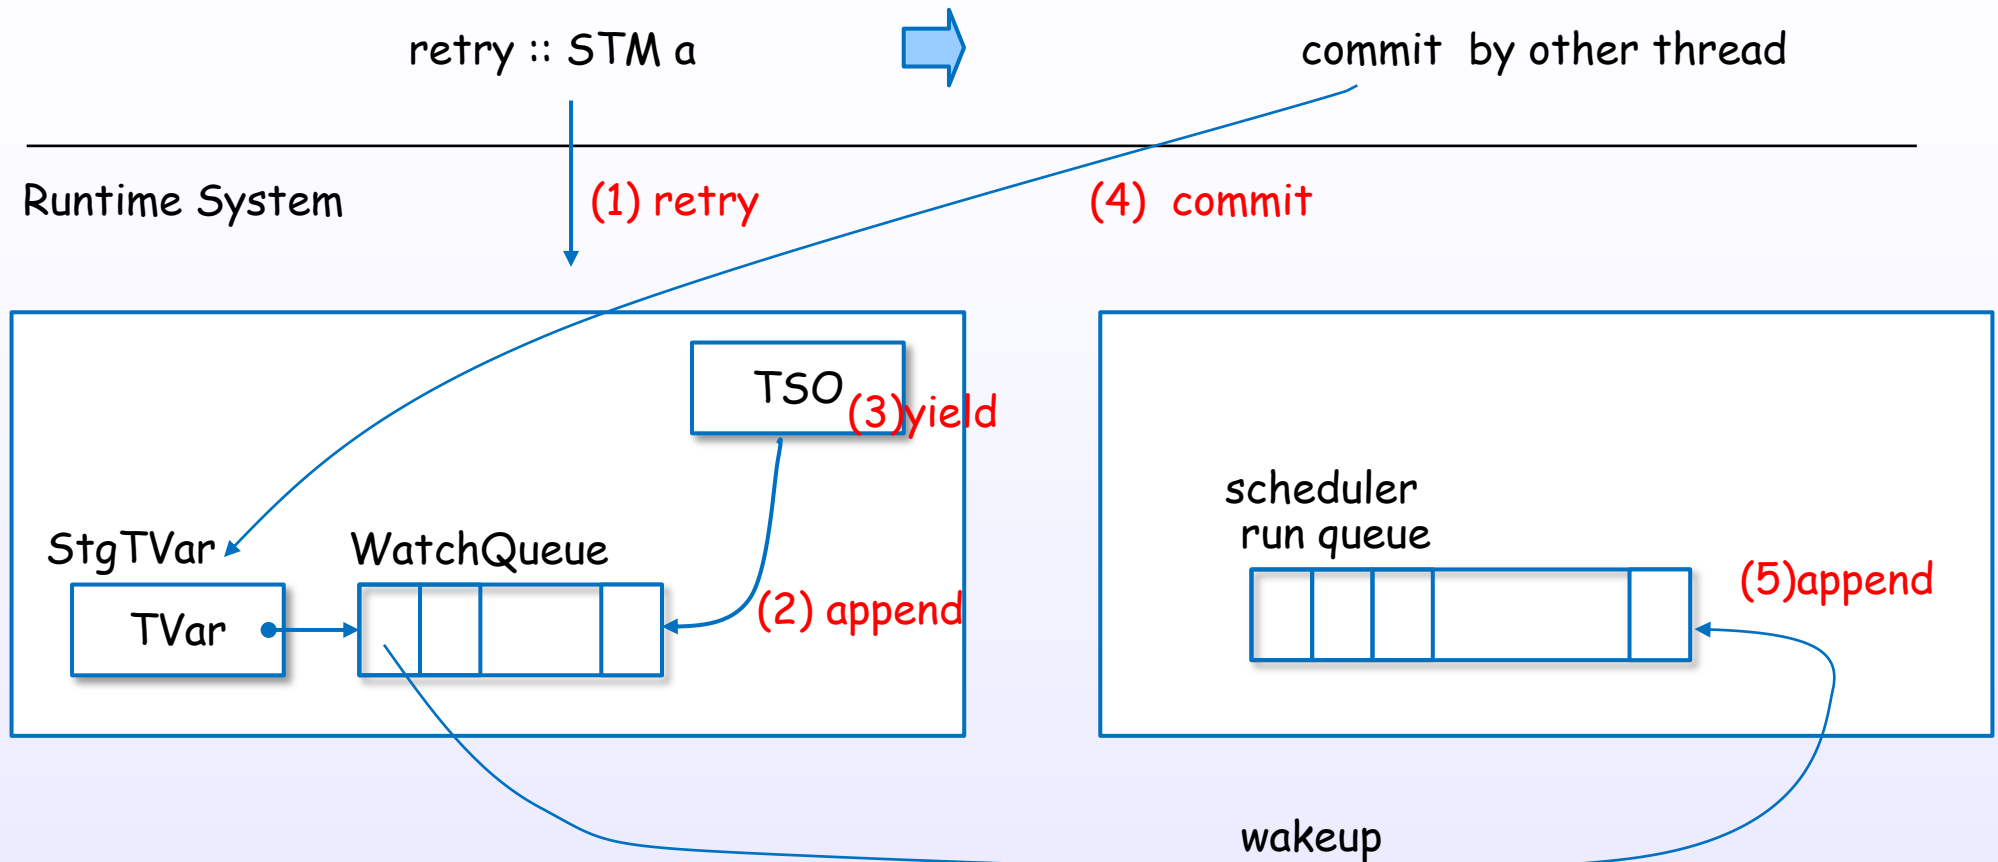
STM, TVar data structure



newTVar, writeTVar, readTVar




retry blocking and wake up




IO and FFI

IO

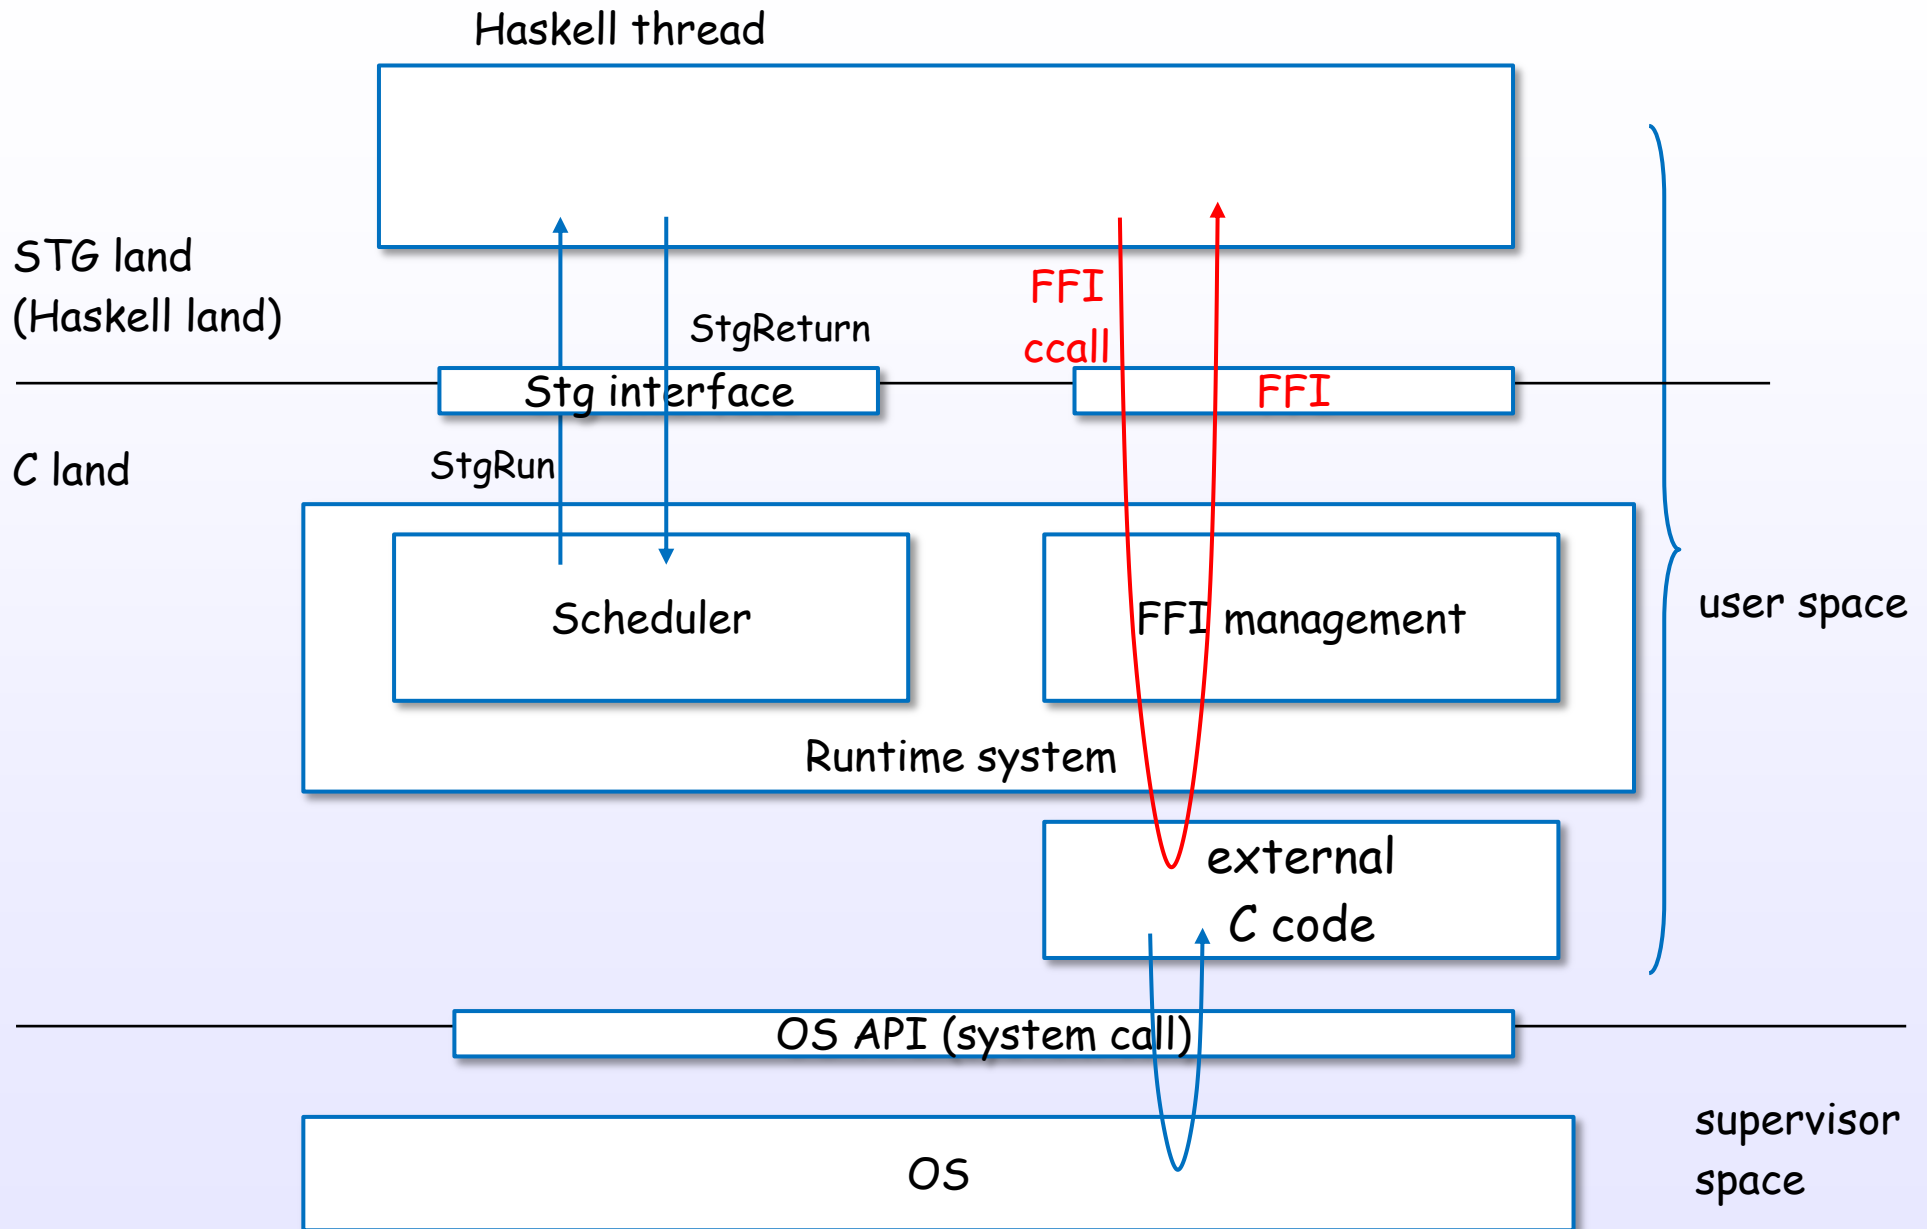
Haskell Thread



IO example: getLine




FFI (Foreign Function Interface)



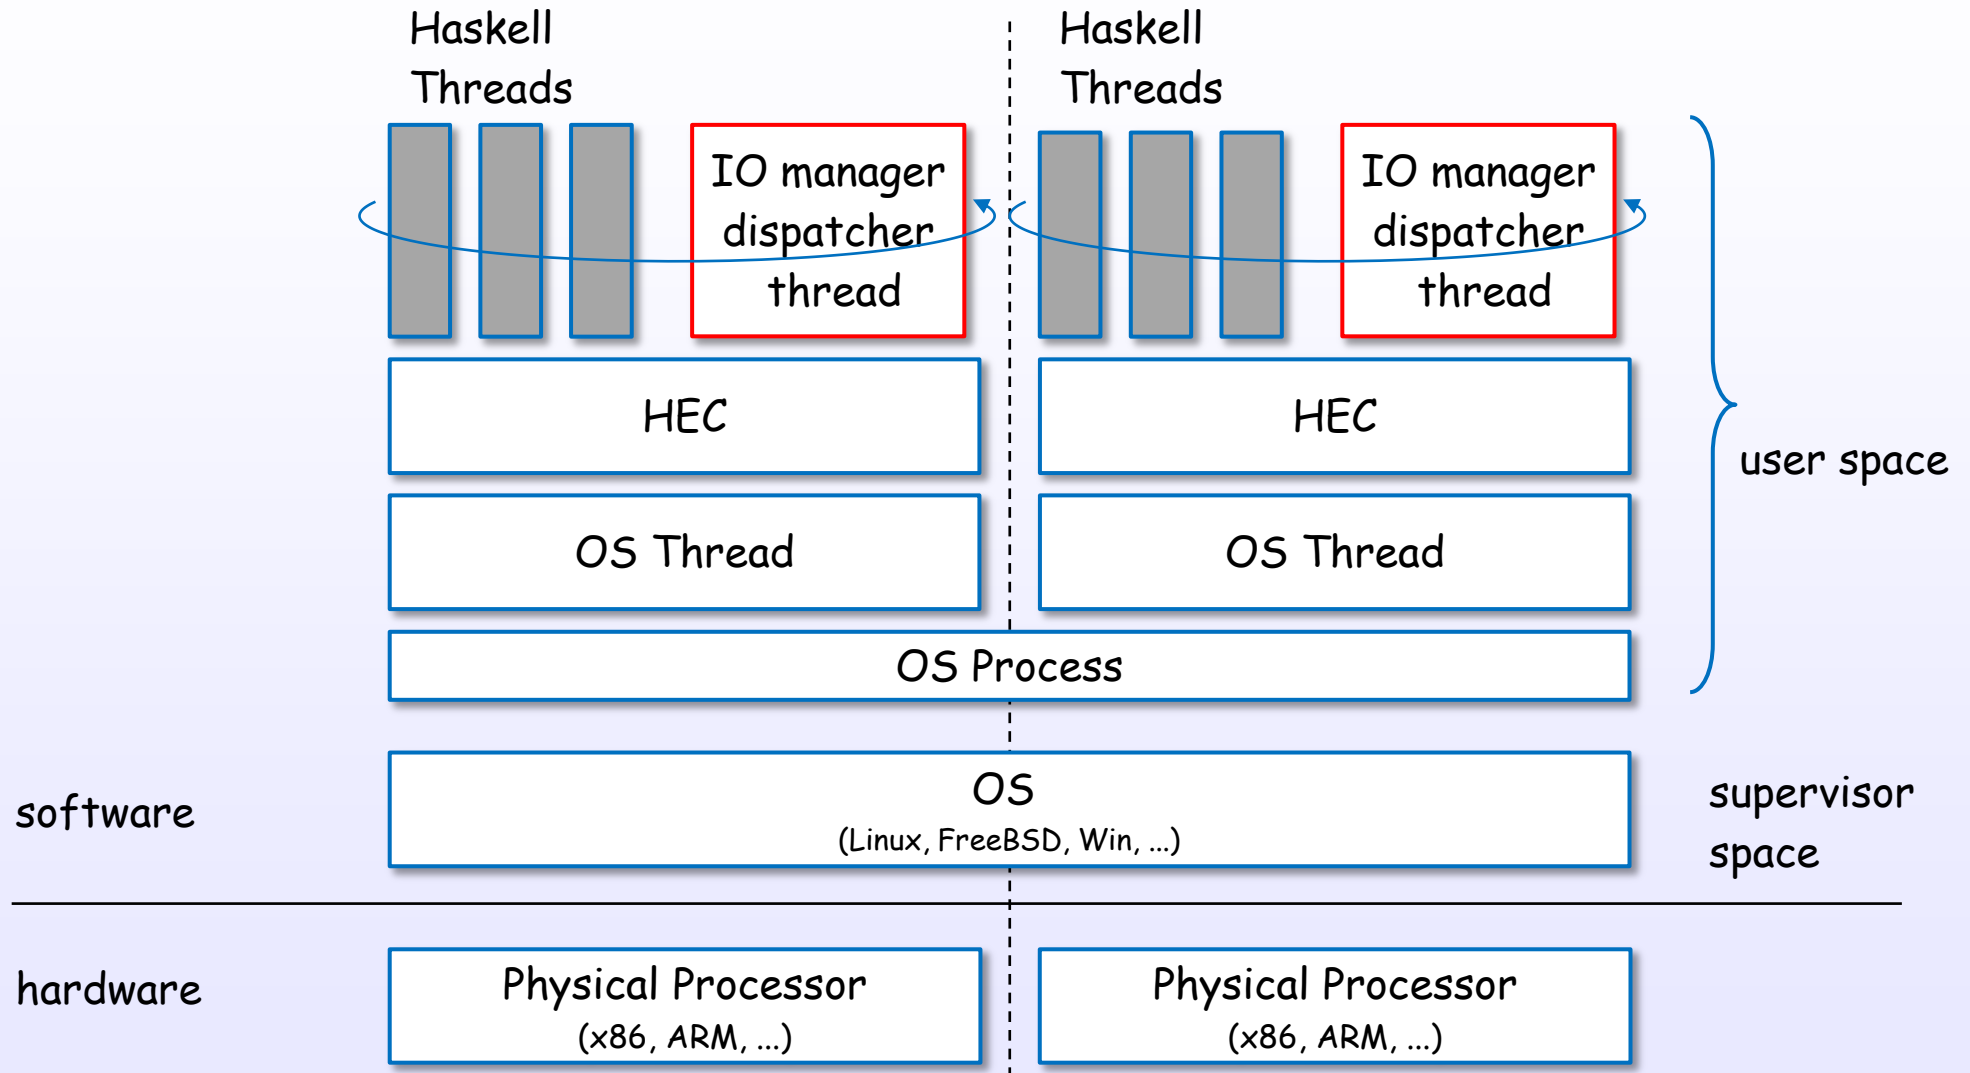
IO manager

IO manager (single core)



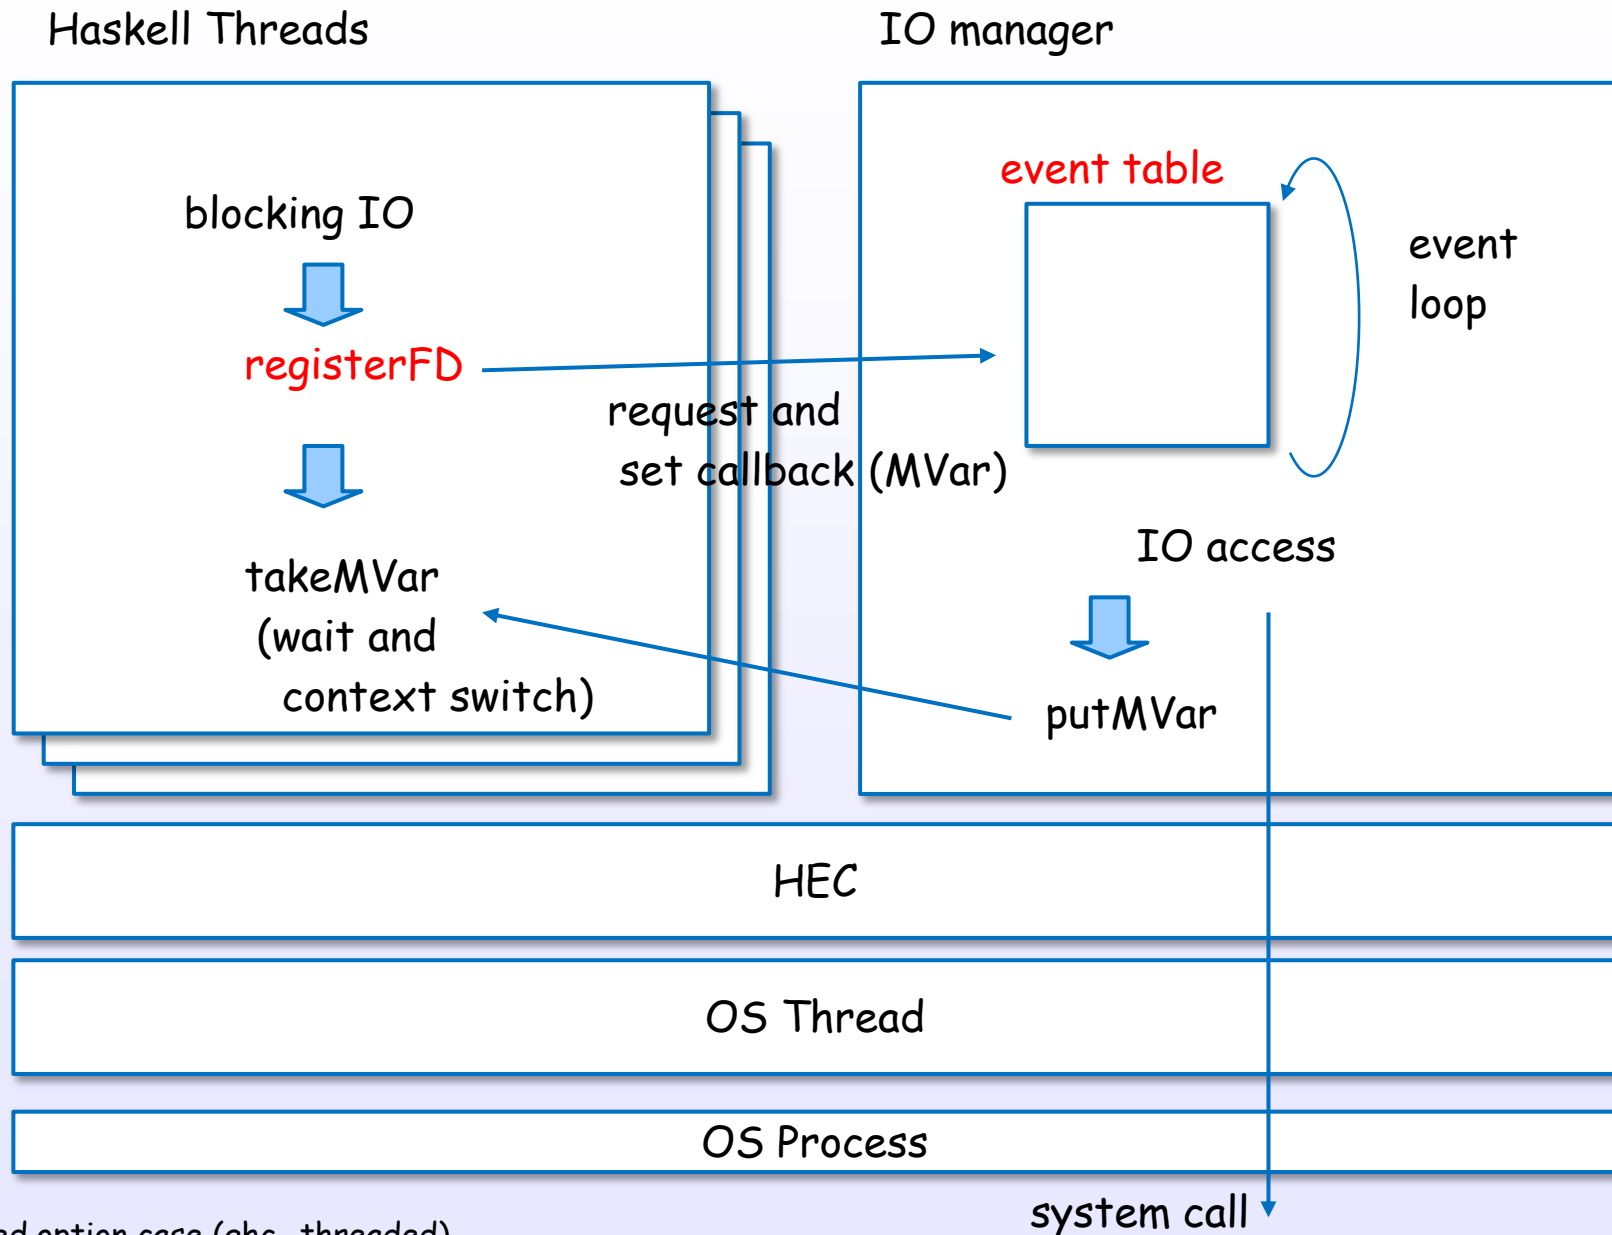
*Threaded option case (ghc -threaded)

IO manager (multi core)



*Threaded option case (ghc -threaded)

IO manager




*Threaded option case (ghc -threaded)


References : [7], [5], [8], [S29], [S30], [S32], [S37], [S35], [S3]

Bootstrap

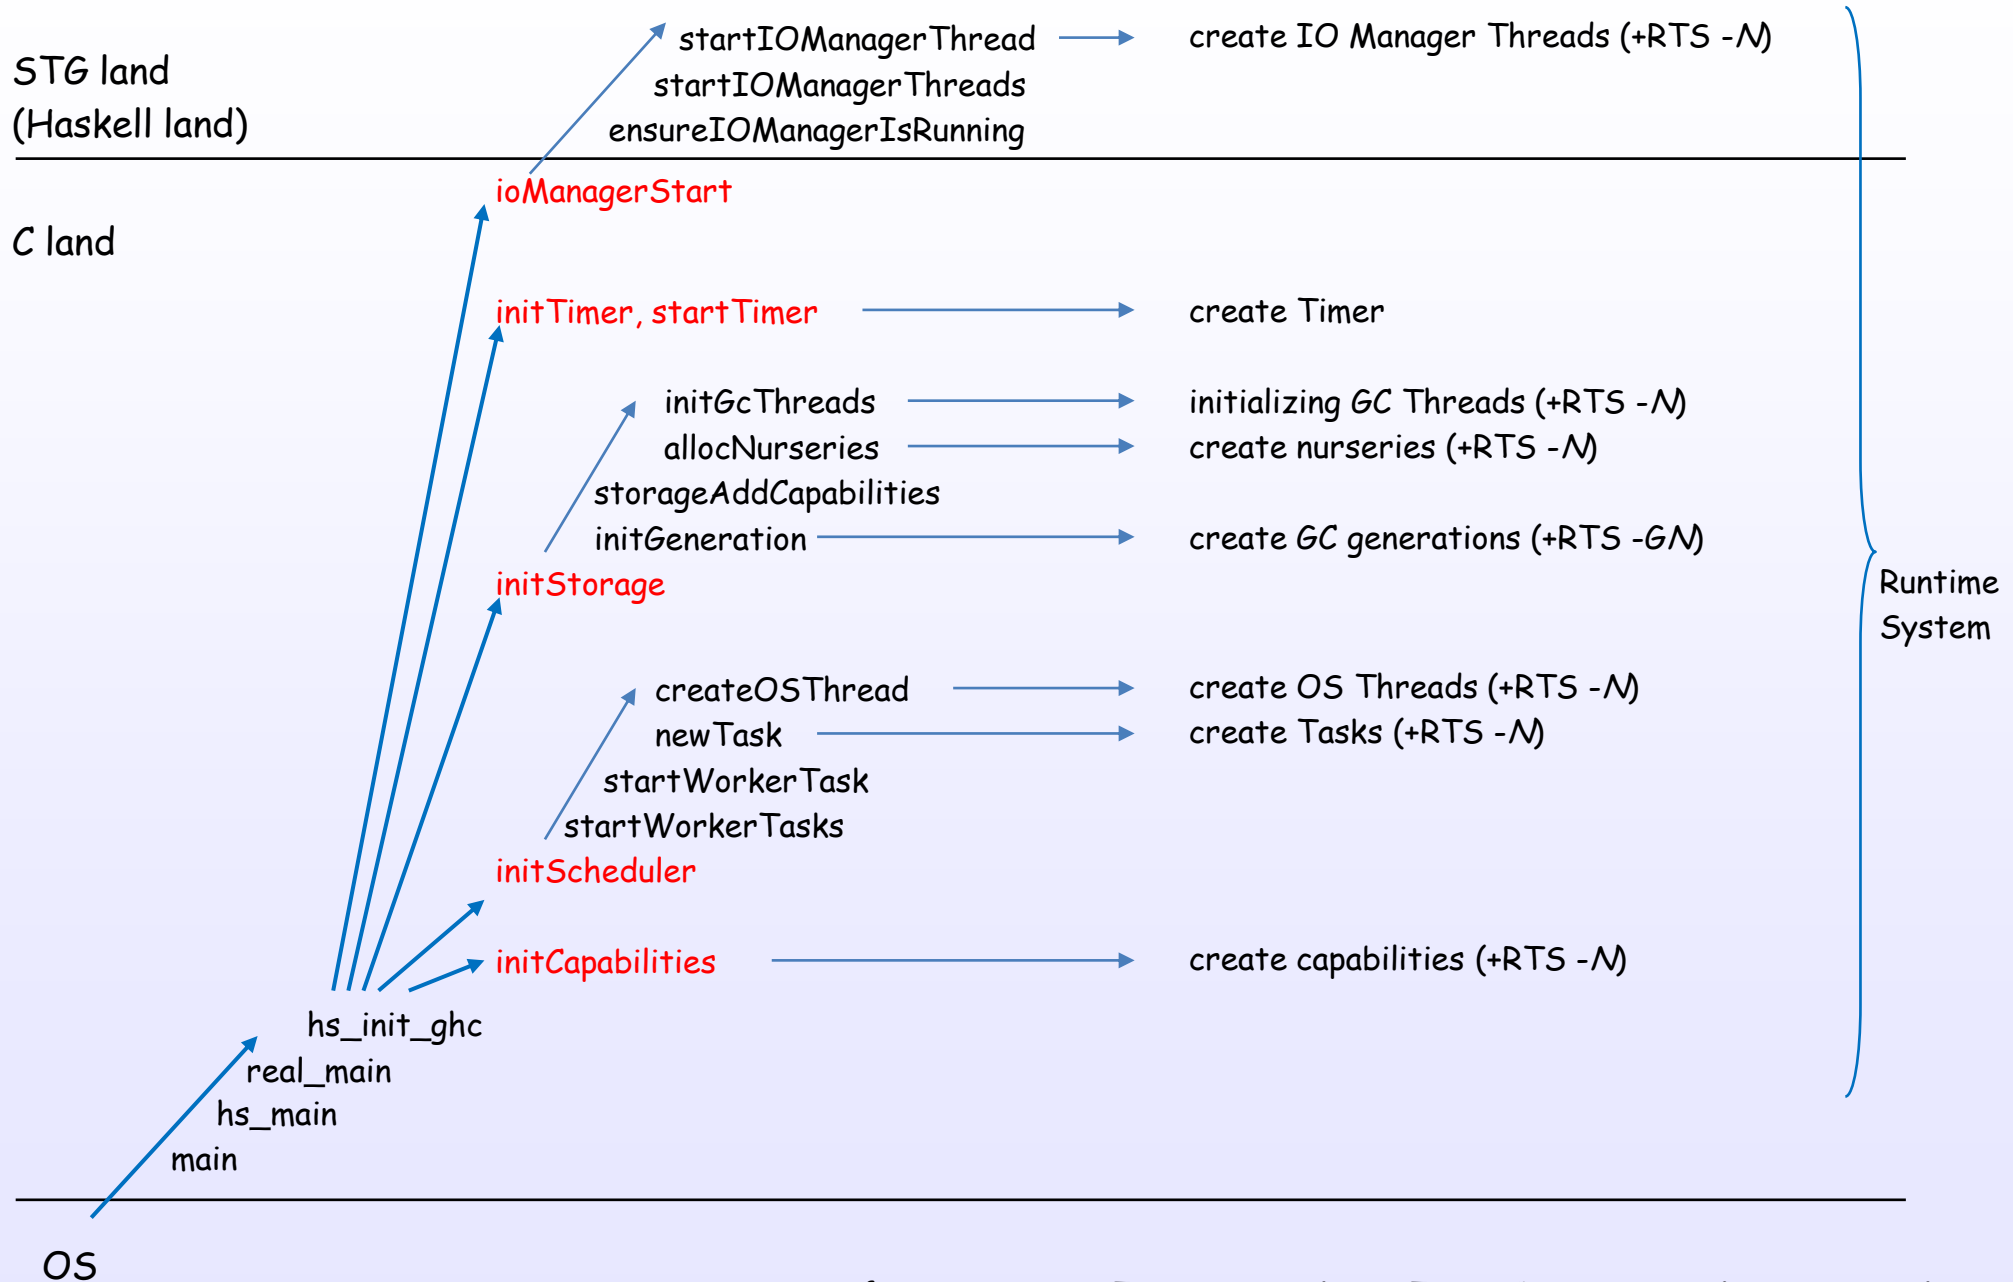
Bootstrap sequence




Exit sequence



Initializing



Create each layers



References

References

- [1] The Glorious Glasgow Haskell Compilation System User's Guide
https://downloads.haskell.org/~ghc/latest/docs/html/users_guide/index.html
- [2] Implementing lazy functional languages on stock hardware: the Spineless Tagless G-machine Version 2.5
<http://research.microsoft.com/en-us/um/people/simonpj/Papers/spineless-tagless-gmachine.ps.gz>
- [3] Making a Fast Curry Push/Enter vs Eval/Apply for Higher-order Languages
<http://research.microsoft.com/en-us/um/people/simonpj/papers/eval-apply/>
- [4] Faster Laziness Using Dynamic Pointer Tagging
<http://research.microsoft.com/en-us/um/people/simonpj/papers/ptr-tag/ptr-tagging.pdf>
- [5] Runtime Support for Multicore Haskell
<http://research.microsoft.com/en-us/um/people/simonpj/papers/parallel/multicore-ghc.pdf>
- [6] Extending the Haskell Foreign Function Interface with Concurrency
<http://community.haskell.org/~simonmar/papers/conc-ffi.pdf>
- [7] Mio: A High-Performance Multicore IO Manager for GHC
<http://haskell.cs.yale.edu/wp-content/uploads/2013/08/hask035-voellmy.pdf>
- [8] The GHC Runtime System
web.mit.edu/~ezyang/Public/jfp-ghc-rts.pdf
- [9] The GHC Runtime System
<http://www.scs.stanford.edu/14sp-cs240h/slides/ghc-rts.pdf>
- [10] Evaluation on the Haskell Heap
<http://blog.ezyang.com/2011/04/evaluation-on-the-haskell-heap/>

References

- [11] IO evaluates the Haskell Heap
<http://blog.ezyang.com/2011/04/io-evaluates-the-haskell-heap/>
- [12] Understanding the Stack
<http://www.well-typed.com/blog/94/>
- [13] Understanding the RealWorld
<http://www.well-typed.com/blog/95/>
- [14] The GHC scheduler
<http://blog.ezyang.com/2013/01/the-ghc-scheduler/>
- [15] GHC's Garbage Collector
http://www.mm-net.org.uk/workshop190404/GHC's_Garbage_Collector.ppt
- [16] Concurrent Haskell
<http://www.haskell.org/ghc/docs/papers/concurrent-haskell.ps.gz>
- [17] Beautiful Concurrency
<https://www.fpcomplete.com/school/advanced-haskell/beautiful-concurrency>
- [18] Anatomy of an MVar operation
<http://blog.ezyang.com/2013/05/anatomy-of-an-mvar-operation/>
- [19] Parallel and Concurrent Programming in Haskell
<http://community.haskell.org/~simonmar/pcph/>
- [20] Real World Haskell
<http://book.realworldhaskell.org/>

References

The GHC Commentary

- [C1] <https://ghc.haskell.org/trac/ghc/wiki/Commentary>
- [C2] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/SourceTree>
- [C3] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler>
- [C4] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/HscMain>
- [C5] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/CoreSynType>
- [C6] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/StgSynType>
- [C7] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/CmmType>
- [C8] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/GeneratedCode>
- [C9] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/SymbolNames>
- [C10] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts>
- [C11] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/HeapObjects>
- [C12] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/Stack>
- [C13] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/GC>
- [C14] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution>
- [C15] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution/Registers>
- [C16] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution/PointerTagging>
- [C17] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Scheduler>
- [C18] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/STM>
- [C19] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Libraries>

References

Source code

- [S1] `includes/stg/Regs.h`
- [S2] `includes/stg/MachRegs.h`
- [S3] `includes/rts/storage/ClosureTypes.h`
- [S4] `includes/rts/storage/Closures.h`
- [S5] `includes/rts/storage/TSO.h`
- [S6] `includes/rts/storage/InfoTables.h`
- [S7] `compiler/main/DriverPipeline.hs`
- [S8] `compiler/main/HscMain.hs`
- [S9] `compiler/cmm/CmmParse.y.source`
- [S10] `compiler/codeGen/StgCmmForeign.hs`
- [S11] `compiler/codeGen/Stg*.hs`
- [S12] `rts/PrimOps.cmm`
- [S13] `rts/RtsMain.c`
- [S14] `rts/RtsAPI.c`
- [S15] `rts/Capability.h`
- [S16] `rts/Capability.c`
- [S17] `rts/Schedule.c`
- [S18] `rts/StgCRun.c`
- [S19] `rts/StgStartup.cmm`
- [S20] `rts/StgMiscClosures.cmm`
- [S21] `rts/HeapStackCheck.cmm`
- [S22] `rts/Threads.c`
- [S23] `rts/Task.c`
- [S24] `rts/Timer.c`
- [S25] `rts/sm/GC.c`
- [S26] `rts/Sparks.c`
- [S27] `rts/WSDeque.c`
- [S28] `rts/STM.h`
- [S29] `rts/posix/Signals.c`
- [S30] `rts/win32/ThrIOManager.c`
- [S31] `libraries/base/GHC/MVar.hs`
- [S32] `libraries/base/GHC/Conc/IO.hs`
- [S33] `libraries/base/GHC/Conc/Sync.lhs`
- [S34] `libraries/base/GHC/Event/Manager.hs`
- [S35] `libraries/base/GHC/Event/Thread.hs`
- [S36] `libraries/base/GHC/IO/BufferedIO.hs`
- [S37] `libraries/base/GHC/IO/FD.hs`
- [S38] `libraries/base/GHC/IO/Handle/Text.hs`
- [S39] `libraries/base/System/IO.hs`
- [S40] `libraries/base/System/Posix/Internals.hs`
- [S41] `AutoApply.o (utils/genapply/GenApply.hs)`

Connect the algorithm and transistor