

$GHC_{(STG, Cmm, asm)}$ illustrated

for hardware person

exploring some mental models and implementations

Takenobu T.

"Any sufficiently advanced technology is
indistinguishable from **magic**."

Arthur C. Clarke

NOTE

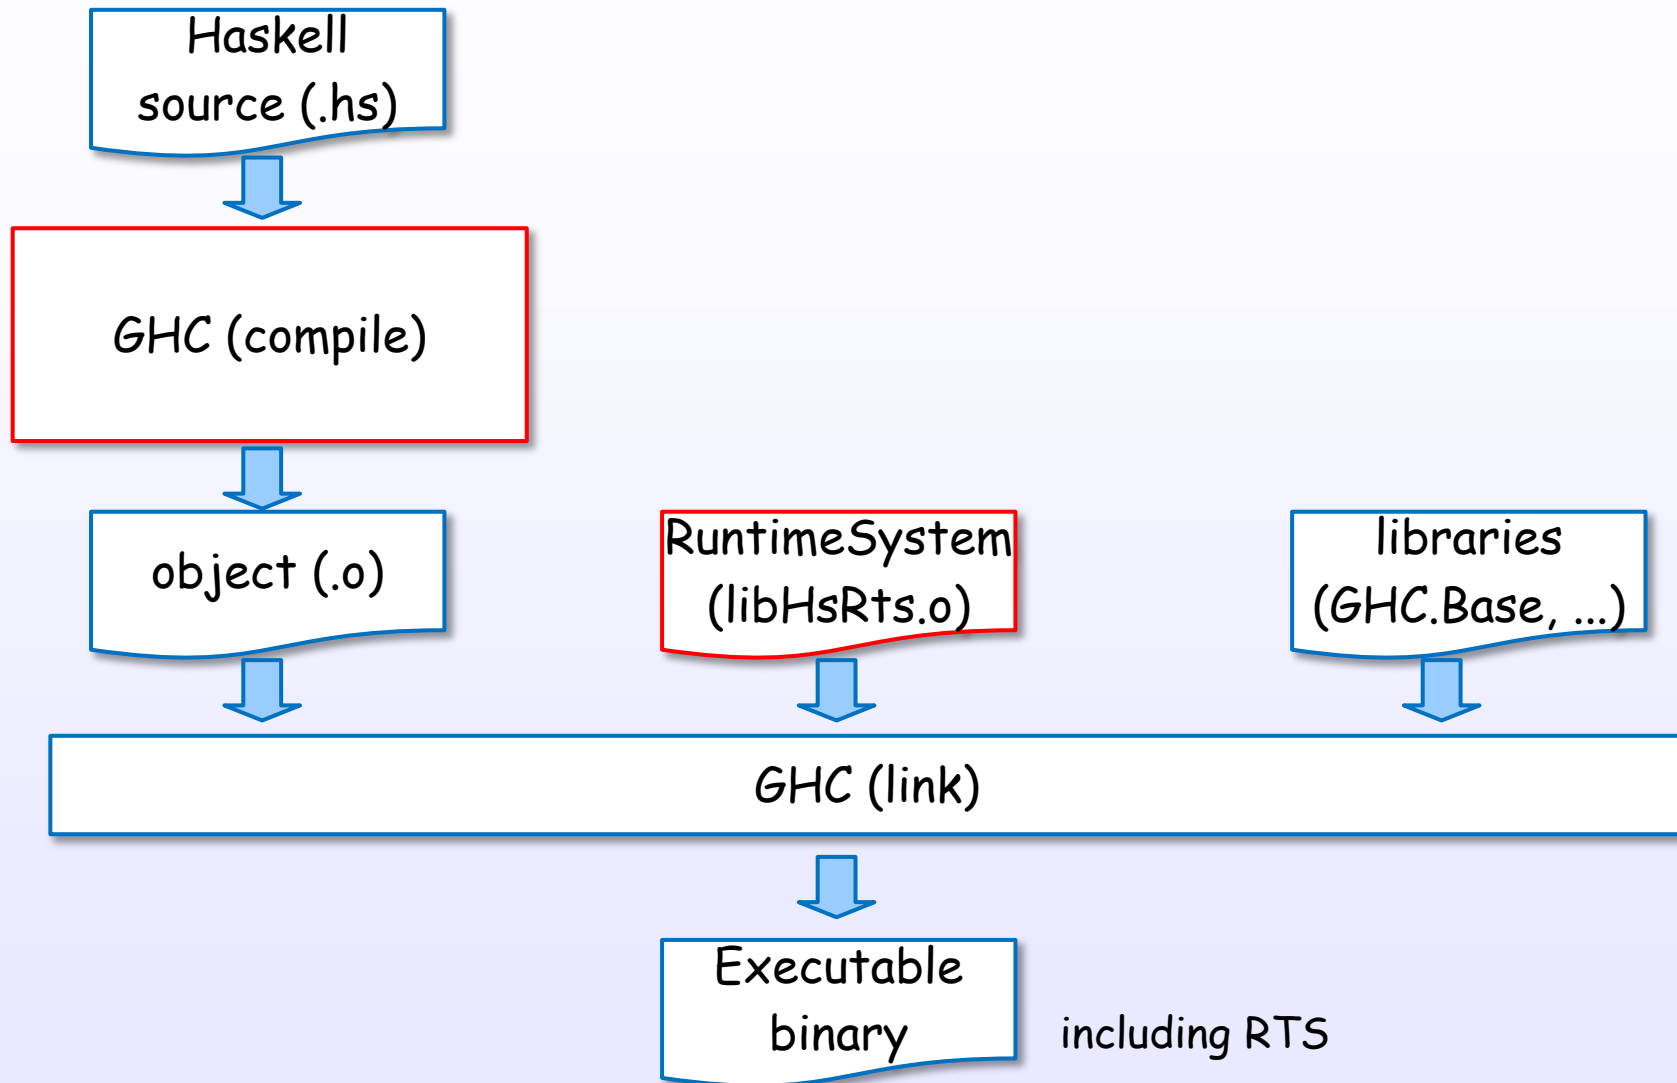
- This is not official document by ghc development team.
- Please don't forget "semantics". It's very important.
- This is written for ghc 7.8.

Contents

- Executable binary
- Compile steps
- Runtime System
- Development languages
- Machine layer/models
- STG-machine
- Heap object in STG-machine
- STG-machine evaluation
- Pointer tagging
- Thunk and update
- Allocate and free(GC) heap objects
- STG - C land interface
- Thread
- Heap and Threads
- Thread context switch
- Creating main thread and forkIO
- Spark
- MVar
- IO and FFI
- IO manager
- Bootstrap
- References

Executable binary

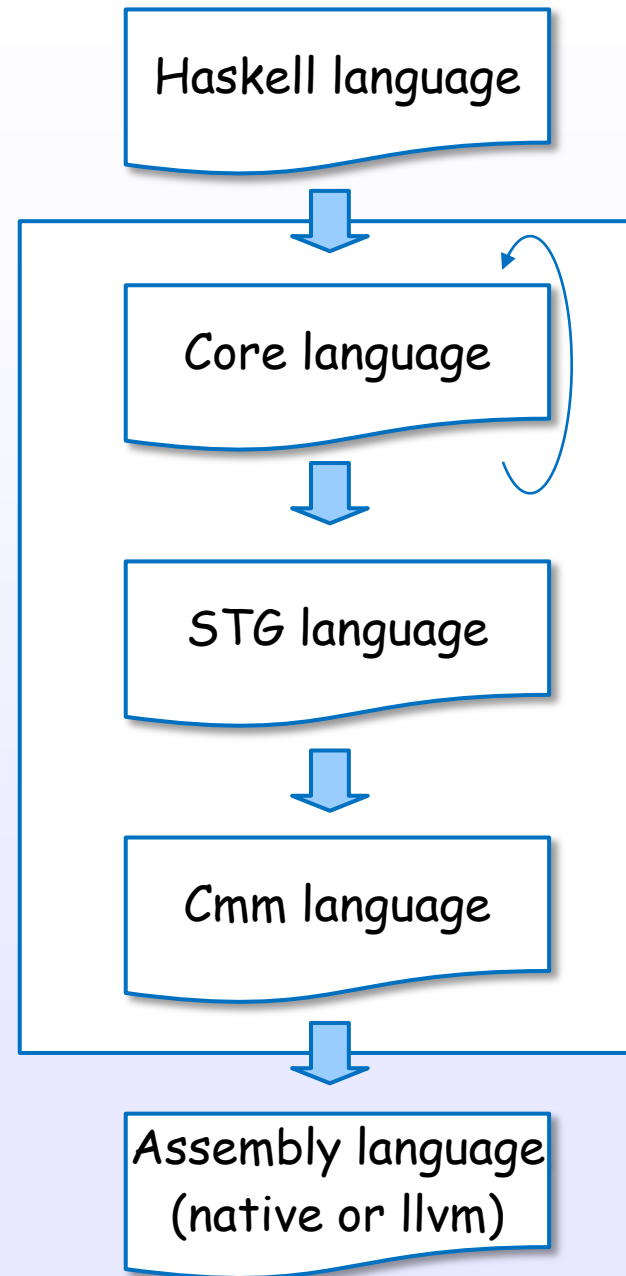
GHC = Compiler + Runtime System (RTS)



Compile steps

GHC transitions between five representations

GHC
compile
steps



each code dumped by

% ghc -ddump-parsed
% ghc -ddump-rn

% ghc -ddump-ds
% ghc -ddump-simpl
% ghc -ddump-prep

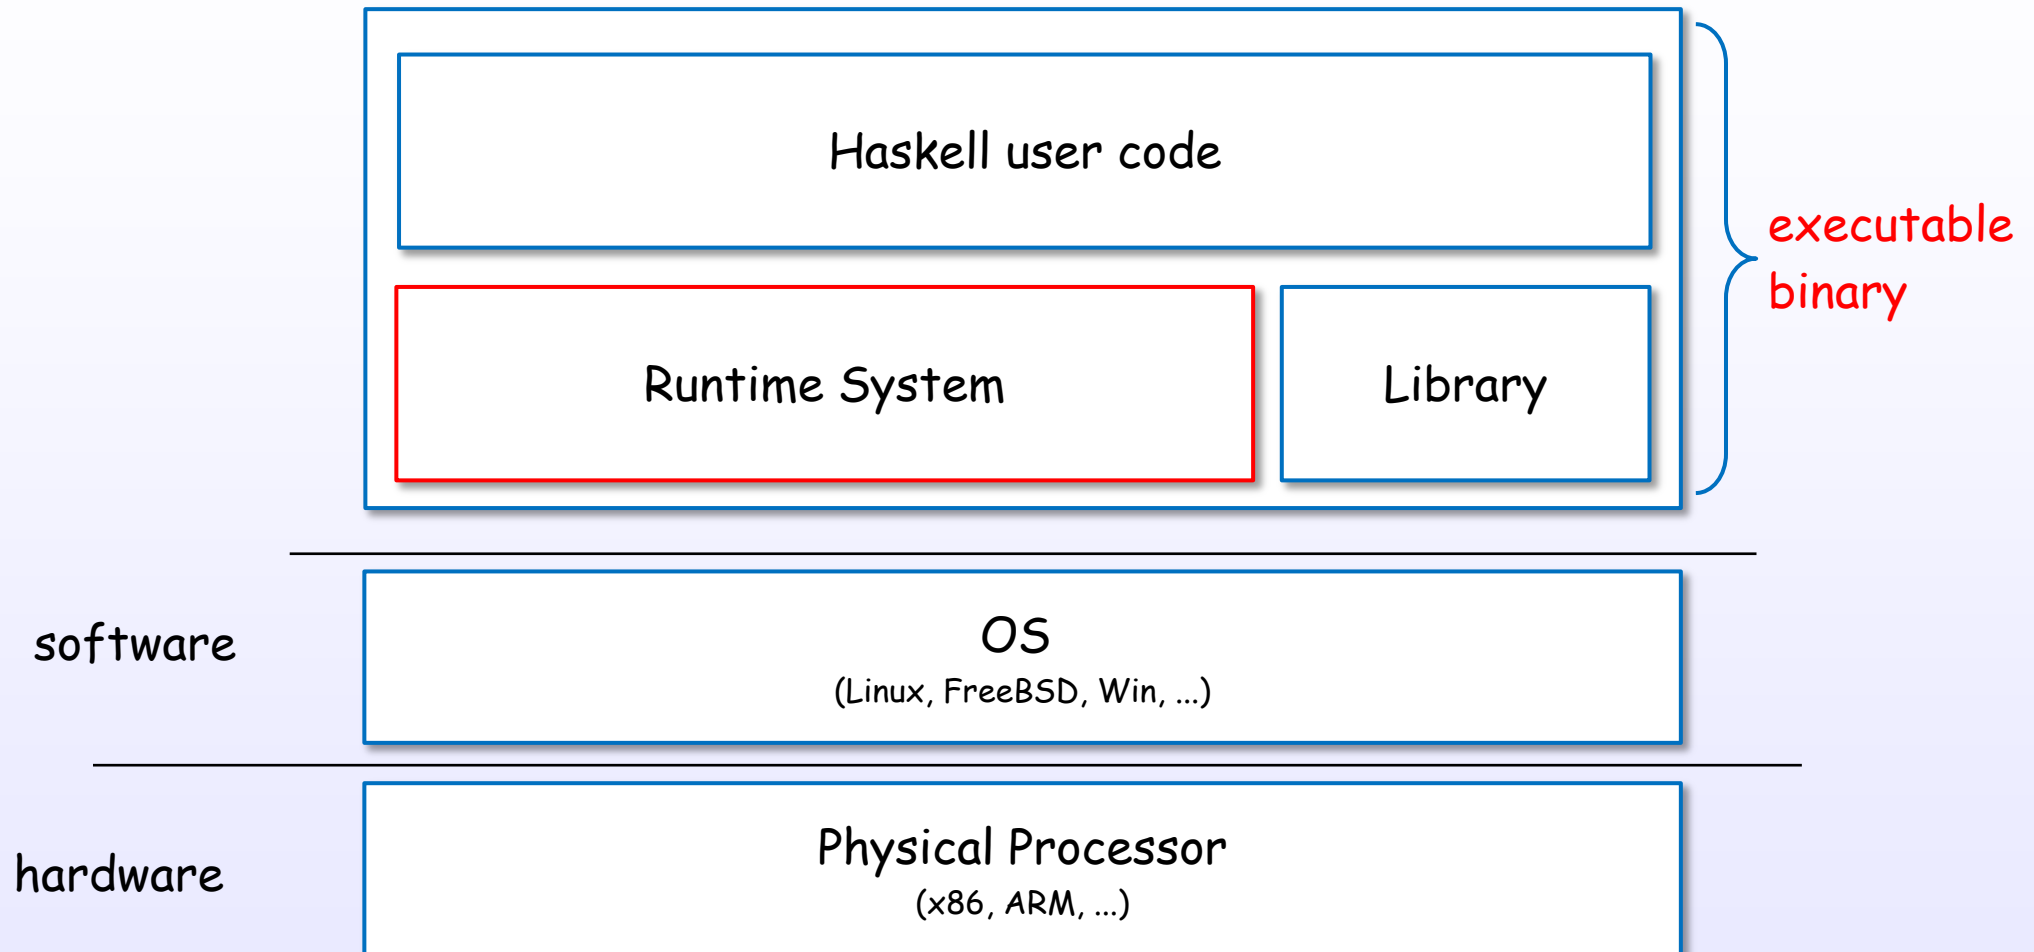
% ghc -ddump-stg

% ghc -ddump-cmm
% ghc -ddump-opt-cmm

% ghc -ddump-llvm
% ghc -ddump-asm

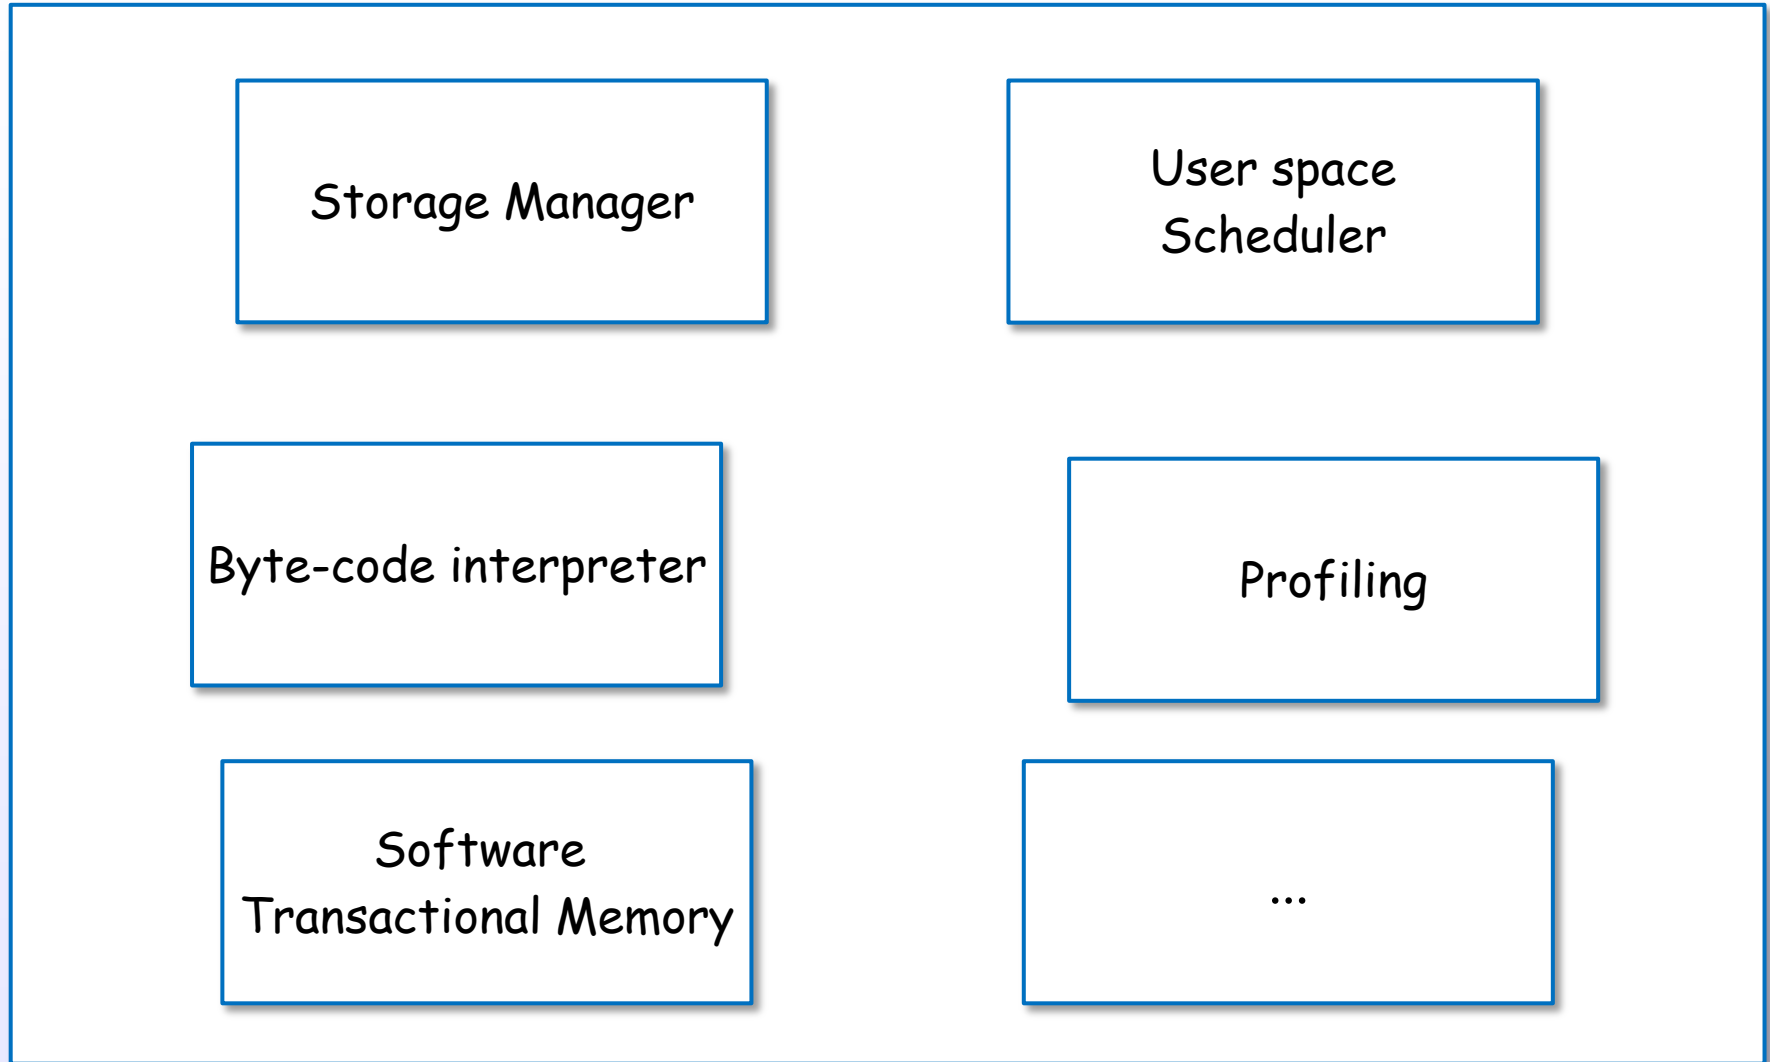
Runtime System

Generated binary includes RTS



Runtime System includes ...

Runtime System



Development languages

GHC developed by some languages

compiler

(\$(TOP)/**compiler**/*)

Haskell

+

Alex (lex)

Happy (yacc)

Cmm (C--)

Assembly

runtime system

(\$(TOP)/**rts**/*)

C

+

Cmm

Assembly

library

(\$(TOP)/**libraries**/*)

Haskell

+

C

Machine layer/models

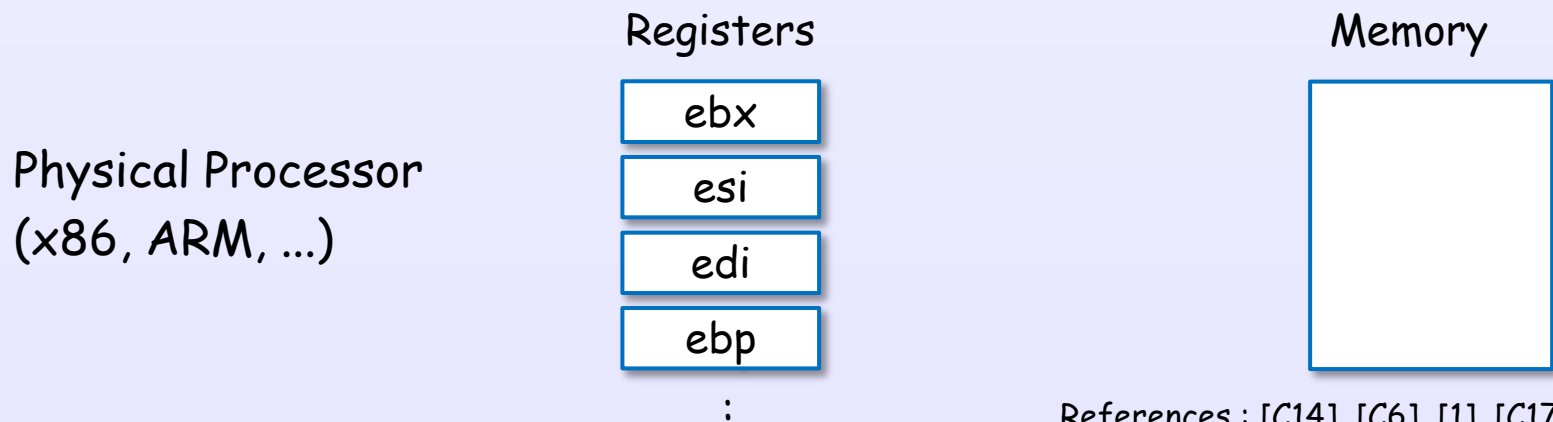
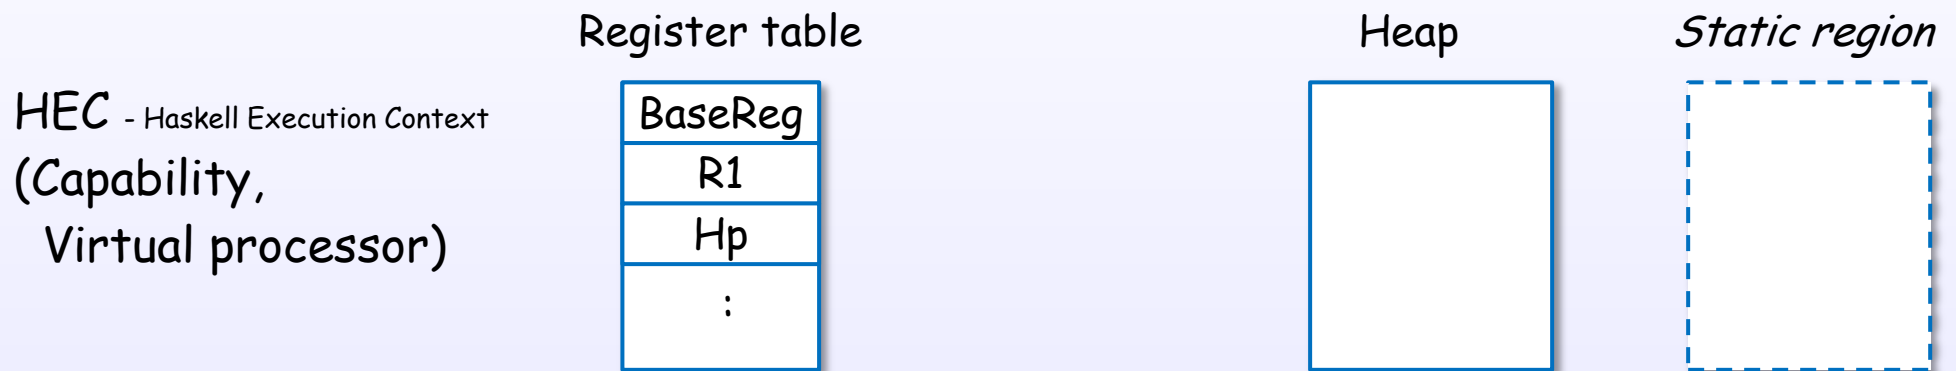
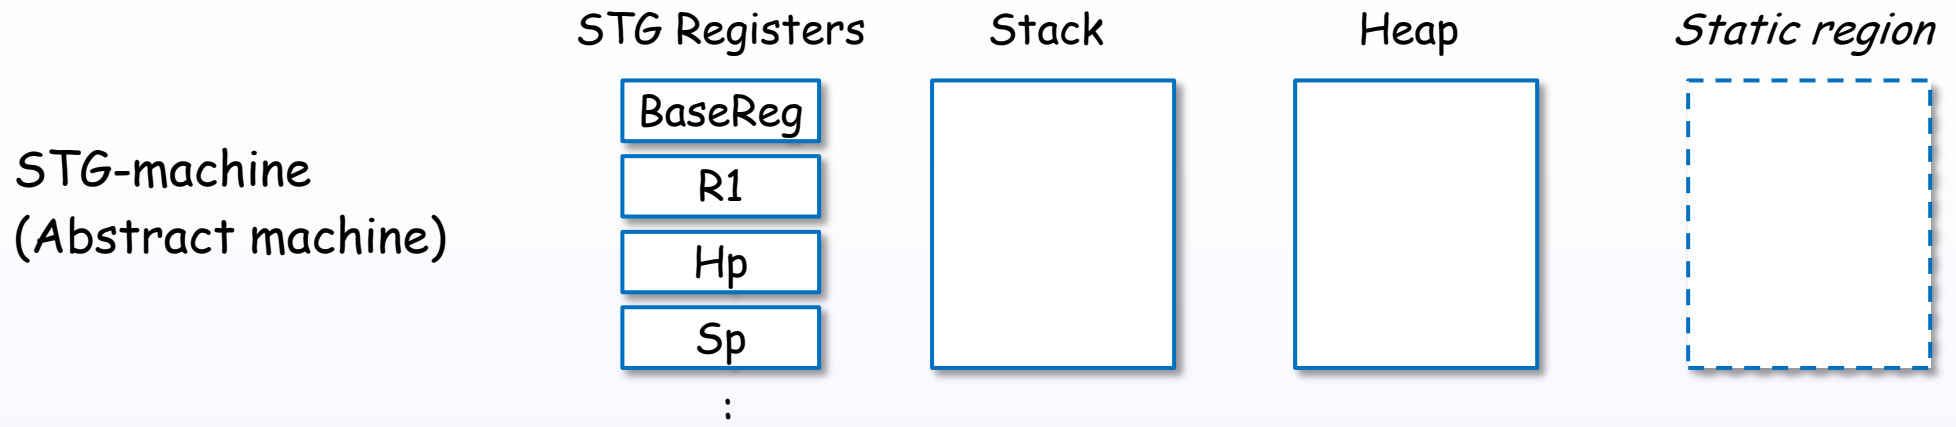
Machine layer

STG-machine
(Abstract machine)

HEC - Haskell Execution Context
(Capability, Virtual processor)

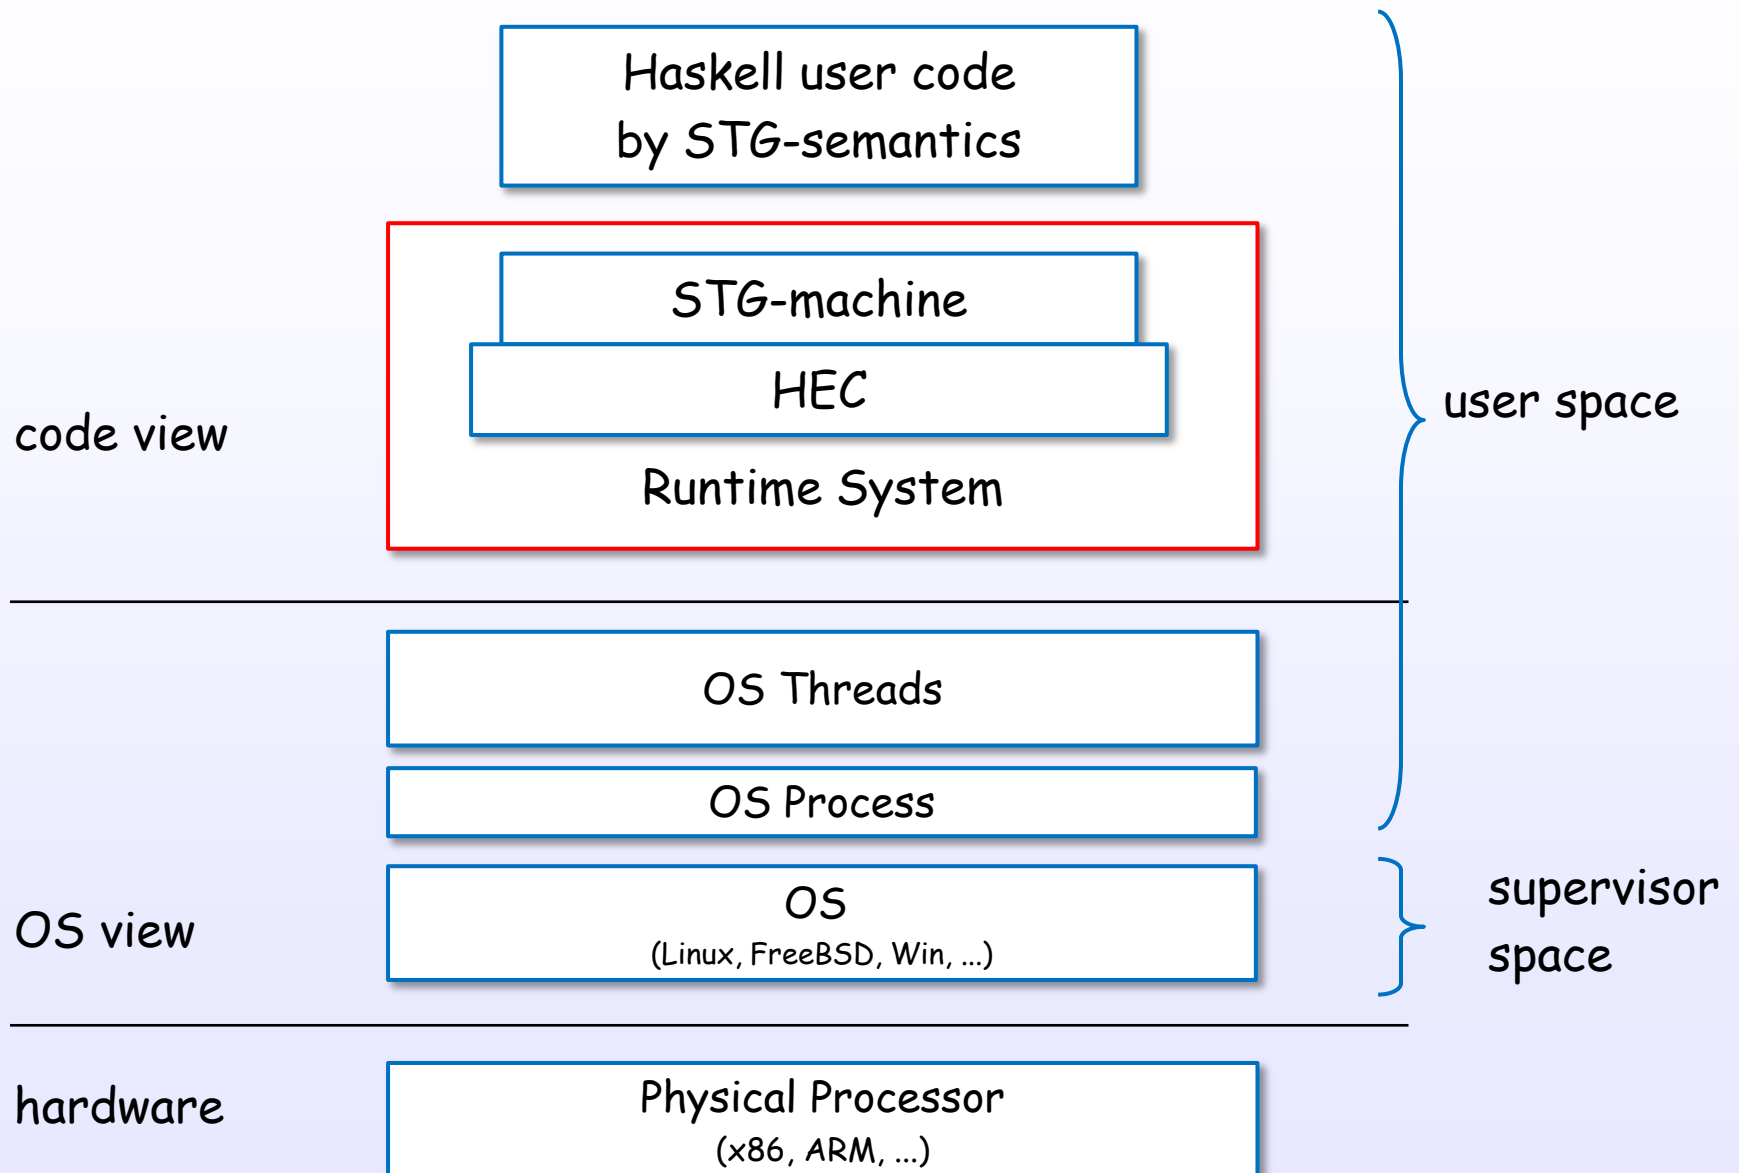
Physical Processor
(x86, ARM, ...)

Machine layer



References : [C14], [C6], [1], [C17], [7], [S15], [S16], [S11]

Runtime system and HEC

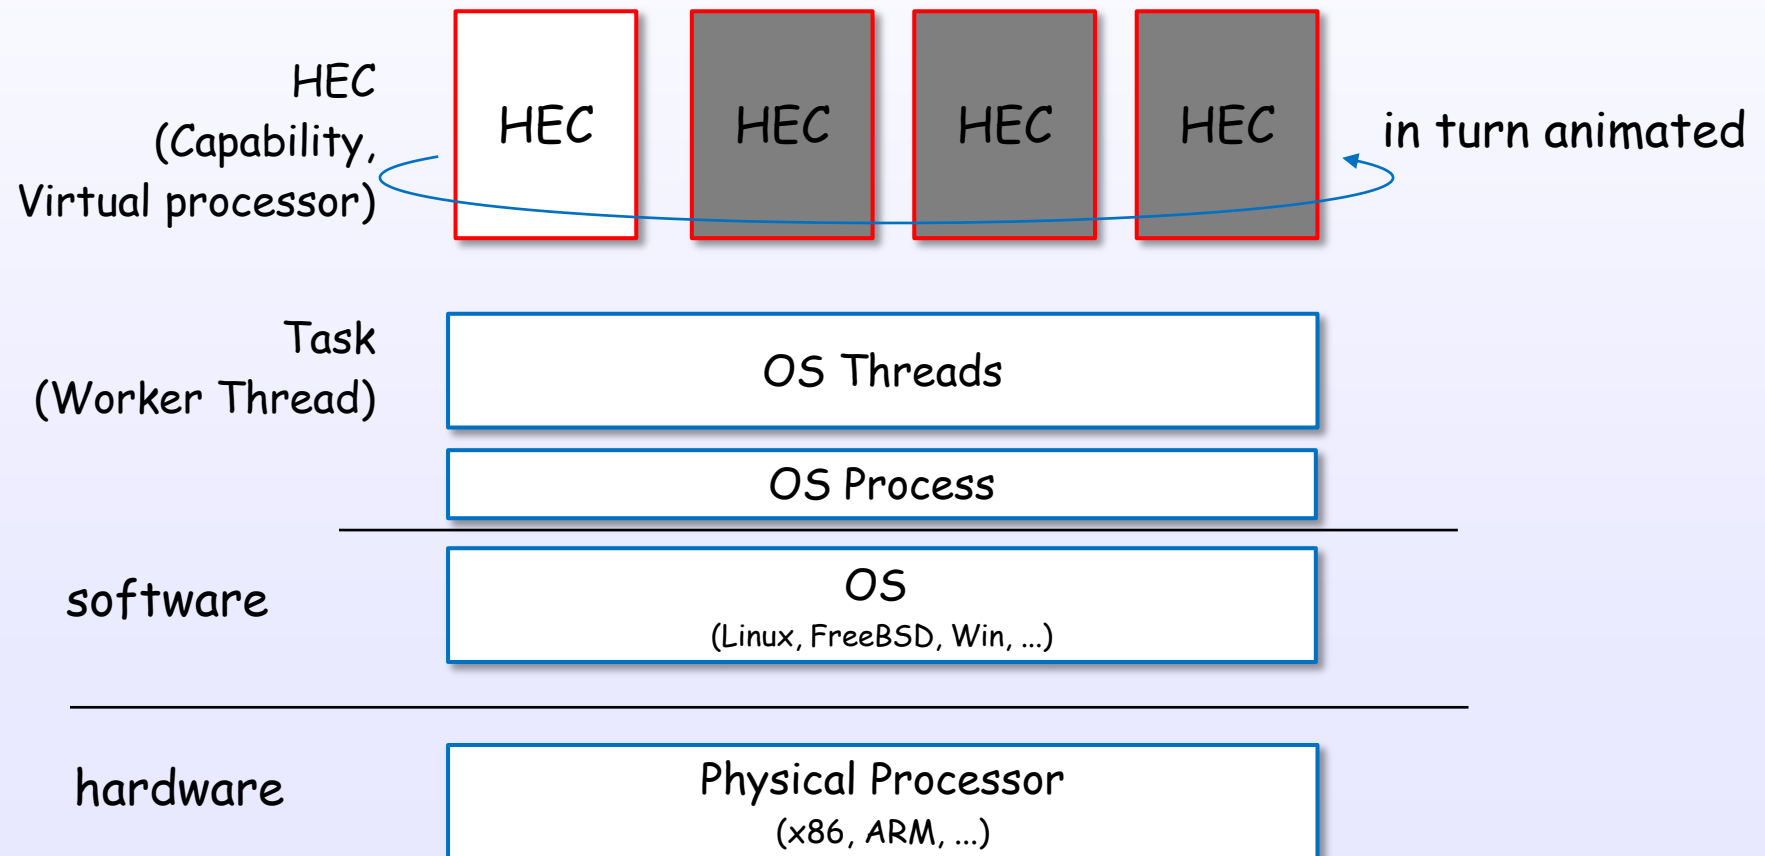


many HECs

multi HEC generated by compile and runtime options :

```
% ghc -rtsopts -threaded
```

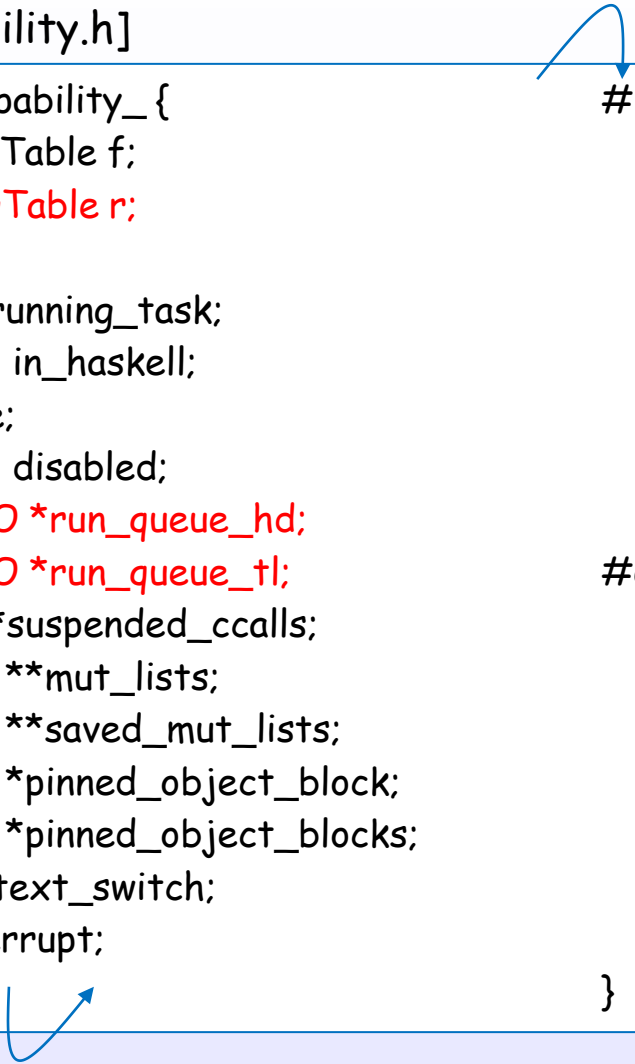
```
% ./xxx +RTS -N4
```



HEC (Capability) data structure

[rts/Capability.h]

```
struct Capability_ {  
    StgFunTable f;  
    StgRegTable r;  
    nat no;  
    Task *running_task;  
    rtsBool in_haskell;  
    nat idle;  
    rtsBool disabled;  
    StgTSO *run_queue_hd;  
    StgTSO *run_queue_tl;  
    InCall *suspended_ccalls;  
    bdescr **mut_lists;  
    bdescr **saved_mut_lists;  
    bdescr *pinned_object_block;  
    bdescr *pinned_object_blocks;  
    int context_switch;  
    int interrupt;  
  
#if defined(THREADED_RTS)  
    Task *spare_workers;  
    nat n_spare_workers;  
    Mutex lock;  
    Task *returning_tasks_hd;  
    Task *returning_tasks_tl;  
    Message *inbox;  
    SparkPool *sparks;  
    SparkCounters spark_stats;  
#endif  
  
    W_ total_allocated;  
    StgTVarWatchQueue *free_tvar_watch_queues;  
    StgInvariantCheckQueue *free_invariant_check_queues;  
    StgTRecChunk *free_trec_chunks;  
    StgTRecHeader *free_trec_headers;  
    nat transaction_tokens;  
}
```



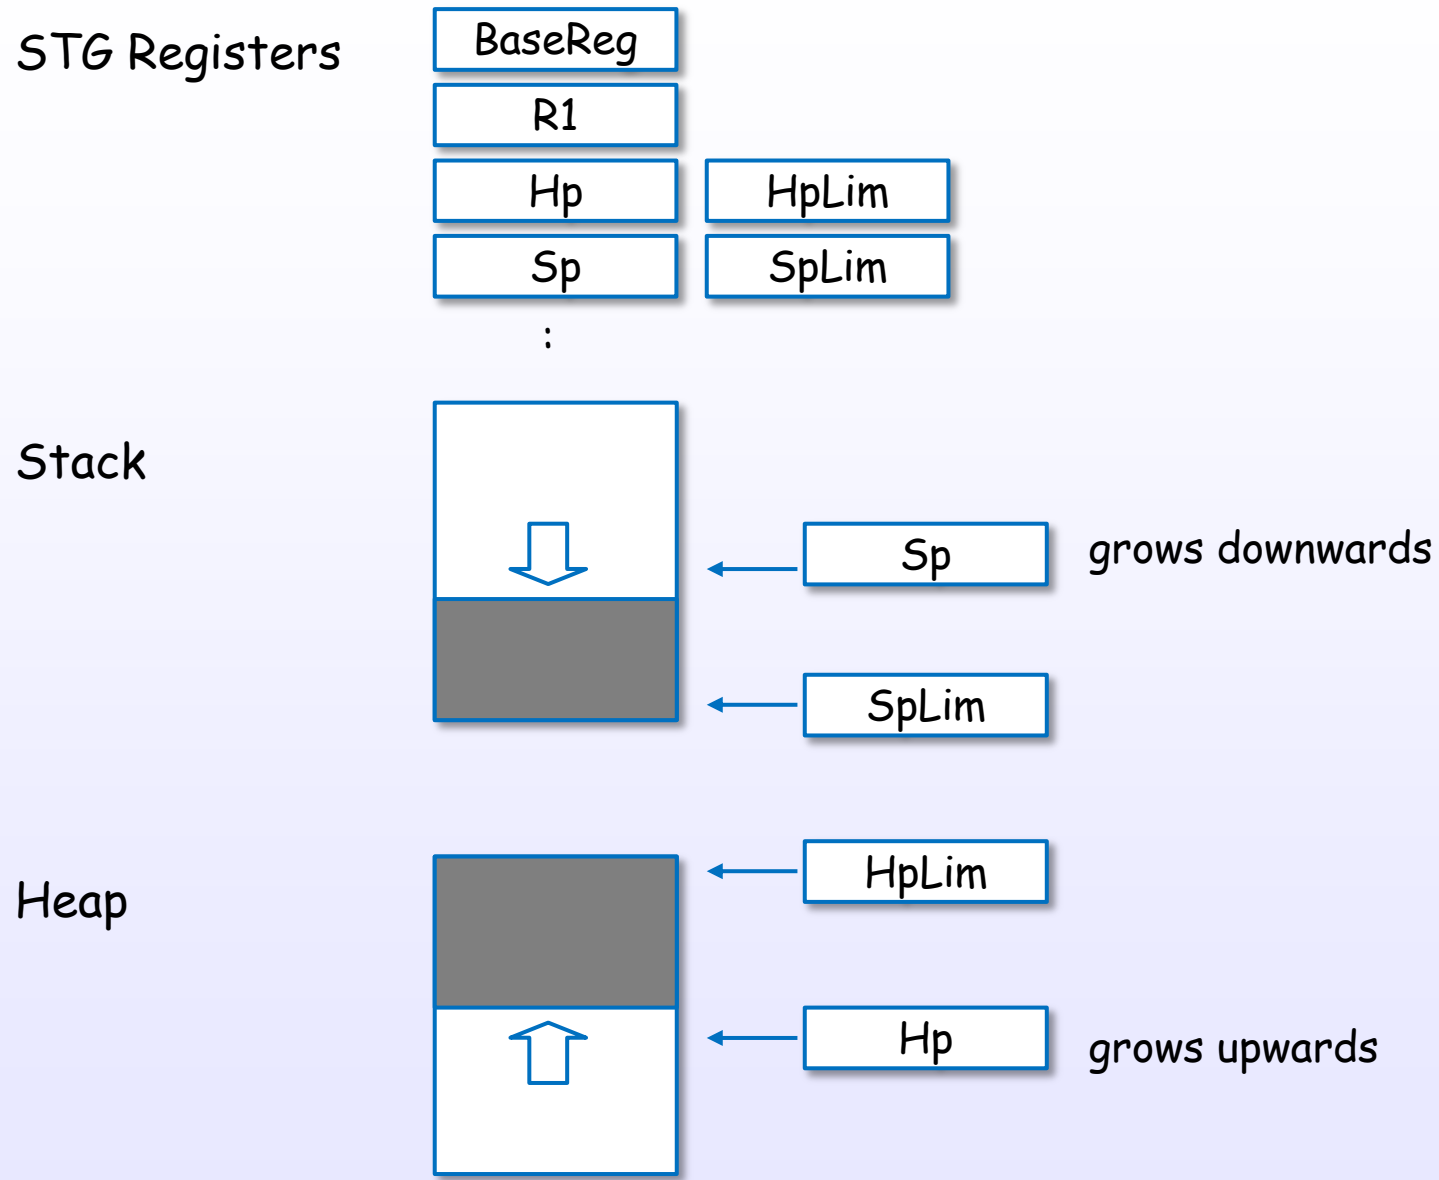
The diagram shows a blue box containing the C code for the Capability_ struct. A blue arrow points from the title 'HEC (Capability) data structure' to the struct definition. Another blue arrow points from the 'StgRegTable r;' field to the text 'HEC (Capabiilty) has Register table and Run queue and ...'. A third blue arrow points from the 'StgTSO *run_queue_hd;' field to the text 'HEC (Capability) is initialized in initCapabilities () [rts/rts/Capability.c]'.

HEC (Capabiilty) has Register table and Run queue and ...

HEC (Capability) is initialized in initCapabilities () [rts/rts/Capability.c]

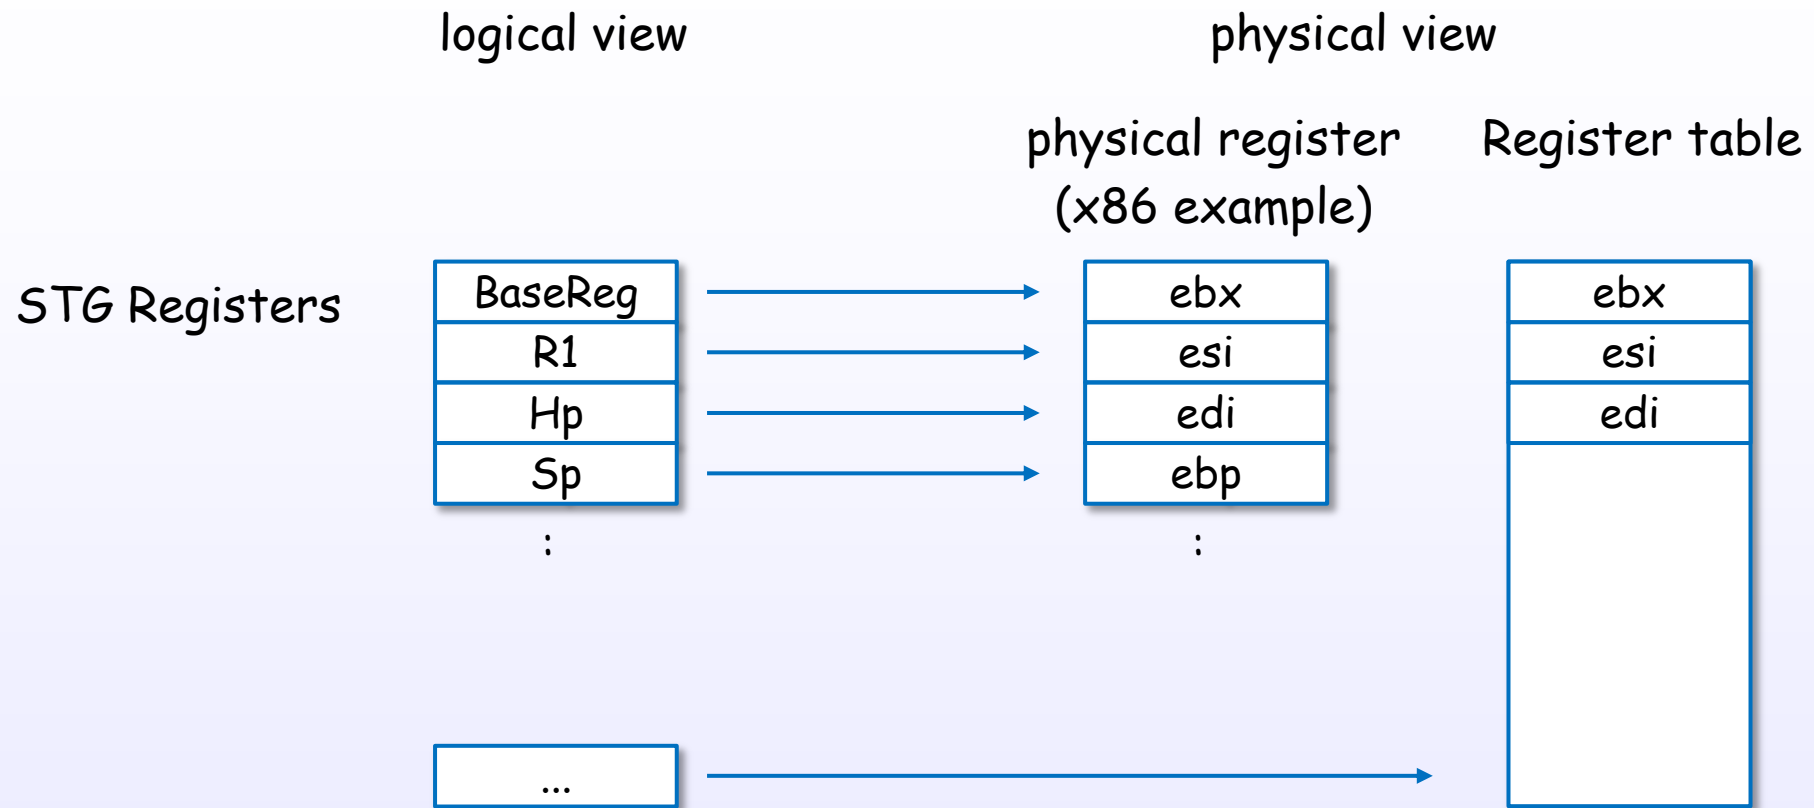
STG-machine

STG-machine

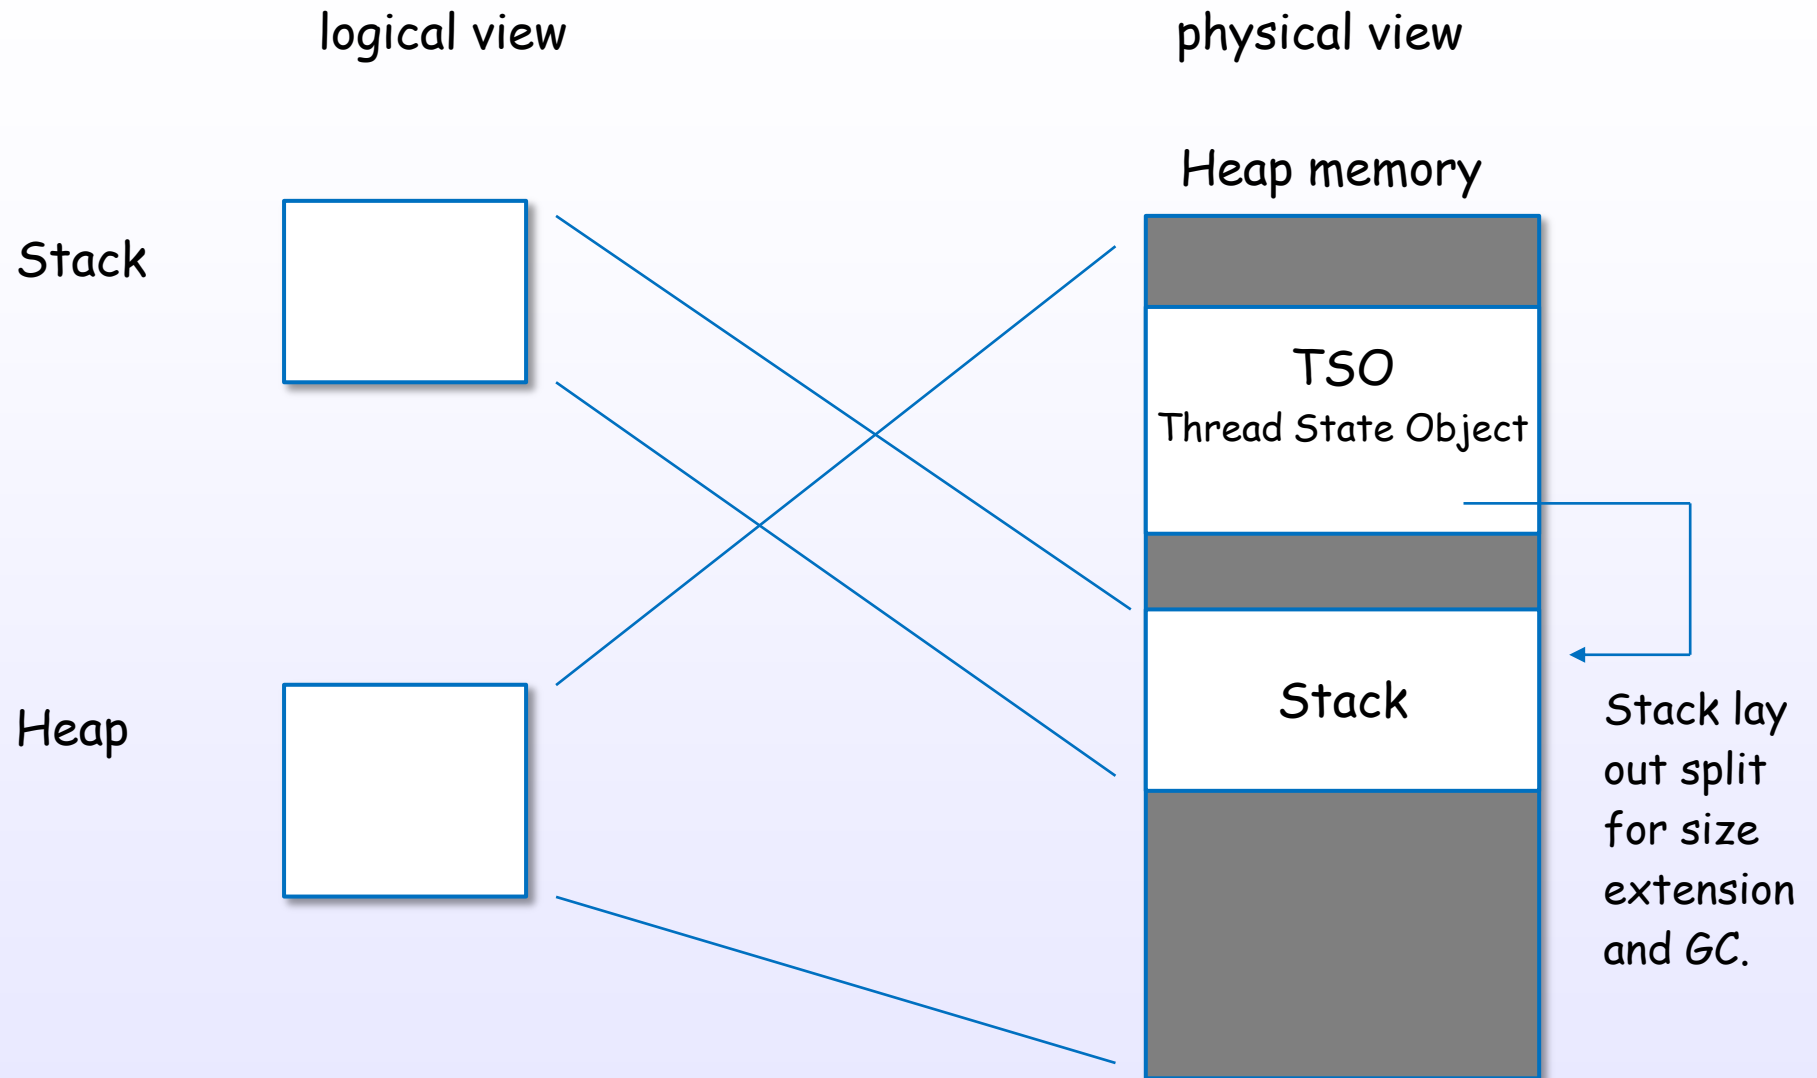


Real Haskell code executed in STG semantics.

STG-machine mapped to physical processor



STG-machine mapped to physical processor



Stack is in TSO object.
TSO object in heap.

TSO object

[includes/rts/storage/TSO.h]

```
typedef struct StgTSO_  
  StgHeader          header;  
  struct StgTSO_*     _link;  
  struct StgTSO_*     global_link;  
  struct StgStack_*   *stackobj;  
  StgWord16          what_next;  
  StgWord16          why_blocked;  
  StgWord32          flags;  
  StgTSOBlockInfo    block_info;  
  StgThreadID        id;  
  StgWord32          saved_errno;  
  StgWord32          dirty;  
  struct InCall_*     bound;  
  struct Capability_* cap;  
  struct StgTRecHeader_* trec;  
  struct MessageThrowTo_* blocked_exceptions;  
  struct StgBlockingQueue_* bq;  
  StgWord32 tot_stack_size;  
} *StgTSOPtr;
```

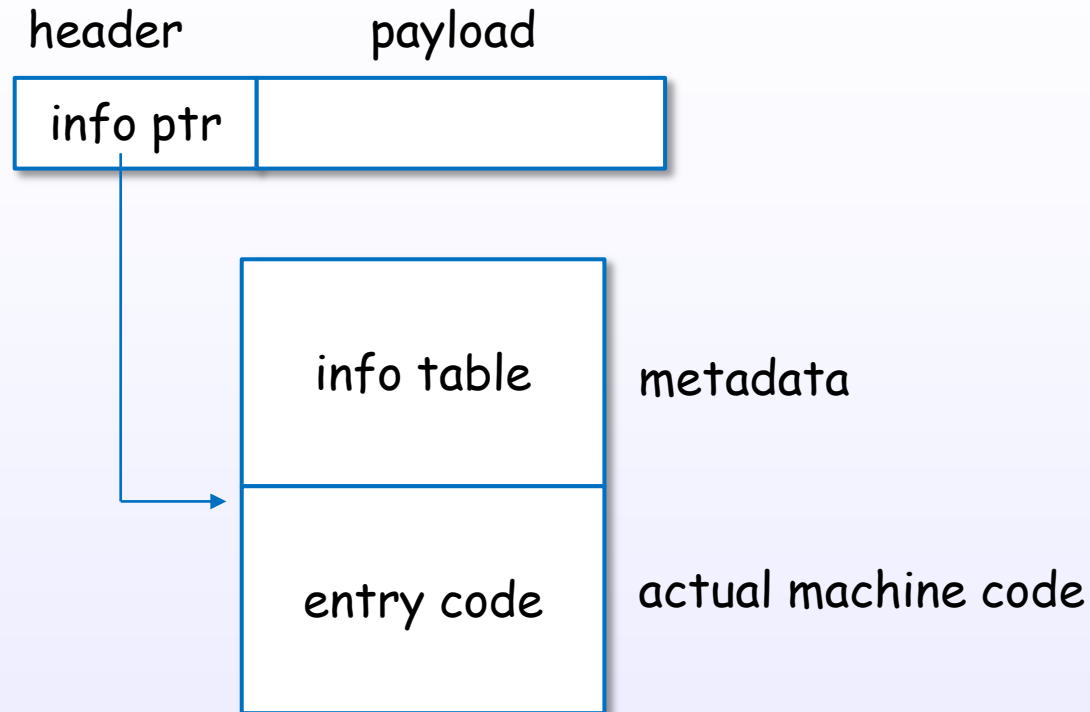
← link to stack object

TSO object is **only ~17words + stack**. Lightweight.

Heap object in STG-machine

Heap object (closure)


logical view




Closure (header + payload) + Info Table + Entry Code

Heap object (closure)

logical view

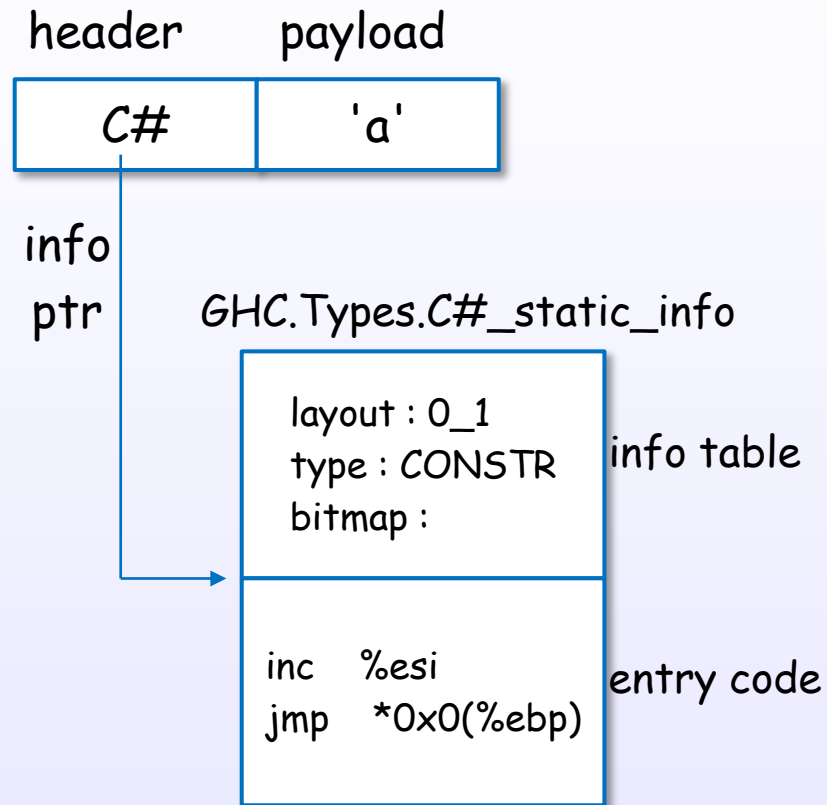


physical view

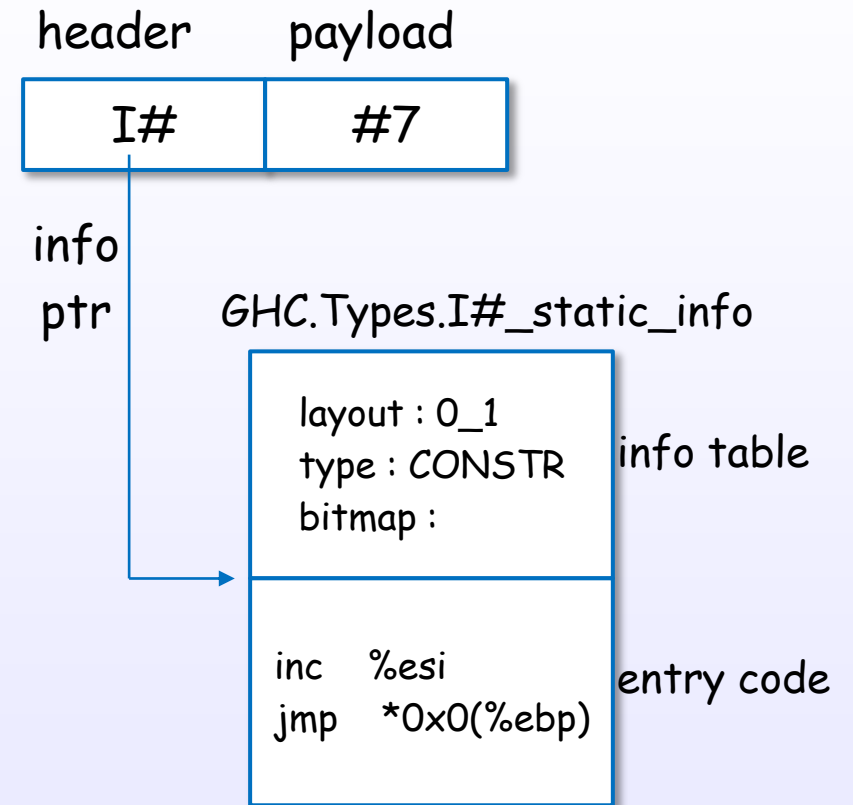


Closure examples : Char, Int

'a' :: Char



7 :: Int



Closure example code

[Example.hs]

```
module Example where
value1 :: Int
value1 = 7
```

STG

[ghc -O -ddump-stg Example.hs]

```
Example.value1 :: GHC.Types.Int
[GblId, Caf=NoCafRefs, Str=DmdType m, Unf=OtherCon []] =
  NO_CCS GHC.Types.I#! [7];
```

Cmm

[ghc -O -ddump-opt-cmm Example.hs]

```
section "data" { __stginit_main:Example:
}
```

```
section "data" {
  Example.value1 closure:
    const GHC.Types.I#_static_info;
    const 7;
}
```

asm

[ghc -O -ddump-asm Example.hs]

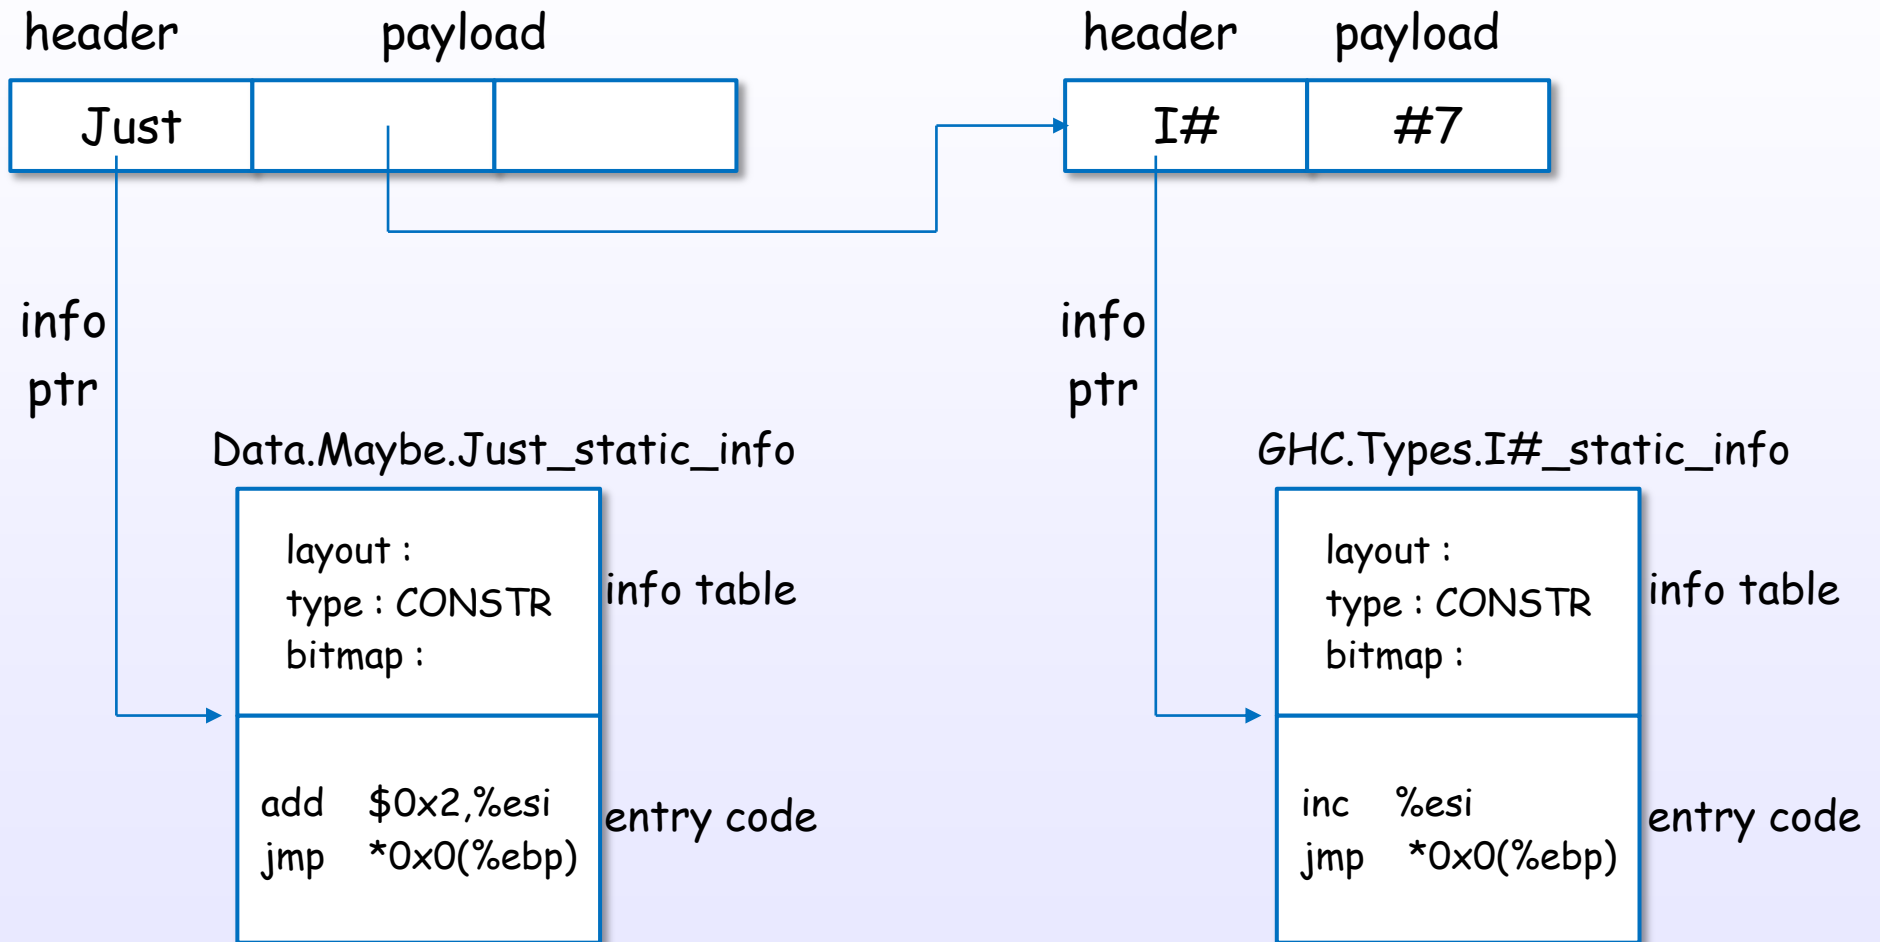
```
.data
    .align 4
.align 1
.globl __stginit_main:Example
__stginit_main:Example:
.data
    .align 4
.align 1
.globl Example.value1_closure
Example.value1_closure:
    .long GHC.Types.I#_static_info
    long 7
```

```
.section .data
    .align 4
.align 1
SMd_srt:
```

header	payload
I#	#7

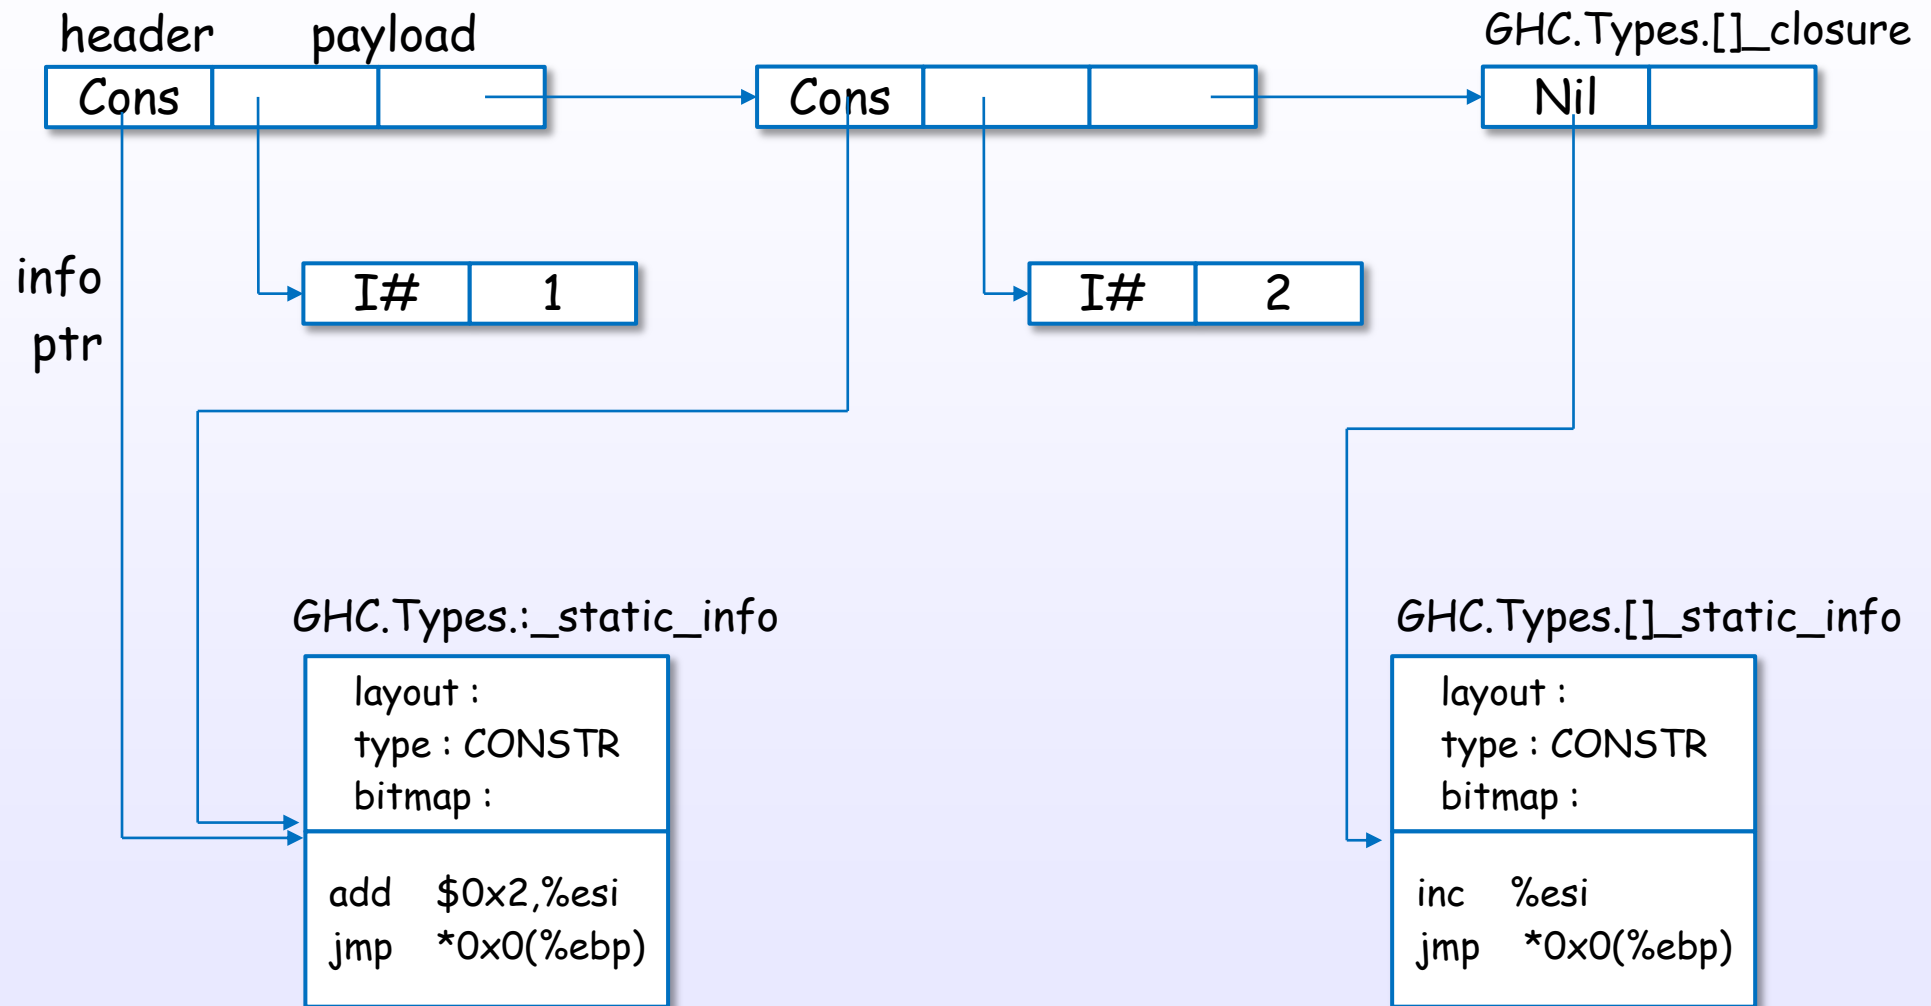
Closure examples : Maybe

Just 7 :: Maybe Int



Closure examples : List

[1, 2] :: [Int]

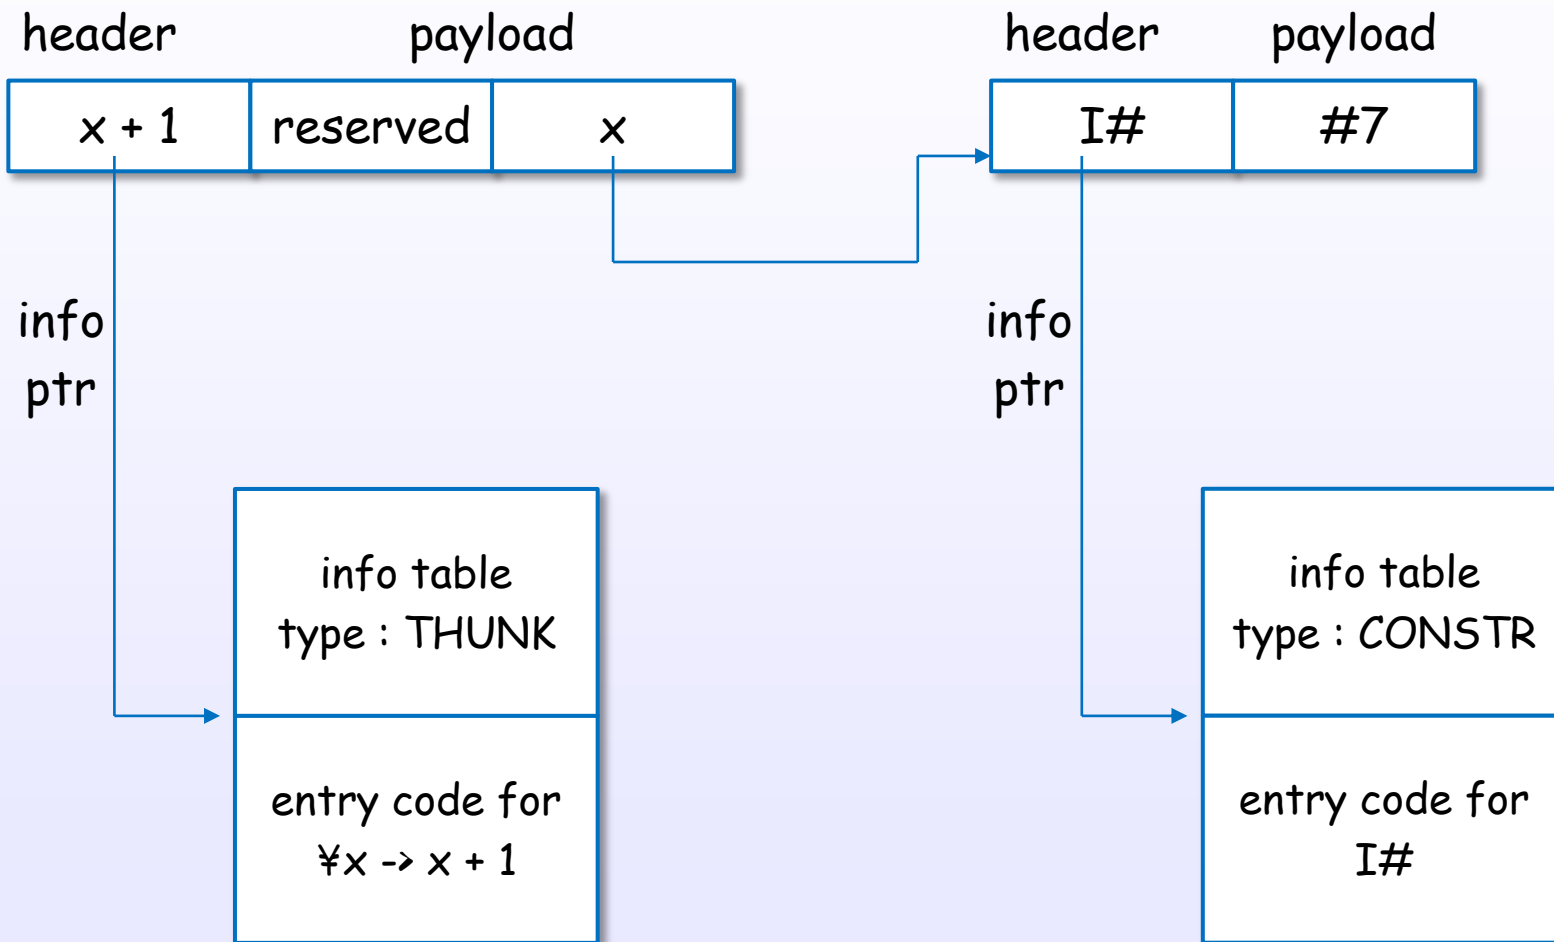


Closure examples : Thunk

"thunk"

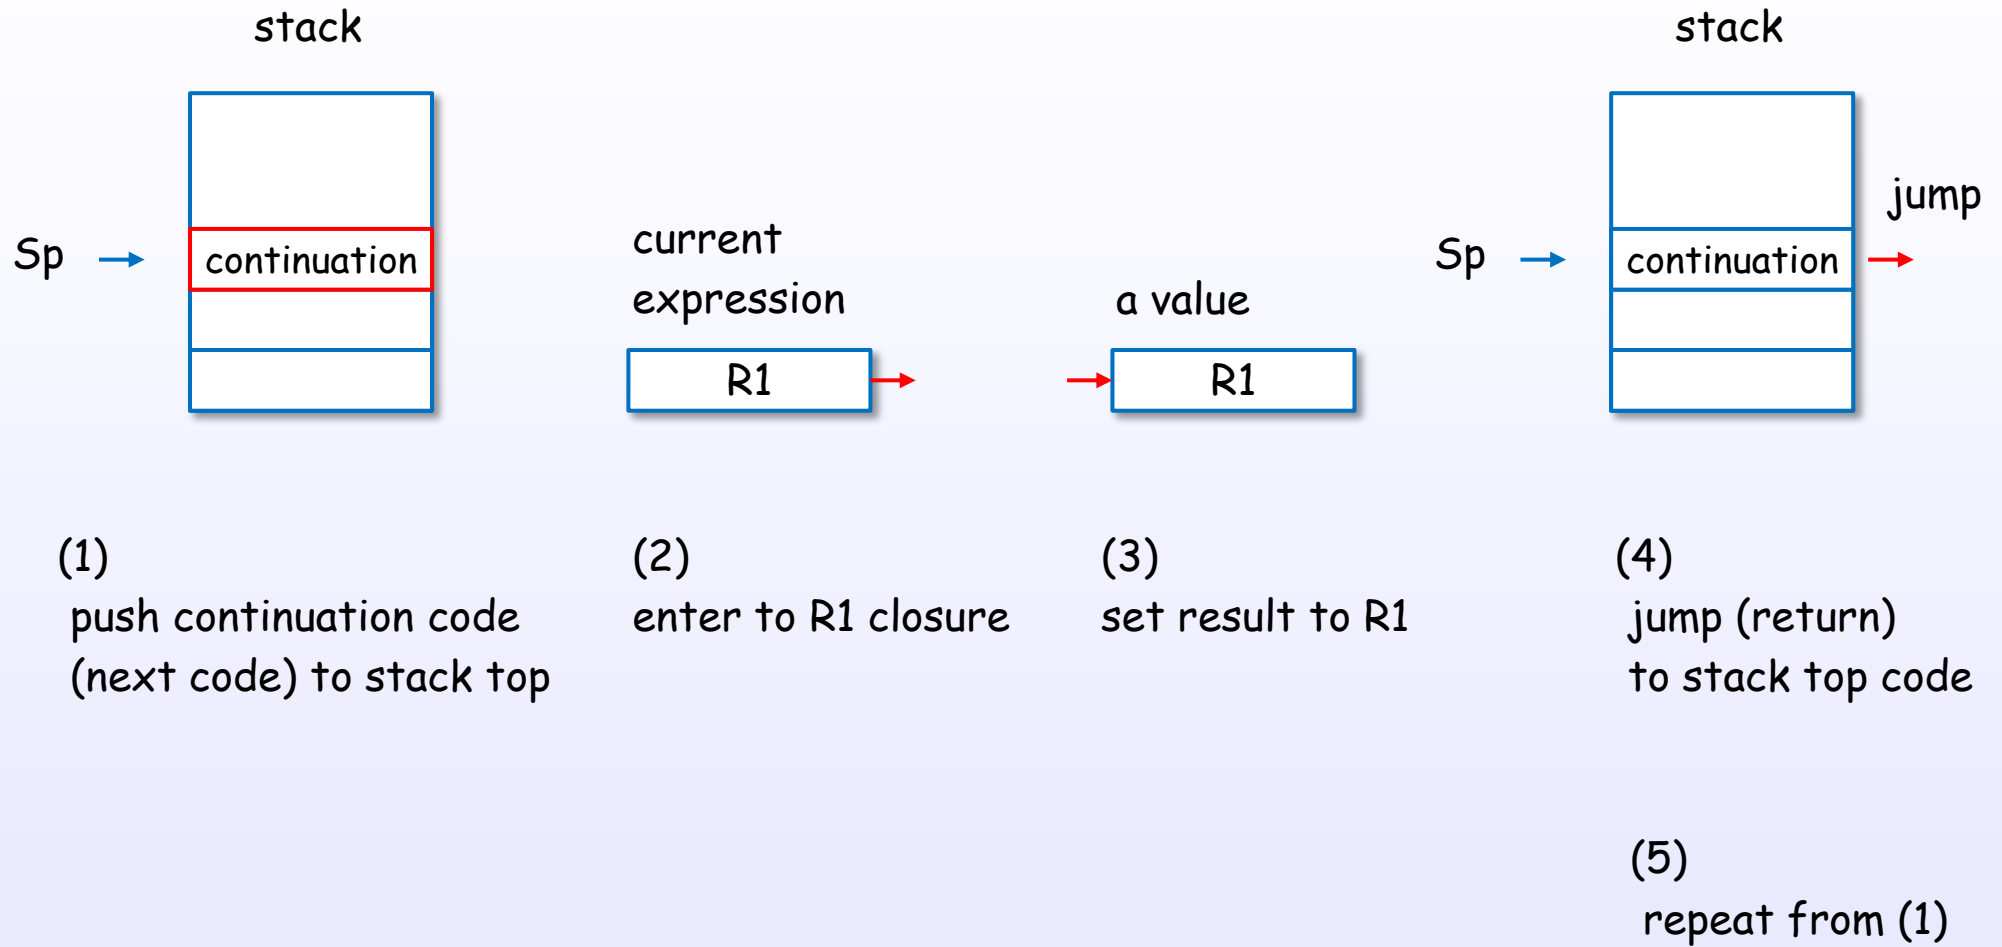
$x + 1 :: \text{Int}$

(free variable : $x = 7$)




STG-machine evaluation


STG evaluation flow



Evaluation is "enter to closure"




↓ set evaluated closure to R1

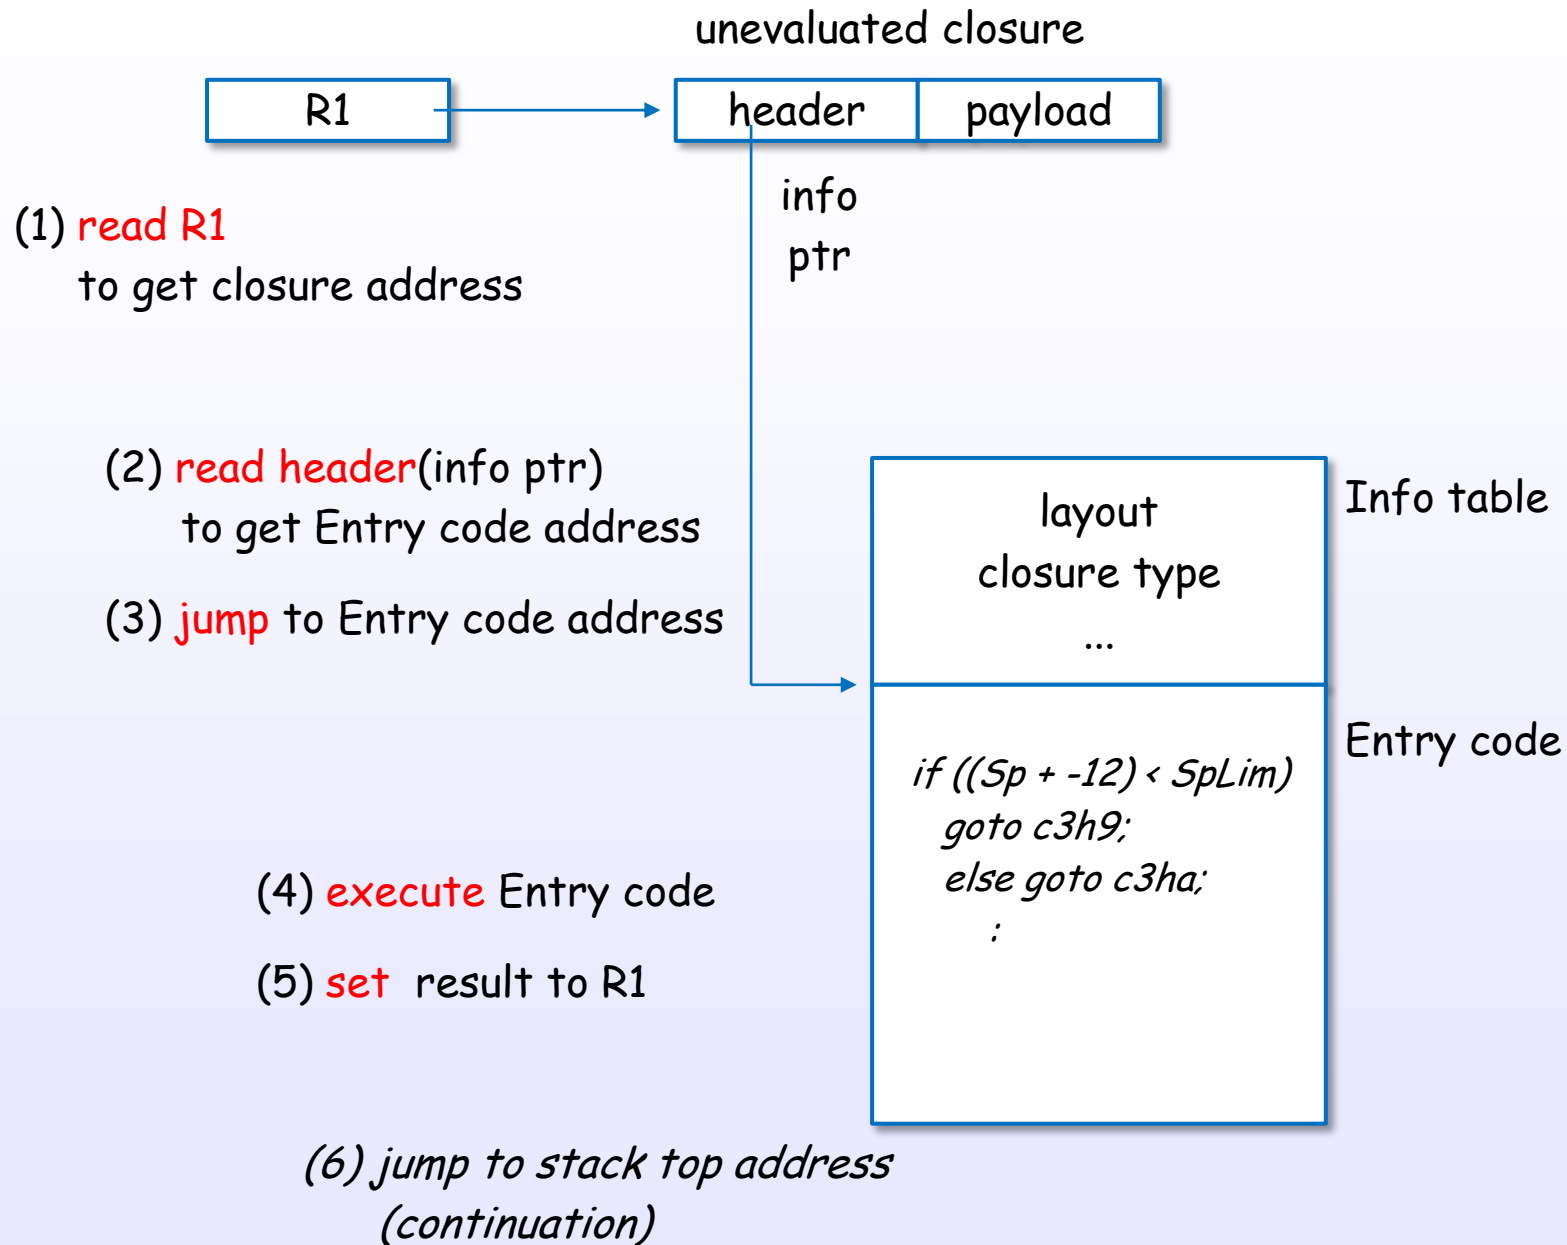


↓ enter to "closure"
STG-machine executing (evaluating) ...

stg_ap_v,




Enter to closure



Pointer tagging

Pointer tagging



pointer



... This closure is unevaluated.



... evaluated closure;
1st constructor value or evaluated.
(for instance: "Nothing")



... evaluated closure; 2nd constructor value.
(for instance: "Just xx")



... evaluated closure; 3rd constructor value.

* 32bit machine case

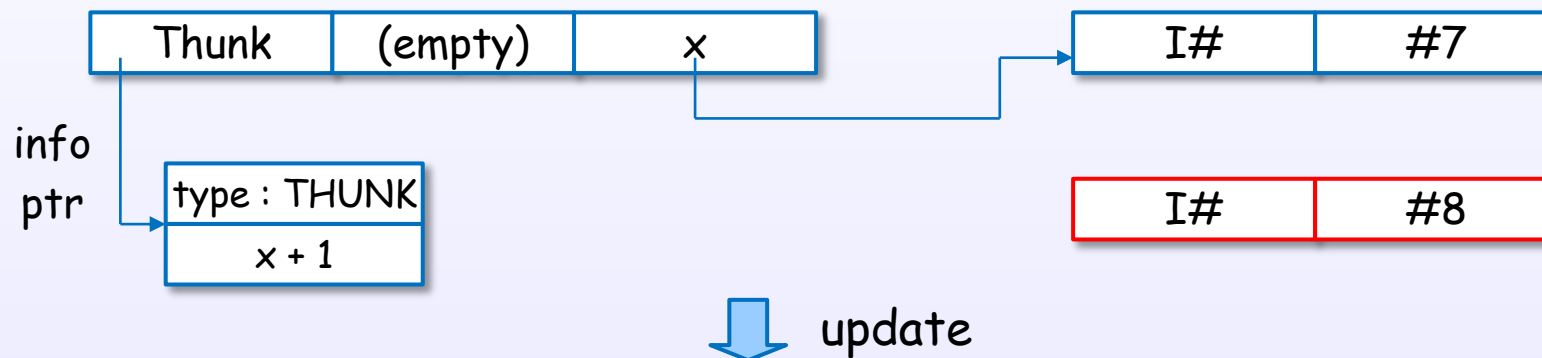

quick judgment!

check only pointer's lower bits without evaluating the closure.

Think and update

Thunk and update

"thunk" $x + 1 :: \text{Int}$ (free variable : $x = 7$)




GC (eliminate Indirect)


lock free

Allocate and free(GC) heap objects

Allocate heap objects

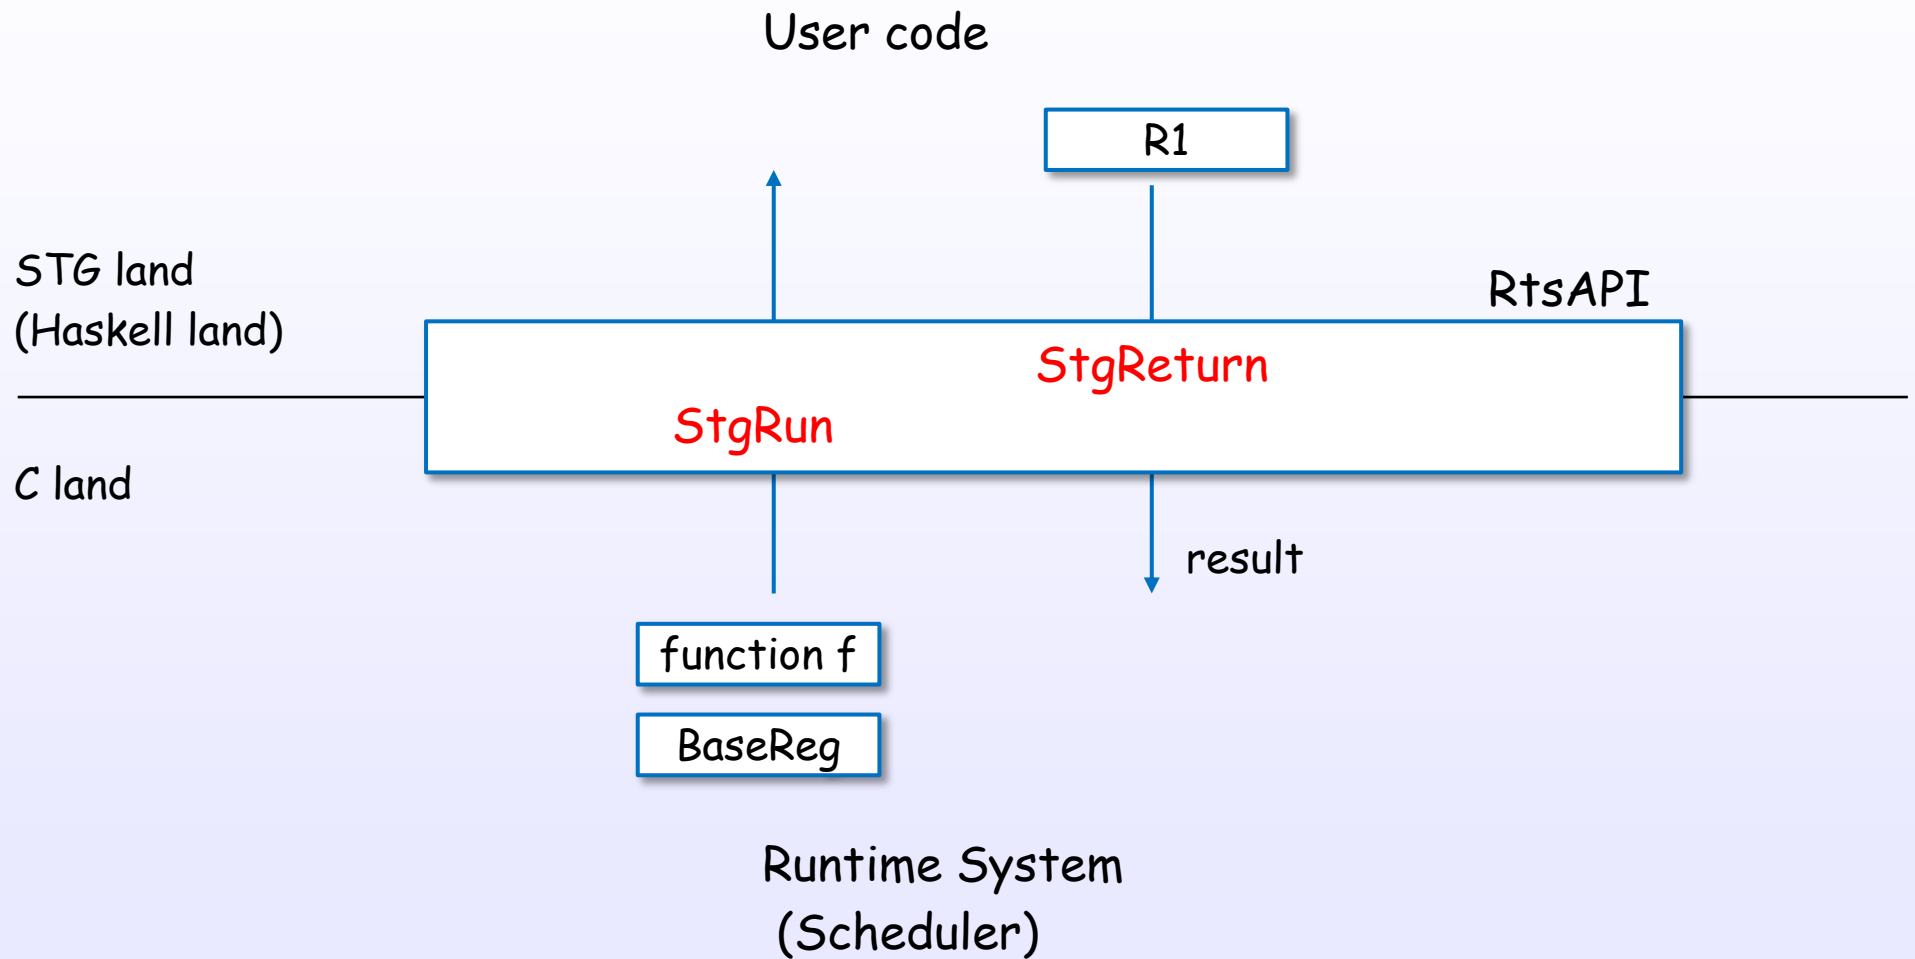


free and collection heap objects



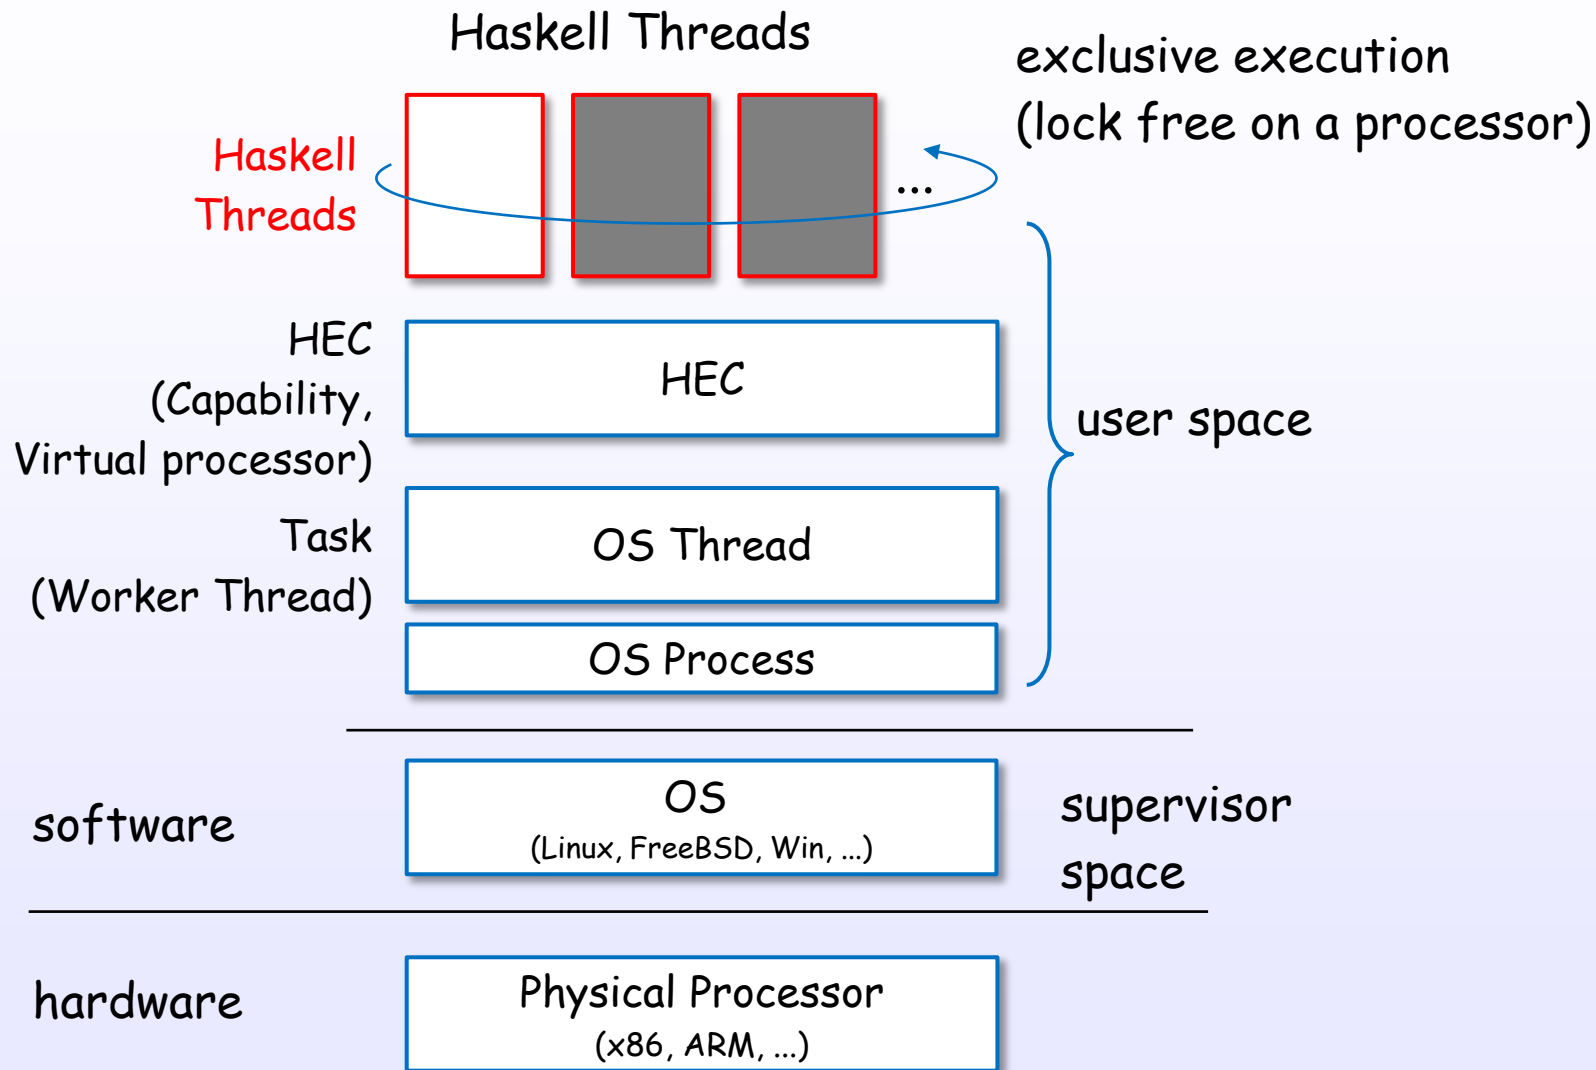
STG - C land interface

STG (Haskell) land - C land interface

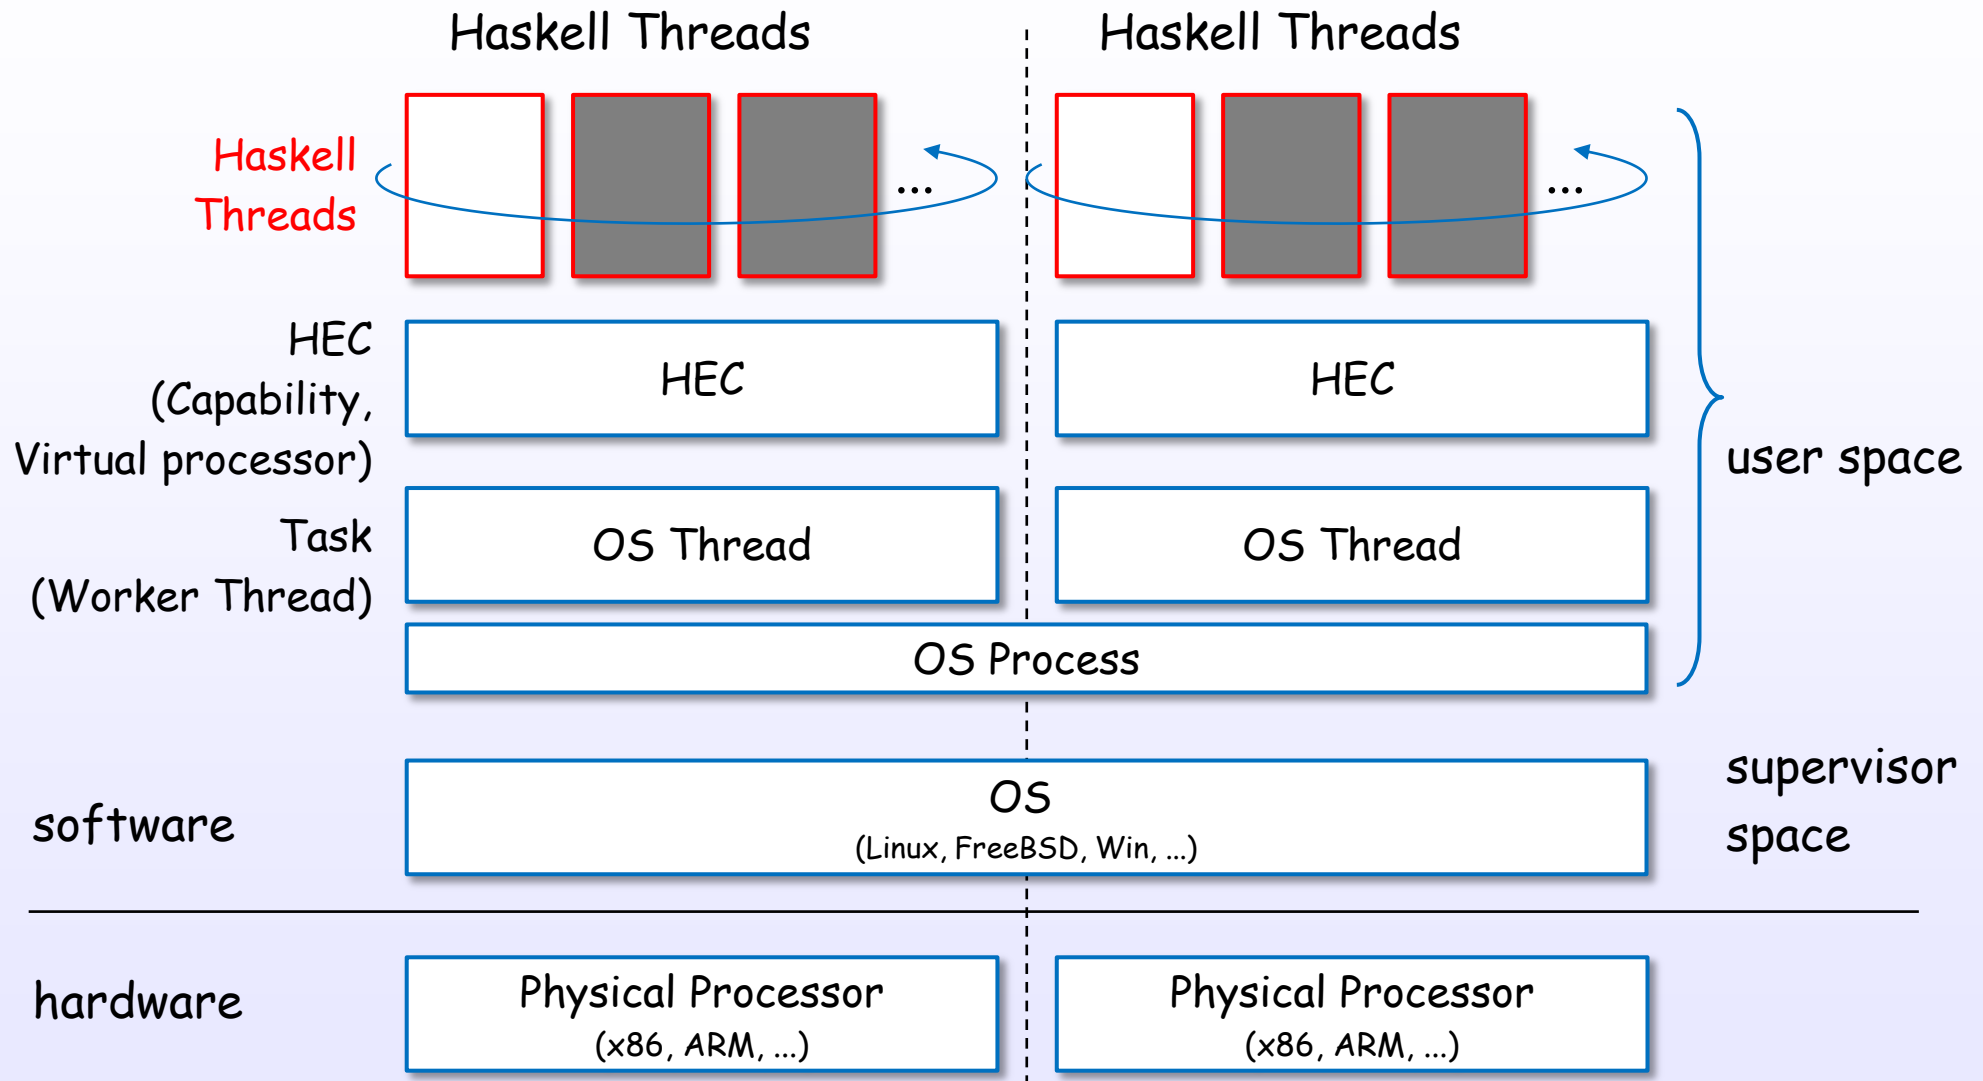


Thread

Thread layer (single core)



Thread layer (multi core)



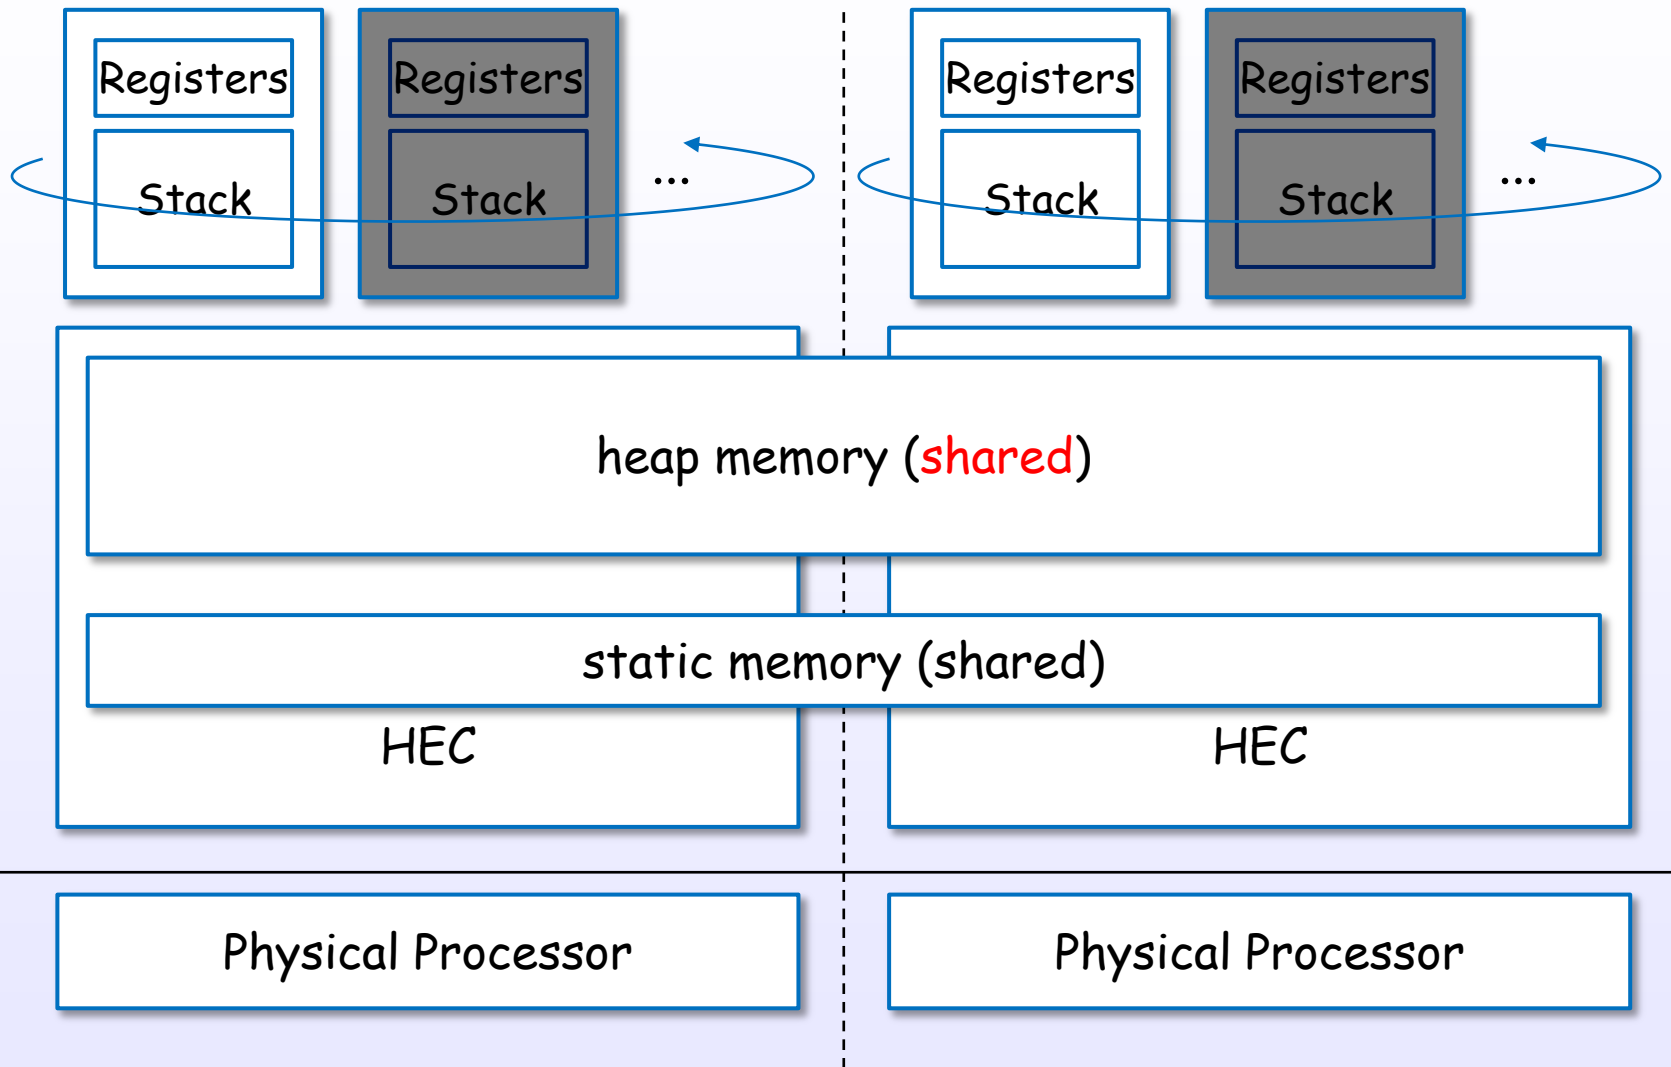
*Threaded option case (ghc -threaded)

References : [4], [7], [8], [13], [C17], [C11], [18], [S17], [S16], [S23], [S22], [S14]

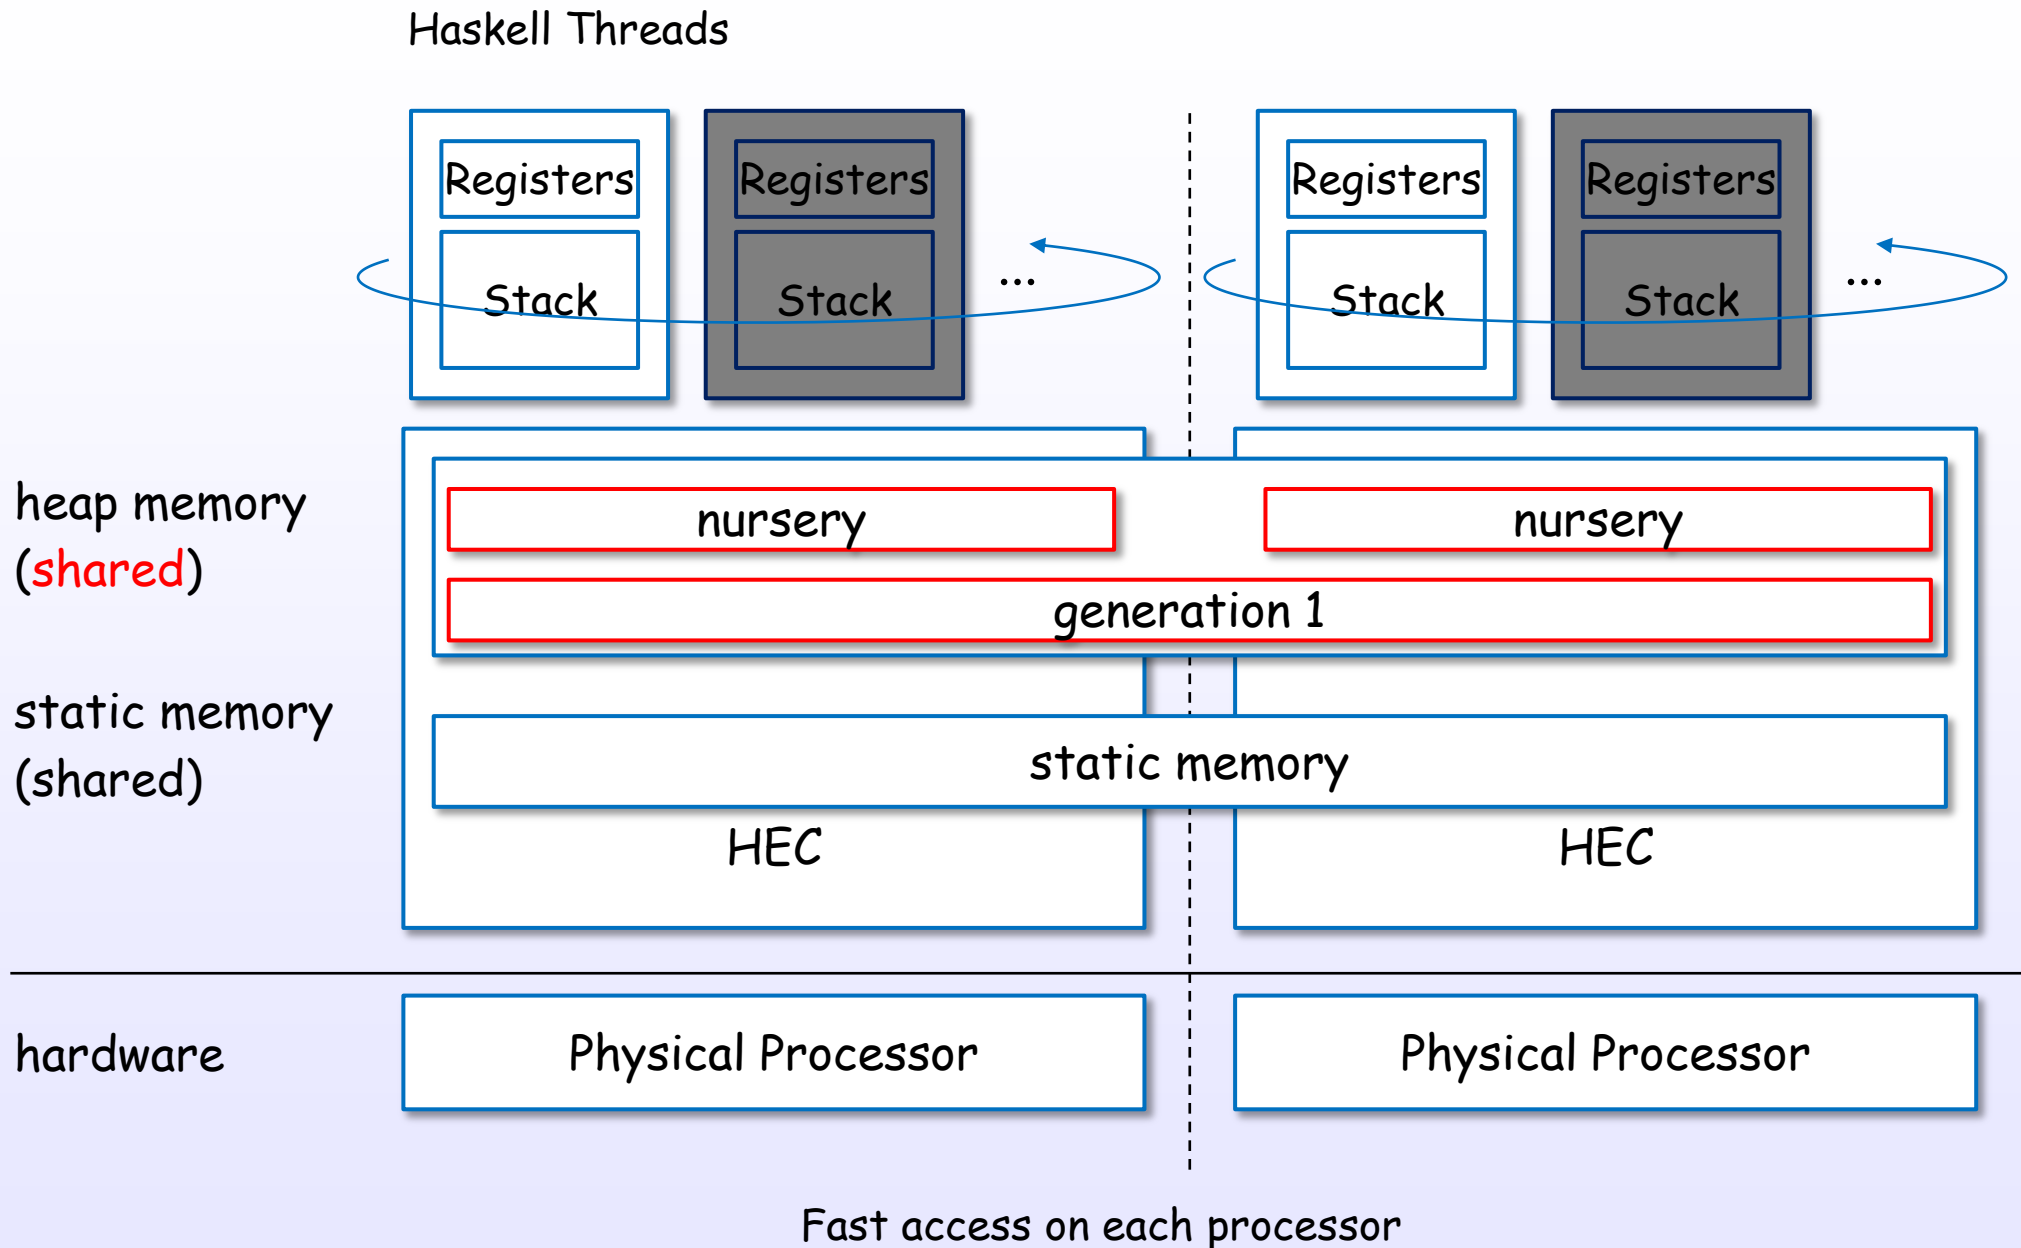
Heap and Threads

Threads and shared heap

Haskell Threads

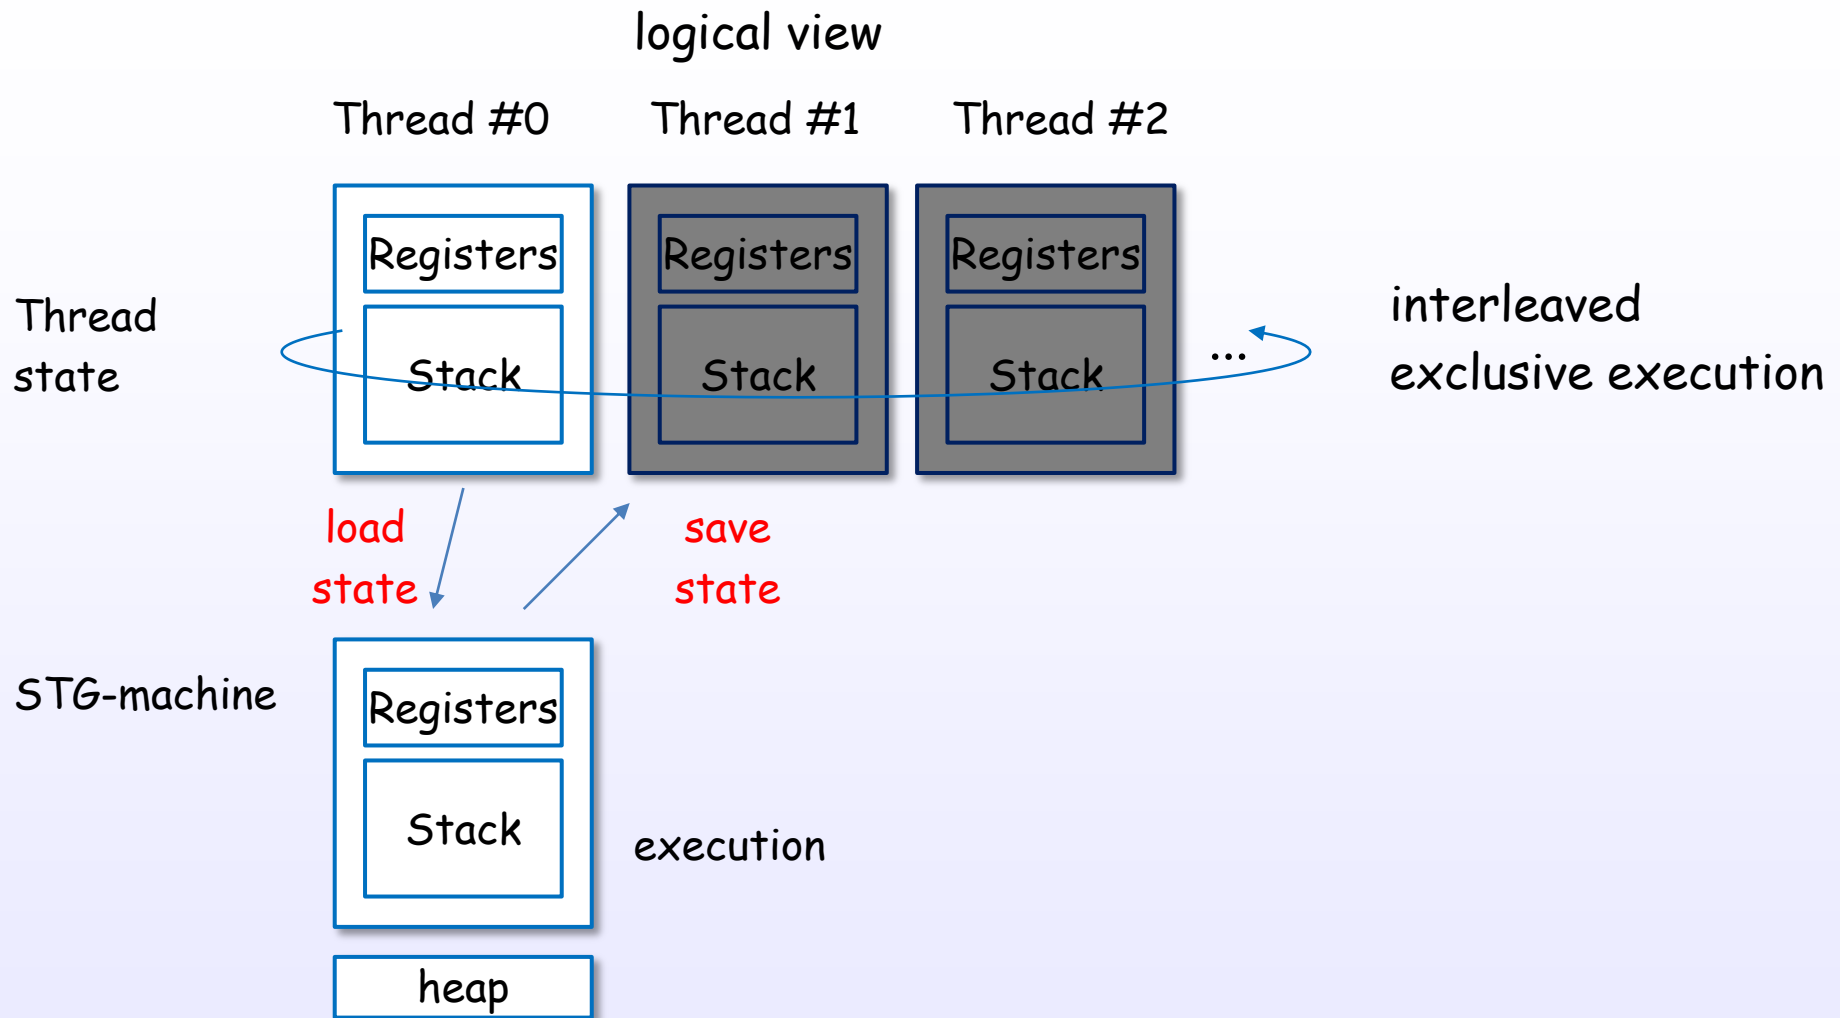


Local heap area

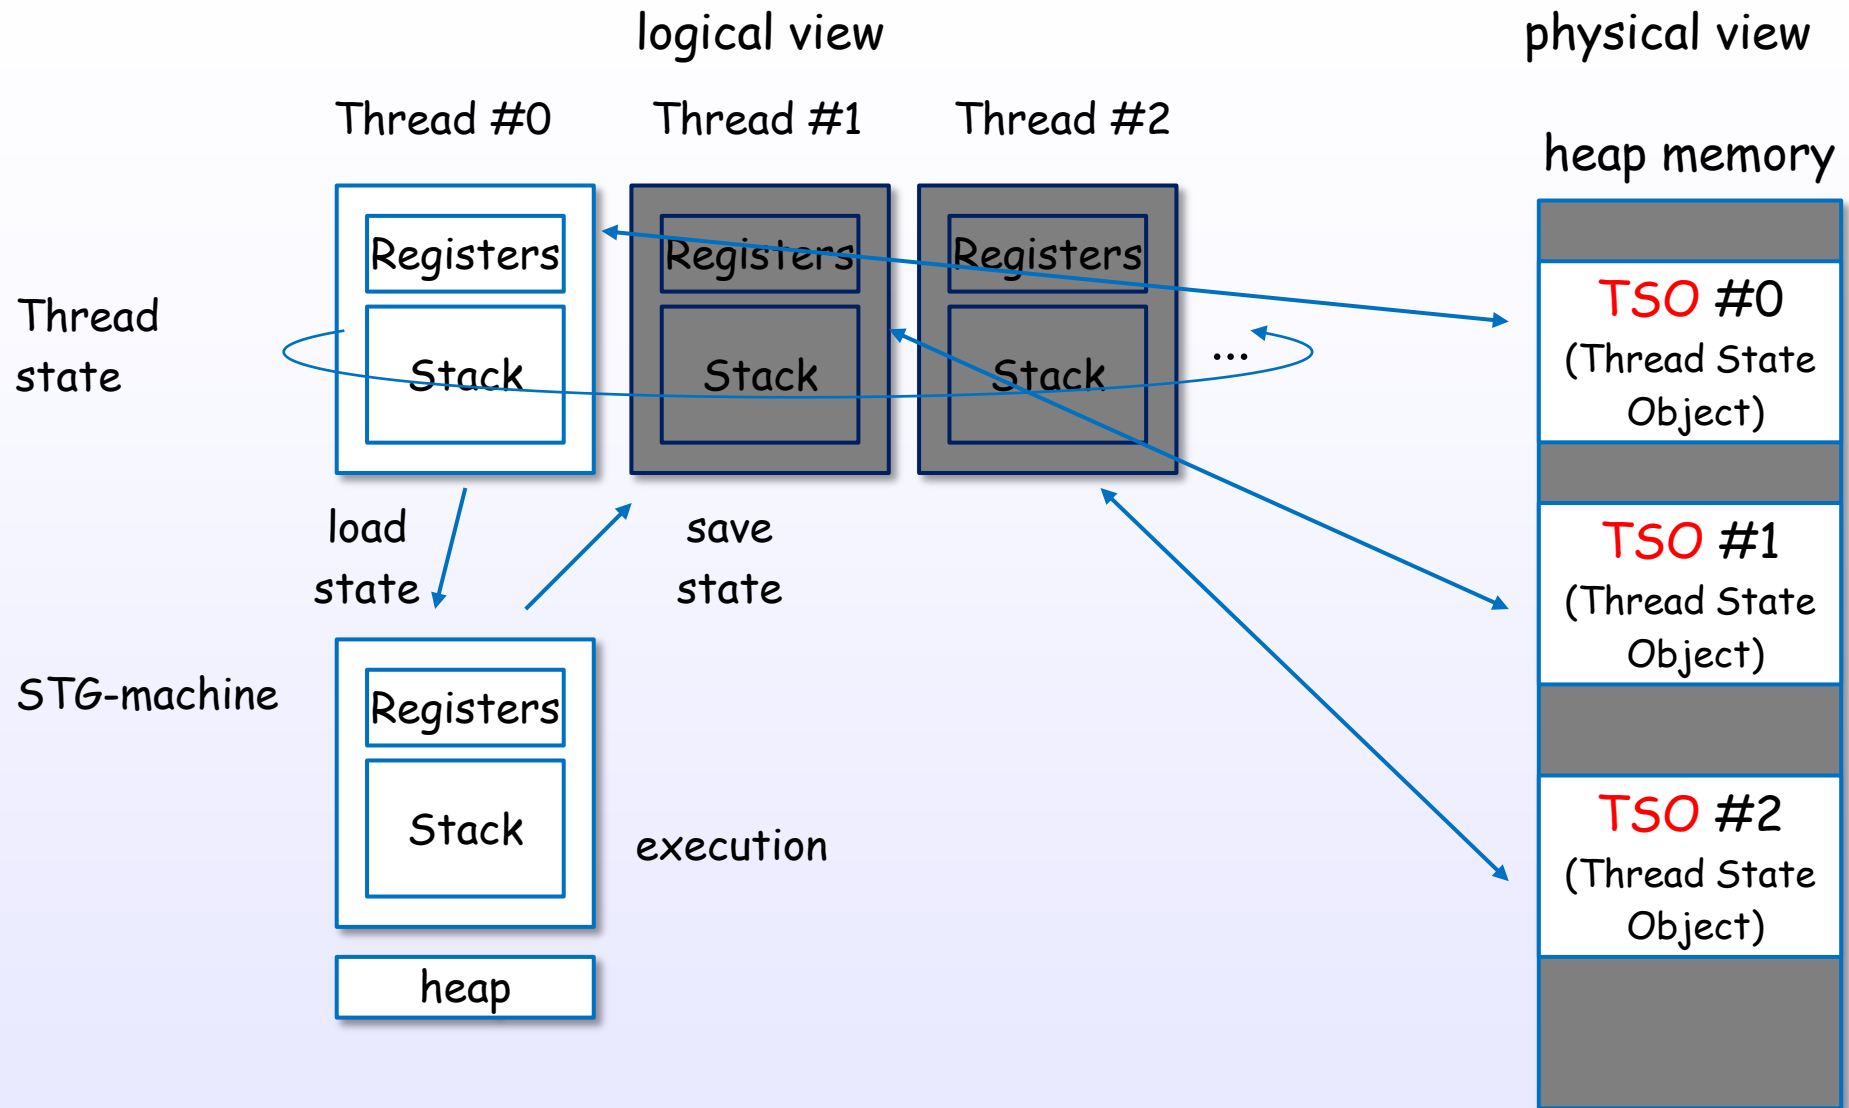


Thread context switch

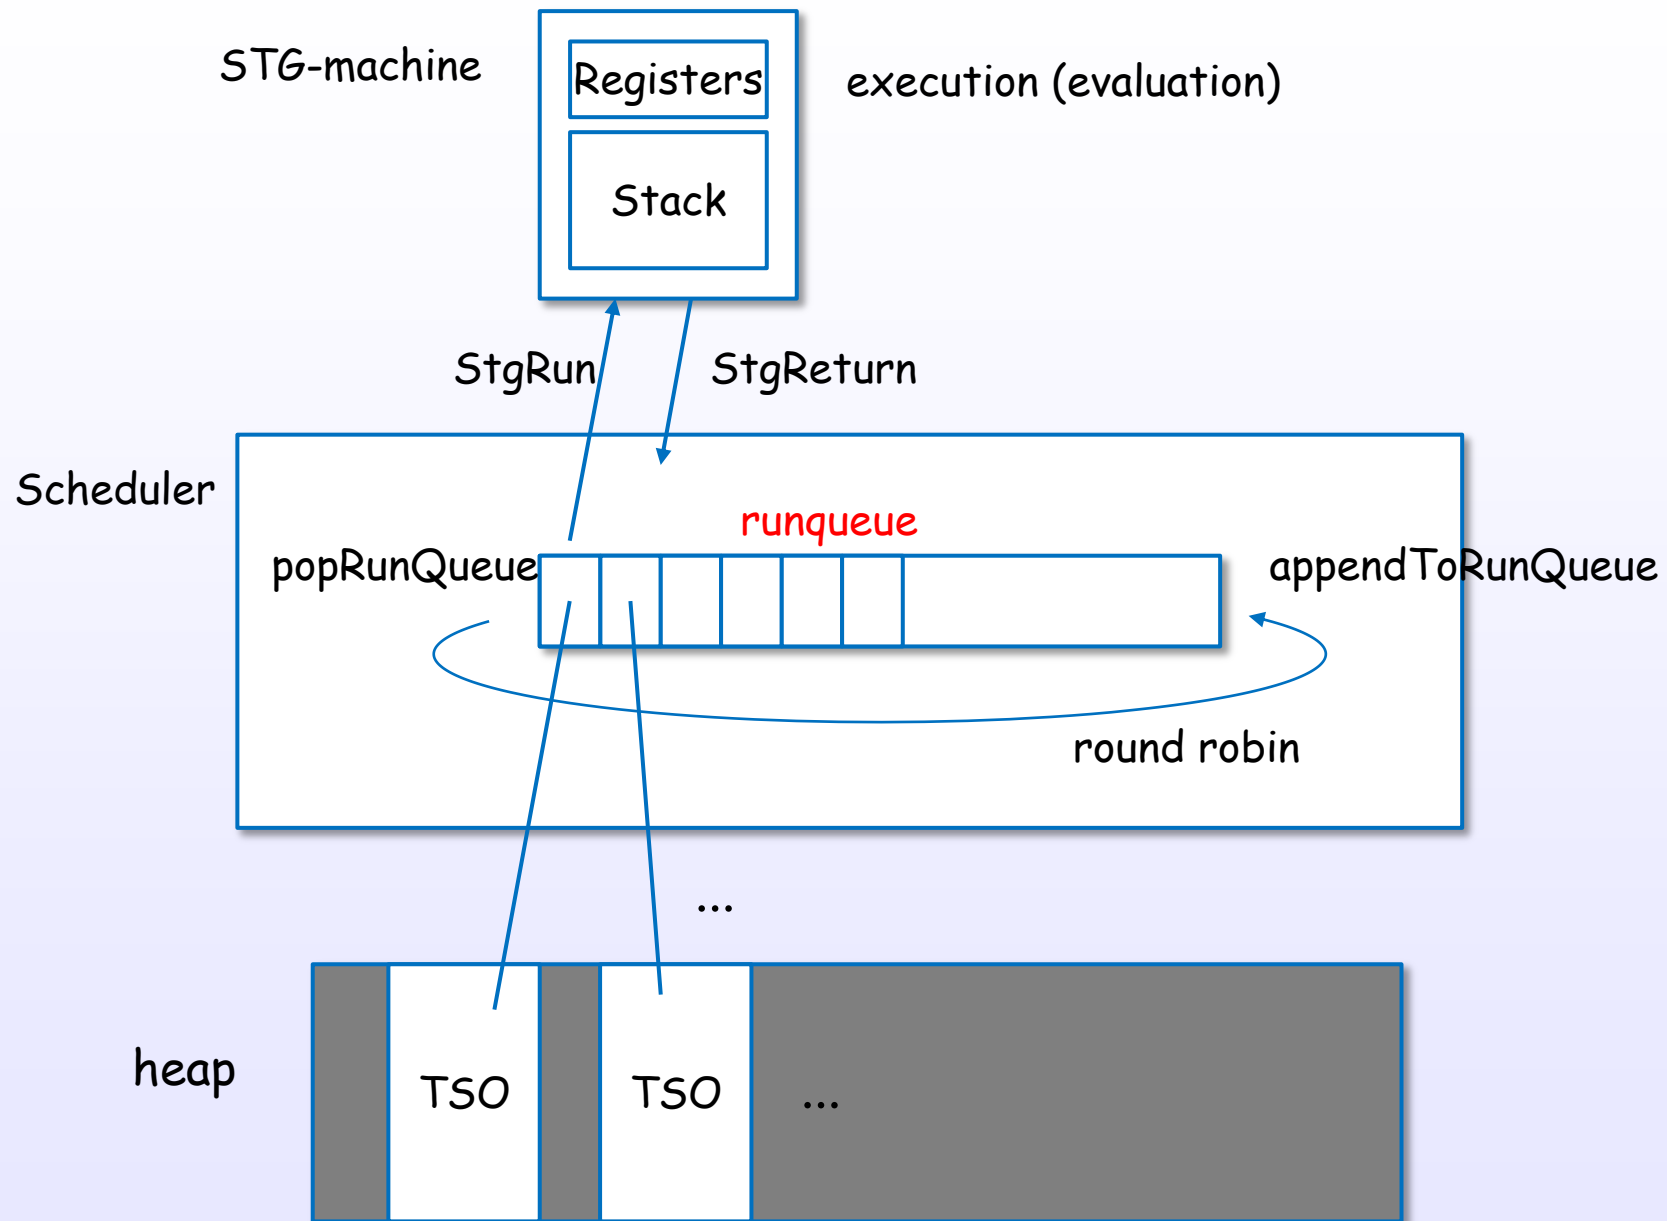
Threads and context switch



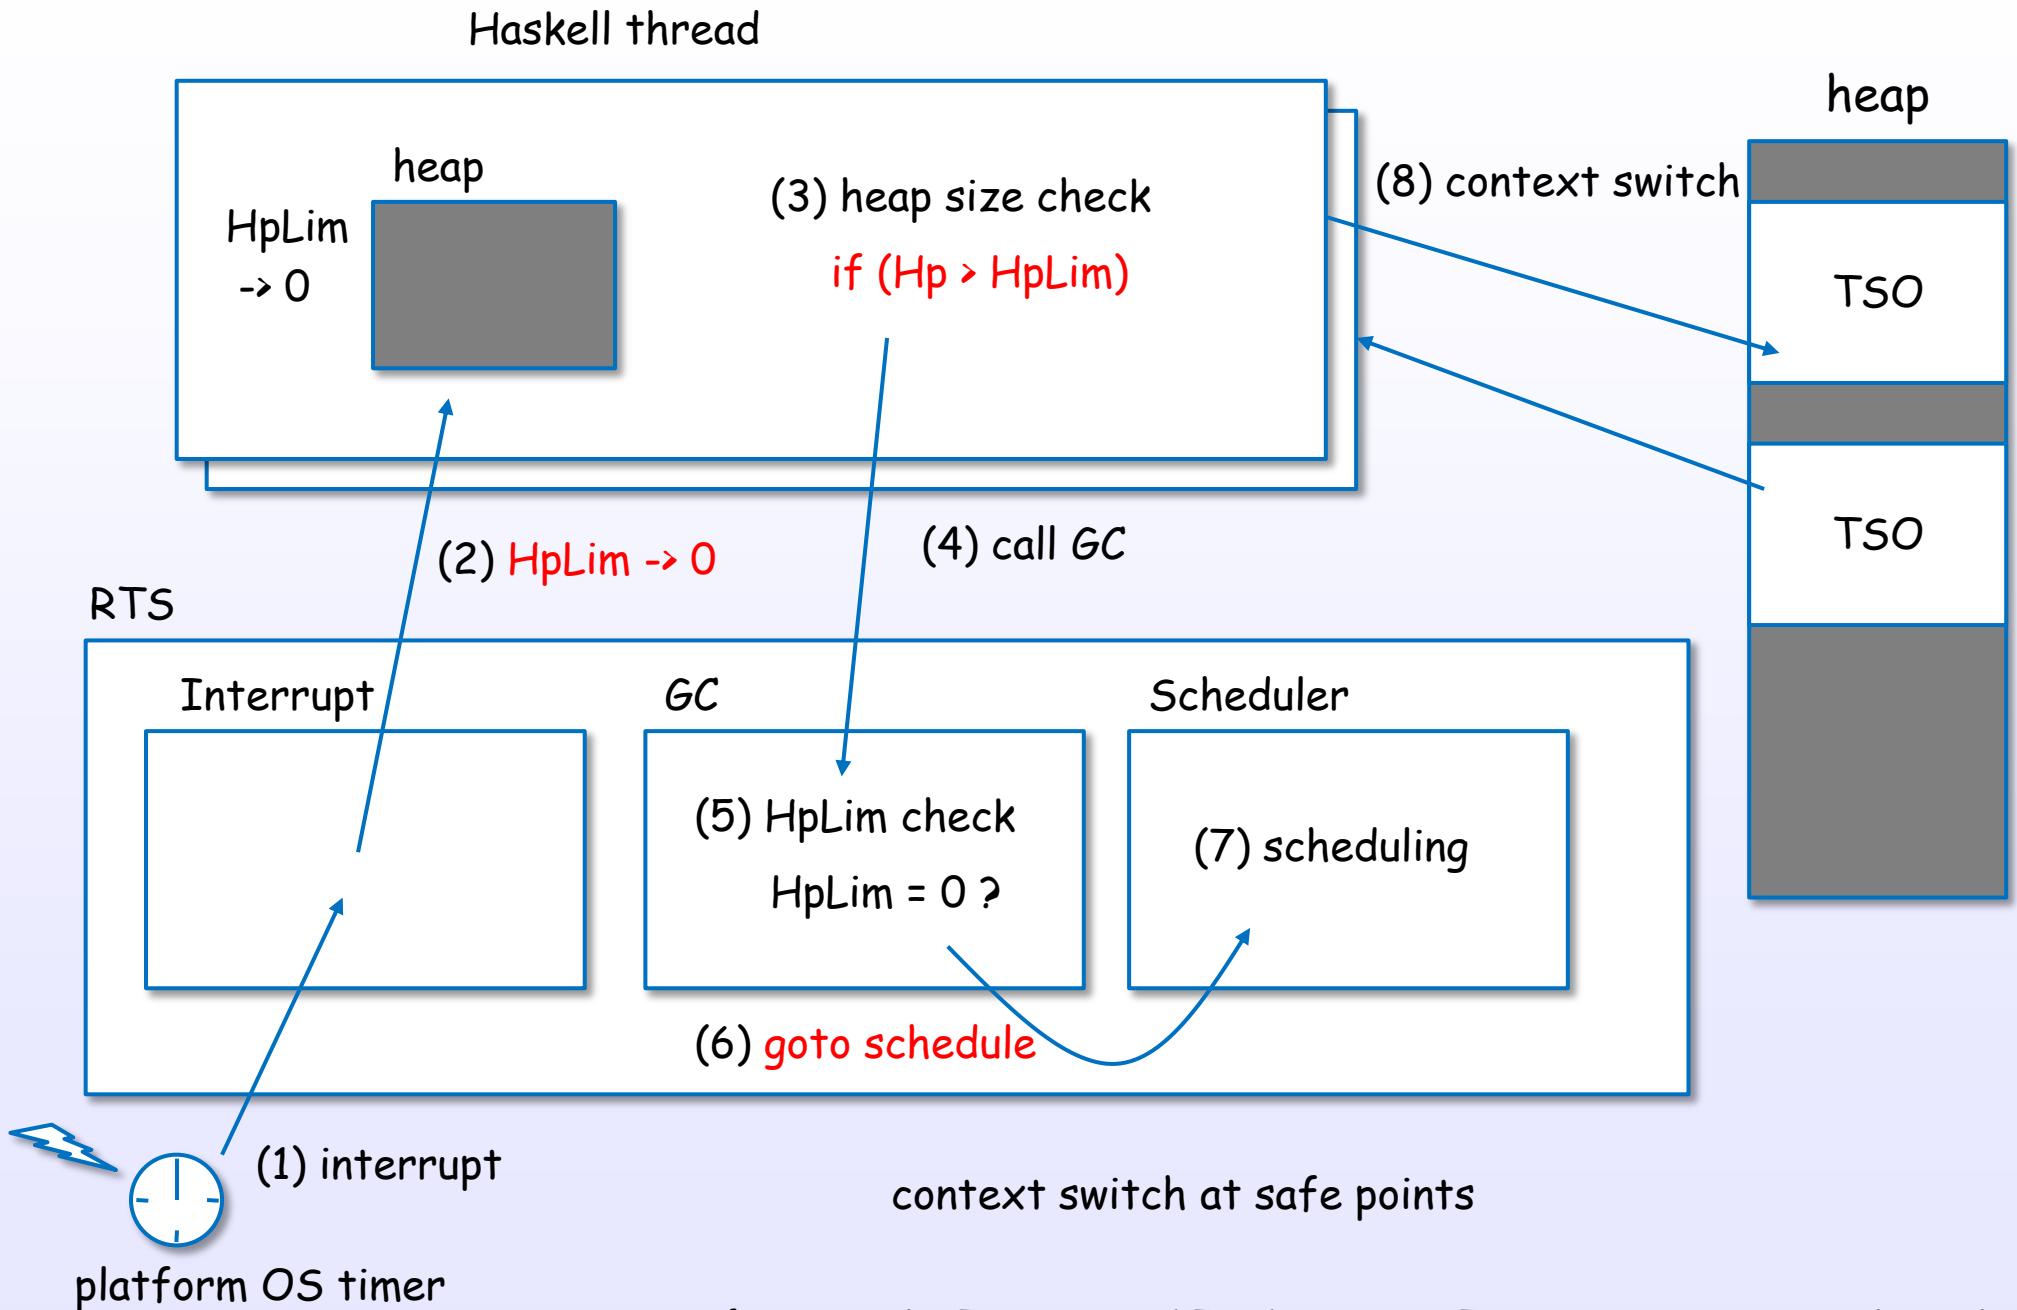
Threads and TSOs



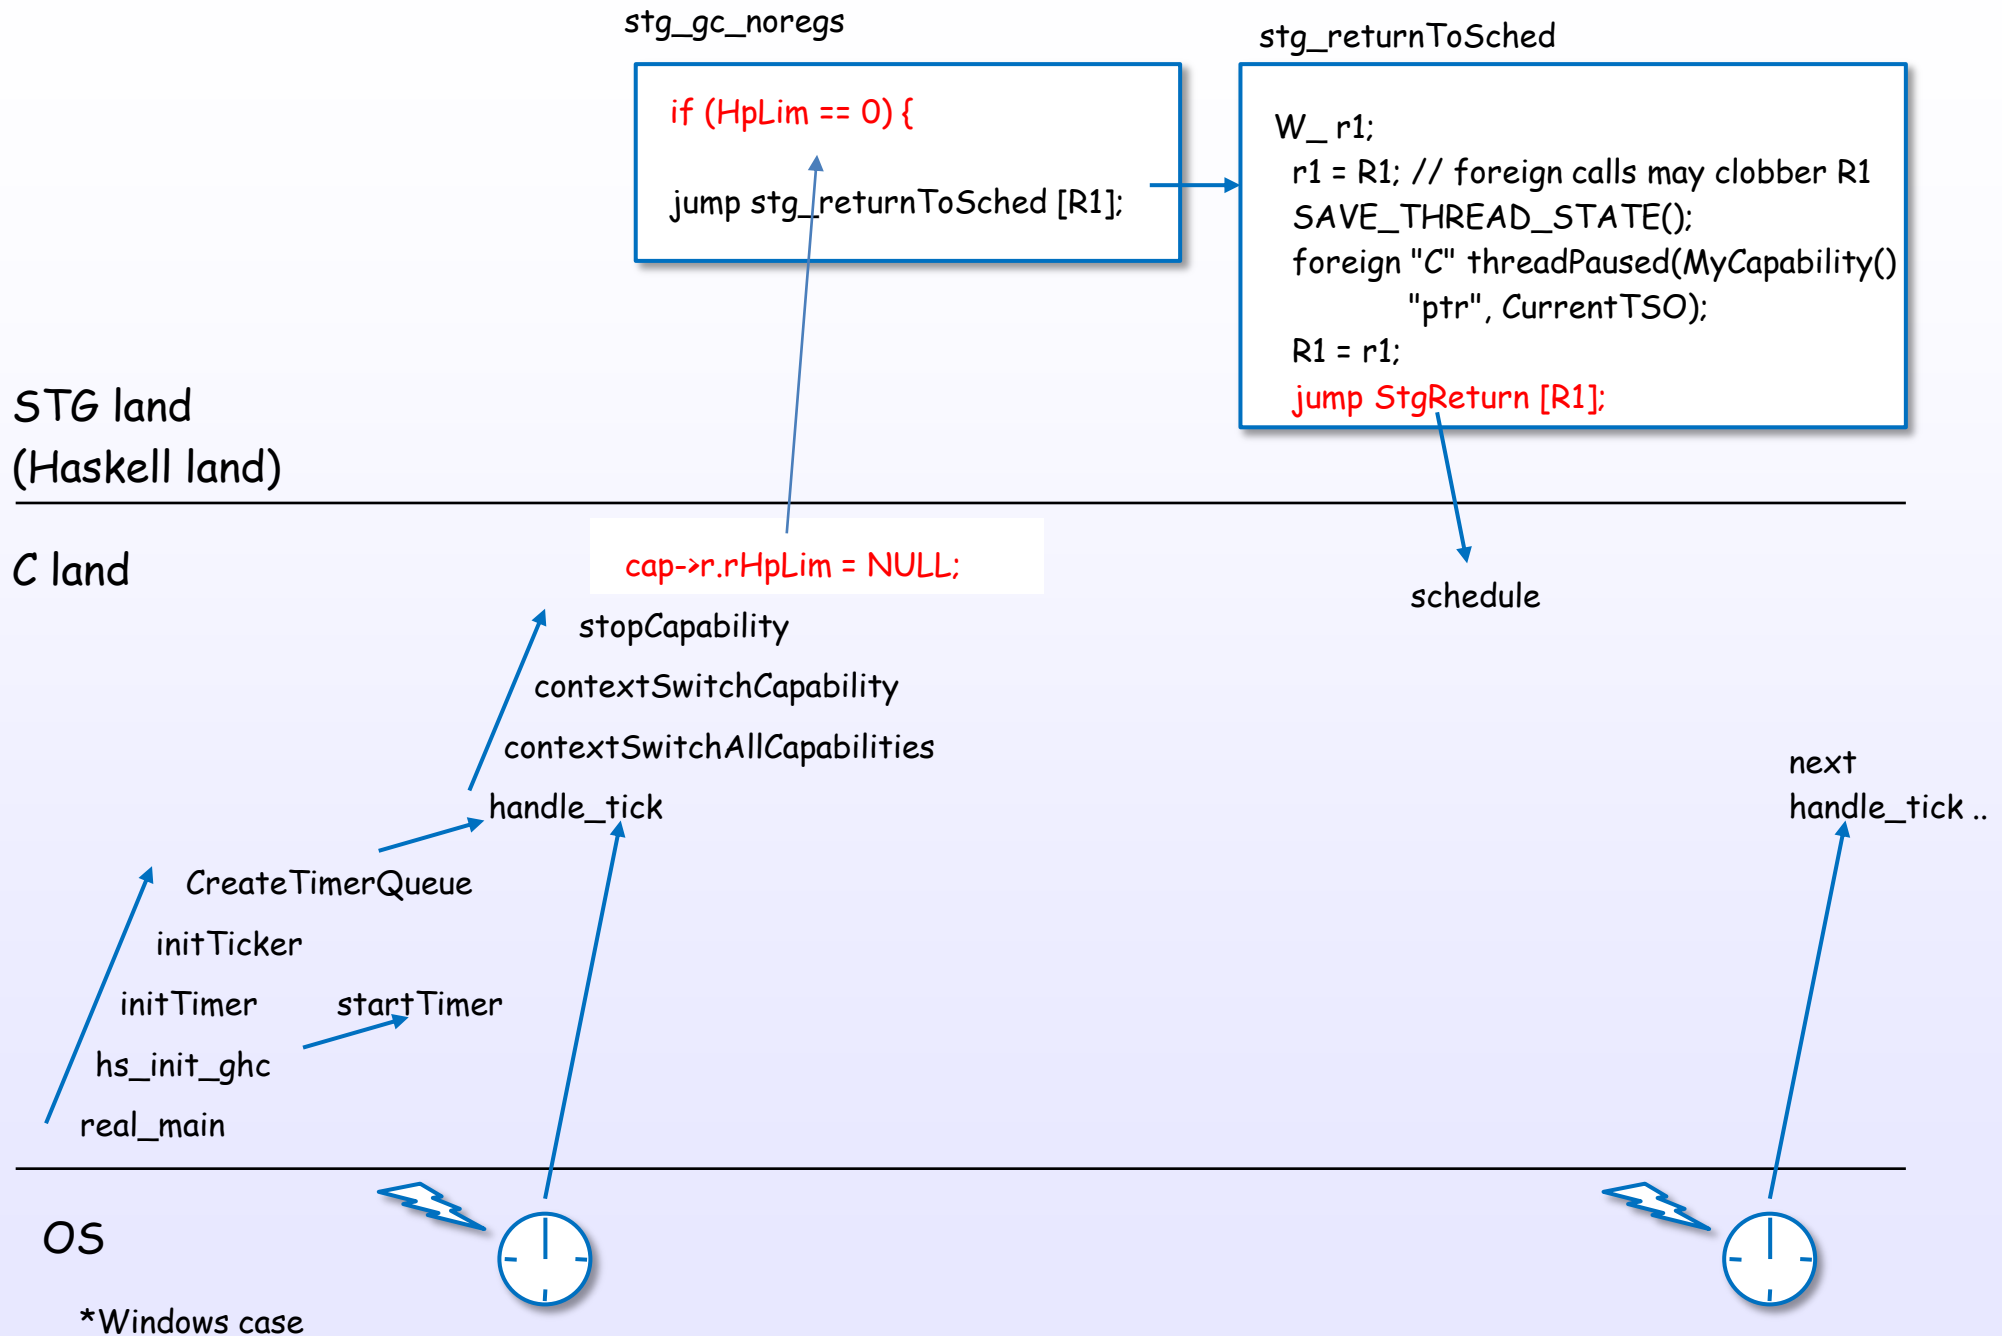
Scheduling threads



Context switch flow



Context switch flow (code)



Creating main thread and forkIO

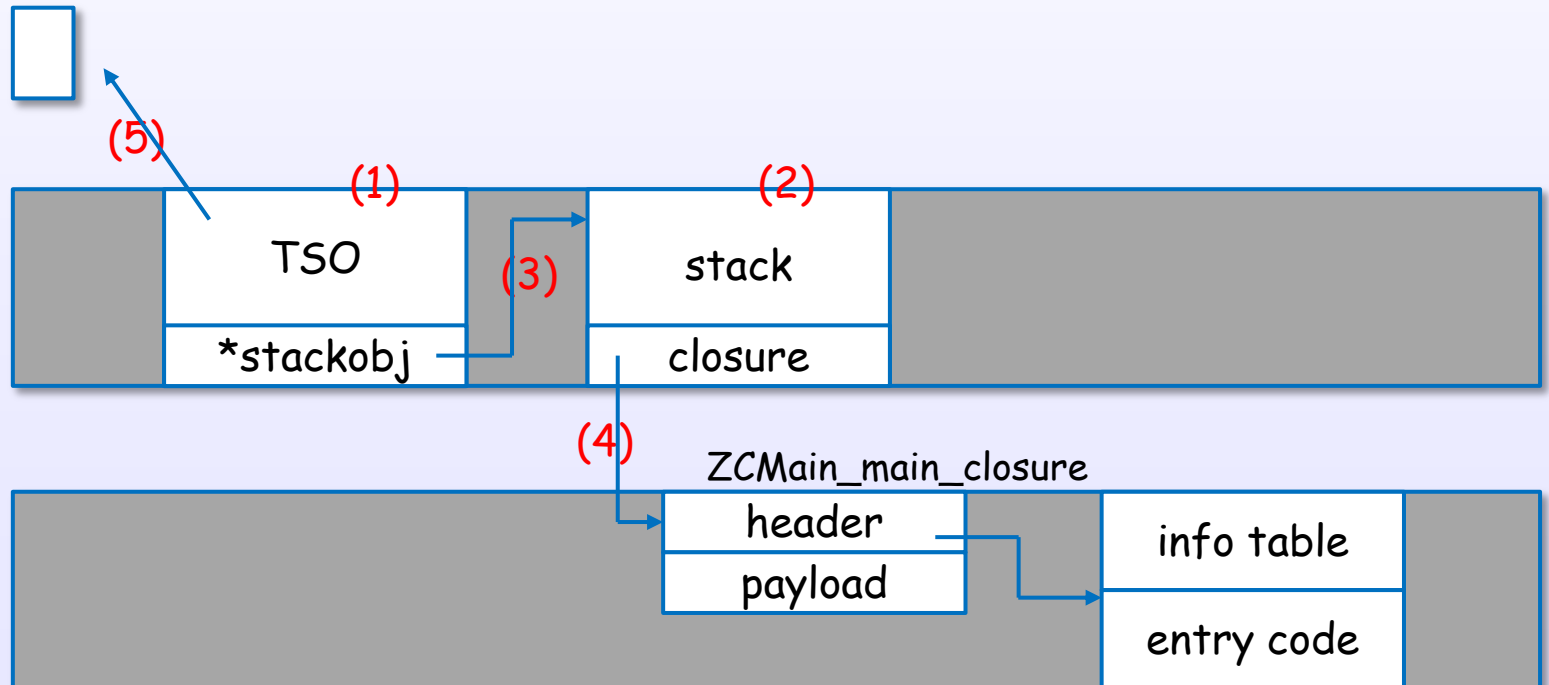
Create main thread

C land

Runtime system bootstrap code [rts/RtsAPI.c]

```
rts_evalLazyIO
  createIOThread
    createThread ... (1), (2), (3)
    pushClosure ... (4)
  scheduleWaitThread
    appendToRunQueue ... (5)
```

scheduler
runqueue



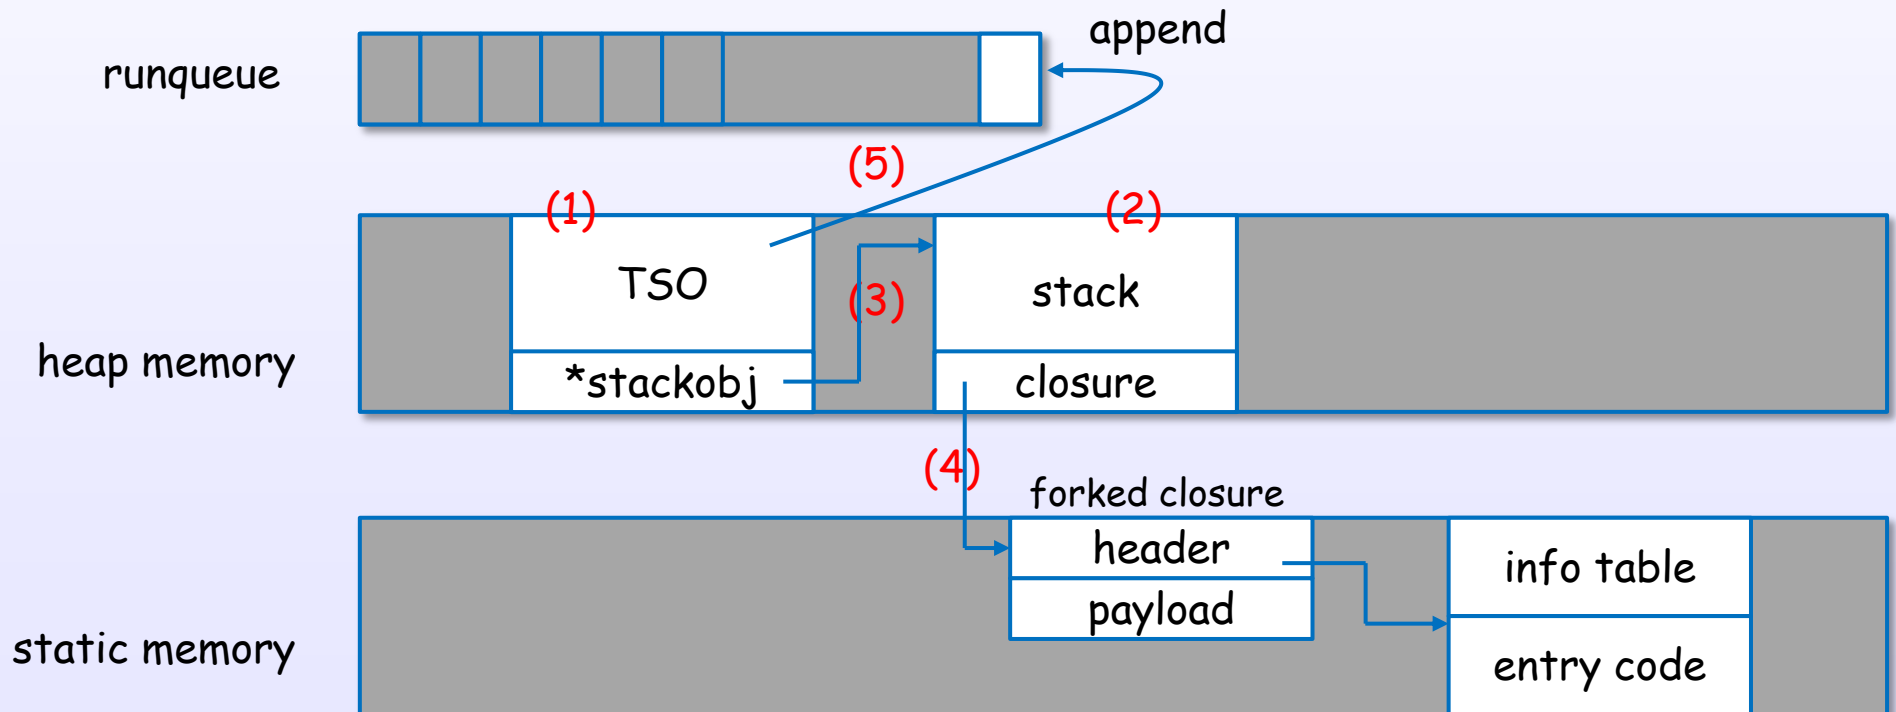
Create sub thread by forkIO

Haskell Threads

```
forkIO
  stg_forkzh
    ccall createIOThread ... (1), (2), (3), (4)
    ccall scheduleThread ... (5)
```

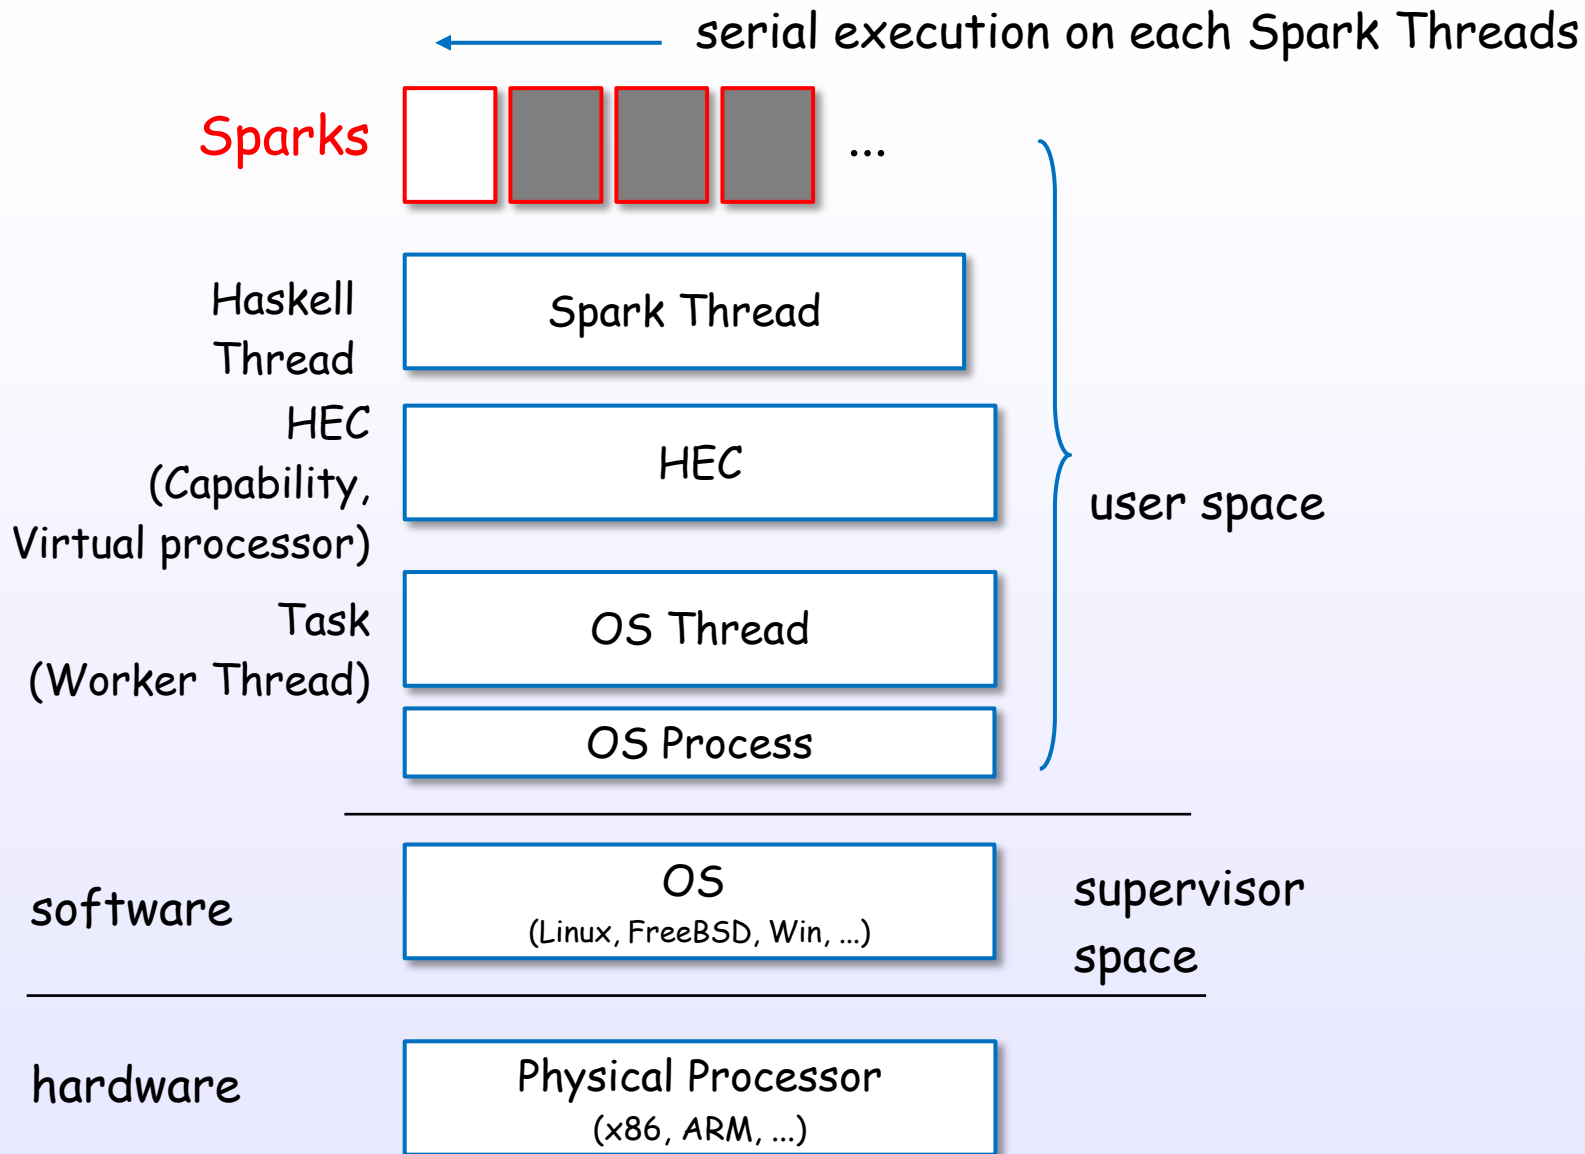
STG land
(Haskell land)

C land



Spark

Spark layer

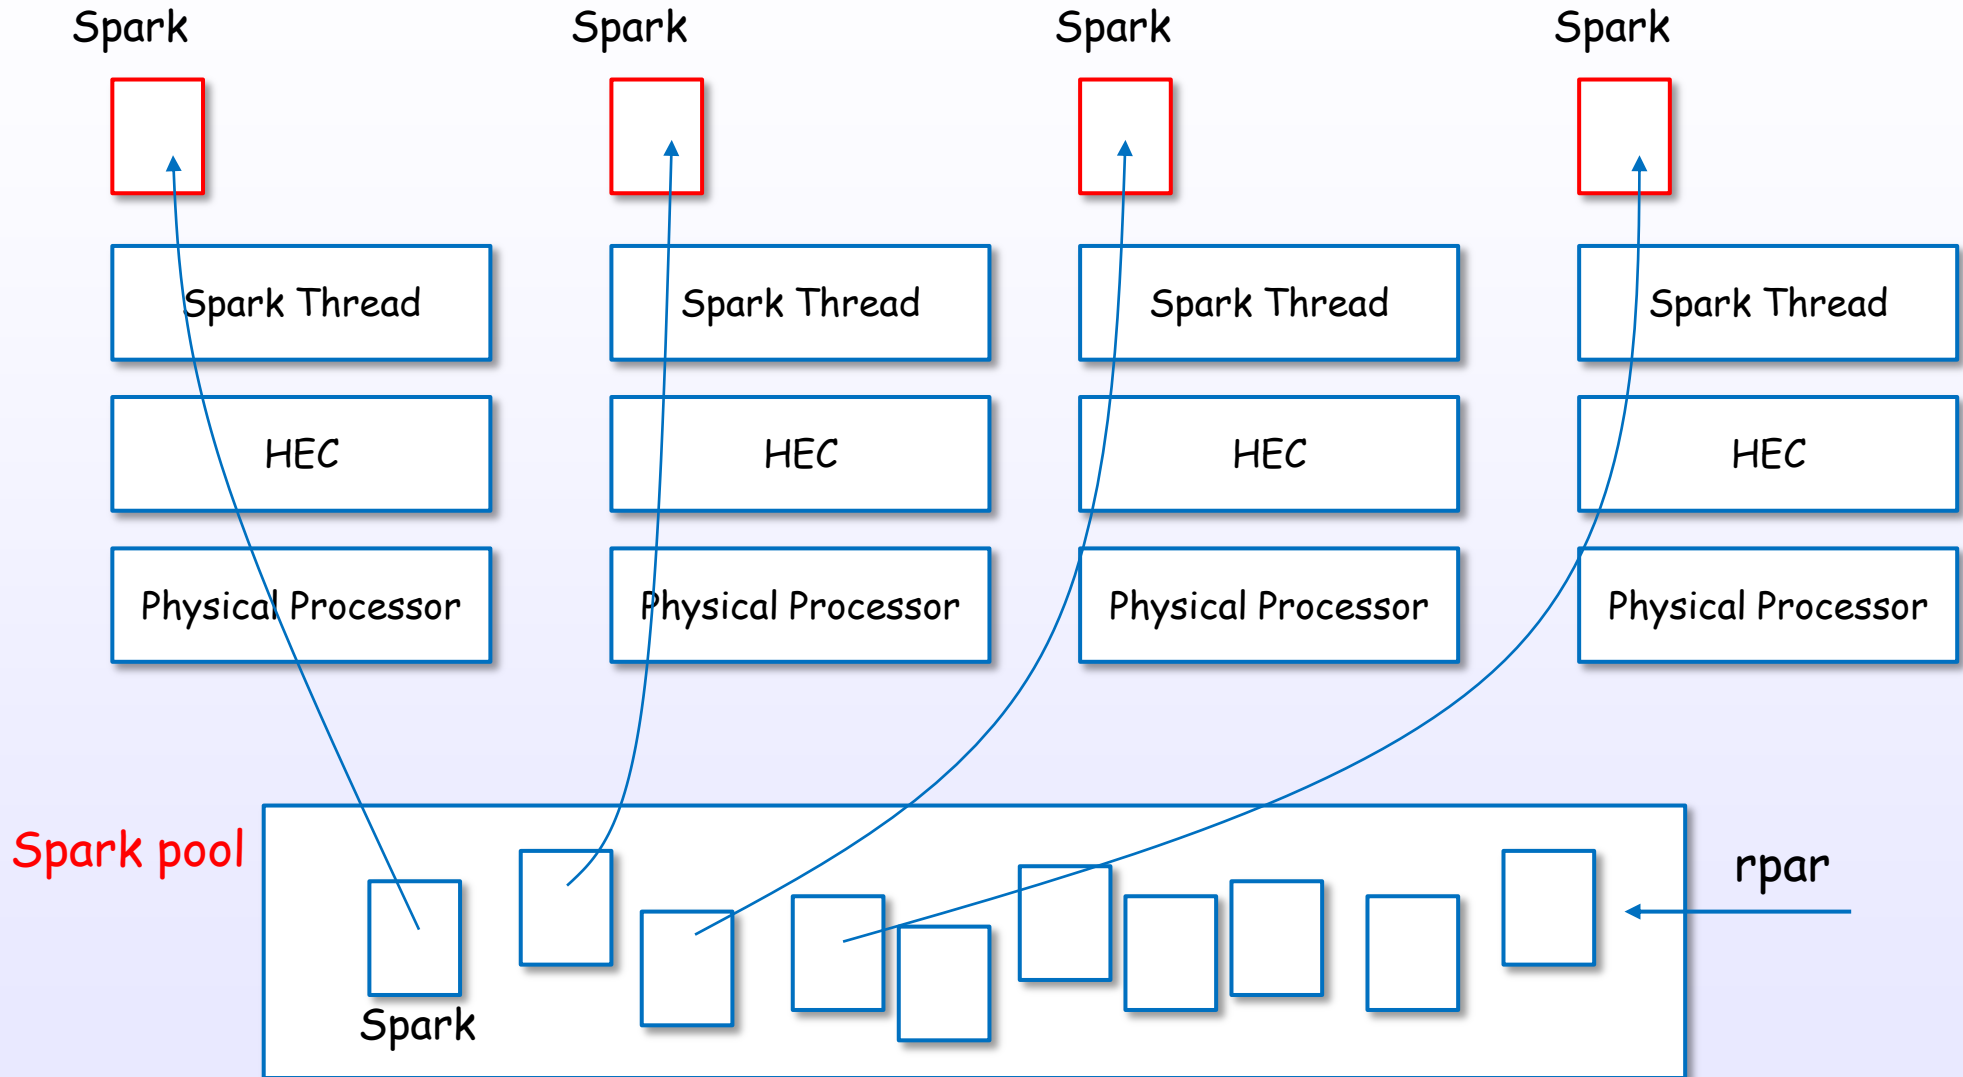


Spark Threads are generated on idle HECs.

References : [C17], [18], [S17], [S26], [S27], [S32], [S12]

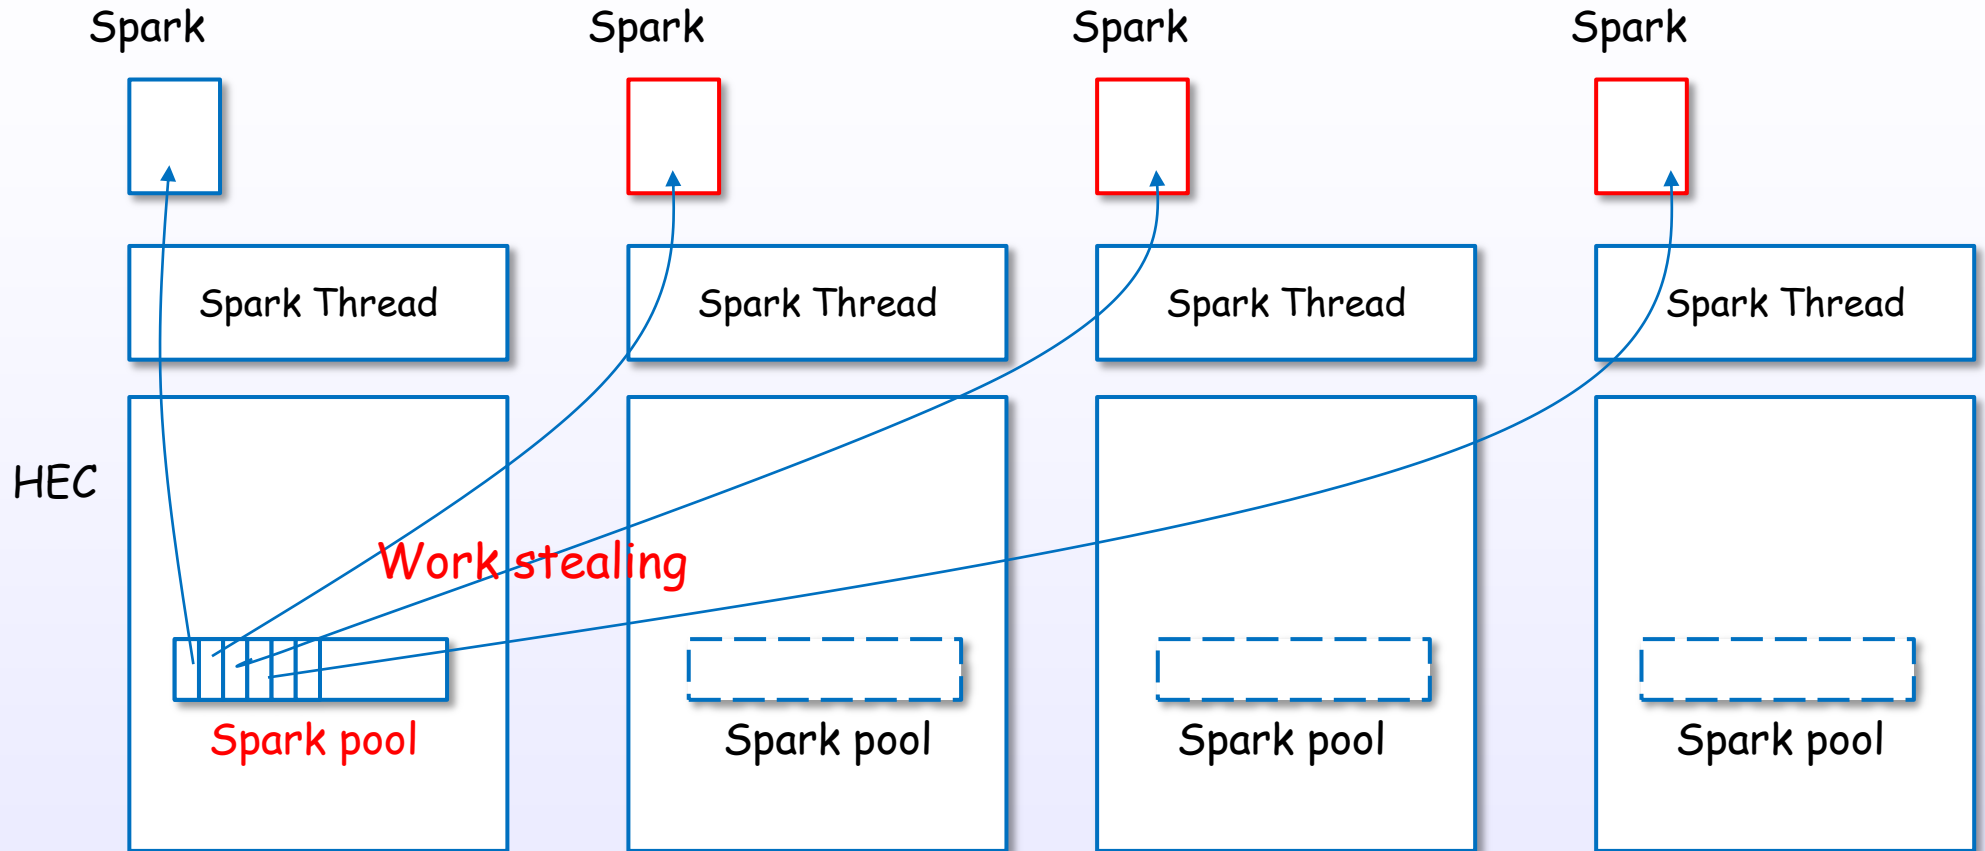
Sparks and Spark pool

logical view




Spark pool and work stealing

physical view



Sparks and closures



MVar


MVar

Haskell Thread #0


Haskell Thread #1



putMVar



takeMVar



MVar

MVar

Haskell Thread



putMVar

BLOCKED
if full



MVar

Haskell Thread



takeMVar


BLOCKED
if empty



MVar

MVar example

Haskell Thread #0




takeMVar

Haskell Thread #1




putMVar




MVar view

User view

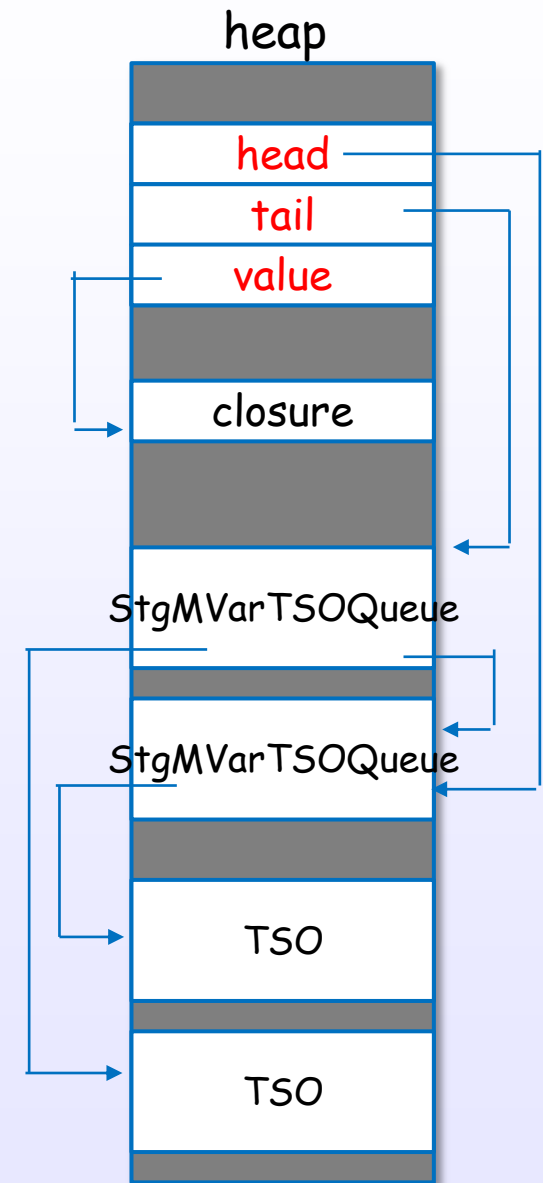
MVar



logical MVar object



physical MVar object



newEmptyMVar

Haskell Threads

```
newEmptyMVar  
newMVar#
```

(1) **call** Runtime primitive

Runtime System


```
stg_newMVarzh  
  ALLOC_PRIM_  
  SET_HDR  
  StgMVar_head  
  StgMVar_tail  
  StgMVar_value
```

(2) **create MVar** object

MVar object

stg_END_TSO_QUEUE_closure

heap



(3) **link** each field

References : [15], [17], [18], [S30], [S12]

takeMVar (empty case)

Haskell Threads

takeMVar
takeMVar#

Runtime System

stg_takeMVarzh

- (1) **create** StgMVarTSOQueue
- (2) **append**
- (3) **StgReturn**

return to scheduler

head

tail

value

append

MVar object

StgMVarTSOQueue

stg_END_TSO_QUEUE

CurrnetTSO


takeMVar (full case)

Haskell Threads


takeMVar
takeMVar#

Runtime System

stg_takeMVarzh
(1) **get** value
(2) **set** empty
(3) **remove** head
(4) **tryWakeupThread**




MVar object



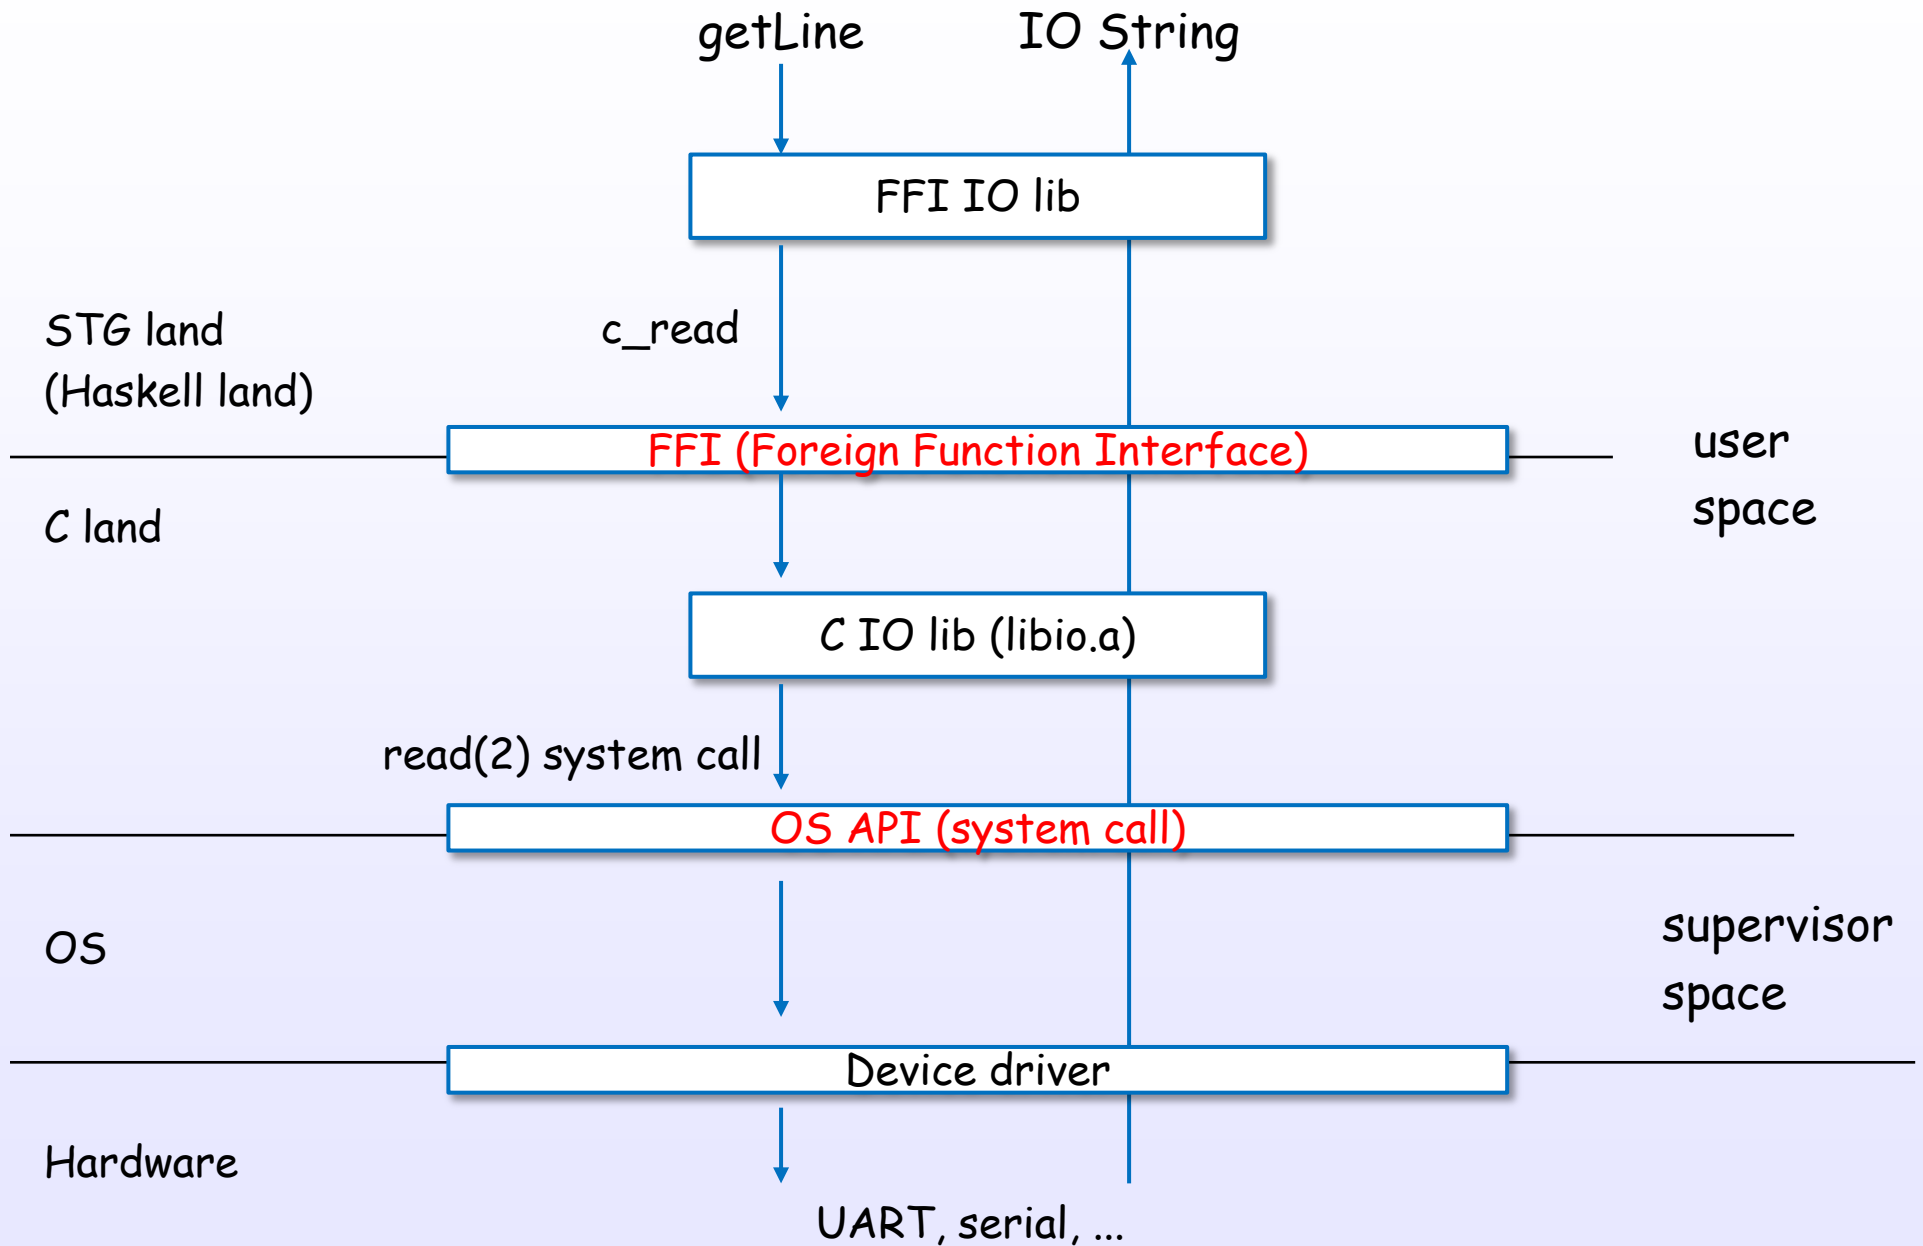
IO and FFI

IO

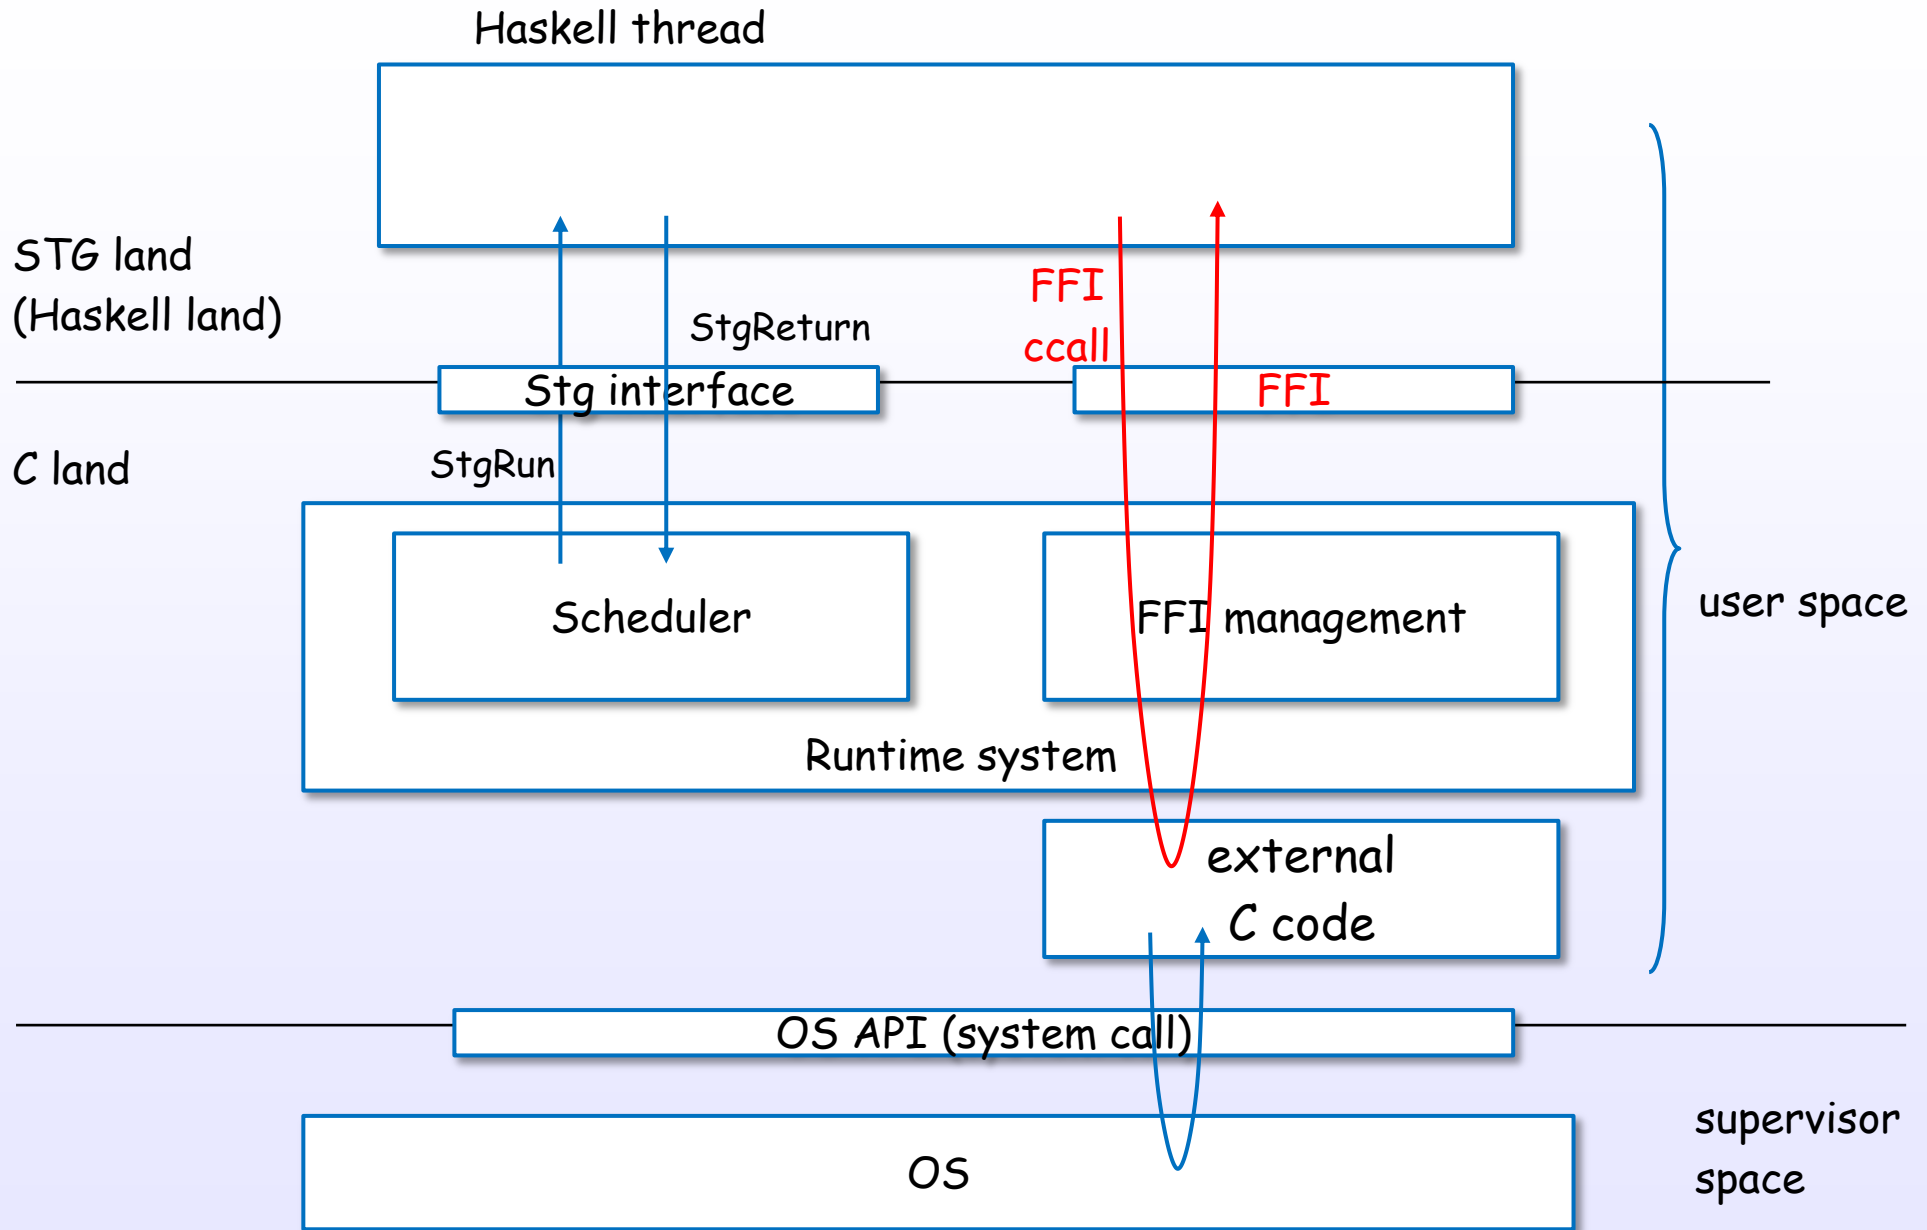
Haskell Thread



IO example: getLine

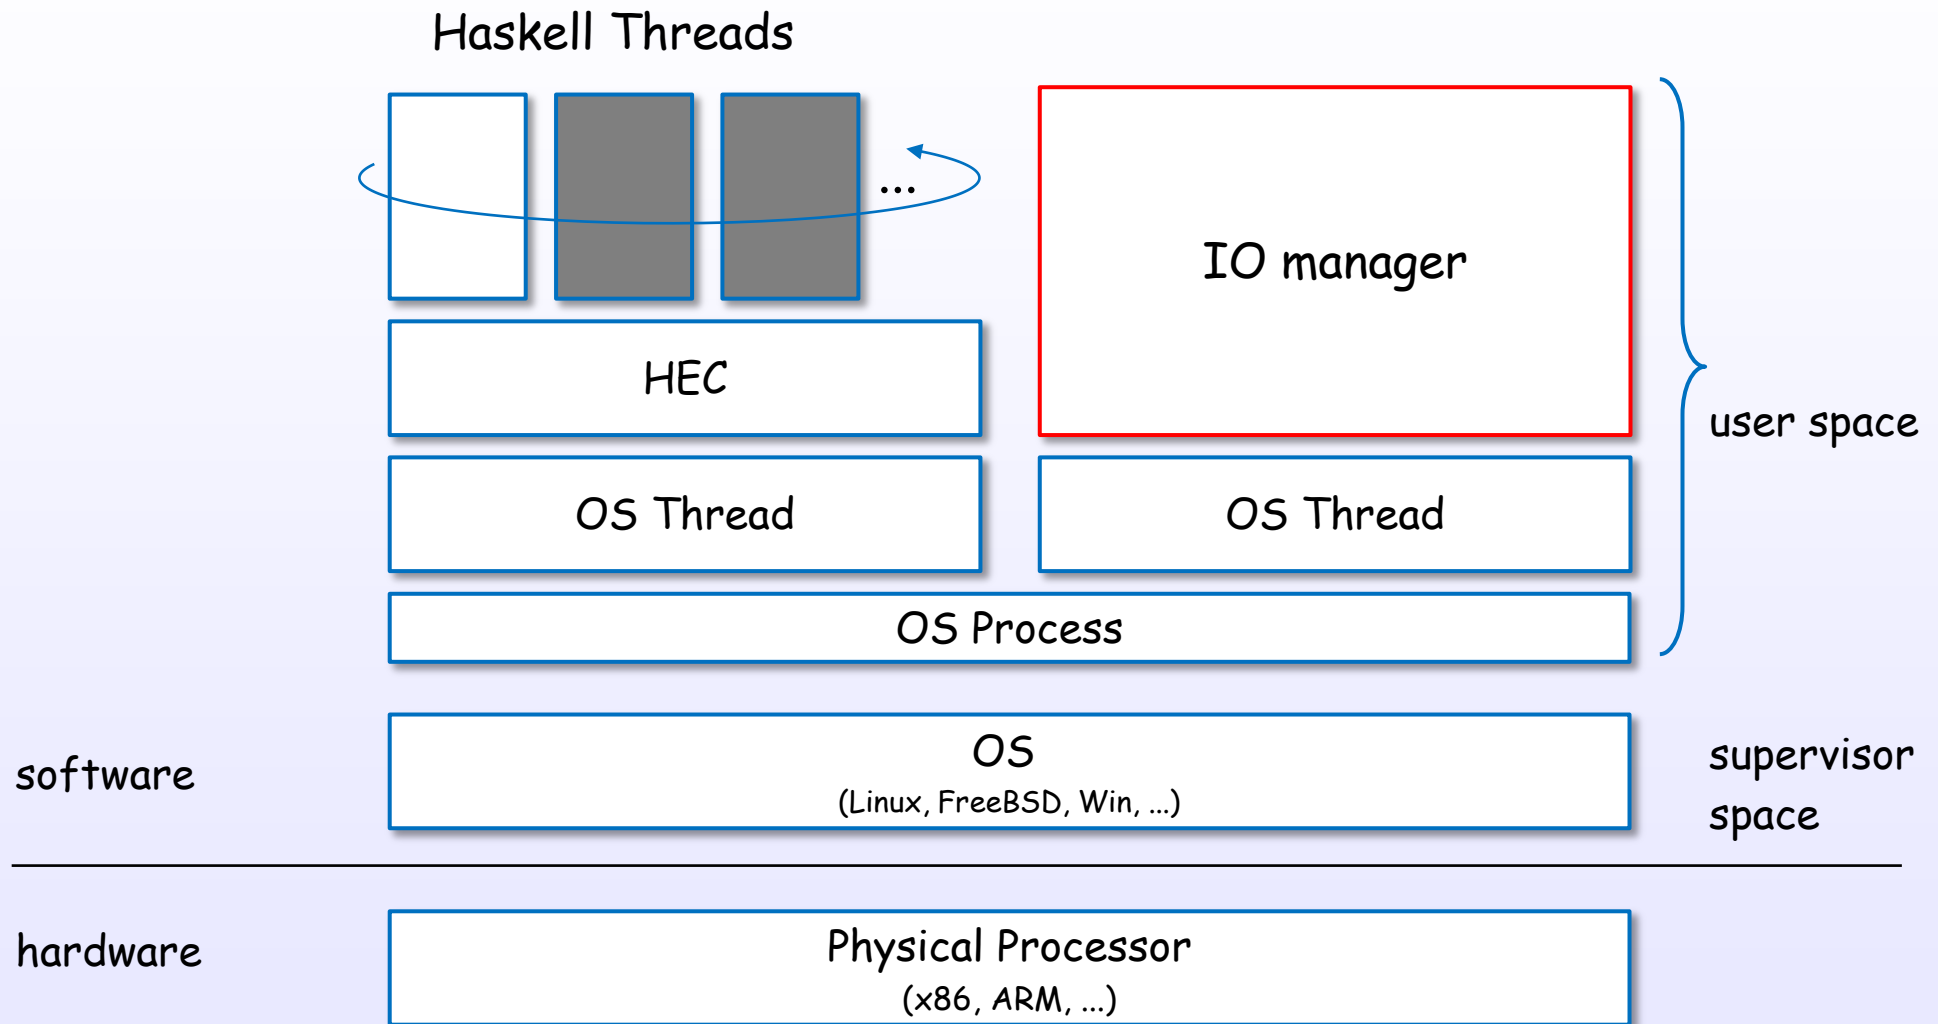


FFI (Foreign Function Interface)



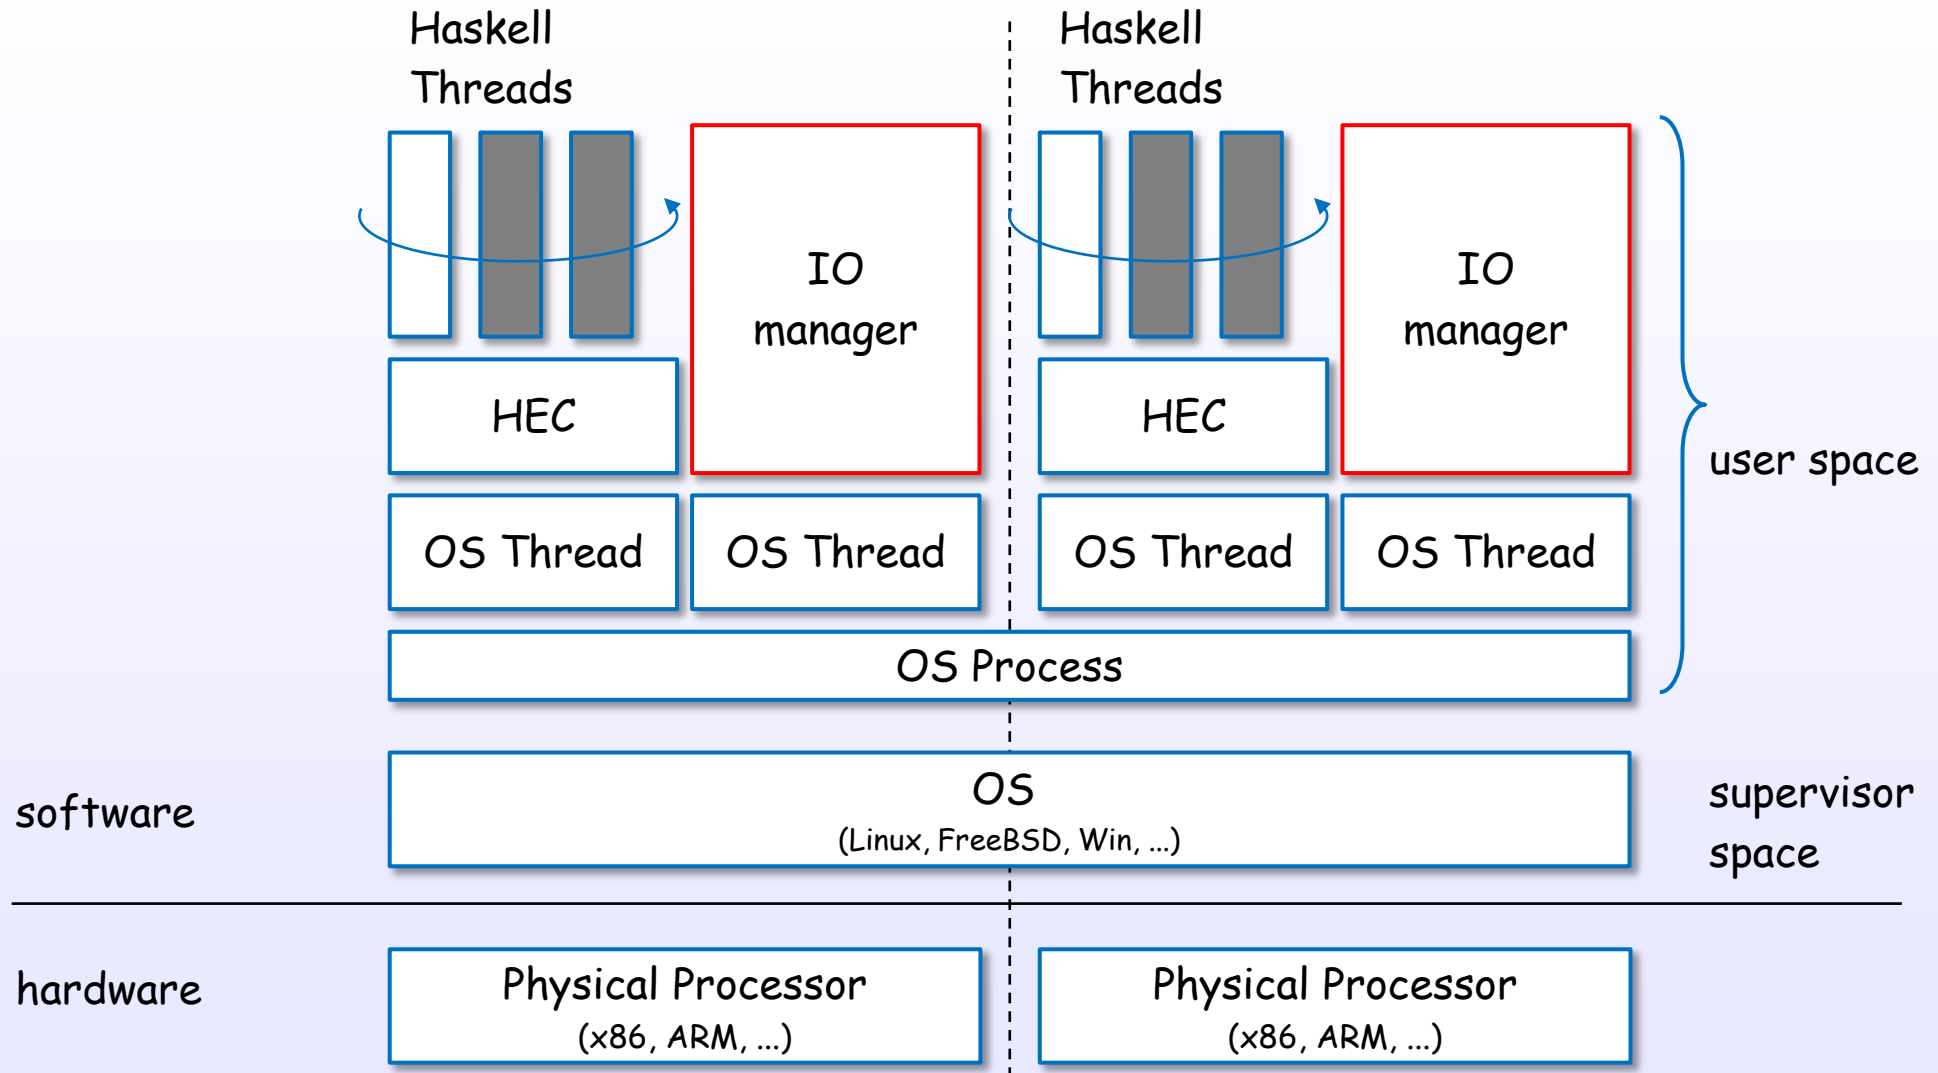
IO manager

IO manager (single core)



*Threaded option case (ghc -threaded)

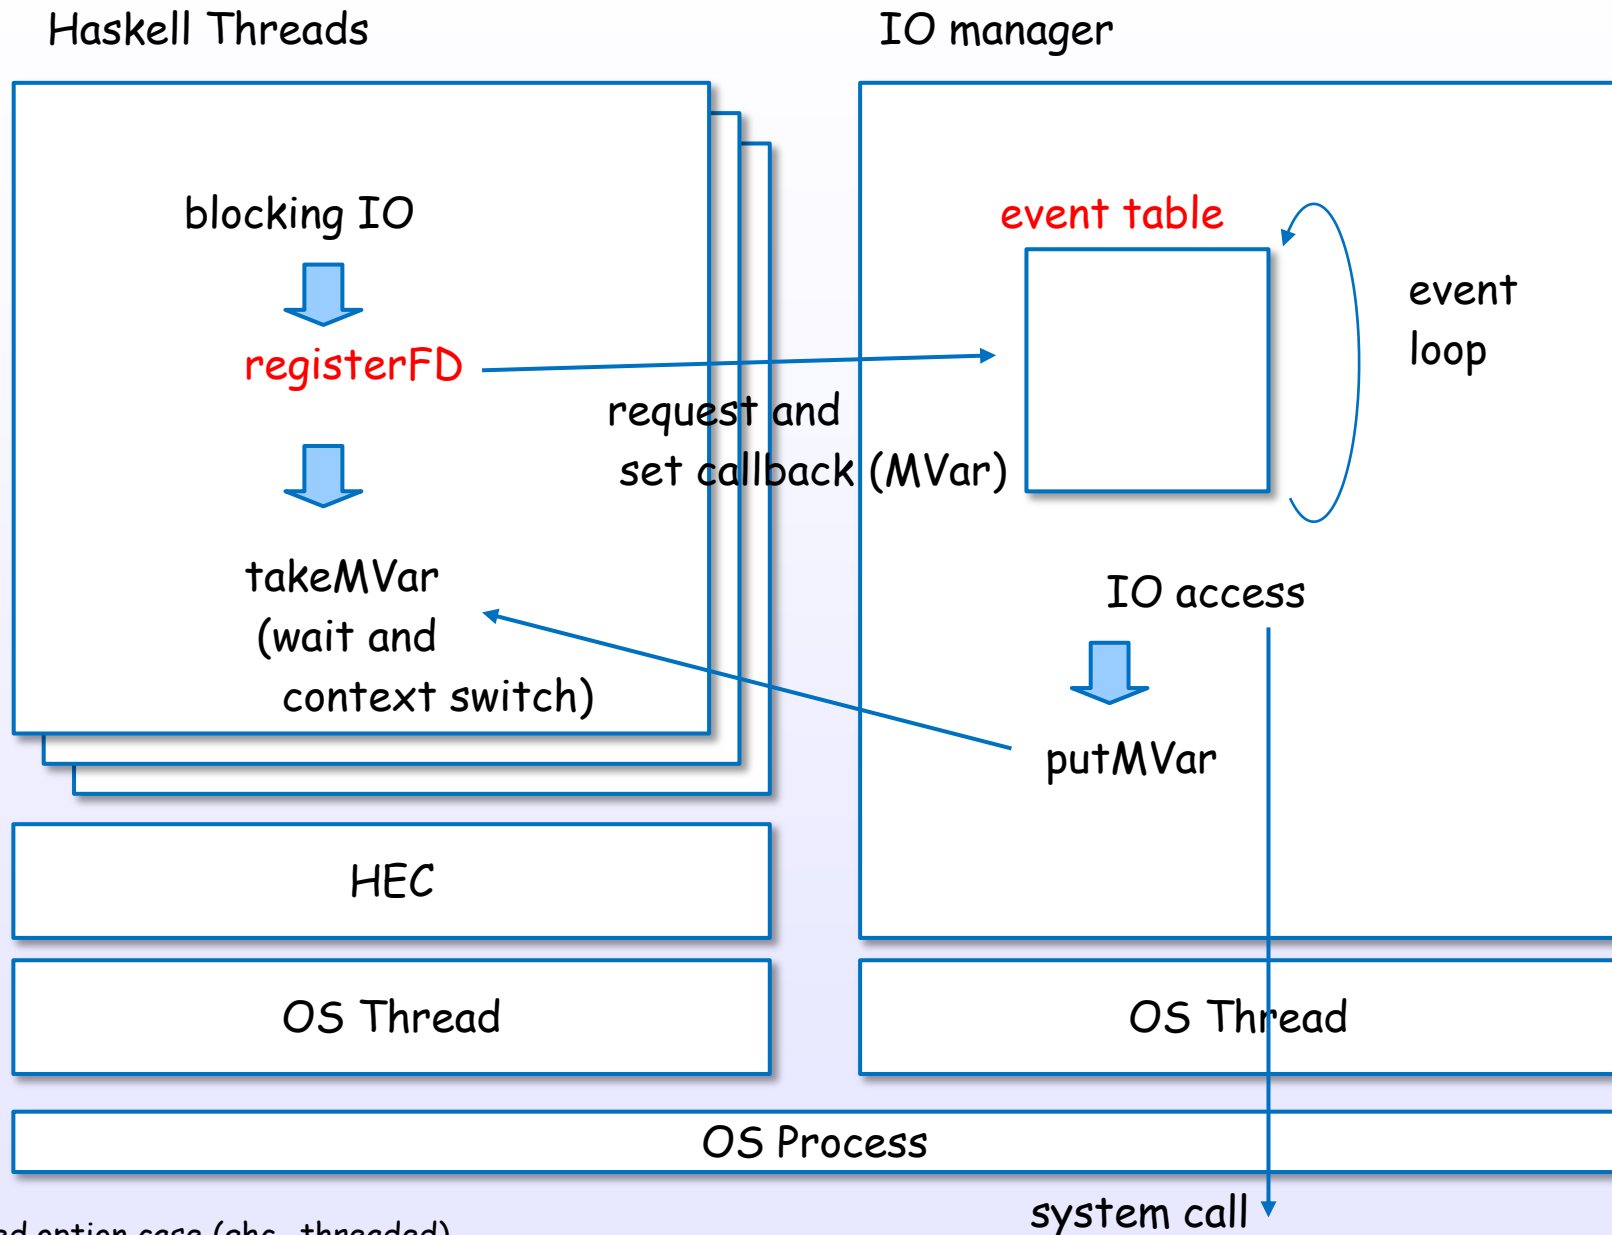
IO manager (multi core)



*Threaded option case (ghc -threaded)

References : [6], [4], [7]

IO manager

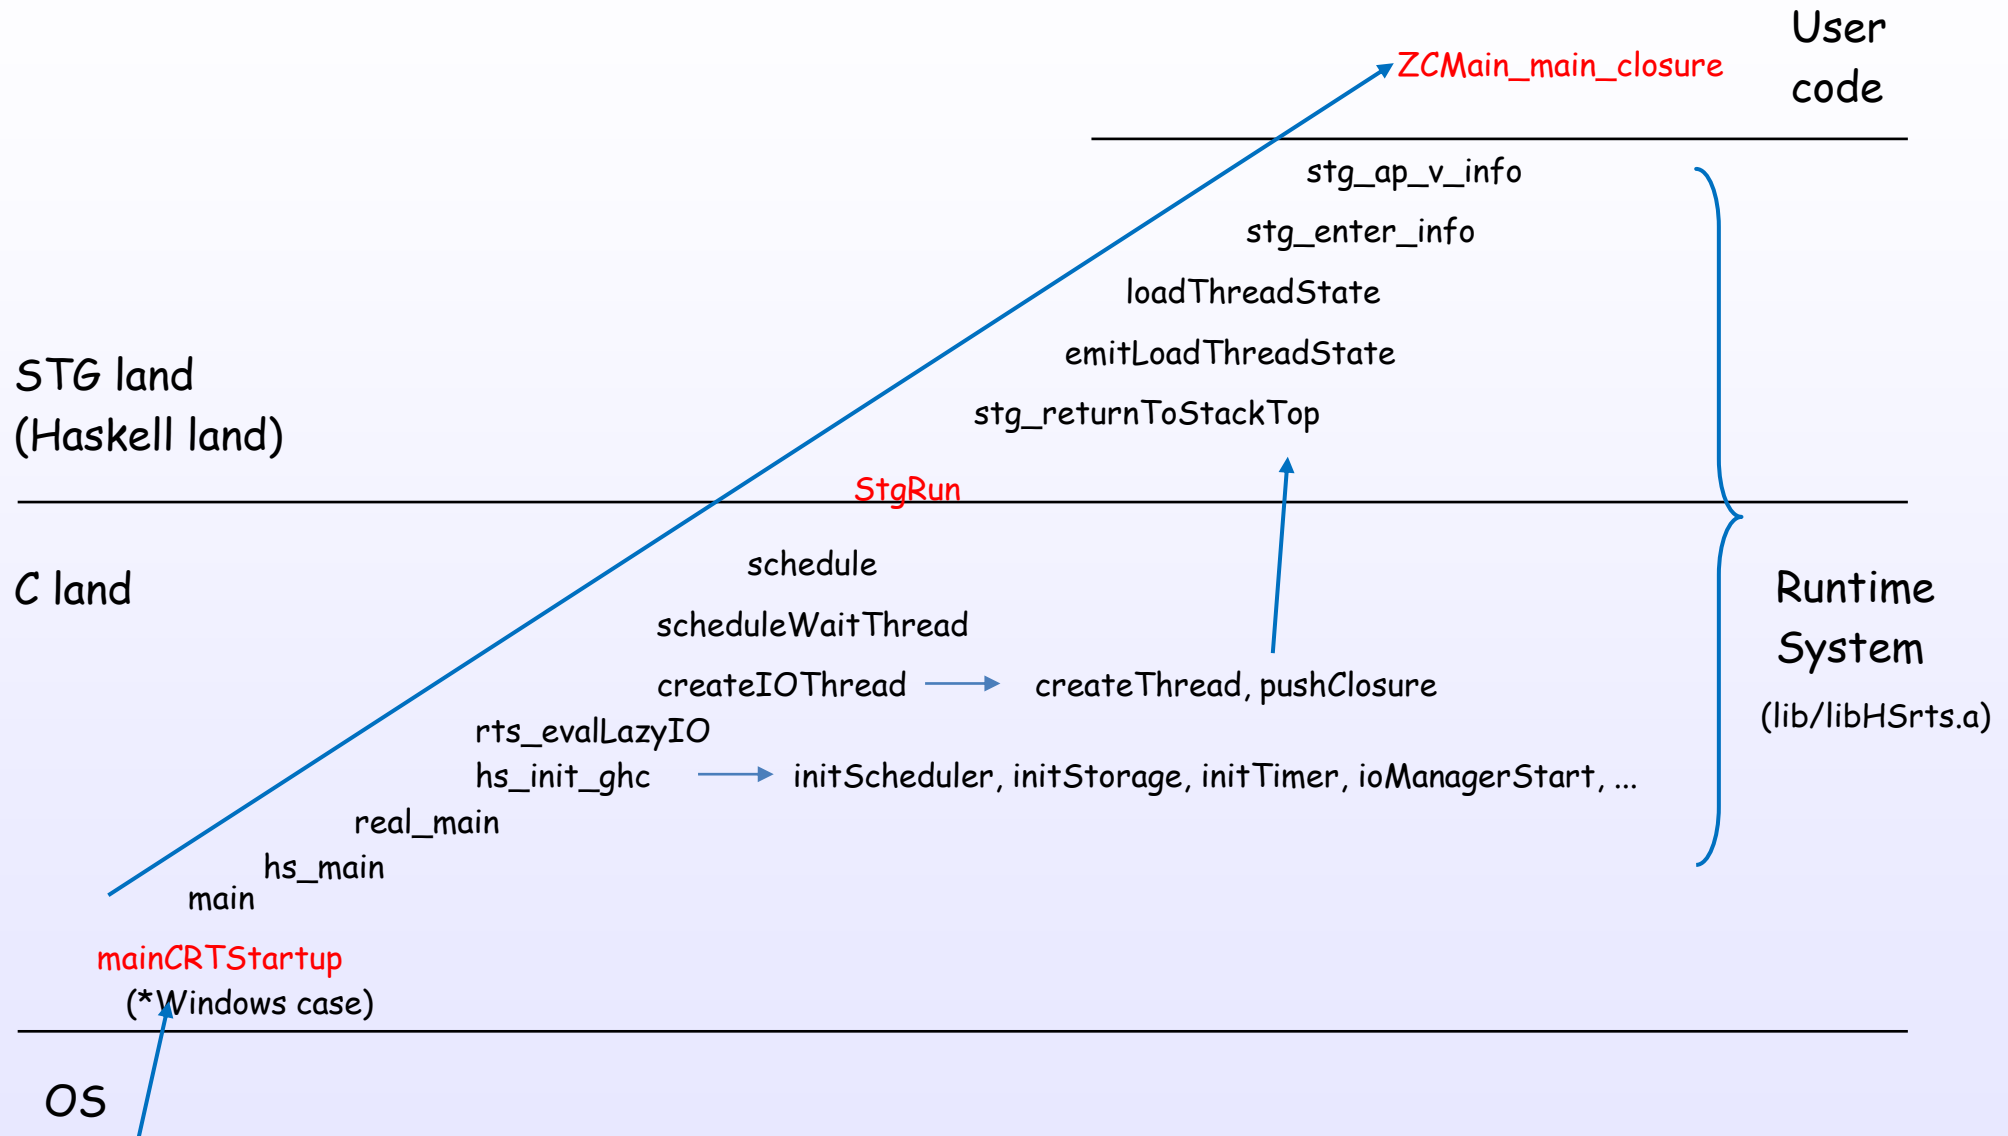


*Threaded option case (ghc -threaded)

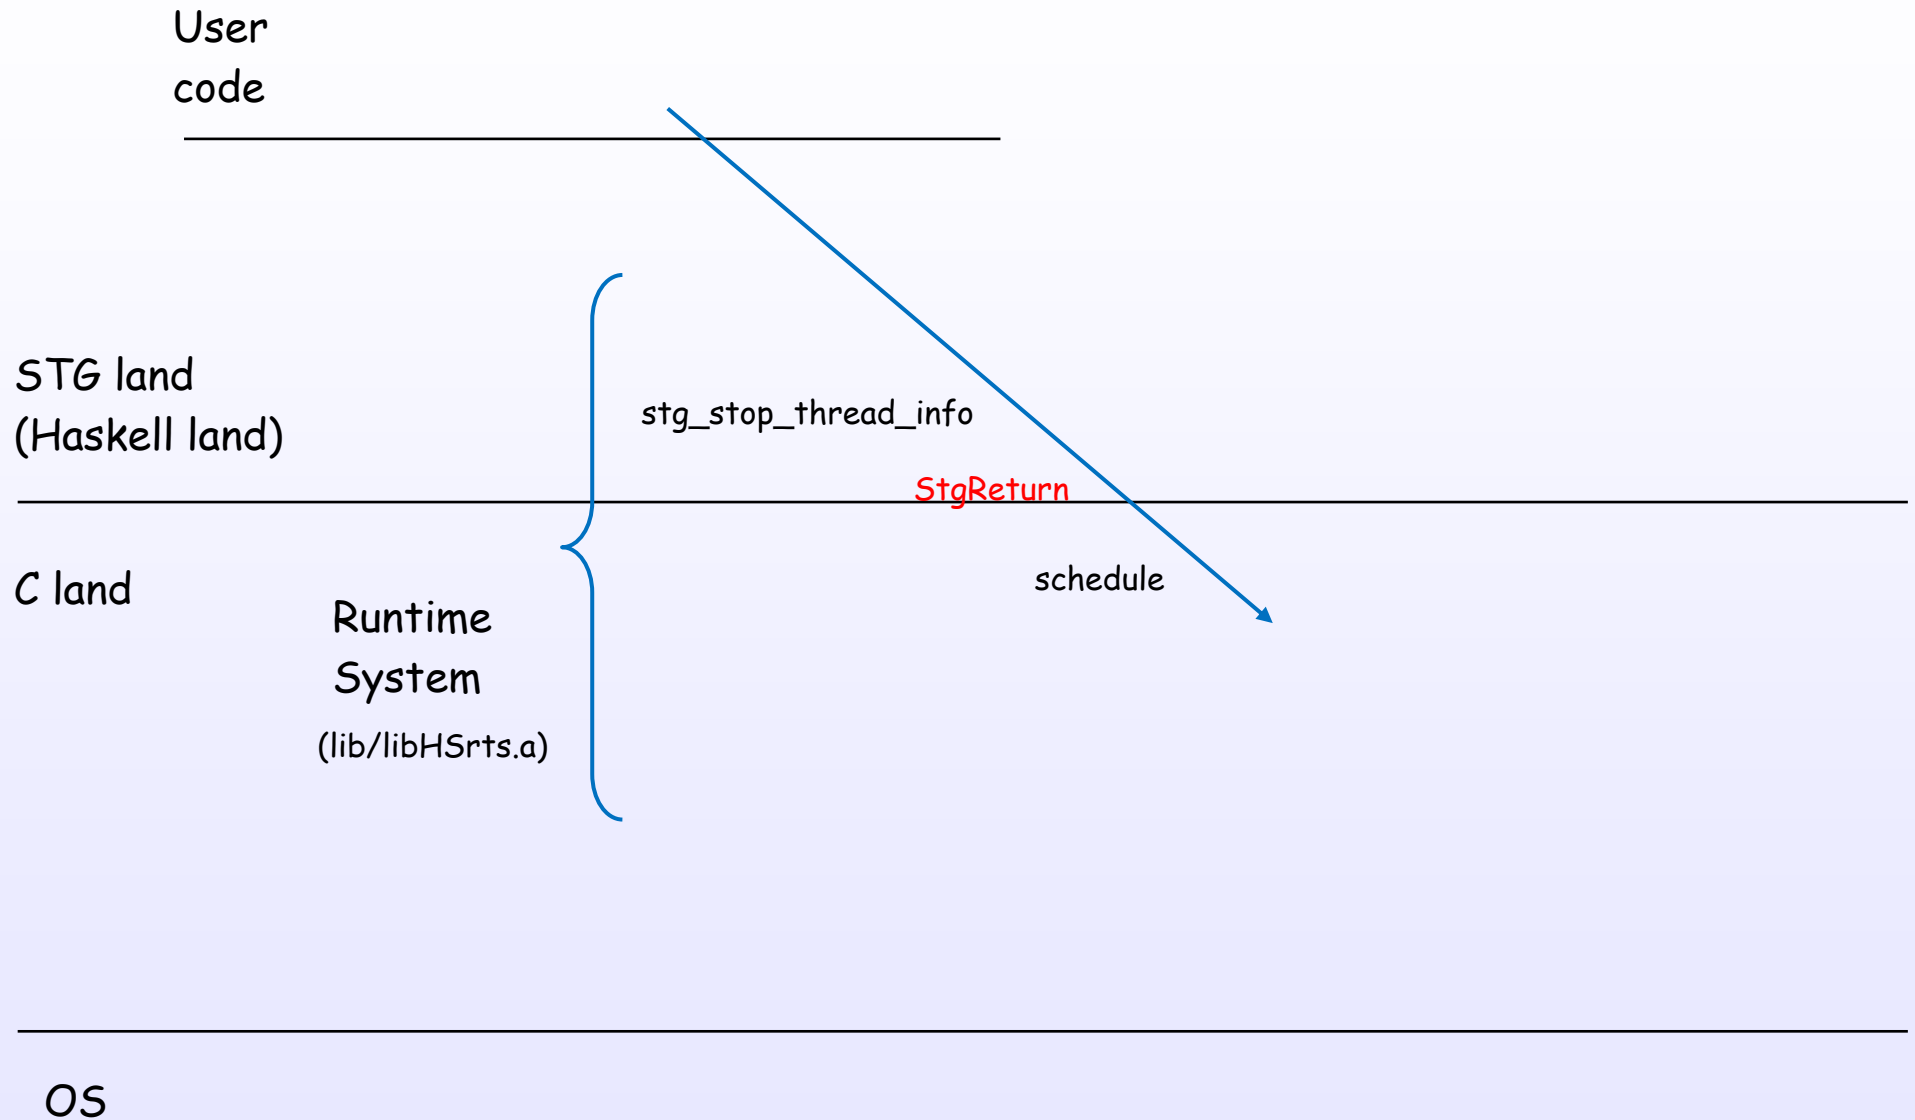
References : [6], [4], [7], [S28], [S29], [S31], [S36], [S34], [S3]

Bootstrap

Bootstrap sequence



Exit sequence



References

References

- [1] Implementing lazy functional languages on stock hardware: the Spineless Tagless G-machine Version 2.5
<http://research.microsoft.com/en-us/um/people/simonpj/Papers/spineless-tagless-gmachine.ps.gz>
- [2] Making a Fast Curry Push/Enter vs Eval/Apply for Higher-order Languages
<http://research.microsoft.com/en-us/um/people/simonpj/papers/eval-apply/>
- [3] Faster Laziness Using Dynamic Pointer Tagging
<http://research.microsoft.com/en-us/um/people/simonpj/papers/ptr-tag/ptr-tagging.pdf>
- [4] Runtime Support for Multicore Haskell
<http://research.microsoft.com/en-us/um/people/simonpj/papers/parallel/multicore-ghc.pdf>
- [5] Extending the Haskell Foreign Function Interface with Concurrency
<http://community.haskell.org/~simonmar/papers/conc-ffi.pdf>
- [6] Mio: A High-Performance Multicore IO Manager for GHC
<http://haskell.cs.yale.edu/wp-content/uploads/2013/08/hask035-voellmy.pdf>
- [7] The GHC Runtime System
web.mit.edu/~ezyang/Public/jfp-ghc-rts.pdf
- [8] The GHC Runtime System
<http://www.scs.stanford.edu/14sp-cs240h/slides/ghc-rts.pdf>
- [9] Evaluation on the Haskell Heap
<http://blog.ezyang.com/2011/04/evaluation-on-the-haskell-heap/>
- [10] IO evaluates the Haskell Heap
<http://blog.ezyang.com/2011/04/io-evaluates-the-haskell-heap/>

References

- [11] Understanding the Stack
<http://www.well-typed.com/blog/94/>
- [12] Understanding the RealWorld
<http://www.well-typed.com/blog/95/>
- [13] The GHC scheduler
<http://blog.ezyang.com/2013/01/the-ghc-scheduler/>
- [14] GHC's Garbage Collector
http://www.mm-net.org.uk/workshop190404/GHC's_Garbage_Collector.ppt
- [15] Concurrent Haskell
<http://www.haskell.org/ghc/docs/papers/concurrent-haskell.ps.gz>
- [16] Beautiful Concurrency
<https://www.fpcomplete.com/school/advanced-haskell/beautiful-concurrency>
- [17] Anatomy of an MVar operation
<http://blog.ezyang.com/2013/05/anatomy-of-an-mvar-operation/>
- [18] Parallel and Concurrent Programming in Haskell
<http://community.haskell.org/~simonmar/pcph/>
- [19] Real World Haskell
<http://book.realworldhaskell.org/>

References

The GHC Commentary

- [C1] <https://ghc.haskell.org/trac/ghc/wiki/Commentary>
- [C2] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/SourceTree>
- [C3] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler>
- [C4] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/HscMain>
- [C5] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/CoreSynType>
- [C6] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/StgSynType>
- [C7] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/CmmType>
- [C8] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/GeneratedCode>
- [C9] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/SymbolNames>
- [C10] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts>
- [C11] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/HeapObjects>
- [C12] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/Stack>
- [C13] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/GC>
- [C14] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution>
- [C15] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution/Registers>
- [C16] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution/PointerTagging>
- [C17] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Scheduler>
- [C18] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Libraries>

References

Source code

- [S1] `includes/stg/Regs.h`
- [S2] `includes/stg/MachRegs.h`
- [S3] `includes/rts/storage/ClosureTypes.h`
- [S4] `includes/rts/storage/Closures.h`
- [S5] `includes/rts/storage/TSO.h`
- [S6] `includes/rts/storage/InfoTables.h`
- [S7] `compiler/main/DriverPipeline.hs`
- [S8] `compiler/main/HscMain.hs`
- [S9] `compiler/cmm/CmmParse.y.source`
- [S10] `compiler/codeGen/StgCmmForeign.hs`
- [S11] `compiler/codeGen/Stg*.hs`
- [S12] `rts/PrimOps.cmm`
- [S13] `rts/RtsMain.c`
- [S14] `rts/RtsAPI.c`
- [S15] `rts/Capability.h`
- [S16] `rts/Capability.c`
- [S17] `rts/Schedule.c`
- [S18] `rts/StgCRun.c`
- [S19] `rts/StgStartup.cmm`
- [S20] `rts/StgMiscClosures.cmm`
- [S21] `rts/HeapStackCheck.cmm`
- [S22] `rts/Threads.c`
- [S23] `rts/Task.c`
- [S24] `rts/Timer.c`
- [S25] `rts/sm/GC.c`
- [S26] `rts/Sparks.c`
- [S27] `rts/WSDeque.c`
- [S28] `rts/posix/Signals.c`
- [S29] `rts/win32/ThrIOManager.c`
- [S30] `libraries/base/GHC/MVar.hs`
- [S31] `libraries/base/GHC/Conc/IO.hs`
- [S32] `libraries/base/GHC/Conc/Sync.lhs`
- [S33] `libraries/base/GHC/Event/Manager.hs`
- [S34] `libraries/base/GHC/Event/Thread.hs`
- [S35] `libraries/base/GHC/IO/BufferedIO.hs`
- [S36] `libraries/base/GHC/IO/FD.hs`
- [S37] `libraries/base/GHC/IO/Handle/Text.hs`
- [S38] `libraries/base/System/IO.hs`
- [S39] `libraries/base/System/Posix/Internals.hs`
- [S40] `AutoApply.o (utils/genapply/GenApply.hs)`

Connect the algorithm and transistor