

$GHC_{(STG, Cmm, asm)}$ illustrated

for hardware person

exploring some mental models and implementations

Takenobu T.

"Any sufficiently advanced technology is
indistinguishable from **magic**."

Arthur C. Clarke

NOTE

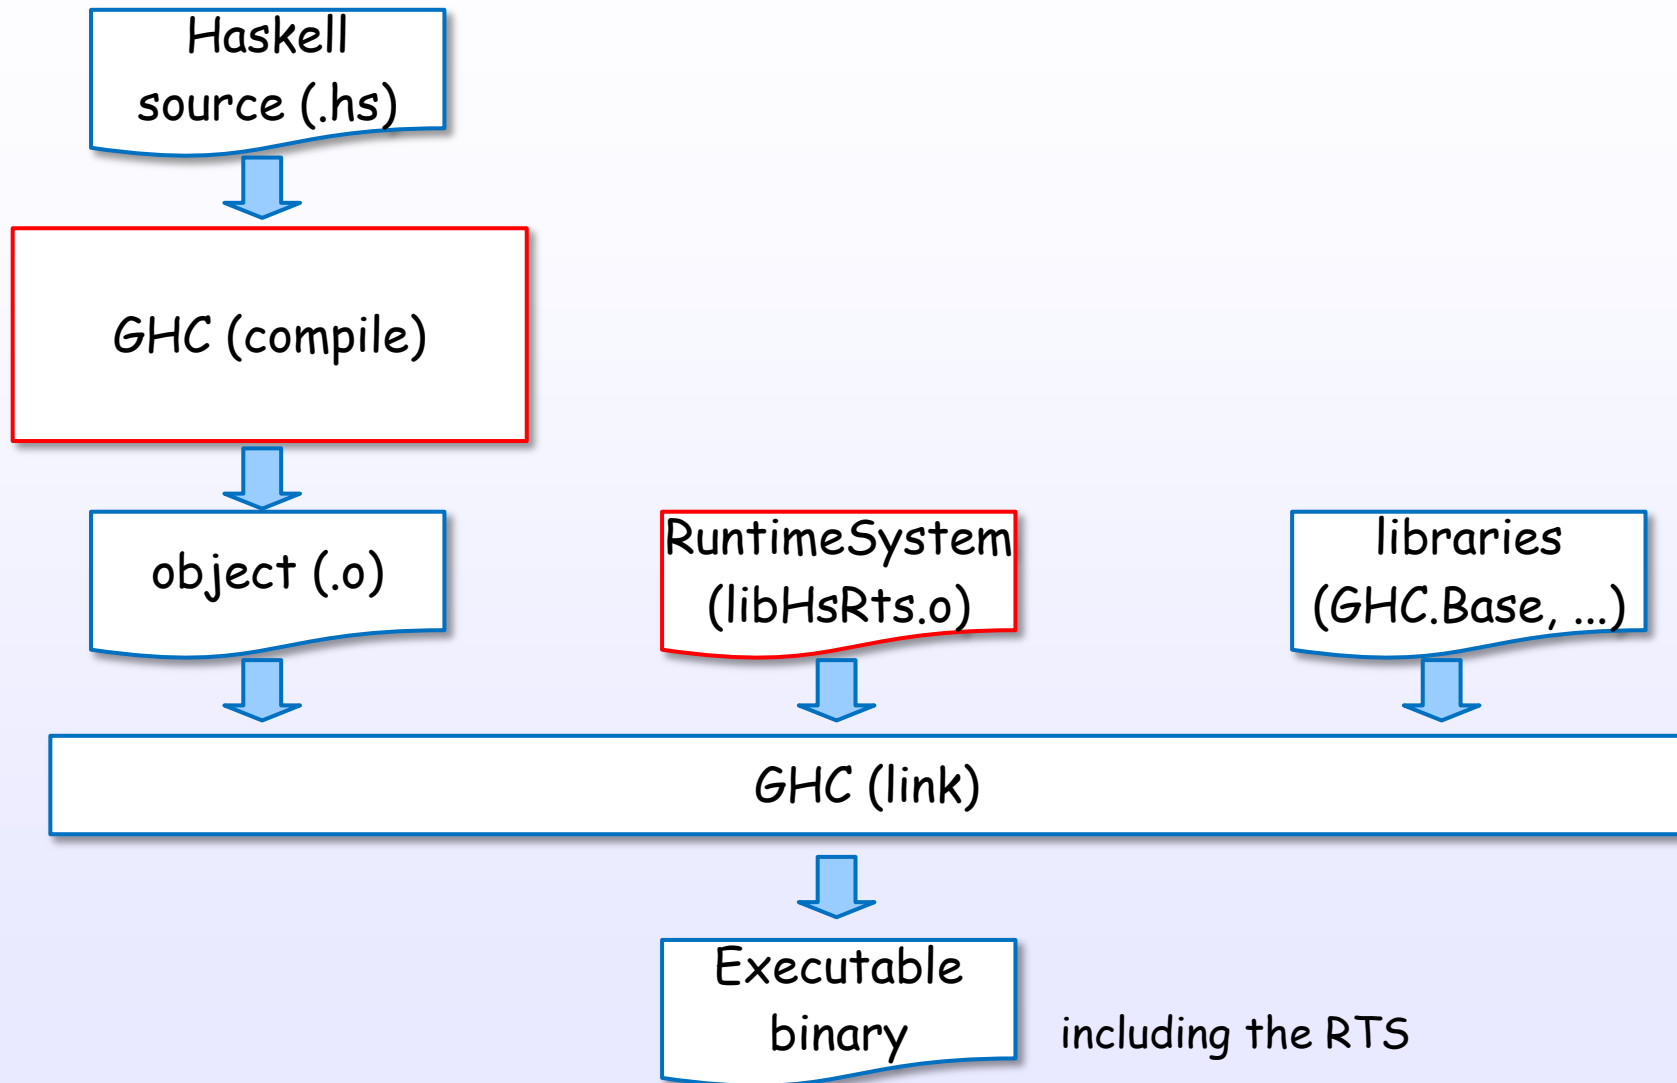
- This is not an official document by the ghc development team.
- Please don't forget "semantics". It's very important.
- This is written for ghc 7.8 (and ghc 7.10).

Contents

- Executable binary
- Compile steps
- Runtime System
- Development languages
- Machine layer/models
- STG-machine
- Heap objects in STG-machine
- STG-machine evaluation
- Pointer tagging
- Thunk and update
- Allocate and free heap objects
- STG - C land interface
- Thread
- Thread context switch
- Creating main and sub threads
- Thread migration
- Heap and Threads
- Threads and GC
- Bound thread
- Spark
- Mvar
- Software transactional memory
- FFI
- IO and FFI
- IO manager
- Bootstrap
- References

Executable binary

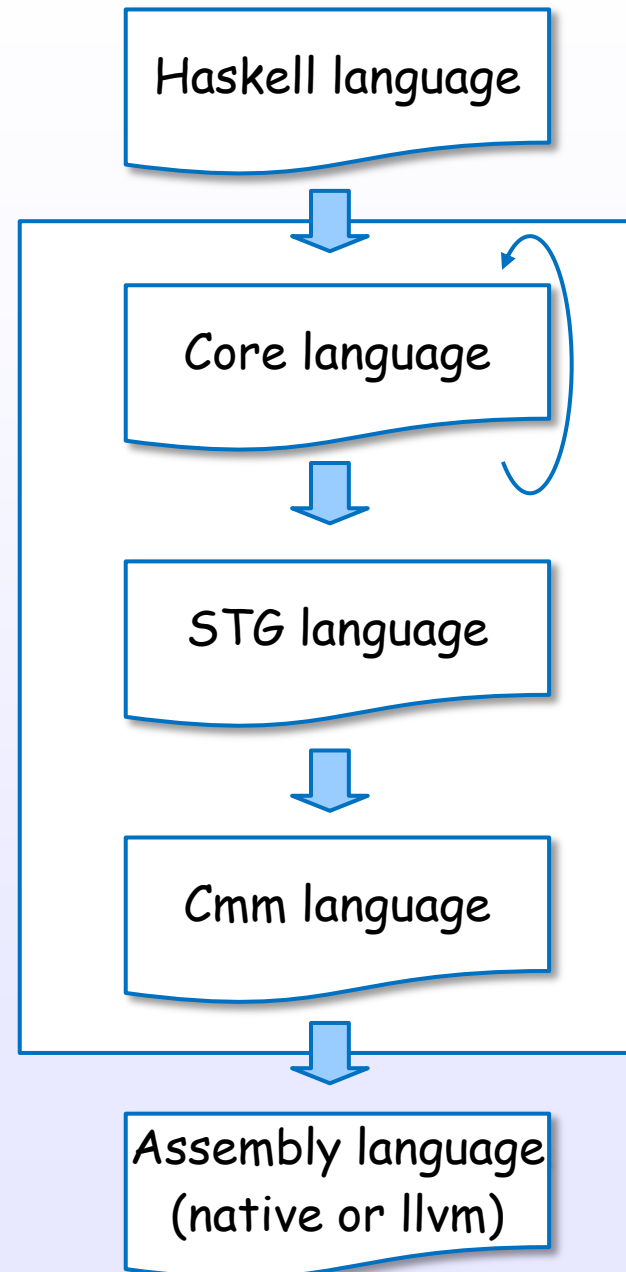
The GHC = Compiler + Runtime System (RTS)



Compile steps

GHC transitions between five representations

GHC
compile
steps



*each intermediate code can
be dumped by :*

`% ghc -ddump-parsed`
`% ghc -ddump-rn`

`% ghc -ddump-ds`
`% ghc -ddump-simpl`
`% ghc -ddump-prep`

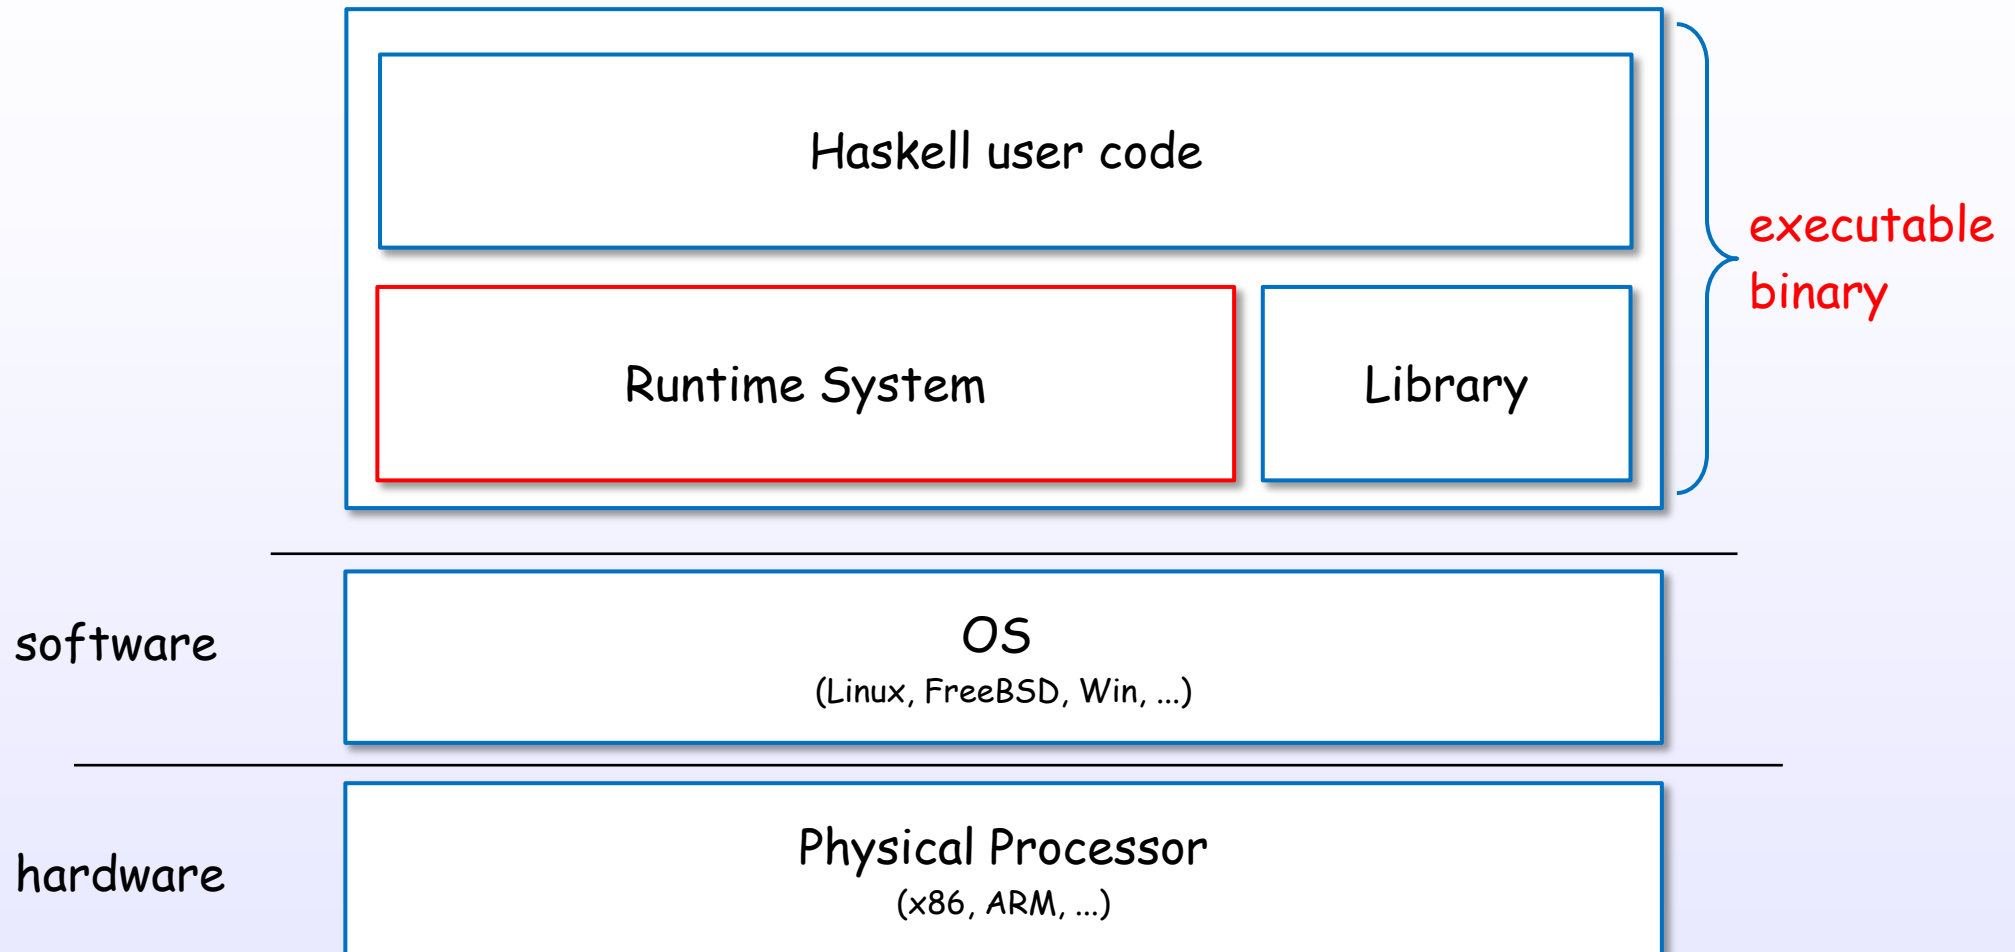
`% ghc -ddump-stg`

`% ghc -ddump-cmm`
`% ghc -ddump-opt-cmm`

`% ghc -ddump-llvm`
`% ghc -ddump-asm`

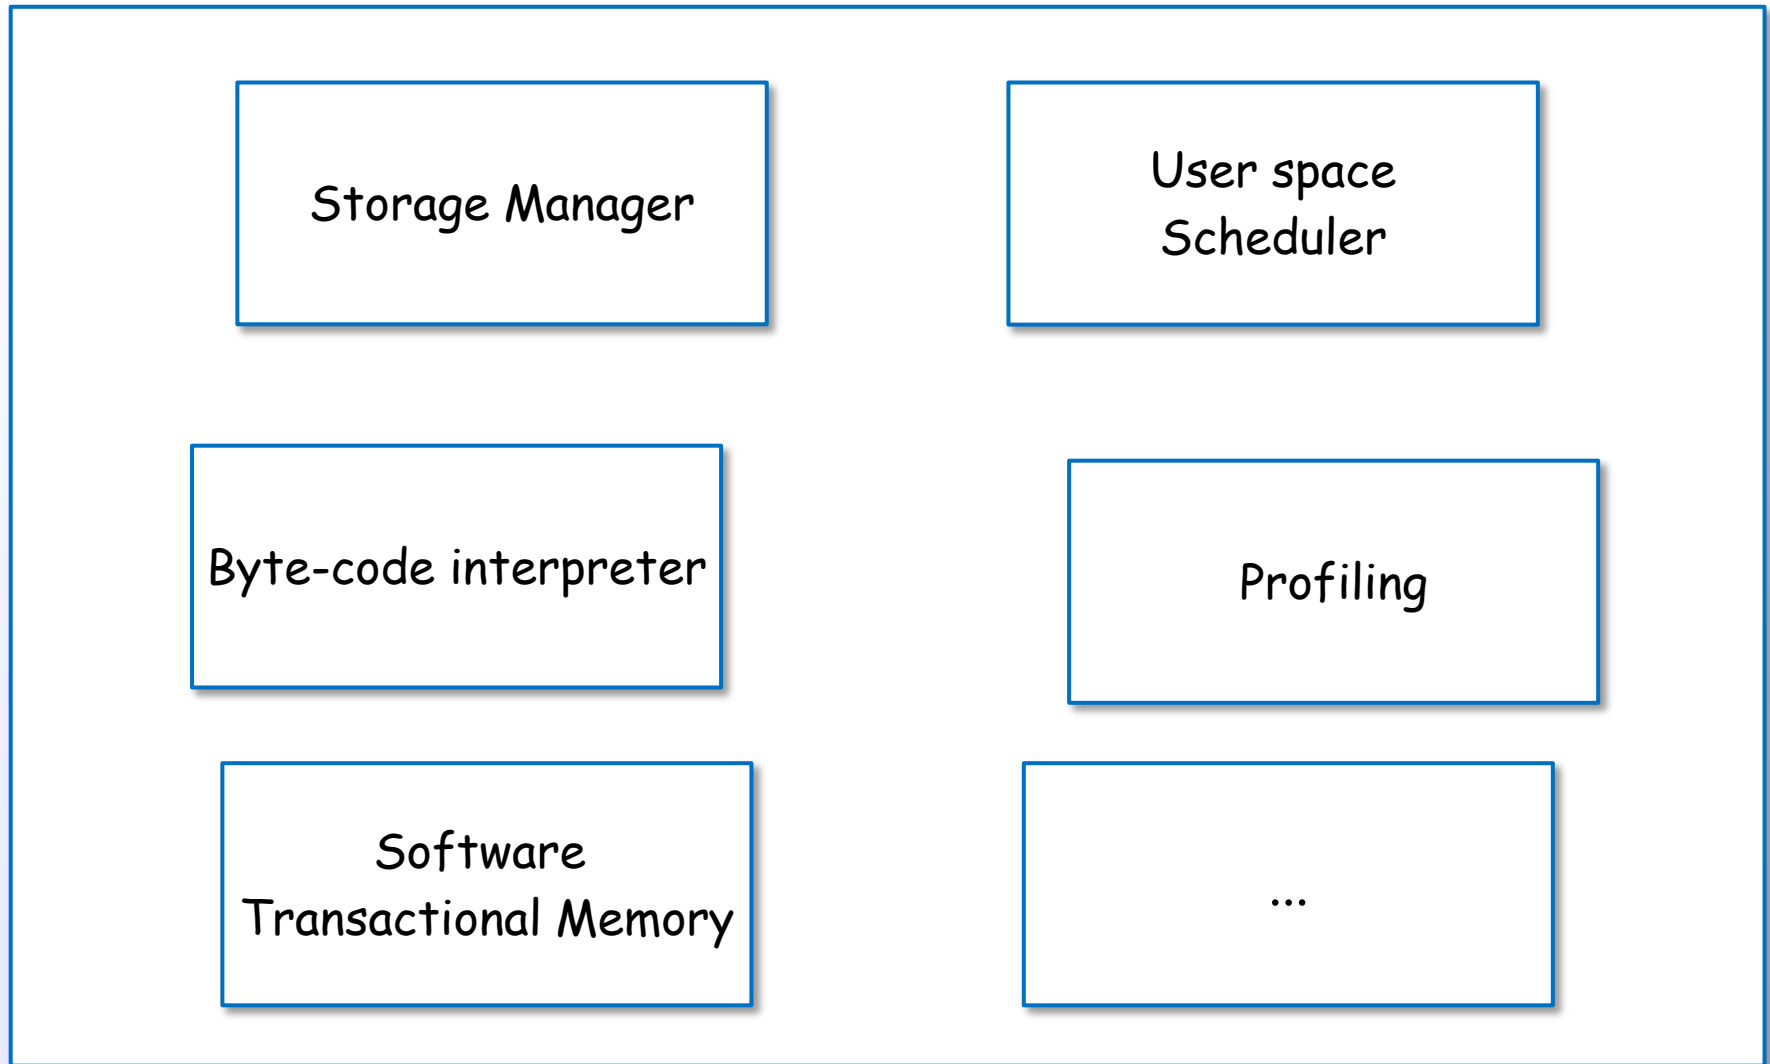
Runtime System

Generated binary includes the RTS



Runtime System includes ...

Runtime System



Development languages

The GHC is developed by some languages

compiler

(\$(TOP)/**compiler**/*)

Haskell

+

Alex (lex)

Happy (yacc)

Cmm (C--)

Assembly

runtime system

(\$(TOP)/**rts**/*)

C

+

Cmm

Assembly

library

(\$(TOP)/**libraries**/*)

Haskell

+

C

Machine layer/models

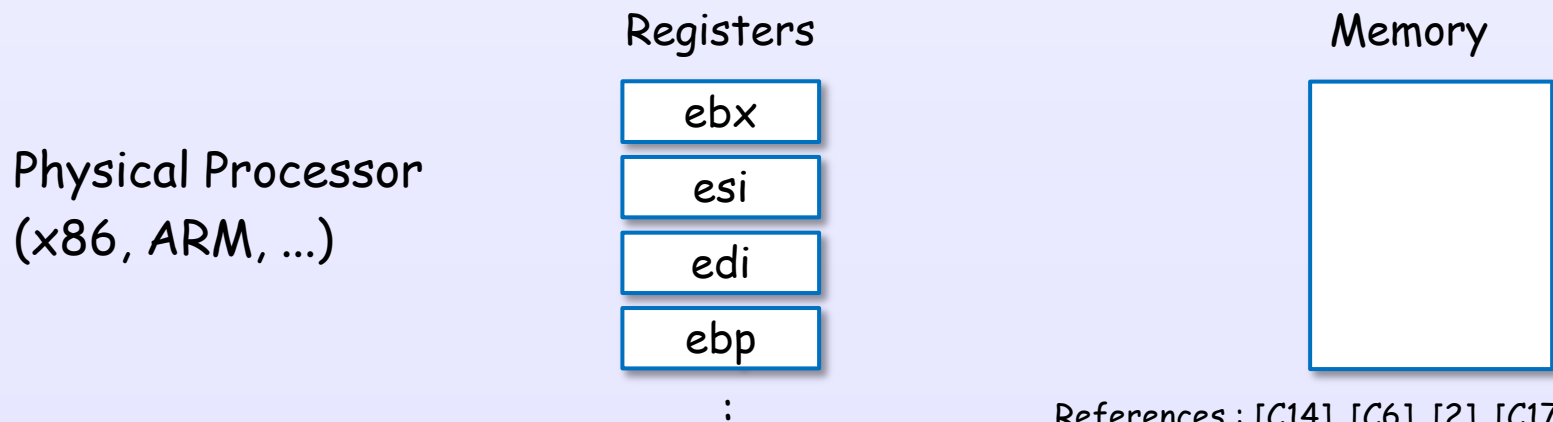
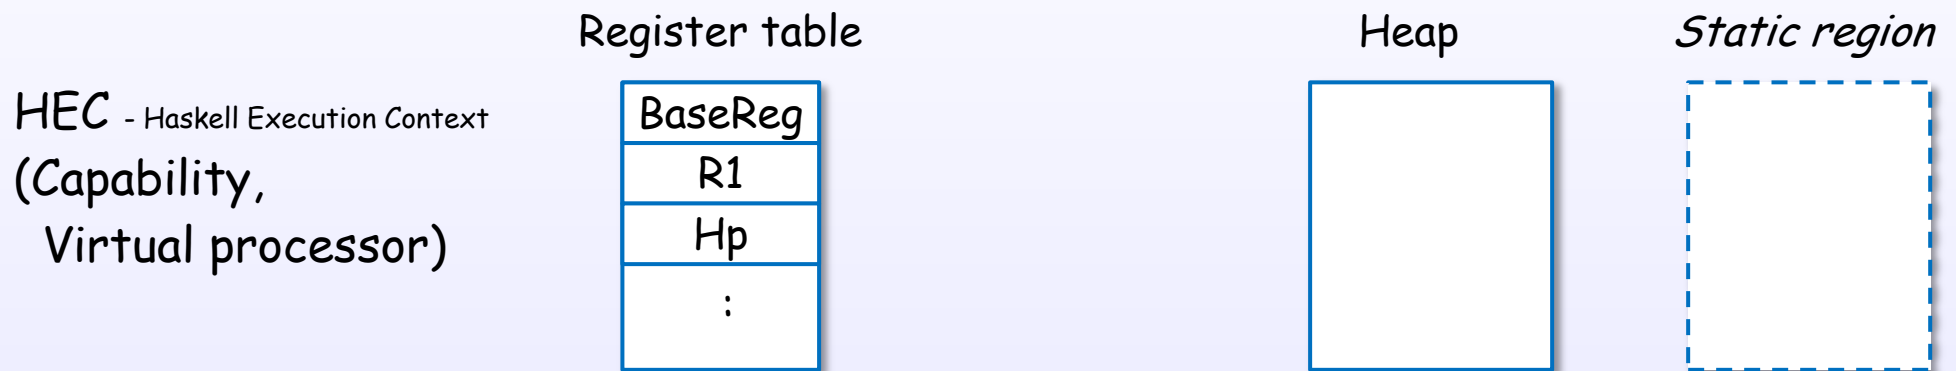
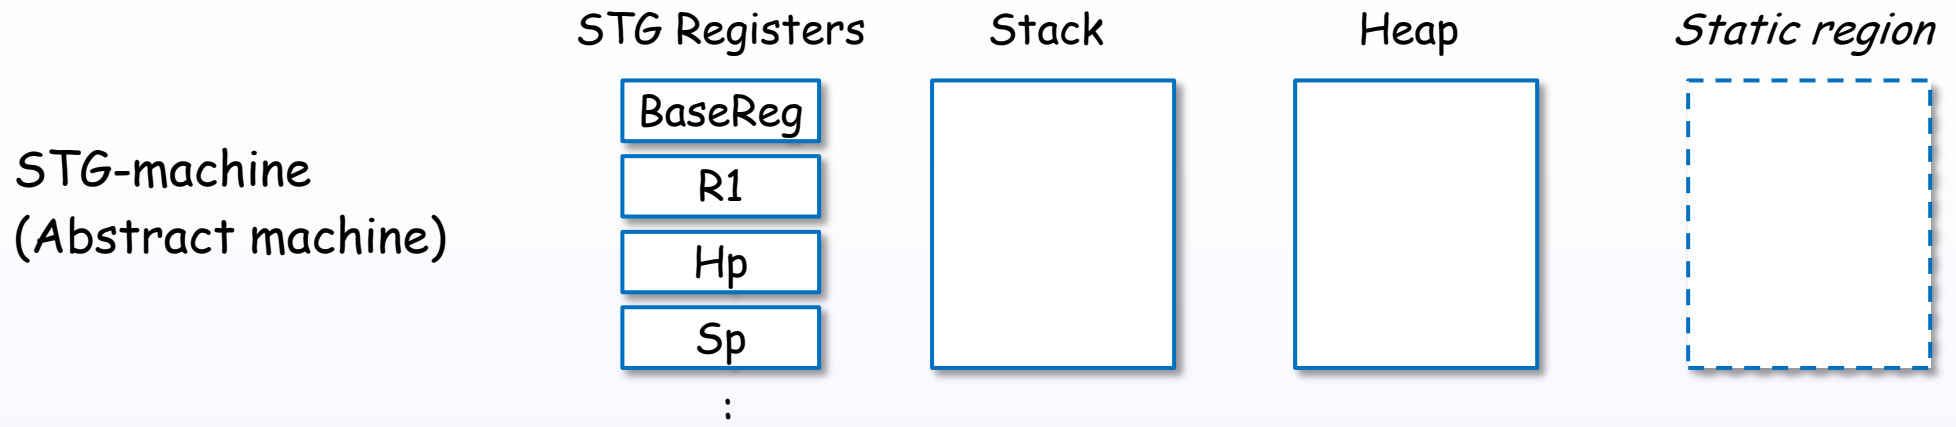
Machine layer

STG-machine
(Abstract machine)

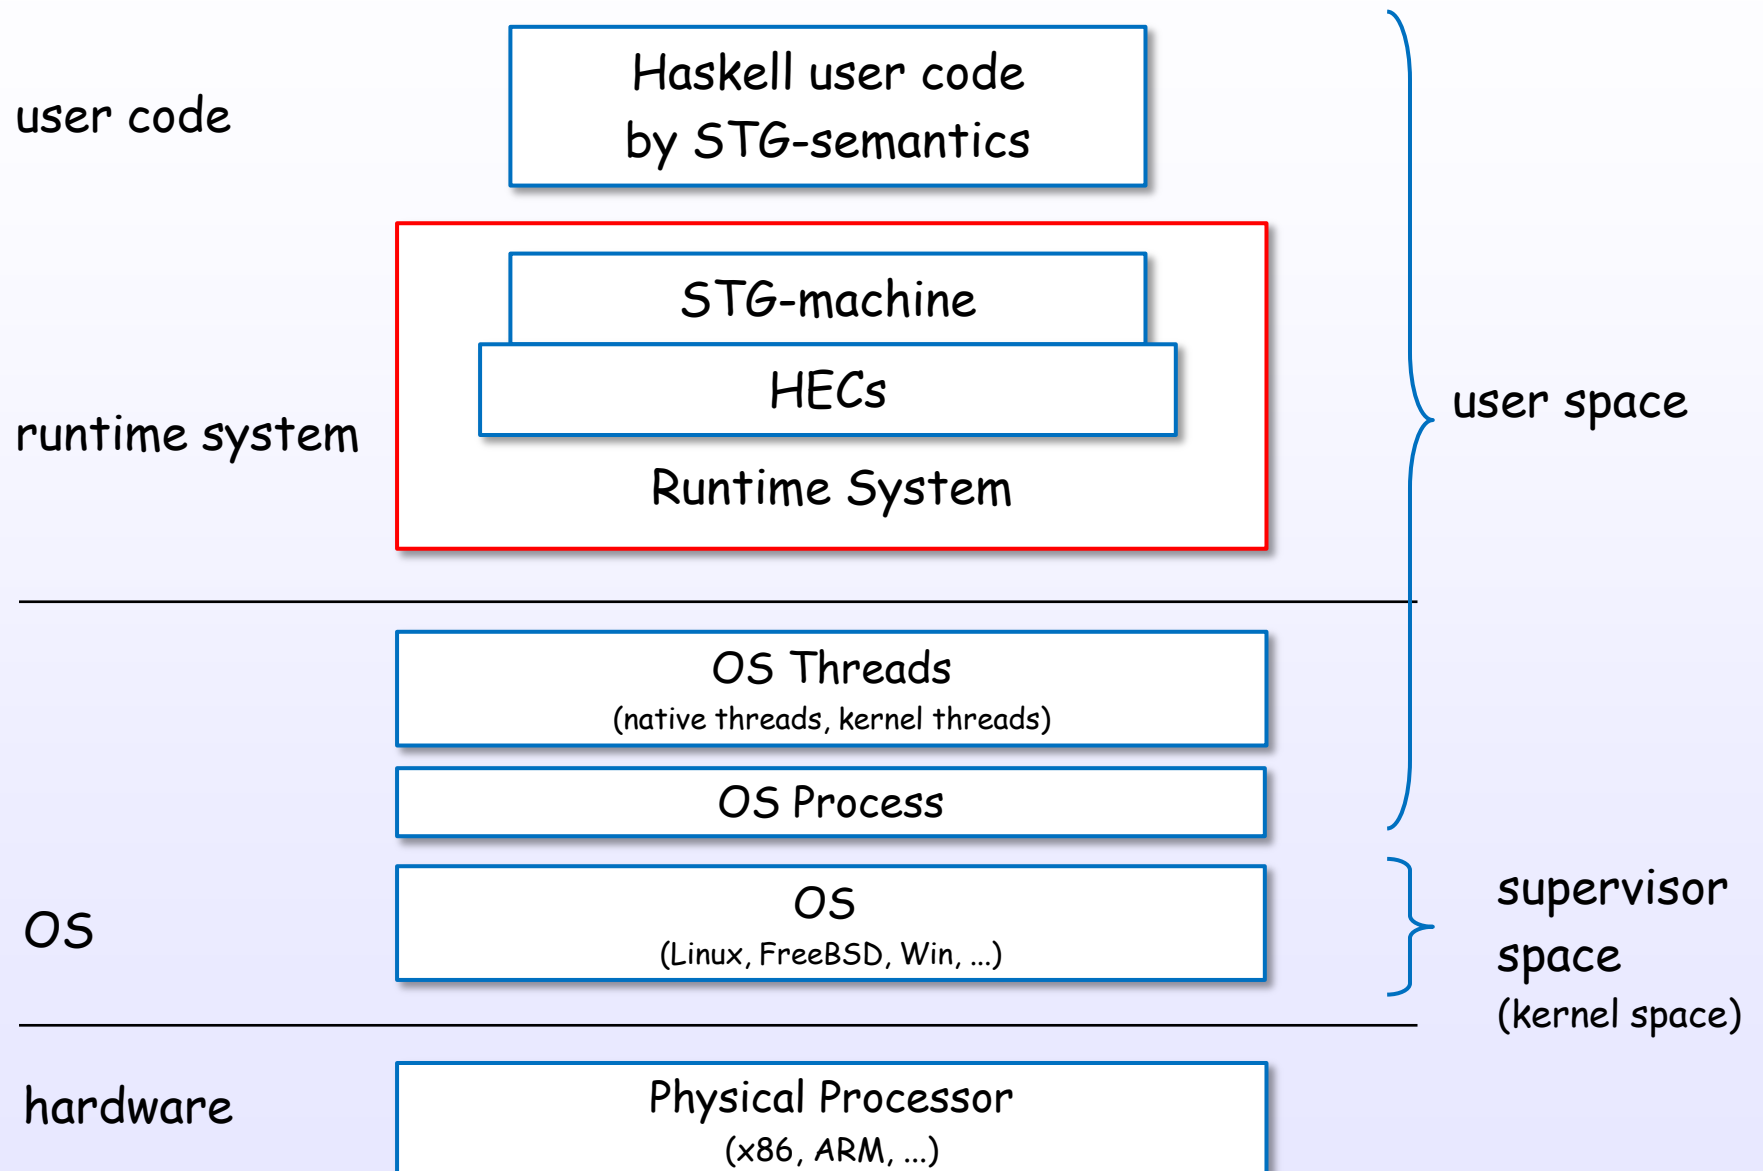
HEC - Haskell Execution Context
(Capability, Virtual processor)

Physical Processor
(x86, ARM, ...)

Machine layer



Runtime system and HEC

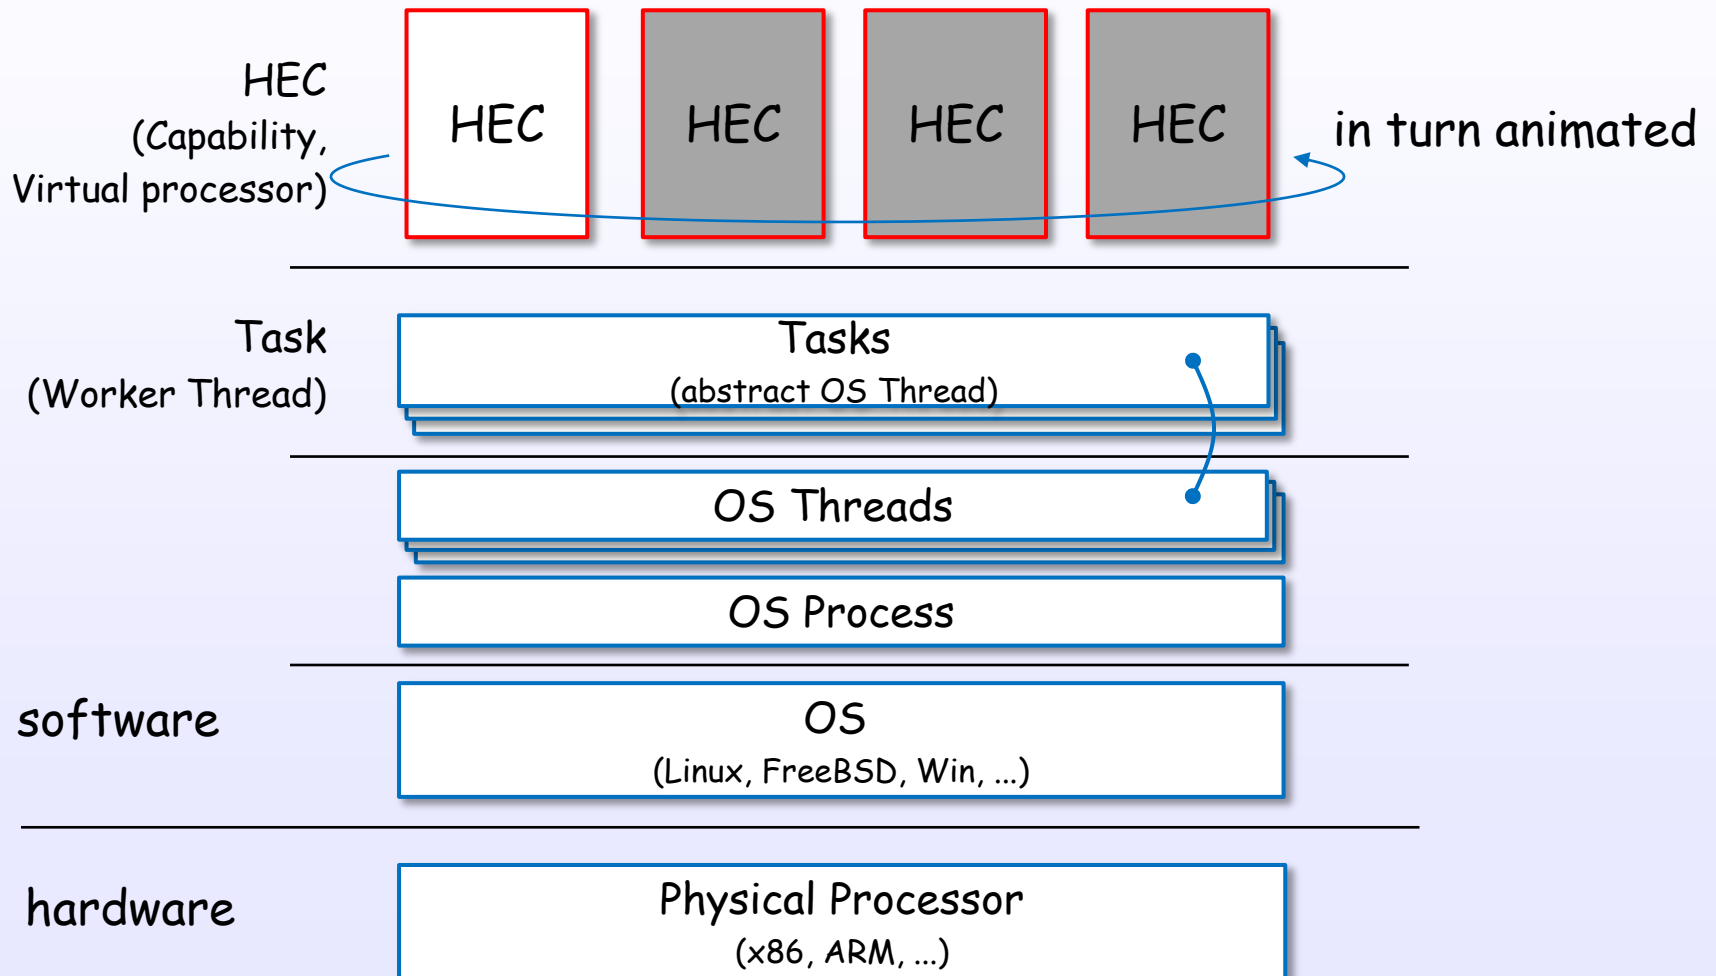


many HECs

Multi HECs can be generated by compile and runtime options :

```
$ ghc -rtsopts -threaded
```

```
$ ./xxx +RTS -N4
```



HEC (Capability) data structure

[rts/Capability.h] (ghc 7.8)

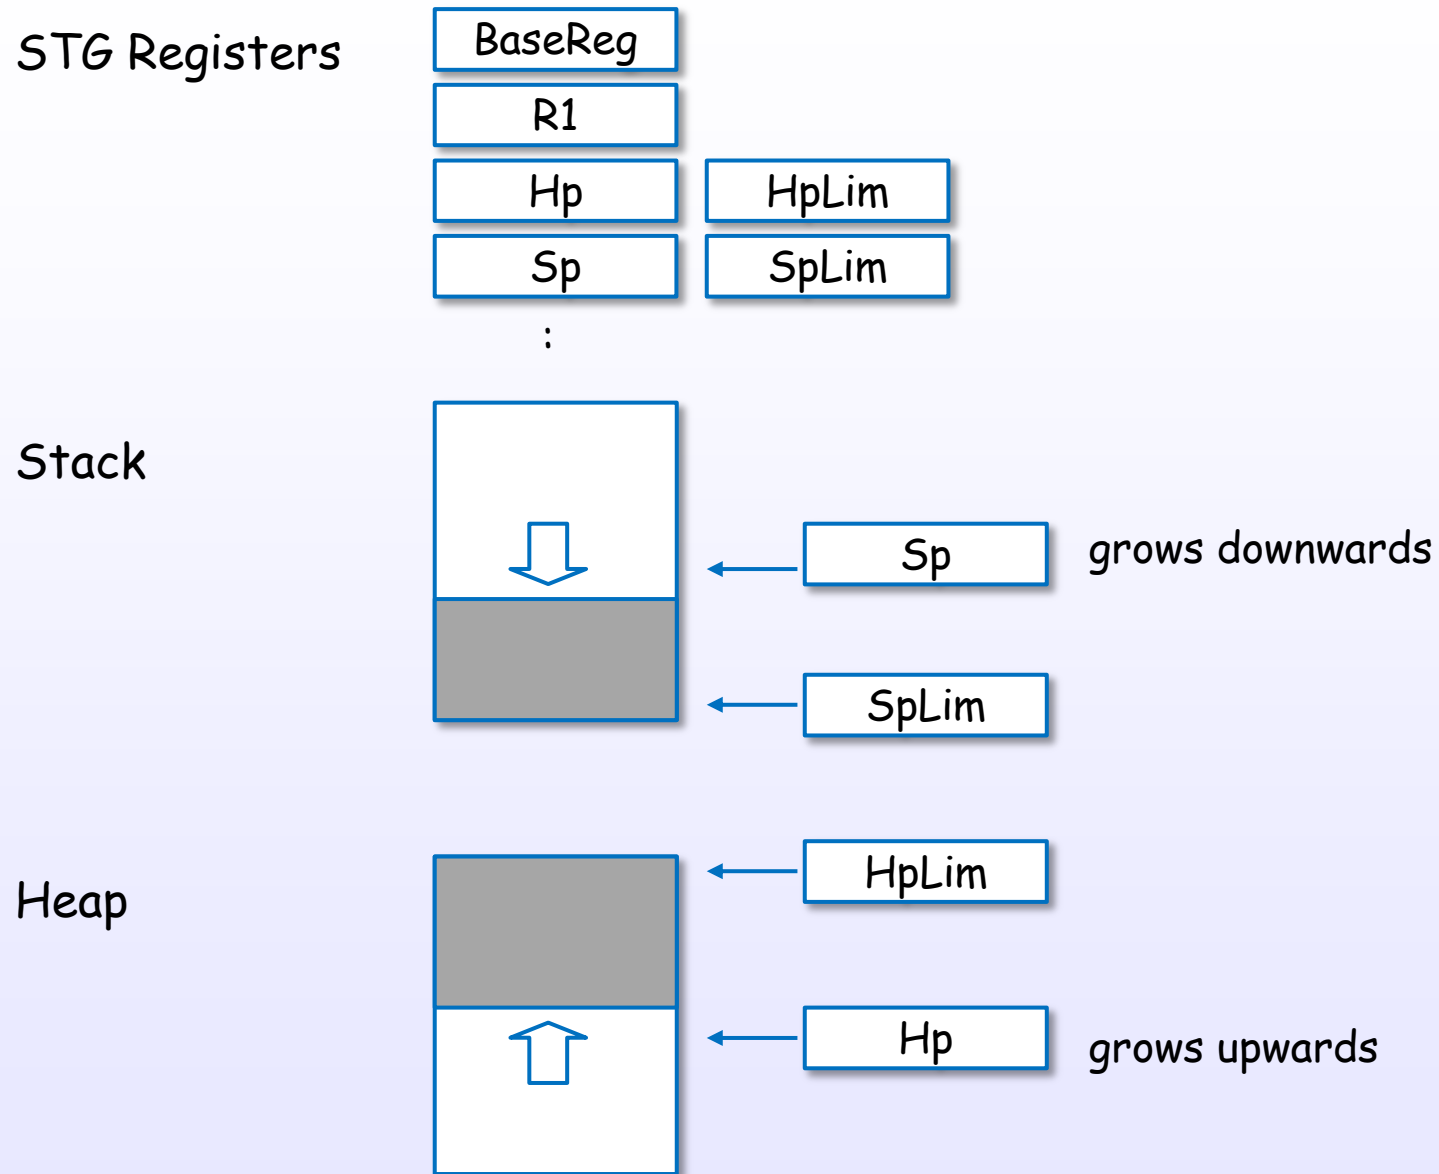
```
struct Capability_ {  
    StgFunTable f;  
    StgRegTable r;  
    nat no;  
    Task *running_task;  
    rtsBool in_haskell;  
    nat idle;  
    rtsBool disabled;  
    StgTSO *run_queue_hd;  
    StgTSO *run_queue_tl;  
    InCall *suspended_ccalls;  
    bdescr **mut_lists;  
    bdescr **saved_mut_lists;  
    bdescr *pinned_object_block;  
    bdescr *pinned_object_blocks;  
    int context_switch;  
    int interrupt;  
  
#if defined(THREADED_RTS)  
    Task *spare_workers;  
    nat n_spare_workers;  
    Mutex lock;  
    Task *returning_tasks_hd;  
    Task *returning_tasks_tl;  
    Message *inbox;  
    SparkPool *sparks;  
    SparkCounters spark_stats;  
#endif  
  
    W_ total_allocated;  
    StgTVarWatchQueue *free_tvar_watch_queues;  
    StgInvariantCheckQueue *free_invariant_check_queues;  
    StgTRecChunk *free_trec_chunks;  
    StgTRecHeader *free_trec_headers;  
    nat transaction_tokens;  
}
```

Each HEC (Capability) has a register table and a run queue and ...

Each HEC (Capability) is initialized at initCapabilities [rts/rts/Capability.c]

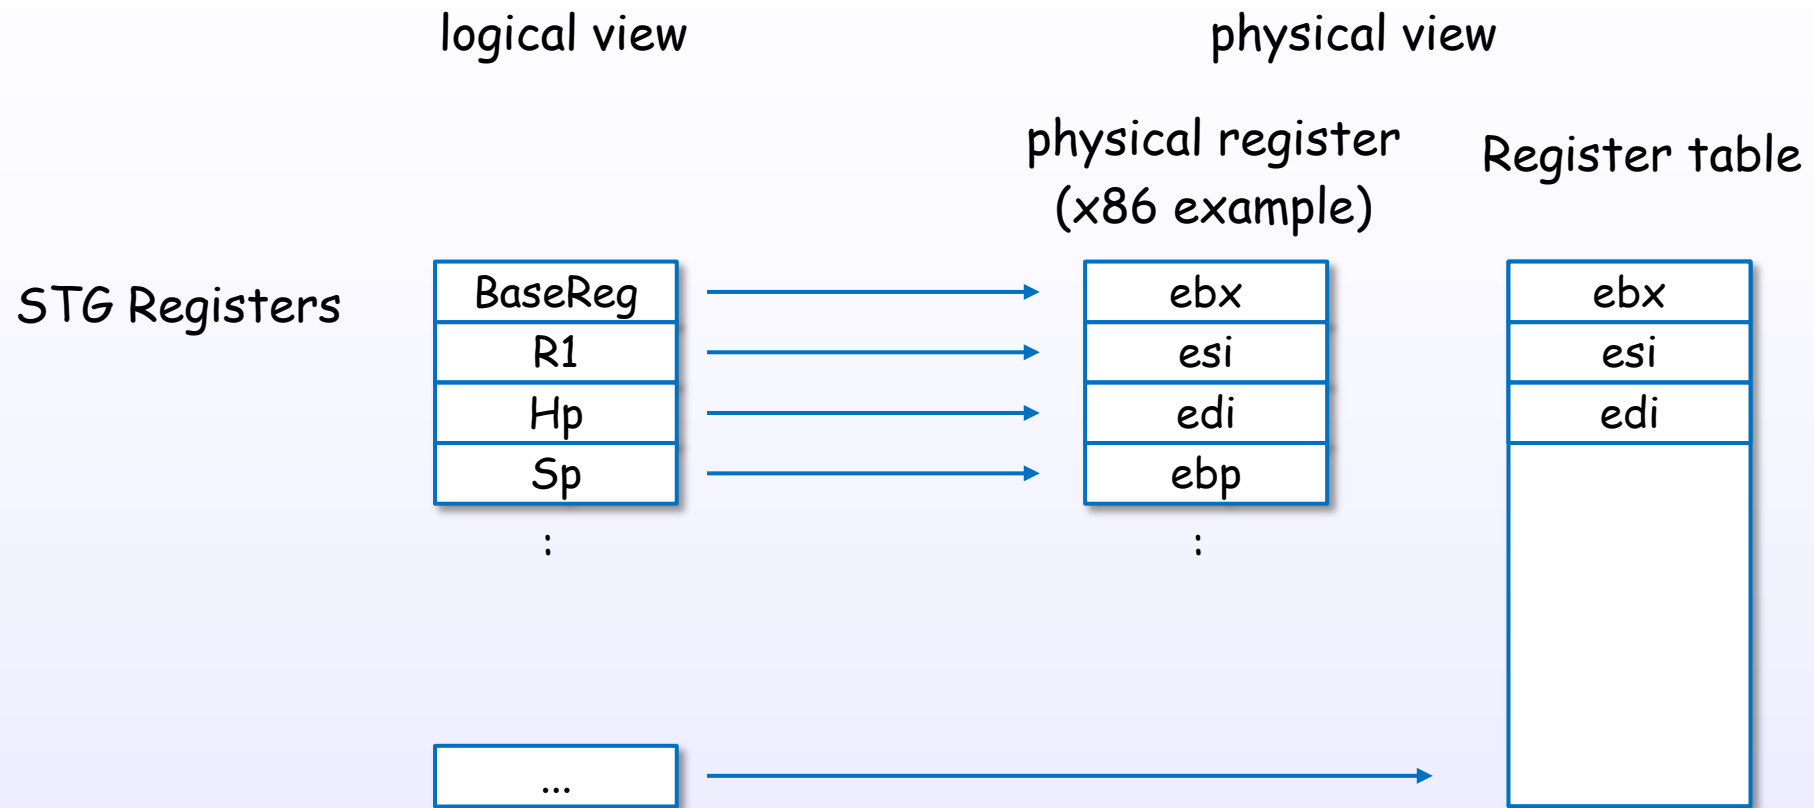
STG-machine

STG-machine

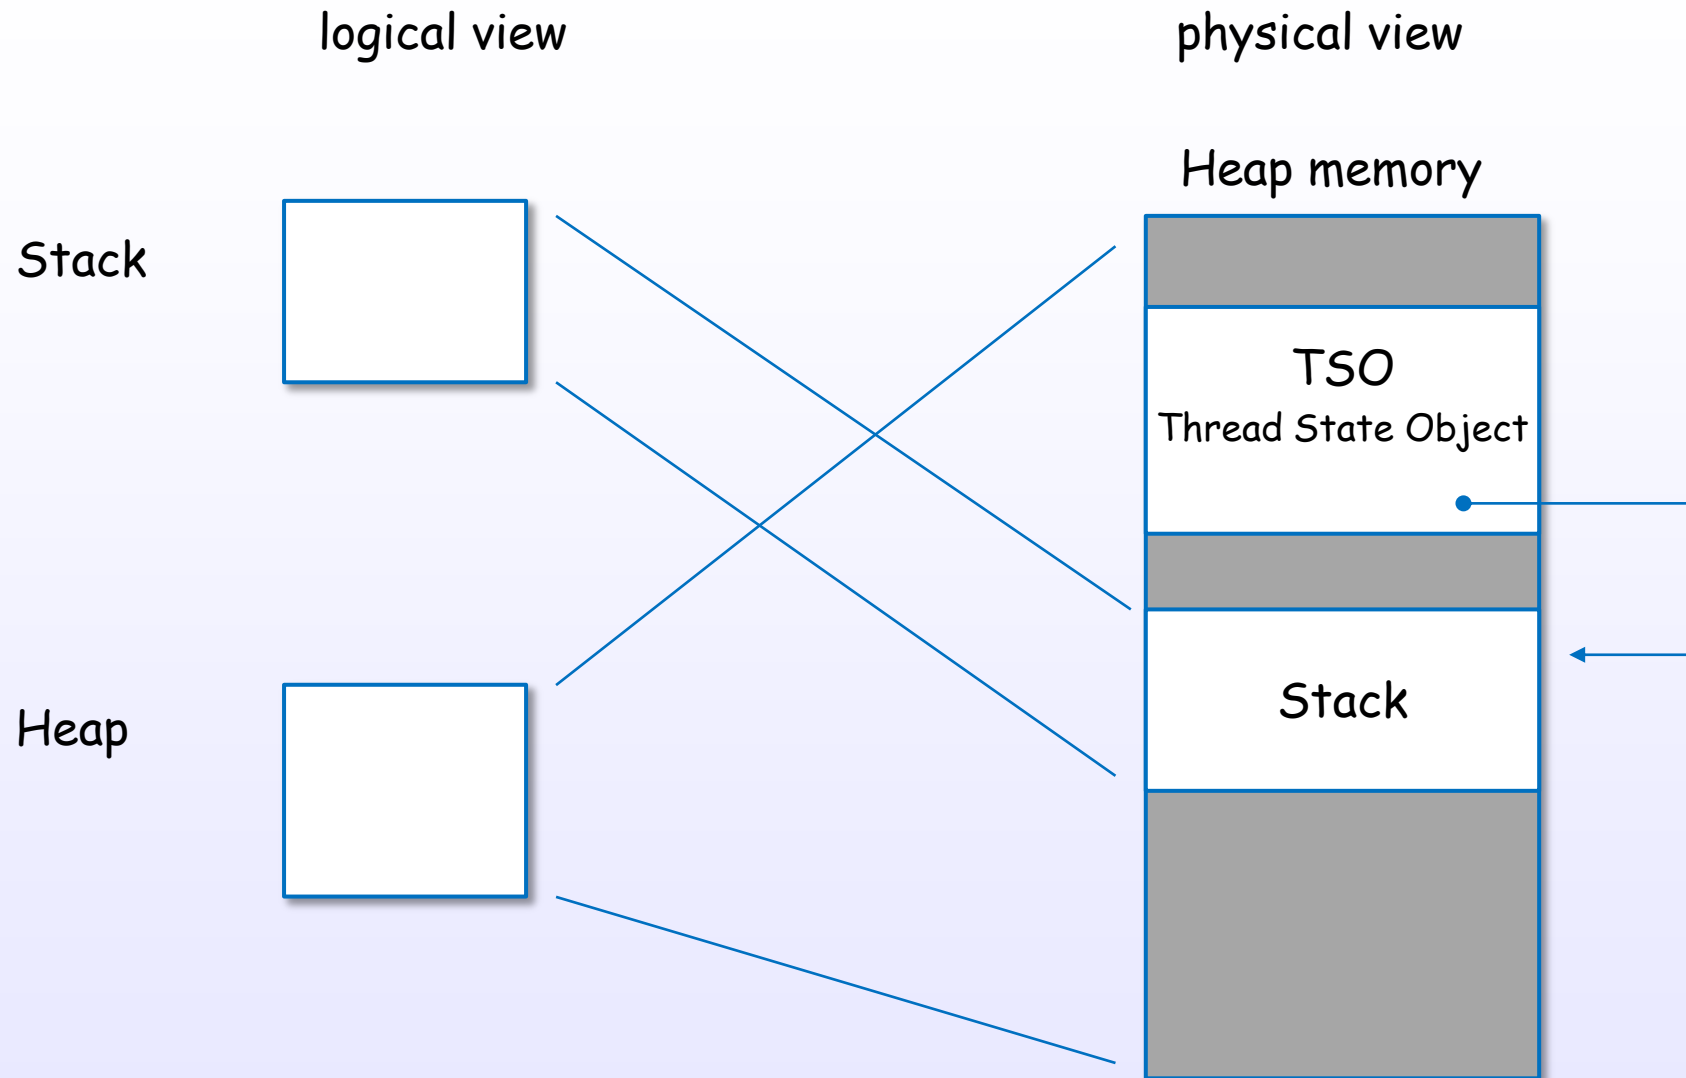


Each real Haskell code is executed in STG semantics.

STG-machine is mapped to physical processor



STG-machine is mapped to physical processor



A stack and a TSO object are in the heap.
The stack is stored separately from the TSO for size extension and GC.

TSO data structure

[includes/rts/storage/TSO.h] (ghc 7.8)

```
typedef struct StgTSO_  
  StgHeader          header;  
  struct StgTSO_*     _link;  
  struct StgTSO_*     global_link;  
  struct StgStack_*   *stackobj;  
  StgWord16           what_next;  
  StgWord16           why_blocked;  
  StgWord32           flags;  
  StgTSOBlockInfo     block_info;  
  StgThreadID         id;  
  StgWord32           saved_errno;  
  StgWord32           dirty;  
  struct InCall_*      bound;  
  struct Capability_*  cap;  
  struct StgTRecHeader_* trec;  
  struct MessageThrowTo_* blocked_exceptions;  
  struct StgBlockingQueue_* bq;  
  StgWord32 tot_stack_size;  
} *StgTSOPtr;
```

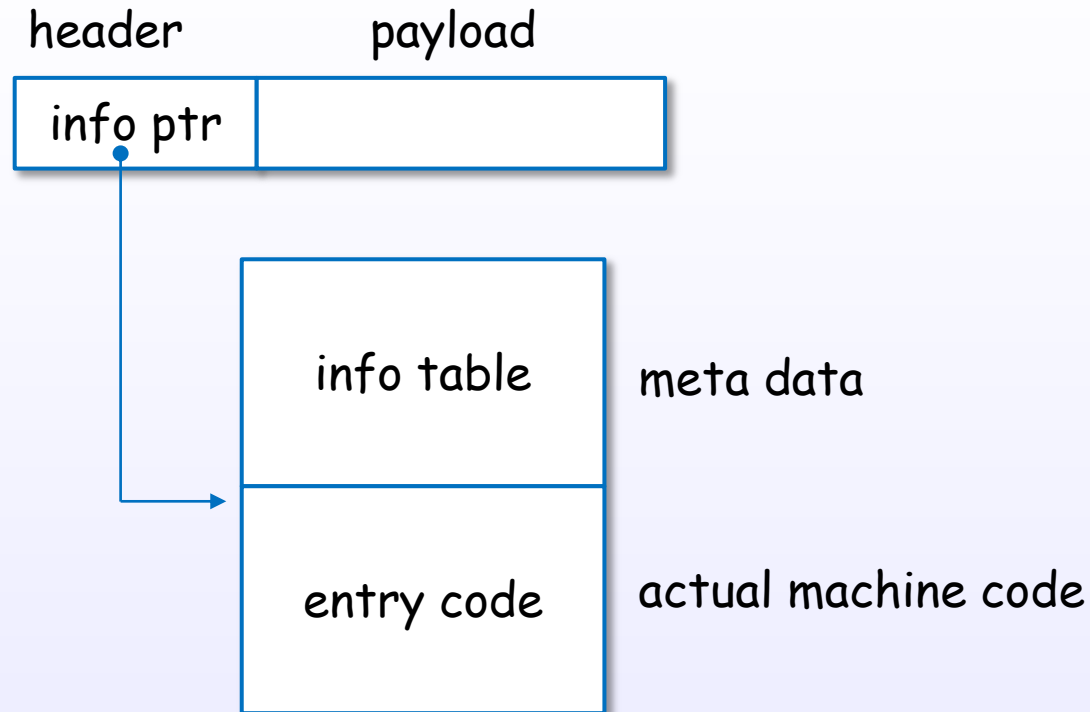
link to stack object

A TSO object is **only ~17words + stack**. Lightweight!

Heap objects in STG-machine

Heap object (closure)

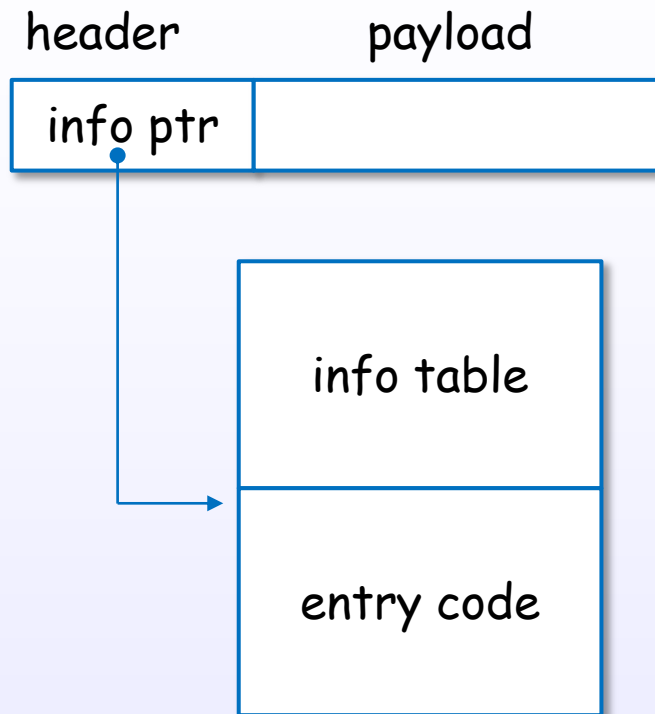
logical view




Closure (header + payload) + Info Table + Entry Code

Heap object (closure)

logical view

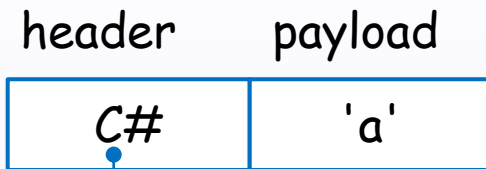


physical view



Closure examples : Char, Int

'a' :: Char



info
ptr

GHC.Types.C#_static_info

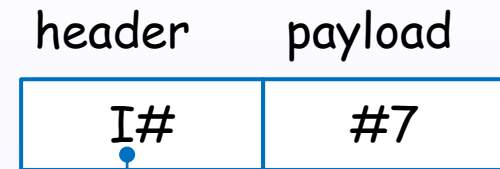
layout : 0_1
type : CONSTR
bitmap :

info table

inc %esi
jmp *0x0(%ebp)

entry code

7 :: Int



info
ptr

GHC.Types.I#_static_info

layout : 0_1
type : CONSTR
bitmap :

info table

inc %esi
jmp *0x0(%ebp)

entry code

Closure example code

[Example.hs]

```
module Example where
value1 :: Int
value1 = 7
```

STG

Cmm

[ghc -O -ddump-stg Example.hs]

```
Example.value1 :: GHC.Types.Int
[GblId, Caf=NoCafRefs, Str=DmdType m, Unf=OtherCon []] =
  NO_CCS GHC.Types.I#! [8];
```

[ghc -O -ddump-opt-cmm Example.hs]

```
section "data" { __stginit_main:Example:
}
```

```
section "data" {
  Example.value1 closure:
  const GHC.Types.I#_static_info;
  const 7;
}
```

asm

[ghc -O -ddump-asm Example.hs]

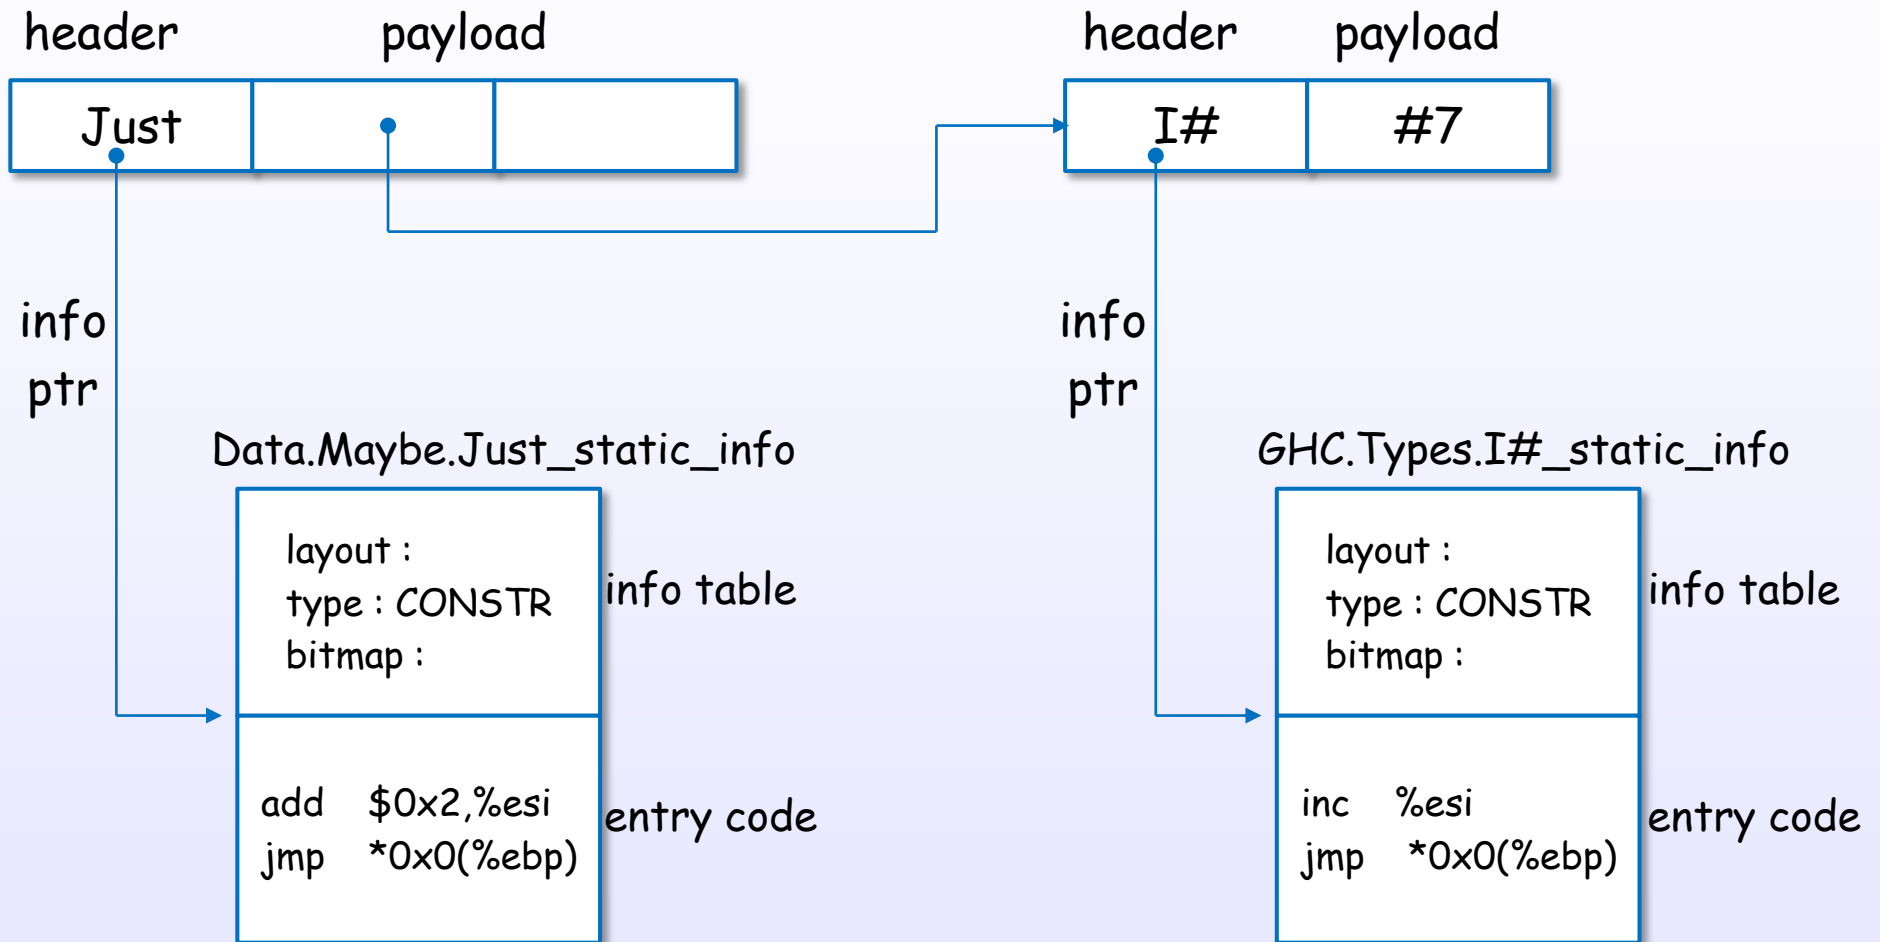
```
.data
    .align 4
    .align 1
    .globl __stginit_main:Example
    __stginit_main:Example:
    .data
        .align 4
        .align 1
        .globl Example.value1_closure
    Example.value1_closure:
        .long GHC.Types.I#_static_info
        long 7
```

```
.section .data
    .align 4
    .align 1
    SMd_srt:
```

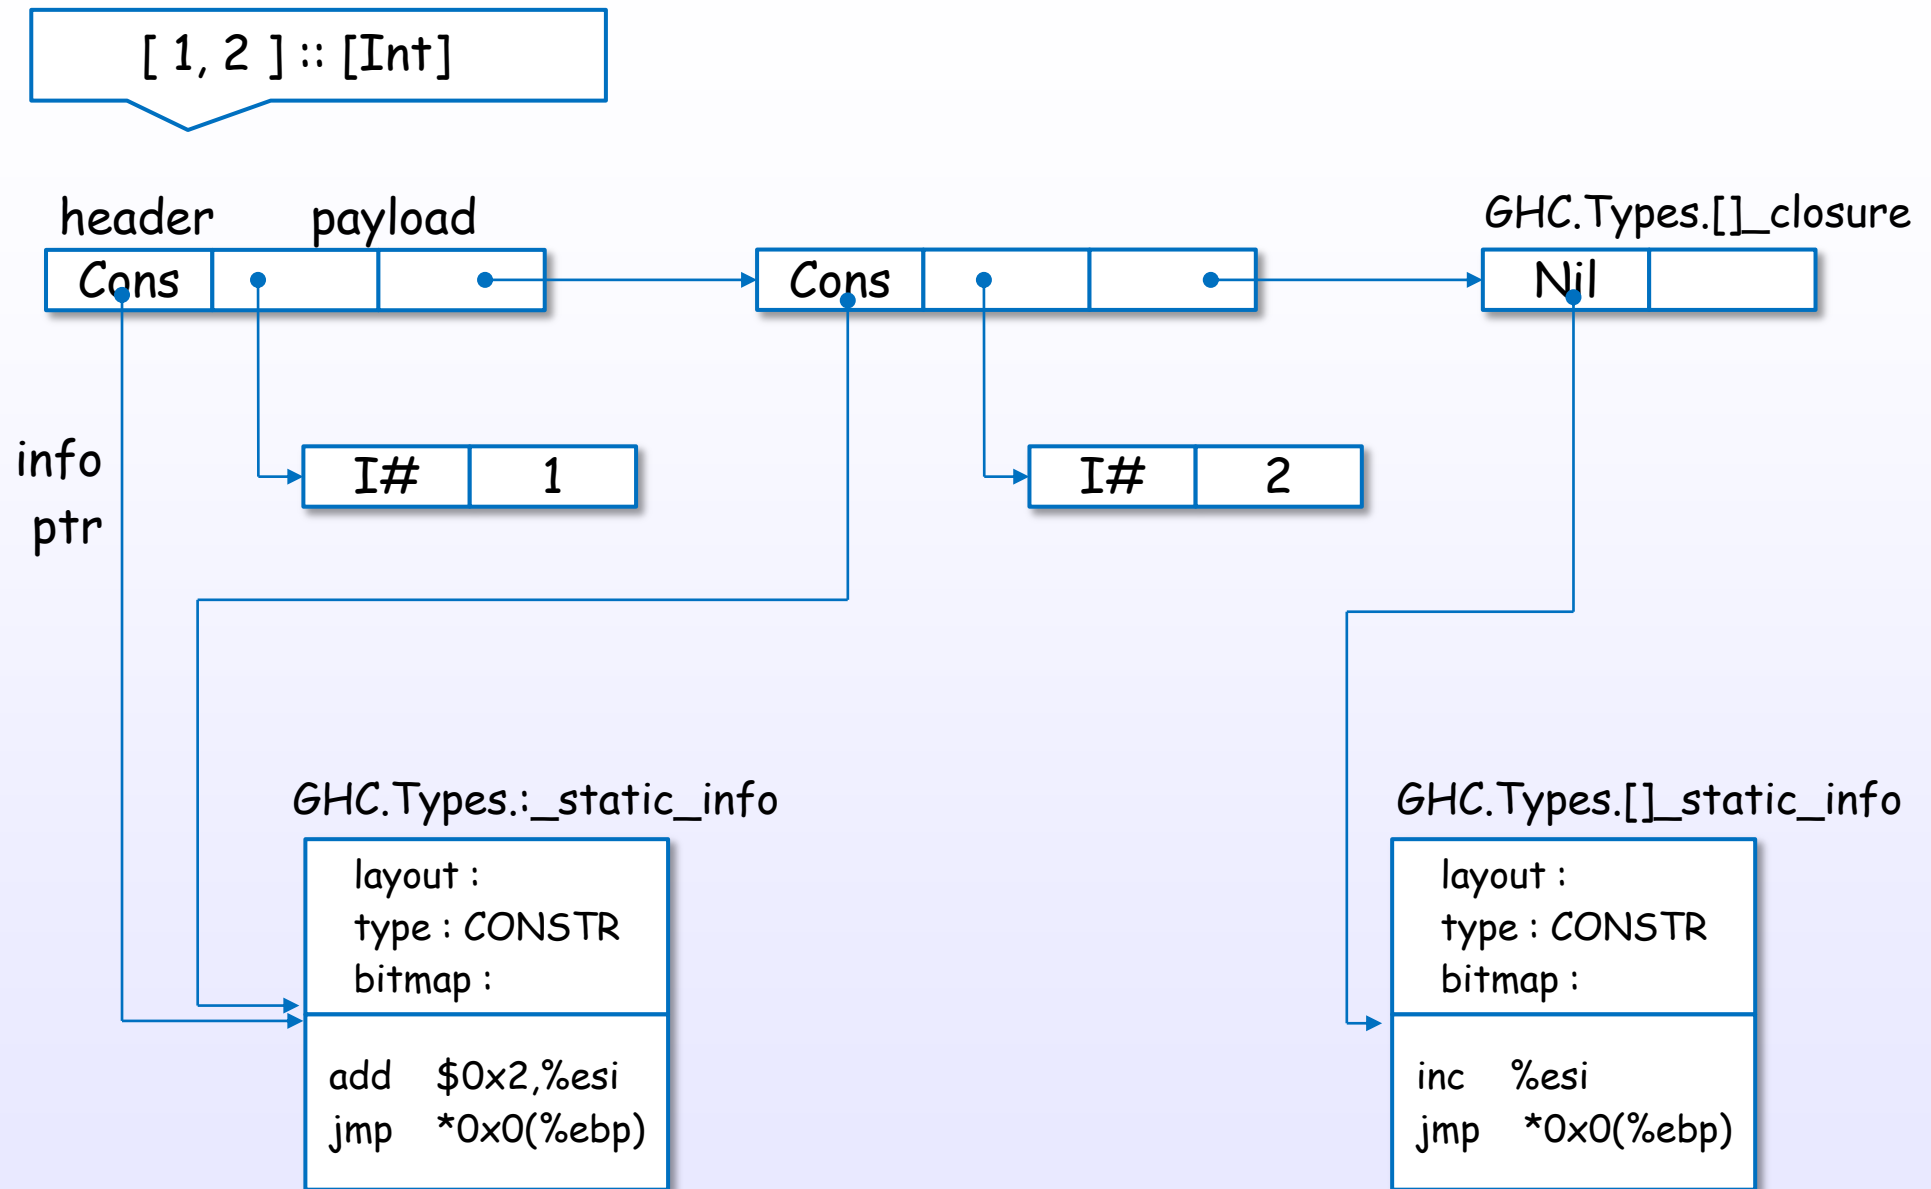
header		payload	
I#		#7	

Closure examples : Maybe

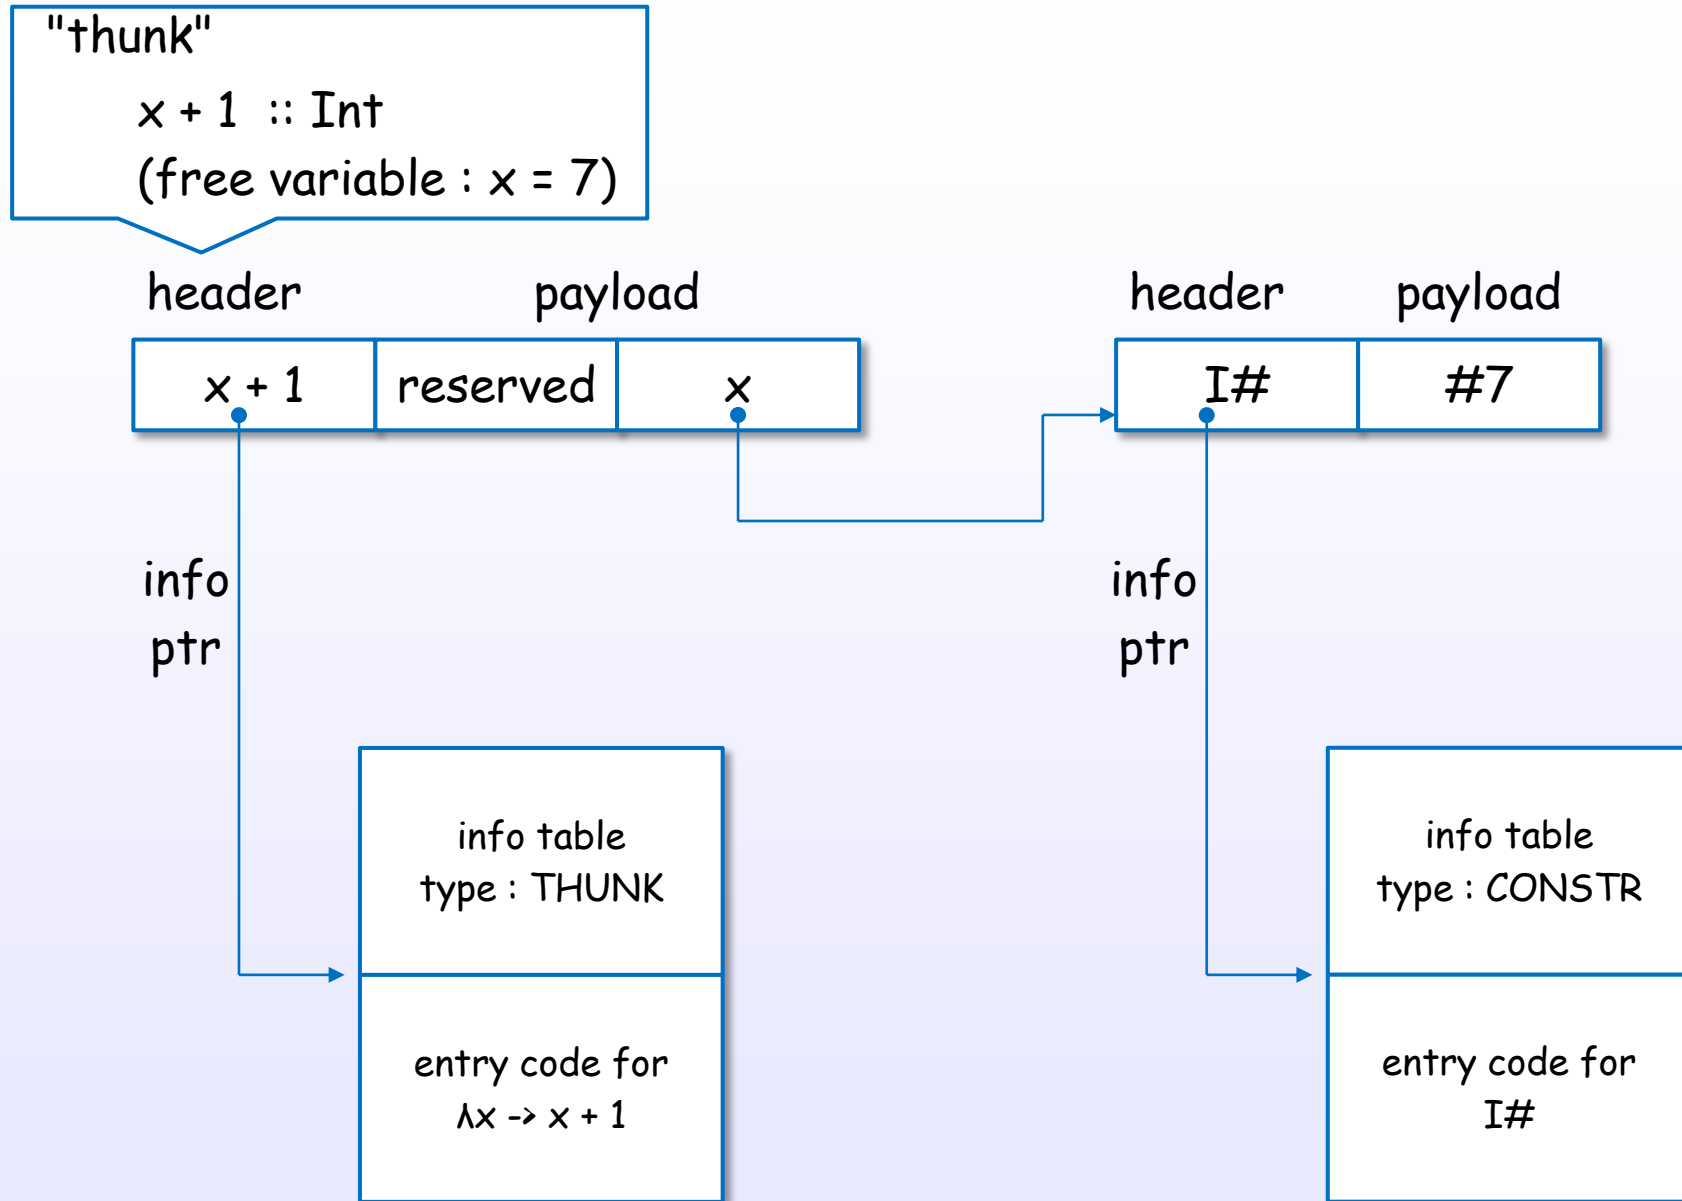
Just 7 :: Maybe Int



Closure examples : List

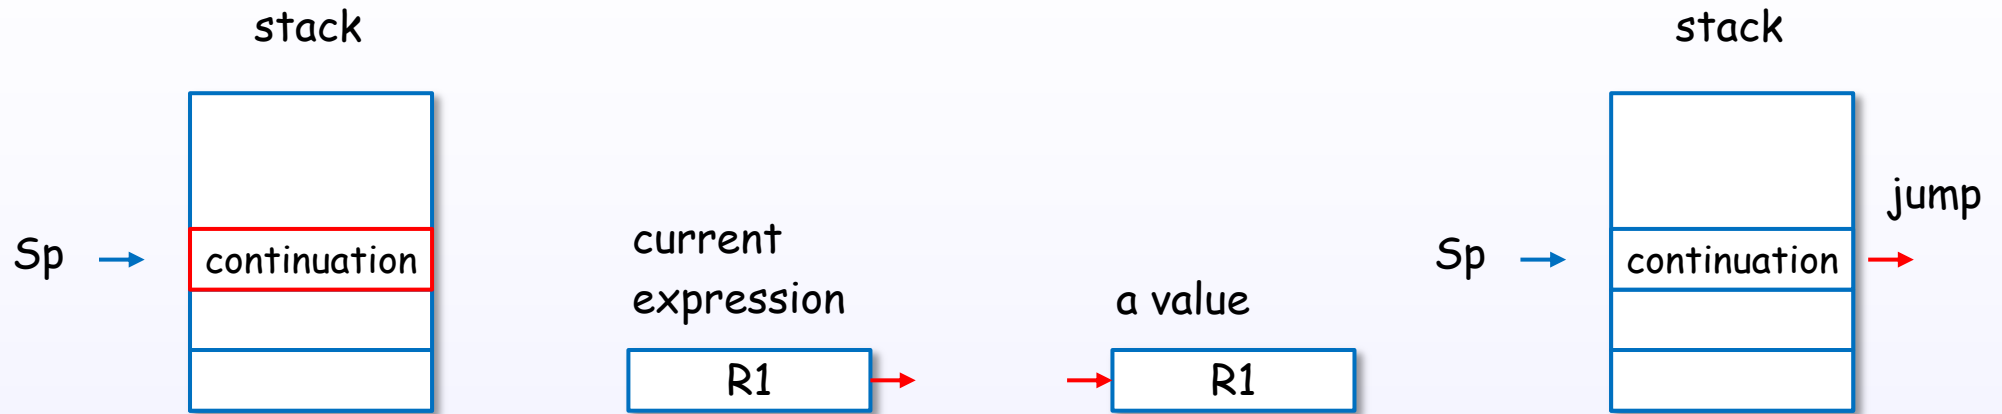


Closure examples : Thunk



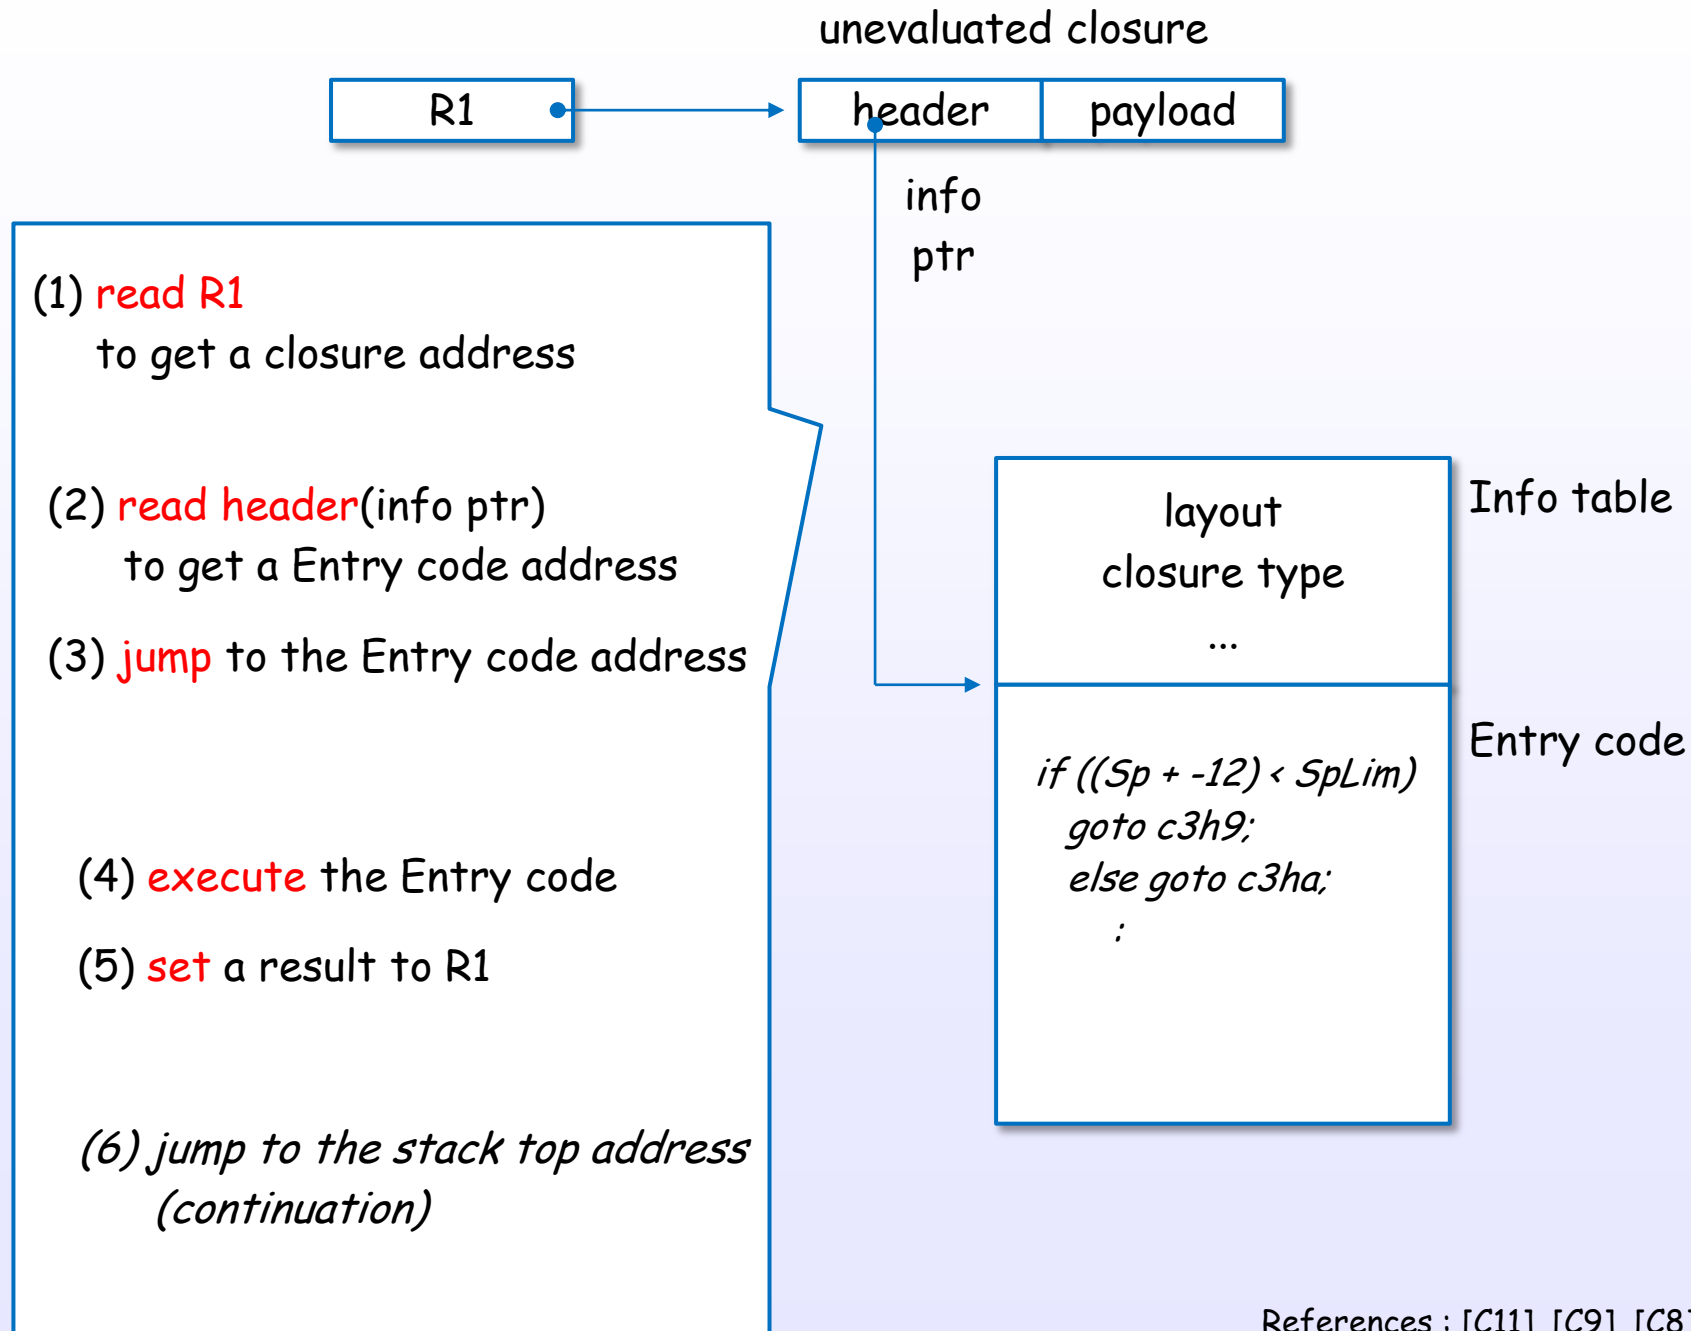
STG-machine evaluation

STG evaluation flow



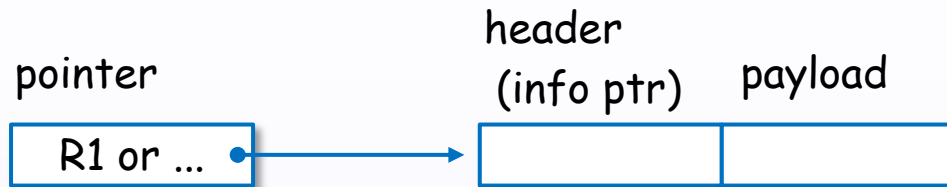
- (1) push a continuation code (next code) to the stack top
- (2) enter to R1 closure
- (3) set a result to R1
- (4) jump (return) to the stack top code
- (5) repeat from (1)

Enter to a closure



Pointer tagging

Pointer tagging



pointer



... an unevaluated closure



... an evaluated closure;
1st constructor value or evaluated.
(for instance: "Nothing")



... an evaluated closure; 2nd constructor value.
(for instance: "Just xx")



... an evaluated closure; 3rd constructor value.

* 32bit machine case

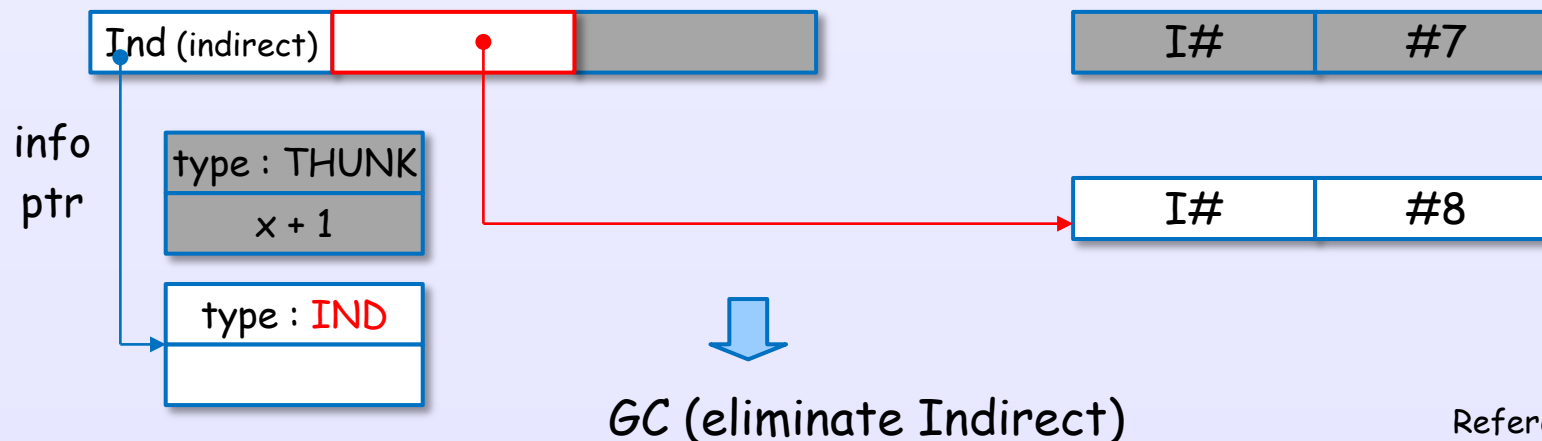

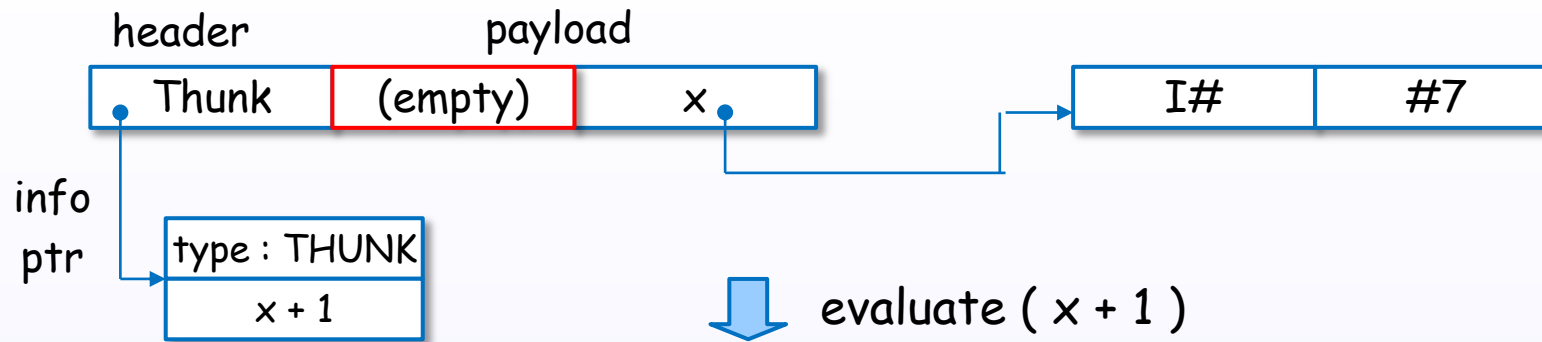
fast judgment!

checking only pointer's lower bits without evaluating the closure.

Think and update

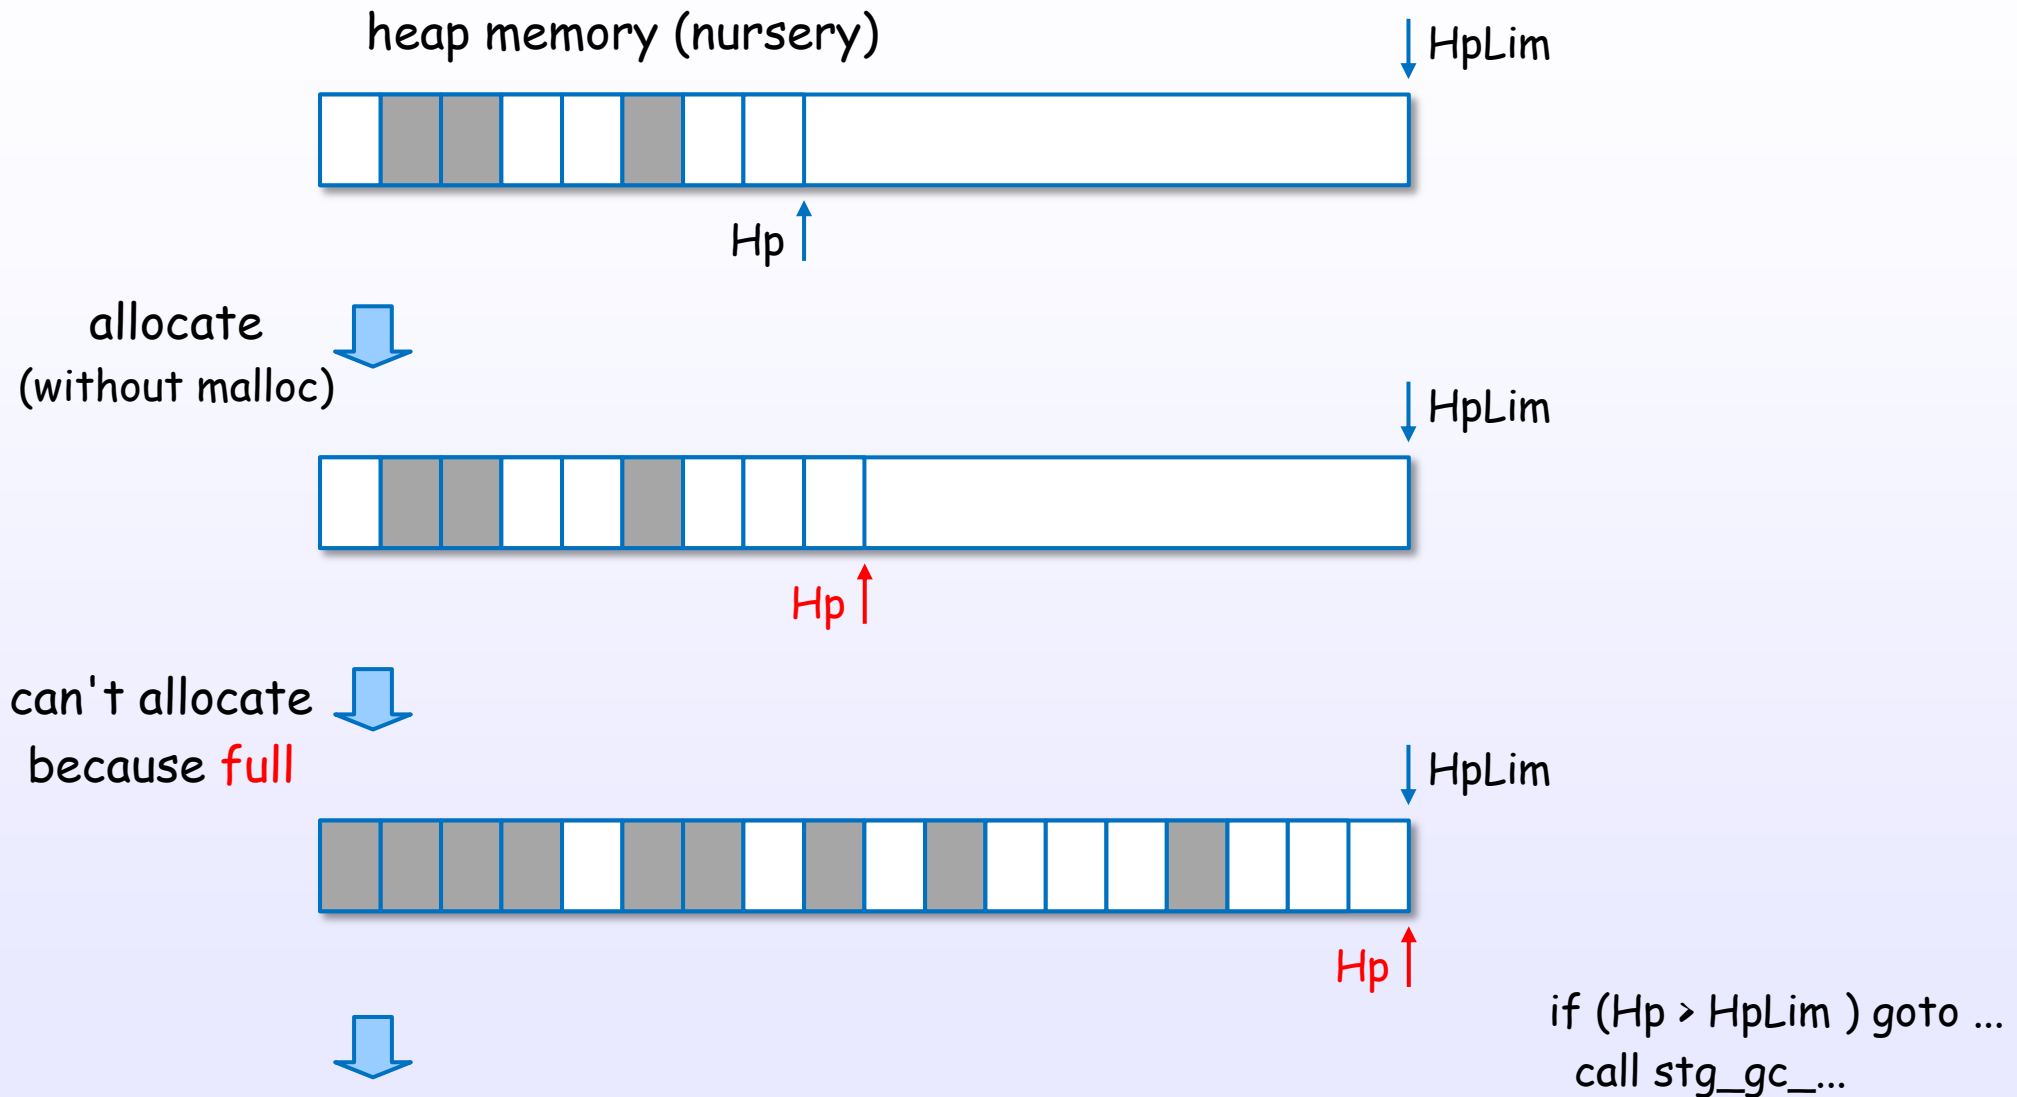
Thunk and update

"thunk" $x + 1 :: \text{Int}$ (free variable : $x = 7$)




Allocate and free heap objects

Allocate heap objects

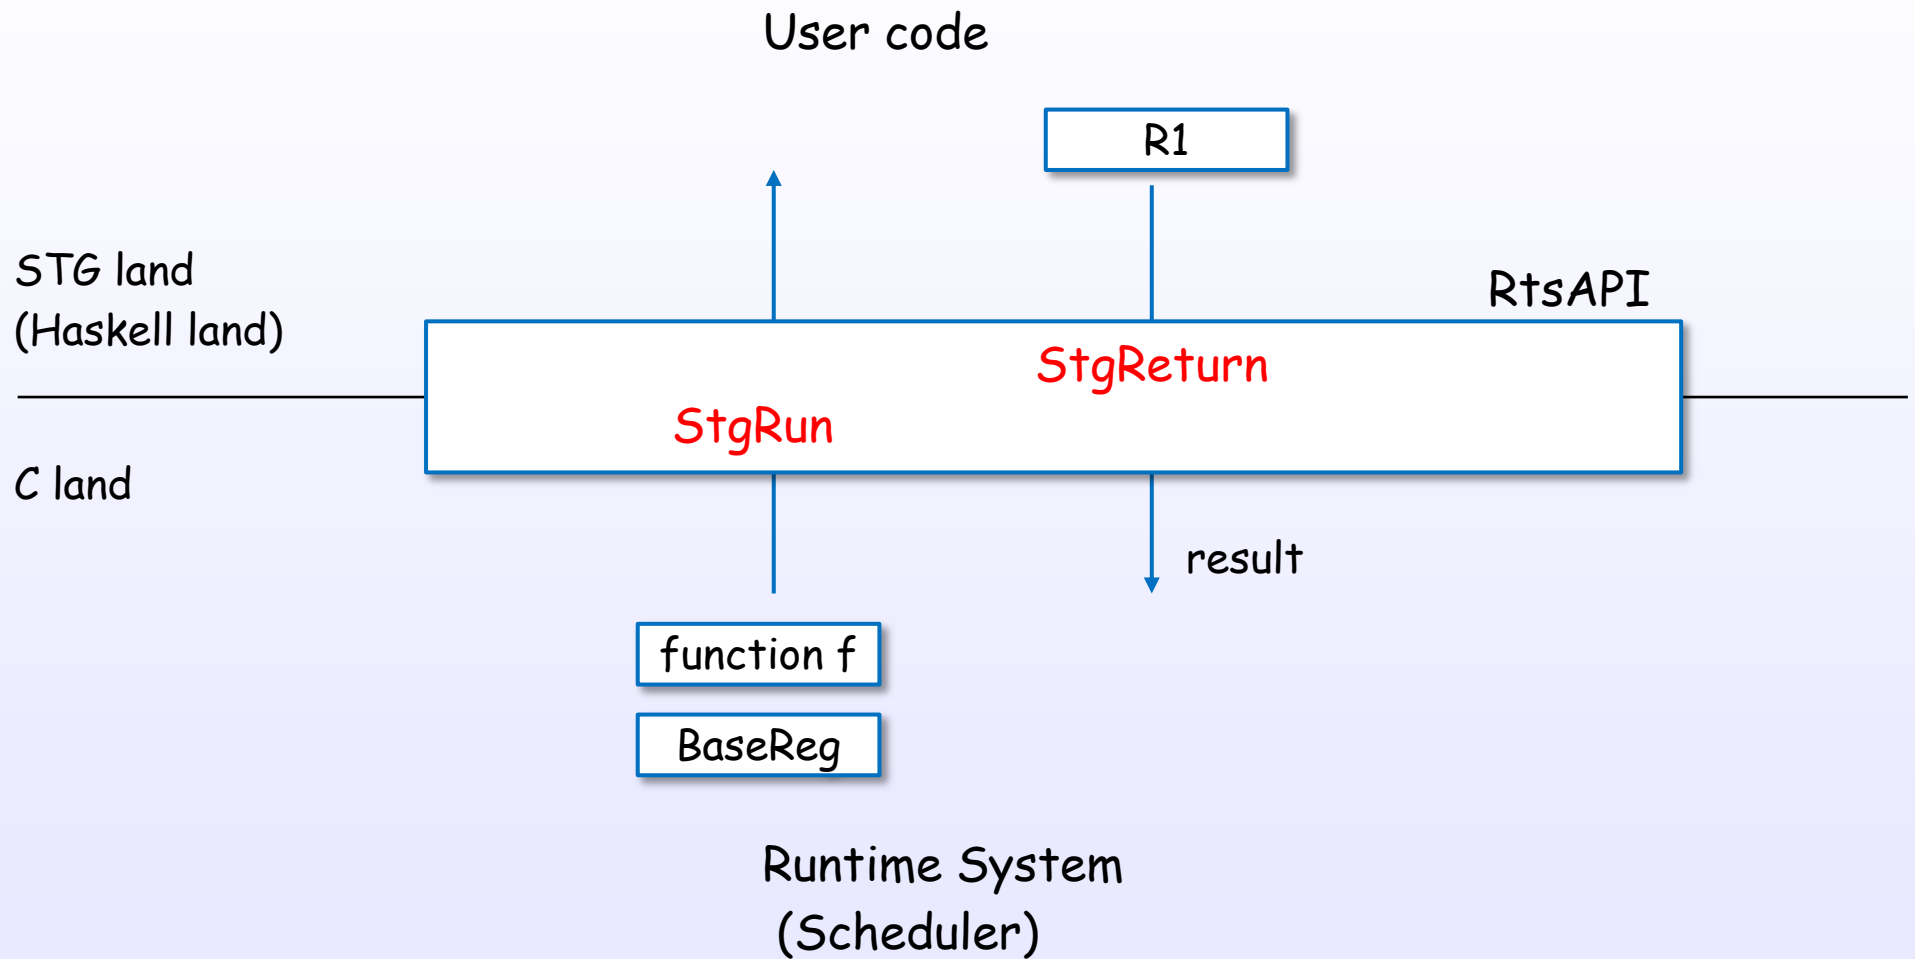


free and collect heap objects



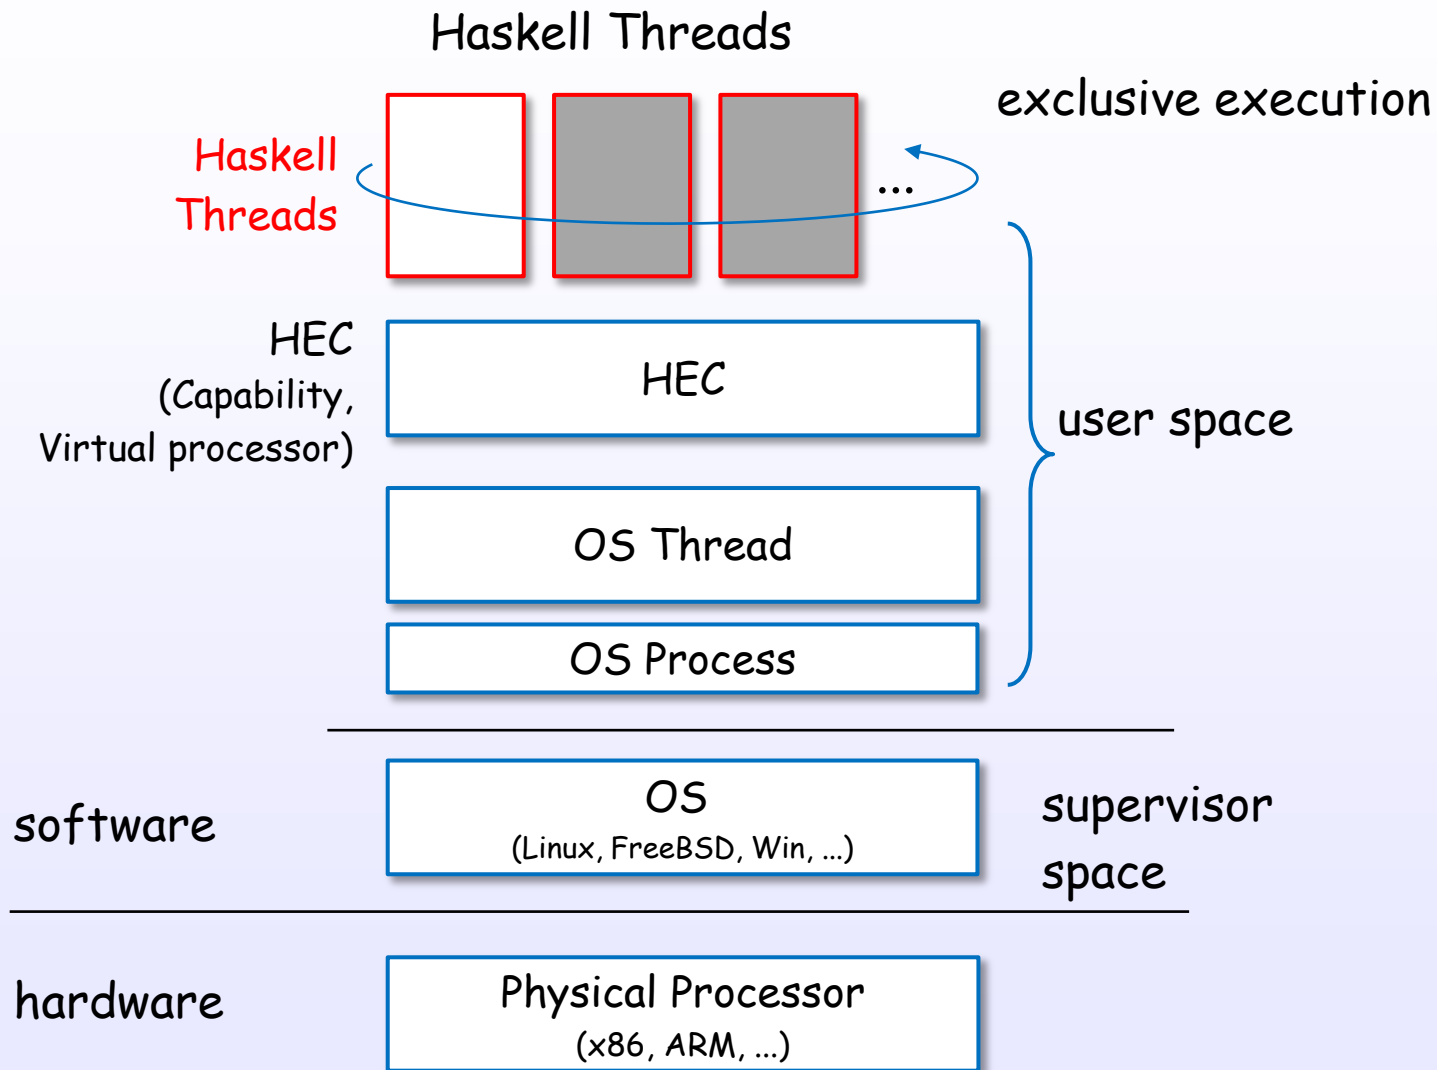
STG - C land interface

STG (Haskell) land - C land interface

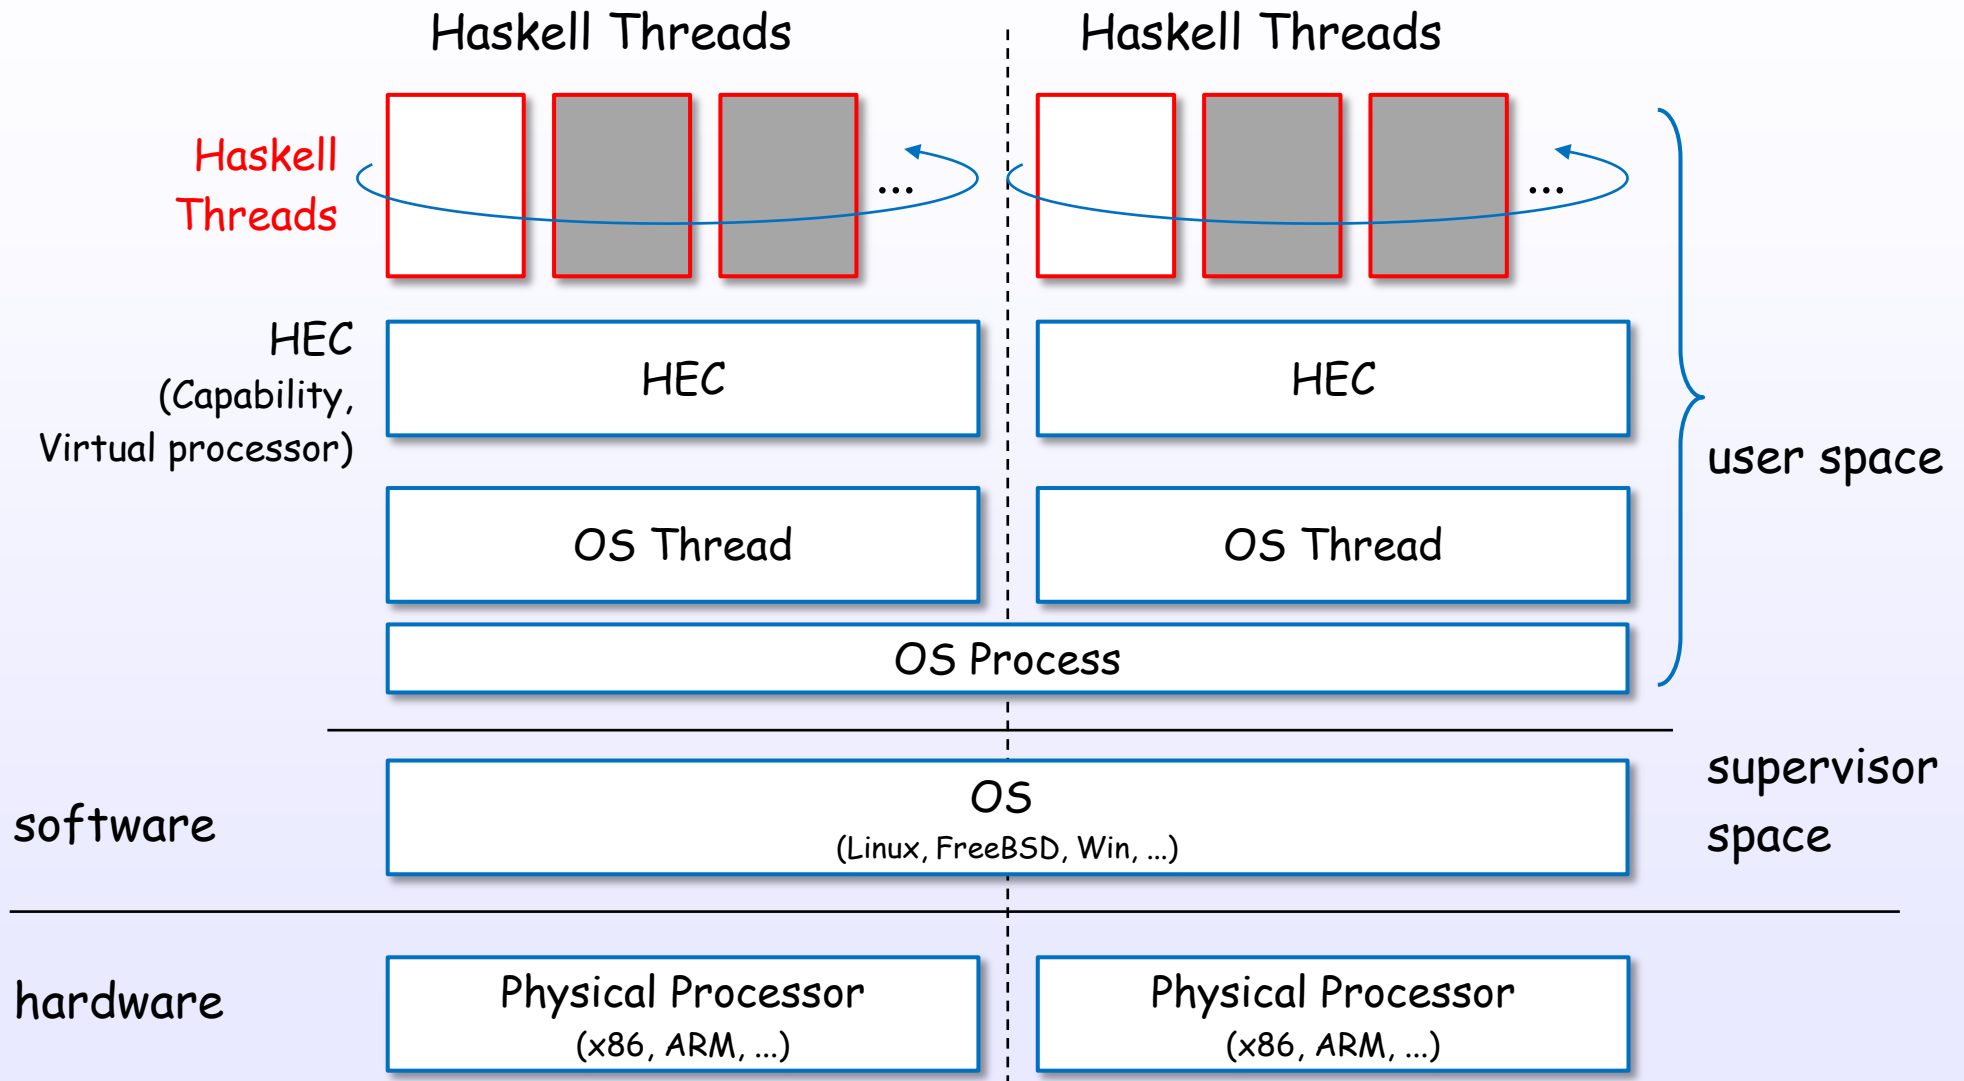


Thread

Thread layer (single core)



Thread layer (multi core)

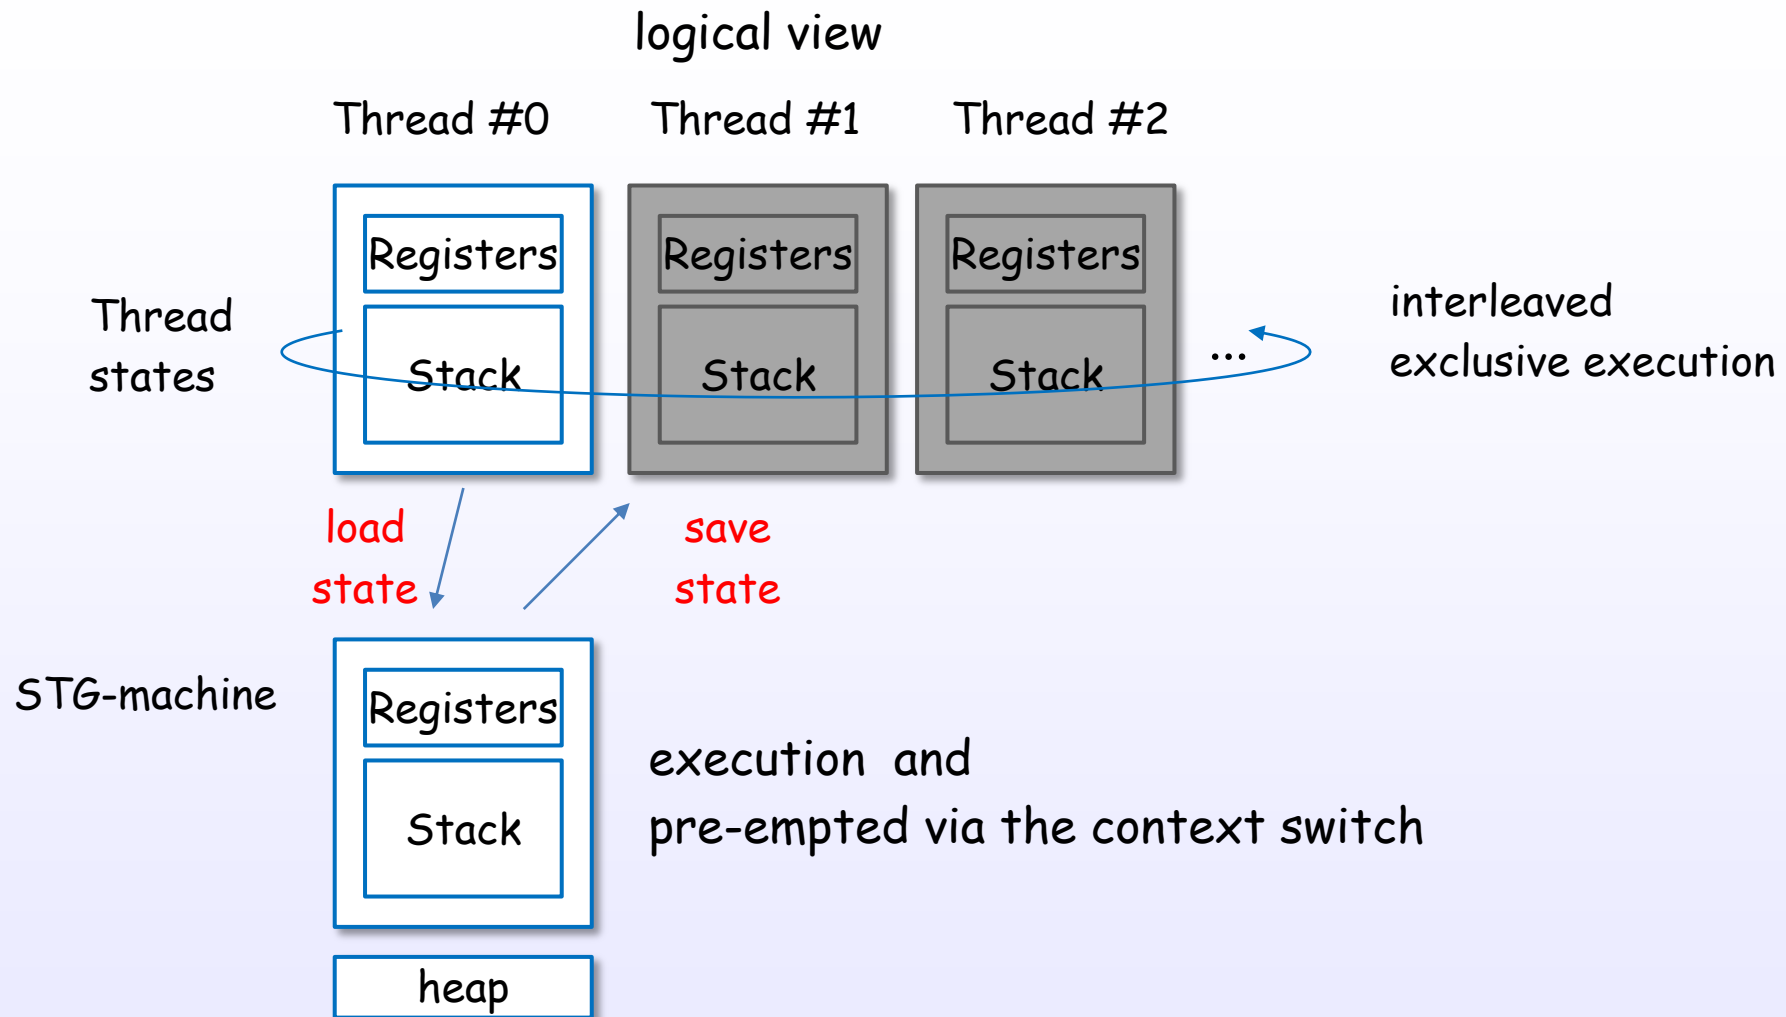


*Threaded option case (ghc -threaded)

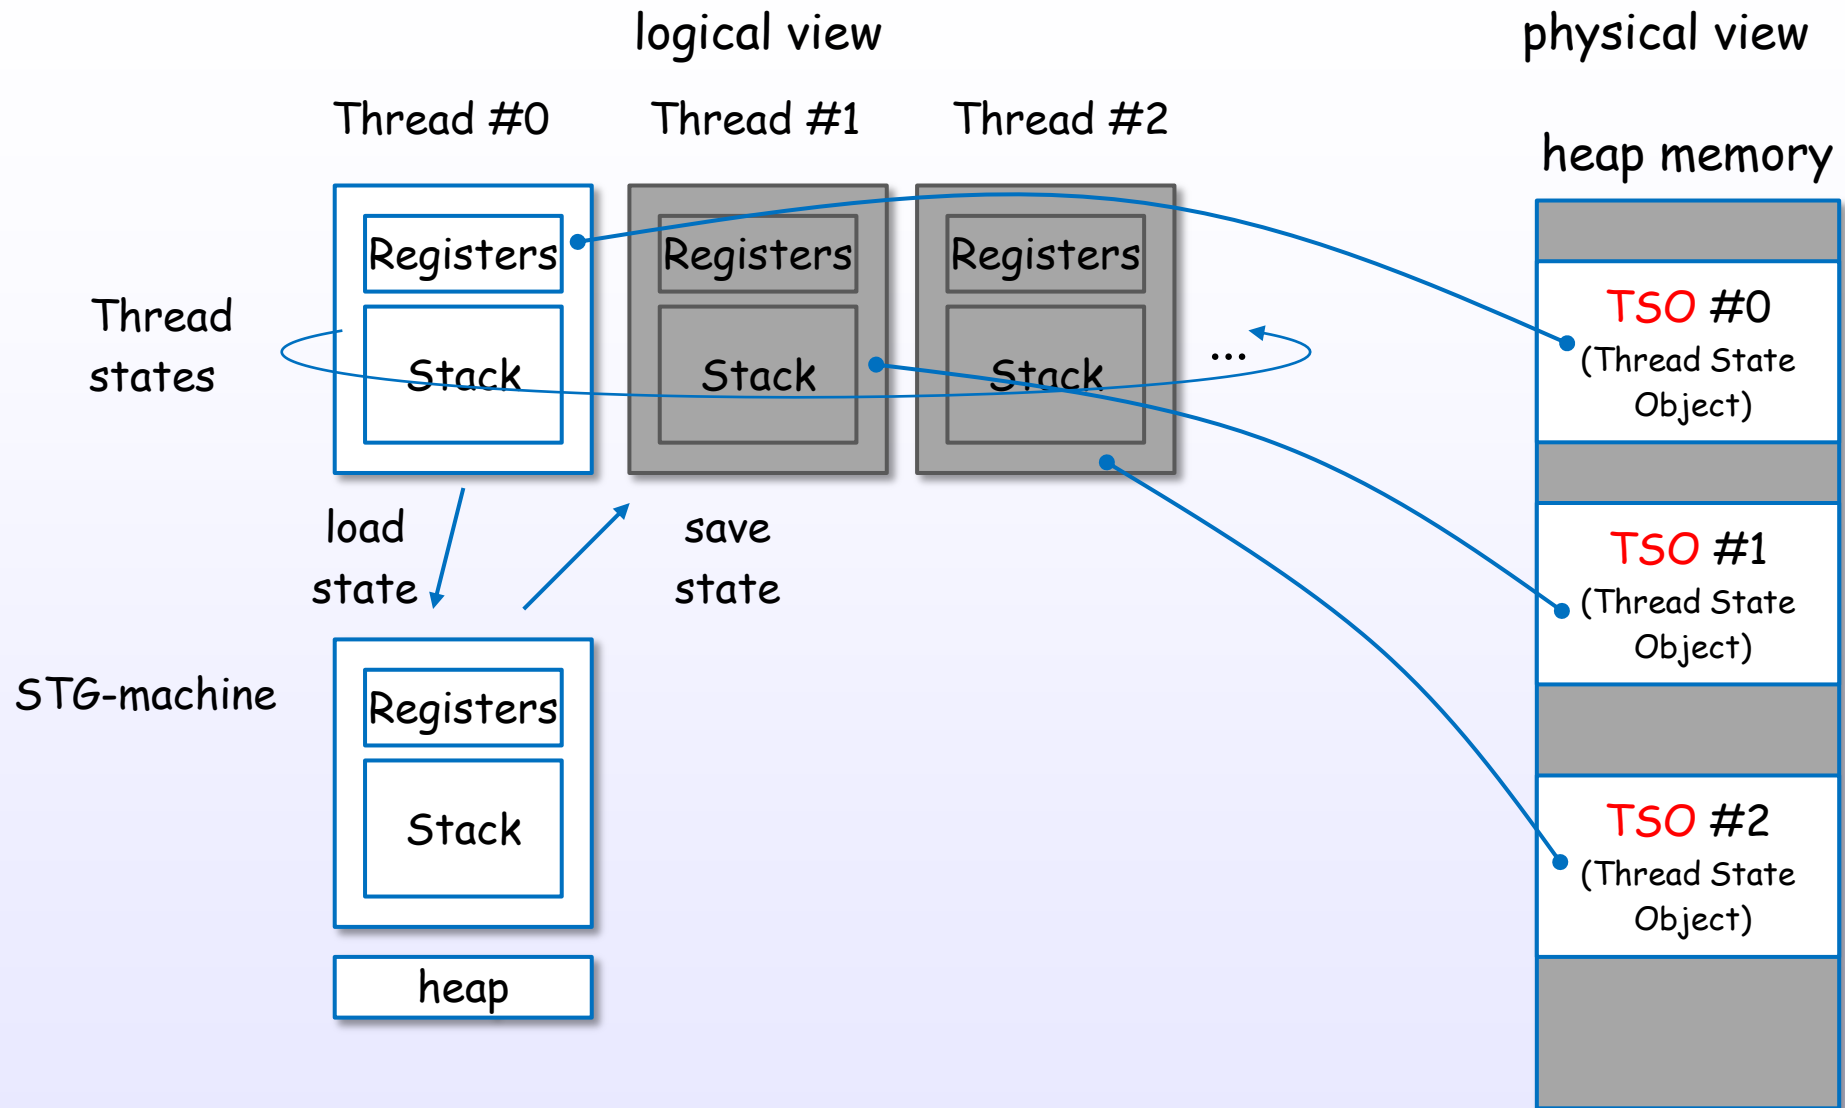
References : [5], [8], [9], [14], [C17], [C11], [19], [S17], [S16], [S23], [S22], [S14]

Thread context switch

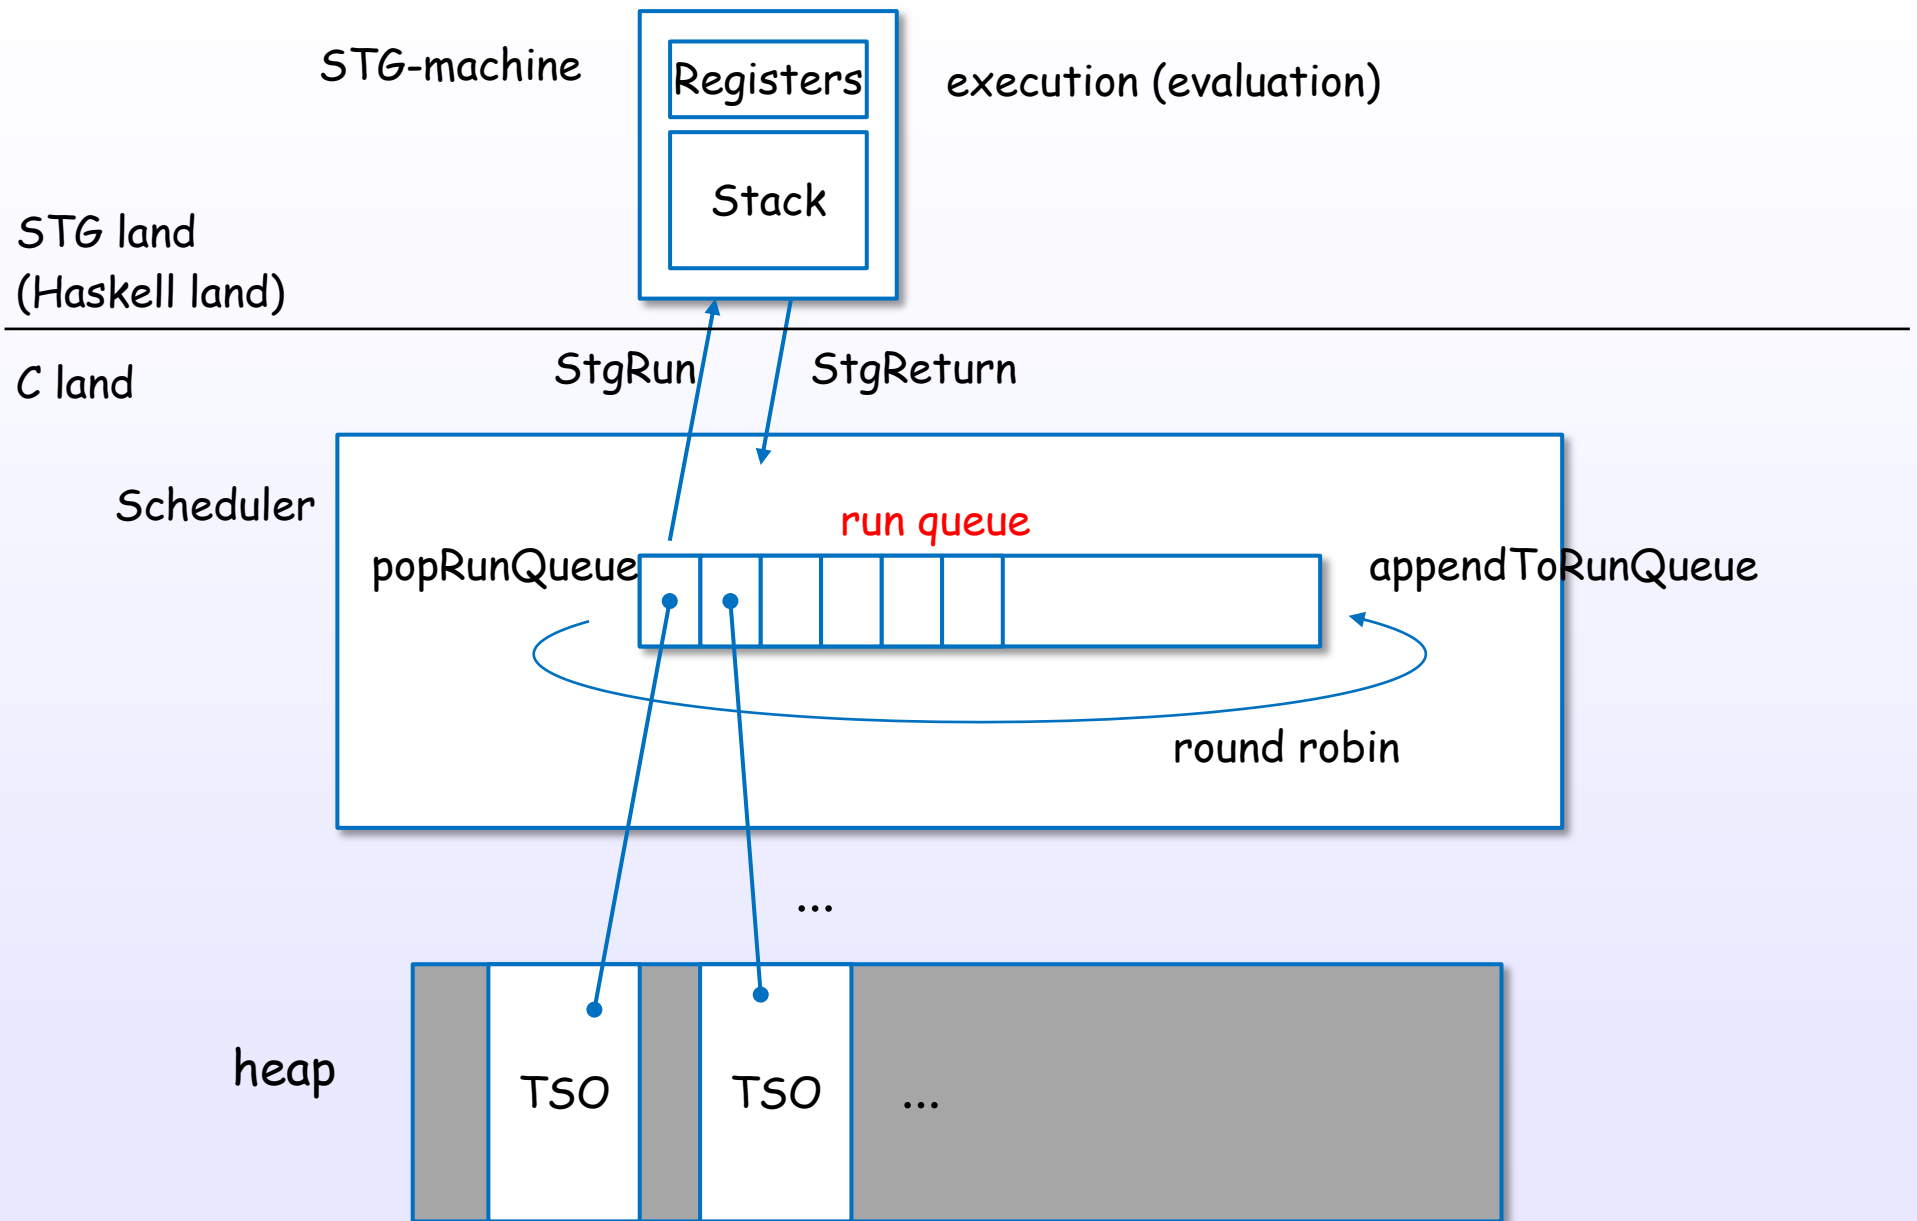
Threads and context switch



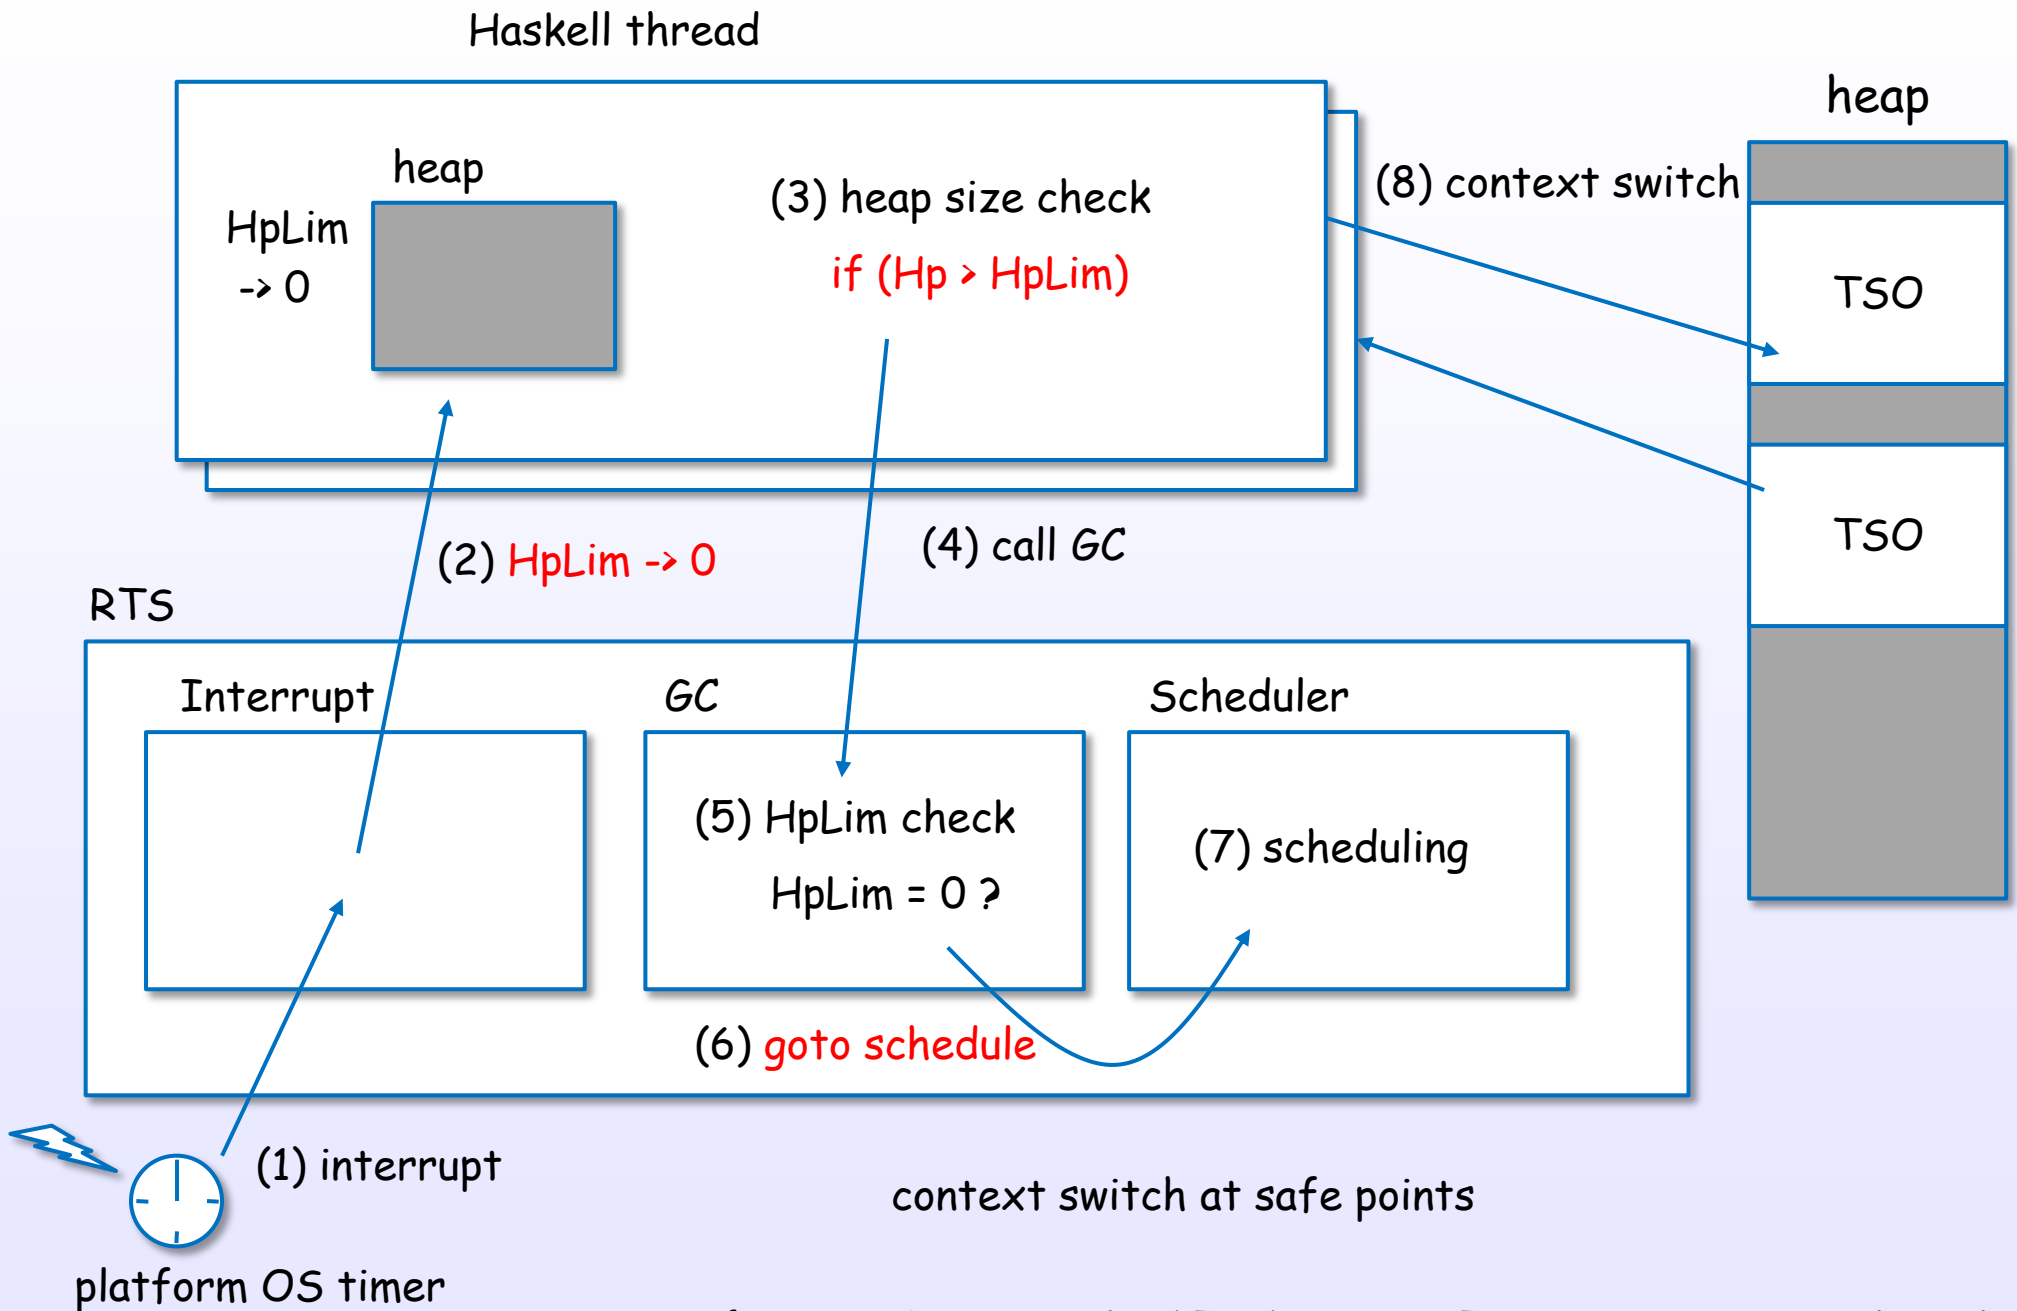
Threads and TSOs



Scheduling by run queue



Context switch flow



Context switch flow (code)

stg_gc_noregs

```
if (HpLim == 0) {  
    jump stg_returnToSched [R1];  
}
```

stg_returnToSched

```
W_ r1;  
r1 = R1; // foreign calls may clobber R1  
SAVE_THREAD_STATE();  
foreign "C" threadPaused(MyCapability()  
    "ptr", CurrentTSO);  
R1 = r1;  
jump StgReturn [R1];
```

STG land
(Haskell land)

C land

```
cap->r.rHpLim = NULL;
```

schedule

stopCapability
contextSwitchCapability
contextSwitchAllCapabilities
handle_tick

CreateTimerQueue

initTicker

initTimer

startTimer

hs_init_ghc

real_main

next
handle_tick ..

OS

*Windows case

References : [5], [8], [9], [14], [C17], [C11], [19], [S17], [S16], [S23], [S22], [S14], [S24]

Creating main and sub threads

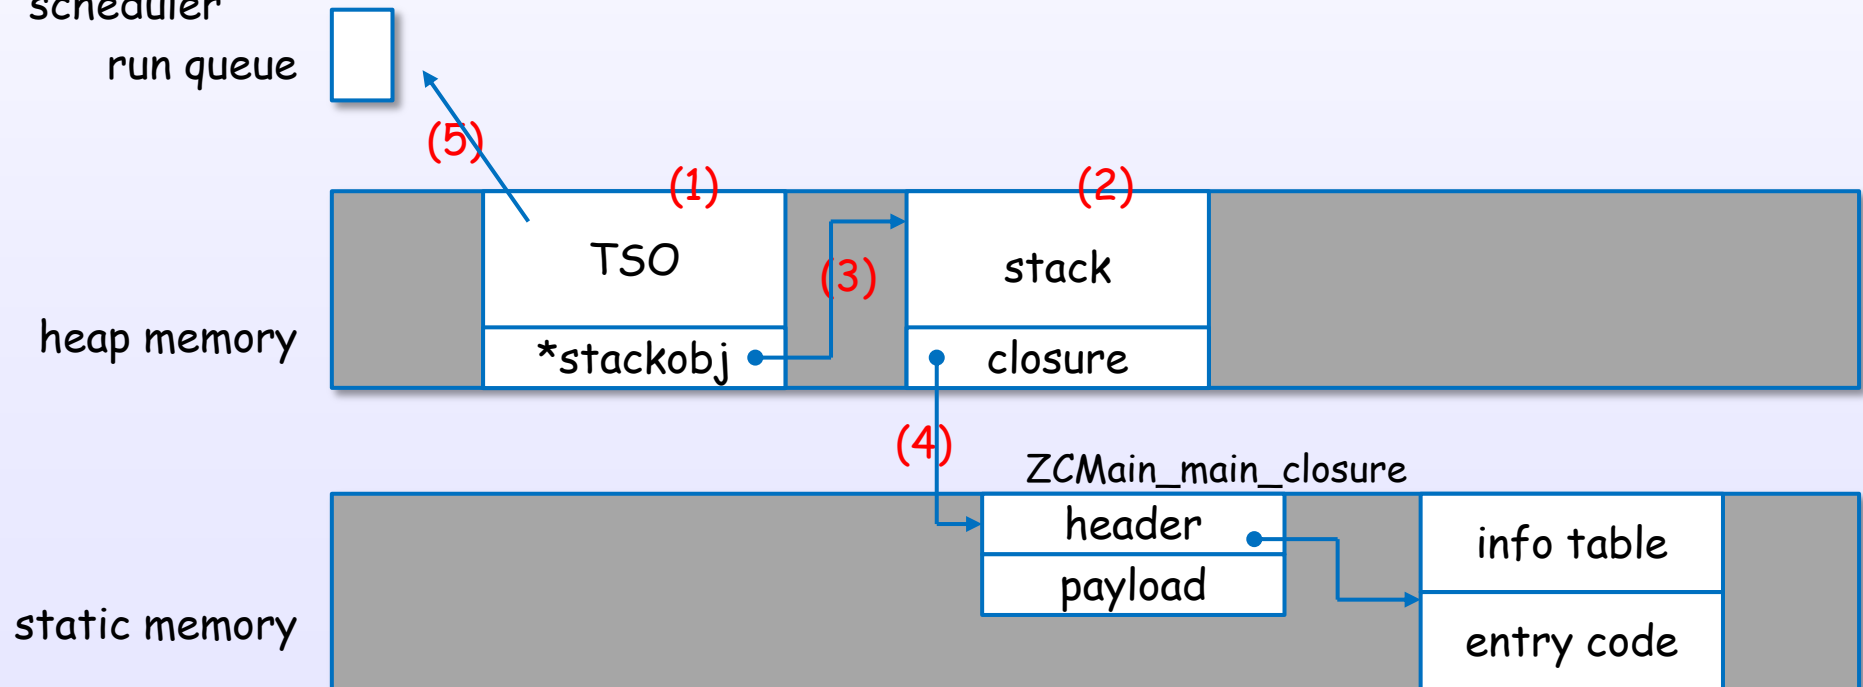
Create a main thread

Runtime
System

Runtime system bootstrap code [rts/RtsAPI.c]

```
rts_evalLazyIO
  createIOThread
    createThread ... (1), (2), (3)
    pushClosure ... (4)
  scheduleWaitThread
    appendToRunQueue ... (5)
```

scheduler
run queue



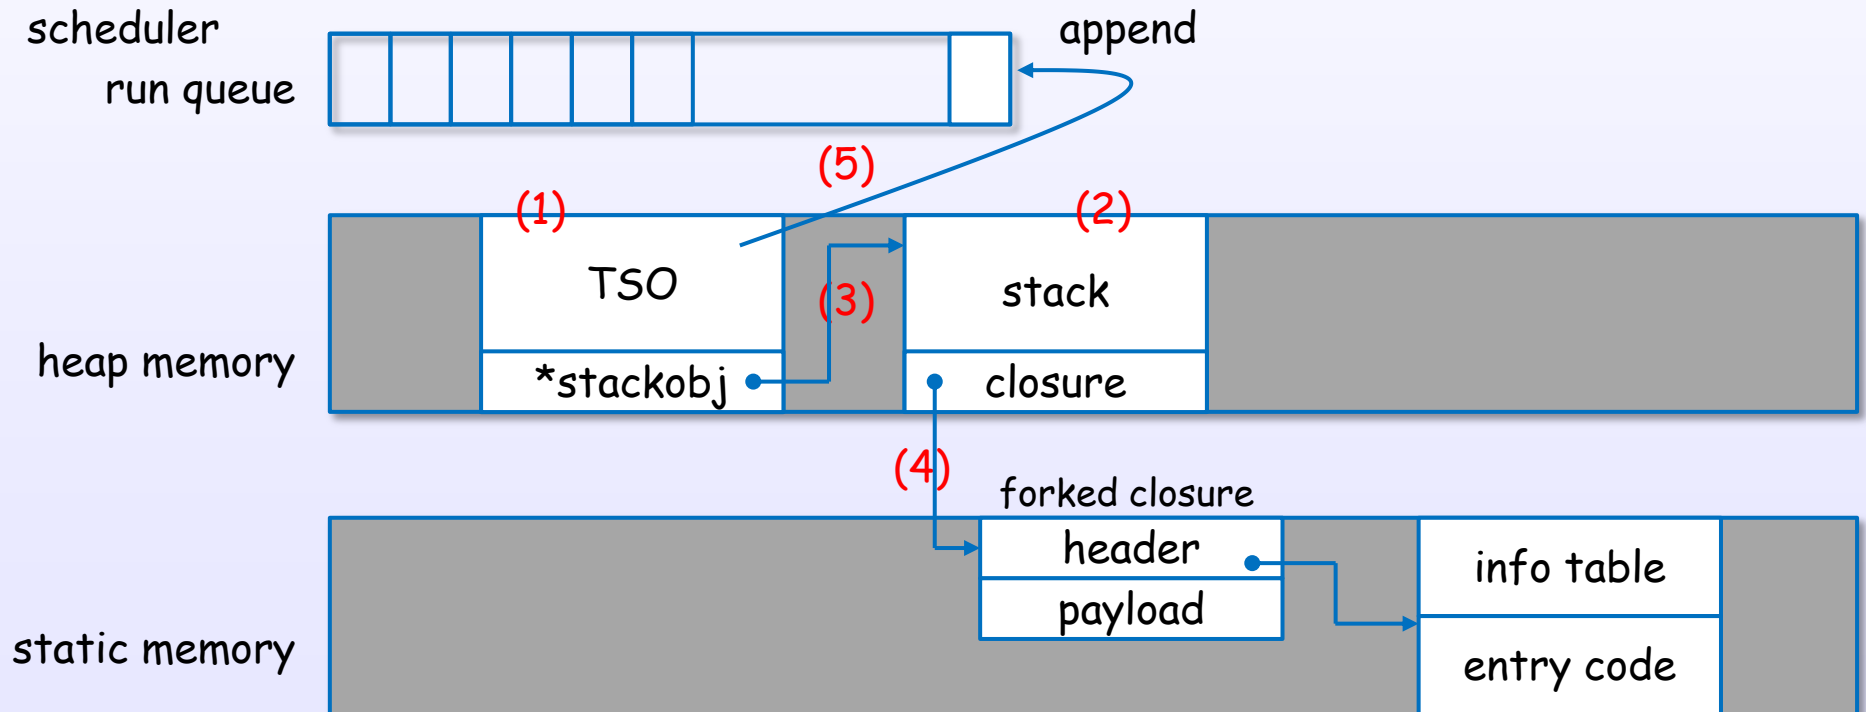
Create a sub thread by forkIO

Haskell Threads

```
forkIO
  stg_forkzh
    ccall createIOThread ... (1), (2), (3), (4)
    ccall scheduleThread ... (5)
```

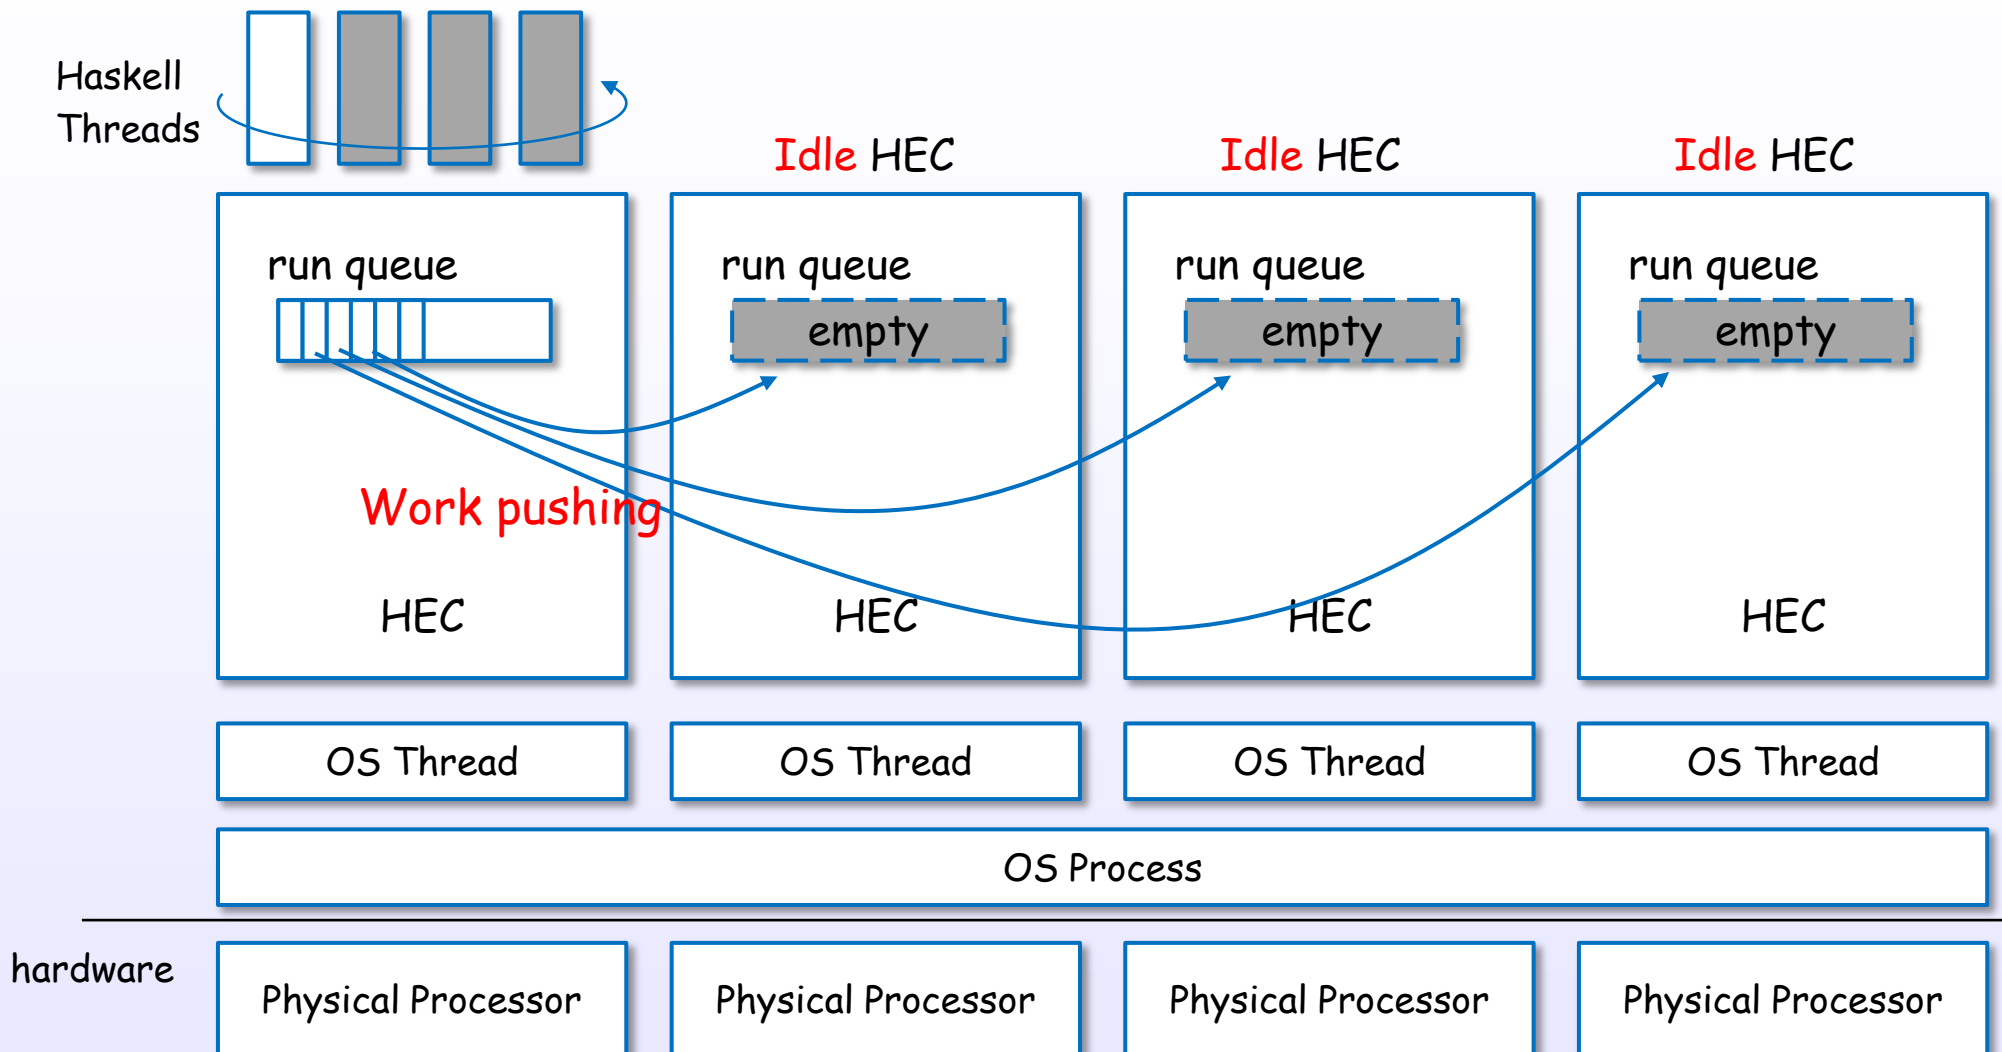
User code

Runtime System



Thread migration

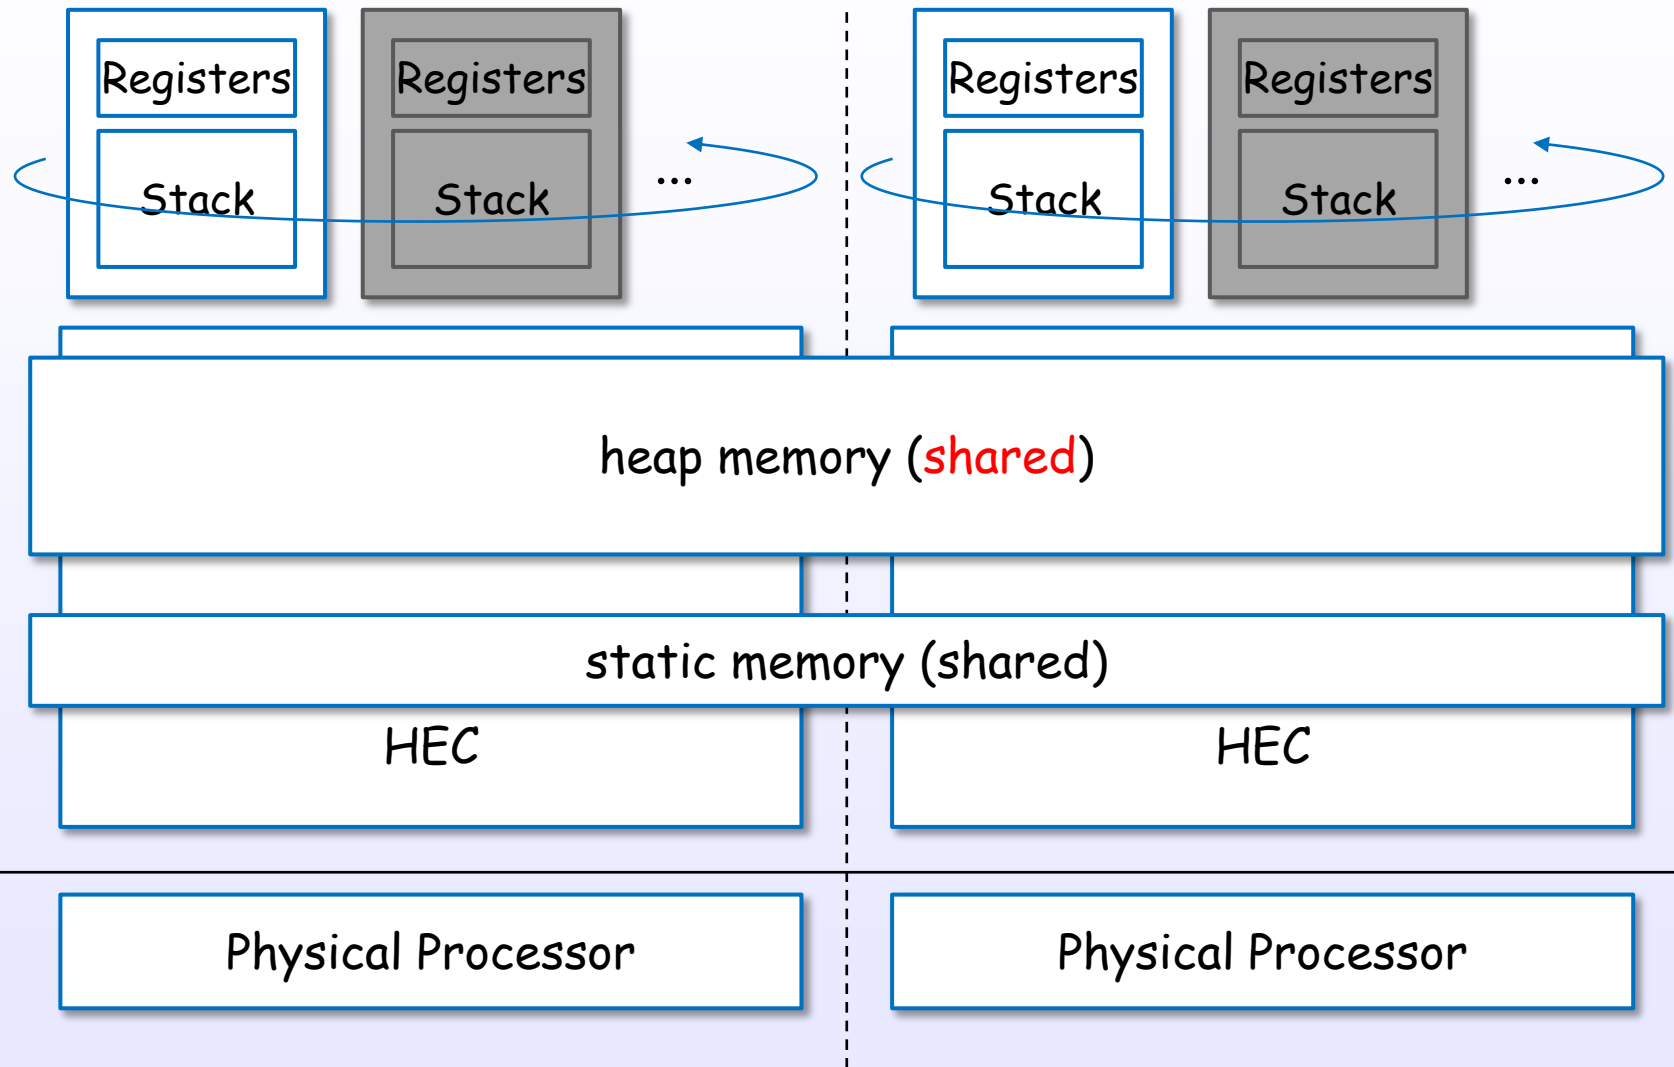
Threads are migrated to idle HECs



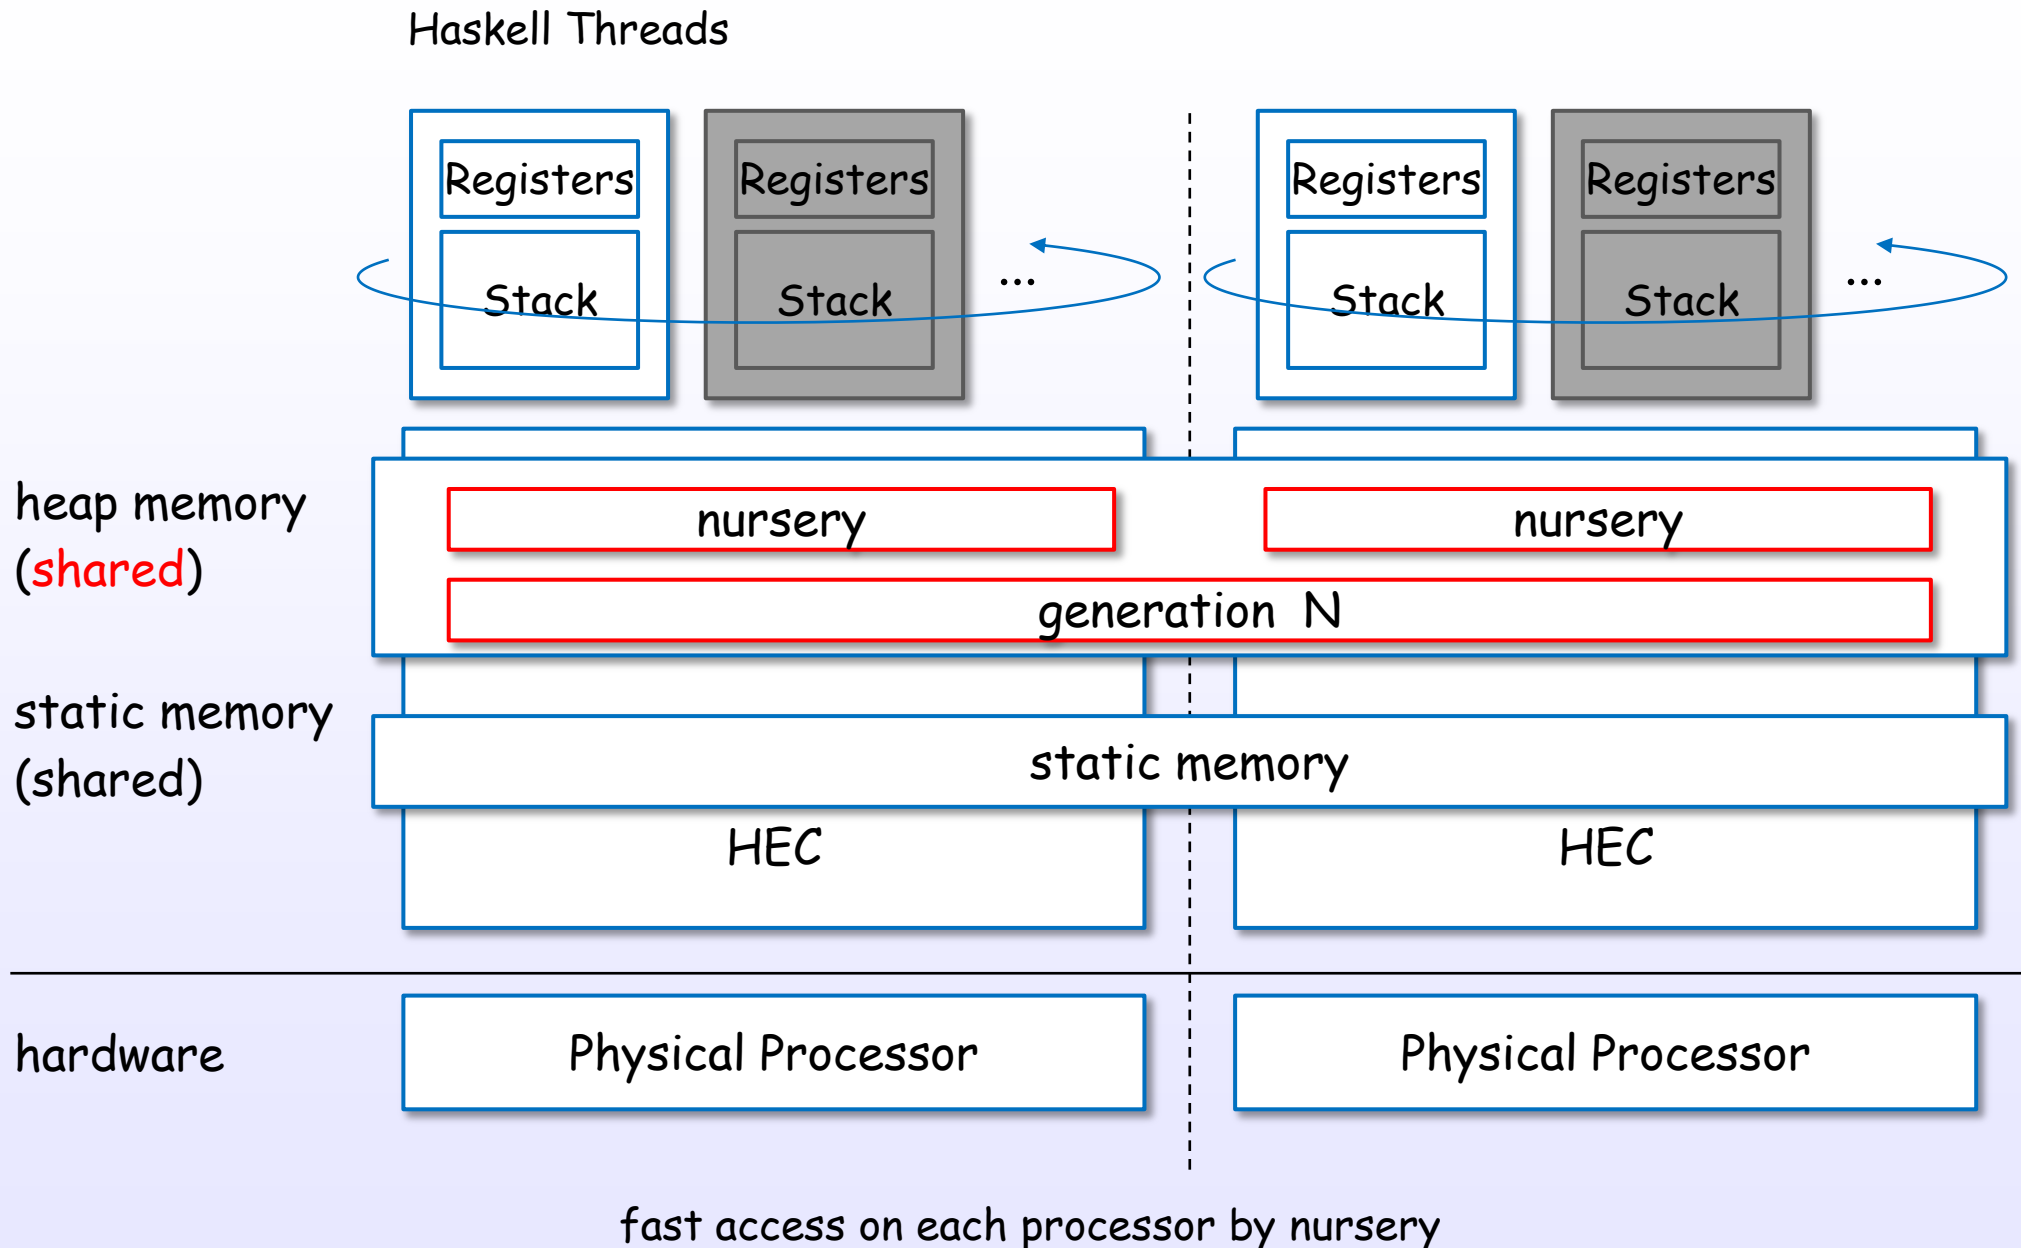
Heap and Threads

Threads and a shared heap

Haskell Threads

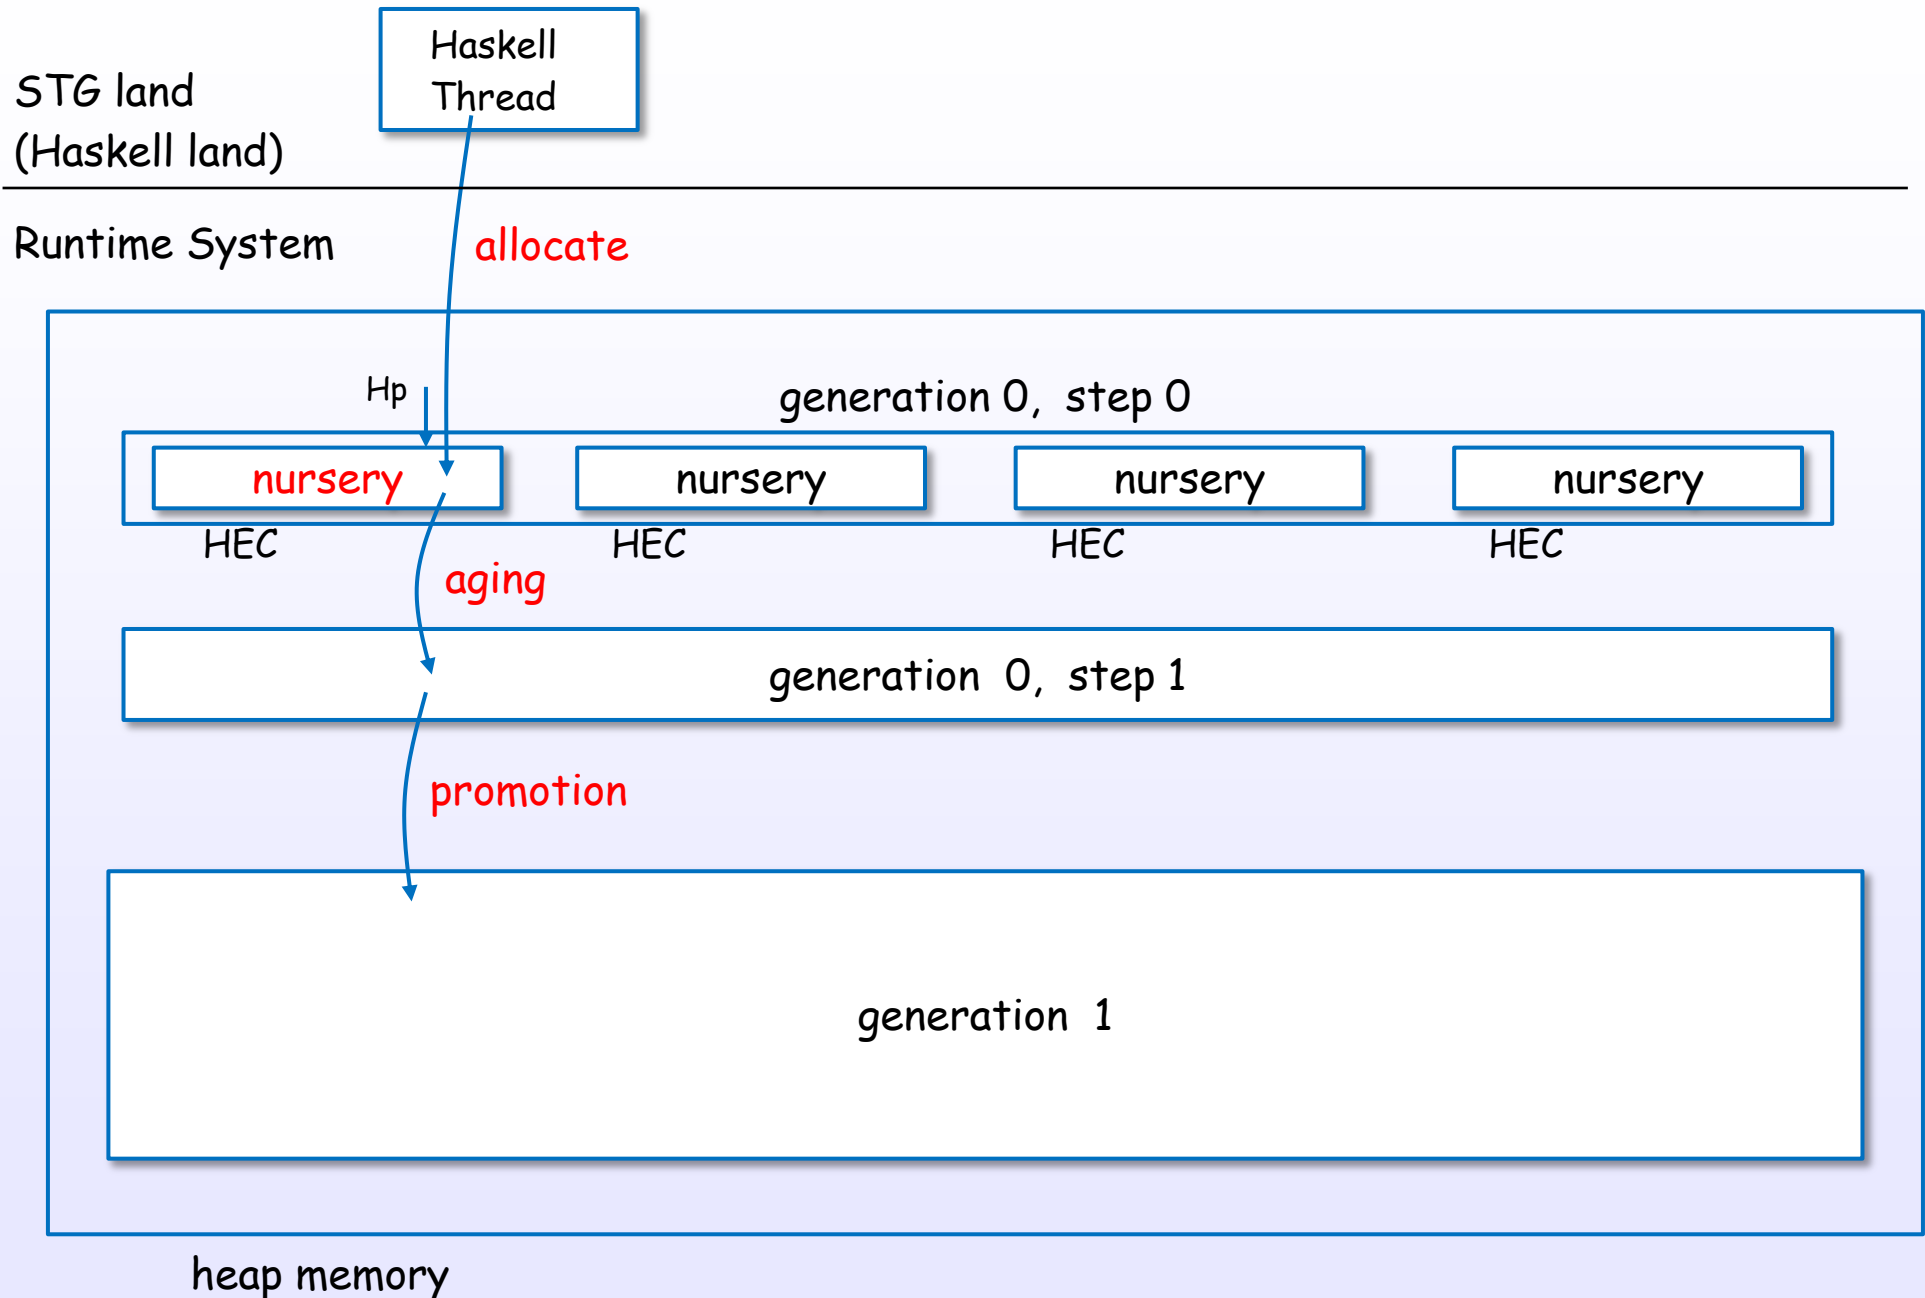


Local allocation area (nursery)



Threads and GC

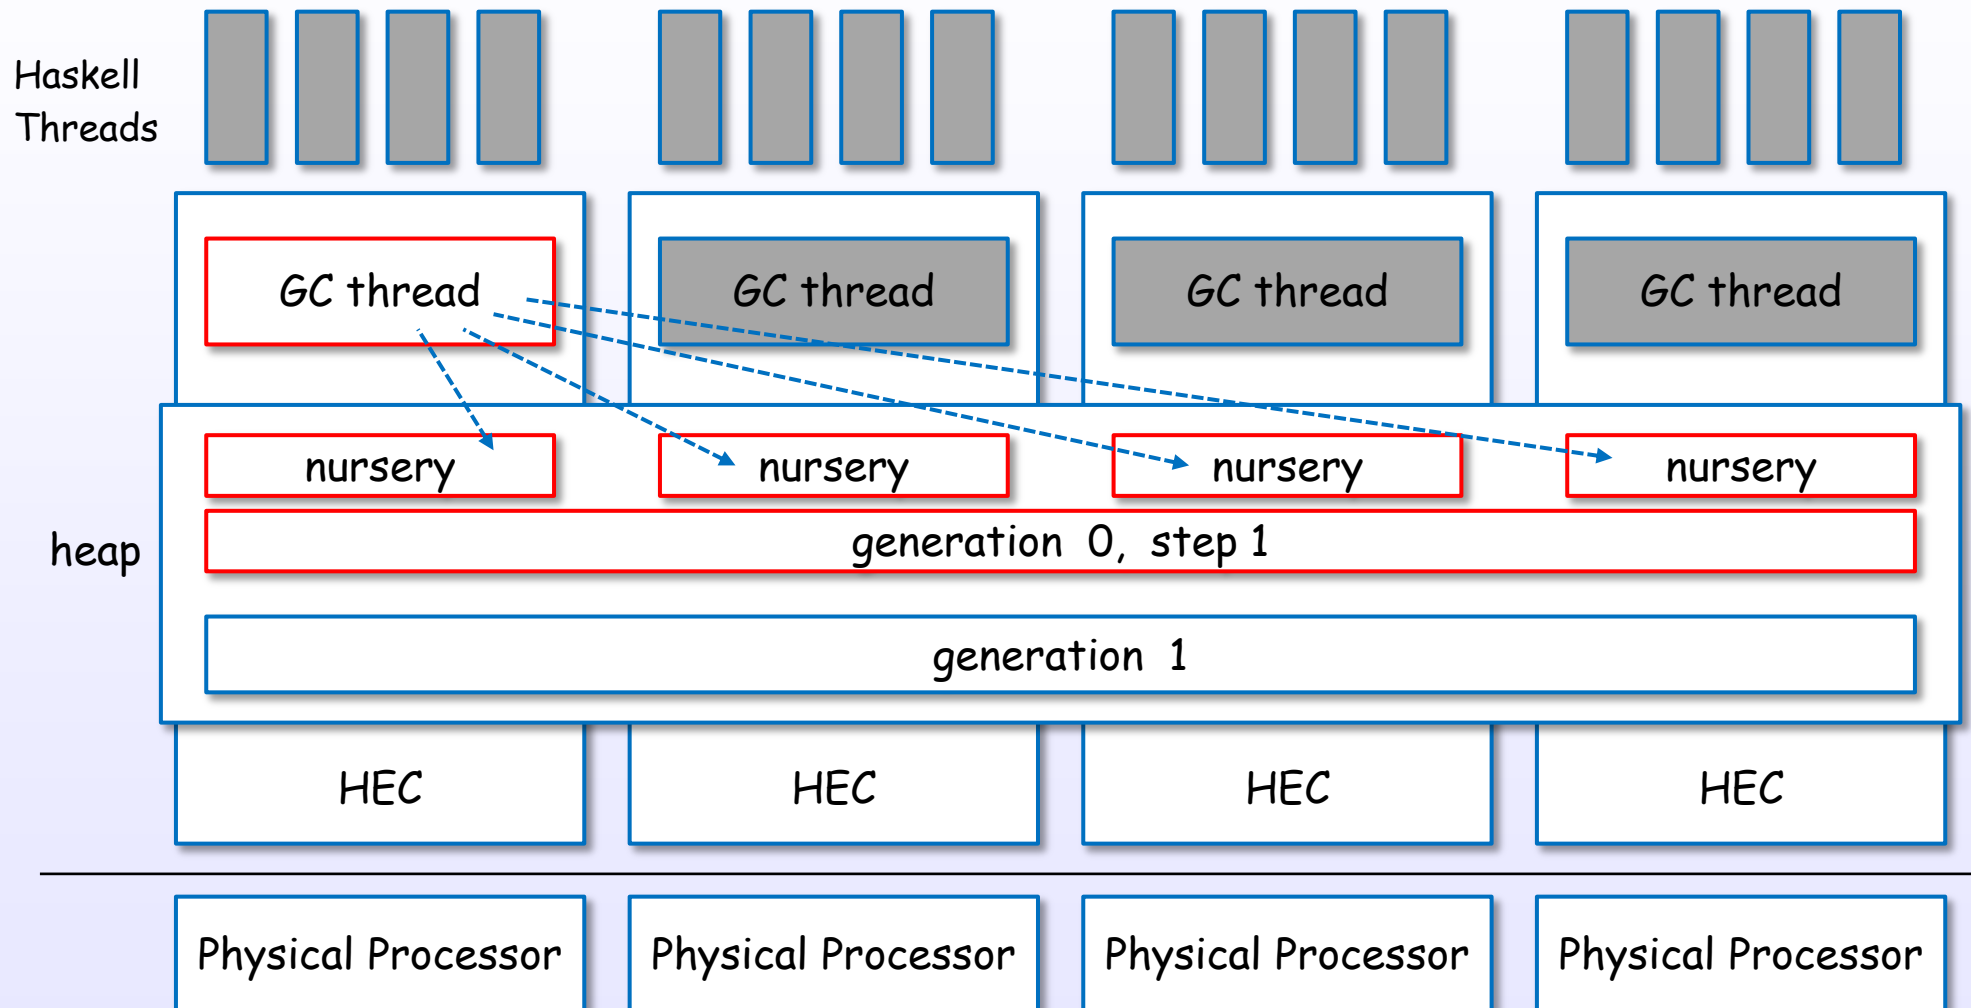
GC, nursery, generation, aging, promotion



Threads and minor GC

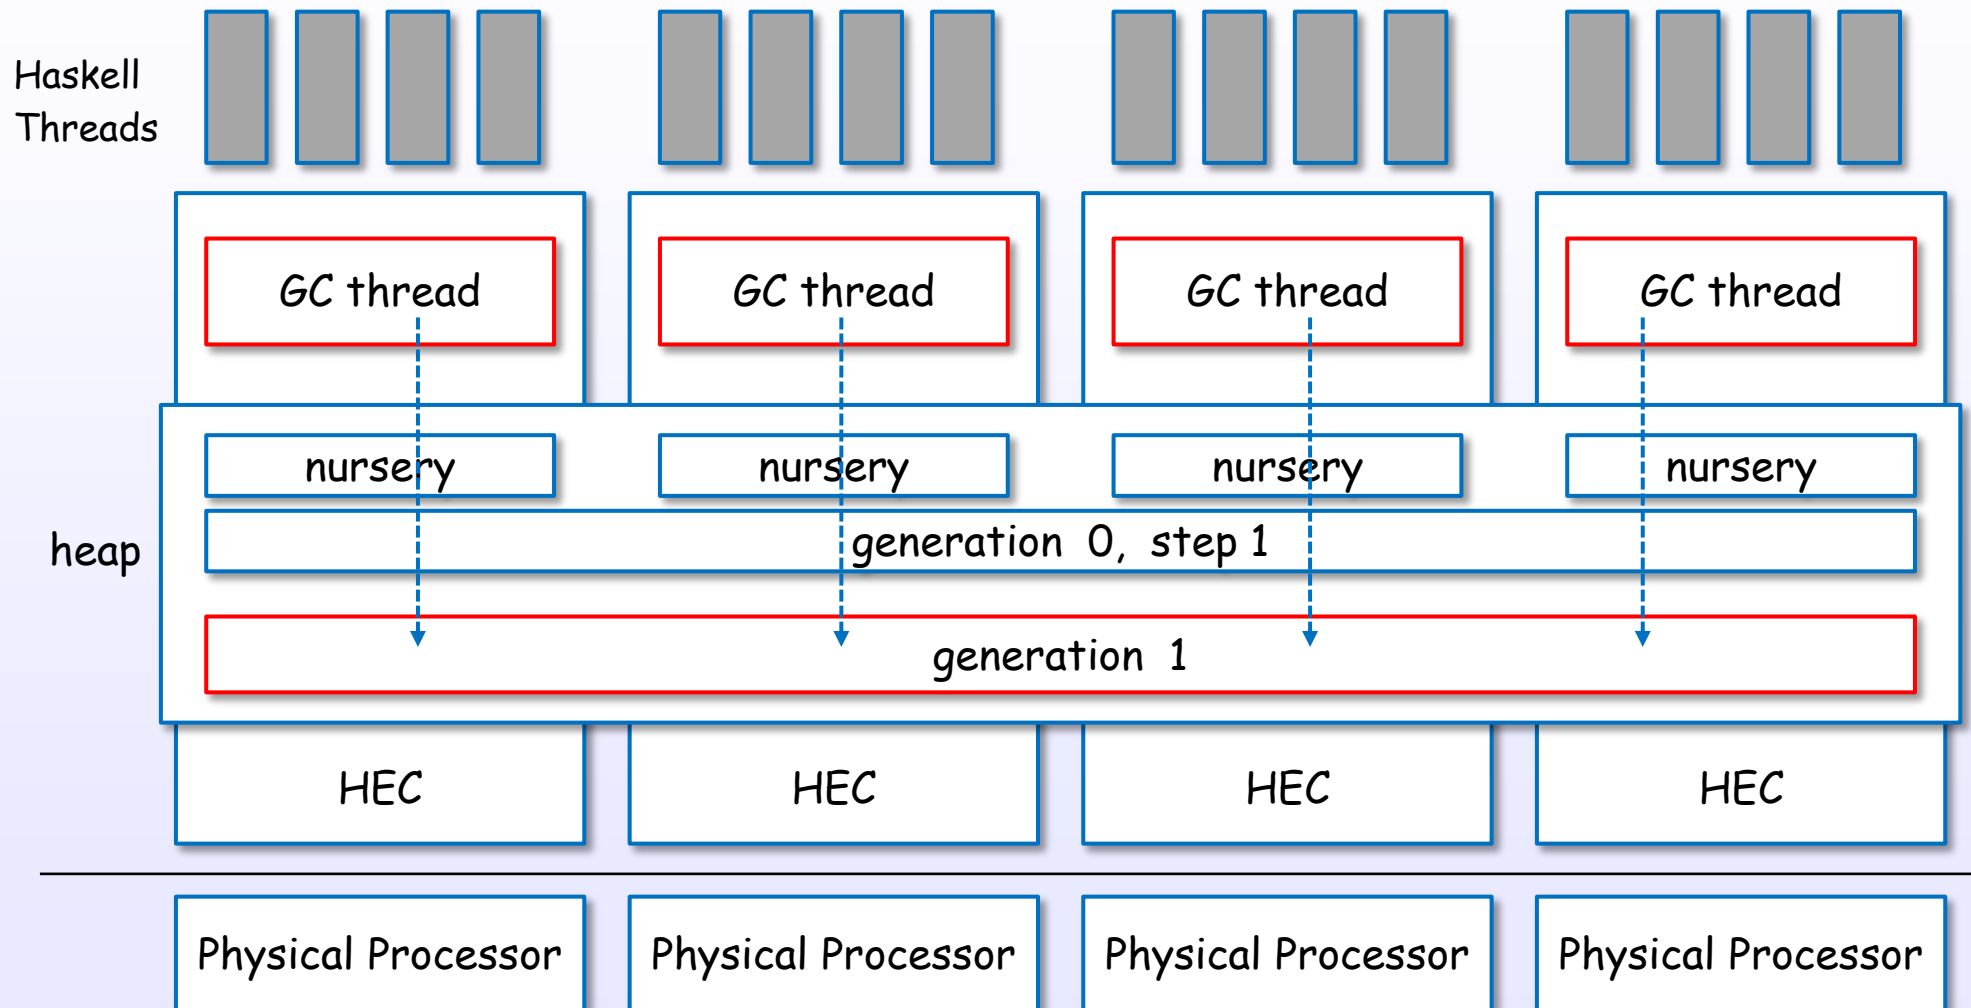
sequential GC for young generation (minor GC)

"stop-the-world" GC



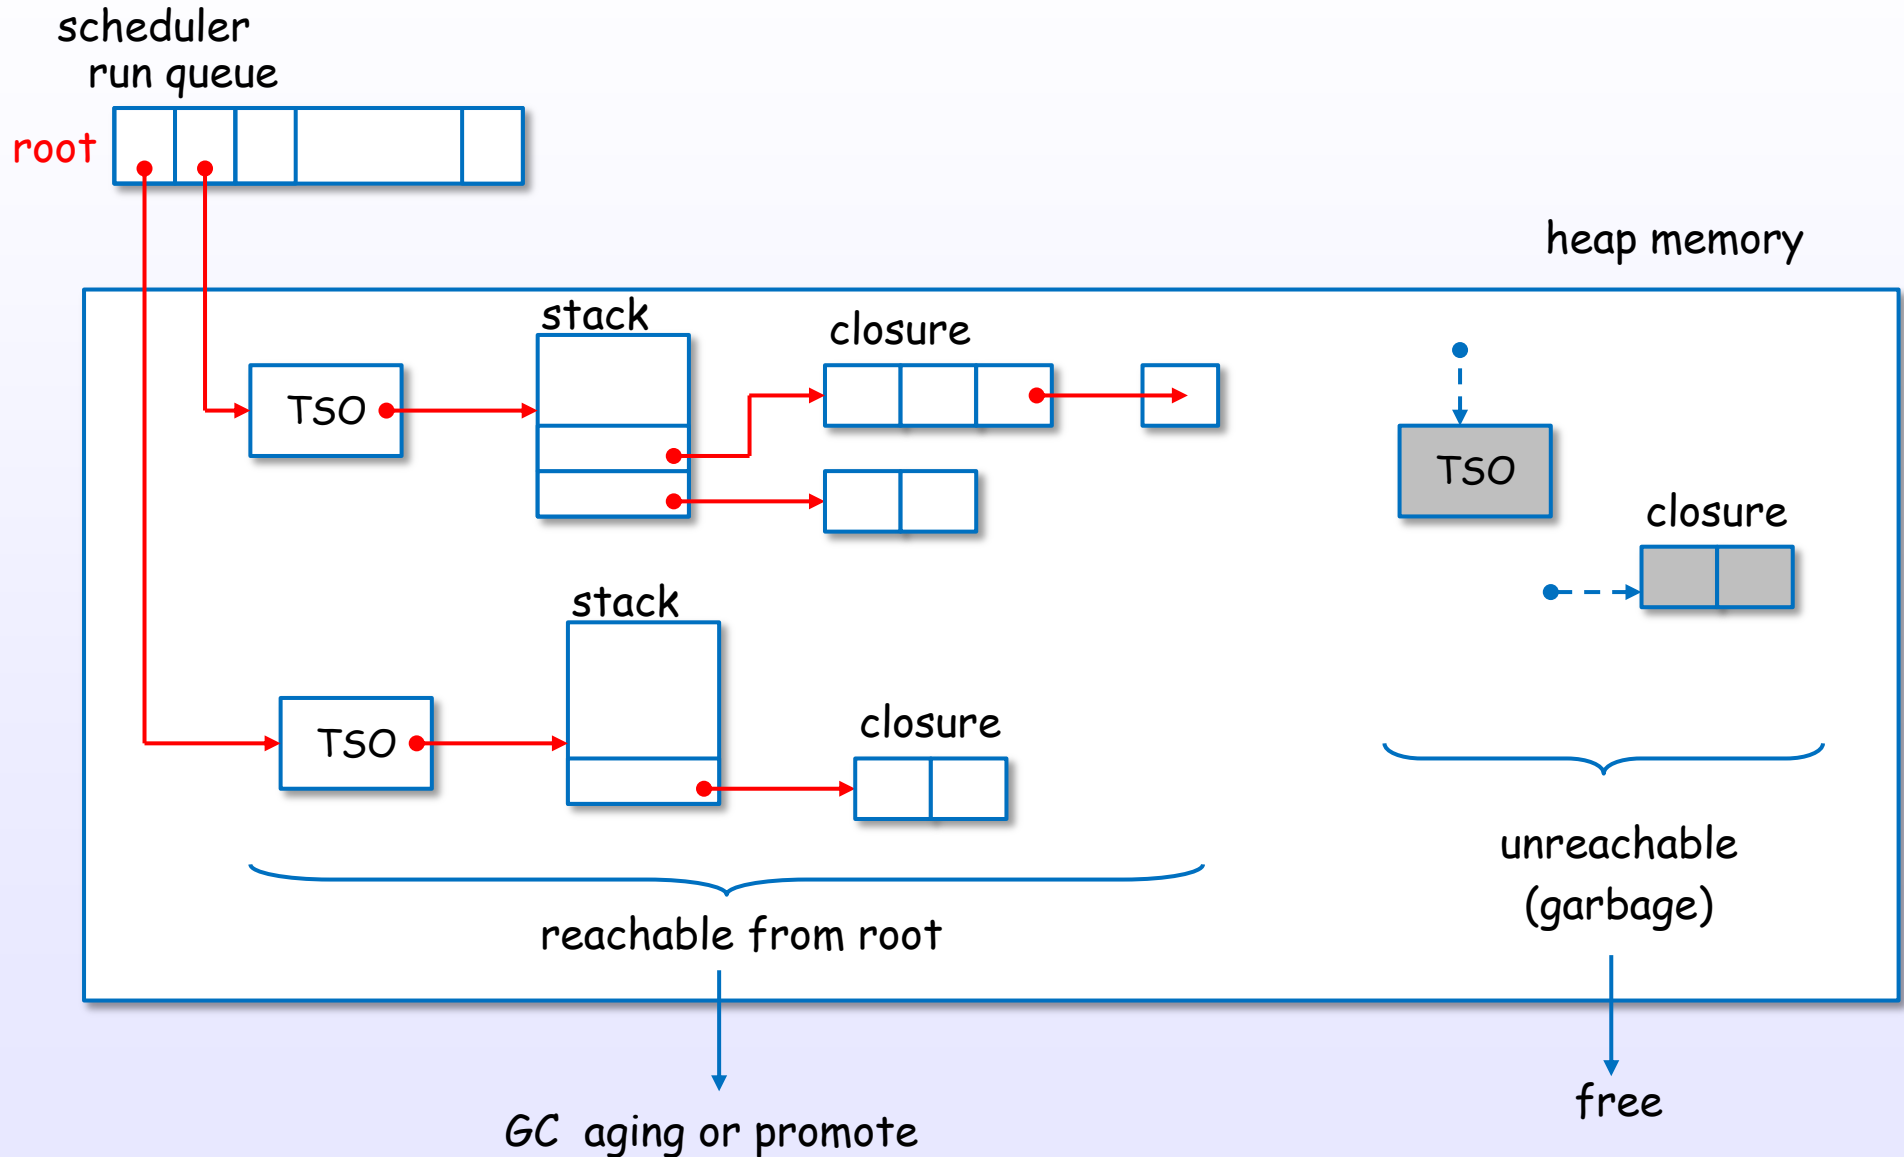
Threads and major GC

parallel GC for oldest generation (major GC)
"stop-the-world" GC



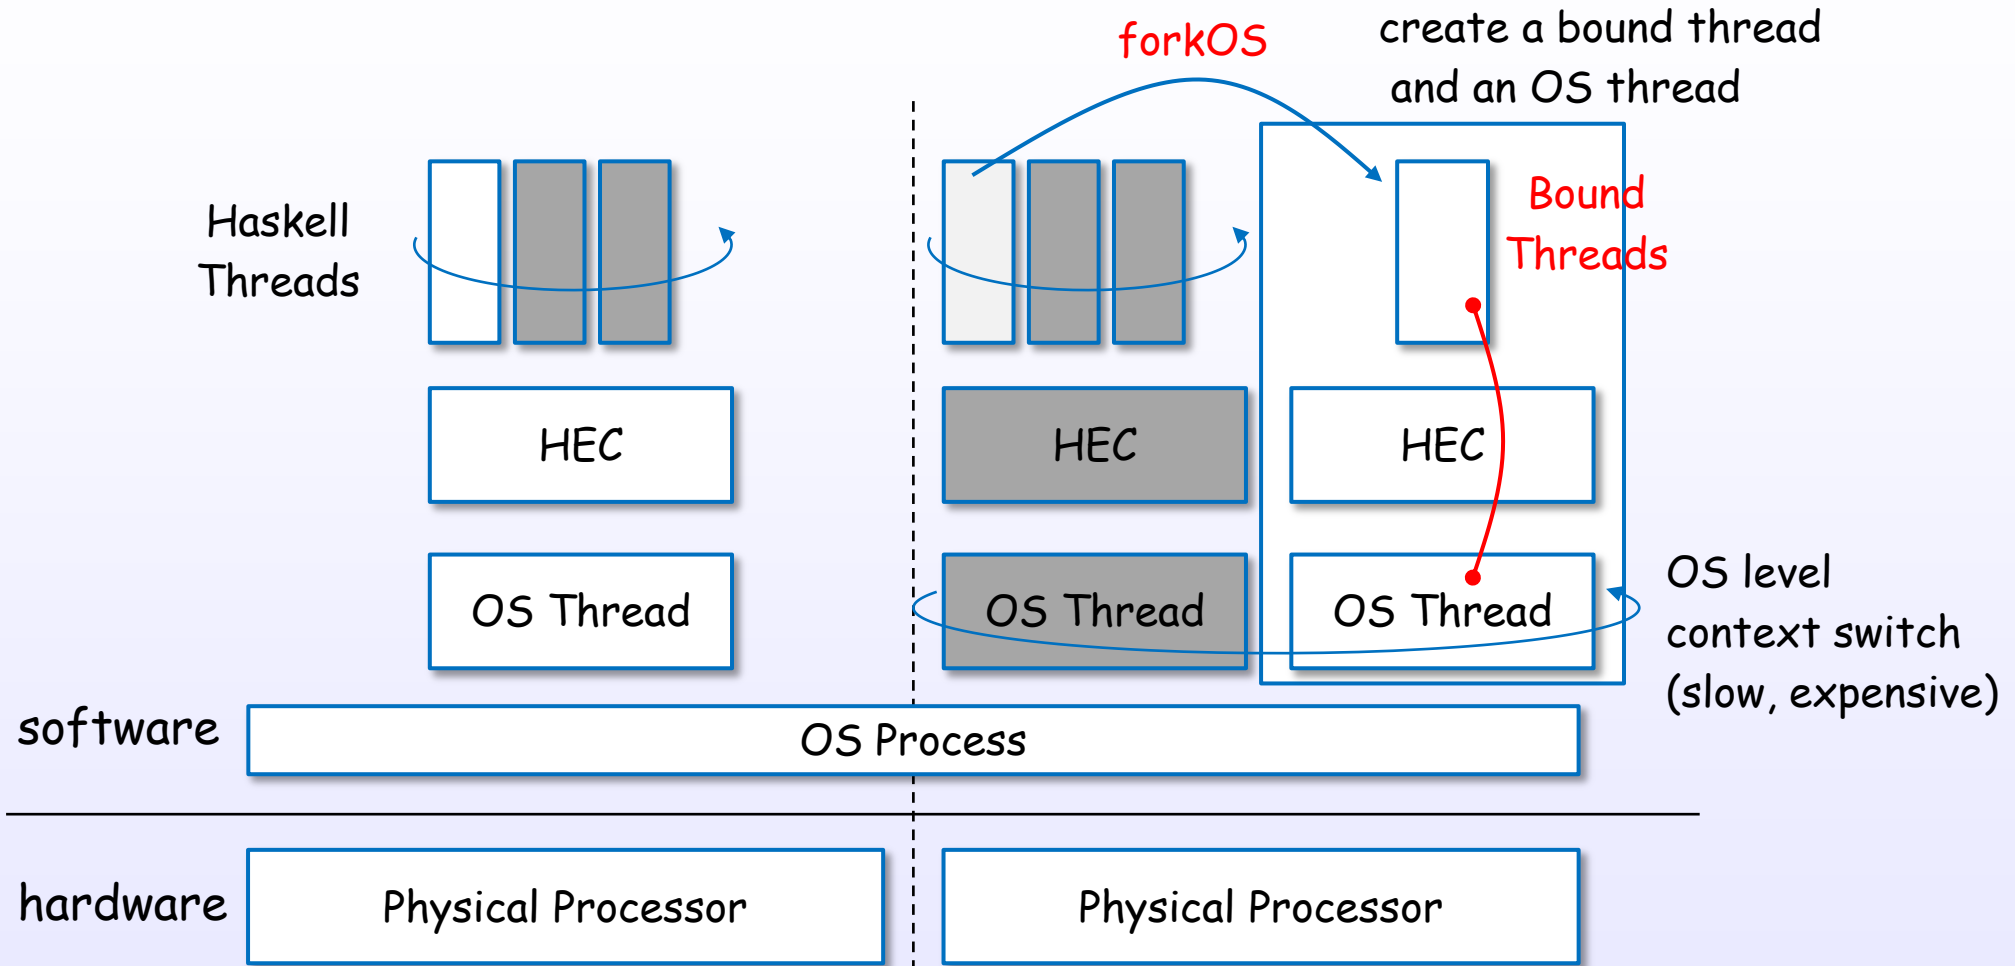
GC discover live objects from the root

Runtime System




Bound thread


Create a bound thread by forkOS




forkIO, forkOn, forkOS



create a haskell thread



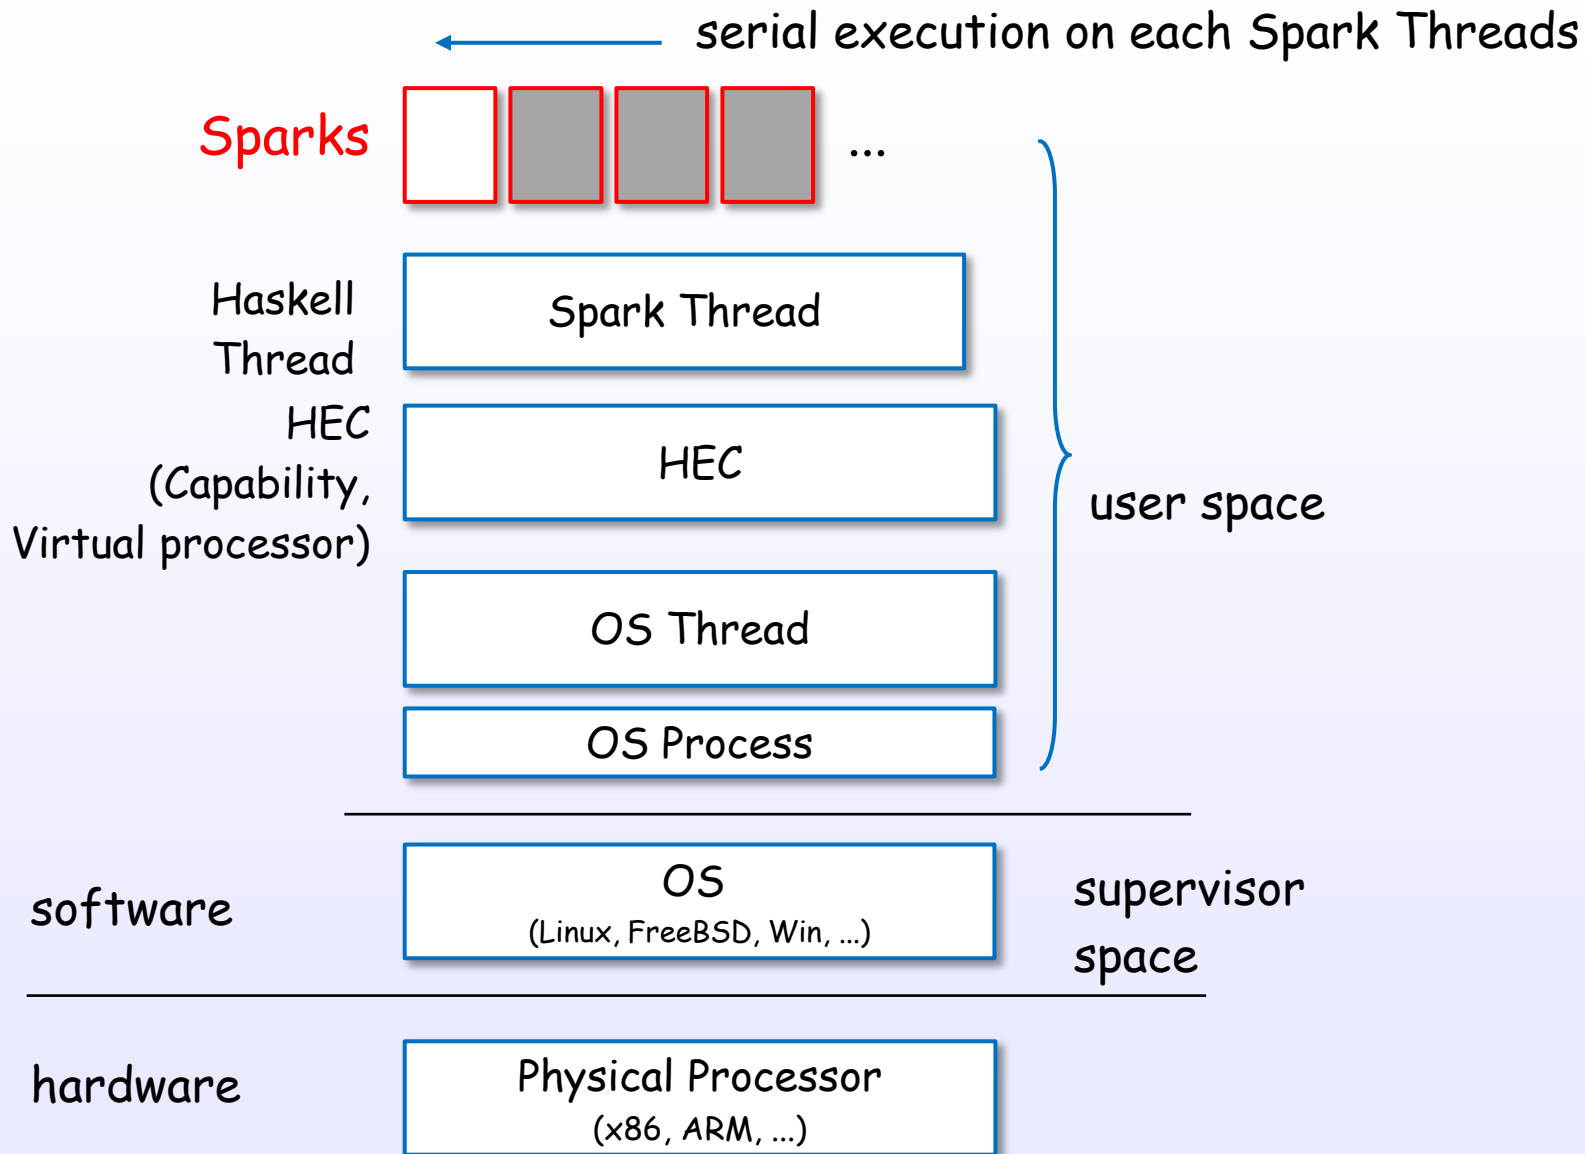
create a haskell thread
on specified HEC



create a haskell thread
and an OS thread

Spark

Spark layer

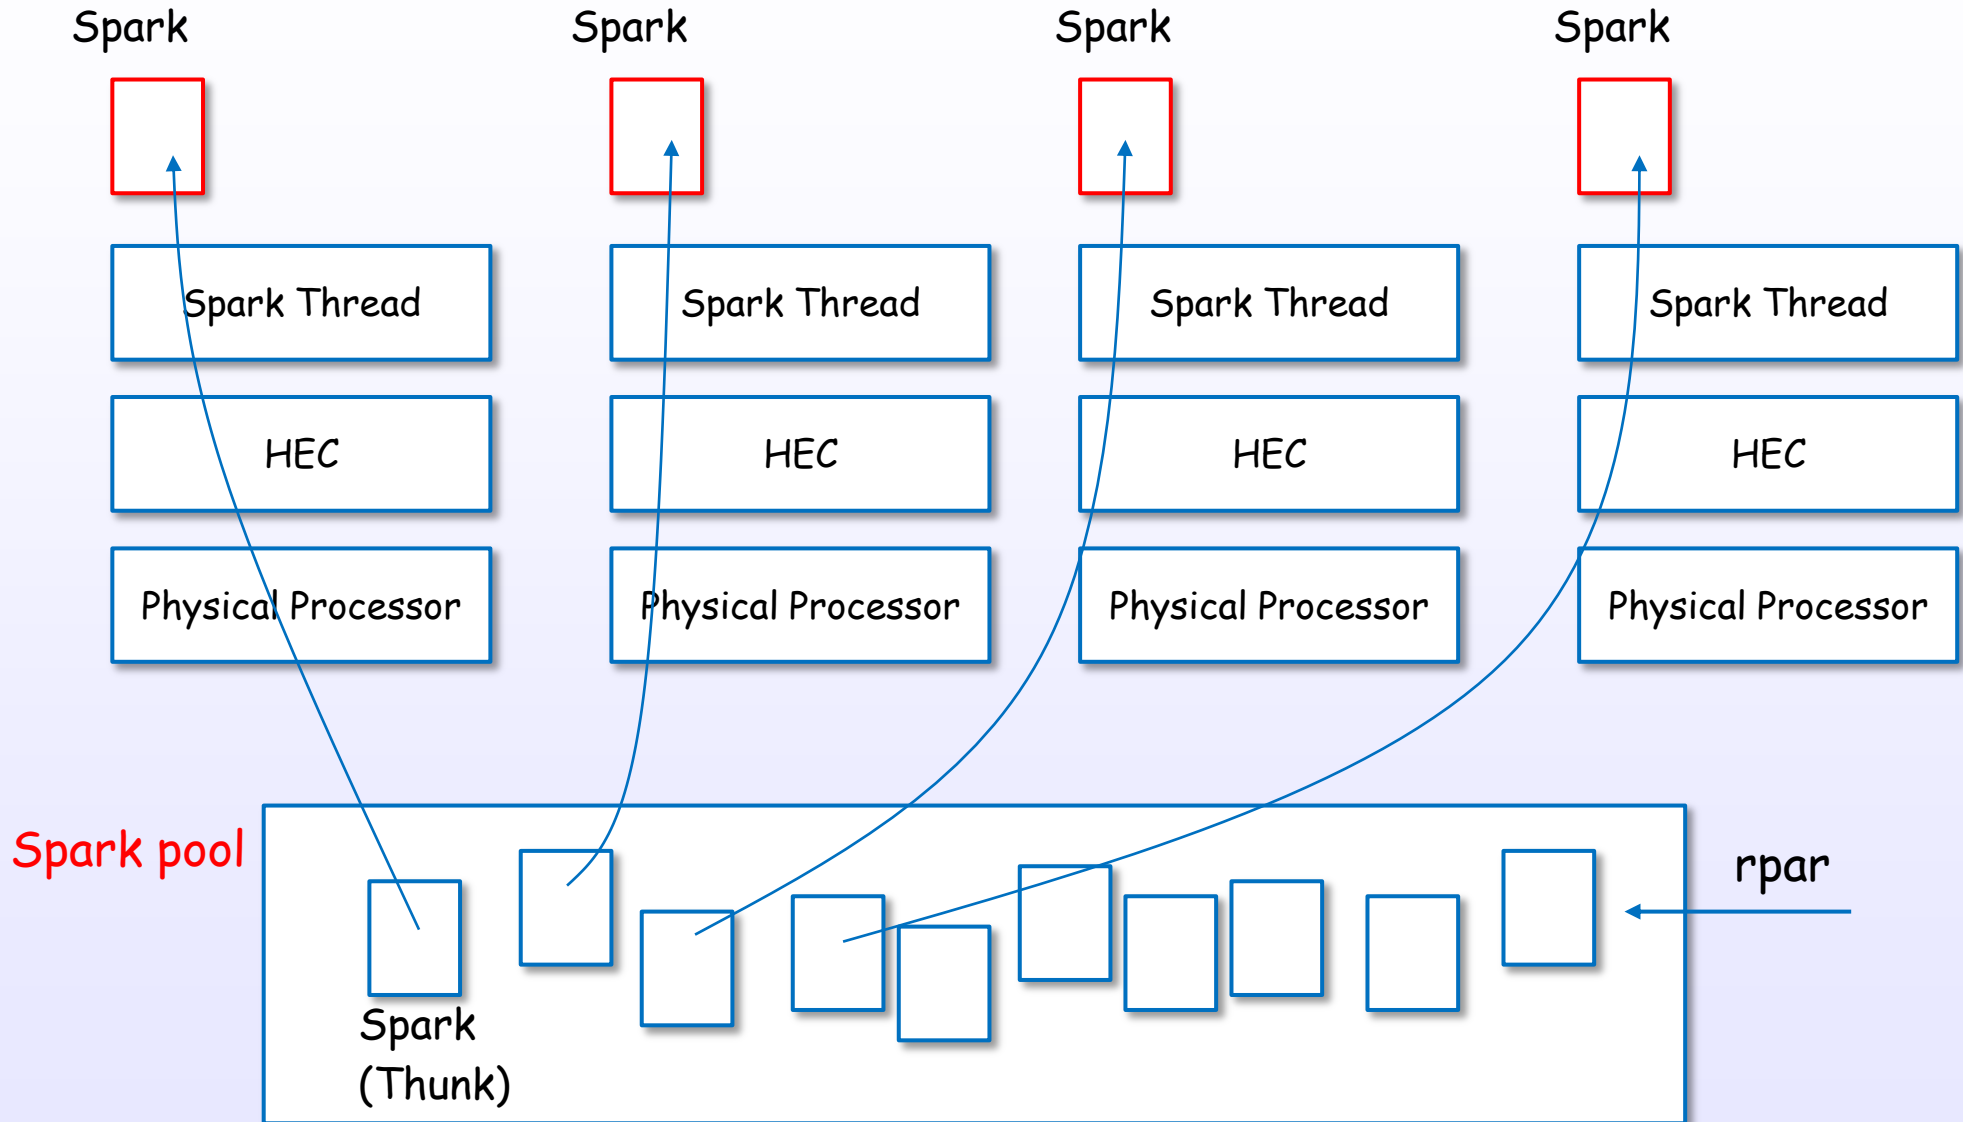


Spark Threads are generated on idle HECs.

References : [C17], [19], [S17], [S26], [S27], [S33], [S12]

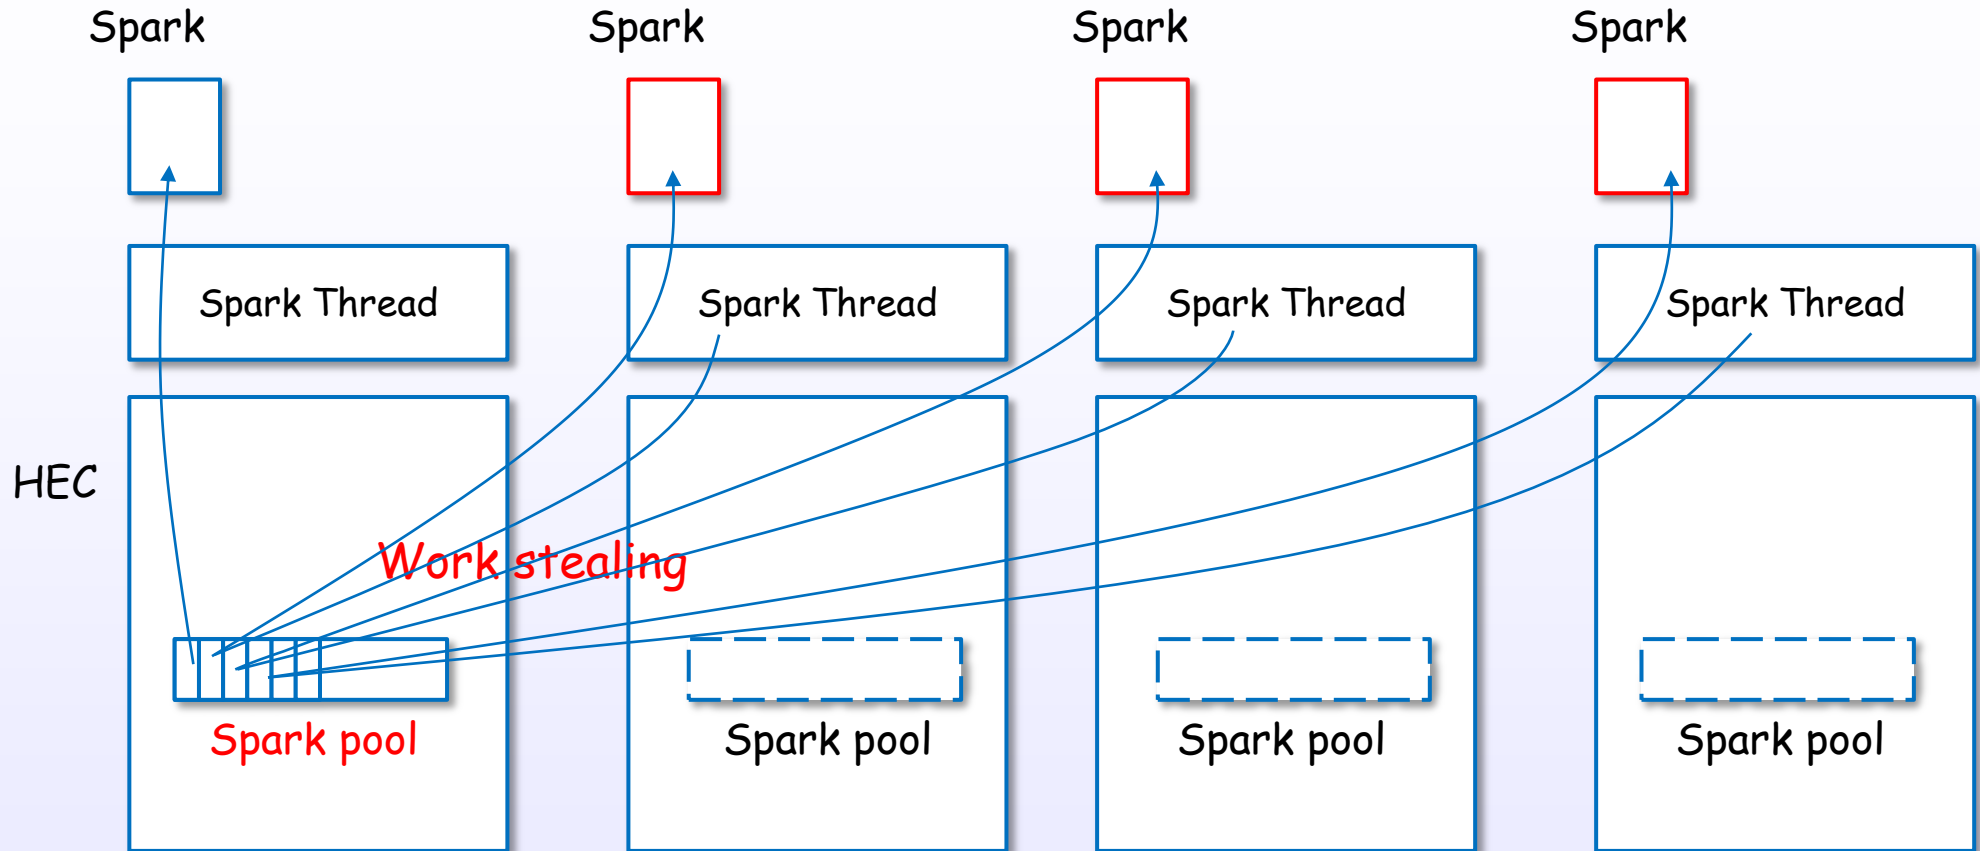
Sparks and Spark pool

logical view

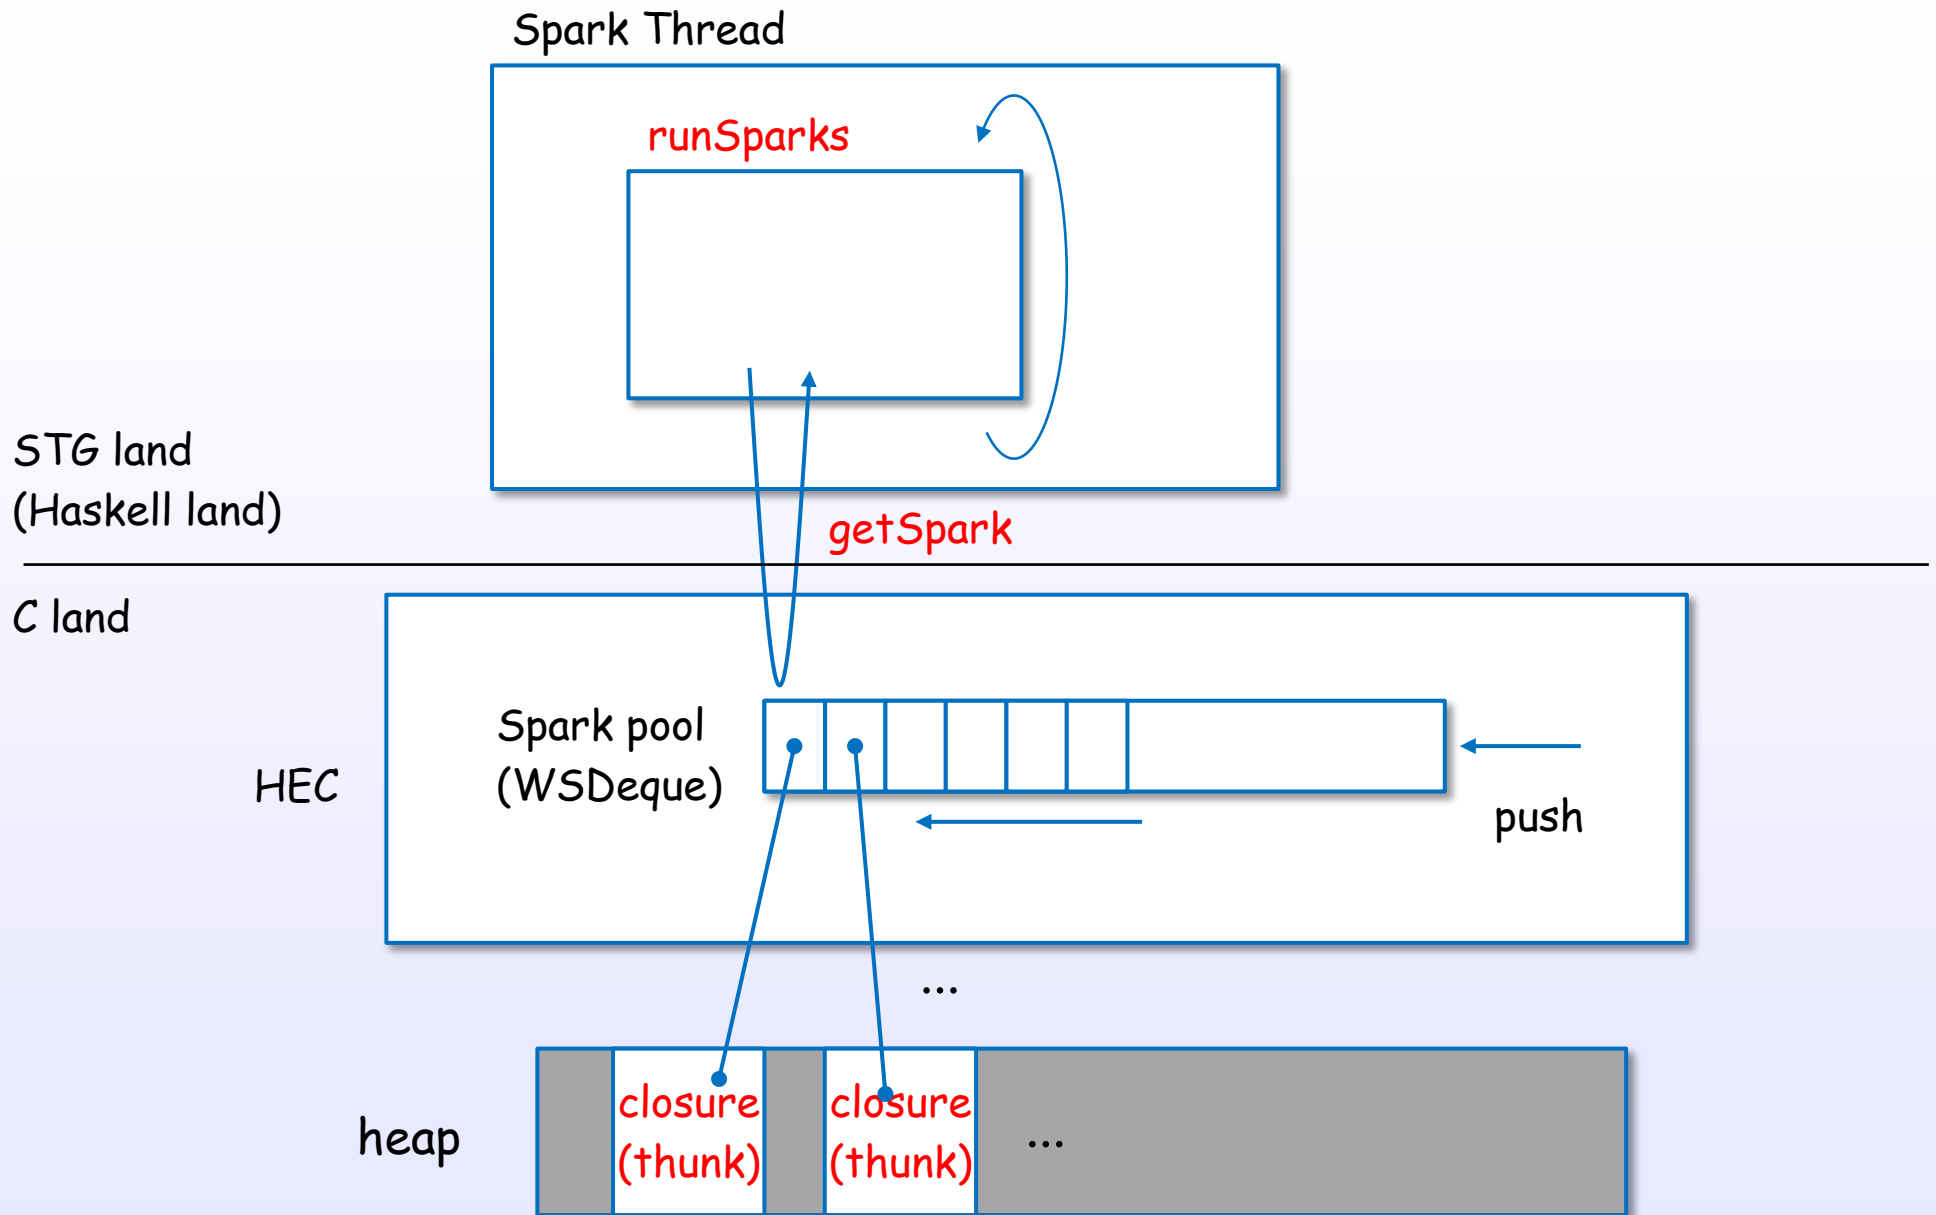


Spark pool and work stealing

physical view



Sparks and closures



(not TSO objects, but closures. therefore very lightweight)

MVar


MVar

Haskell Thread #0




putMVar

Haskell Thread #1



takeMVar



MVar

MVar and blocking

Haskell Thread



putMVar



BLOCKED
if full



MVar

Haskell Thread



takeMVar

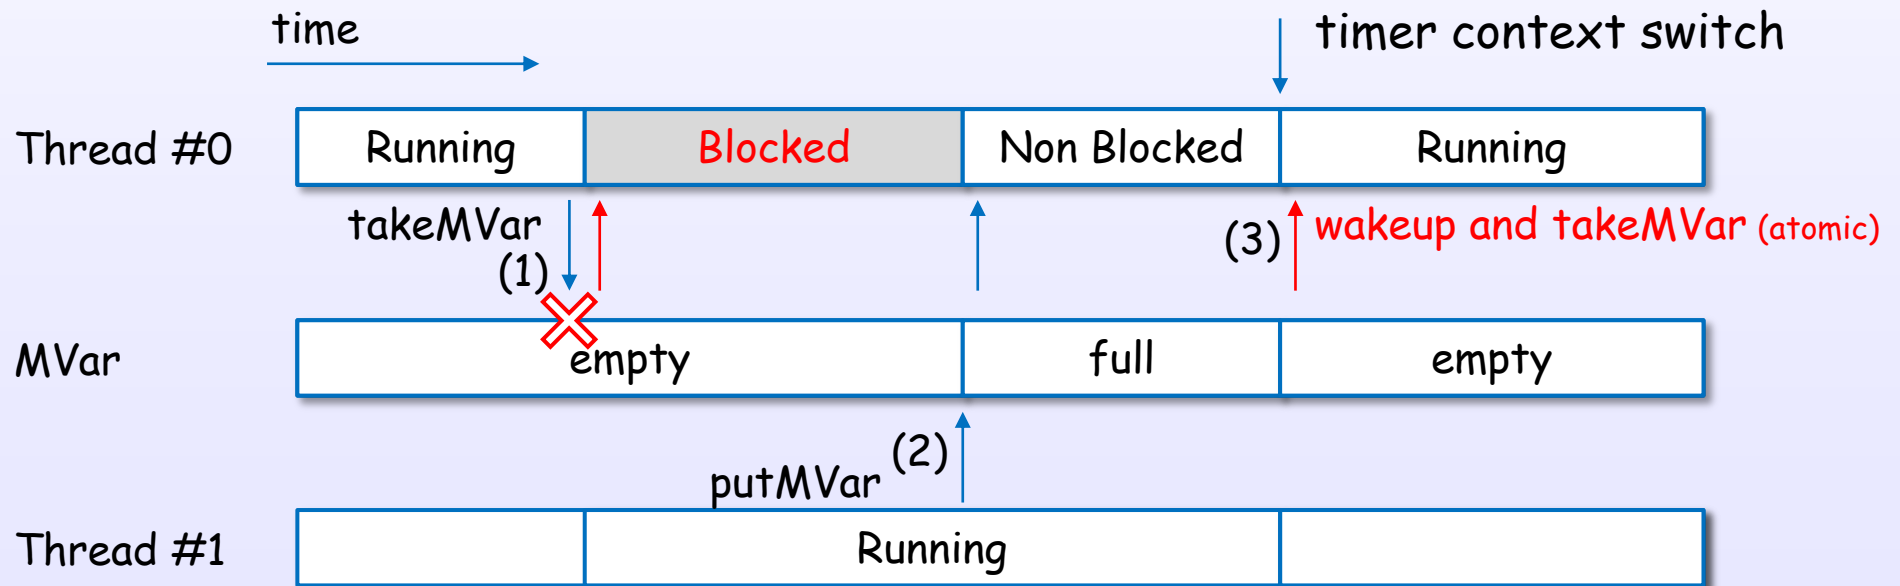



BLOCKED
if empty



MVar


MVar example




MVar object view

User view

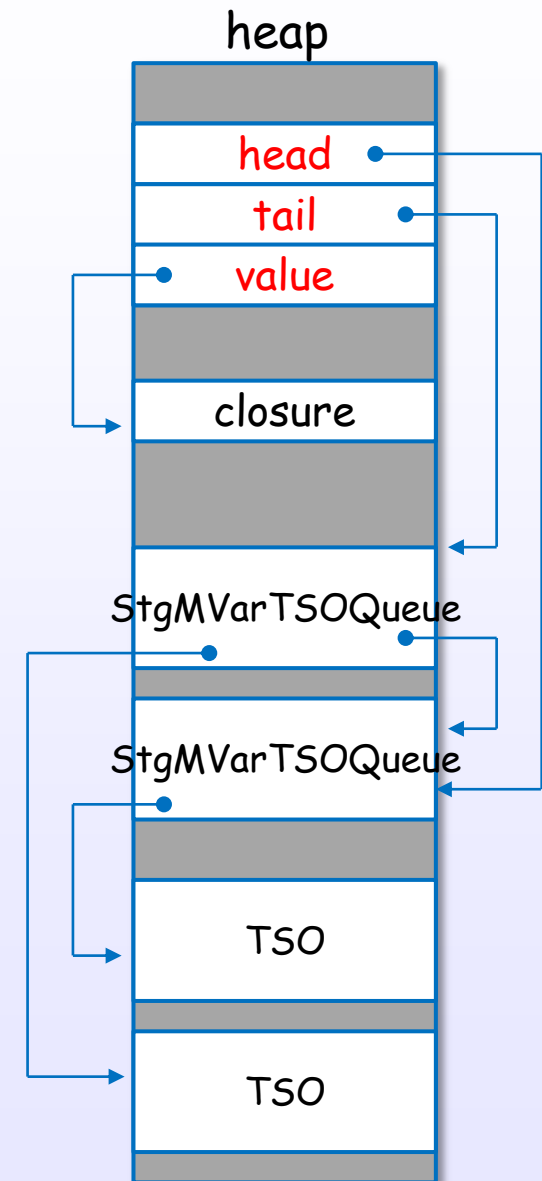
MVar



logical MVar object



physical MVar object



newEmptyMVar

Haskell Threads


```
newEmptyMVar  
newMVar#
```

(1) **call** the Runtime primitive

Runtime System

```
stg_newMVarzh  
  ALLOC_PRIM_  
  SET_HDR  
  StgMVar_head  
  StgMVar_tail  
  StgMVar_value
```

(2) **create a MVar** object



(3) **link** each fields

References : [16], [18], [19], [S31], [S12]

takeMVar (empty case)


Haskell Threads

```
takeMVar  
takeMVar#
```

Runtime System

```
stg_takeMVarzh  
  create StgMVarTSOQueue ... (1)  
  append      ... (2)  
  StgReturn   ... (3)
```

(3) return to the scheduler



MVar object

(1) create

StgMVarTSOQueue

(2) append

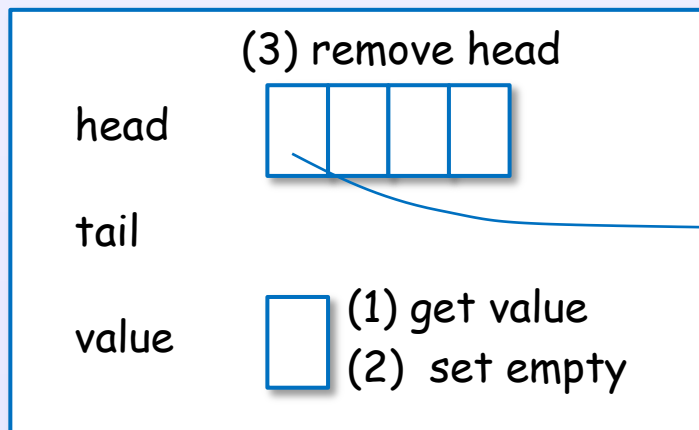
takeMVar (full case)

Haskell Threads

```
takeMVar  
takeMVar#
```

Runtime System

```
stg_takeMVarzh  
(1) get value  
(2) set empty  
(3) remove head  
(4) tryWakeupThread
```



MVar object

scheduler

run queue

fairness round robin



append

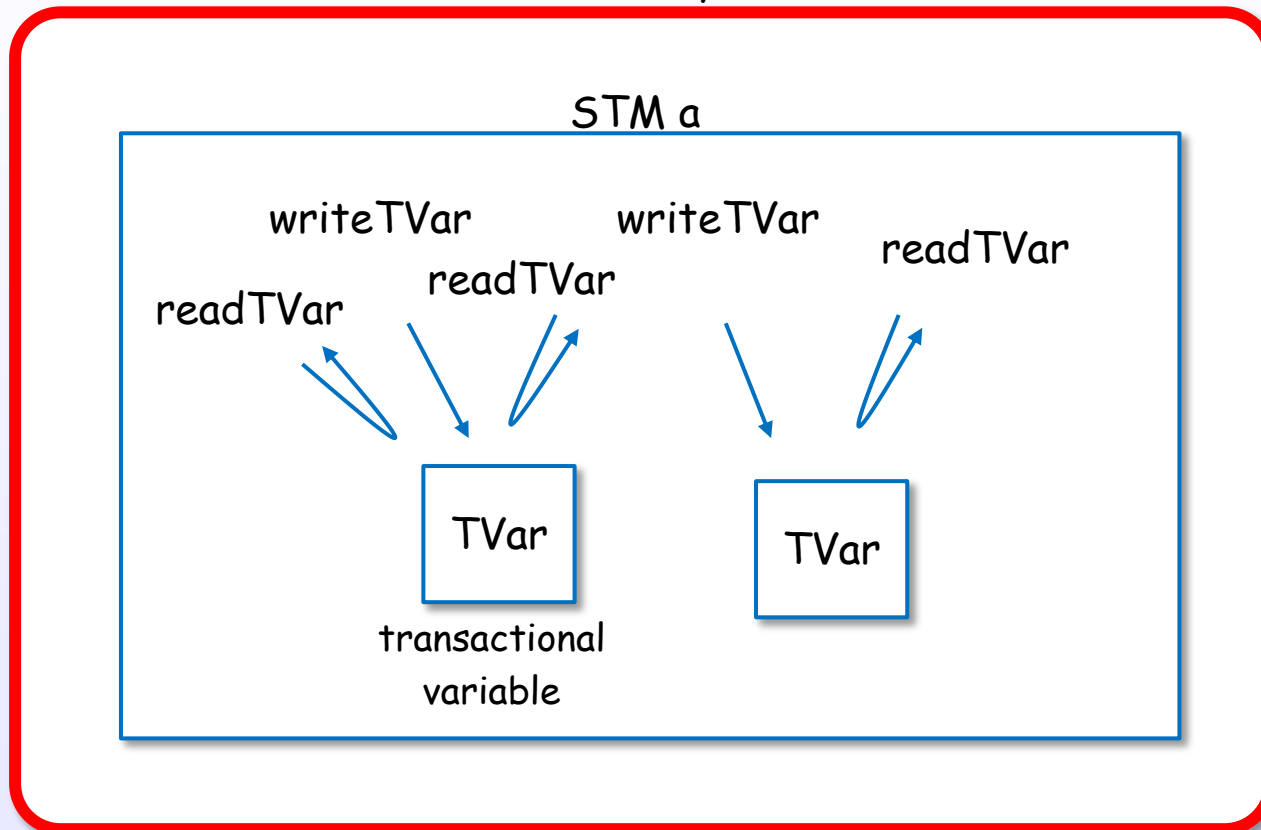
(4) wakeup the thread
blocked by putMVar

Software transactional memory

Create a atomic block by atomically

atomically :: STM a -> IO a

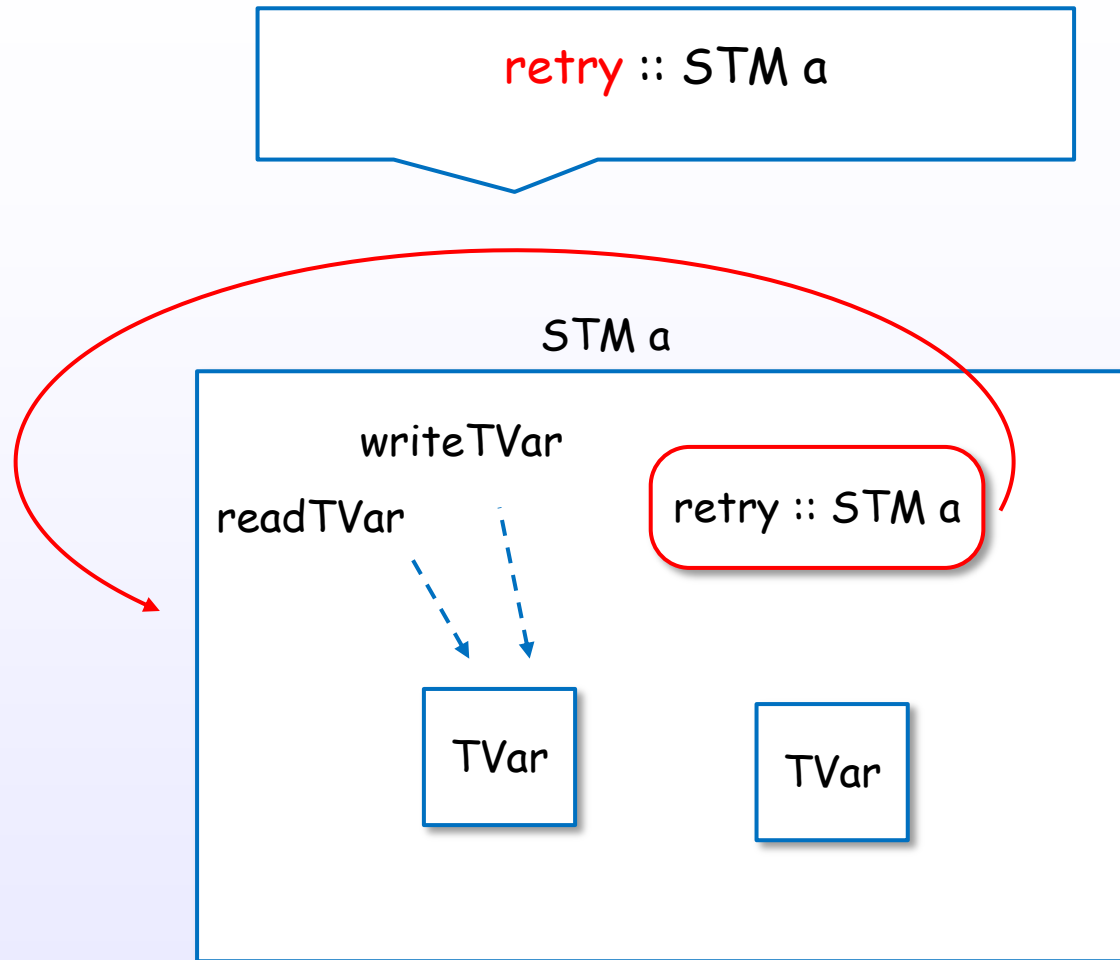
atomically



Create and evaluate a **composable** "atomic block"

Atomic block = All or Nothing

Rollback and blocking control by retry

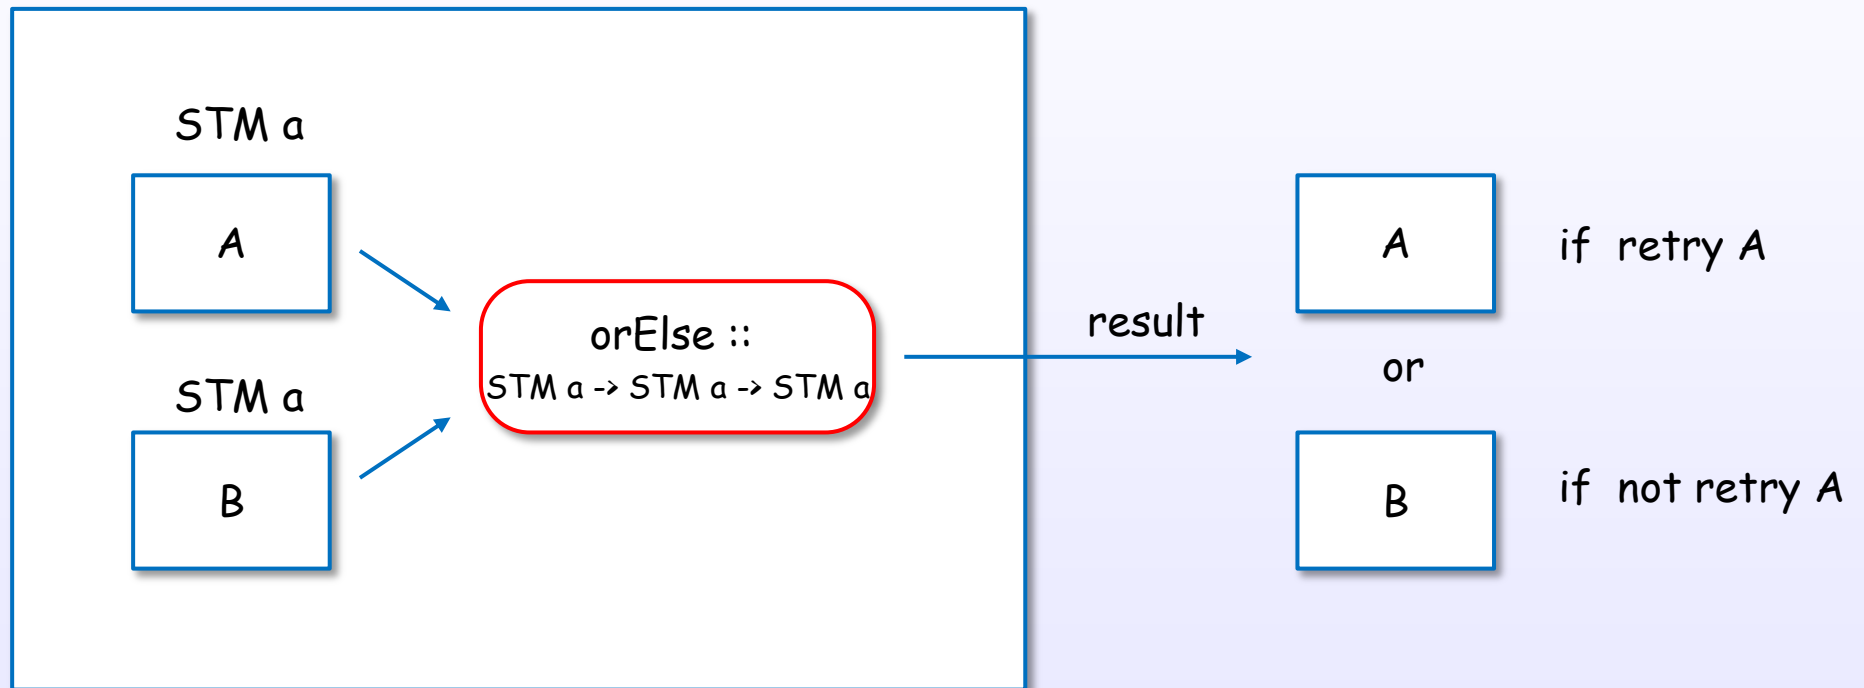


Discard, blocking and try again

Compose OR case by orElse

orElse :: STM a -> STM a -> STM a

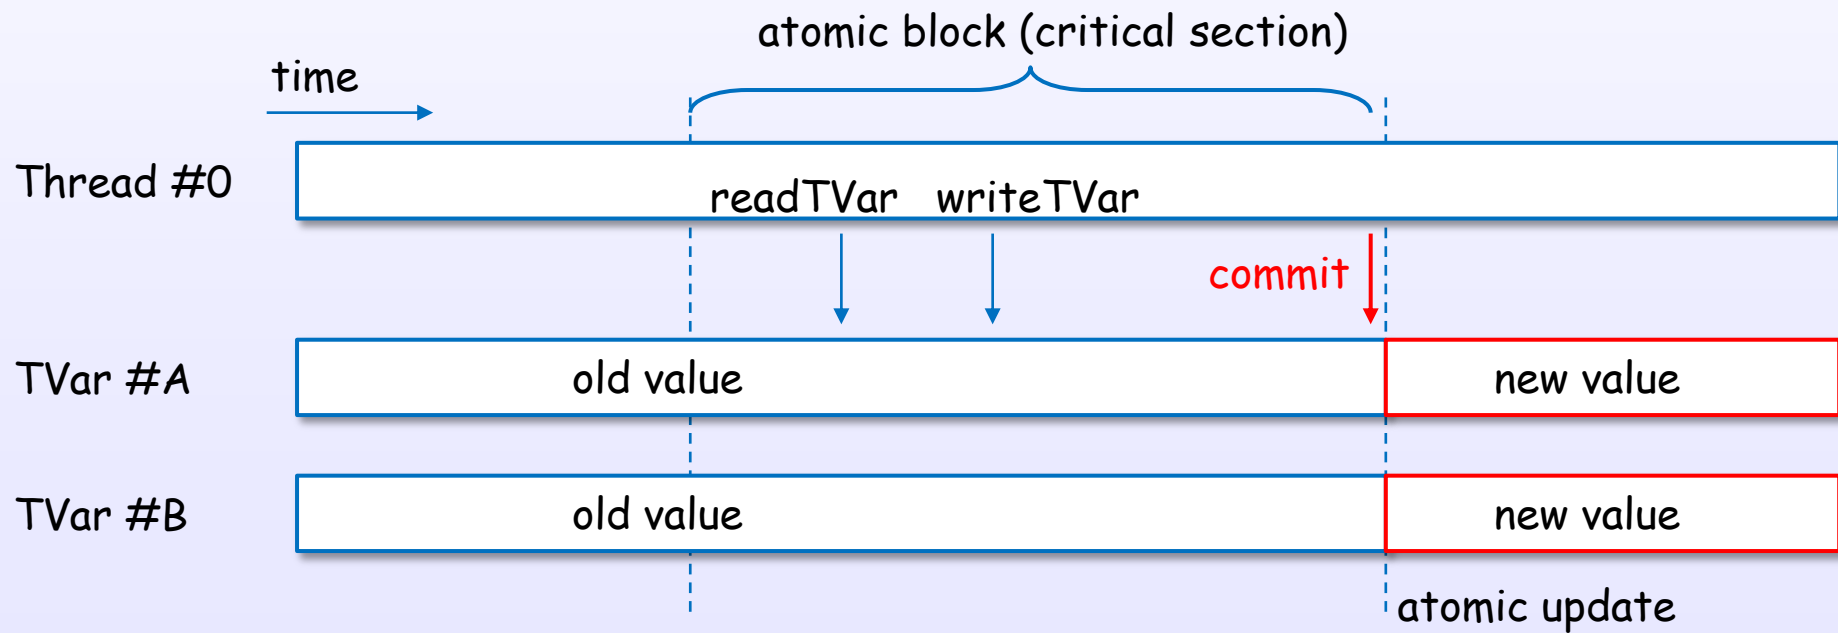
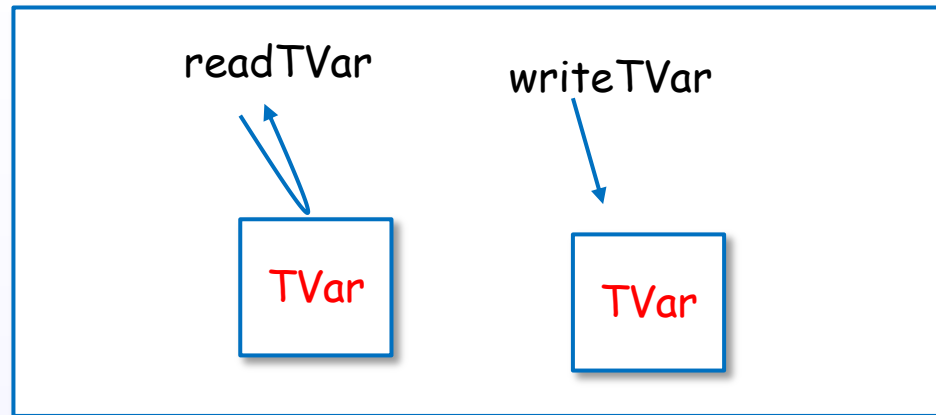
STM a




A **or** B or Nothing

STM, TVar example (normal case)

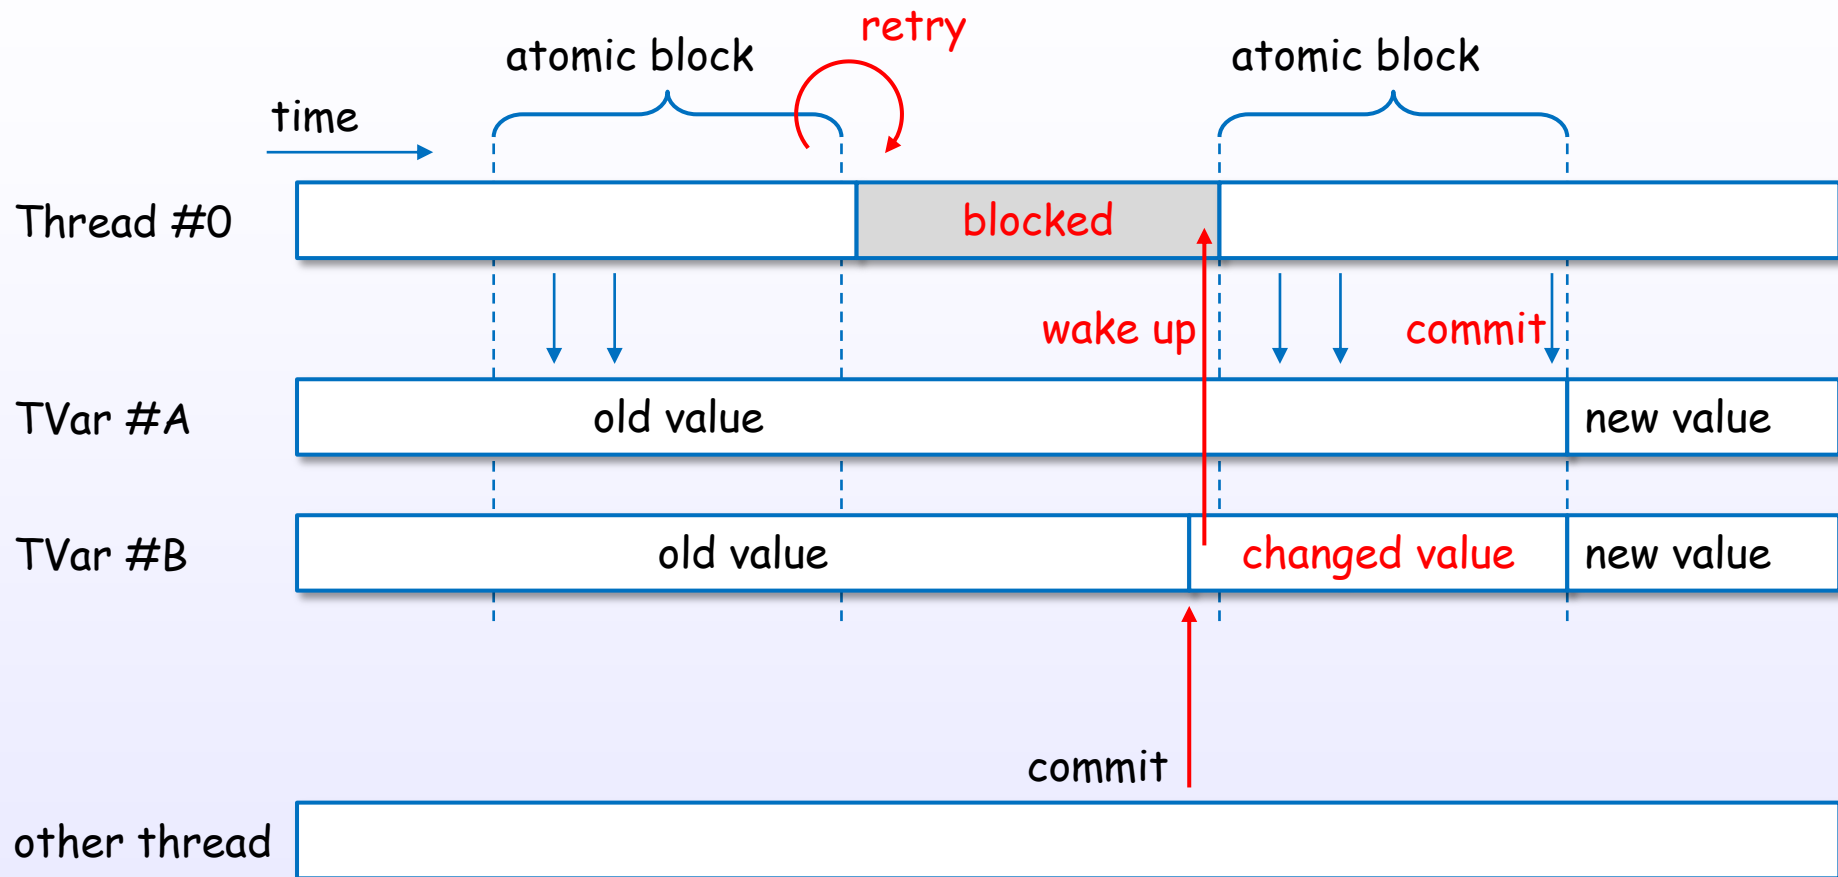
STM a




STM, TVar example (conflict case)



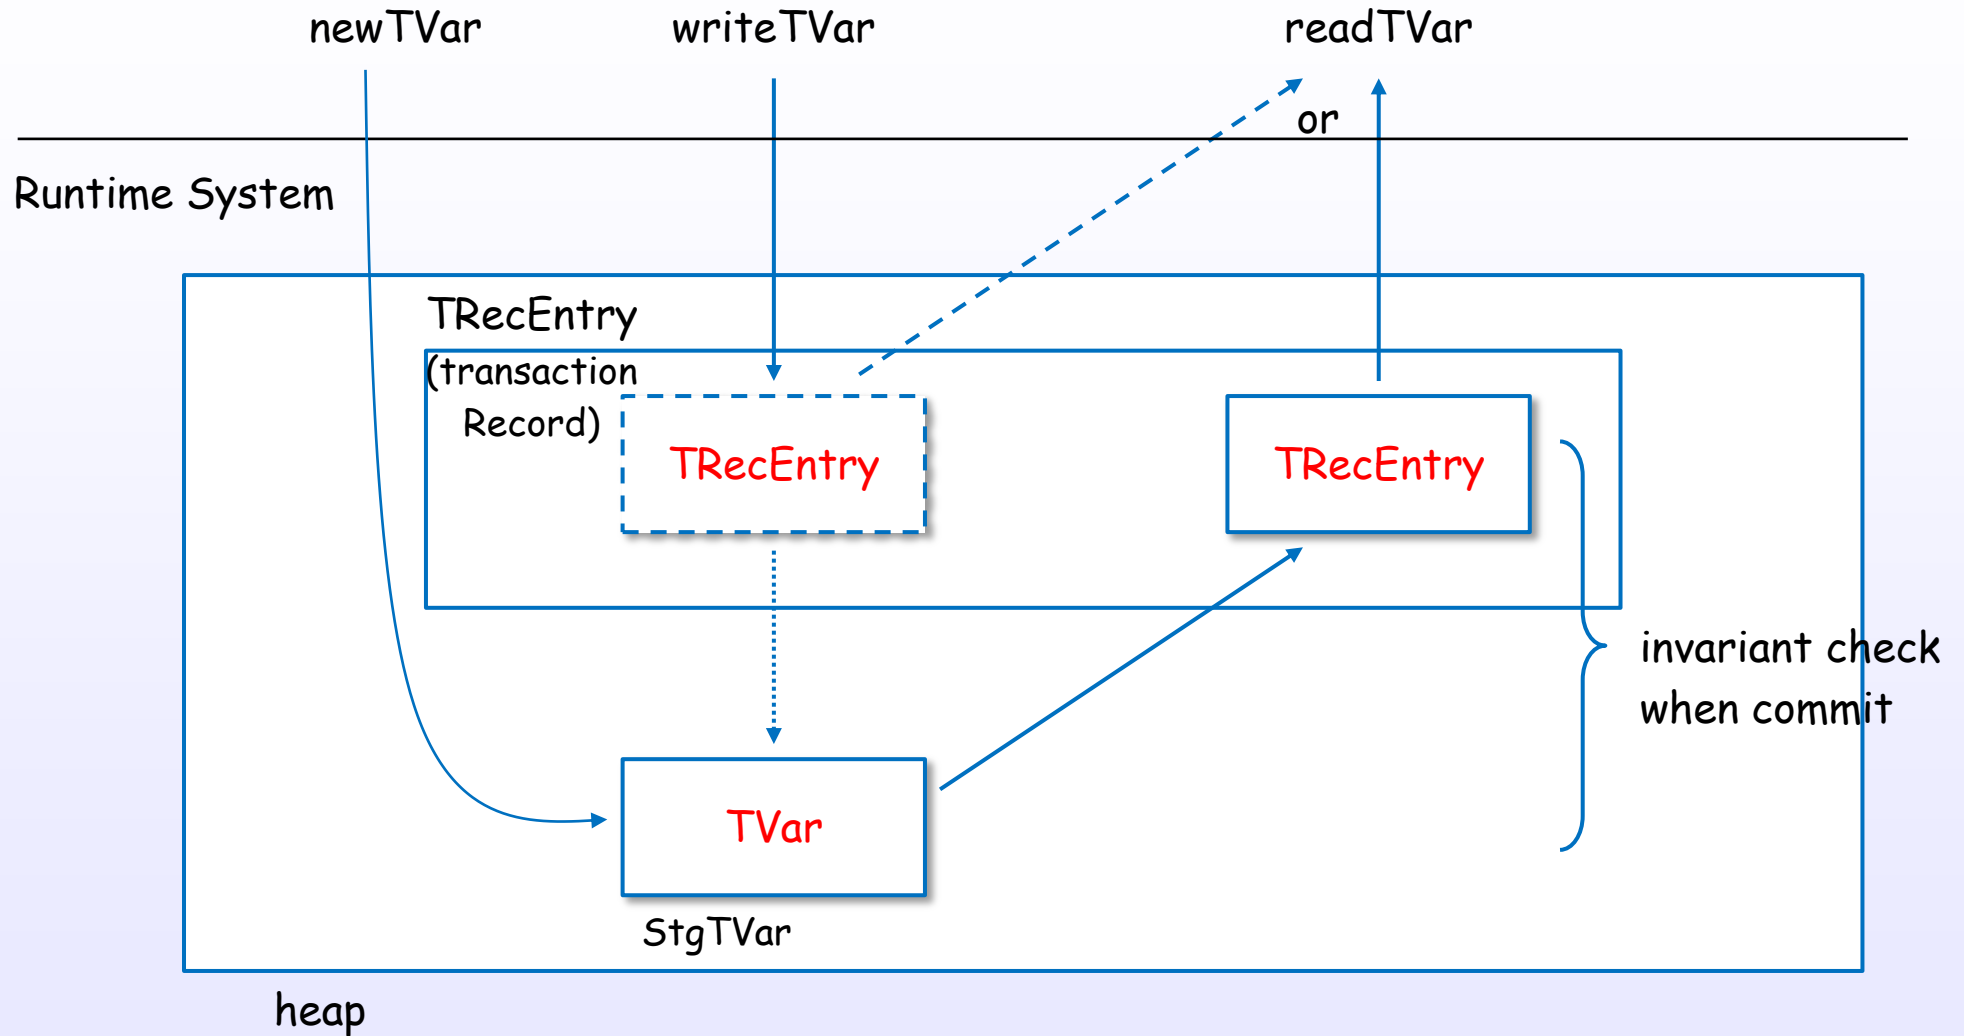
retry example



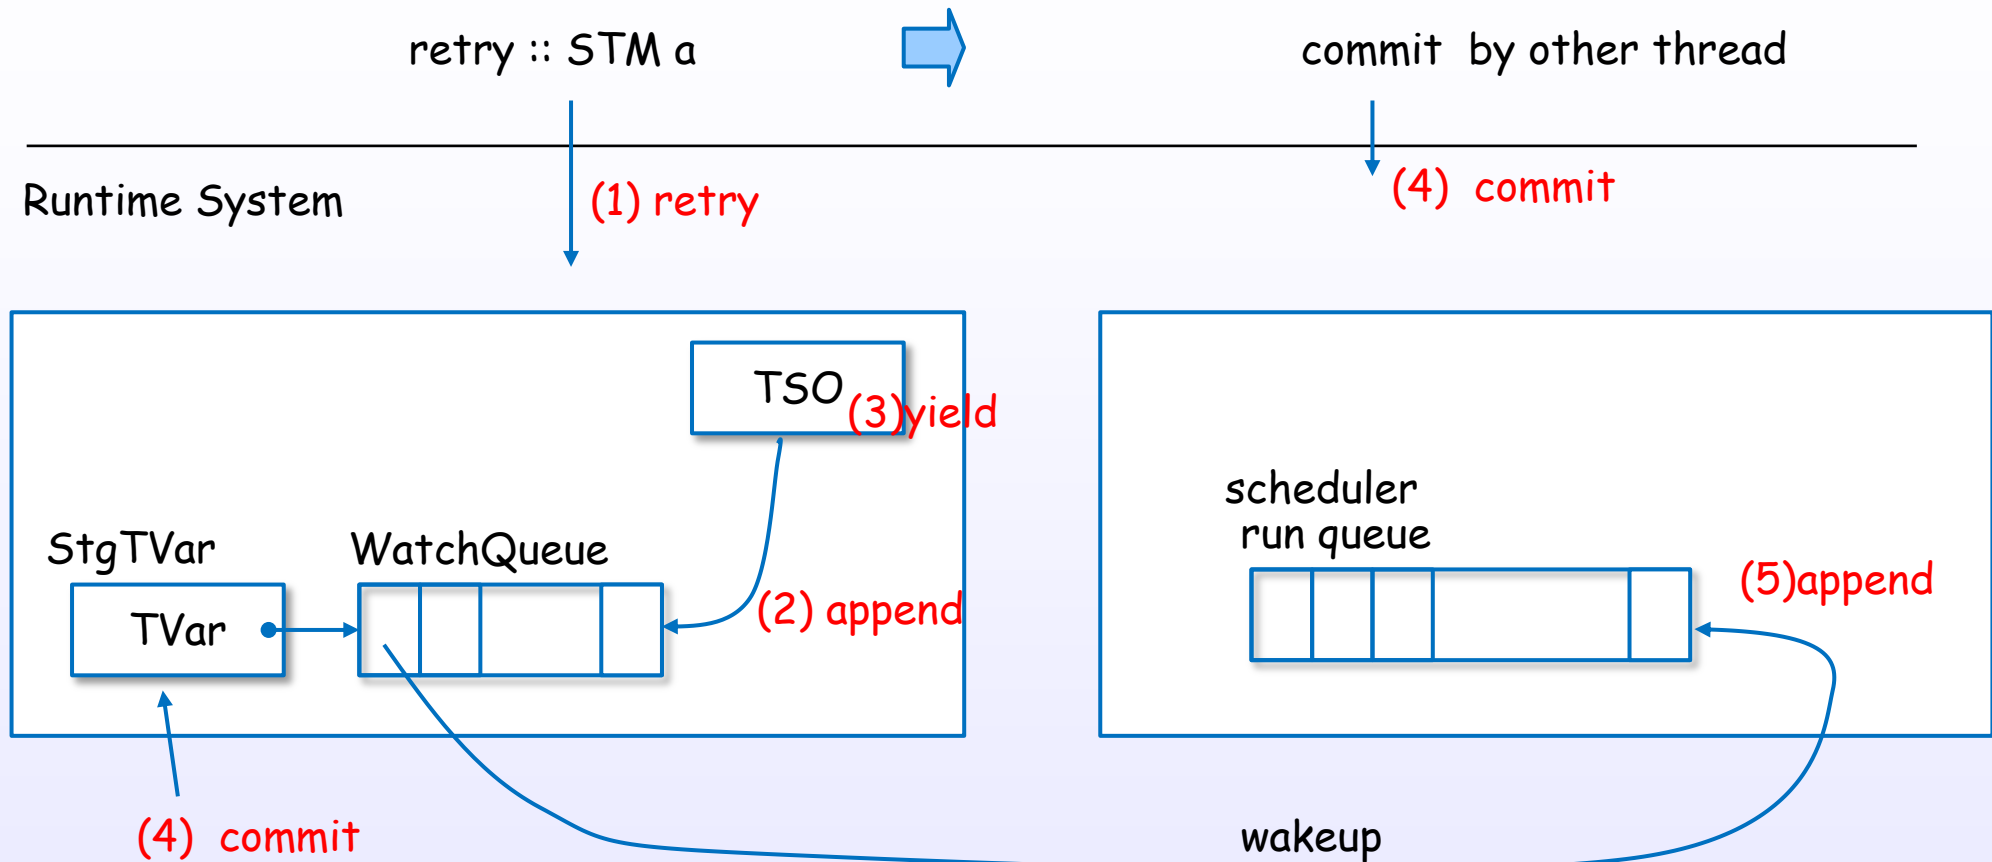
STM, TVar data structure



newTVar, writeTVar, readTVar

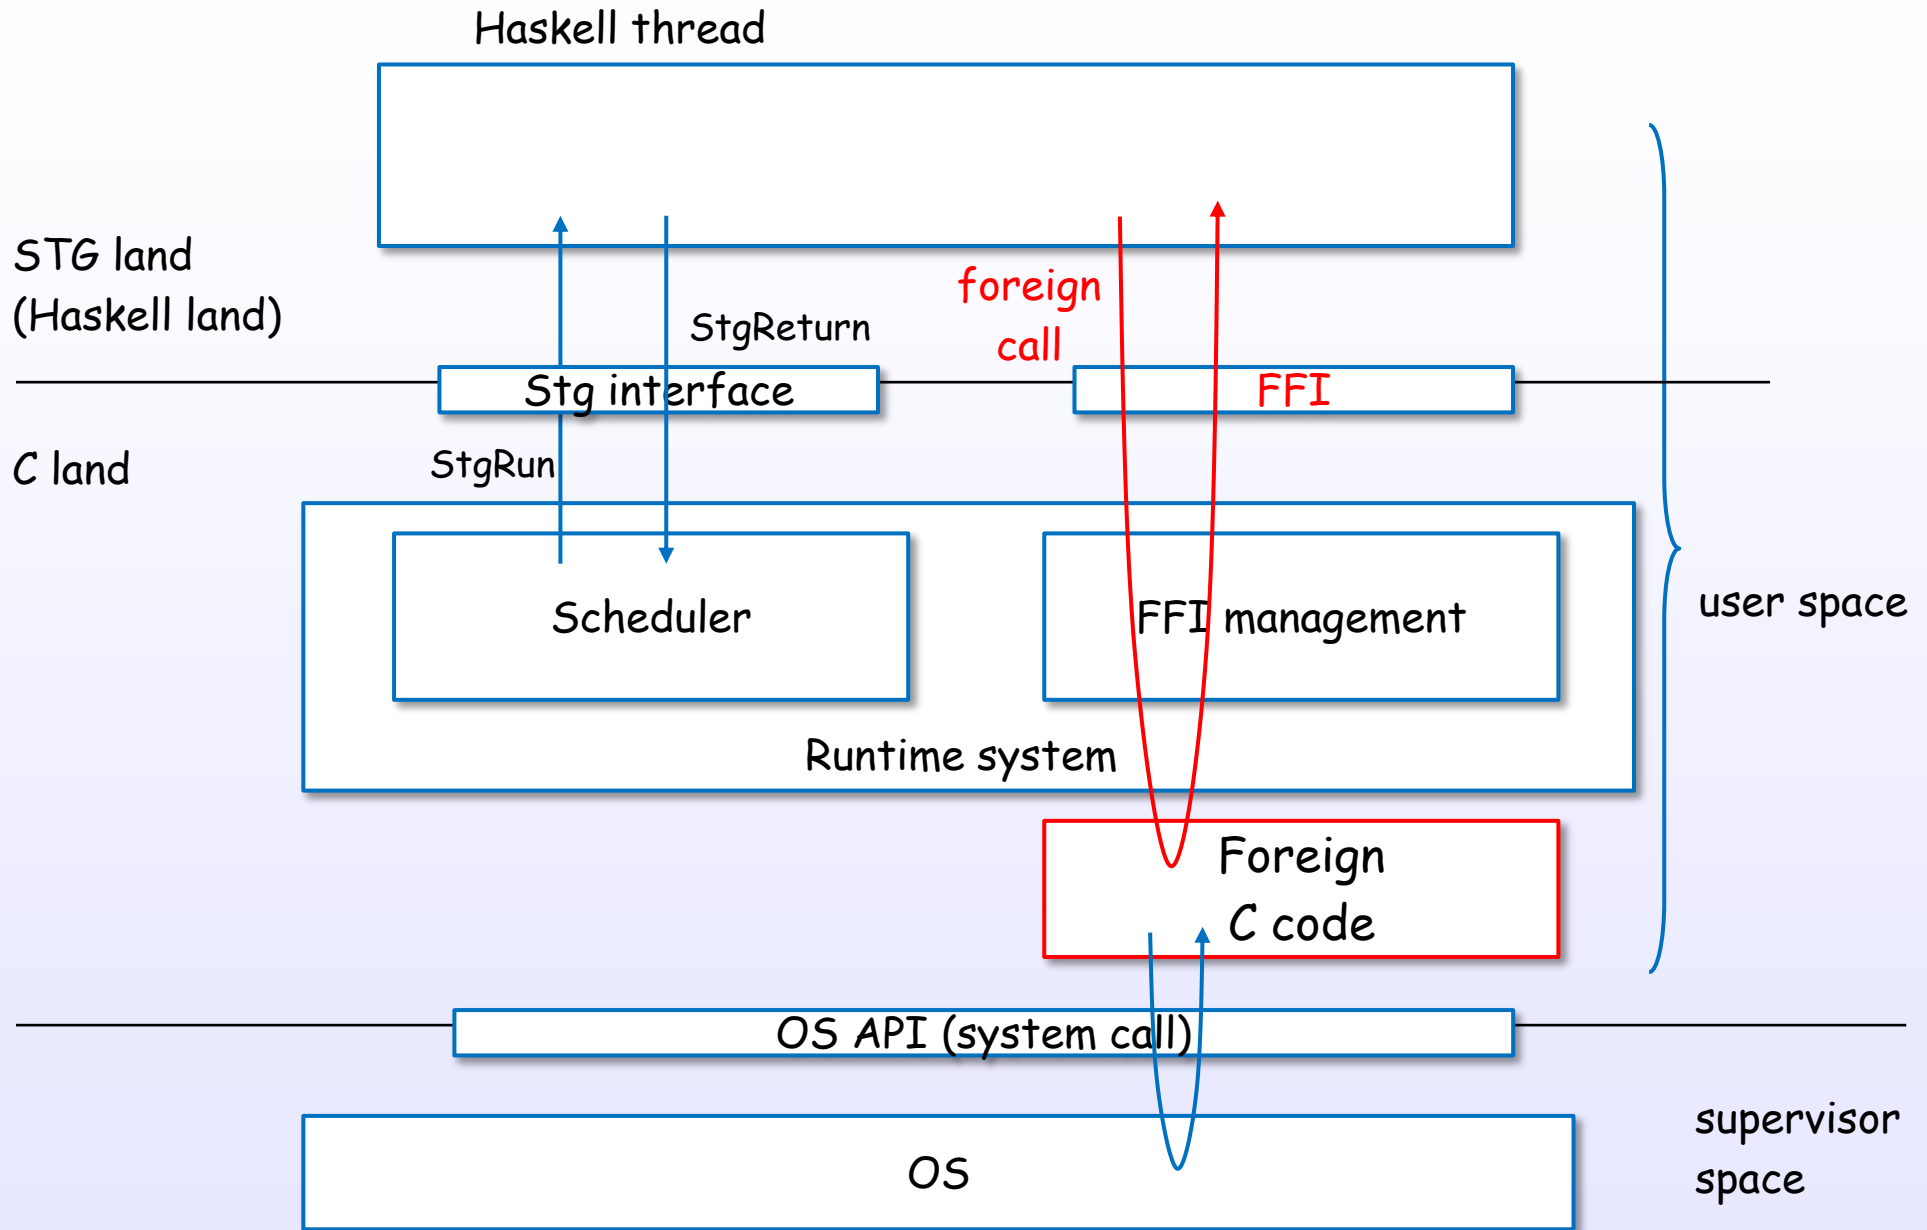


block by retry, wake up by commit

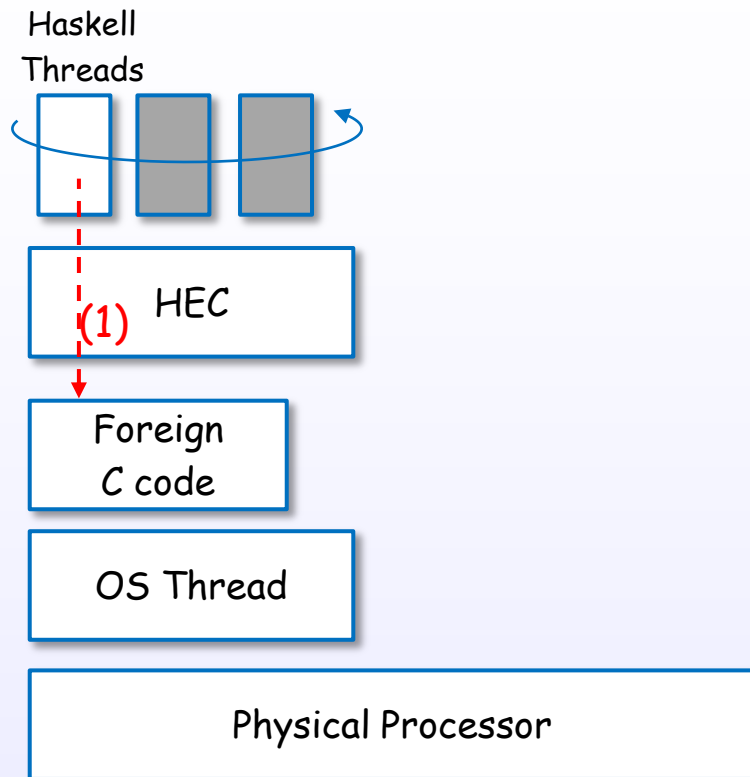


FFI

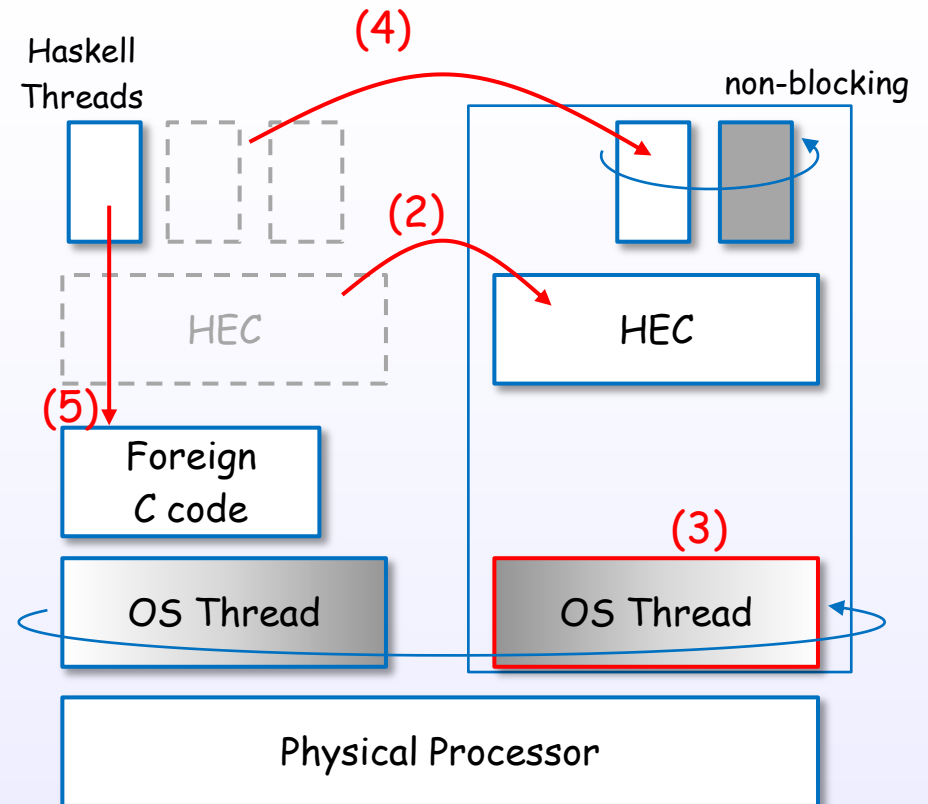
FFI (Foreign Function Interface)



FFI and OS Threads



(1) a safe foreign call (FFI)



(2) move the HEC to other OS thread

(3) spawn or draw an OS thread

(4) move Haskell threads

(5) call the foreign C code

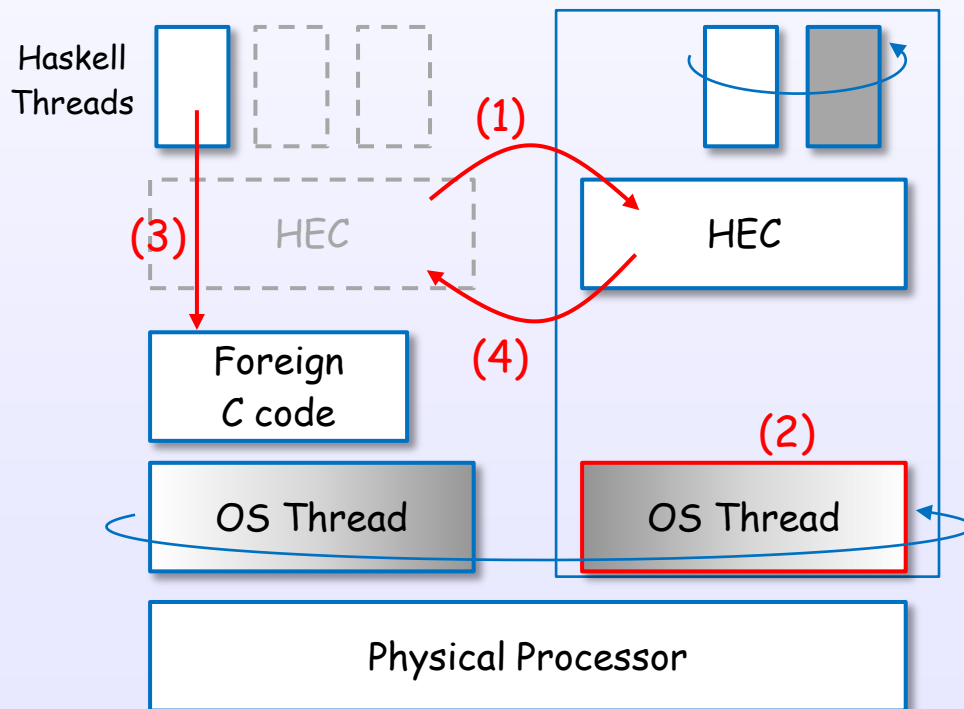
A safe foreign call (code)

Haskell Threads

```
ccall suspendThread  
ccall FOREIGN_C_CODE ... (3)  
ccall resumeThread
```

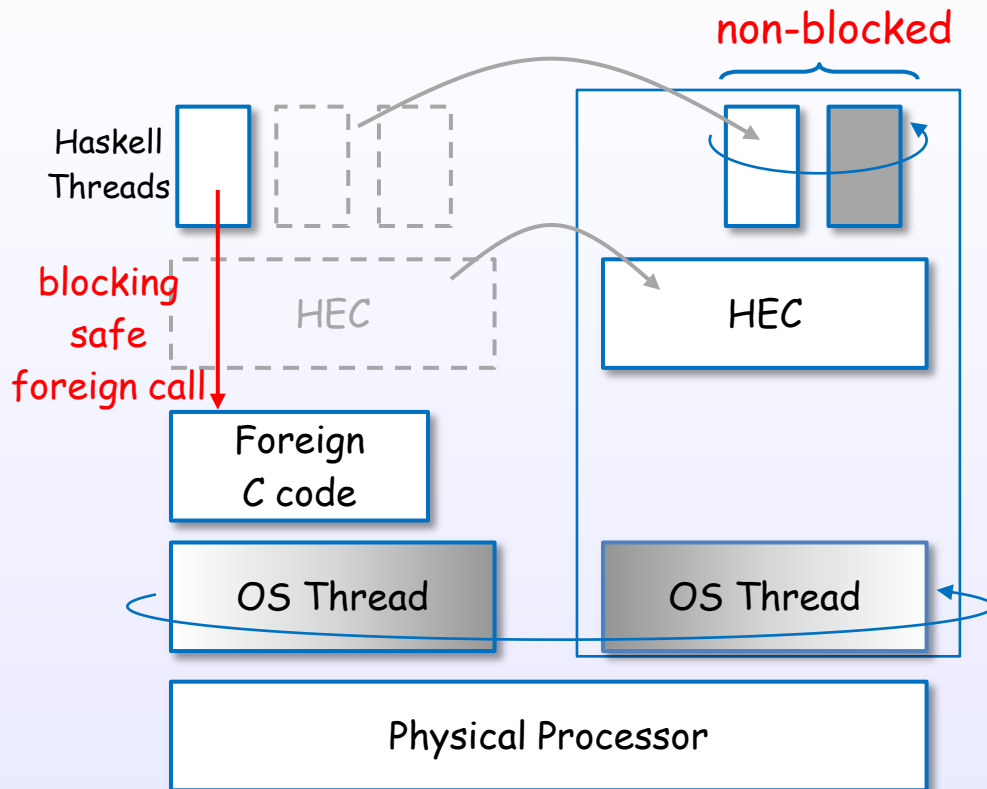
```
releaseCapability_  
giveCapabilityToTask ... (1)  
startWorkerTask  
createOSThread ... (2)
```

```
waitForReturnCapability ... (4)
```

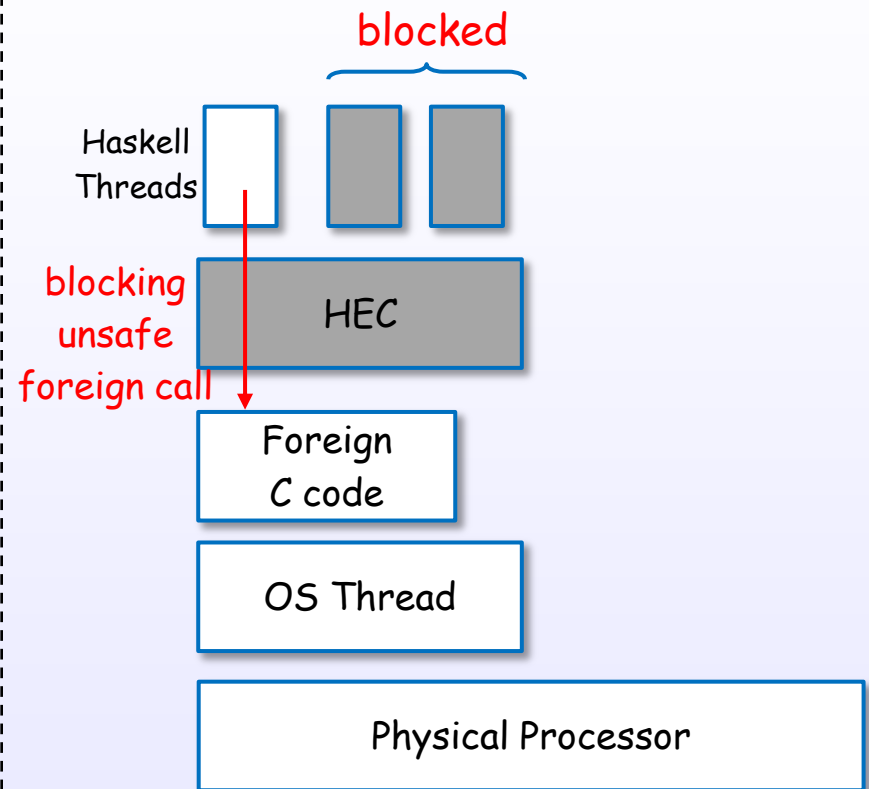


a safe and an unsafe foreign call

a **safe** foreign call



an **unsafe** foreign call




faster,
but blocking to the other Haskell threads

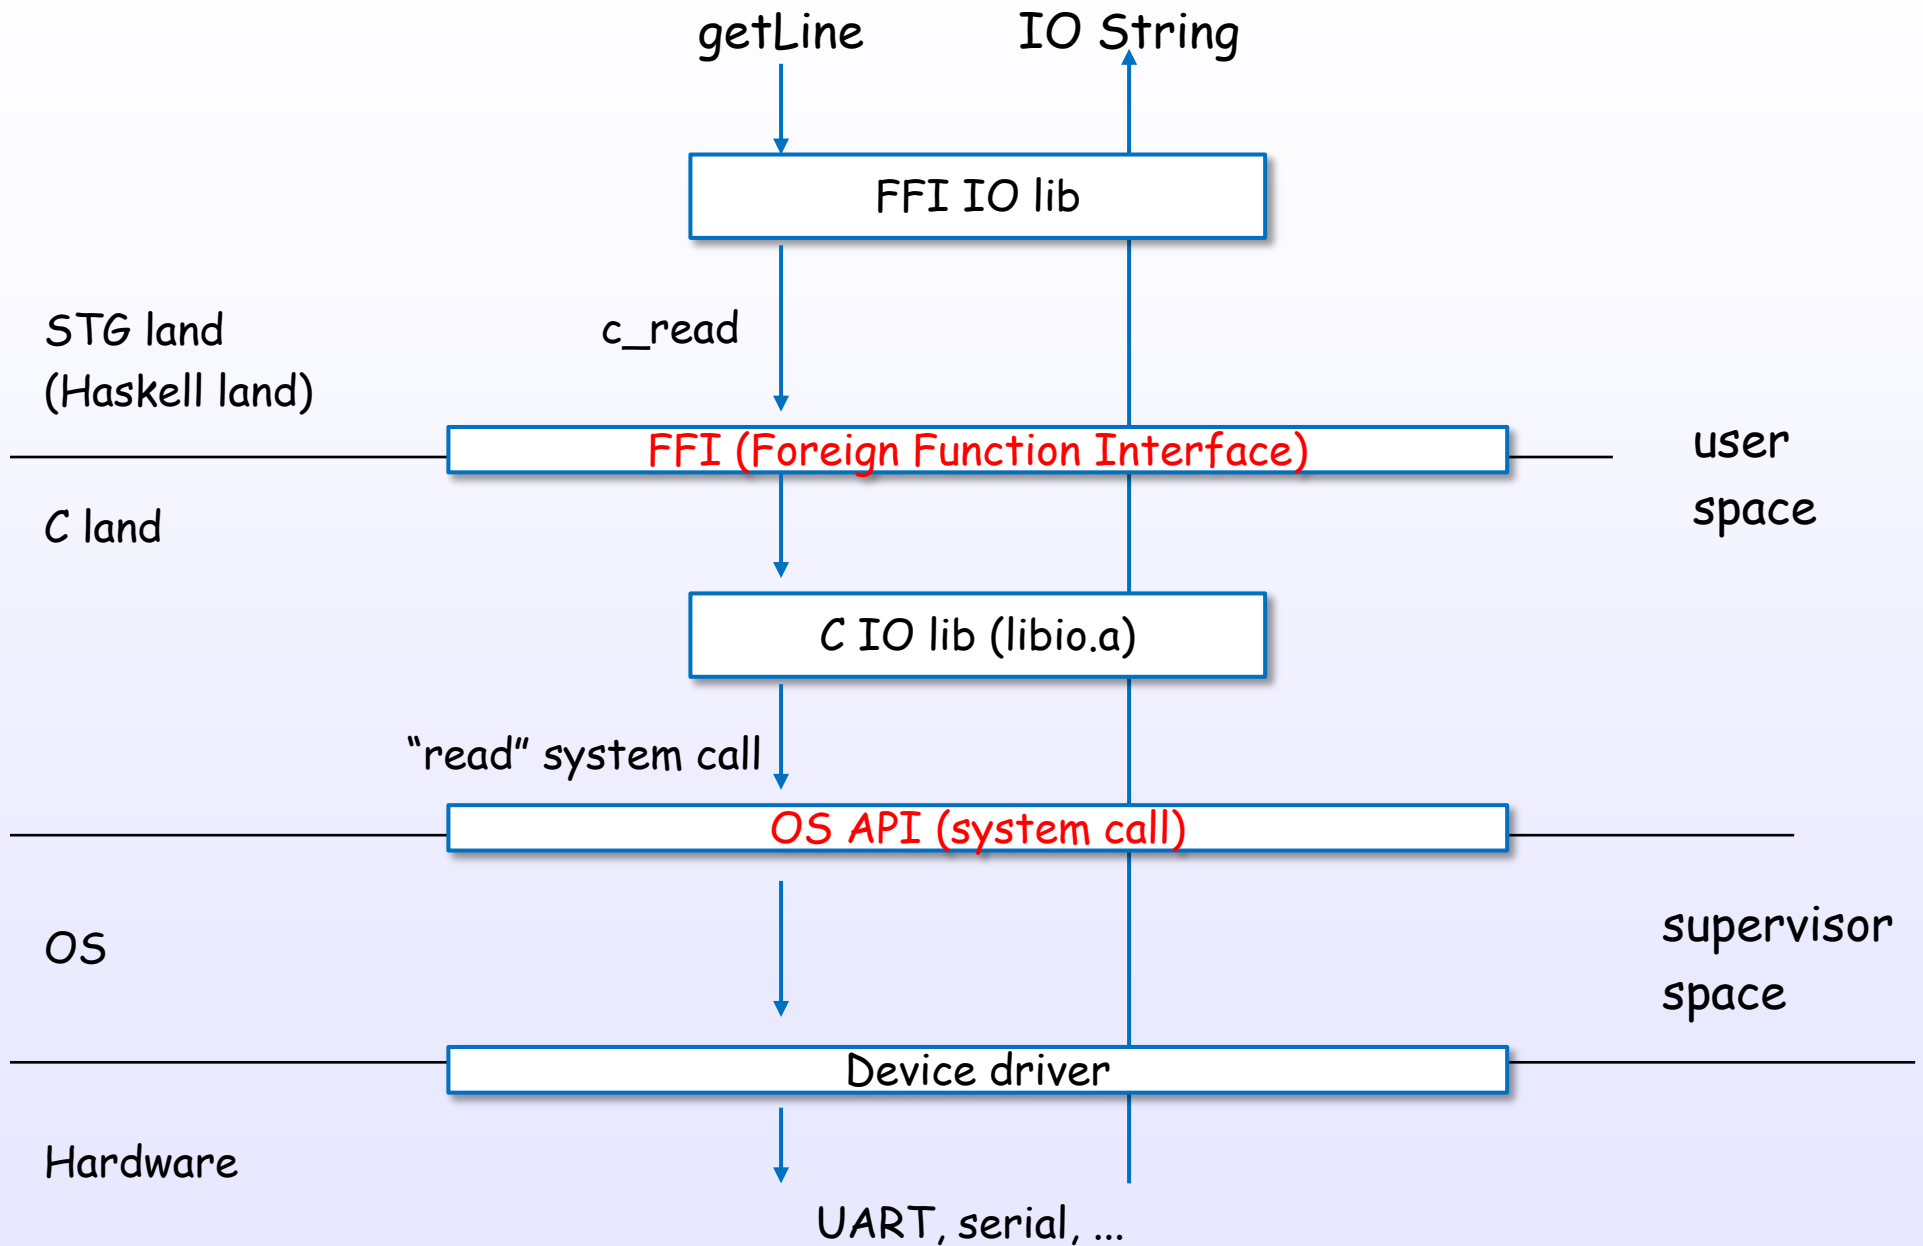
IO and FFI

IO

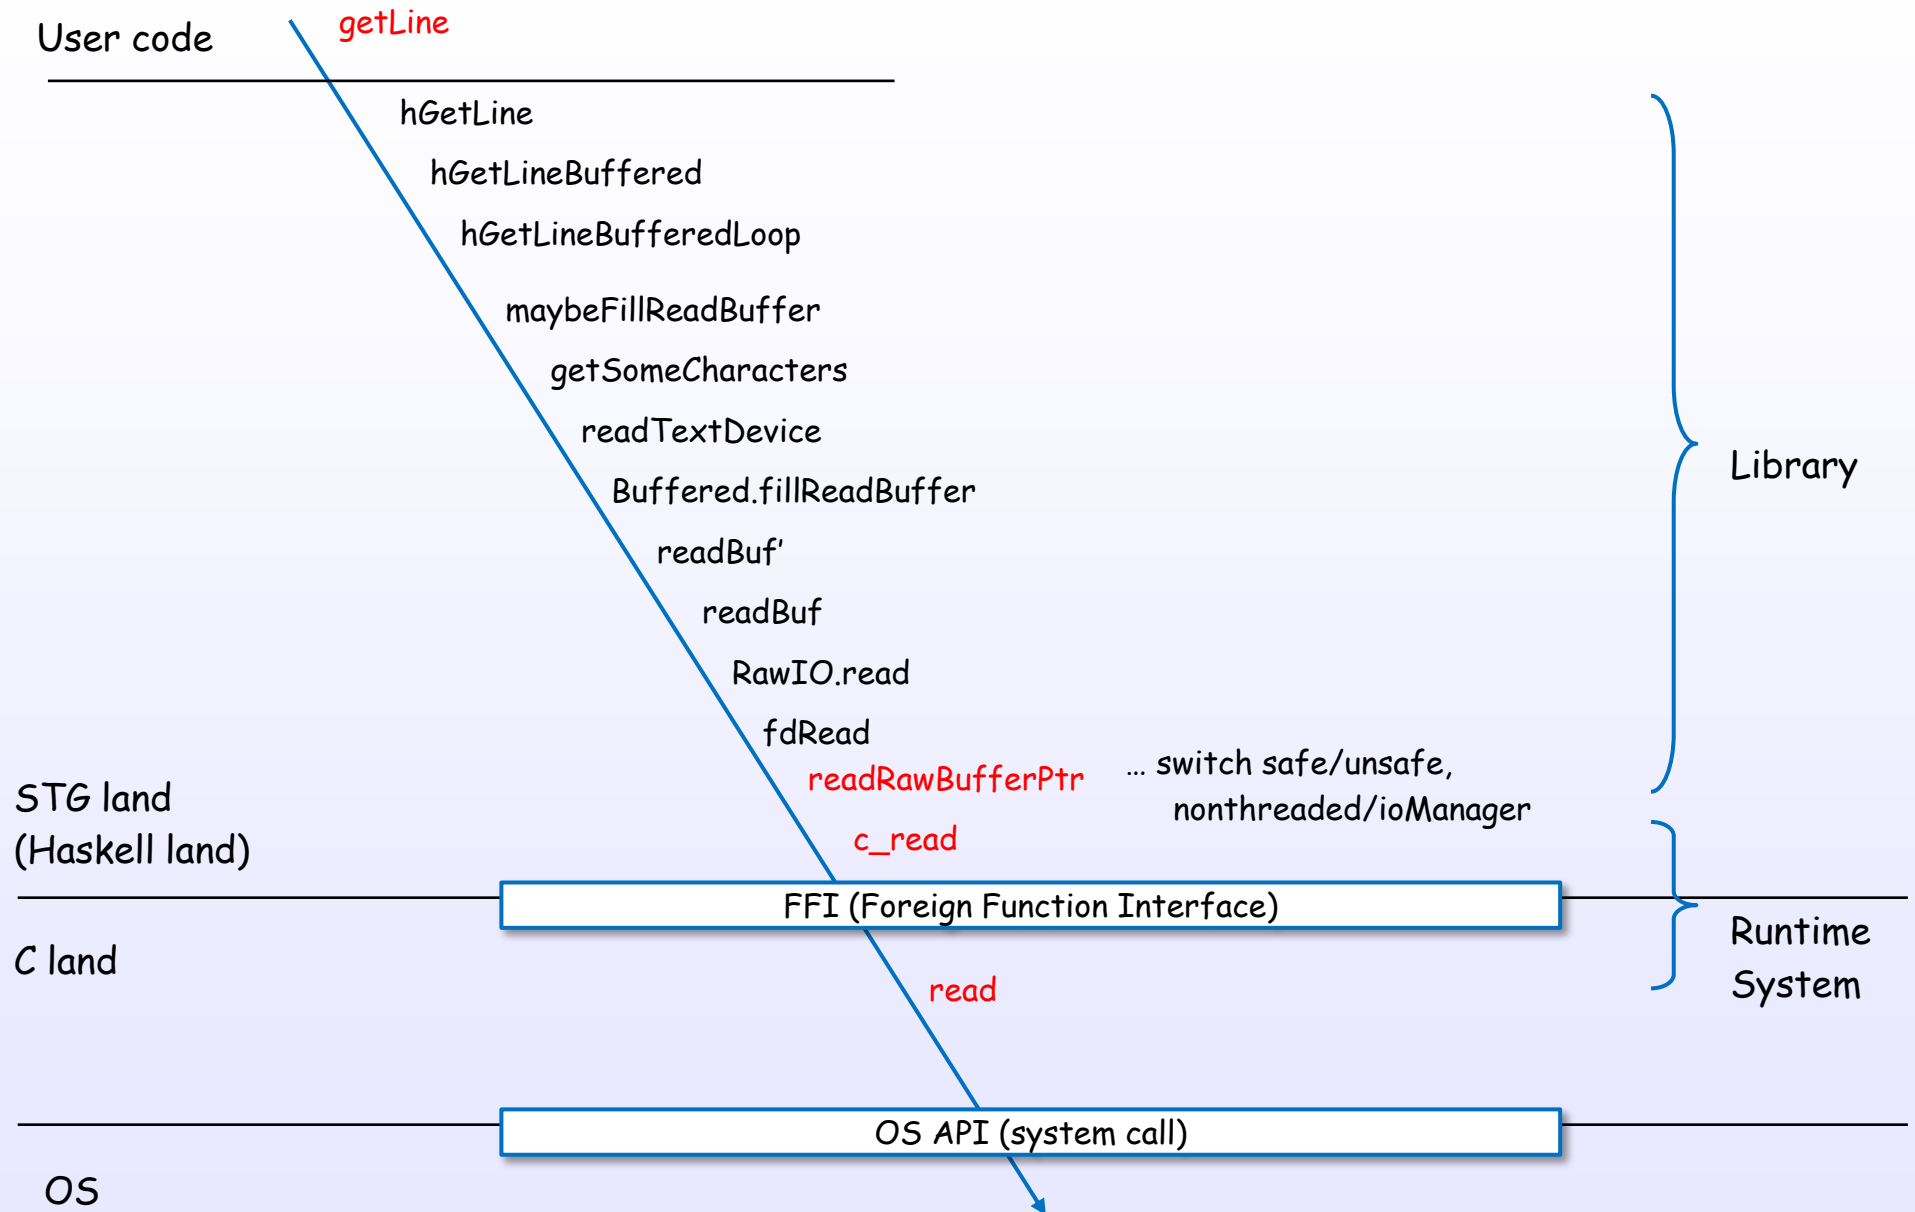
Haskell Thread



IO example: getLine

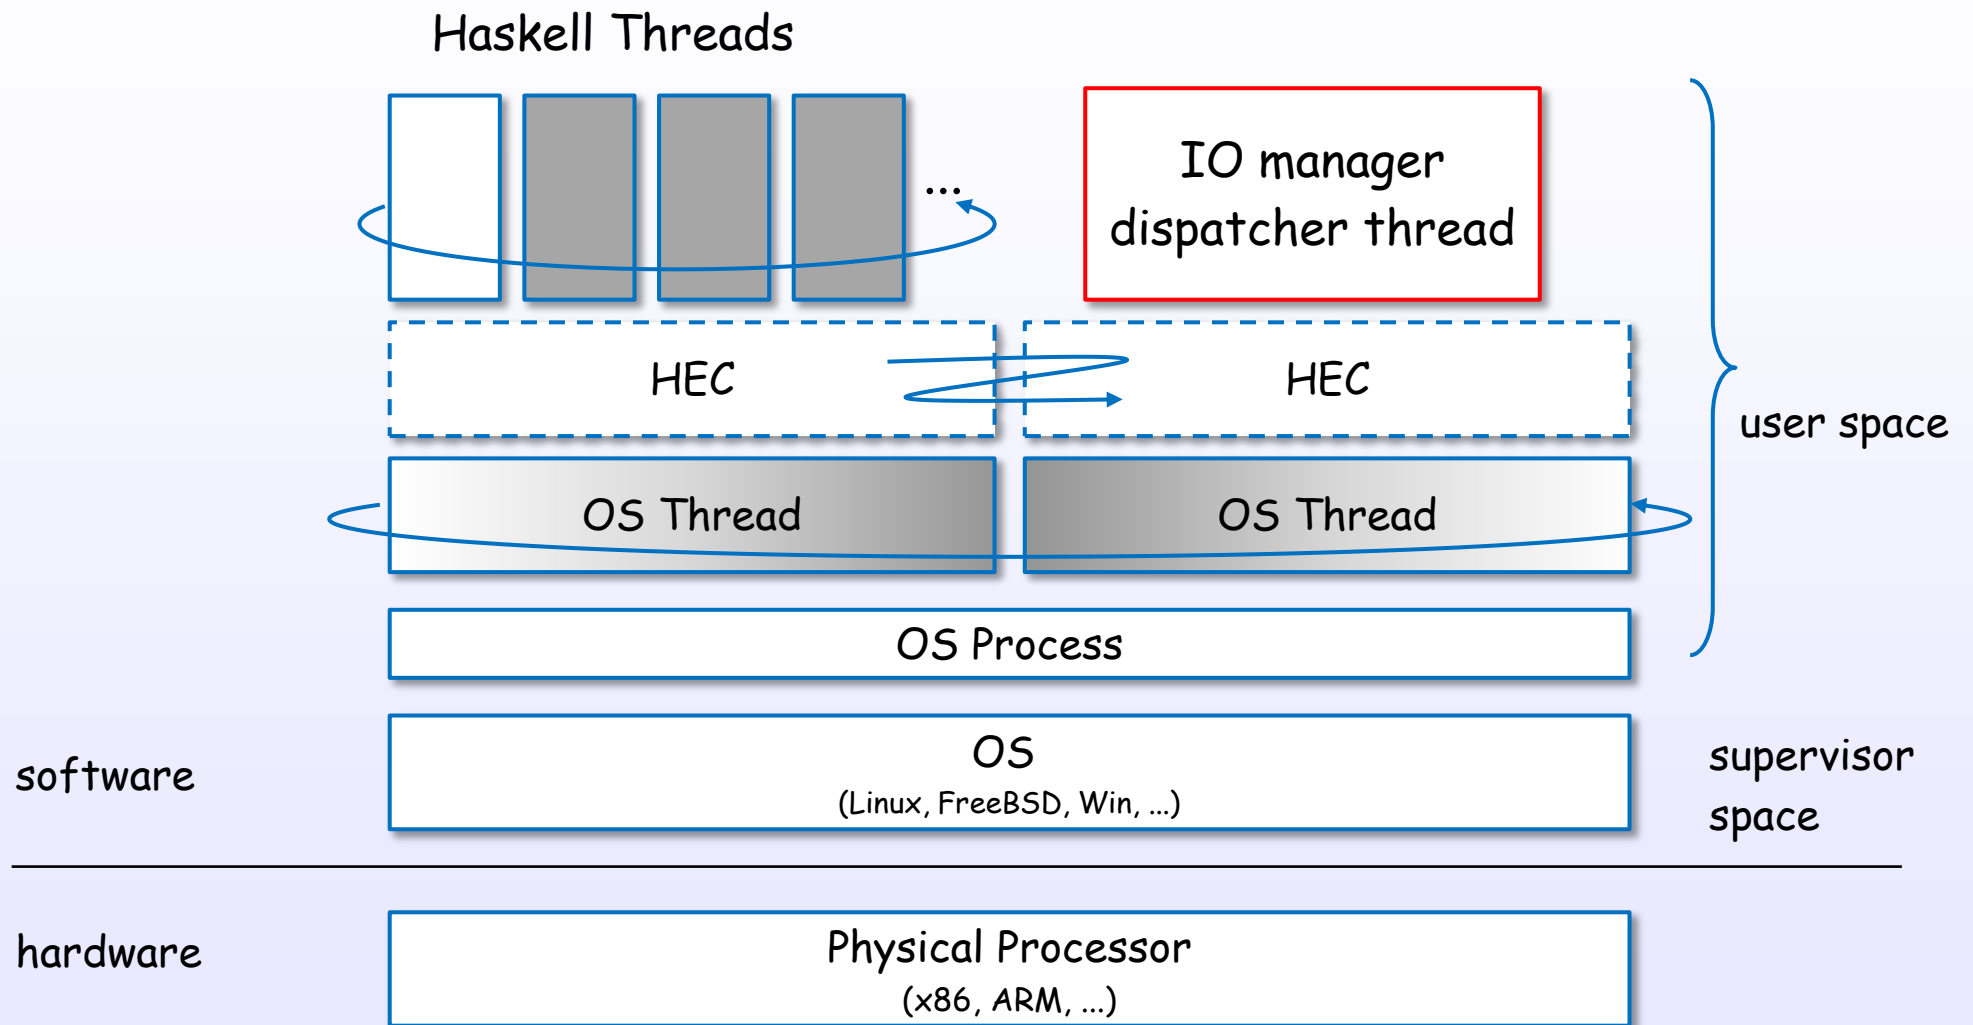


IO example: `getLine` (code)



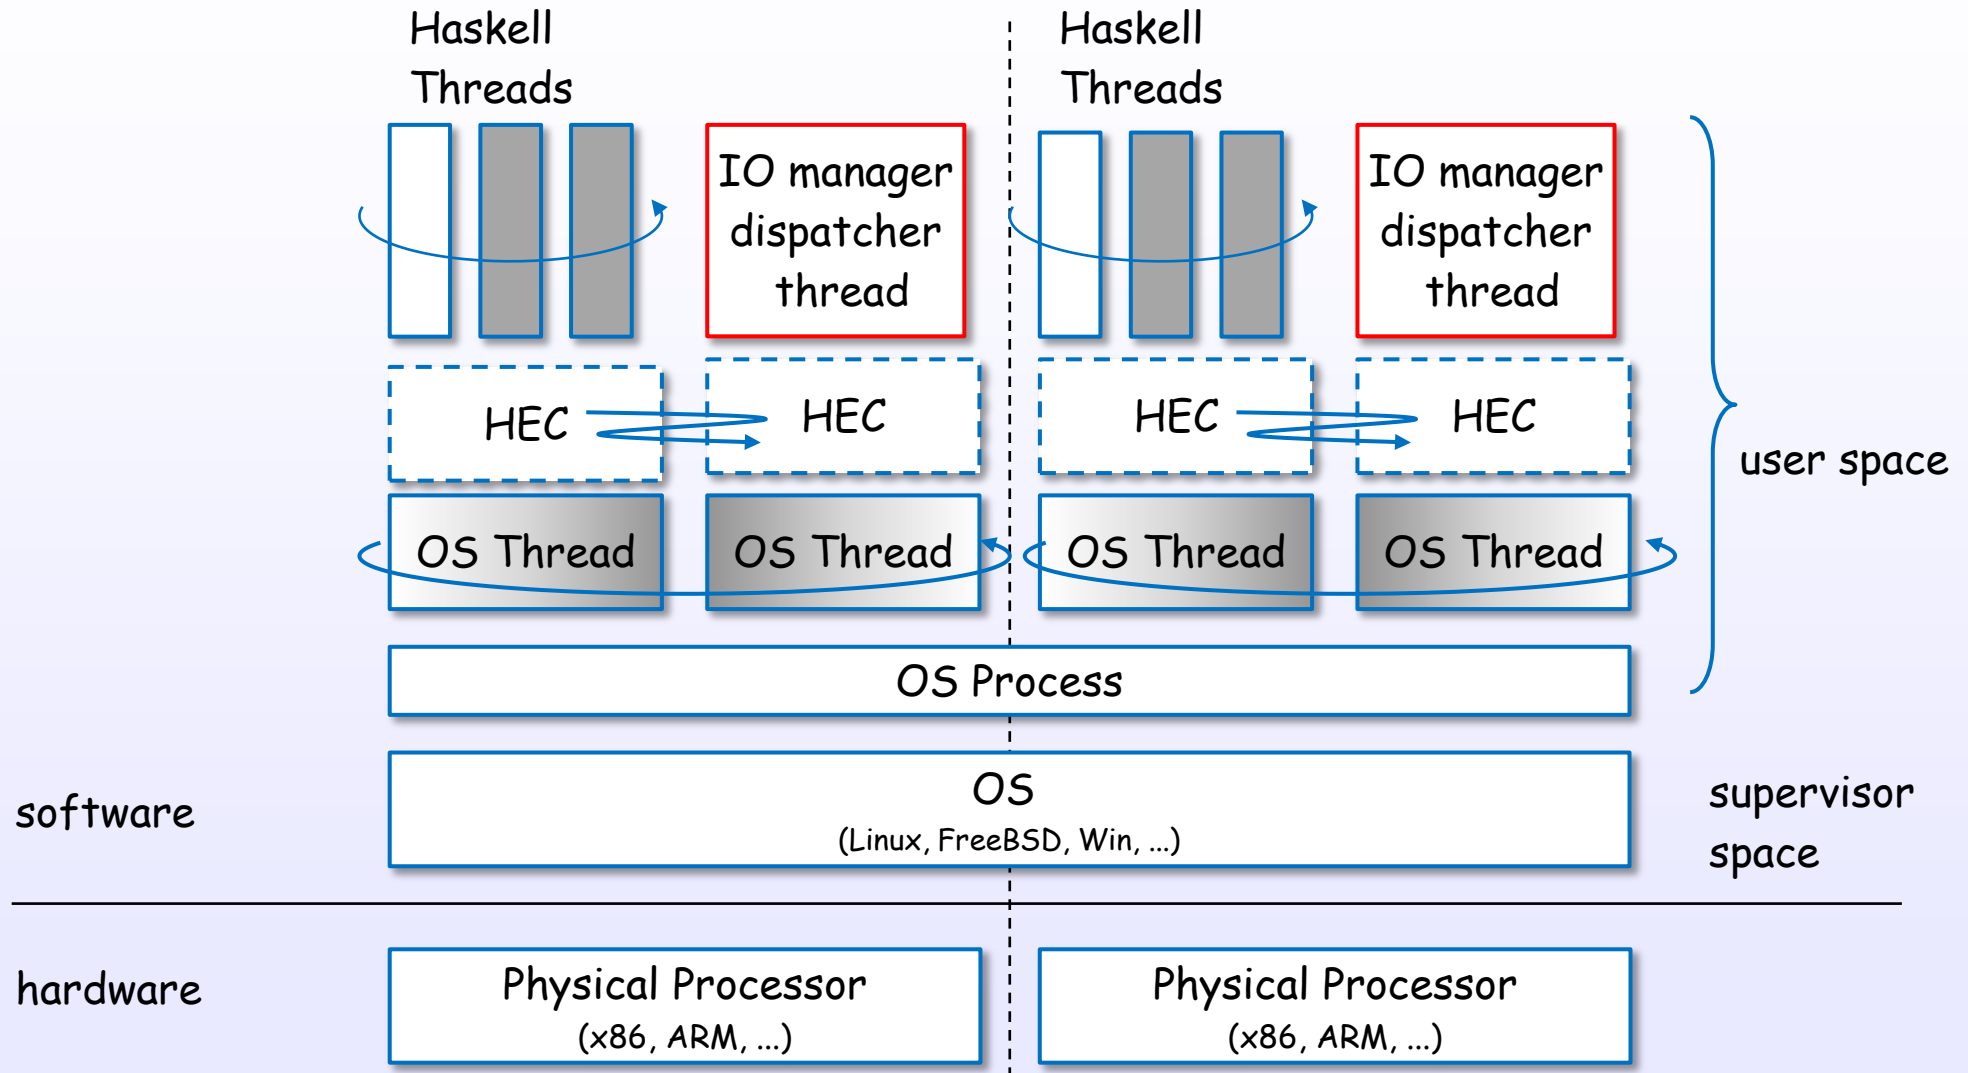
IO manager

IO manager (single core)



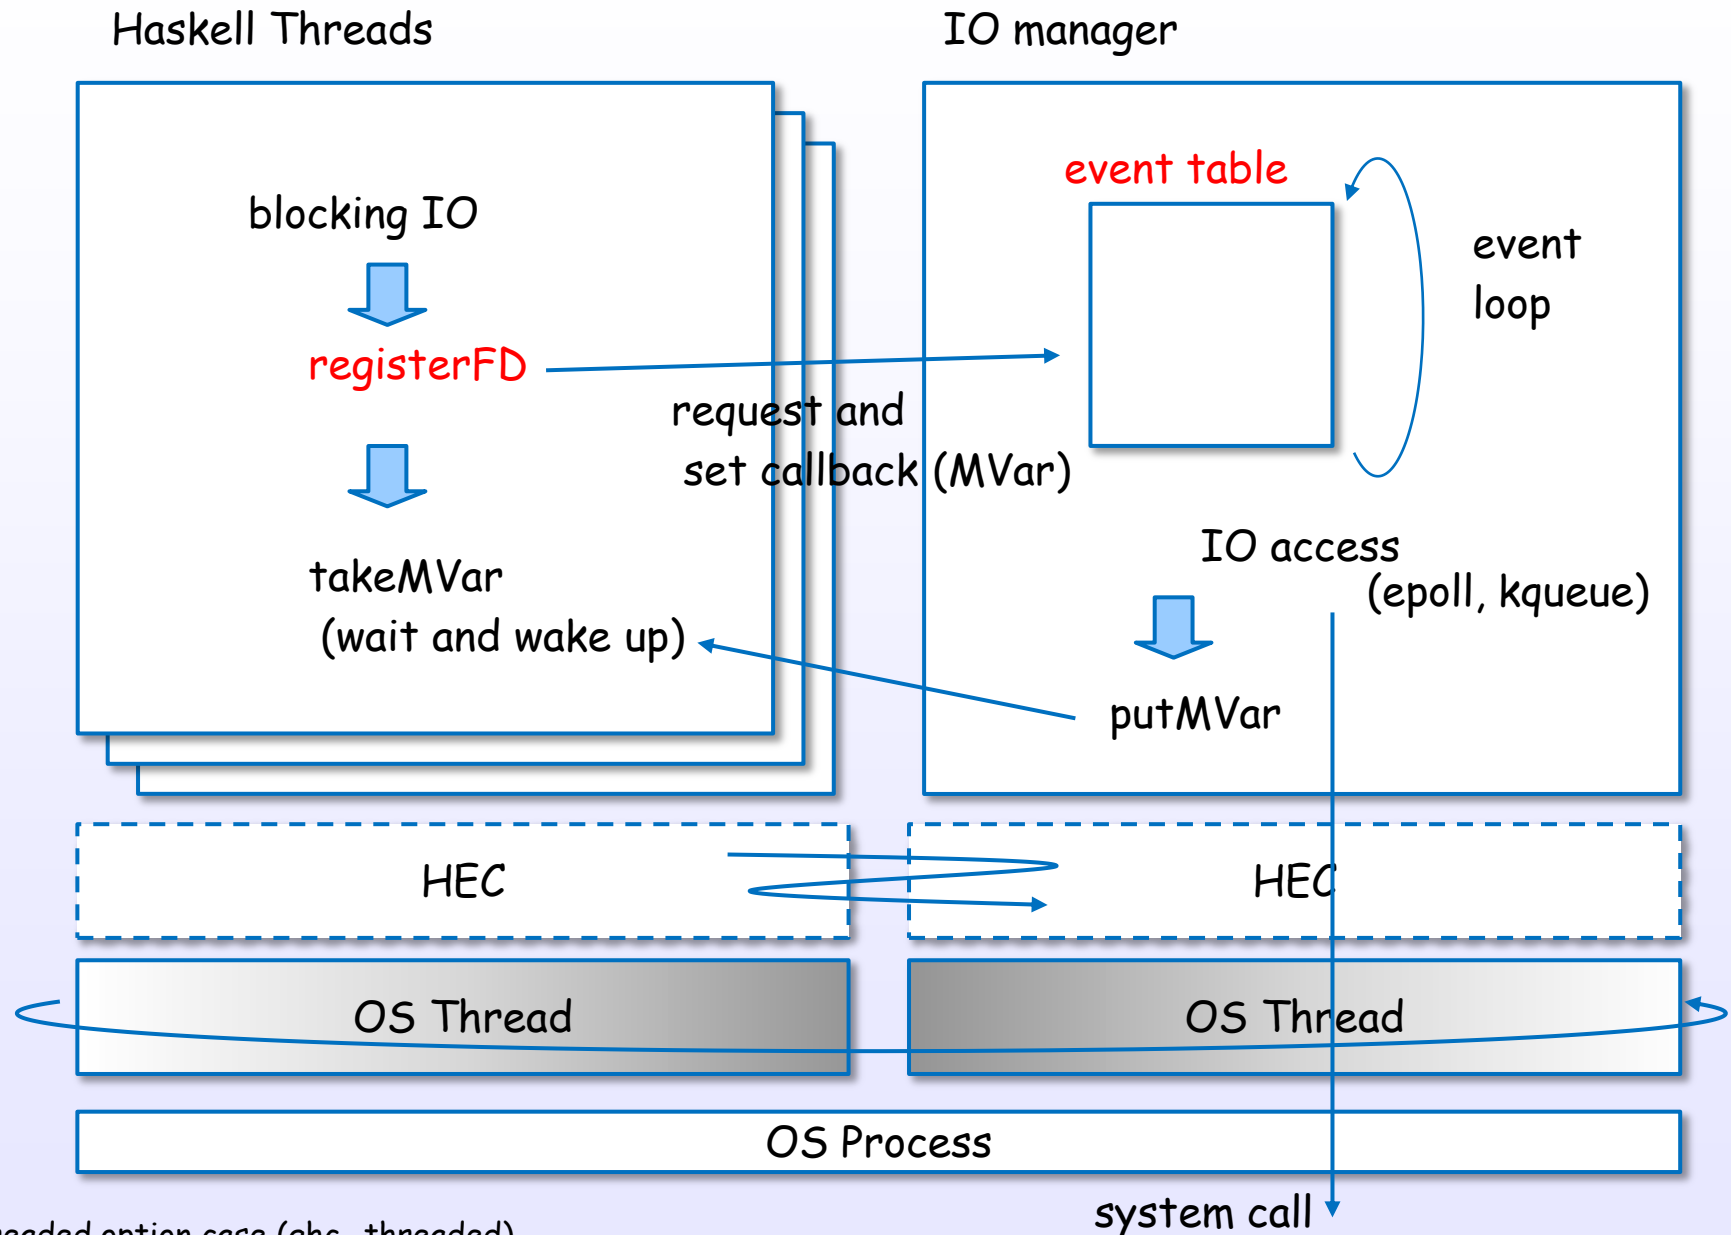
*Threaded option case (ghc -threaded)

IO manager (multi core)



*Threaded option case (ghc -threaded)

IO manager

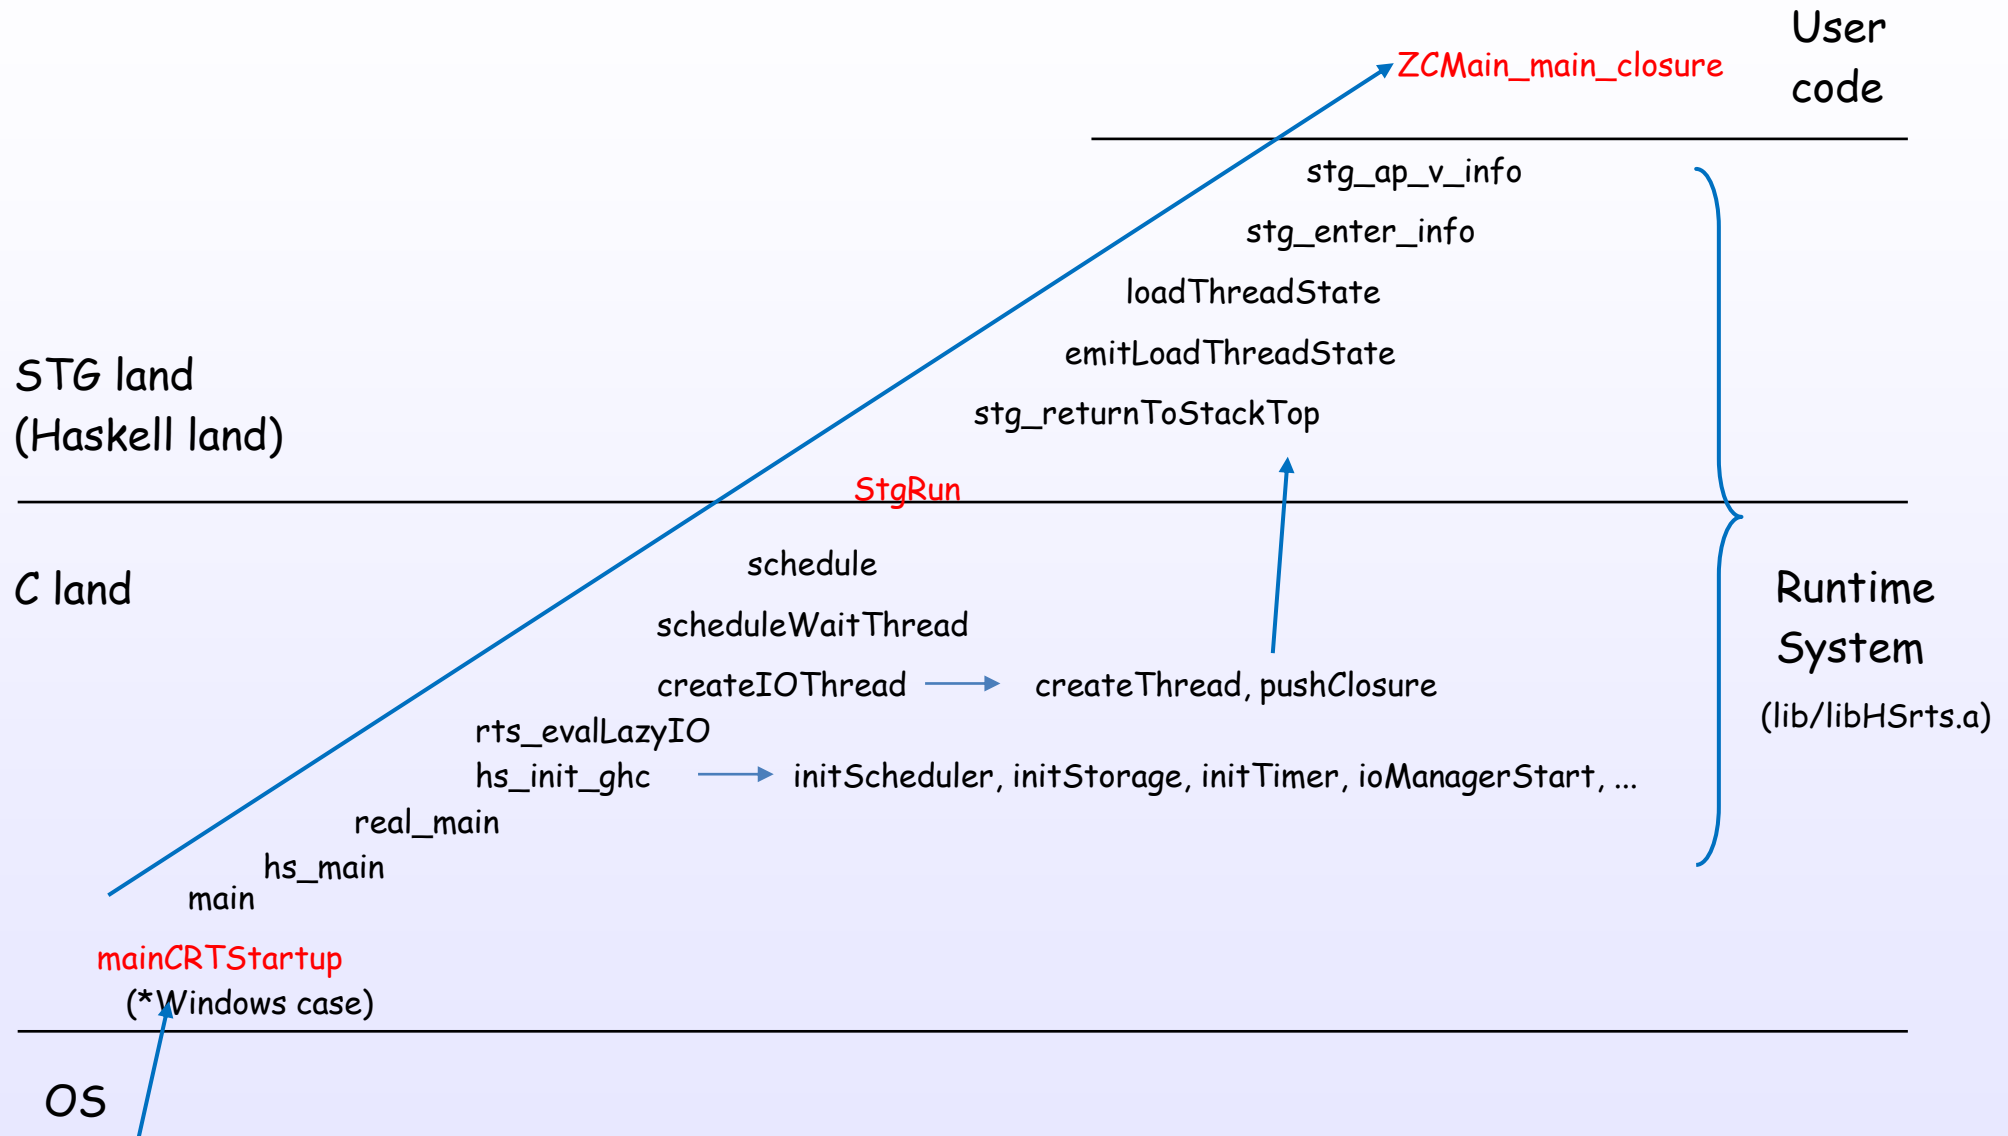


*Threaded option case (ghc -threaded)


References : [7], [5], [8], [S29], [S30], [S32], [S37], [S35], [S3]

Bootstrap

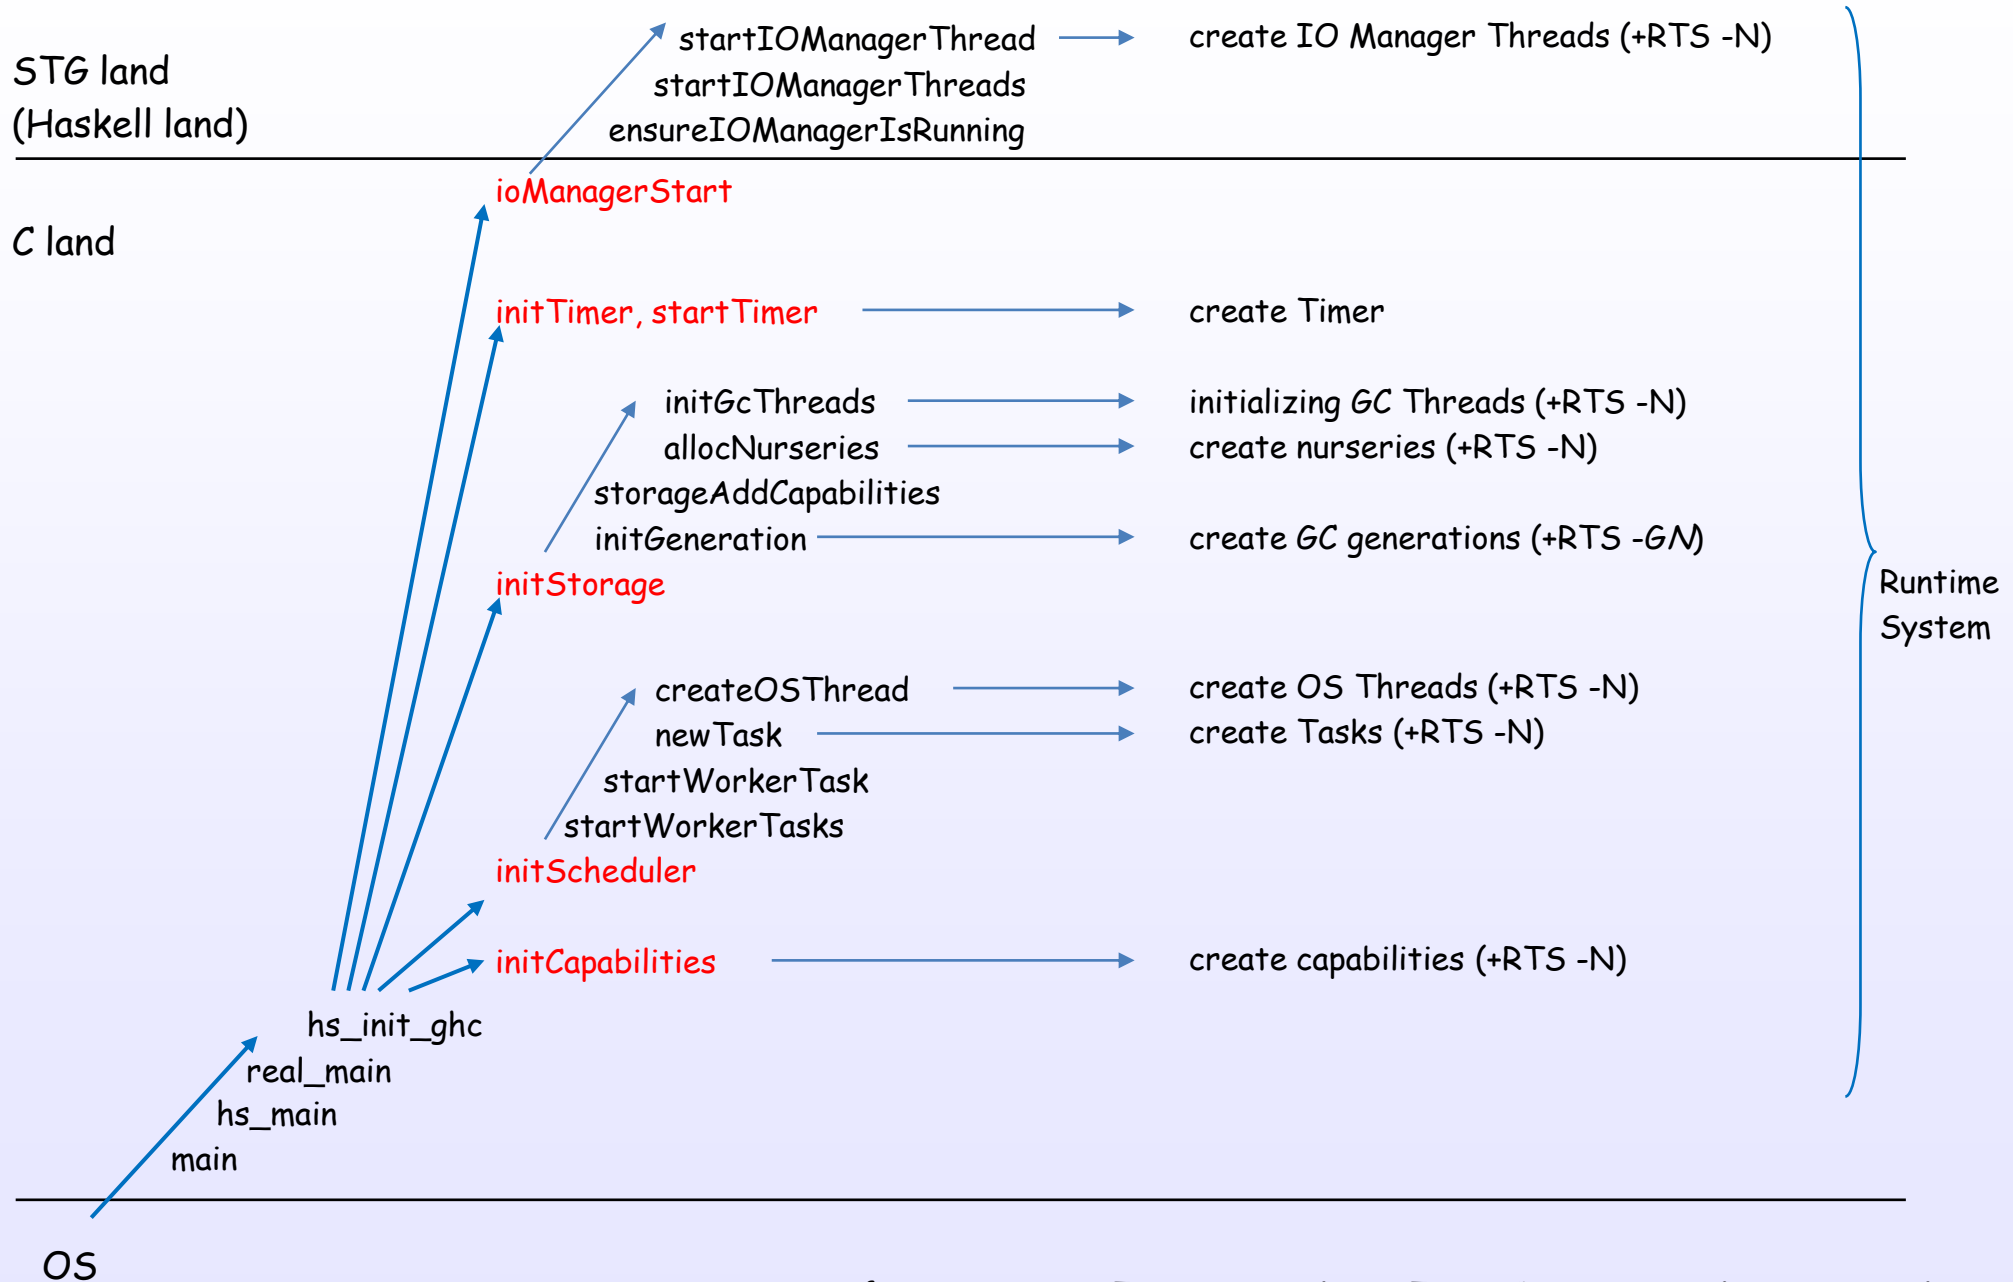
Bootstrap sequence



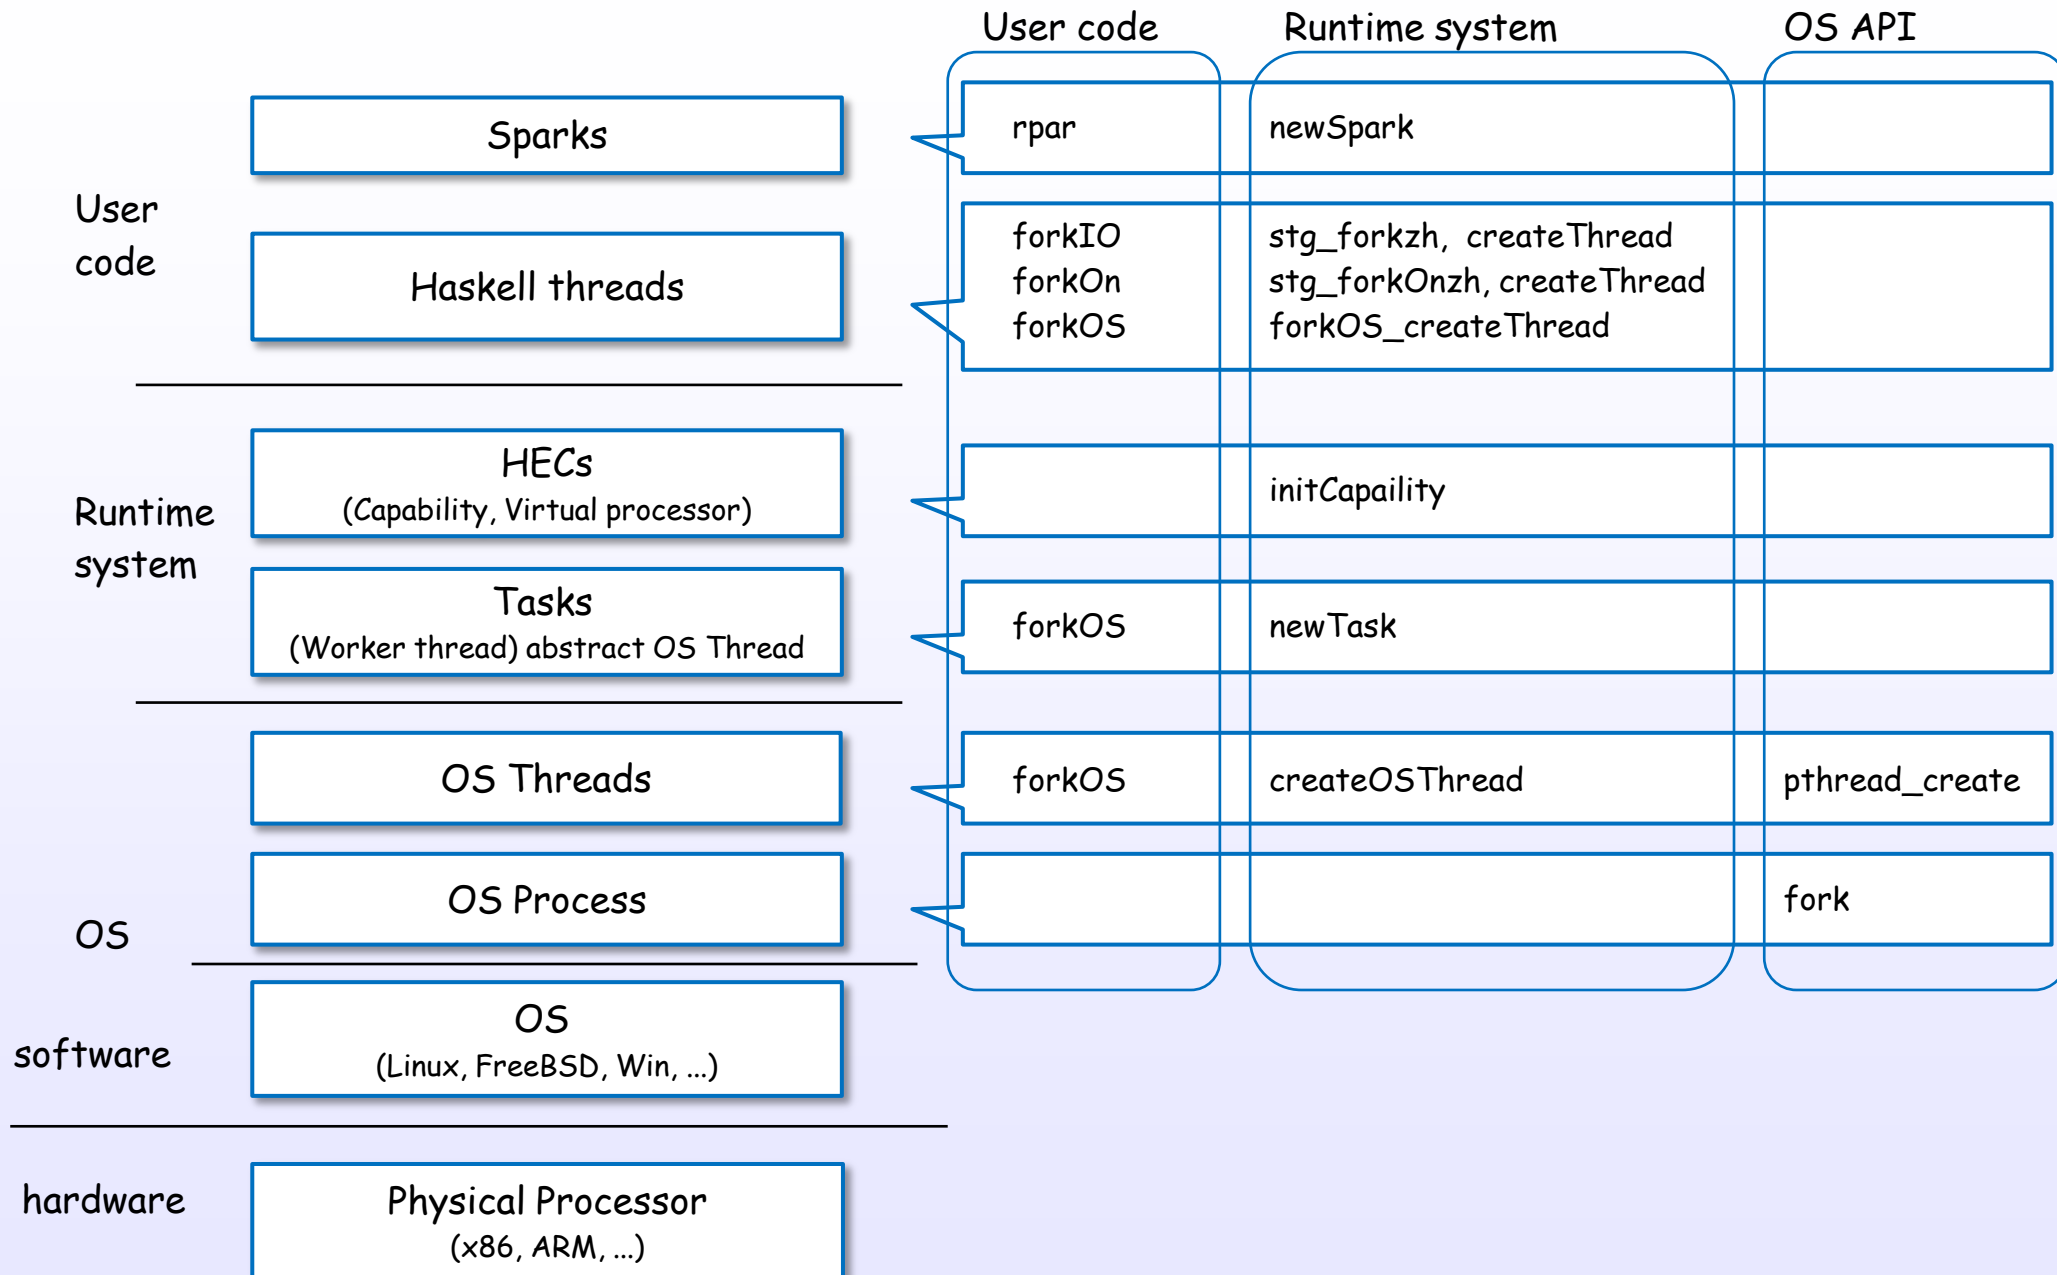
Exit sequence



Initializing



Create each layers



References

References

- [1] The Glorious Glasgow Haskell Compilation System User's Guide
https://downloads.haskell.org/~ghc/latest/docs/html/users_guide/index.html
- [2] Implementing lazy functional languages on stock hardware: the Spineless Tagless G-machine Version 2.5
<http://research.microsoft.com/en-us/um/people/simonpj/Papers/spineless-tagless-gmachine.ps.gz>
- [3] Making a Fast Curry Push/Enter vs Eval/Apply for Higher-order Languages
<http://research.microsoft.com/en-us/um/people/simonpj/papers/eval-apply/>
- [4] Faster Laziness Using Dynamic Pointer Tagging
<http://research.microsoft.com/en-us/um/people/simonpj/papers/ptr-tag/ptr-tagging.pdf>
- [5] Runtime Support for Multicore Haskell
<http://research.microsoft.com/en-us/um/people/simonpj/papers/parallel/multicore-ghc.pdf>
- [6] Extending the Haskell Foreign Function Interface with Concurrency
<http://community.haskell.org/~simonmar/papers/conc-ffi.pdf>
- [7] Mio: A High-Performance Multicore IO Manager for GHC
<http://haskell.cs.yale.edu/wp-content/uploads/2013/08/hask035-voellmy.pdf>
- [8] The GHC Runtime System
web.mit.edu/~ezyang/Public/jfp-ghc-rts.pdf
- [9] The GHC Runtime System
<http://www.scs.stanford.edu/14sp-cs240h/slides/ghc-rts.pdf>
- [10] Evaluation on the Haskell Heap
<http://blog.ezyang.com/2011/04/evaluation-on-the-haskell-heap/>

References

- [11] IO evaluates the Haskell Heap
<http://blog.ezyang.com/2011/04/io-evaluates-the-haskell-heap/>
- [12] Understanding the Stack
<http://www.well-typed.com/blog/94/>
- [13] Understanding the RealWorld
<http://www.well-typed.com/blog/95/>
- [14] The GHC scheduler
<http://blog.ezyang.com/2013/01/the-ghc-scheduler/>
- [15] GHC's Garbage Collector
http://www.mm-net.org.uk/workshop190404/GHC's_Garbage_Collector.ppt
- [16] Concurrent Haskell
<http://www.haskell.org/ghc/docs/papers/concurrent-haskell.ps.gz>
- [17] Beautiful Concurrency
<https://www.fpcomplete.com/school/advanced-haskell/beautiful-concurrency>
- [18] Anatomy of an MVar operation
<http://blog.ezyang.com/2013/05/anatomy-of-an-mvar-operation/>
- [19] Parallel and Concurrent Programming in Haskell
<http://community.haskell.org/~simonmar/pcph/>
- [20] Real World Haskell
<http://book.realworldhaskell.org/>

References

The GHC Commentary

- [C1] <https://ghc.haskell.org/trac/ghc/wiki/Commentary>
- [C2] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/SourceTree>
- [C3] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler>
- [C4] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/HscMain>
- [C5] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/CoreSynType>
- [C6] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/StgSynType>
- [C7] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/CmmType>
- [C8] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/GeneratedCode>
- [C9] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/SymbolNames>
- [C10] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts>
- [C11] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/HeapObjects>
- [C12] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/Stack>
- [C13] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/GC>
- [C14] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution>
- [C15] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution/Registers>
- [C16] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution/PointerTagging>
- [C17] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Scheduler>
- [C18] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/STM>
- [C19] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Libraries>

References

Source code

- [S1] `includes/stg/Regs.h`
- [S2] `includes/stg/MachRegs.h`
- [S3] `includes/rts/storage/ClosureTypes.h`
- [S4] `includes/rts/storage/Closures.h`
- [S5] `includes/rts/storage/TSO.h`
- [S6] `includes/rts/storage/InfoTables.h`
- [S7] `compiler/main/DriverPipeline.hs`
- [S8] `compiler/main/HscMain.hs`
- [S9] `compiler/cmm/CmmParse.y.source`
- [S10] `compiler/codeGen/StgCmmForeign.hs`
- [S11] `compiler/codeGen/Stg*.hs`
- [S12] `rts/PrimOps.cmm`
- [S13] `rts/RtsMain.c`
- [S14] `rts/RtsAPI.c`
- [S15] `rts/Capability.h`
- [S16] `rts/Capability.c`
- [S17] `rts/Schedule.c`
- [S18] `rts/StgCRun.c`
- [S19] `rts/StgStartup.cmm`
- [S20] `rts/StgMiscClosures.cmm`
- [S21] `rts/HeapStackCheck.cmm`
- [S22] `rts/Threads.c`
- [S23] `rts/Task.c`
- [S24] `rts/Timer.c`
- [S25] `rts/sm/GC.c`
- [S26] `rts/Sparks.c`
- [S27] `rts/WSDeque.c`
- [S28] `rts/STM.h`
- [S29] `rts/posix/Signals.c`
- [S30] `rts/win32/ThrIOManager.c`
- [S31] `libraries/base/GHC/MVar.hs`
- [S32] `libraries/base/GHC/Conc/IO.hs`
- [S33] `libraries/base/GHC/Conc/Sync.lhs`
- [S34] `libraries/base/GHC/Event/Manager.hs`
- [S35] `libraries/base/GHC/Event/Thread.hs`
- [S36] `libraries/base/GHC/IO/BufferedIO.hs`
- [S37] `libraries/base/GHC/IO/FD.hs`
- [S38] `libraries/base/GHC/IO/Handle/Text.hs`
- [S39] `libraries/base/System/IO.hs`
- [S40] `libraries/base/System/Posix/Internals.hs`
- [S41] `AutoApply.o (utils/genapply/GenApply.hs)`

Connect the algorithm and transistor